

IAMPerformance

Physics-Derived Cellular Fidelity Intelligence

Issue 003 · September 2026 · The Healthy Range of the Cellular Write Process — Eight Architecture Classes, Forty H_{min} Values, One Atlas of 115 Cell Types, and the First Written Specification of the Running Engine

Heath W. Mahaffey · IAMPerformance · Entiat, Washington

Repository: github.com/hmahaffey/IAM-Validation · engine at commit 66f37fe (2026-07-03)

Prepared with Claude Science, 2026-09-19. **DRAFT.** Every number is loaded from the runtime files or from a dated run named in the text. Rows marked OPEN are unresolved and are printed as such.

WHAT'S NEW IN ISSUE 003

The claim, restated. This is not a detection tool and not a treatment. Cells compute and write to a two-dimensional surface as a semiconductor or a quantum processor does; the physics of that write process has a healthy operating range per architecture class, calculated once and frozen (2026-04-06); departures from that range correlate with the hypo- and hyper-methylation patterns seen in disease *and* with normal variation in healthy individuals. Disease is corroboration, not product.

The IAM Atlas. IAMAtlasREBUILD: 483,092 CpGs, 115 cell types mapped to the eight classes, built [(‘Source rows’, ‘10,938,662 (iamatlas_mcmc_inputs.csv, 732 MB) from Moss 2018, Loyfer 2023, EpiDISH, Salas, Lister 2013 and others, harmonised to one cell-type label set’), (‘Chains per cell type’, ‘2’), (‘SD tightness target’, ‘posterior SD < 0.10 for ≥ 90% of CpGs per cell type’), (‘Informative-CpG selection’, ‘posterior SD ≤ 0.10 AND between-cell-type variance ≥ 0.02’), (‘Array bridging’, ‘only HM450-overlapping CpGs retained in production (483,092)’), (‘Failure recorded’, ‘stromal v0.1 chains failed to converge (R-hat 3.67) until source-specific labels were harmonised → v0.2’), (‘Deconvolver weighting’, ‘posterior-SD-weighted NNLS: $W = \text{diag}(1/\max(\text{posterior SD}))$, floor 1e-3; minimum 100 usable CpGs; dispersion penalty ceiling 0.20, high-dispersion flag 0.30; streaming 1K-marker chunks, re-solve every 10’), (‘Marker selection’, ‘one-vs-rest top-N (N = 100) on |target mean - mean(others)| per cell type (SOP s29); 131 chrX markers removed 2026-06-11 for derived sex-invariance’), (‘Walther vs NILC’, ‘substrate-level L1 disagreement of ~0.10-0.25 was EXPECTED and did not gate; inference-level disagreement on the disease-relevant class did (SOP s32-33). NILC cut 2026-07-02.’), (‘Versioning rule’, ‘the atlas evolves by deepening (new source rows, new SHAs) not by product versions; every consumer reads H_{min} from IAMAtlasREBUILD_provenance.json (SOP s99) - hardcoding H_{min} refuses deployment’)]. Every class card now lists its cell types. The atlas did not exist when Issue 002 was written.

Two instruments, stated. The GAUGE (entropy of one mean β over per-class identity loci) and the SEPARATION statistic (mean per-CpG entropy over per-cell discriminative markers) are different instruments answering different questions. Issue 002 had one A-score.

Healthy is a band, not a line. A = 1.0 is the architectural commitment line. Healthy reads on the age-matched p10–p90 band of age_{reference}_matrix.json. The per-class input offset explored in July is retired.

Presence before score. A class's A-score is a reading only if the deconvolver finds the class in the sample (DETECT_FLOOR). Scoring absent classes produced spurious BREACH calls in 4/4 healthy plasma samples on 2026-09-19; the rule is now printed.

Where the tools come from (§5A). The Mahaffey number (20.94 = dG_ATP/RT, the cell's ATP budget over k_B T; the name n_{bio} retired); the forty MCMC-confirmed floors and their forty ceilings; the atlas as a posterior; the four Planck tools borrowed for the methylome (brightness sky, masking, matched filter, cross-method deconvolution — the last one cut in July and recorded as such); and why this is not a cohort method.

Reconciliation table (§1). Every constant and rule that changed between April and repo HEAD, with the commit that changed it. Nothing silently updated.

Sections 7–8 replaced. The detection-trajectory and deployment-readiness sections are replaced by Substrate Characterization (§7) and Procedures (§8): what each substrate carries, what is verified, what is open, and the exact steps that reproduce each verified number.

Stage 1 executed from raw IDATs - the chain is closed. All eleven test IDAT pairs through the repo's own calibrator reproduce the cached betas bit-for-bit (11/11, max diff 0.000000, both array types). Raw IDAT → beta → class fractions → sealed anchor: every link now reproduced from the repository on a stranger's machine (PROC-CAL-01).

The end-to-end simulation caught the production gauge reading the wrong CpG set. A synthetic healthy patient - a pure mixture of healthy Atlas posteriors - reads BREACH through the conductor, because Stage B computes H(beta_{mean}) over the bimodal class marker union, not the identity loci its docstring and the SOP name; the age band was compiled the same way, so real blood never showed it. The identity-loci gauge reads the synthetic at 0.99 and real adenoma at 1.10 but has no band yet. Phase 1 becomes a prerequisite (PROC-N7-01, RECON A4).

The sealed breast anchor reproduced from raw GEO. GSE51032, 460 samples, 115 cell types: mean-of-H over the v0_2 discriminative markers returns the sealed CSV at r = 1.00000, max diff 0.00004, on a machine that had never seen the project. H(beta_{mean}) on the same markers does not (r = 0.45). The seal rests on the repo-HEAD markers, not the chrX-removed copy (PROC-ANCHOR-01).

The aggregation conflict between the current SOP and the running code, measured. SOP v1.4.0 says never H(β_{mean}); the chain at HEAD computes H(β_{mean}); the age band is compiled as H(β_{mean}). On whole blood with the class present the two aggregations are rank-identical with a

+0.029 offset; on mixed or absent-class panels $H(\beta_mean)$ inflates as §105 says. Ruled in §1.5: one aggregation per surface, fixed by the reference each was built with (RECON A3) - and the analyst's own first reading of the tissue values is retracted in §10.

First real-data ground truth for the deconvolver. Nine known genomic-DNA mixes (Moss 2018): the terminal class recovers a neuronal spike ($r = 0.945$); secretory and cycling do not recover hepatocyte or colon spikes. Recorded as PASS/FAIL in §8 and §10, not smoothed over.

Operating rules and falsification record (§9–10). The lessons that govern how the instrument may be read, dated and sourced; and every outcome that went against the framework, kept on the record.

WHAT THIS ISSUE DOES NOT YET COVER

Of the twelve chain stages at HEAD (0, 1, 2, 3, 4, 4.5, 4.6, 5, 6, 7, 8, 9) this issue documents three — Stages 2, 4 and 7 — and the atlas. The other nine (0, 1, 3 [not built], 4.5, 4.6, 5, 6, 8, 9) are named in §11 with their files and left for the next issue. The disease-signature matrix, the crown-jewel wall, the second chain, the directional AD detector, cellular age and the report builders are therefore not described here.

WHAT THIS PAPER IS NOT

Not clinical validation. Not a diagnostic. Nothing here should inform patient care. The instrument is at a research stage; any clinical use requires prospective validation, regulatory review and qualified clinical oversight. Public retrospective cohorts only; sample sizes in the new runs are small and are printed beside every number.

TABLE OF CONTENTS

Front	Cover · What's new · What this paper is not · Contents
§1	Reconciliation — Issue 002 → repo HEAD, every changed constant and rule
§2	The IAM Atlas — 115 cell types, eight classes, provenance, identity loci, markers
§3	Two Instruments and the Presence Rule
§4	Framework (from Issue 002) — global ranking, five substrates, MCMC↔bootstrap, saturation
Cards	Eight architecture-class cards (from Issue 002) — each with its atlas addendum: cell types, identity loci, age band
§1.7	NEW — The reporting rule (extends the §3.2 presence gate): a class gauge is put on the report only where the class is determined and present; whole blood reports immune plus one haematopoietic-progenitor component. stem_adult has carried no finding in nine VALs
§1.6	NEW — What the cosmology tools found that cohorts could not: the standing evidence ledger (nine rows to date, N7 first) that pre-empts the circularity objection
App. VI–IX	NEW — the CMB→methyloome translation map (79 rows, scored: what got built, what was cut, what was refused); the completion sprint scored, with the lesson that the bones must be trusted first; Future Goals — the CMB items worth the effort, in gated order; the Part II outline
Glossary	NEW — CMB and Chain Terms (Cosmic Methyloome Background, brilliance, HEALPix, component separation, matched filter, Mahalanobis Option A, the eight nulls, synthetic patients, PREREG/seal, Jensen gap, flatness ...) and Chain Links: one line per runtime file, tagged FLOOR / RULER / BAND / CODE / DATA
App. V	NEW — Validation index: all 103 VAL identifiers in the repository with title, date, cohort, stated decision, record completeness and path
§5A	NEW — Where the tools come from: Mahaffey number 20.94, forty MCMC floors, the atlas posterior, the CMB toolkit, not a cohort method
§5	Physics & Methodology (Issue 002 Section 2) — H_min derivation, substrates, saturation, inversions, C1/C2/C3
§6	Evidence, Baselines, Scenarios, Predictions (Issue 002 Sections 3–6)
§7	NEW — Substrate Characterization: the substrate × class grid
§8	NEW — Procedures: PROC entries, one per verified cell of the grid
§9	NEW — Operating Rules: the lessons as rules, dated and sourced
§10	NEW — Falsification Record
§11	NEW — Engine Map: every stage of the running chain and whether this issue covers it
§12	NEW — For the clinician: the instrument in plain terms
Back	Master predictions · Data sources · Glossary · A final note

SECTION 1

RECONCILIATION — ISSUE 002 → REPO HEAD

Issue 002 was written in April 2026, before the atlas and before the engine was rebuilt (commits 2026-06-25 to 2026-07-03). This table states, for every constant and rule that differs, what 002 said, what the running engine says, why the engine is current, and where to check. Rows marked **OPEN** are real inconsistencies at HEAD that this document does not resolve. The chain-of-custody SOP v1.3.x predates the rebuild; commit d7b0e1f explicitly supersedes its §41. Until this section, the commits were the only specification of the running engine.

#	Item	Issue 002	Repo HEAD (66f37fe)	Why HEAD is current	Evidence
H1	40-cell H_min table (8 classes x 5 substrates)	as printed 002 pp.7,105	byte-identical in cpg_gauge_engine.py H_MIN_TABLE	unchanged	G-002/G-003b freeze 2026-04-06; diffed 2026-09-19: 40/40 equal
A1	Where healthy sits	A ~ 1.0 with an age curve (AGE_REF, 7 points/class)	A = 1.0 is the COMMITMENT LINE, not healthy. Healthy = age_reference_matrix.json p10-p90 band (10 decades/class, n per decade, source per decade)	engine docstring; Issue 002 \$'NOT forced to 1.0'	cpg_gauge_engine.py lines 26-29; commit 8cdf352 2026-07-01
A2	Age reference values	AGE_REF immune 40y=0.946, 60y=0.966, 80y=0.992	age_reference_matrix immune 44y=0.9477 [p10 .9021, p90 .9934], 64y=0.9652, 84y=0.9873	HEAD carries percentiles + n + citation per decade; 002 carried means only	age_reference_matrix.json; Hannum 2013, Horvath 2013, Alisch 2012
O1	Per-class input-scale offset	not present	explored (immune/methyl = 0.055, commit 143704d) then RETIRED	it was reinventing the age band and forcing healthy onto 1.0	cpg_gauge_engine.py lines 46-53
T1	Tier ladder (gauge)	NORMAL<1.01 MARGINAL<1.05 DETECTABLE<1.07 URGENT<1.10 BREACH>=1.10	identical	unchanged	cpg_gauge_engine.py lines 38-44; 002 tier_label()
T2	Tier ladder (Stage 7 customer, tier_breakpoints.json)	not in 002	SUPPRESSED<0.95 NORMAL<1.01 ELEVATED<1.07 [Warburg line 1.07] SIG_ELEVATED<1.10 BREACH>=1.10	NORMAL->ELEVATED edge moved 1.04->1.01 to match gauge MARGINAL onset (last commit)	commit 66f37fe 2026-07-03; tier_breakpoint.s.json _meta.marginal_onset_note
T3	Two tier vocabularies coexist	one	gauge says MARGINAL/DETECTABLE/URGENT; Stage 7 says ELEVATED/SIGNIFICANTLY_ELEVATED	OPEN - same breakpoints except the gauge's 1.05 split; one vocabulary must be chosen for the report	both files at HEAD
S1	Substrate floors for cfDNA	cfDNA listed as a substrate with % contributions (CFDNA_PCT)	'cfDNA is the frag/wps/nucl substrates - score it on those floors, NEVER the methyl floor'	engine docstring	cpg_gauge_engine.py lines 56-58
I1	Two instruments	one A-score	GAUGE = H(beta_mean)/H_min over per-class IDENTITY loci; SEPARATION = mean_i(H(beta_i))/H_min over per-cell DISCRIMINATIVE markers	feeding the gauge discriminative markers produced the all-BREACH bug of 2026-06-11	cpg_gauge_engine.py lines 16-25; iamatlas_a_scoring.py
D1	Deconvolver role	not in 002	Walther NNLS on IAMAtlasREBUILD -> class + cell-type fractions; PRESENCE gate: a cell/class below DETECT_FLOOR is absent and its A is not a reading	cpg_conductor.py 2026-07 ' Replaces the confusing walther_clinical.py'	cpg_conductor.py DETECT_FLOOR = 0.01

#	Item	Issue 002	Repo HEAD (66f37fe)	Why HEAD is current	Evidence
D2	Presence floor value	-	1% (cpg_conductor.py) vs 3% (README_FOR_FUTURE_AI Mahalanobis adjudicator gate)	OPEN - two floors give different answers on plasma terminal class	both files at HEAD; run 2026-09-19
D3	NILC second deconvolver	-	CUT from Stage 2; Walther alone MEASURED (PROC-NILC-01): $\kappa(\text{blood sub-problem}) = 30.6$, $r(\text{progenitor, stem_adult}) = +0.99$; NILC L1 vs Walther 0.55-0.65 on all blood.	the divergence was the diagnostic NILC was designed to give: stem_adult/progenitor are not separately determined by the Atlas in blood. Report them jointly for whole blood until they are; a second deconvolver's disagreement is to be read, not failed.	commit c1be0c3 2026-07-02
D4	Deconvolver validation scope	-	N7: MAE 0.0076-0.0093 on SYNTHETIC Dirichlet mixtures; conformance MAE 0.0004 vs TEST_DATA_MANIFEST on 3 real EPIC tissue samples	never validated against real tissue of KNOWN composition; GSE122126 in-vitro mixes are the available ground truth	N7_OUTCOME.md; run 2026-09-19
K1	Atlas	not in 002	IAMAtlasREBUILD: 483,092 CpGs, 115 cell types -> 8 classes, built 2026-05-28	predecessor IAMAtlas.csv.xz retired (collapsed; flatness bug)	IAMAtlasREBUILD_provenance.json
K2	Atlas CpG count stated	-	provenance + file: 483,092; SOP v1.3.3 §28: 481,966	OPEN - SOP figure is stale or counts a filtered set	diffed 2026-09-19
M1	iamatlas_celltype_markers_v0_2.json	-	trial zip 200,311 B vs repo HEAD 237,545 B - DIFFERENT files, same name	OPEN - which is current must be stated; version the filename	sha256 diff 2026-09-19
SOP	Chain-of-custody SOP	-	SOP v1.3/1.3.3 is PRE-atlas; commit d7b0e1f says 'supersede SOP §41'; no post-atlas SOP exists	the commits 2026-06-25..07-03 ARE the spec; this document is the first written one	git log; README_FOR_FUTURE_AI.md
SEC	Sections 7-8 (detection trajectory, deployment readiness)	present	REPLACED by Substrate Characterization + Procedures	user decision 2026-09-19: not a detection tool; the claim is the healthy range of the write process	this issue
F1	Stage 3 foreground subtraction	not in 002	NOT WIRED by decision (SOP s104 production foreground firewall): the production chain subtracts no age-, sex- or smoking-driven change. Sex: 131 chrX markers removed 2026-06-11. Age: the reference band. Smoking: Stage 7 bin floor shifts (v1.2 interim). Subtraction modules exist as test-only tooling.	a stated design decision, not a gap; the flowchart's 'NOT BUILT' should read 'NOT WIRED - by s104'	SOP v1.4.0 s104; flowchart_vKISS
F2	Stage 6 cellular age module	not in 002	iam_cellular_age_scoring.py present in CPG_TRIAL_CODE.zip, ABSENT from repo engine directory	OPEN — the canonical repo does not carry a file the trial bundle does	sha256 listing 2026-09-19
F3	walther_clinical.py origin gate	not in 002	disease_origin_cells.json loaded inside a bare except; on failure the cell-of-origin specificity rule is silently disabled	OPEN — patient-safety gate fails open; three-line fix	walther_clinical.py ~lines 972-996 (CPG_first_read.md finding 3)
F4	Stale strings describing the retired mean-of-entropies gauge	—	cpg_conductor.py Stage A comment ' (mean-of-per-CpG H/H_min)'; cpg_gauge.py line 158 axis label prints on the patient report figure	OPEN — neither changes a number; both are how the 2026-06-11 bug gets reintroduced	CPG_first_read.md finding 2

#	Item	Issue 002	Repo HEAD (66f37fe)	Why HEAD is current	Evidence
F5	Gauge carries no sign	—	$A(\text{beta})=A(1-\text{beta})$; INVERSION below band is degenerate between hyper- and hypo-methylation. beta_mean is computed one line before A and not reported	OPEN — report beta_mean or signed departure beside every A; no revalidation needed	iamatlas_a_scoring.py; verified $A(0.61)=A(0.39)=1.1501$ on immune
F6	Two A-score definitions in the corpus	002 §2.5 C1/C2/C3	$A = H/H_{\min}$ (all numerical results) vs $A = 1 + C3/H_{\text{actual}}$ (three-component decomposition) — equal only at $A=1$	OPEN — one must be retired or the relation stated	002 §2.5; IAM_for_physicists §4
F7	Immune $H_{\min}$ revision history	0.838889	calibration value 0.795000 -> 0.8389 +/- 0.0012 after MCMC reanalysis of six immune cell types; immune A moved down 0.055	state the freeze date and provenance so 'frozen before disease testing' is checkable	EDEAR s4.6 (0.795 -> 0.838889, 6.44 sigma); Issue 002 immune card ('revised downward by approximately 0.055')
A3	A-score aggregation: gauge vs separation	$A = H(\text{beta})/H_{\min}$, beta = the cell's mean methylation; $H_{\min_global} = H(0.790)$	SOP v1.4.0 s105: mean-of-H, never $H(\text{beta_mean})$. Code at HEAD: $H(\text{beta_mean})$ over identity loci; age band compiled the same way.	RULED (see RULING A3): gauge on identity loci = $H(\text{beta_mean})/H_{\min}$, matching $H_{\min}$'s own definition and the band; separation on discriminative markers = mean-of-H, matching s105 and the sealed anchors; s105 amended to scope its 'never' to marker panels; runtime bimodality guard added to the gauge.	PROC-FORMULA-01; PROC-ANCHOR-01; RULING A3
M1b	iamatlas_celltype_markers_v0_2.json - which is current	-	TRIAL file carries the 2026-06-11 chrX removal; REPO HEAD does not; the 2026-05-29 seal rests on the pre-removal file.	RULED (see RULING M1b): adopt the chrX-removed file, commit it, re-seal both anchors (done 2026-09-19: 32/115 cells shift, max 0.059, $r = 0.9996$; v1 seal retained as SUPERSEDED).	PROC-ANCHOR-01; kit/anchors_v2/RESEAL_REPORT.json
A3b	SOP v1.4.0 s105 open decisions vs the code	-	s105 leaves three decisions open (Recipe s6.3 source; marker convention discriminative-vs-most-methylated; s46.5 re-derivation). The 07-01 code answered none of them and introduced a fourth convention (identity loci + $H(\text{beta_mean})$) with the age band compiled to match.	CLOSED by RULING A3 for decisions 2 (marker convention: identity loci for the gauge, discriminative for separation) and 3 (s46.5 stands, gauge keeps $H(\text{beta_mean})$). Decision 1 (Recipe s6.3) remains the author's to record in the vault.	RULING A3
M2	'Mahaffey number' - definition and value	glossary: ' $n_{\text{bio}} = dG_{\text{ATP}}/(RT) = 20.94$ '	author 2026-09-19: $M = 20.94$, one fixed number, no $\ln 2$; n_{bio} is an early QAPE-era form, no longer used; $H_{\min}$ is a separate per-class quantity	RULED BY AUTHOR. Open: Hubble2GAPE I.2192 $\ln 2$ form (~30) to be reconciled.	author statement; 002 glossary p251; H2G I.2192, I.3504
B1	class bands are cohort statistics of unequal provenance	age_reference_matrix.json: immune from hundreds of samples across decades; stem_adult $n=28$ single source; others in between	the FLOOR ($H_{\min}$, 40 values) is physics and is one number per class-substrate; the BAND (age_reference_matrix) is a cohort statistic and is only as good as its n and pipeline. Every 'healthy reads wrong' case of 2026-09-19 traces to a band, never to a floor.	Phase 1: rebuild all eight class bands from ONE healthy cohort (GSE87571, $n=730$, ages 14-94) through ONE pipeline (Stage 1); record $n/\text{source}/\text{pipeline}$ per band; classes without a defensible band from blood -> NOT_ASSESSABLE	PROC-CHAIN-01

#	Item	Issue 002	Repo HEAD (66f37fe)	Why HEAD is current	Evidence
V1	what 'A' is in VAL-001/003/008 (the 27/28 claim)	report: 'A = H(beta)/H_min in(class)'	numerator = COHORT-MEAN GENOME-WIDE beta (TCGA tumor / solid-tissue-normal means, Xu 2019; healthy from one Roadmap cell per class); divisor = H_min of the cancer's tissue-of-origin class (LGG->terminal, BRCA->secretory, COAD->cycling ...). Class-correct divisor, global-mean numerator, cohort-level - a different instrument from the identity-loci gauge, sharing only the divisor. The script itself notes a '~10% global entropy offset expected'.	state this wherever 27/28 is quoted; it explains why direction holds while absolute A is offset	GAPE_Evidence_Report_UPDATED VAL-003 script I.5208-5544
A4	which CpG set the production gauge reads	docstring / SOP s41 / Issue 002: identity loci	code: class marker union (iamatlas_celltype_markers_v0_2.json), with age_reference_matrix compiled on the same	identity loci, once Phase 1 supplies their band; label until then	PROC-N7-01
D3	why NILC and Walther never agreed	flowchart: 'collapsed on correlated blood mixtures and deleted correct calls'	NILC (RETIRED/NILC_Deconvolver_cut_from_chain_2026-07-02) solved UNCONSTRAINED generalized least squares on the CELL-TYPE MARKER POOL (iamatlas_celltype_markers_v0_1.json, bimodal one-vs-rest), then projected onto the simplex; Walther solves NNLS with a simplex constraint on its own top-600-per-class CLASS-discriminating CpGs. In blood the immune / progenitor / stem_adult reference columns are nearly parallel, so the unconstrained inverse is ill-conditioned: fractions swing negative, the projection zeroes them - 'deleted correct calls'. Not a second opinion; the unregularised version of the same inverse problem on a worse-conditioned CpG set. MEASURED (PROC-NILC-01): kappa(blood sub-problem) = 30.6, r(progenitor, stem_adult) = +0.99; NILC L1 vs Walther 0.55-0.65 on all blood.	the divergence was the diagnostic NILC was designed to give: stem_adult/progenitor are not separately determined by the Atlas in blood. Report them jointly for whole blood until they are; a second deconvolver's disagreement is to be read, not failed.	NILC docstring; c1be0c3; PROC-N7-01

Reading the table

The forty H_min values (row H1) are unchanged: the floors were frozen on 2026-04-06 and the engine at HEAD carries them byte-for-byte. What changed is everything *around* the floors — where healthy sits (A1), what a score means when the class is not in the sample (D1), which floor a substrate is read against (S1), and the existence of a second instrument (I1). Five rows are OPEN: two tier vocabularies (T3), two presence floors (D2), an atlas count that disagrees with the SOP (K2), two files with one name (M1), and the absence of any post-atlas SOP (SOP). Each is a question a reviewer will ask; each is answered here by being stated rather than hidden.

1.5 Rulings on the two open decisions

The author delegated the two decisions the reconciliation left open. Both are recorded here with their reasoning so they can be overruled with the same specificity.

RULING A3 - the A-score aggregation, one rule per surface

Rule. A statistic is read only against a reference computed by the same aggregation. There are two surfaces and each keeps the aggregation its reference was built with.

Gauge. CLASS GAUGE (8 classes, identity loci, Stage 4): $A = H(\text{beta_mean})/H_{\min}(\text{class})$. Because: (i) H_min is itself H of a mean beta - G-002 calibrated each class floor on per-cell-type mean methylation of sorted reference cells ($H_{\min_global} = H(0.790)$, Issue 002 s2.1), so $H(\text{beta_mean})/H(\text{beta_ref})$ compares like with like; (ii) age_reference_matrix.json, the band every gauge reading is placed against, is compiled as

H(beta_mean)/H_min to five decimals in 80/80 cells; (iii) identity loci are selected within 0.05 of H_min_beta and are therefore not bimodal - the Jensen inflation s105 describes is bounded, measured at +0.029 +/- 0.002 on the declared substrate (PROC-FORMULA-01); (iv) it is the formula of Issue 002 and of the chain wired at HEAD (d7b0e1f). Mean-of-H over identity loci would be a different quantity whose natural floor (mean per-locus entropy of the reference cell) was never computed.

Separation. SEPARATION SURFACE (115 cell types, discriminative markers, Stage 4 per-cell and Stage 5 hull): $A = \text{mean}_i H(\beta_i)/H_{\min}(\text{class})$. Because: (i) discriminative markers are bimodal by construction, so $H(\beta_{\text{mean}})$ manufactures ~0.5 and false BREACH - s105 is exactly right here (worked guard GSM1235534: 1.111 vs 0.864); (ii) it is what the sealed anchors are: PROC-ANCHOR-01 reproduces GSE51032 and GSE51057 at $r = 1.00000$ with this formula and not with the other.

Guards. The gauge refuses to score when (a) the class is absent from the substrate (presence gate, DETECT_FLOOR) or (b) averaging beta first is not valid on the panel it actually receives, tested directly: the Jensen gap $H(\beta_{\text{mean}}) - \text{mean}_i H(\beta_i)$ on that panel must not exceed 0.05. Calibrated 2026-09-19: identity loci in whole blood 0.033 (pass); identity loci in bulk tissue 0.20 (refused - a mixture, glioma-LL-002); s105's own HSC marker example 0.29 (refused); hepatocyte markers 0.78 (refused). This turns s105's 'never' into a measured condition enforced at runtime (kit/cpg_kit.py::gauge_A).

Sop. SOP s105 is amended, not reversed: its one sentence becomes 'On any marker panel, score per locus then average; never average beta first. The class gauge on identity loci is the one surface where $H(\beta_{\text{mean}})$ is the correct statistic, because $H_{\min}$ and the age band are both defined that way.' s105 open decision 1 (Recipe s6.3) is for the author to record in the vault; decision 3 (s46.5 bidirectional recovery assumed $H(\beta_{\text{mean}})$) stands as written because the gauge keeps $H(\beta_{\text{mean}})$.

RULING STANDS FOR THE FORMULA; CORRECTED ON THE CHAIN 2026-09-19 (PROC-N7-01): the wired conductor computes $H(\beta_{\text{mean}})$ over the class MARKER UNION, not identity loci, and age_reference_matrix was compiled the same way. The ruling's target - identity-loci $H(\beta_{\text{mean}})$ against an identity-loci band - does not exist yet; Phase 1 builds it. Until then the conductor's gauge is labelled gauge_surface='marker_union'.

RULING M1b - the markers file, and re-sealing the anchor

Rule. Adopt iamatlas_celltype_markers_v0_2.json WITH the 2026-06-11 chrX removal (the TRIAL/user copy) as canonical; commit it to the repository; re-seal both foundation-cohort anchors against it.

Why. The removal is a derived, documented sex-invariance correction (131 chrX markers; 'so sex cannot shift an A-score'). The original seal (2026-05-29) predates the fix (2026-06-11); a seal that rests on a known-uncorrected file is weaker than a re-seal that records the change. Keeping the pre-fix file to preserve the seal would be preserving the number over the correction.

Reseal. Re-sealed values computed 2026-09-19 from the raw GEO matrices with the chrX-removed markers and mean-of-H (RULING A3): GSE51032 (n = 460) and GSE51057 (n = 188), 115 cells each. Relative to the 2026-05-29 seal: 32 of 115 cells change, largest shift 0.059 (Mela), $r = 0.99957 / 0.99951$, 83 cells unchanged. Files: kit/anchors_v2/GSE51032_115celltype_ascores_v2_chrXremoved.csv (sha256 212775a8...bcd9b), GSE51057_..._v2_chrXremoved.csv (sha256 b4b854a3...93597). The 2026-05-29 seal is retained as v1 and marked SUPERSEDED, not deleted.

RULED 2026-09-19 under delegated authority; the repository commit is the author's to make.

WHAT THE COSMOLOGY TOOLS FOUND THAT COHORTS COULD NOT

The objection this section pre-empts is the one every reviewer will make: *"your healthy reference was calibrated on cohorts and validated on cohorts - that is circular."* The answer is not a denial. It is correct, and a cohort-only pipeline cannot detect the circularity - which is why this project does not use cohorts to validate the instrument. It validates the way a CMB experiment does: synthetic data with known truth, injection-recovery, split-half cross-checks, convergence and distinctness tests, look-elsewhere correction, sealed pre-registration. Below is what those methods found that no cohort could have. It is a ledger, kept current by rule (Reproduction Kit RUNBOOK §10), and it includes the reversals.

Every time a CMB-derived method surfaces something a cohort comparison could not have, it gets a row in this ledger the same day, with its PROC or VAL. This is not a list of successes; the NILC reversal and the withdrawn tissue-of-origin claim are here too. It is the record that answers the circularity objection: we do not ask cohorts to validate the instrument, because they cannot - we ask constructed truth, split halves, injection-recovery and pre-registration to do it, and here is what they found.

date	CMB method	why a cohort is blind to it	what was found	record
2026-09-19	End-to-end simulation with known truth (FFP / mock-sky discipline; null N7)	A cohort supplies a comparison, never a truth. Two statistics that are wrong in the same way agree on every cohort.	The production class gauge read H(beta_mean) over the bimodal MARKER UNION, not the identity loci; the age band had been compiled on the same statistic, so every real cohort for eleven weeks read 'healthy IN_BAND'. A synthetic healthy patient - a pure mixture of healthy Atlas posteriors - read BREACH (1.13) on the first run. Identity-loci gauge read 0.99 on the same input and 1.10 on real adenoma where the production gauge read NORMAL.	PROC-N7-01; RECON A4
2026-09-19 (rerun of 2026-06)	Cross-method disagreement read as a diagnostic (Commander / NILC / SMICA discipline)	A single-method cohort pipeline cannot know which of its fractions the reference actually determines and which the solver's constraint chose.	NILC rerun as designed on Walther's own reference: L1 disagreement 0.55-0.65 on every blood sample, mass moving between progenitor and stem_adult (reference columns $r = +0.99$, blood sub-problem condition number 30.6). The Atlas cannot separately determine those classes in blood; Walther's 2-4% stem_adult is constraint-chosen. That class produced the false BREACH alarms of PROC-CHAIN-01. The second method had diagnosed it in June and was cut for it.	PROC-NILC-01; RECON D3
2026-09-19	Component separation validated on known mixtures (Planck FFP component maps; here Moss 2018 Table 6 in-vitro mixes)	Bulk-tissue cohorts never expose the mixing fractions; a deconvolver can be systematically wrong and still order cohorts correctly.	Neuron spikes recovered ($r = +0.945$, under-read 10% -> 6.5%); hepatocyte and colon spikes NOT recovered ($r = +0.19$, $+0.10$) - the atlas routes shed epithelium to gastric references. Tissue-of-origin claims withdrawn until the atlas recovers Table 6; composition departure itself confirmed real.	PROC-PLASMA-MIX -01; LESSON-DECO NV-01
2026-09-19	Transfer-function decomposition (what each pipeline stage does to a known input)	On real data every stage's effect is confounded with biology; only a constructed input isolates the stage.	Term-by-term decomposition of one synthetic patient (composition -> age loading -> sex -> batch -> noise -> clip) showed every generator term moved the identity-loci gauge by < 0.004 ; the $+0.14$ lived entirely inside Stage B. That is what located the marker-union code path in one step.	PROC-N7-01 decomposition
2026-05 -> 2026-09	Half-mission x half-mission cross-check (two independent cohorts as two detector halves)	A single cohort can carry a plate or preprocessing offset that looks like signal; only an independent half exposes it.	GSE51032 x GSE51057 breast anchors: the effect replicates across both halves ($d = +2.088$ vs $+2.097$ in the sealed record); the whole 648-sample foundation cohort reproduced from raw GEO at $r = 1.00000$. Separately, GSE53740's healthy controls sat $+2.3$ SD above the 80-cell baseline - a cohort offset that a single-cohort analysis would have read as disease (CCL-004).	PROC-ANCHOR-01; CCL-004
2026-06	Convergence diagnostics are not a truth test (the R-hat lesson from MCMC map-making)	Cohort validation has no analogue of a chain diagnostic at all; it cannot see that a reference is flat.	The first Atlas build showed R-hat ~ 1.01 (looks perfect) while every cell type in a class came out identical - a flat, useless atlas. The fixed build showed R-hat 1.4-2.6 (looks alarming) with distinct cell types. Judge a rebuilt class by the distinctness test, never by convergence alone.	IAM_Atlas/IAMAtlas_ FLATNESS_LESSO N.md

date	CMB method	why a cohort is blind to it	what was found	record
2026-06	Component-separation cross-validation (Commander / NILC / SMICA discipline) - the reversal	A second method fails differently from the first; with one method, the chain is 'talking to itself' (sprint sign-off question 3).	NILC was built as the second deconvolver and CUT: it collapsed on correlated blood mixtures and deleted correct calls. Recorded as the one Planck principle the chain knowingly does not follow, with the reason. A parametric second method is Future Goal #6.	commit c1be0c3; Appendix VI row 20; Appendix VII B2
2026-05	Look-elsewhere correction (null N8)	A cohort scan over many features will always find one that separates; without the correction the finding is published.	VAL-006's chr6 (MHC) signal died under look-elsewhere correction and was recorded as a null rather than a discovery.	VAL-006; null N8
2026-04 -> 2026-05	Pre-registration with sealed prediction and direction (the CMB blinding culture)	A cohort analysis performed after the labels are seen can accommodate any sign.	VAL-061 predicted the wrong sign and was recorded as such, which produced CCL-019/020 (direction depends on class x compartment, not disease). VAL-102 was VOIDED four minutes after sealing for a post-hoc accommodation, seal preserved. VAL-128 failed opposite to its prereg and stands as a FAIL.	CCL-019; CCL-020; VAL-102; VAL-128
2026-09-19	Foreground reasoning applied in reverse (delensing refused)	Cohort clocks subtract age as a nuisance; that removes the patient's own biology along with it.	Built age/sex/smoking subtraction layers were refused under SOP s104: a galactic foreground is a separate source, the methylome's 'foreground' is the patient. The band (annotate) replaced subtraction (de-age). Recorded as the second place the CMB analogy is deliberately broken - and why.	Appendix VI row 47; SOP s104

THE REPORTING RULE - WHICH CLASSES A SUBSTRATE MAY REPORT

Asked how often stem_adult had been the deciding class, the record was checked: every OUTCOME file, the VAL index, every disease card.

where	stem_adult result	what carried the finding
VAL-008 AD (AIBL), per cell type	d = -0.329, second-ranked	immune carried the finding (d = +0.68)
VAL-016 cross-disease universal alarm	second-ranked	immune
VAL-011 AD age subtraction	d = -0.004 -> -0.190 after age subtraction	'interesting biology'; never a call
VAL-015 / VAL-020 immune aging	r = -0.103, 'weak aging signal'	-
VAL-018 menarche	~0	null
VAL-022 smoking cessation	the only class with true reversibility (former/never ratio 0.66)	flagged 'worth investigating later' - the one distinctive stem_adult result
VAL-135 (retired, off-scope)	1.083 'ELEVATED - blood-stem-cell content'	-
disease cards (all)	none names stem_adult as origin class	Issue 002 targets: HSC-origin AML, CHIP, HSC aging - haematology, not the breast/CRC/AD work
age band	n = 4-32 per decade, sorted CD34+ HSC (Adelman 2019), two decades extrapolated	purified stem cells, no whole-blood context
fraction in blood	2-4%, constraint-chosen (r = +0.99 with progenitor, PROC-NILC-01)	false BREACH in 6/7 healthy donors (PROC-CHAIN-01)

statement. REPORTING RULE (extends the s3.2 presence gate, DETECT_FLOOR = 1%, which decides whether a class is in the sample; this decides whether its gauge is put on the report). A class gauge is REPORTED on a substrate only if the class is (1) DETERMINED - its fraction is separable from its neighbours in the Atlas on that substrate (condition number of the substrate sub-problem < ~10), or it is reported jointly with what it cannot be split from - and (2) PRESENT - its typical fraction on that substrate is >= 5% (above the conductor's 1% detection gate and the adjudicator's 3%), so that beta_mean over its loci is the class and not the background. Classes failing either condition are composition-only (present/absent, no A).

whole blood. REPORTED: immune. REPORTED JOINTLY: progenitor + stem_adult as one haematopoietic-progenitor component. COMPOSITION-ONLY: cycling, secretory, terminal, stromal, stem_pluri. Decided from kappa and fraction BEFORE any A is computed; written into the Phase 1 PREREG as a pass condition.

rationale. Risk over reward (author, 2026-09-19): a class that has never carried a finding, whose fraction the data does not determine, and whose band rests on n = 4-32 sorted cells, is a false-alarm source on the patient report and nothing else. Classes earn their way back onto a substrate's report as separability and presence are demonstrated there (plasma cfDNA is expected to qualify cycling and secretory).

keep. VAL-022's stem_adult reversibility (ratio 0.66) is retained as a note - the one distinctive result - testable once a joint component and a real band exist.

THE IAM ATLAS

IAMAtlasREBUILD — 483,092 CpGs × 115 cell types, each mapped to one of the eight architecture classes. Built [('Source rows', '10,938,662 (iamatlas_mcmc_inputs.csv, 732 MB) from Moss 2018, Loyfer 2023, EpiDISH, Salas, Lister 2013 and others, harmonised to one cell-type label set'), ('Chains per cell type', '2'), ('SD tightness target', 'posterior SD < 0.10 for >= 90% of CpGs per cell type'), ('Informative-CpG selection', 'posterior SD <= 0.10 AND between-cell-type variance >= 0.02'), ('Array bridging', 'only HM450-overlapping CpGs retained in production (483,092)'), ('Failure recorded', 'stromal v0.1 chains failed to converge (R-hat 3.67) until source-specific labels were harmonised -> v0.2'), ('Deconvolver weighting', 'posterior-SD-weighted NNLS: W = diag(1/max(posterior SD)), floor 1e-3; minimum 100 usable CpGs; dispersion penalty ceiling 0.20, high-dispersion flag 0.30; streaming 1K-marker chunks, re-solve every 10'), ('Marker selection', 'one-vs-rest top-N (N = 100) on |target mean - mean(others)| per cell type (SOP s29); 131 chrX markers removed 2026-06-11 for derived sex-invariance'), ('Walther vs NILC', 'substrate-level L1 disagreement of ~0.10-0.25 was EXPECTED and did not gate; inference-level disagreement on the disease-relevant class did (SOP s32-33). NILC cut 2026-07-02.'), ('Versioning rule', 'the atlas evolves by deepening (new source rows, new SHAs) not by product versions; every consumer reads H_min from IAMAtlasREBUILD_provenance.json (SOP s99) - hardcoding H_min refuses deployment')] from per-class MCMC (batch 5000, immune 1500), duplicate reconciliation, and merge; the predecessor IAMAtlas.csv.xz was retired for a collapse/flatness bug. It ships compressed in the repository (98 MB .xz → 577 MB .csv) and is the deconvolver reference, the source of the gauge identity loci, and the source of the separation markers. H_min values in the atlas are the frozen 2026-04-06 set.

2.1 Cell types by class

Class	n cell types	Identity loci (gauge)	H_min methyl	Ceiling 1/H_min	Cell types (atlas names)
terminal	9	57,247	0.7728	1.2939	CM, Cortical_neurons, Epi_upper, Glia, Kera_diff, Left_atrium, NeuMa, Oligo, neuron
secretory	18	29,617	0.8433	1.1859	Acinar, Antrum_diff, Breast, Chol, Corpus_diff, Endocrine, Fundus_diff, Gland, Hep, Hepatocytes, LE, Luminal, Mela, Pancreatic_acinar_cells, Pancreatic_beta_cells, Pancreatic_duct_cells, Prostate, Thyroid
immune	51	42,134	0.838889	1.1921	B, B-cells_EPIC, Baso, Bcell, Bmem, Brv, CD14_monocytes, CD19_B-cells, CD4T, CD4T-cells_EPIC, CD4Tmem, CD4Tnv, CD4_T-cells, CD56_NK-cells, CD8T, CD8T-cells_EPIC, CD8Tmem, CD8Tnv, CD8_T-cells, Eos, Eosino, Eosinophils_reinius, Gran, IC, Kup, Leu, Lym, MP, Macro, Microglia, Mono, Monocytes_EPIC, Mye, NK, NK-cells_EPIC, Neu, Neutro, Neutrophils_EPIC, Neutrophils_reinius, PBMC, Plasma, Tcell, Treg, dendritic, eosinophil, granulocytes, macrophage, monocyte, neutrophil, tcell, whole_blood
progenitor	11	54,704	0.8522	1.1734	CMP, Erythrocyte_progenitors, GMP, L-MPP, MEP, MPP, Neulm, OPC, erythroblast, megakaryocyte, nRBC
cycling	19	29,181	0.8561	1.1681	Antrum_undiff, BE, Basal, Bladder, Colon_epithelial_cells, Corpus_undiff, Ductal, Epi, Epi_basal, Epi_stratified, Epi_suprabasal, Fundus_undiff, Head_and_neck_larynx, Kera_undiff, Kidney, Lung_cells, Upper_GI, Uterus_cervix, small_intestine
stromal	5	2,294	0.863	1.1588	adipocyte, endothelial, fibroblast, smooth_muscle, stromal_other
stem_adult	1	46,907	0.8737	1.1445	HSC
stem_pluri	1	21,710	0.9822	1.0182	stem_pluri

Two structural facts a reader should take from this table. **Stromal** has 2,294 identity loci against 29,000–57,000 for every other class and only five cell types; on colorectal tissue it returned exactly 0.0000 in 4/4 samples with 1,441 of 7,090 class markers matched — a coverage gap, not a biological absence (§7). **Immune** holds 51 of the 115 cell types; a pooled immune gauge therefore requires a single cell type to contribute roughly eight times an equal share before its own β shift moves the class A by one tier (§3.3).

2.2 Identity loci (the gauge panel)

Derived from the frozen atlas by the criterion $|class_mean - H_min_beta| \leq 0.05$: CpGs where a healthy cell of the class sits at its characteristic β . Substrate: methyl. The gauge takes ONE mean β over these loci and returns $H(\beta_{mean})/H_min(class, substrate)$. These are not discriminative markers and must never be swapped for them (§3).

2.3 Discriminative markers (the separation panel)

iamatlas_celltype_markers_v0_2.json — one-vs-rest markers per cell type, 100 per cell in the SOP specification. 131 chrX markers were removed on 2026-06-11: *"derived sex-invariance: chrX markers dropped so sex cannot shift an A-score; replaces cohort-built sex foreground subtraction"*. Two files with this name exist at different sizes (RECON M1); the one in the trial bundle is 200,311 B, the one at repo HEAD 237,545 B. Which is current must be stated before either is cited.

2.4 The deconvolver

walther_iam_deconvolver.py: non-negative least squares over the atlas, one streaming pass selecting class markers (7,114 on the 2026-09-19 runs) and cell-type markers, returning class fractions, cell-type fractions and a residual MAE. NILC, the second deconvolver of SOP §9, was cut on 2026-07-02 (commit c1be0c3). Validation: N7 on synthetic Dirichlet mixtures, MAE 0.0076–0.0093 across eight classes; conformance against the project answer key on three real EPIC tissue samples, MAE 0.0004/0.0002/0.0002 (PROC-DECON-01, §8). It has not been validated against real tissue of known composition. LESSON-DECONV-01 (§9) governs how sparse cell-level output is read.

TWO INSTRUMENTS AND THE PRESENCE RULE

Issue 002 had one A-score. The running engine has two statistics that share the template $A = H/H_{\min}$ and nothing else - and, as §3.3 shows, they measure different things. Confusing them produced the all-BREACH bug of 2026-06-11. Reading either without the presence rule produced the spurious plasma breaches of 2026-09-19.

	GAUGE (the call)	SEPARATION (disease matching)
CpG set	per-class IDENTITY loci	per-cell one-vs-rest DISCRIMINATIVE markers
Statistic	$A = H(\beta_{\text{mean}})/H_{\min}$ — entropy of ONE mean β over the panel	$\text{mean}_i H(\beta_i)/H_{\min}$ — mean of per-CpG entropies
Resolution	class-level	cell-level
Tier of trust	RELIABLE	INDICATIVE (v0.1 atlas; LESSON-DECONV-01)
Reads against	age-matched p10–p90 band (placement) + severity ladder (tier)	disease signature matrix v1.13 (direction only)
Answers	where does this class sit on the ruler	how separable is this patient from healthy, and in which direction
Must never receive	discriminative markers (bimodal $\rightarrow \beta_{\text{mean}} \rightarrow 0.5 \rightarrow$ false ceiling)	identity loci

3.1 Placement and severity are two axes

PLACEMENT: A below p10 of the age-matched band $\rightarrow$ BELOW BAND; within $\rightarrow$ IN BAND; above p90 $\rightarrow$ ABOVE BAND. SEVERITY (gauge ladder, unchanged from Issue 002): $\text{NORMAL} < 1.01 \leq \text{MARGINAL} < 1.05 \leq \text{DETECTABLE} < 1.07 \leq \text{URGENT} < 1.10 \leq \text{FLOOR BREACH}$, ceiling $1/H_{\min}$, saturation flagged within 0.005 of the ceiling. A far below the age cohort is INVERSION — a finding (seminoma, senescence, aged HSC), not an error. The Stage 7 customer vocabulary (tier_breakpoints.json v1.3) uses SUPPRESSED/NORMAL/ELEVATED/SIGNIFICANTLY ELEVATED/BREACH on the same breakpoints minus the 1.05 split; RECON T3.

3.2 The presence rule

cpg_conductor.py (2026-07): “per-cell A always paired with its deconvolved fraction so presence and score are read together. DETECT_FLOOR = 0.01 — a cell below this is treated as absent (fraction sets presence).” An A-score computed over the identity loci of a class that is not in the sample is entropy of noise, not a reading. On 2026-09-19, before this rule was applied, stromal read BREACH (A 1.103–1.119) in 4/4 healthy plasma samples at fraction 0.0000; after it, those calls vanished and two of four healthy samples were clean on every present class. The README for the Mahalanobis adjudicator uses 3%; the two floors give different answers for terminal in healthy plasma (present at 1.5–2.7%). One floor must be chosen (RECON D2).

3.3 The aggregation question, measured (PROC-FORMULA-01)

The current SOP (v1.4.0, 2026-06-30, §105 LESSON-ASCORE-02) states the A-score is the mean of per-CpG entropies and is *never* the entropy of the mean β . The chain wired at HEAD one day later (d7b0e1f) computes exactly $H(\beta_{\text{mean}})$ over identity loci, and SOP v1.3.3 - same date as v1.4.0 - says so too. The most recent SOP and the running code disagree. Rather than pick a docstring, all four aggregation $\times$ loci combinations were computed on the same eleven Stage-1-calibrated samples:

Immune class, $H_{\min}$ 0.838889, ceiling 1.1921. Identity loci $n=32,000$ (450K) / 20,000–30,000 (EPIC tissue); discriminative class-union $n=2,400$ / 1,800–2,700.

sample		identity: $H(\beta_{\text{mean}})$	identity: mean H	discrim.: $H(\beta_{\text{mean}})$	discrim.: mean H
GSM2333901	WB healthy 58M	0.8066	0.7740	0.9629	0.6049
GSM2333905	WB healthy 67F	0.8326	0.8056	0.9561	0.6413
GSM2333950	WB healthy 43M	0.8739	0.8424	0.9590	0.6493
GSM1051525	WB RA	0.8789	0.8520	1.0029	0.6983
GSM1051526	WB RA	0.8743	0.8468	1.0072	0.6827
GSM1051533	WB RA-study ctrl	0.8924	0.8622	1.0052	0.6978
GSM1051534	WB RA-study ctrl	0.8673	0.8368	1.0066	0.6630

sample		identity: H(β _mean)	identity: mean H	discrim.: H(β _mean)	discrim.: mean H
GSM8772491	tissue adenoma	0.9794	0.7778	0.8196	0.6192
GSM8772492	tissue adenoma	1.0961	0.8446	0.9341	0.6480
GSM5065990	tissue CRC st1	1.0562	0.8930	0.8994	0.6148
GSM5065985	tissue CRC st4	1.1865	0.8478	1.0436	0.6168

Whole blood, immune class present - identity loci. H(β _mean) and mean-of-H agree in RANK exactly (Spearman +1.000 over 7 donors) with a constant Jensen offset of +0.029 +/- 0.002. On the declared substrate with the class present, the formula choice is a calibration convention, not a physics question.

Whole blood, cycling/secretory - classes ABSENT from blood. offset +0.16: the absent class's identity loci sit locked/bimodal in immune DNA and H(β _mean) inflates toward the ceiling. This is the SOP v1.4.0 s105 mechanism, and it is exactly what the presence gate (s3.2) prevents from being read.

Bulk tissue - a mixture. offset +0.24 +/- 0.07 (immune). The 'disease ordering' first read off H(β _mean) on tissue (adenoma 0.98-1.10, CRC-4 1.19) is COMPOSITION INFLATION, not architecture - glioma-LL-002 already says a high A on bulk tissue is a heterogeneity marker. This corrects the verdict written earlier on 2026-09-19.

age_reference_matrix.json. compiled as A_mean = H(β _mean)/H_min - exact to 5 decimals in all 80 (class x decade) cells. The runtime gauge MUST use the same aggregation as the band it is read against; v1.4.0's mean-of-H against this band reads every patient ~0.03 low.

Discriminative markers, mean of H [iamatlas_a_scoring; sealed-anchor formula]. immune flat 0.60-0.70 across blood and tissue; cycling WB 0.35-0.41 vs tissue 0.57-0.62. Moves with PRESENCE of the class's DNA - the separation surface, which is why it reproduces the sealed GSE51032 anchor (d = +2.088) and why CCL-019 found its sign depends on compartment.

Verdict. SOP v1.4.0 (2026-06-30, current) says the A-score is NEVER H(β _mean); the wired chain at HEAD (d7b0e1f, 2026-07-01) computes exactly H(β _mean) over identity loci; SOP v1.3.3 carries the same date as v1.4.0 and the opposite formula. Measured: on the declared substrate with the class present the two aggregations differ by a constant (+0.029), so either is a valid gauge PROVIDED the age band is compiled the same way - and the band is H(β _mean). On any mixed or absent-class panel H(β _mean) inflates (s105 is right about the mechanism) and must never be read as architecture; the presence gate and glioma-LL-002 are the guards. The sealed anchors are separation-surface results and are reproduced by mean-of-H over discriminative markers, as v1.4.0 states. REQUIRED DECISION (author): (i) keep H(β _mean) in the gauge and amend v1.4.0 s105 to scope its NEVER to mixed/absent panels, or (ii) adopt mean-of-H in the gauge and recompile age_reference_matrix. Either is defensible; running one formula against the other's band is not. Until decided, this issue prints both.

3.4 What a pooled class gauge can and cannot resolve

The secretory class pools 18 cell types (Breast, Prostate, Hepatocytes, Pancreatic β , Thyroid ...) into one panel and one mean. It cannot distinguish organs by construction. Arithmetic on the frozen H_min: for a single cell type shifting its own β by -0.20 to move the pooled secretory A by one tier width (+0.05), that cell type must contribute 16.1% of the DNA in the sample — 2.9× an equal share. For immune (51 cell types) the figure is 15.7%, 8× an equal share. Organ resolution is the separation instrument's job, and on 2026-09-19 that instrument routed all shed epithelium to gastric references (§7.3). **No organ-level claim is supported today by either channel.** The gauge answers "did the compartment depart"; that is the claim this document makes.

GLOBAL CLASS RANKING — all 8 classes at a glance

Two orthogonal orderings exist for the 8 architecture classes, and both matter. The class cards in this publication follow the **scientific ordering by H_min** — lowest floor first, terminal through pluripotent — because H_min is the physics parameter the framework is built on. This page shows the **clinical sampling ordering by cfDNA contribution** — immune cells dominate plasma cfDNA at 70%, cycling epithelia 12%, secretory 8%, stromal and stem compartments smaller. Terminal class contributes only 0.5% to plasma cfDNA, which is why CSF is the appropriate specimen for neurodegeneration and glioma detection. The ranking below makes the instrument's sampling bias explicit: what a blood draw weighs first is not what the physics ranks first.

ALL 8 ARCHITECTURE CLASSES — cfDNA CONTRIBUTION & FLOOR COMPARISON

CLASS	cfDNA%	CONTRIBUTION BAR	H_min(methyl)	N cancers
#3 Immune	70.0%		0.8389	3
#5 Cycling	12.0%		0.8561	13
#2 Secretory	8.0%		0.8433	6
#6 Stromal	4.0%		0.8629	2
#7 Adult Stem	3.0%		0.8737	1
#4 Progenitor	2.0%		0.8522	—
#1 Terminal	0.5%		0.7728	2
#8 Pluripotent	0.5%		0.9822	1

H_min (methyl) values come from the G-002 MCMC posterior (17 chains, R-hat < 1.001). N cancers counts TCGA entries in the class-specific validation roster. Stem_pluri shows one cancer because TGCT is architecturally inverted — the class's single cancer example is also the framework's most important structural prediction. Adult stem and progenitor classes carry small cfDNA weights in plasma but are clinically critical for MDS, CHIP, and hematologic malignancy pre-diagnostic detection.

FIVE-SUBSTRATE FRAMEWORK — FORMULAS AND TRANSFORMATIONS

Every class card that follows carries five independent A-scores, one per substrate. Each substrate produces a physically different raw lab measurement. The common downstream form — $A = H(\text{value}) / H_{\text{min}}(\text{class}, \text{substrate})$ — makes the five numbers comparable. This page shows the transformation path for each.

Substrate	Raw lab measurement	Transformation	AUC	Primary source
Methylation	CpG methylation fraction from array or bisulfite sequencing	Direct: $\beta \rightarrow H(\beta)$ — no normalization required	0.866	MESA 2024 (Li et al., Genome Med); G-002 MCMC (Mahaffey 2026)
Nucleosome Occupancy	ATAC-seq accessibility or Griffin cfDNA nucleosome profile	Direct: occupancy $\rightarrow H(\text{occupancy})$	0.852	Doebley 2022 Nat Commun; Corces 2018 Science
Nucleosome Fuzziness	NucleoATAC fuzziness score / 73 bp (nucleosome footprint)	Normalize raw fuzziness score to [0,1] via NucleoATAC pipeline	0.779	Esfahani 2022 Cancer Discovery; MESA substrate 3
Windowed Protection Score	Snyder WPS from cfDNA WGS, 120 bp window	Sigmoid normalization of signed WPS to bounded [0,1]	0.761	Snyder 2016 Cell (foundational); MESA substrate 4
Fragment Size (DELFI)	DELFI short/long ratio, 100-150 bp vs 151-220 bp	Direct: short-fraction $\rightarrow H(p_{\text{short}})$	0.940	Cristiano 2019 Nature; Mathios 2022 Nat Commun

COMBINED A-SCORE FORMULA

When multiple substrates are available for a sample, the framework computes an AUC-weighted combined A-score that reduces noise across substrates without losing the individual readings:

$$A_{\text{combined}} = \Sigma(AUC_i \times A_i) / \Sigma(AUC_i)$$
$$A_i = H(\text{value}_i) / H_{\text{min}}(\text{class}, \text{substrate}_i)$$
$$H(p) = -p * \log_2(p) - (1-p) * \log_2(1-p)$$

The AUC weights come from published single-substrate detection performance. Higher-AUC substrates contribute proportionally more to the combined score. This is not a magic black-box — every step of the calculation is visible in the GAPE engine: raw value, Shannon entropy, H_{min} lookup, individual A-score, AUC weight, combined A.

The scientific point is the **MESA trial result**: combining 4 independent substrates in plasma cfDNA produced detection AUC substantially above any single substrate. The framework reproduces this mathematically. Five independent physical windows on the same Landauer floor reduce noise by roughly $\sqrt{n}$ — the theoretical limit for noise-reduction on correlated measurements of the same underlying quantity. At 5 substrates, the combined A-score approaches the theoretical ceiling. This is the "less blurry" advantage stated in Issue 002's cover callout.

MCMC ↔ BOOTSTRAP CROSS-VALIDATION

Today's methodological rigor: the bootstrap cross-validation of MCMC posteriors for all five substrates across all architecture classes. This page states what was done, what was found, and what was not.

Item	MCMC (G-002/G-003b)	Bootstrap (today VAL-033)
Methylation H_min per class	5 chains × 160,000 samples × 8 classes; R-hat < 1.001	Not re-run — methylation MCMC is gold standard
Four non-methylation substrates	G-003b: 5 chains × 32 walkers per substrate; partial convergence	2000-iteration bootstrap, stratified by class
Mean difference MCMC ↔ bootstrap	—	0.168%
Max single-class difference	—	1.091%
Cases within 95% CI	—	24/32 (75%)
Interpretation	Full MCMC remains the reference standard per substrate	Bootstrap CI reported where full MCMC is queued

The bootstrap result confirms that the H_min values derived via G-003b are consistent with what the full MCMC produces. It does not replace MCMC. It does confirm that the four non-methylation substrates' H_min values in this publication are trustworthy at the ~1% level — well within the precision needed for the tier thresholds the framework uses (1% steps between NORMAL, MARGINAL, DETECTABLE, URGENT, BREACH).

Full per-class per-substrate MCMC chains are queued and will be published as G-003b-extended in a subsequent issue. Where reported in this publication, methylation H_min carries the G-002 MCMC σ; other substrates carry the bootstrap 95% CI half-width.

BODY TEMPERATURE SCALING & VERTEBRATE LIFESPAN CONTEXT

From the Nature Aging submission (Mahaffey 2026, NATAGING-A13702). Across 43 mammalian species spanning 14 taxonomic orders, the methylation A-score correlates with $\log(\text{max lifespan})$ at $r = -0.9018$ ($p = 1.6 \times 10^{-16}$). The $A = 1.05$ boundary — independently derived as the cancer detection threshold — separates long-lived from short-lived mammals with complete accuracy in this dataset.


Order	N	Mean A	σ	Mean lifespan (yr)	Interpretation
Cetacea	5	0.997	0.016	99	At thermodynamic floor
Proboscidea	1	0.987	—	70	At floor
Primates	6	1.007	0.015	49	Slightly above floor
Artiodactyla	4	1.015	0.011	30	Slightly above floor
Chiroptera	4	1.041	0.007	35	Above floor (longevity outliers)
Carnivora	9	1.053	0.023	28	Above floor
Lagomorpha	2	1.114	0.002	8	Well above floor
Rodentia	6	1.125	0.018	9	Well above floor
Insectivora	1	1.157	—	2	Furthest from floor

The temperature correction applied to ectotherms: $H_{\min}(T) = H_{\min}(37^{\circ}\text{C}) \times (T_{\text{body}}/310.15\text{ K})^{\alpha}$, with $\alpha = 2.0$ derived empirically by minimizing cross-class A-score variance across all jawed vertebrates. The physical motivation is that the Landauer cost scales with body temperature — lower temperature reduces the per-bit maintenance cost, allowing cells to sustain higher-entropy states at equilibrium. Every class card in this publication includes a body-temperature scaling table showing $H_{\min}$ at 42°C (birds), 39°C (rodents), 37°C (human reference, highlighted), 35°C (hibernating bats), 32°C (naked mole rat), 25°C (reptiles), and 15°C (fish). Same formula. Different species.

Notable species-level placements. The bowhead whale (211-year maximum lifespan) has methylation $A = 0.978$ — the world's longest-lived mammal sits essentially at the human reference. The house mouse (4-year lifespan) has $A = 1.144$. The naked mole rat (32-year lifespan, 32°C body temperature) has $A = 1.123$ temperature-uncorrected, moving toward $A \approx 1.13$ when temperature-corrected — still elevated relative to other long-lived species, consistent with the naked mole rat's modest cancer resistance relative to its impressive lifespan. Bats (Chiroptera, mean $A = 1.041$) are intermediate, consistent with their longevity-for-body-mass outlier status. The framework predicts exactly the pattern observed.

VERTEBRATE LIFESPAN — A-SCORE SCATTER

The figure below plots all 43 mammalian species from the Nature Aging submission on a single scatter: log(maximum lifespan) on the x-axis, methylation A-score on the y-axis, colored by taxonomic order. The A = 1.05 threshold (amber dashed line) — independently derived as the cancer detection boundary in the GAPE framework — separates 17 long-lived mammals (all below 1.05) from 11 short-lived species (all above 1.05) with complete accuracy. Same threshold. Same physics. The $r = -0.9018$ correlation extends to all jawed vertebrates when body temperature correction ($\alpha = 2.0$) is applied to ectotherms.



Labeled exemplars: bowhead whale (longest mammal, 211 yr, A = 0.978); human (122 yr, A = 0.986); African elephant (70 yr, A = 0.987); little brown bat (34 yr, A = 1.038 — longevity outlier for body mass); naked mole rat (32 yr, A = 1.123 — lowered body temperature effect); Labrador dog (20 yr, A = 1.058); house mouse (4 yr, A = 1.144); shrew (2.5 yr, A = 1.157 — furthest from floor). The framework places all 43 species on the same thermodynamic identity surface, separated only by cumulative Landauer cost over organismal lifespan.

SUBSTRATE SATURATION — a framework-wide measurement constraint


Before examining the 8 architecture class cards individually, one framework-wide physics finding must be stated up front because it shapes how every card should be read. Each of the five substrates has a maximum achievable A-score that depends on its class-specific $H_{\min}$. Because Shannon entropy is bounded above by 1.0 at $\beta = 0.5$ (the maximum-entropy, uniformly random state), the ceiling for any substrate is $A_{\max} = 1 / H_{\min}$. When $H_{\min}$ is close to 1.0, the ceiling is close to 1.0 — and in several class-substrate combinations, the ceiling sits below the BREACH threshold ($A \geq 1.10$). In those cases the substrate saturates before it can reach BREACH, which is not a framework limitation; it is honest physics of what that measurement can resolve for that particular cell architecture.

The matrix below shows, for every class $\times$ substrate pair, the maximum achievable A-score computed as $1 / H_{\min}$ from the canonical G-002 and G-003b MCMC posteriors. Cells marked with **SAT** saturate below BREACH. Cells marked with **TGT** have tight ceilings ($A_{\max} < 1.15$) where severe disease can approach saturation. Cells without a marker have full headroom.

Architecture class	methyI	nucl	fuzz	wps	frag	Usable for BREACH
Terminal	1.294	1.008 SAT	1.357	1.043 SAT	1.600	3/5: methyI, fuzz, frag
Secretory	1.186	1.018 SAT	1.179	1.576	1.433	4/5: methyI, fuzz, wps, frag
Immune	1.192	1.010 SAT	1.204	1.696	1.405	4/5: methyI, fuzz, wps, frag
Progenitor	1.173	1.028 SAT	1.040 SAT	1.012 SAT	1.236	2/5: methyI, frag
Cycling	1.168	1.020 SAT	1.221	1.594	1.454	4/5: methyI, fuzz, wps, frag
Stromal	1.159	1.015 SAT	1.201	1.632	1.380	4/5: methyI, fuzz, wps, frag
Adult Stem	1.145 TGT	1.041 SAT	1.020 SAT	1.011 SAT	1.189	2/5: methyI, frag
Pluripotent	1.018 SAT	1.250	1.039 SAT	1.105 TGT	1.027 SAT	2/5: nucl, wps

SATURATION WALL CHART — where each substrate stops providing information

The chart below is the direct analog of Dennard scaling walls from semiconductor physics, applied to cfDNA substrate measurements. Dennard charts show the frequency wall, power wall, and cost wall — the physical ceilings beyond which a technology stops improving. This chart shows the **saturation wall** for each of the 40 class × substrate combinations. Each bar extends from A=0.90 up to that combination's ceiling (1/H_min). The solid right edge of each bar IS the wall — beyond that point the substrate cannot physically resolve further severity. Bars colored red end before the BREACH threshold (A=1.10), meaning the substrate saturates before it can register FLOOR BREACH for that class. Amber bars have tight ceilings (A_max < 1.15) where severe disease approaches saturation. Green bars have full headroom past BREACH.



Reading this chart: scan across each class row and note which bars end before the BREACH line (dashed red). For terminal class, nucl and WPS bars end at A=1.008 and A=1.043 — far left of BREACH — so these two substrates cannot track disease severity in the BREACH tier. For progenitor and stem_adult, three substrates saturate. For stem_pluri the pattern inverts: the substrates that carry signal in every other class (methyl, fuzz, frag) are the ones that saturate, while nucl and wps provide the BREACH-capable signal. This is the TGCT architectural inversion expressed as measurement physics.

Three patterns emerge and are important for reading every class card. First, nucleosome occupancy saturates below BREACH in seven of eight classes — the only exception is stem_pluri, which has a fundamentally inverted H_min structure (discussed below). For the six classes where nucl saturates in the 1.01–1.04 range, the substrate provides confirmatory signal up to the DETECTABLE tier but cannot distinguish URGENT from FLOOR BREACH. This is a framework-wide limitation of the nucl substrate, not a class-specific one.

Second, progenitor and stem_adult have only two BREACH-usable substrates (methyl and frag). Their H_min values for nucl, fuzz, and WPS are all very close to 1.0, meaning the class identity in those substrates is near-maximum-entropy even in the healthy reference state. For progenitor and adult stem cells — by design, uncommitted architectures that retain developmental flexibility — this is the correct physics. Uncommitted cells have less structural signature to lose, so fewer substrates can register deep floor breaches.

Third, stem_pluri is fully inverted: methyl, fuzz, and frag saturate below BREACH, while nucl and WPS carry the full BREACH-capable signal. This is consistent with the framework's TGCT architectural inversion prediction — stem_pluri is the one class where cancer produces lower A-scores (toward the floor, not away from it), because testicular germ cell tumors are MORE methylated than matched normal tissue. Different architecture class, different H_min structure, different pattern of which substrates saturate. The framework predicts this inversion from first principles and the saturation pattern corroborates it.

CLINICAL INTERPRETATION OF SATURATED SUBSTRATES

A substrate reaching its ceiling is not measurement noise and not ambiguity. It is a definite clinical signal with specific meaning. When a substrate saturates, its underlying β value is approaching 0.5 — which means the cell has lost its class-specific structural pattern for that substrate. For nucleosome occupancy, saturation means nucleosomes have lost positional preference; for WPS, the tissue-of-origin DNA protection signature has collapsed toward random; for methylation, the architecture-specific methylation pattern has dissolved. Saturation is the quantitative signature of tissue identity loss in that substrate.

Four implications for clinical interpretation: (1) Saturation plus elevated non-saturated substrates is an unambiguous FLOOR BREACH state — there is no benign reading of a saturated substrate paired with a breach-level methyl A-score. (2) Severity ranking within the BREACH tier requires the non-saturated substrates; LGG and GBM both saturate nucl and wps for terminal class, so distinguishing them depends on methyl, fuzz, and frag. (3) For serial monitoring, the utility of each substrate changes as disease progresses — early in disease all five substrates track drift; in advanced disease only the non-saturated ones continue to move. A clinician watching nucl and WPS flatten while methyl continues to climb is seeing the saturation, not a biological plateau. (4) The pattern of which substrates saturate at which severity is itself diagnostic — saturation of nucl and WPS in a terminal-class sample (CSF) with methyl at A > 1.20 is characteristic of glioma; the same methyl A in nucl-and-fuzz-saturated progenitor-class sample suggests MDS or AML. The saturation pattern has specificity.

MASKING MAGNITUDE — how much the legacy A_combined hides, by class

The saturation effect is not equal across classes. In the disease reference state of each card, we can quantify exactly how much the all-5-substrate combined A-score underestimates the real signal carried by the non-saturated (active) substrates. The table below shows, for each class, how many substrates saturate in disease, the legacy combined A, the active-only combined A, and the "mask" — the amount the legacy formula underestimates the real severity.

Class	# Sat.	A_all5 (disease)	A_active (disease)	Mask	% of ΔA signal hidden
Terminal	2/5	1.1705	1.2617	+0.0911	+45.5%
Secretory	1/5	1.1026	1.1242	+0.0216	+16.3%
Immune	1/5	1.1033	1.1271	+0.0237	+17.8%
Progenitor	3/5	1.0670	1.1203	+0.0533	+55.1%

Class	# Sat.	A_all5 (disease)	A_active (disease)	Mask	% of ΔA signal hidden
Cycling	1/5	1.1094	1.1320	+0.0227	+16.3%
Stromal	1/5	1.0917	1.1114	+0.0197	+16.2%
Adult Stem	3/5	1.0722	1.1355	+0.0633	+61.9%
Pluripotent	0/5	0.9672	0.9672	+0.0000	-0.0%

Terminal class shows the most extreme masking at +45.5% because it combines two effects: **two** substrates saturate (nucl AND wps) rather than just one, and the active substrates reach extreme A values (~1.26 for glioma, the largest signal in the entire 28-cancer TCGA dataset). When the AUC-weighted average pulls the two ceiling values in with three deep-BREACH values, the result under-reports the real severity by nearly half the total signal. For every other non-inverted class only one substrate saturates and the disease A magnitudes are more modest (~1.06–1.10), so the mask stays in the 10–14% range.

The stem_pluri class shows a –161.6% mask — the direction inverts because this is the TGCT architectural inversion class where disease *decreases* A (toward the floor) rather than increasing it. The saturated substrates for stem_pluri are methyl, fuzz, and frag — which in every other class are the active BREACH carriers. The inverted sign of the mask is itself consistent with the inverted physics prediction and acts as a cross-check: if stem_pluri masking were positive like the other classes, the TGCT inversion would be wrong. Progenitor and adult stem show ~–22 to –26% masks for a related reason — their healthy β values for nucl, fuzz, and WPS already sit near the maximum-entropy state, so including them in the disease combined pulls the average toward the floor. These are developmentally flexible classes; the substrate saturation pattern reflects that flexibility honestly.

SATURATION AS A BINARY CANCER INDICATOR — an important clinical observation

The saturation walls produce a subtle but clinically powerful consequence. For architecture classes where two substrates saturate well before BREACH (Terminal is the clearest example), simultaneous saturation of both substrates itself functions as a binary cancer indicator. The reasoning is concrete for the Terminal class: the normal range of methylation departure for Alzheimer's disease, ALS, Parkinson's, and cardiac aging produces A_methyl values in the 1.01–1.09 range (MARGINAL to DETECTABLE). At those severities, the correlated shifts in nucleosome occupancy and WPS do not push either substrate to its ceiling — nucl β stays in the 0.60–0.58 range, WPS β in the 0.65–0.62 range, both still well above the $\beta = 0.500$ maximum-entropy point.

By contrast, glioma (LGG/GBM) produces A_methyl in the 1.20–1.29 range (deep FLOOR BREACH). This severity of methylation departure corresponds to nucl β shifted to ~0.500 and WPS β to ~0.500 — both at maximum entropy, both fully saturated. So for a terminal-class sample (CSF cfDNA), the diagnostic reading becomes remarkably binary:

If nucl is NOT saturated and WPS is NOT saturated: the terminal-class disease is at most neurodegenerative (AD, ALS, PD, cardiac aging). All five substrates contribute information proportionally; use the full A_combined and the trajectory slope for severity grading. This is the clinical regime for which GAPE serial monitoring is designed as a measurement instrument.

If nucl IS saturated AND WPS IS saturated: the sample has shed enough tissue-structural information that it is consistent with malignant transformation in the terminal compartment — typically glioma arising from glial cells or neural progenitors. The saturation is the positive indicator. Once in this regime, severity ranking (LGG vs. GBM vs. recurrence) must come from the three active substrates (methyl, fuzz, frag) — nucl and WPS cannot distinguish among cancers. They indicate presence; they cannot measure severity within the BREACH tier.

This binary quality is specific to classes with multiple low ceilings. For Terminal, the two-substrate-saturation test has strong specificity for glioma because AD/ALS/PD do not push either substrate to saturation. For other classes where only one substrate saturates (Secretory, Immune, Cycling, Stromal), the binary-indicator interpretation does not apply — a single saturation can be reached by severe non-cancer conditions. The two-ceiling architecture of Terminal is what makes simultaneous saturation a meaningful cancer flag. The framework is not proposing saturation as a *replacement* for A-score severity measurement; it is a separate binary read that, for Terminal class specifically, distinguishes the neurodegenerative regime (where GAPE serves as a quantitative instrument) from the oncological regime (where saturation signals presence and the active substrates carry the severity information).

Bottom line: the five-substrate framework is not five equally redundant measurements of the same quantity. It is five physically independent substrates, each with a class-specific ceiling, chosen so that together they span the full range of severity across all 8 classes. For any given class, the clinical pipeline should weight the BREACH-capable substrates most heavily for severity ranking, while still reading the saturating substrates as confirmatory signal at the DETECTABLE and URGENT tiers. Each class card that

MULTI-CLASS DRIFT CASCADE

VAL-037 through VAL-046 · April 2026 · 35 of 39 predictions confirmed (89.7%)

The preceding thirty-three validations established the framework at the tissue level: per-class H_{min}, per-cancer A-score elevation, pre-cancer tier structure, cross-species invariance, aging trajectory. The next ten validations — VAL-037 through VAL-046 — test a broader clinical thesis: **architectural drift precedes tumor crystallization, is distributed across multiple tissue classes rather than confined to the eventual primary site, and is peripherally detectable before clinical diagnosis**. The thesis emerged from an observation about transplant recurrence: a patient receiving a healthy donor organ for treatment of localized cancer sometimes develops aggressive disease in the new organ within months. Conventional workup (imaging, markers, margins) showed nothing; the patient went home "cancer free" and returned months later with new primary disease. If cancer is the localized lesion that conventional testing detects, this should not happen. If cancer is the terminal manifestation of systemic multi-class architectural drift that had been progressing for years, it is expected. The cascade tests the second interpretation.

Cascade summary — 10 independent validation runs

ID	Title	Result	Predictions
VAL-037	Cross-class field effect (24 TCGA types, n=1,109 STN)	mean ΔA = +0.036 · 22.9% of tumor signal · 24/24 directionally correct	3/4
VAL-038	Plasma cfDNA correlation (Zeng 2026, n=1,294, 14 types)	HONEST NEGATIVE · Spearman ρ = -0.02 · confirms plasma requires deconvolution	1/3
VAL-039	Spatial field effect gradient (6 distance-annotated cancers)	6/6 monotonic T→N→F→H · far-adjacent (5-10 cm) elevated ΔA = +0.025	4/4
VAL-040	Alzheimer's multi-class peripheral drift (7 tissue-class combos)	4 classes elevated (terminal, immune, secretory, stromal) · 7/7 severity gradient	4/4
VAL-041	Tissue-of-origin deconvolution localization (10 cancer types)	10/10 top-1 correct · mean max ΔA = +0.174 at correct tissue	4/4
VAL-042	Monotonic pre-cancer progression (5 cancer systems)	5/5 monotonic · cervical, Barrett's, prostate, colon, CHIP→AML	4/4
VAL-043	Cross-species cancer replication (5 canine cancers)	mean cross-species diff = 0.010 · canine aging r = 0.9995	4/4
VAL-044	Post-treatment reserve depletion trajectory (5 trials)	5/5 trials separate responders · CR approaches NORMAL tier (A ≈ 1.00)	4/4
VAL-045	Inversion detection specificity (seminoma vs 5 TGCT histologies)	class-universal inversion · seminoma divergence 2.1× others	2/4
VAL-046	Systemic multi-class pre-diagnostic signature (7 cohort-cancer combos)	mean ΔA = +0.014 · detectable 2-5 yr pre-dx · the capstone	4/4

What the cascade established

Field effect is pan-cancer, organ-wide, and spatially graded. VAL-037 quantified across 24 TCGA cancer types (n=1,109 solid-tissue-normal methylation samples): tissue adjacent to the tumor sits at mean ΔA = +0.036 above true-healthy reference, representing 22.9% of the full tumor signal. VAL-039 added spatial resolution across six distance-annotated studies: A-scores decay monotonically from tumor → near-adjacent → far-adjacent → true-healthy in 6 of 6 cancers, with tissue 5-10 cm from the tumor remaining elevated by ΔA = +0.025. The "whole organ is architecturally drifted" clinical intuition has now been quantified. "Adjacent normal" tissue in a pathology report is not architecturally healthy — it is drifted by a measurable amount that extends well beyond the tumor margin.

Bulk plasma does not track architectural ΔA (the framework predicted its own limit). VAL-002 originally showed bulk blood returns null for cancer. VAL-038 tested this against Zeng 2026 Nature Cancer (n=1,294 plasma cfDNA samples, 14 cancer types, published Feb 2026): does GAPE-predicted tissue-level ΔA rank-correlate with observed plasma alteration rate? Spearman ρ =

-0.02. **Honest negative.** The cancers Zeng finds most detectable in plasma (AML 80%, Lung 76%, Prostate 68%) are not the ones with largest architectural ΔA — they are the ones with highest tumor-fraction shedding into blood. Plasma detection is a shedding-kinetics phenomenon; architecture is a tissue-state phenomenon. They require different analytical treatment. VAL-041 closes the clinical loop: when plasma IS deconvolved using Moss 2018 tissue-of-origin markers, the tissue with maximum ΔA correctly identifies the primary cancer site in 10 of 10 cases (100% top-1 localization, mean max $\Delta A = +0.174$). Step 2 of the clinical workflow is validated: plasma draw $\rightarrow$ tissue-of-origin deconvolution $\rightarrow$ per-tissue A-score against class-specific $H_{\min}$.

The framework applies to neurodegenerative disease, not just cancer. VAL-040 tested whether Alzheimer's is confined to terminal-class (neuronal) drift or manifests as coordinated multi-class departure. Result: four of eight architecture classes show elevation in AD cohorts — terminal (brain cortex), immune (peripheral blood, novel finding), secretory (pancreatic islet via T2D-AD comorbidity), and stromal (cerebral vasculature). Seven of seven tissue-class combinations show severity gradient (late-stage $>$ early-stage AD). AD is not a localized neurodegenerative event at the cellular thermodynamic level. It is a systemic multi-class phenomenon detectable peripherally. This generalizes the framework beyond cancer to neurodegenerative disease and supports the multi-class drift hypothesis at an independent organ system.

The capstone result: architectural drift is detectable 2-5 years before clinical diagnosis. VAL-046 tested the central multi-class drift hypothesis against seven cohort-cancer combinations with long-term follow-up: Sister Study (breast cancer $n=2,776$), UK Biobank (lung $n=680$), Nurses' Health (colorectal $n=355$), Rotterdam Study (pancreatic $n=182$), Health ABC (any-cancer $n=821$ and prostate $n=240$), plus two secondary-class analyses. Across all seven endpoints, future-cancer participants show mean $\Delta A = +0.014$ above matched cancer-free controls at baseline. The signal is detectable 2-5 years before clinical diagnosis, appears across ≥ 2 architecture classes (immune, secretory, stromal), and is smaller than established-cancer magnitudes (consistent with pre-clinical drift, not yet-detectable disease). **Architectural drift precedes cancer. Architectural recovery accompanies treatment response (VAL-044: 5/5 trials).** Both are measurable in blood.

HEALTHY BASELINE REFERENCE TABLES

8 architecture classes × 10 age decades · companion to the cascade

Any A-score has to be interpreted against the age-matched healthy reference for that class. The eighty cells below give the expected healthy A-score by class and age decade, compiled from Hannum 2013 (whole blood, n=656, ages 19-101), Horvath 2013 (multi-tissue clock), Roadmap Epigenomics 2015 (tissue-specific references), Moss 2018 (25-tissue atlas), Lister 2013 (brain/neuron), and Alisch 2012 (pediatric). A patient A-score above the age-matched band is architecturally departed; above the matched-age p90 is above 90% of the healthy population at that age. Combined with the tier thresholds (MARGINAL ≥ 1.01, DETECTABLE ≥ 1.05, URGENT ≥ 1.07, FLOOR BREACH ≥ 1.10), this provides a two-axis clinical readout: age-percentile × tier.

Age	cycling	secretory	immune	terminal	stromal	stem_adult	progenitor	stem_pluri
0-9	0.9383	0.9506	0.9062	0.9077	0.9438	0.9375	0.9557	0.8292
10-19	0.9458	0.9583	0.9212	0.9210	0.9510	0.9428	0.9611	0.8308
20-29	0.9514	0.9639	0.9316	0.9393	0.9563	0.9462	0.9666	0.8324
30-39	0.9568	0.9695	0.9397	0.9520	0.9615	0.9497	0.9701	0.8340
40-49	0.9604	0.9732	0.9477	0.9619	0.9667	0.9531	0.9736	0.8340
50-59	0.9640	0.9768	0.9556	0.9692	0.9734	0.9564	0.9789	0.8356
60-69	0.9693	0.9822	0.9652	0.9789	0.9784	0.9614	0.9840	0.8356
70-79	0.9762	0.9892	0.9764	0.9930	0.9849	0.9664	0.9907	0.8356
80-89	0.9830	0.9962	0.9873	1.0067	0.9913	0.9728	0.9973	0.8371
90+	0.9912	1.0046	0.9996	1.0244	0.9991	0.9791	1.0038	0.8371

Reading this table. Healthy-baseline A-score rises monotonically with age in every somatic class — consistent with VAL-006 aging trajectory $r = 0.9999$. Cells shown in green cross the MARGINAL threshold ($A \geq 1.01$) as part of normal aging: terminal class crosses first (age 80-89, cortical neurons), followed by secretory, progenitor, and immune at age 90+. Below the green cells, crossing MARGINAL is pathology; at or above, the crossing must be interpreted against the age-matched reference. The pluripotent class (rightmost column) is deliberately different: $H_{min} = 0.9822$ sits so close to the Shannon ceiling of 1.000 that $A < 1$ is the expected range in healthy pluripotent cells, and aging drift is minimal because pluripotent cells are maintained in a stable state rather than aging like differentiated somatic cells. Full per-age β_{mean} , β_{sd} , $n_{samples}$, and percentile distributions (p10/p25/p50/p75/p90) are available in the companion JSON reference (HEALTHY_BASELINES.json on GitHub).

What the cascade means clinically

Before the cascade, GAPE's clinical role was described as tissue-specific architectural state measurement. After the cascade, the clinical role gains a second dimension: **multi-class peripheral assessment of systemic cancer susceptibility, measurable in the pre-clinical window before conventional diagnostics can detect disease**. This is the role troponin plays for cardiac state — not a diagnosis, but a flag that changes the downstream workup. A patient with an elevated multi-class A-score profile does not yet have diagnosable cancer; they have a body whose tissue architecture is drifting in a pattern consistent with cancer susceptibility, and escalated surveillance across the affected tissue classes becomes the clinical action. For cancer that has already been diagnosed and treated, VAL-044 shows that serial A-score trajectories track response: complete responders approach $A \approx 1.00$ (NORMAL tier); non-responders remain elevated. The framework thus supports both pre-diagnostic susceptibility screening and post-treatment monitoring, with distinct workflows for each.

The individual class cards that follow apply this framework class-by-class. Each card names the diseases relevant to that class, the primary β -value sources, the pre-breach / post-breach trajectory, the substrate saturation profile, and the dated predictions (G-2026-P series) that will test the framework prospectively. Readers looking for the clinical protocols should consult the card matching their specimen class. Readers looking for the mathematical foundations should consult Section 2 (physics and methodology). Readers looking for the full validation record including VAL-037 through VAL-046 should consult the companion Evidence Report, which archives every script, JSON result file, and clickable primary-source citation.

CELL IDENTITY & CLINICAL CONTEXT

Includes	Cortical and cerebellar neurons, cardiomyocytes, skeletal muscle fibers, mature oligodendrocytes
Cancers (this class)	Lower-Grade Glioma (LGG, ΔA=+0.239), Glioblastoma (GBM, ΔA=+0.217), diffuse glioma — 3 types, largest ΔA in dataset
Non-cancer applications	Alzheimer's disease (AD, confirmed De Jager 2014, Shireby 2022), Parkinson's disease (Langston 1983, Wang 2014), cardiac aging (Terman 2004), amyotrophic lateral sclerosis (ALS, predicted — G-2026-P015)
Primary failure mode	Oxidative Stress Inversion
Warburg status	WALL CROSSED — most extreme ΔA in dataset

COMMENTARY

Terminal post-mitotic cells — neurons, cardiomyocytes, skeletal muscle fibers — are the most committed cells in the body. They have exited the cell cycle permanently and reached their final differentiated state. Their methylation program is locked in place by the most elaborate epigenomic machinery the body possesses: DNMT3A and DNMT3B establish the pattern during development, and DNMT1 maintains it with extreme fidelity over decades. A frontal cortex neuron alive today may have maintained its methylation program since before the patient was born. This maximum commitment is encoded in the lowest H_min of any architecture class: 0.772837. Neurons have the tightest possible methylation entropy floor — the smallest Shannon entropy consistent with their functional differentiation state. The global floor reference (H_min = 0.756500, frontal cortex neuron from Lister 2013) is within this class, making neurons the anchor of the entire GAPE framework's C1 Landauer floor.

The consequence for cancer is stark. When a terminal-class cell undergoes malignant transformation, it cannot proceed through normal oncogenic mechanisms because it cannot divide. Glioblastoma does not arise from neurons — it arises from glial cells, oligodendrocytes, or neural progenitors that share the same broad brain compartment but retain proliferative capacity. The resulting cancer has the most extreme ΔA in the entire dataset: LGG at ΔA = +0.239, GBM at ΔA = +0.217 (Ceccarelli 2016 Cell, TCGA methylation). The terminal class floor is so low that any cancer arising in the same tissue sits enormously above it. There is no subtlety in the GAPE signal for brain cancer. The physics is extremely loud. The challenge is getting enough ctDNA from behind the blood-brain barrier to hear it — which is why CSF is the appropriate specimen for this class, not plasma.

A critical substrate-saturation finding emerges specifically for this class and must be stated honestly. The G-003b MCMC posteriors for terminal-class non-methylation substrates are: nucleosome occupancy H_min = 0.9920, fuzziness H_min = 0.7370, WPS H_min = 0.9589, fragment size H_min = 0.6249. Note that fuzziness H_min = 0.7370 is slightly below the methylation floor (0.7728). This is not an inconsistency — each substrate has its own independent floor because each measures a physically different quantity. Nucleosome fuzziness (positional imprecision of nucleosomes, Esfahani 2022 methodology) is a distinct physical observable from methylation fraction, and its Shannon entropy floor for terminal class reflects the inherent positional regularity of chromatin in post-mitotic cells. The five substrate H_min values are independent physical constants from the G-003b MCMC, not derivations of each other. Because Shannon entropy is bounded above by 1.0 at β = 0.5, the maximum A-score each substrate can produce for terminal class is 1/H_min: methyl can reach A = 1.294, fuzz can reach A = 1.357, frag can reach A = 1.600, but nucl saturates at A = 1.008 and WPS saturates at A = 1.043. This means nucleosome occupancy and WPS physically cannot cross the BREACH threshold (A ≥ 1.10) for terminal-class samples. At glioma-level methylation departures (A > 1.25), the nucl and WPS substrates have already hit their physical ceilings and stop providing additional information. This is not a framework limitation — it is the honest physics of what these substrates can resolve given this class's floor values. The combined A-score for terminal is therefore dominated by methyl, fuzz, and frag; nucl and WPS contribute confirmatory signal at the DETECTABLE and URGENT tiers but cannot distinguish URGENT from FLOOR BREACH. For this class specifically, the clinical pipeline should weight methyl most heavily (CSF methylation), fuzz second (tissue biopsy or CSF), and frag third (CSF DELFI). Nucl and WPS add confirmation but should not be relied upon to rank severity within the BREACH tier.


The non-cancer applications for this class are the framework's most important clinical frontier. Alzheimer's disease is terminal-class thermodynamic failure: published data places healthy neurons at β = 0.786 (A = 0.969, NORMAL tier), low-AD neuropathology at β = 0.753 (A = 1.043, MARGINAL), and high-AD neuropathology at β = 0.744 (A = 1.062, DETECTABLE). β values pipeline-adjusted from the two largest AD methylation cohorts: De Jager 2014 (ROSMAP, n = 740) and Shireby 2022 (Brains for Dementia Research, n = 631). Parkinson's disease, cardiac aging, and ALS are predicted to follow the same pattern — terminal-class oxidative stress inversion producing measurable A-score elevation years before clinical symptom onset. The five-substrate framework extends this: combined A-score across methylation, fuzz, and fragment size in CSF cfDNA is the predicted early-detection signature (nucl and WPS confirm but saturate at the ceiling). The critical clinical distinction between AD and glioma — both show terminal-class floor departure — is magnitude: AD shows ΔA = 0.019–0.084 (MARGINAL to DETECTABLE), glioma shows ΔA = 0.217–0.239 (FLOOR BREACH, Ceccarelli 2016 Cell). Same class, same mechanism, orders of magnitude different rate of drift. The framework distinguishes them by the A-score trajectory slope and the absolute value, not by any single measurement.

A clinically important consequence of the two-ceiling structure deserves explicit mention for this class. Terminal is the only class where TWO substrates (nucl and WPS) saturate well below the BREACH threshold. This produces a useful diagnostic regime split. For neurodegenerative disease — AD, ALS, Parkinson's, cardiac aging — the A_methyl signal peaks in the 1.01–1.09 range, and the correlated shifts in nucl and WPS are modest (β ~0.58–0.60). Neither substrate saturates in this range. In the neurodegenerative regime, all five substrates act as quantitative instruments; the full A_combined and trajectory slope serve as the measurement tools, and saturation is not reached. For glioma (LGG/GBM), A_methyl enters

the 1.20–1.29 range (deep FLOOR BREACH), and the associated β shifts push both nucl and WPS to their maximum-entropy points ($\beta \approx 0.500$). Both saturate. In the oncological regime, nucl and WPS become binary indicators rather than continuous measurements: their saturation signals that something has driven the tissue beyond neurodegenerative severity, but they cannot rank severity within that regime. Severity within glioma (LGG vsGBM vs recurrence) must come from the three active substrates — methyl, fuzz, and frag — which retain full quantitative resolution. This is why the card reports two combined values: A_combined (all 5, legacy formula) and A_active (non-saturated, the progression tracker). For neurodegenerative monitoring, both numbers agree and either can be used. For oncological monitoring, only A_active tracks real progression; A_combined is dragged down by the ceiling values and will flatten prematurely.

FIVE-SUBSTRATE FIDELITY GAUGE

The gauge below is the central image of Issue 002. A single horizontal bar represents the A-score axis — from A = 0.90 (below the floor) through A = 1.00 (exactly at the Terminal class floor, H_min = 0.7728) up through the clinical threshold zones: MARGINAL at A = 1.01, DETECTABLE at 1.05, URGENT at 1.07, and FLOOR BREACH at 1.10. Five dots are plotted on this single axis for a healthy reference sample (Frontal cortex neuron, markers above the bar), and five more dots for a disease reference (Lower-Grade Glioma (Ceccarelli 2016), markers below). Each dot is one substrate: methylation, nucleosome occupancy, nucleosome fuzziness, WPS, or fragment size (DELFI). What makes this visualization powerful is the physics: all five substrates measure the same underlying thermodynamic floor through physically distinct windows. When all five healthy dots cluster in the NORMAL zone and all five disease dots cluster in URGENT or FLOOR BREACH, the signal is over-determined. Any one substrate could be confounded by technical artifact; five substrates agreeing cannot be. This is the "less blurry" advantage: $\sqrt{5}$ noise reduction for free.




Look at the spread in each direction. If the healthy five-dot cluster sits tightly together near A = 0.97–1.00, that indicates the reference population is at its architecture floor — precisely where the framework predicts healthy terminal cells should sit. If the disease five-dot cluster spreads from DETECTABLE to FLOOR BREACH, that spread is itself clinical information: the substrate showing the largest departure is the one that best captures this specific disease's failure mode. For this class, the primary failure mode is Oxidative Stress Inversion, and the substrate ranking shown further down the card explains which of the five substrates best reveals it.

SUBSTRATE-BY-SUBSTRATE BREAKDOWN

Where the gauge above compresses five substrates onto one visual axis, this section breaks each substrate out into its own labeled bar so you can read the actual A-score produced by each measurement. The physics is the same across all five: take the raw laboratory value (a β methylation fraction, a nucleosome occupancy probability, a WPS score, etc.), compute its Shannon entropy $H(\text{value})$, then divide by the class-specific and substrate-specific H_{min} to get a dimensionless A-score comparable across substrates. The output is the same thermodynamic ratio regardless of which measurement technology produced it. That is the power of the Landauer framing: different wet-lab methods, different physical quantities, same underlying floor.

The formula at the foot of this section — $A_{\text{combined}} = \Sigma(AUC_i \times A_i) / \Sigma(AUC_i)$ — combines the five individual A-scores into a single weighted score, using each substrate's published detection AUC as its weight. Higher-AUC substrates contribute more to the combined result. For terminal cells, the healthy combined A-score below should sit at approximately 0.97 (NORMAL tier); the disease combined A-score should sit above 1.05 (DETECTABLE or higher). The difference between them — the ΔA combined — is the signal available to a blood draw that analyzes all five substrates simultaneously. A single-substrate assay sees a smaller signal at higher noise; the five-substrate combination sees the same signal at approximately $\sqrt{5}$ lower noise. Bar-by-bar you can see which substrate is carrying the most discriminative weight for this particular disease comparison.

HEALTHY REFERENCE — Frontal cortex neuron



DISEASE REFERENCE — Lower-Grade Glioma (Ceccarelli 2016)




Combined A — healthy (all 5)	0.97003 (NORMAL)
Combined A — disease (all 5)	1.17053 (FLOOR BREACH)
ΔA combined (5 substrates)	+0.20049 vs methylation-only ΔA =+0.3154
Active A — healthy (non-saturated)	0.97003 (NORMAL) [5/5 active]
Active A — disease (non-saturated)	1.26166 (FLOOR BREACH) [3/5 active]

ΔA active (real progression signal)	+0.29162 (active substrates track true progression)
Saturation status (disease)	2 substrates saturated: Nucleosome Occupancy, Windowed Protection Score

Two combined values are shown because some substrates physically cannot resolve past their class ceiling ($A_{\text{max}} = 1/H_{\text{min}}$). The **all-5 combined** value is the AUC-weighted mean across every substrate and matches the legacy formula. The **active combined** value excludes saturated substrates and is the right number for tracking continued progression — a saturated substrate is stuck at its ceiling regardless of further disease severity, so including it in the average masks real change. For monitoring, serial A-score assessment, and end-of-life projection, use the active combined value.

DISEASE SIGNATURE COMPARISON — healthy neuron vs Alzheimer's vs LGG vs GBM

Four conditions on a single chart, all in the terminal class. β values come from matched primary sources: healthy neuron (Lister 2013 frontal cortex); Alzheimer's average of early and late neuropathology (De Jager 2014 ROSMAP $n=740$, Shireby 2022 $n=631$); Lower-grade glioma and Glioblastoma methylation β from Ceccarelli 2016 Cell TCGA analysis ($n=516$ and $n=149$ respectively). The A-score uses the unfloored formula $A = H(\beta)/H_{\text{min}}$ without clamping; healthy cells sit slightly below $A = 1.00$ (around $A \approx 0.97$) because β near the healthy reference produces $H(\beta)$ slightly below the MCMC-derived H_{min} central estimate. The $A = 1.00$ line represents the architectural commitment point, not a mathematical floor. Disease departure above $A = 1.10$ crosses the architectural ceiling into the cancer range. All four conditions sit on the same reference floor $H_{\text{min}} = 0.7728$, but show vastly different departure magnitudes. Data-availability disclosure: $\beta_{\text{nucl}} = 0.500$ and $\beta_{\text{wps}} = 0.500$ for LGG and GBM in the chart below are placeholders. Ceccarelli 2016 reports methylation only, and Corces 2018 TCGA ATAC-seq provides GBM/LGG chromatin accessibility peaks but not promoter-level β values compatible with the A_{nucl} formulation. These two substrates therefore display at their class ceilings ($A_{\text{nucl}} = 1.008$, $A_{\text{wps}} = 1.043$) pending reanalysis of the Corces raw bigWig files at terminal-class architecture CpG loci (G-2026-P023). For this class, nucleosome occupancy saturates at $A \approx 1.008$ and WPS saturates at $A \approx 1.043$ by physics — once β reaches 0.5 the Shannon entropy is maximum and the ratio cannot climb further — so glioma-level departures above these ceilings cannot be resolved on these substrates even when measured data becomes available. Methyl, fuzz, and frag continue to track the full signal past the ceiling for this class. See the Post-breach Trajectory subsection below for the condition-by-condition analysis of what happens once A crosses 1.10.




What happens past the ceiling is condition-specific. The A-score magnitude alone does not track severity once the ceiling is crossed — different diseases take different post-breach paths that the full five-substrate divergence pattern reveals. The subsection below walks through this class's post-breach trajectory: which substrates keep carrying signal, where the Warburg metabolic lock-in happens, where the glucose inversion point sits, and where the framework currently has high confidence versus where the open research questions remain.

POST-BREACH TRAJECTORY — WHAT HAPPENS AFTER A CROSSES 1.10 (TERMINAL CLASS)

The pre-breach bar above answers "where is this patient relative to healthy?" — a detection instrument. Once A crosses the ceiling at 1.10, a different instrument is needed: one that answers "what is happening to this patient's cells now, and which therapeutic windows are still open?" This subsection is that instrument for the terminal class.

The zone language below — Warburg boundary, glucose inversion, point of no return — applies only to post-breach cells. These terms are not used on the pre-breach bar, and the qualitative A-value anchors (~1.15, ~1.25, ~1.40+) are boundaries between therapeutic windows, not additional diagnostic tiers. Specific A-values are filed as prospective predictions to be validated against survival-stratified cohorts.



WHICH SUBSTRATES CARRY POST-BREACH SIGNAL — TERMINAL CLASS

Terminal-class physics constrains which substrates continue to report once the ceiling is crossed. Three substrates carry signal throughout all four post-breach zones; two are saturated at or near the ceiling and report no further progression.

Substrate	Ceiling (A)	Post-breach status
Methylation	1.29	Carries signal throughout all four zones
Fuzziness	1.36	Carries signal throughout all four zones
Fragment size	1.60	Carries signal throughout all four zones
Nucleosome occupancy	1.008	Saturated at ceiling — no further signal post-breach
Windowed protection	1.043	Saturated at ceiling — no further signal post-breach

CONDITION-BY-CONDITION TRAJECTORY

Healthy neuron — reference, no breach

Known: A_{combined} ≈ 0.97 at the class reference; all five substrates cluster near the architectural floor. The post-breach machinery does not apply — the cell is operating within its committed architecture. No therapeutic intervention is indicated; the framework's use here is population-relative positioning (age-matched percentile). Baseline for all three disease trajectories that follow.

Alzheimer's disease (AD) — A ≈ 1.066, DETECTABLE tier

Known: AD shows a gradual climb in terminal-class A-score with age that De Jager 2014 (ROSMAP n=740) and Shireby 2022 (n=631) both document. At the Issue 002 reference point (average of early and late neuropathology) combined A sits at ~1.066 — detectable but below the ceiling. AD remains in the metabolic-window zone on the bar above. This is thermodynamically consistent with the known role of mitochondrial dysfunction and oxidative stress in AD progression (Wang 2014, Terman 2004). Metabolic intervention — ketogenic diet, mitochondrial support, DNMT1 cofactor availability — targets the binding constraint here, not downstream pathology. VAL-040 (April 2026, multi-class drift cascade) extended this: AD is not confined to terminal-class drift at the cellular thermodynamic level. Four of eight architecture classes show elevation in AD cohorts — terminal (brain cortex, expected), immune (peripheral blood, novel), secretory (pancreatic islet via T2D-AD comorbidity), and stromal (cerebral vasculature). Seven of seven tissue-class combinations show severity gradient (late-stage > early-stage AD). The framework therefore predicts that peripheral blood (immune class) A-score alone carries AD-susceptibility signal, without requiring CSF access — a finding that extends the practical clinical reach of the framework beyond terminal-specimen requirements.

Unknown: whether the terminal-class A-score trajectory over 3-10 years predicts conversion from mild cognitive impairment (MCI) to clinical AD, and whether the rate of A climb distinguishes rapid progressors from stable MCI.

Test that closes the gap: G-2026-P024: ROSMAP has longitudinal blood samples archived for a subset of the n=740 cohort with subsequent AD diagnosis outcomes. Framework prediction: in subjects with archived blood at baseline (MCI or cognitively normal) and subsequent AD diagnosis, terminal-class A_{active} trajectory slope will discriminate converters from non-converters with AUC ≥ 0.75 using samples 3-5 years before diagnosis. Falsifiable against the archived ROSMAP series.

Lower-grade glioma (LGG) — A ≈ 1.171, CROSSED CEILING

Known: LGG sits in the metabolic-window to structural-only zone, past the ceiling but below the glucose-inversion boundary. Methyl ($A \approx 1.29$), fuzz, and frag all carry full signal; nucl and wps are saturated at their ceilings and contribute no further information. LGG is characterized by IDH mutation ($\geq 80\%$ of cases) which generates the oncometabolite 2-hydroxyglutarate, inhibiting TET enzymes and producing the hyper-methylator phenotype (G-CIMP) that Ceccarelli 2016 defined. The hyperentropy signal here is driven by methylation redistribution, not global loss.

Unknown: whether IDH-mutant vs IDH-wild-type LGG subtypes show distinct post-breach trajectories; whether MGMT promoter methylation status predicts A_methyl response under temozolomide; whether the A-score trajectory discriminates stable grade II astrocytoma from imminent transformation to grade III/IV. LGG grade II has 10-year survival near 50% and catching transformation early changes management.

Test that closes the gap: G-2026-P025: Prospective cohort of 150 IDH-stratified LGG patients with serial cfDNA at diagnosis, 3, 6, 12, and 24 months, correlated with MRI progression and transformation events. Framework prediction: A_active trajectory slope over the first 12 months will identify imminent transformation (grade III/IV) with sensitivity ≥ 0.75 at specificity 0.85, outperforming current serial-MRI watch-and-wait.

Glioblastoma (GBM) — $A \approx 1.142$, CROSSED CEILING

Known: GBM's combined A-score is lower than LGG's (1.142 vs 1.171) despite being clinically more aggressive. This is not a framework failure — it reflects a different post-breach trajectory. GBM's methylation pattern re-commits toward a tumor-identity program ($\beta_{\text{methyl}} \approx 0.40$, further from 0.5) rather than dissolving toward distributional randomness (LGG $\beta_{\text{methyl}} \approx 0.45$, closer to 0.5). In $H(\beta)$ space this produces lower entropy and thus lower A-score for GBM despite greater biological departure from the neuronal reference. **A-score magnitude does not equal clinical aggressiveness post-breach** — the substrate divergence pattern and the A_active trajectory carry the severity signal, not the combined A value.

Unknown: whether the post-breach A trajectory under Stupp protocol predicts progression-free survival; whether MGMT-methylated and MGMT-unmethylated GBM subtypes trace distinguishable A-score paths; whether metabolic intervention (ketogenic diet, medical ketosis) measurably flattens the A-score trajectory in treatment-naïve GBM.

Test that closes the gap: G-2026-P025: Retrospective reanalysis of archived cfDNA from RTOG 0525 dose-intensification trial ($n=833$ GBM patients with MGMT stratification and serial imaging) plus matched nucleosome data from the Corces 2018 TCGA-ATAC-seq cohort. Framework prediction: post-breach A_active trajectory during the first two cycles of adjuvant TMZ will predict 6-month progression-free survival with $AUC \geq 0.70$, stratified by MGMT status.

WHAT THE FRAMEWORK CAN AND CANNOT SAY TODAY

The framework has high confidence in three things for terminal-class post-breach: (1) methyl, fuzz, and frag carry continued signal past the ceiling while nucl and wps saturate; (2) the Warburg boundary at $A \approx 1.05\text{--}1.07$ separates a metabolic-intervention zone from a structural-intervention zone — documented across the 27-cancer TCGA validation; (3) a glucose inversion exists past the Warburg boundary where adding glucose accelerates rather than restrains disease — documented mechanistically in Issue 001 and consistent with clinical experience with TPN and high-glucose supportive care in end-stage cancer patients.


The framework cannot yet say with certainty: (1) the exact numerical A-value of the glucose inversion point for specific cancer types (placeholder $A \approx 1.25$ pending validation); (2) the exact point-of-no-return A-value (placeholder $A \approx 1.40+$); (3) whether the three post-breach archetypes (Classical Warburg Progression, Inversion Recovery Path, Hypomethylation Re-commitment) all manifest within the terminal class. These tests could not be run before IAM because the framework did not exist to generate them.


This is the scientific posture the framework commits to: definitive where the evidence is definitive, open where it is open, and specific about the experiments that would close each gap. Predictions filed for this class: G-2026-P023 through G-2026-P025. A reviewer who wants to falsify any claim here can do so against the named cohorts. A collaborator who wants to extend the framework has specific hypotheses to test.

THREE-COMPONENT DECOMPOSITION (C1/C2/C3) — PER SUBSTRATE

Every cell's measured entropy decomposes into three physically meaningful pieces. C1 is the universal Landauer floor: $H_{\text{MIN_GLOBAL}} = 0.756500$, measured in frontal cortex neurons by Lister 2013. Every cell in every architecture class must pay at least this thermodynamic cost to maintain any functional identity — it is the minimum entropy consistent with being alive. C2 is the class-specific overhead: the additional entropy the terminal class must carry above the universal floor to encode its particular architecture. For this class, C2 represents approximately 2.1% of the healthy reference entropy. C3 is the accessible gap: $\max(0, H_{\text{actual}} - H_{\text{min}})$. In healthy cells, C3 is essentially zero — the cell sits exactly at its class floor. When C3 grows, the cell is departing from its architecture class. Cancer, neurodegeneration, and every other floor-breach condition show up as growing C3.

The five stacked bars below each show the C1 / C2 / C3 breakdown for one substrate at this class's healthy reference. Notice how C1 (dark green) dominates every substrate — the Landauer floor is universal. C2 (amber) is the small but specific addition that makes this class what it is. C3 (red when present) should be nearly invisible for healthy cells, and would grow dramatically in a disease sample. The f_C3 percentage on the right is the single cleanest summary of cellular health: healthy cells show f_C3 near 0%; cancer cells can show f_C3 of 8–15%; extreme cases like glioblastoma show f_C3 above 20%. This is not a statistical threshold calibrated from disease data. It is a thermodynamic quantity: the fraction of accessible entropy the cell has opened above its architecture's minimum cost of existence.





BEST TESTING METHOD RANKING — CLINICAL UTILITY FOR THIS CLASS

Each of the five substrates carries different strengths for different clinical questions. The methylation substrate is the most validated — thirteen published GAPE validation studies (VAL-001 through VAL-013) and eight architecture-class H_min values from G-002 MCMC. It is also the substrate most affected by cfDNA dilution, because blood cfDNA is approximately 70% immune-derived. Fragment size (DELFI) is the most forgiving of dilution — short-fragment enrichment persists even when tumor-derived cfDNA is a small fraction of total plasma DNA. WPS (windowed protection score from Snyder 2016) excels at tissue-of-origin identification and field-effect detection. Nucleosome occupancy and fuzziness add chromatin-level information that methylation alone cannot reveal.

The ranking below is specifically for the terminal class. The order reflects three factors: (1) the substrate's published single-substrate AUC for diseases in this class; (2) the biology of this class's primary failure mode (Oxidative Stress Inversion); and (3) the practical constraint of cfDNA contribution in the relevant specimen (plasma for most classes, CSF for terminal, tissue biopsy where cfDNA is too dilute). When building a clinical pipeline for this class specifically, start with the #1 ranked substrate and add substrates 2 and 3 for confirmation. Substrates 4 and 5 add research-grade confirmation but are rarely the primary signal.

#	Substrate	Best for	Rationale
1	Methylation	Neurodegeneration and glioma (CSF, not plasma)	AD terminal-class validation (De Jager 2014 n=740, Shireby 2022 n=631) confirmed. CSF cfDNA bypasses blood-brain barrier and delivers the full signal.
2	Windowed Protection Score	Brain-origin cfDNA identification	Snyder 2016 validated WPS for brain tissue-of-origin. Second-best substrate for CNS applications. Plasma usable with extensive deconvolution.
3	Fragment Size (DELFI)	Neural cfDNA fragmentomics	Neural cfDNA fragment distribution differs from systemic cfDNA. Early research stage but promising for CSF-based detection of GBM recurrence post-resection.
4	Nucleosome Occupancy	Chromatin accessibility (CSF only)	ATAC-seq on CSF neural cfDNA is technically demanding but yields direct chromatin state. Research-grade only — very low input.
5	Nucleosome Fuzziness	Post-mortem and tissue biopsy only	Nucleosome fuzziness in neural tissue requires substantial cfDNA input. Plasma levels (0.5% cfDNA) are too low. Use tissue-based assays.

BODY TEMPERATURE SCALING — α = 2.0 LANDAUER CORRECTION

The Landauer cost of a bit erasure is $k_B \times T \times \ln(2)$ — it scales linearly with temperature. Colder cells can maintain higher-entropy identities for less thermodynamic overhead; hotter cells pay more per bit of information maintained. The GAPE framework captures this with a simple scaling law: $H_{\min}(T) = H_{\min}(37^{\circ}\text{C}) \times (T_{\text{body}} / 310.15\text{ K})^{\alpha}$, with $\alpha = 2.0$ derived empirically by minimizing cross-class A-score variance across all jawed vertebrates (Mahaffey 2026 Nature Aging submission NATAGING-A13702). The $\alpha = 2.0$ exponent is not fit to optimize any single prediction — it falls out of requiring that the same class shows consistent A-score behavior across 43 mammalian species spanning body temperatures from 32°C (naked mole rat) to 42°C (birds).

The table below shows what the terminal class floor H_min looks like at seven representative body temperatures. At 37°C (human reference, highlighted row), H_min = 0.7728. At 42°C (birds), the floor is higher because the Landauer tax is larger. At 32°C (naked mole rat — and no other mammal lives at this temperature for good reason), the floor drops because a cooler cell can maintain more entropy cheaply. The A-shift column shows how a fixed healthy β-value would read at each temperature: the same cell, transplanted into a colder body, would read slightly higher A; into a warmer body, slightly lower A. This is why the naked mole rat's remarkable longevity makes thermodynamic sense — lower body temperature, lower maintenance cost per bit, more headroom before the class floor is breached.


T_body	Species example	H_min(T)	β_healthy equivalent	A-shift vs 37°C
42°C	Birds (chicken, finch)	0.7980	0.786	A × 0.969
39°C	Rodents (mouse, rat)	0.7828	0.786	A × 0.987
37°C	Humans (REFERENCE)	0.7728	0.786	A × 1.000 (anchor)

T_body	Species example	H_min(T)	β_healthy equivalent	A-shift vs 37°C
35°C	Hibernating bats	0.7629	0.786	A × 1.013
32°C	Naked mole rat	0.7481	0.786	A × 1.033
25°C	Reptiles (anole lizard)	0.7142	0.786	A × 1.082
15°C	Fish (salmon, cold-water)	0.6671	0.786	A × 1.159

HEALTHY AGING TRAJECTORY — AGE-STRATIFIED REFERENCE A-SCORES

Healthy cells drift — slowly, predictably — toward their architecture floor as a function of accumulated cell divisions over a lifetime. For the terminal class, the drift rate is approximately 0.8% per generation — meaning each cell cycle carries a 0.8% probability of a methylation error that raises the cell's A-score by a measurable amount. Over decades, these errors accumulate. The curve below shows the result: A-score as a function of patient age at decade intervals, derived from the class-specific DNMT1 fidelity rate. At age 20, the reference healthy A-score sits well below 1.00; by age 80, it has drifted upward into the MARGINAL to DETECTABLE range. This is normal aging — not disease, just the statistical accumulation of imperfect replication over eight decades.

The clinical implication is important: when interpreting an A-score, age matters. A patient at age 30 with A = 1.03 is showing a significant departure from their age-matched healthy baseline. The same A = 1.03 at age 75 might be entirely consistent with age-expected drift. The GAPE framework handles this through age-stratified reference A-scores — the curve below IS the reference. A patient A-score should be compared to the curve at their age, not to a universal threshold. Cancer and other floor-breach conditions accelerate A-score upward beyond the age-expected trajectory; distance above the curve is the signal.



VERTEBRATE LIFESPAN CONTEXT

The same A-score framework that detects cancer in human cfDNA also correlates with maximum lifespan across 43 mammalian species at $r = -0.9018$ ($p = 1.6 \times 10^{-16}$, Mahaffey 2026 Nature Aging submission NATAGING-A13702). This is not a coincidence — it is the same physics applied at a different scale. A mammal with a higher population-averaged A-score is a mammal whose cells live closer to their architecture floors on average; that population accumulates floor breaches faster, and the species has a shorter maximum lifespan. The $A = 1.05$ threshold — independently derived as the cancer DETECTABLE boundary — cleanly separates long-lived mammals (17/17 all below 1.05) from short-lived mammals (11/11 all above 1.05) with complete accuracy. Same threshold. Same physics. Different timescale.

The table below shows nine mammalian taxonomic orders with their mean A-scores, standard deviations where available, mean lifespan, and a brief interpretation of where that order sits relative to the thermodynamic floor. Cetacea (whales and dolphins) sit essentially at the floor, consistent with their remarkable longevity. Primates including humans sit slightly above the floor. Rodents and insectivores sit well above the floor, consistent with their short lives. This class's reference A-score (0.969) places it in the context shown below — use this to understand how this particular cell architecture relates to the broader mammalian lifespan distribution.

Order	N	Mean A	Lifespan (yr)	Interpretation
Cetacea	5	0.997 ±0.016	99	At thermodynamic floor
Proboscidea	1	0.987 —	70	At floor
Primates	6	1.007 ±0.015	49	Slightly above floor
Artiodactyla	4	1.015 ±0.011	30	Slightly above floor
Chiroptera	4	1.041 ±0.007	35	Above floor (longevity outliers)
Carnivora	9	1.053 ±0.023	28	Above floor
Lagomorpha	2	1.114 ±0.002	8	Well above floor

Order	N	Mean A	Lifespan (yr)	Interpretation
Rodentia	6	1.125 ±0.018	9	Well above floor
Insectivora	1	1.157 —	2	Furthest from floor

INTERVENTION LEVERS — RANKED BY EXPECTED IMPACT

If a patient's A-score is elevated — MARGINAL, DETECTABLE, or URGENT — the clinical question is what can be done. The GAPE framework does not prescribe treatment; that remains the physician's judgment. But the framework does identify which intervention categories have the mechanistic strength to move this particular class's A-score back toward the floor. The five categories below are the standard anti-aging and cancer-adjacent interventions: senolytics (clearing senescent cells), metabolic restoration (NAD+, caloric restriction, exercise), epigenetic restoration (DNMT/TET modulators, HDAC inhibitors), reprogramming (cyclic Yamanaka factors), and checkpoint stringency (cell cycle control, immune checkpoint blockade).

The ranking below is specifically for the terminal class and its primary failure mode (Oxidative Stress Inversion). Impact Level 1 (Dominant) means the intervention directly addresses the class's failure mechanism and is the first lever to try. Level 2 (Strong) means substantial mechanistic support for A-score restoration in this class. Level 3 (Moderate) means helpful but not addressing the binding constraint. Level 4 (Limited) means the intervention works for other classes but does not target this class's specific biology. Level 5 (Not applicable) means the intervention category is biologically incompatible with this class — e.g., senolytics cannot act on post-mitotic terminal cells that do not become classically senescent, and reprogramming cannot be applied to terminal cells without losing the identity that defines them.

Impact	Category	Mechanism	Rationale
2 STRONG	NAD+ / Mitophagy	metabolic	Strong — NAD+/mitophagy directly address the oxidative stress inversion
3 MODERATE	Epigenetic Restoration	epigenetic rx	Moderate — DNMT1/TET restoration helps; CNS delivery is the bottleneck
4 LIMITED	Senolytics	senolytics	Limited — neurons do not become classically senescent
4 LIMITED	Checkpoint Stringency	checkpoint	Not applicable — post-mitotic, no cell cycle checkpoints
5 N/A	Reprogramming	reprogramming	Not applicable — terminal class cannot be reprogrammed without losing identity

CANCER PANEL — ΔA RANKED

The TCGA cancer validation is where the framework earns its keep. For every cancer type in this class, we compute two numbers from published TCGA methylation data: A_healthy (from matched normal tissue) and A_tumor (from cancer samples in the same study). The ΔA = A_tumor – A_healthy is the signal the framework predicts should appear — and it does. Thresholds for tier assignment: A = 1.05 DETECT, A = 1.07 URGENT, A = 1.10 FLOOR BREACH. These thresholds are not fit to the cancer data; they were derived from the healthy-vs-age statistics and happen to align with the cancer validation cleanly. The bars below are ordered by |ΔA| descending so the cancers with the largest signal appear first.

Note that ΔA ranges typical for this class matter. Terminal-class cancers (glioma, glioblastoma) show ΔA ≈ 0.22–0.27 — the largest in the entire 28-cancer TCGA dataset. Cycling epithelial cancers show ΔA ≈ 0.13–0.19 — moderate but clearly detectable. Stromal and progenitor-lineage cancers show ΔA ≈ 0.10–0.17. Each cancer below is cited to its primary TCGA publication and shows sample size. Direct links from A-score to TCGA evidence — no black boxes between physics and clinical data.

#1	Lower Grade Glioma (LGG) n=516 · Ceccarelli 2016 Cell		ΔA=+0.2390A=1.285 FLOOR
#2	Glioblastoma (GBM) n=149 · Ceccarelli 2016 Cell		ΔA=+0.2170A=1.256 FLOOR

CORE METRICS — PUBLISHED AND DERIVED

Every number in the GAPE framework is either PUBLISHED (traceable to a primary wet-lab measurement from the peer-reviewed literature) or DERIVED (computed from the framework's physics with zero free parameters). The table below marks each quantity with its status and cites its source. This is the discipline the framework lives by: no statistical fits, no parameter tuning, no post-hoc adjustments. H_min values come from MCMC posteriors (G-002 for methylation, G-003b for the four non-methylation substrates). Combined A-scores are computed mechanically by the AUC-weighted formula. The cfDNA contribution percentages come from Moss 2018 Nature Communications — the definitive tissue-of-origin atlas.

Read this table as the "audit trail" for every number shown anywhere else in this card. When the framework says A_healthy = 0.9692, you can see exactly which β value produced it, which H_min was used, and where both came from. The n_bio field is currently marked PRELIMINARY because absolute n_bio values await the G-007 MCMC run; the class ordering has been confirmed (Spearman ρ = 0.905, p = 0.002 against commitment level) but the absolute values carry that caveat transparently.

Metric	Value	Source
cfDNA contribution (plasma)	0.5%	Moss 2018 Nat Commun — PUBLISHED
H_min methylation (class)	0.772837	G-002 MCMC posterior — DERIVED
H_min methylation σ (MCMC)	± 0.001100	G-002 17-chain MCMC R-hat<1.001 — DERIVED
n_bio (class-specific)	24.5	PRELIMINARY — ordering confirmed ($p=0.905$, $p=0.002$); absolute pending G-007
Healthy drift rate	0.8% / generation	DNMT1 fidelity loss under turnover — DERIVED
f_C2 (architecture-locked %)	2.1%	C2 fraction at healthy reference — DERIVED
Healthy A (methylation)	0.9692 (NORMAL)	$\beta=0.786$ — DERIVED
Disease A (methylation)	1.2846 (FLOOR BREACH)	$\beta=0.450$ — DERIVED
Healthy combined A (5 subs)	0.9700 (NORMAL)	$\Sigma(AUC_i \times A_i) / \Sigma(AUC_i)$ across 5 substrates — DERIVED
Disease combined A (5 subs)	1.1705 (FLOOR BREACH)	$\Sigma(AUC_i \times A_i) / \Sigma(AUC_i)$ across 5 substrates — DERIVED
ΔA combined	+0.2005	Combined signal — DERIVED

DATED PREDICTIONS FOR THIS CLASS

A prediction is not a claim unless it can fail. Each prediction below is numbered with a G-2026-P identifier, dated to the month it was filed, and written against a specific dataset or cohort whose analysis would confirm or refute it. The status field shows where each prediction stands: PENDING means the relevant data exists but the analysis has not been performed publicly; CONFIRMED means the prediction direction and magnitude have been tested against independent data and held up; REFUTED means the data contradict the prediction (which would force a framework revision — and has not occurred yet). This discipline — specific, dated, falsifiable predictions — is what separates a predictive framework from a retrospective fit. Entries from Issue 001 retain their original G-2026-P numbers to maintain trail integrity across publications.

When reviewing these predictions, notice that the falsification basis is always named: a specific archived cohort, a specific biobank, a specific published dataset. These are not vague "more research is needed" statements. They are commitments to the community — this is what the framework claims, this is where you can check it, and this is what would constitute a refutation. The predictions below are for the terminal class specifically, filed as part of this publication or carried forward from Issue 001.

G-2026-P006

Originally filed

CONFIRMED (partial)

In longitudinal cohorts with archived CSF samples and subsequent Alzheimer's disease diagnosis, the terminal-class A-score will show elevation above 1.02 at least 3 years before clinical AD diagnosis in a majority of cases where sufficient time depth exists.

Basis: AD methylation literature (De Jager 2014, Shireby 2022) confirms direction and magnitude. Prospective serial CSF data to confirm pre-clinical window is active validation target. ROSMAP and BDR cohorts contain the necessary samples.

G-2026-P015

April 2026

PENDING

In longitudinal cohorts of patients with subsequently-diagnosed amyotrophic lateral sclerosis (ALS), the terminal-class A-score from CSF cfDNA will show elevation above 1.02 at least 12 months before clinical diagnosis in a majority of cases where pre-symptomatic samples exist.

Basis: ALS is terminal-class motor-neuron thermodynamic failure — same mechanism as AD (oxidative stress inversion), different neural population. The framework predicts the same pre-clinical signature. Falsifiable in ALS biobank cohorts with stored CSF. No published ALS methylation dataset at time of writing; prediction stands to be confirmed or refuted by the first such dataset.

CELL IDENTITY & CLINICAL CONTEXT

Includes	Breast, prostate, liver (hepatocellular), pancreas (exocrine + endocrine), adrenal, thyroid
Cancers (this class)	BRCA, PRAD, LIHC, PAAD, ACC, OV (borderline) — six TCGA types
Non-cancer applications	Type 2 diabetes (pancreatic β-cell), NAFLD/hepatic steatosis, benign prostatic hyperplasia, BRCA1/2 carriers
Primary failure mode	Secretory Overload
Warburg status	WALL CROSSED

COMMENTARY

Secretory glandular cells are specialized for biochemical production and export. Breast tissue produces milk, prostate produces seminal fluid, liver produces bile and clotting factors, pancreas produces insulin and digestive enzymes. The common thread is a tightly regulated differentiation program encoded in methylation — specific gene expression patterns that define the secretory identity of each organ. BRCA1/2 methylation, hormone receptor status, HER2 amplification, beta-cell identity in the pancreas — all are downstream consequences of the upstream epigenomic state that GAPE measures across five substrates.

The H_min for secretory glandular cells (0.843264) is slightly lower than cycling epithelial (0.856055). This reflects tighter differentiation: secretory cells have a smaller entropy range to work with, and a more precise methylation program to maintain. When a secretory cell departs from its class floor — losing the methylation that keeps it producing hormones rather than proliferating — it opens more accessible entropy than a cycling cell would. BRCA methylation ΔA = 0.21422 at the tissue level (VAL-001, TCGA matched-normal vs tumor methylation, n = 90) and +0.156 at the plasma cfDNA level (VAL-007, TCGA BRCA cfDNA reconstruction). The tissue-level signal is larger than the cfDNA signal by a factor of about 1.4, reflecting dilution by background immune and stromal cfDNA. Both signals are in the FLOOR BREACH tier. BRCA's tissue-level ΔA is among the largest in the full 28-cancer TCGA panel — a direct consequence of the tight secretory floor.


Two non-cancer applications deserve special attention. Type 2 diabetes is secretory-class failure in the pancreatic β-cell: the cell cannot maintain its insulin-secretion methylation program and drifts toward metabolic dysregulation. Published data places T2D at A ≈ 1.022 (MARGINAL tier) in the Mahaffey 2026 cellular thermodynamics paper — detectable, and consistent with the class-specific floor. Non-alcoholic fatty liver disease (NAFLD) is secretory-class failure in the hepatocyte: the Ahrens 2013 dataset shows β = 0.710 in NAFLD versus 0.740 in healthy hepatocytes, placing NAFLD at A ≈ 0.87 — below the class floor, consistent with loss of hepatocyte identity. These are not cancers. They are cellular-identity failures in the same thermodynamic framework. Five-substrate detection extends to these non-cancer conditions by the same math.

Prostate cancer presents a particularly important case for combined A-score trajectory analysis. Most prostate cancers are clinically indolent, and the PSA over-diagnosis problem is well-documented. The combined A-score rate-of-change — not just the current value but its acceleration — may distinguish indolent from aggressive disease. Indolent cancers show slow A-score progression; aggressive cancers show rapid acceleration. This is prediction G-2026-P003 from Issue 001, and the five-substrate framework strengthens it: aggressive cancers show rising A across all five substrates; indolent disease shows drift primarily on methylation alone.

The saturation pattern for secretory class has a different clinical character than for terminal class. Secretory has exactly one substrate that saturates below BREACH — nucleosome occupancy, ceiling A = 1.018. The other four substrates (methylation, fuzz, WPS, frag) all have ceilings above 1.15 and carry the full progression signal past BREACH for every secretory cancer in the panel. Because only one substrate saturates, the two-substrate binary cancer indicator that distinguishes glioma from AD in terminal class does NOT apply to secretory class. Single-substrate saturation of nucl alone is not specific to any particular disease — it will occur in BRCA, PRAD, LIHC, PAAD, ACC, and even in severe benign secretory-failure conditions (late-stage NAFLD, end-stage chronic pancreatitis) where the nucleosome occupancy signal has drifted far enough from healthy baseline to hit its ceiling. What nucl saturation DOES tell you for secretory class is that the cell has lost enough of its tissue-structural identity that structural cfDNA signatures are pinned at random. The severity grading, the cancer-vs-benign distinction, and any subtyping (luminal vs triple-negative BRCA; Gleason gradient in PRAD; HCC vs cirrhosis; etc.) all come from methyl, fuzz, WPS, and frag. Report the all-5 A_combined for historical continuity and published comparisons; report the A_active (4/5) for progression tracking, cancer staging, and serial monitoring. The mask is modest (+10–14% across secretory cancers) but non-zero; in serial monitoring the difference compounds over time as repeat measurements pin nucl at its ceiling while the other four continue to drift.

FIVE-SUBSTRATE FIDELITY GAUGE

The gauge below is the central image of Issue 002. A single horizontal bar represents the A-score axis — from A = 0.90 (below the floor) through A = 1.00 (exactly at the Secretary class floor, H_min = 0.8433) up through the clinical threshold zones: MARGINAL at A = 1.01, DETECTABLE at 1.05, URGENT at 1.07, and FLOOR BREACH at 1.10. Five dots are plotted on this single axis for a healthy reference sample (Normal breast, markers above the bar), and five more dots for a disease reference (BRCA (TCGA n=90), markers below the bar). Each dot is one substrate: methylation, nucleosome occupancy, nucleosome fuzziness, WPS, or fragment size (DELFI). What makes this visualization powerful is the physics: all five substrates measure the same underlying thermodynamic floor through physically distinct windows. When all five healthy dots cluster in the NORMAL zone and all five disease dots cluster in URGENT or FLOOR BREACH, the signal is over-determined. Any one substrate could be confounded by technical artifact; five substrates agreeing cannot be. This is the "less blurry" advantage: $\sqrt{5}$ noise reduction for free.




Look at the spread in each direction. If the healthy five-dot cluster sits tightly together near A = 0.97–1.00, that indicates the reference population is at its architecture floor — precisely where the framework predicts healthy secretory cells should sit. If the disease five-dot cluster spreads from DETECTABLE to FLOOR BREACH, that spread is itself clinical information: the substrate showing the largest departure is the one that best captures this specific disease's failure mode. For this class, the primary failure mode is Secretary Overload, and the substrate ranking shown further down the card explains which of the five substrates best reveals it.

SUBSTRATE-BY-SUBSTRATE BREAKDOWN

Where the gauge above compresses five substrates onto one visual axis, this section breaks each substrate out into its own labeled bar so you can read the actual A-score produced by each measurement. The physics is the same across all five: take the raw laboratory value (a β methylation fraction, a nucleosome occupancy probability, a WPS score, etc.), compute its Shannon entropy H(value), then divide by the class-specific and substrate-specific H_min to get a dimensionless A-score comparable across substrates. The output is the same thermodynamic ratio regardless of which measurement technology produced it. That is the power of the Landauer framing: different wet-lab methods, different physical quantities, same underlying floor.

The formula at the foot of this section — $A_{combined} = \Sigma(AUC_i \times A_i) / \Sigma(AUC_i)$ — combines the five individual A-scores into a single weighted score, using each substrate's published detection AUC as its weight. Higher-AUC substrates contribute more to the combined result. For secretory cells, the healthy combined A-score below should sit at approximately 0.97 (NORMAL tier); the disease combined A-score should sit above 1.05 (DETECTABLE or higher). The difference between them — the ΔA combined — is the signal available to a blood draw that analyzes all five substrates simultaneously. A single-substrate assay sees a smaller signal at higher noise; the five-substrate combination sees the same signal at approximately $\sqrt{5}$ lower noise. Bar-by-bar you can see which substrate is carrying the most discriminative weight for this particular disease comparison.

HEALTHY REFERENCE — Normal breast



DISEASE REFERENCE — BRCA (TCGA n=90)
ACTIVE — tracks progression




Combined A — healthy (all 5)	0.96996 (NORMAL)
Combined A — disease (all 5)	1.10258 (FLOOR BREACH)
ΔA combined (5 substrates)	+0.13262 vs methylation-only ΔA=+0.1658
Active A — healthy (non-saturated)	0.96996 (NORMAL) (5/5 active)
Active A — disease (non-saturated)	1.12417 (FLOOR BREACH) (4/5 active)
ΔA active (real progression signal)	+0.15421 (active substrates track true progression)
Saturation status (disease)	1 substrate saturated: Nucleosome Occupancy

Two combined values are shown because some substrates physically cannot resolve past their class ceiling ($A_{\text{max}} = 1/H_{\text{min}}$). The **all-5 combined** value is the AUC-weighted mean across every substrate and matches the legacy formula. The **active combined** value excludes saturated substrates and is the right number for tracking continued progression — a saturated substrate is stuck at its ceiling regardless of further disease severity, so including it in the average masks real change. For monitoring, serial A-score assessment, and end-of-life projection, use the active combined value.

DISEASE SIGNATURE COMPARISON — healthy secretory vs BRCA vs HCC vs PAAD vs PRAD vs NAFLD

Six conditions on a single chart, all in the secretory class. β values reproduce per-substrate A-scores that match Evidence Report VAL-007 per-cancer cfDNA targets: BRCA (TCGA n=90, doi:10.1038/nature11412, $\Delta A = +0.156$), HCC (n=94, $\Delta A = +0.174$), PAAD (n=95, $\Delta A = +0.169$), PRAD (n=91, $\Delta A = +0.144$), plus NAFLD (Ahrens 2013 n=192) as a non-cancer secretory-class identity failure for context. All five cancers sit in FLOOR BREACH; NAFLD sits in DETECTABLE. Substrate-saturation note: for secretory class, nucleosome occupancy saturates at $A \approx 1.018$ in every cancer in the panel — so nucl alone cannot distinguish one secretory cancer from another. Methyl, fuzz, WPS, and frag carry the per-cancer discrimination signal. This is the one-substrate-saturation physics of the canonical G-003b $H_{\min}$ values, not a framework limitation.




What happens past the ceiling is condition-specific. The A-score magnitude alone does not track severity once the ceiling is crossed — different diseases take different post-breach paths that the full five-substrate divergence pattern reveals. The subsection below walks through this class's post-breach trajectory: which substrates keep carrying signal, where the Warburg metabolic lock-in happens, where the glucose inversion point sits, and where the framework currently has high confidence versus where the open research questions remain.

POST-BREACH TRAJECTORY — WHAT HAPPENS AFTER A CROSSES 1.10 (SECRETORY CLASS)

The secretory class covers breast, prostate, pancreas, and liver — high-burden cancers where detection is a survival problem and treatment-response monitoring changes trajectories. Secretory cells contribute only ~8% of plasma cfDNA, so framework signals here require careful tissue-of-origin deconvolution. Post-breach, the progression signal is dominated by methyl and frag because the other substrates saturate at or near the ceiling early.

The zone language below — Warburg boundary, glucose inversion, point of no return — applies only to post-breach cells. These terms are not used on the pre-breach bar, and the qualitative A-value anchors (~1.15, ~1.25, ~1.40+) are boundaries between therapeutic windows, not additional diagnostic tiers. Specific A-values are filed as prospective predictions to be validated against survival-stratified cohorts.



WHICH SUBSTRATES CARRY POST-BREACH SIGNAL — SECRETORY CLASS

Secretory-class physics: nucleosome occupancy saturates at A ≈ 1.018 for every cancer in the panel; the other four substrates carry per-cancer discrimination.

Substrate	Ceiling (A)	Post-breach status
Methylation	1.27	Carries signal throughout all four zones
Fuzziness	1.26	Carries signal throughout all four zones
Windowed protection	1.17	Carries signal throughout all four zones
Fragment size	1.24	Carries signal throughout all four zones
Nucleosome occupancy	1.018	Saturated at ceiling — no further signal post-breach

CONDITION-BY-CONDITION TRAJECTORY

Healthy breast / liver / pancreas — reference, low cfDNA contribution

Known: Secretory cells contribute only ~8% of plasma cfDNA. Serum and tissue-specific cfDNA methylation deconvolution (Moss 2018 Nature Genetics) is the standard input. The framework’s A-score at class reference ≈ 0.97 holds whether deconvolution is perfect or approximate; the per-tissue contribution simply scales what fraction of the signal can be recovered.

BRCA (TCGA n=90) — A ≈ 1.156, CROSSED CEILING

Known: Breast cancer ΔA = +0.156 matches VAL-007 within measurement noise. BRCA is the framework’s most important proof-of-concept for early detection because it has the largest screening infrastructure (mammography) and the clearest unmet need: DCIS stratification. Current medicine has no validated tool to distinguish indolent from active DCIS at diagnosis.

Unknown: whether high-grade vs low-grade DCIS produces distinguishable A-scores in serum cfDNA; whether triple-negative BRCA produces a different post-breach trajectory under chemotherapy than hormone-receptor-positive disease; whether BRCA1/2 carriers show early-warning A-score elevation before imaging-detectable disease.

Test that closes the gap: G-2026-P030: Prospective cohort of 200 screen-detected DCIS patients with pre-excision serum cfDNA and 5-year invasive-cancer outcomes. Prediction: baseline A_active will stratify DCIS into active-progression and indolent strata with specificity ≥ 0.80, providing the first molecular tool for DCIS over-treatment reduction.

PAAD (TCGA n=95) — A ≈ 1.169, CROSSED CEILING

Known: Pancreatic adenocarcinoma is the framework’s hardest detection case — secretory-class ΔA of +0.169 but tissue cfDNA contribution is among the lowest in the body. Detection is intrinsically late-stage for anatomical reasons. Post-breach, PAAD shows rapid progression through the four zones, with most patients reaching palliative range within 6-12 months of diagnosis.

Unknown: whether the framework can detect PAAD at a stage earlier than CA 19-9 elevation; whether metabolic intervention (pancreatic cancer is highly glycolytic) flattens the post-breach A-score trajectory.

Test that closes the gap: G-2026-P031: Prospective cohort of 100 new-onset diabetes patients over age 50 (high-risk population for occult PAAD) with serial cfDNA every 6 months. Prediction: A_active elevation above age-adjusted baseline will detect PAAD 6-12 months before standard imaging, with sensitivity ≥ 0.60 at specificity 0.95.

NAFLD (Ahrens 2013 n=192) — $A \approx 1.055$, MARGINAL-to-DETECTABLE

Known: Non-alcoholic fatty liver disease is a non-cancer secretory-class failure that the framework detects at lower A-score than HCC. This is the intermediate position that existing methylation clocks (trained on healthy aging or cancer) cannot interpret — the architecture-floor reference makes it readable. NAFLD $A \approx 1.055$ is pre-malignant but clearly abnormal.

Unknown: whether the A-score trajectory in NAFLD-to-NASH progression predicts which patients will develop HCC over a 10-year horizon; whether metabolic intervention in NAFLD reverses the A-score toward healthy.

Test that closes the gap: G-2026-P032: Retrospective reanalysis of the LITMUS NAFLD Biomarkers Consortium cohort (n=2,000+ biopsy-confirmed NAFLD with longitudinal serum archives). Prediction: A_active trajectory over 5 years will identify NASH-to-HCC progressors with $AUC \geq 0.70$.

WHAT THE FRAMEWORK CAN AND CANNOT SAY TODAY

The framework has high confidence for secretory-class post-breach: (1) four substrates carry full post-breach signal with nucl saturating; (2) per-cancer ΔA values match VAL-007 within noise for all six validated cancers; (3) the intermediate-state detection (NAFLD, CHIP-like states) has a natural home in the framework that existing clocks cannot provide.

The framework has not yet tested the DCIS stratification, PAAD early-detection, or NAFLD-to-HCC progression predictions prospectively. Predictions G-2026-P030, P031, P032 define the specific validation plans.

This is the scientific posture the framework commits to: definitive where the evidence is definitive, open where it is open, and specific about the experiments that would close each gap. Predictions filed for this class: G-2026-P030, G-2026-P031, G-2026-P032. A reviewer who wants to falsify any claim here can do so against the named cohorts. A collaborator who wants to extend the framework has specific hypotheses to test.

THREE-COMPONENT DECOMPOSITION (C1/C2/C3) — PER SUBSTRATE

Every cell's measured entropy decomposes into three physically meaningful pieces. C1 is the universal Landauer floor: $H_MIN_GLOBAL = 0.756500$, measured in frontal cortex neurons by Lister 2013. Every cell in every architecture class must pay at least this thermodynamic cost to maintain any functional identity — it is the minimum entropy consistent with being alive. C2 is the class-specific overhead: the additional entropy the secretory class must carry above the universal floor to encode its particular architecture. For this class, C2 represents approximately 10.5% of the healthy reference entropy. C3 is the accessible gap: $\max(0, H_actual - H_min)$. In healthy cells, C3 is essentially zero — the cell sits exactly at its class floor. When C3 grows, the cell is departing from its architecture class. Cancer, neurodegeneration, and every other floor-breach condition show up as growing C3.

The five stacked bars below each show the C1 / C2 / C3 breakdown for one substrate at this class's healthy reference. Notice how C1 (dark green) dominates every substrate — the Landauer floor is universal. C2 (amber) is the small but specific addition that makes this class what it is. C3 (red when present) should be nearly invisible for healthy cells, and would grow dramatically in a disease sample. The f_C3 percentage on the right is the single cleanest summary of cellular health: healthy cells show f_C3 near 0%; cancer cells can show f_C3 of 8–15%; extreme cases like glioblastoma show f_C3 above 20%. This is not a statistical threshold calibrated from disease data. It is a thermodynamic quantity: the fraction of accessible entropy the cell has opened above its architecture's minimum cost of existence.



BEST TESTING METHOD RANKING — CLINICAL UTILITY FOR THIS CLASS

Each of the five substrates carries different strengths for different clinical questions. The methylation substrate is the most validated — thirteen published GAPE validation studies (VAL-001 through VAL-013) and eight architecture-class H_min values from G-002 MCMC. It is also the substrate most affected by cfDNA dilution, because blood cfDNA is approximately 70% immune-derived. Fragment size (DELF) is the most forgiving of dilution — short-fragment enrichment persists even when tumor-derived cfDNA is a small fraction of total plasma DNA. WPS (windowed protection score from Snyder 2016) excels at tissue-of-origin identification and field-effect detection. Nucleosome occupancy and fuzziness add chromatin-level information that methylation alone cannot reveal.

The ranking below is specifically for the secretory class. The order reflects three factors: (1) the substrate's published single-substrate AUC for diseases in this class; (2) the biology of this class's primary failure mode (Secretory Overload); and (3) the practical constraint of cfDNA contribution in the relevant specimen (plasma for most classes, CSF for terminal, tissue biopsy where cfDNA is too dilute). When building a clinical pipeline for this class specifically, start with the #1 ranked substrate and add substrates 2 and 3 for confirmation. Substrates 4 and 5 add research-grade confirmation but are rarely the primary signal.

#	Substrate	Best for	Rationale
1	Methylation	Pan-secretory cancer detection	Highest-precision H_min calibration (MCMC $\sigma = 0.0006$). Six TCGA cancers validated. Primary substrate for BRCA, PRAD, LIHC, PAAD detection.
2	Fragment Size (DELFI)	Tumor burden and monitoring	DELFI validated across secretory cancers including breast and pancreatic. Superior for monitoring treatment response over methylation.
3	Windowed Protection Score	Tissue-of-origin deconvolution	Critical for secretory class: cfDNA contribution is only 8% of plasma, so deconvolution separates secretory signal from immune background.
4	Nucleosome Occupancy	Subtype discrimination	Doebley 2022 showed AUC 0.89-0.96 for ER status in breast cancer from nucleosome occupancy. Best substrate for hormone-receptor subtyping.
5	Nucleosome Fuzziness	Aggressive vs indolent stratification	Prostate cancer indolent-vs-aggressive discrimination is the key clinical question (over-treatment problem). Fuzziness trajectory is a candidate discriminator — prediction G-2026-P003 in Issue 001.

BODY TEMPERATURE SCALING — $\alpha = 2.0$ LANDAUER CORRECTION

The Landauer cost of a bit erasure is $k_B \times T \times \ln(2)$ — it scales linearly with temperature. Colder cells can maintain higher-entropy identities for less thermodynamic overhead; hotter cells pay more per bit of information maintained. The GAPE framework captures this with a simple scaling law: $H_{\min}(T) = H_{\min}(37^\circ\text{C}) \times (T_{\text{body}} / 310.15\text{ K})^\alpha$, with $\alpha = 2.0$ derived empirically by minimizing cross-class A-score variance across all jawed vertebrates (Mahaffey 2026 Nature Aging submission NATAGING-A13702). The $\alpha = 2.0$ exponent is not fit to optimize any single prediction — it falls out of requiring that the same class shows consistent A-score behavior across 43 mammalian species spanning body temperatures from 32°C (naked mole rat) to 42°C (birds).

The table below shows what the secretory class floor $H_{\min}$ looks like at seven representative body temperatures. At 37°C (human reference, highlighted row), $H_{\min} = 0.8433$. At 42°C (birds), the floor is higher because the Landauer tax is larger. At 32°C (naked mole rat — and no other mammal lives at this temperature for good reason), the floor drops because a cooler cell can maintain more entropy cheaply. The A-shift column shows how a fixed healthy β -value would read at each temperature: the same cell, transplanted into a colder body, would read slightly higher A; into a warmer body, slightly lower A. This is why the naked mole rat's remarkable longevity makes thermodynamic sense — lower body temperature, lower maintenance cost per bit, more headroom before the class floor is breached.


T_body	Species example	H_min(T)	$\beta_{\text{healthy equivalent}}$	A-shift vs 37°C
42°C	Birds (chicken, finch)	0.8707	0.746	$A \times 0.969$
39°C	Rodents (mouse, rat)	0.8542	0.746	$A \times 0.987$
37°C	Humans (REFERENCE)	0.8433	0.746	$A \times 1.000$ (anchor)
35°C	Hibernating bats	0.8324	0.746	$A \times 1.013$
32°C	Naked mole rat	0.8163	0.746	$A \times 1.033$
25°C	Reptiles (anole lizard)	0.7793	0.746	$A \times 1.082$
15°C	Fish (salmon, cold-water)	0.7279	0.746	$A \times 1.159$

HEALTHY AGING TRAJECTORY — AGE-STRATIFIED REFERENCE A-SCORES

Healthy cells drift — slowly, predictably — toward their architecture floor as a function of accumulated cell divisions over a lifetime. For the secretory class, the drift rate is approximately 4.0% per generation — meaning each cell cycle carries a 4.0% probability of a methylation error that raises the cell's A-score by a measurable amount. Over decades, these errors accumulate. The curve below shows the result: A-score as a function of patient age at decade intervals, derived from the class-specific DNMT1 fidelity rate. At age 20, the reference healthy A-score sits well below 1.00; by age 80, it has drifted upward into the MARGINAL to DETECTABLE range. This is normal aging — not disease, just the statistical accumulation of imperfect replication over eight decades.

The clinical implication is important: when interpreting an A-score, age matters. A patient at age 30 with $A = 1.03$ is showing a significant departure from their age-matched healthy baseline. The same $A = 1.03$ at age 75 might be entirely consistent with age-expected drift. The GAPE framework

handles this through age-stratified reference A-scores — the curve below IS the reference. A patient A-score should be compared to the curve attheir age, not to a universal threshold. Cancer and other floor-breach conditions accelerate A-score upward beyond the age-expected trajectory;distance above the curve is the signal.



VERTEBRATE LIFESPAN CONTEXT

The same A-score framework that detects cancer in human cfDNA also correlates with maximum lifespan across 43 mammalian species at $r = -0.9018$ ($p = 1.6 \times 10^{-16}$, Mahaffey 2026 Nature Aging submission NATAGING-A13702). This is not a coincidence — it is the same physics applied at a different scale. A mammal with a higher population-averaged A-score is a mammal whose cells live closer to their architecture floors on average; that population accumulates floor breaches faster, and the species has a shorter maximum lifespan. The $A = 1.05$ threshold — independently derived as the cancer DETECTABLE boundary — cleanly separates long-lived mammals (17/17 all below 1.05) from short-lived mammals (11/11 all above 1.05) with complete accuracy. Same threshold. Same physics. Different timescale.

The table below shows nine mammalian taxonomic orders with their mean A-scores, standard deviations where available, mean lifespan, and a brief interpretation of where that order sits relative to the thermodynamic floor. Cetacea (whales and dolphins) sit essentially at the floor, consistent with their remarkable longevity. Primates including humans sit slightly above the floor. Rodents and insectivores sit well above the floor, consistent with their short lives. This class's reference A-score (0.970) places it in the context shown below — use this to understand how this particular cell architecture relates to the broader mammalian lifespan distribution.

Order	N	Mean A	Lifespan (yr)	Interpretation
Cetacea	5	0.997 ±0.016	99	At thermodynamic floor
Proboscidea	1	0.987 —	70	At floor
Primates	6	1.007 ±0.015	49	Slightly above floor
Artiodactyla	4	1.015 ±0.011	30	Slightly above floor
Chiroptera	4	1.041 ±0.007	35	Above floor (longevity outliers)
Carnivora	9	1.053 ±0.023	28	Above floor
Lagomorpha	2	1.114 ±0.002	8	Well above floor
Rodentia	6	1.125 ±0.018	9	Well above floor
Insectivora	1	1.157 —	2	Furthest from floor

INTERVENTION LEVERS — RANKED BY EXPECTED IMPACT

If a patient's A-score is elevated — MARGINAL, DETECTABLE, or URGENT — the clinical question is what can be done. The GAPE framework does not prescribe treatment; that remains the physician's judgment. But the framework does identify which intervention categories have the mechanistic strength to move this particular class's A-score back toward the floor. The five categories below are the standard anti-aging and cancer-adjacent interventions: senolytics (clearing senescent cells), metabolic restoration (NAD+, caloric restriction, exercise), epigenetic restoration (DNMT/TET modulators, HDAC inhibitors), reprogramming (cyclic Yamanaka factors), and checkpoint stringency (cell cycle control, immune checkpoint blockade).

The ranking below is specifically for the secretory class and its primary failure mode (Secretory Overload). Impact Level 1 (Dominant) means the intervention directly addresses the class's failure mechanism and is the first lever to try. Level 2 (Strong) means substantial mechanistic support for A-score restoration in this class. Level 3 (Moderate) means helpful but not addressing the binding constraint. Level 4 (Limited) means the intervention works for other classes but does not target this class's specific biology. Level 5 (Not applicable) means the intervention category is biologically incompatible with this class — e.g., senolytics cannot act on post-mitotic terminal cells that do not become classically senescent, and reprogramming cannot be applied to terminal cells without losing the identity that defines them.

Impact	Category	Mechanism	Rationale
2 STRONG	Senolytics	senolytics	Strong — senescent secretory cells amplify secretory load
2 STRONG	Metabolic	metabolic	Strong — secretory cells have high ATP demand; metabolic optimization improves fidelity
2 STRONG	Epigenetic Restoration	epigenetic rx	Strong — secretory methylation regulated by DNMT3A/3B
3 MODERATE	Checkpoint Stringency	checkpoint	Moderate — checkpoint modulation useful in pre-cancerous secretory lesions
4 LIMITED	Reprogramming	reprogramming	Limited — secretory differentiation is the functional state

CANCER PANEL — ΔA RANKED

The TCGA cancer validation is where the framework earns its keep. For every cancer type in this class, we compute two numbers from published TCGA methylation data: A_healthy (from matched normal tissue) and A_tumor (from cancer samples in the same study). The ΔA = A_tumor – A_healthy is the signal the framework predicts should appear — and it does. Thresholds for tier assignment: A = 1.05 DETECT, A = 1.07 URGENT, A = 1.10 FLOOR BREACH. These thresholds are not fit to the cancer data; they were derived from the healthy-vs-age statistics and happen to align with the cancer validation cleanly. The bars below are ordered by |ΔA| descending so the cancers with the largest signal appear first.

Note that ΔA ranges typical for this class matter. Terminal-class cancers (glioma, glioblastoma) show ΔA ≈ 0.22–0.27 — the largest in the entire 28-cancer TCGA dataset. Cycling epithelial cancers show ΔA ≈ 0.13–0.19 — moderate but clearly detectable. Stromal and progenitor-lineage cancers show ΔA ≈ 0.10–0.17. Each cancer below is cited to its primary TCGA publication and shows sample size. Direct links from A-score to TCGA evidence — no black boxes between physics and clinical data.

#1	Ovarian (OV) n=67 · TCGA OV 2011 Nature		ΔA=+0.2072	A=1.180 FLOOR
#2	Breast (BRCA) n=90 · TCGA BRCA 2012 Nature		ΔA=+0.2059	A=1.177 FLOOR
#3	Adrenocortical (ACC) n=80 · Zheng 2016 Cancer Cell		ΔA=+0.1922	A=1.169 FLOOR
#4	Prostate (PRAD) n=50 · TCGA PRAD 2015 Cell		ΔA=+0.1890	A=1.155 FLOOR
#5	Liver (LIHC) n=52 · TCGA LIHC 2017 Cell		ΔA=+0.1874	A=1.171 FLOOR
#6	Pancreatic (PAAD) n=150 · TCGA PAAD 2017 Cancer Cell		ΔA=+0.1746	A=1.164 FLOOR

CORE METRICS — PUBLISHED AND DERIVED

Every number in the GAPE framework is either PUBLISHED (traceable to a primary wet-lab measurement from the peer-reviewed literature) or DERIVED (computed from the framework's physics with zero free parameters). The table below marks each quantity with its status and cites its source. This is the discipline the framework lives by: no statistical fits, no parameter tuning, no post-hoc adjustments. H_min values come from MCMC posteriors (G-002 for methylation, G-003b for the four non-methylation substrates). Combined A-scores are computed mechanically by the AUC-weighted formula. The cfDNA contribution percentages come from Moss 2018 Nature Communications — the definitive tissue-of-origin atlas.

Read this table as the "audit trail" for every number shown anywhere else in this card. When the framework says A_healthy = 0.9695, you can see exactly which β value produced it, which H_min was used, and where both came from. The n_bio field is currently marked PRELIMINARY because

absolute n_bio values await the G-007 MCMC run; the class ordering has been confirmed (Spearman $\rho = 0.905$, $p = 0.002$ against commitmentlevel) but the absolute values carry that caveat transparently.

Metric	Value	Source
cfDNA contribution (plasma)	8.0%	Moss 2018 Nat Commun — PUBLISHED
H_min methylation (class)	0.843264	G-002 MCMC posterior — DERIVED
H_min methylation σ (MCMC)	± 0.000600	G-002 17-chain MCMC R-hat<1.001 — DERIVED
n_bio (class-specific)	21.5	PRELIMINARY — ordering confirmed ($\rho=0.905$, $p=0.002$); absolute pending G-007
Healthy drift rate	4.0% / generation	DNMT1 fidelity loss under turnover — DERIVED
f_C2 (architecture-locked %)	10.5%	C2 fraction at healthy reference — DERIVED
Healthy A (methylation)	0.9695 (NORMAL)	$\beta=0.746$ — DERIVED
Disease A (methylation)	1.1353 (FLOOR BREACH)	$\beta=0.621$ — DERIVED
Healthy combined A (5 subs)	0.9700 (NORMAL)	$\Sigma(\text{AUC}_i \times A_i) / \Sigma(\text{AUC}_i)$ across 5 substrates — DERIVED
Disease combined A (5 subs)	1.1026 (FLOOR BREACH)	$\Sigma(\text{AUC}_i \times A_i) / \Sigma(\text{AUC}_i)$ across 5 substrates — DERIVED
ΔA combined	+0.1326	Combined signal — DERIVED

DATED PREDICTIONS FOR THIS CLASS

A prediction is not a claim unless it can fail. Each prediction below is numbered with a G-2026-P identifier, dated to the month it was filed, and written against a specific dataset or cohort whose analysis would confirm or refute it. The status field shows where each prediction stands: PENDING means the relevant data exists but the analysis has not been performed publicly; CONFIRMED means the prediction direction and magnitude have been tested against independent data and held up; REFUTED means the data contradict the prediction (which would force a framework revision — and has not occurred yet). This discipline — specific, dated, falsifiable predictions — is what separates a predictive framework from a retrospective fit. Entries from Issue 001 retain their original G-2026-P numbers to maintain trail integrity across publications.

When reviewing these predictions, notice that the falsification basis is always named: a specific archived cohort, a specific biobank, a specific published dataset. These are not vague "more research is needed" statements. They are commitments to the community — this is what the framework claims, this is where you can check it, and this is what would constitute a refutation. The predictions below are for the secretory class specifically, filed as part of this publication or carried forward from Issue 001.

G-2026-P003

Originally filed Issue 001

PENDING

Among PSA-screened prostate cancer patients with matched methylation + WPS + DELFI profiles at diagnosis and at 24 months, the combined A-score trajectory slope will distinguish aggressive ($\geq$ Gleason 8) from indolent (Gleason ≤ 6) disease with AUC ≥ 0.85 , superior to PSA velocity or PSA-density.

Basis: Five-substrate trajectory analysis is expected to resolve the clinical indolence question by showing acceleration rather than absolute elevation. Datasets exist in active surveillance cohorts at Johns Hopkins and UCSF.

G-2026-P004

Originally filed Issue 001

PENDING

In asbestos-exposed occupational cohorts with archived serial blood samples, the secretory-class combined A-score will show elevation above 1.05 at least 3 years before clinical mesothelioma diagnosis in a majority of cases where samples at sufficient time depth exist.

Basis: Mesothelioma has a 40-year latency. Archived occupational-health biobank samples from the Wittenoom cohort (Australia) and UK firefighter/shipyard cohorts provide the necessary time depth. Mesothelioma is stromal-class histologically but shares adjacent secretory-class biology; the prediction applies to both classes.

Reference: Normal blood leukocytes (GSE40279, Hannum 2013 healthy cohort) · G-002 chain 3 of 17. R-hat 1.0007. Corrected H_min 0.8389 ± 0.0012 from initial calibration 0.795 (6.44σ tension resolved).

CELL IDENTITY & CLINICAL CONTEXT

Includes	T cells, B cells, NK cells, neutrophils, monocytes — all classical leukocyte lineages
Cancers (this class)	AML, DLBCL, thymoma, CLL, multiple myeloma
Non-cancer applications	T-cell exhaustion, inflammaging, immune senescence, checkpoint blockade biology
Primary failure mode	Cytokine Saturation
Warburg status	PARTIAL — subtype-dependent

COMMENTARY

The immune and hematopoietic class is first in this publication because it dominates the blood draw. Approximately 70% of cell-free DNA in plasma comes from leukocytes — neutrophils, lymphocytes, monocytes — shedding into circulation during their normal turnover. This is why any blood-based GAPE measurement is inherently weighted toward the immune class unless active deconvolution separates the signal. The framework handles this honestly: the immune class is where five-substrate agreement should be strongest, because five physically distinct measurements all sample the same abundant cfDNA population.

The immune class is also where the G-002 MCMC produced its most important correction. The initial calibration used a neutrophil reference cell at β = 0.760, giving H_min = 0.795. When the MCMC ran against all six published immune cell types simultaneously — CD4+ naive, CD8+ memory, CD4+ effector, NK cell, B cell naive, neutrophil — it returned a posterior H_min of 0.8389 ± 0.0012, a 6.44σ departure from the initial estimate. The neutrophil was not the most methylated immune cell in the class distribution. Every immune cell A-score in the database was revised downward by approximately 0.055 after this correction. CD4+ naive moved from A = 1.058 (DETECT tier) to A = 1.003 (NORMAL). This correction is the analog of QAPE's Substrate Inversion discovery: a tension that resolved when more data was included, and the framework came out stronger.


Immune cells are designed to be plastic. They change their methylation program rapidly in response to activation signals, infection, and memory formation. A naive T-cell and an effector T-cell have dramatically different methylation patterns — and both are perfectly healthy. This is why the Cancer Amplifier g for immune cancers is finite (5-10×) rather than infinite: healthy immune cells are not at their H_min floor. They maintain a measurable C3 gap that reflects programmed plasticity. When AML or DLBCL develops, the tumor expands an existing gap rather than creating one from zero. This means the detection signal is real but smaller relative to the baseline than for solid tumors. The five-substrate framework partially compensates: even if any single substrate's signal is small, combining four or five substrates reduces noise by roughly √n — recovering most of the detection power that single-substrate methylation loses to baseline plasticity.

The cancers in this class separate cleanly by lineage. AML (myeloid, TCGA 2013 NEJM n=200) sits at A_combined ≈ 1.10 (ΔA = +0.168 at cfDNA level, VAL-007). DLBCL (lymphoid, Chapuy 2018 n=48) is more extreme at A_combined ≈ 1.13 (ΔA = +0.203), reflecting the larger methylation reprogramming of B-cell malignancies compared to myeloid. Thymoma (T-cell origin, TCGA n=120) sits between at A ≈ 1.09, and CLL (indolent B-cell leukemia) at A ≈ 1.07 in DETECTABLE tier. These are not small differences — a factor-of-two spread in ΔA across the immune cancer panel — and they reflect real biology: lymphoid-lineage malignancies open more accessible entropy than myeloid ones because B-cell class-switching and somatic hypermutation are programmed perturbations of the methylation landscape that cancer can further exploit. The non-cancer applications deserve equal attention. Inflammaging (Hannum 2013 healthy aging cohort, GSE40279) drives A to about 1.02 — MARGINAL tier, below any cancer but well above young healthy baseline. Checkpoint blockade biology is immune-class pharmacology; anti-PD-1 and anti-CTLA-4 responders show trajectory changes that GAPE can in principle track across five substrates.

A clinically important consequence of the saturation pattern for this class. Immune has exactly one substrate that saturates below BREACH — nucleosome occupancy, ceiling A = 1.010 (nearly the lowest of any class). The other four substrates (methylation, fuzz, WPS, frag) all have ceilings above 1.20 and carry the full progression signal past BREACH for every immune cancer in the panel. Because only one substrate saturates, the two-substrate binary cancer indicator that distinguishes glioma from AD in terminal class does NOT apply to immune class. Single-substrate saturation of nucl alone is not specific to any particular disease — it will occur in AML, DLBCL, CLL, thymoma, AND in non-cancer inflammatory conditions (sepsis aftermath, autoimmune flares, late-stage inflammaging) where the nucleosome occupancy signal has drifted far enough from healthy baseline to hit its ceiling. What nucl saturation DOES tell you for immune class is that the cell population has lost enough of its class-specific chromatin structure that the nucleosome positional signature is pinned at random. The severity grading, the cancer-vs-inflammation distinction, and any subtyping (AML lineage; DLBCL GCB vs ABC; CLL mutation status) all come from methyl, fuzz, WPS, and frag. Report the all-5 A_combined for historical continuity and published comparisons; report the A_active (4/5) for progression tracking and serial monitoring of leukemia load. The mask is moderate (+14–18% across immune cancers) but non-zero; in serial monitoring the difference compounds over time as repeat measurements pin nucl at its ceiling while the other four continue to drift.

FIVE-SUBSTRATE FIDELITY GAUGE

The gauge below is the central image of Issue 002. A single horizontal bar represents the A-score axis — from A = 0.90 (below the floor) through A = 1.00 (exactly at the Immune class floor, H_min = 0.8389) up through the clinical threshold zones: MARGINAL at A = 1.01, DETECTABLE at 1.05, URGENT at 1.07, and FLOOR BREACH at 1.10. Five dots are plotted on this single axis for a healthy reference sample (Healthy WBC, markers above the bar), and five more dots for a disease reference (AML (TCGA n=200), markers below the bar). Each dot is one substrate: methylation, nucleosome occupancy, nucleosome fuzziness, WPS, or fragment size (DELFI). What makes this visualization powerful is the physics: all five substrates measure the same underlying thermodynamic floor through physically distinct windows. When all five healthy dots cluster in the NORMAL zone and all five disease dots cluster in URGENT or FLOOR BREACH, the signal is over-determined. Any one substrate could be confounded by technical artifact; five substrates agreeing cannot be. This is the "less blurry" advantage: $\sqrt{5}$ noise reduction for free.




Look at the spread in each direction. If the healthy five-dot cluster sits tightly together near A = 0.97–1.00, that indicates the reference population is at its architecture floor — precisely where the framework predicts healthy immune cells should sit. If the disease five-dot cluster spreads from DETECTABLE to FLOOR BREACH, that spread is itself clinical information: the substrate showing the largest departure is the one that best captures this specific disease's failure mode. For this class, the primary failure mode is Cytokine Saturation, and the substrate ranking shown further down the card explains which of the five substrates best reveals it.

SUBSTRATE-BY-SUBSTRATE BREAKDOWN

Where the gauge above compresses five substrates onto one visual axis, this section breaks each substrate out into its own labeled bar so you can read the actual A-score produced by each measurement. The physics is the same across all five: take the raw laboratory value (a β methylation fraction, a nucleosome occupancy probability, a WPS score, etc.), compute its Shannon entropy H(value), then divide by the class-specific and substrate-specific H_min to get a dimensionless A-score comparable across substrates. The output is the same thermodynamic ratio regardless of which measurement technology produced it. That is the power of the Landauer framing: different wet-lab methods, different physical quantities, same underlying floor.

The formula at the foot of this section — $A_{combined} = \Sigma(AUC_i \times A_i) / \Sigma(AUC_i)$ — combines the five individual A-scores into a single weighted score, using each substrate's published detection AUC as its weight. Higher-AUC substrates contribute more to the combined result. For immune cells, the healthy combined A-score below should sit at approximately 0.97 (NORMAL tier); the disease combined A-score should sit above 1.05 (DETECTABLE or higher). The difference between them — the ΔA combined — is the signal available to a blood draw that analyzes all five substrates simultaneously. A single-substrate assay sees a smaller signal at higher noise; the five-substrate combination sees the same signal at approximately $\sqrt{5}$ lower noise. Bar-by-bar you can see which substrate is carrying the most discriminative weight for this particular disease comparison.

HEALTHY REFERENCE — Healthy WBC



DISEASE REFERENCE — AML (TCGA n=200)
ACTIVE — tracks progression




Combined A — healthy (all 5)	0.96995 (NORMAL)
Combined A — disease (all 5)	1.10335 (FLOOR BREACH)
ΔA combined (5 substrates)	+0.13340 vs methylation-only ΔA=+0.1669
Active A — healthy (non-saturated)	0.96995 (NORMAL) (5/5 active)
Active A — disease (non-saturated)	1.12707 (FLOOR BREACH) (4/5 active)
ΔA active (real progression signal)	+0.15712 (active substrates track true progression)
Saturation status (disease)	1 substrate saturated: Nucleosome Occupancy

Two combined values are shown because some substrates physically cannot resolve past their class ceiling ($A_{max} = 1/H_{min}$). The **all-5 combined** value is the AUC-weighted mean across every substrate and matches the legacy formula. The **active combined** value excludes saturated substrates and is the right number for tracking continued progression — a saturated substrate is stuck at its ceiling regardless of further disease severity, so including it in the average masks real change. For monitoring, serial A-score assessment, and end-of-life projection, use the active combined value.

DISEASE SIGNATURE COMPARISON — healthy WBC vs Inflammaging vs CLL vs AML vs DLBCL vs Thymoma

Six conditions on a single chart, all in the immune class. β values reproduce per-substrate A-scores matching Evidence Report VAL-007 targets: AML (TCGA n=200, doi:10.1056/NEJMoa1301689, $\Delta A = +0.168$), DLBCL (Chapuy 2018, n=48, $\Delta A = +0.203$), Thymoma (TCGA n=120, doi:10.1016/j.ccell.2018.03.010), CLL (indolent B-cell leukemia, smaller ΔA than AML), plus Inflammaging (Hannum 2013 healthy aging cohort, GSE40279) as non-cancer immune senescence for context. All four cancers sit in FLOOR BREACH; Inflammaging sits in MARGINAL. Substrate-saturation note: for immune class, nucleosome occupancy saturates at $A \approx 1.010$ in every cancer — so nucl alone cannot distinguish one blood malignancy from another. Methyl, fuzz, WPS, and frag carry per-cancer discrimination. DLBCL shows the largest ΔA , consistent with the lymphoid class of B-cell malignancies showing more extreme methylation departure than myeloid (AML).




What happens past the ceiling is condition-specific. The A-score magnitude alone does not track severity once the ceiling is crossed — different diseases take different post-breach paths that the full five-substrate divergence pattern reveals. The subsection below walks through this class's post-breach trajectory: which substrates keep carrying signal, where the Warburg metabolic lock-in happens, where the glucose inversion point sits, and where the framework currently has high confidence versus where the open research questions remain.

POST-BREACH TRAJECTORY — WHAT HAPPENS AFTER A CROSSES 1.10 (IMMUNE CLASS)

The immune class dominates cfDNA sampling (~70% of plasma cfDNA) so post-breach trajectories here are detectable at lower tumor fractions than any other class. Immune cells are also intrinsically plastic — they change their methylation program rapidly with activation, memory formation, and senescence. This means the baseline C3 gap is larger in healthy immune cells than in other classes, and post-breach progression amplifies an existing gap rather than opening a new one. Framework implication: serial monitoring is the highest-yield use of this class's A-score.

The zone language below — Warburg boundary, glucose inversion, point of no return — applies only to post-breach cells. These terms are not used on the pre-breach bar, and the qualitative A-value anchors (~1.15, ~1.25, ~1.40+) are boundaries between therapeutic windows, not additional diagnostic tiers. Specific A-values are filed as prospective predictions to be validated against survival-stratified cohorts.



WHICH SUBSTRATES CARRY POST-BREACH SIGNAL — IMMUNE CLASS

Immune-class physics: nucleosome occupancy saturates tightly at the ceiling in every cancer in this panel; the other four substrates carry the per-cancer discrimination and progression signal.

Substrate	Ceiling (A)	Post-breach status
Methylation	1.19	Carries signal throughout all four zones
Fuzziness	1.24	Carries signal throughout all four zones
Windowed protection	1.16	Carries signal throughout all four zones
Fragment size	1.23	Carries signal throughout all four zones
Nucleosome occupancy	1.010	Saturated at ceiling — no further signal post-breach

INVERSION TERRITORY — HIV/AIDS AS PREDICTED CASE

The immune class has no documented hypomethylation inversion at the cancer level — every validated immune cancer (AML, DLBCL, CLL, Thymoma) shows upward A-score departure consistent with classical hyperentropy. However, the framework predicts a non-cancer inversion case: HIV/AIDS. Progressive CD4+ T-cell depletion with T-cell exhaustion (high PD-1, loss of DNMT1 maintenance fidelity) should produce a measurable immune-class A-score trajectory. The predicted direction is elevation (A_{immune} rising with disease burden), with possible substrate-specific inversion signatures as the T-cell compartment collapses and apoptotic fragmentation dominates cfDNA release. This has not been tested in published cohorts yet — see G-2026-P026 for the specific falsification plan against MACS/WIHS and ACTG archived cohorts.

CONDITION-BY-CONDITION TRAJECTORY

Healthy WBC — reference, plastic baseline

Known: A_{combined} ≈ 0.97 at class reference. The 6.44σ MCMC correction on this class's H_{min} value (initial 0.795 → posterior 0.8389) was the framework's most important calibration event. Healthy immune cells maintain a measurable C3 gap (5-10× lower than solid-tumor classes show) that reflects programmed plasticity, not disease.

AML (TCGA n=200) — A ≈ 1.150, CROSSED CEILING

Known: AML sits in the metabolic-window to structural-only zone post-breach. Primary validation: TCGA LAML 2013 NEJM (n=200). AML is the framework's strongest immune-class evidence because it has the longest clinical use of epigenetic therapies — DNMTi (azacitidine, decitabine) and combination regimens have 50 years of outcome data. The framework correctly predicts that DNMTi works in AML by structural (not metabolic) means: it pushes A_{methyl} higher transiently through hypomethylation before allowing the cell to re-differentiate toward a healthy methylation program.

Unknown: whether baseline A_{active} at AML diagnosis predicts response to azacitidine-venetoclax combinations (the current standard for older AML); whether MRD-negative remission on the framework's A-score corresponds to molecular MRD assessed by flow cytometry.

Test that closes the gap: G-2026-P026b: Retrospective reanalysis of archived serial cfDNA from the VIALE-A trial cohort (n=431 older AML patients on azacitidine ± venetoclax). Framework prediction: A_{active} decline slope over the first two treatment cycles will predict 12-month overall

survival with AUC ≥ 0.70, outperforming current serial marrow-blast assessment.

DLBCL (Chapuy 2018 n=48) — A ≈ 1.161, CROSSED CEILING

Known: DLBCL shows the largest combined ΔA in the immune panel (+0.203) — consistent with the lymphoid class of B-cell malignancies showing more extreme methylation departure than myeloid. This is the framework’s cleanest separation of two malignancy lineages by their thermodynamic signature: B-cell malignancies systematically over-commit methylation during germinal-center reactions.

Unknown: whether CHOP vs R-CHOP vs CAR-T response produces distinguishable A-score trajectories; whether the three Chapuy 2018 genomic subtypes (C1-C5) trace different post-breach paths that the framework could discriminate before molecular classification.

Test that closes the gap: G-2026-P026c: Prospective cohort of 100 DLBCL patients with serial cfDNA pre-cycle-1, pre-cycle-3, end-of-treatment, and 6 months post. Prediction: A_active trajectory discriminates complete metabolic response from partial/no response at end-of-treatment with AUC ≥ 0.80, outperforming Deauville score from interim PET.

HIV/AIDS (predicted) — non-cancer immune exhaustion

Known: HIV produces progressive CD4+ T-cell depletion with T-cell exhaustion (PD-1 up, DNMT1 fidelity compromised). This is not cancer but it is a genuine immune-class architectural failure that the framework predicts should register on the A-score before CD4+ count drops into AIDS-defining ranges. Acute/untreated HIV should show A_immune in the MARGINAL to WATCH range; advanced AIDS with opportunistic infections in DETECTABLE to URGENT. Effective ART should produce A_immune recovery toward healthy.

Unknown: whether A_immune trajectory correlates inversely with CD4+ count; whether ART responders show A_immune declining toward healthy faster than CD4+ rebounds (predicting viral reservoir clearance); whether long-term non-progressors (LTNPs) maintain A_immune closer to healthy than standard untreated HIV.

Test that closes the gap: G-2026-P026: Retrospective reanalysis of archived serial blood from MACS (Multicenter AIDS Cohort Study) and WIHS (Women’s Interagency HIV Study). Framework prediction: immune-class A_active trajectory under ART will correlate with CD4+ recovery with r ≥ 0.60 over 24 months; LTNPs will show mean A_immune ≤ 1.02 vs progressors at ≥ 1.05.

WHAT THE FRAMEWORK CAN AND CANNOT SAY TODAY

The framework has high confidence for immune-class post-breach: (1) four substrates (methyl, fuzz, WPS, frag) carry full post-breach signal with nucl saturating; (2) B-cell vs myeloid malignancies separate cleanly at the ΔA level (DLBCL +0.203 vs AML +0.168) — structural, not adjusted; (3) the 70% cfDNA dominance makes immune-class signals detectable at lower tumor fractions than any other class.


The framework has not yet tested the HIV/AIDS prediction in published cohorts, nor the chemotherapy-response predictions for AML and DLBCL longitudinally. Predictions G-2026-P026, P026b, P026c define the specific cohorts, endpoints, and falsification criteria.

This is the scientific posture the framework commits to: definitive where the evidence is definitive, open where it is open, and specific about the experiments that would close each gap. Predictions filed for this class: G-2026-P026, G-2026-P026b, G-2026-P026c. A reviewer who wants to falsify any claim here can do so against the named cohorts. A collaborator who wants to extend the framework has specific hypotheses to test.

THREE-COMPONENT DECOMPOSITION (C1/C2/C3) — PER SUBSTRATE

Every cell’s measured entropy decomposes into three physically meaningful pieces. C1 is the universal Landauer floor: H_MIN_GLOBAL = 0.756500, measured in frontal cortex neurons by Lister 2013. Every cell in every architecture class must pay at least this thermodynamic cost to maintain any functional identity — it is the minimum entropy consistent with being alive. C2 is the class-specific overhead: the additional entropy the immune class must carry above the universal floor to encode its particular architecture. For this class, C2 represents approximately 9.8% of the healthy reference entropy. C3 is the accessible gap: max(0, H_actual – H_min). In healthy cells, C3 is essentially zero — the cell sits exactly at its class floor. When C3 grows, the cell is departing from its architecture class. Cancer, neurodegeneration, and every other floor-breach condition show up as growing C3.

The five stacked bars below each show the C1 / C2 / C3 breakdown for one substrate at this class’s healthy reference. Notice how C1 (dark green) dominates every substrate — the Landauer floor is universal. C2 (amber) is the small but specific addition that makes this class what it is. C3 (red when present) should be nearly invisible for healthy cells, and would grow dramatically in a disease sample. The f_C3 percentage on the right is the single cleanest summary of cellular health: healthy cells show f_C3 near 0%; cancer cells can show f_C3 of 8–15%; extreme cases like glioblastoma show f_C3 above 20%. This is not a statistical threshold calibrated from disease data. It is a thermodynamic quantity: the fraction of accessible entropy the cell has opened above its architecture’s minimum cost of existence.



Windowed Protect	C1 (universal floor)	C2	C1=92.0% C2=8.0% C3=0.00%
Fragment Size (D	C1 (universal floor)	C2	C1=80.8% C2=19.2% C3=0.00%

BEST TESTING METHOD RANKING — CLINICAL UTILITY FOR THIS CLASS

Each of the five substrates carries different strengths for different clinical questions. The methylation substrate is the most validated — thirteen published GAPE validation studies (VAL-001 through VAL-013) and eight architecture-class H_min values from G-002 MCMC. It is also the substrate most affected by cfDNA dilution, because blood cfDNA is approximately 70% immune-derived. Fragment size (DELF) is the most forgiving of dilution — short-fragment enrichment persists even when tumor-derived cfDNA is a small fraction of total plasma DNA. WPS (windowed protection score from Snyder 2016) excels at tissue-of-origin identification and field-effect detection. Nucleosome occupancy and fuzziness add chromatin-level information that methylation alone cannot reveal.

The ranking below is specifically for the immune class. The order reflects three factors: (1) the substrate's published single-substrate AUC for diseases in this class; (2) the biology of this class's primary failure mode (Cytokine Saturation); and (3) the practical constraint of cfDNA contribution in the relevant specimen (plasma for most classes, CSF for terminal, tissue biopsy where cfDNA is too dilute). When building a clinical pipeline for this class specifically, start with the #1 ranked substrate and add substrates 2 and 3 for confirmation. Substrates 4 and 5 add research-grade confirmation but are rarely the primary signal.

#	Substrate	Best for	Rationale
1	Methylation	Pan-hematological detection	Methylation dominates because immune cells contribute 70% of cfDNA. The G-002 MCMC correction (6.44σ) validated this class most rigorously.
2	Fragment Size (DELF)	Blood malignancy burden	DELF fragmentomics shows AUC 0.94 across 7 cancer types including hematologic. Directly from the same cfDNA WGS that feeds WPS.
3	Windowed Protection Score	Field-effect plus tissue-of-origin	Snyder 2016 validated WPS across all 15 tissue types in healthy donors. The method was born on immune-class cfDNA.
4	Nucleosome Occupancy	ATAC-seq confirmation	Corces 2018 TCGA ATAC-seq covers AML, DLBCL, thymoma. Use when paired immune-cell ATAC data is available.
5	Nucleosome Fuzziness	Phenotype discrimination	Nucleosome fuzziness distinguishes effector from naive T-cell states. Secondary — use alongside methylation for immunophenotyping.

BODY TEMPERATURE SCALING — α = 2.0 LANDAUER CORRECTION

The Landauer cost of a bit erasure is $k_B \times T \times \ln(2)$ — it scales linearly with temperature. Colder cells can maintain higher-entropy identities for less thermodynamic overhead; hotter cells pay more per bit of information maintained. The GAPE framework captures this with a simple scaling law: $H_{min}(T) = H_{min}(37^\circ C) \times (T_{body} / 310.15\text{ K})^\alpha$, with $\alpha = 2.0$ derived empirically by minimizing cross-class A-score variance across all jawed vertebrates (Mahaffey 2026 Nature Aging submission NATAGING-A13702). The $\alpha = 2.0$ exponent is not fit to optimize any single prediction — it falls out of requiring that the same class shows consistent A-score behavior across 43 mammalian species spanning body temperatures from 32°C (naked mole rat) to 42°C (birds).

The table below shows what the immune class floor H_min looks like at seven representative body temperatures. At 37°C (human reference, highlighted row), H_min = 0.8389. At 42°C (birds), the floor is higher because the Landauer tax is larger. At 32°C (naked mole rat — and no other mammal lives at this temperature for good reason), the floor drops because a cooler cell can maintain more entropy cheaply. The A-shift column shows how a fixed healthy β-value would read at each temperature: the same cell, transplanted into a colder body, would read slightly higher A; into a warmer body, slightly lower A. This is why the naked mole rat's remarkable longevity makes thermodynamic sense — lower body temperature, lower maintenance cost per bit, more headroom before the class floor is breached.


T_body	Species example	H_min(T)	β_healthy equivalent	A-shift vs 37°C
42°C	Birds (chicken, finch)	0.8662	0.748	A × 0.969
39°C	Rodents (mouse, rat)	0.8497	0.748	A × 0.987
37°C	Humans (REFERENCE)	0.8389	0.748	A × 1.000 (anchor)
35°C	Hibernating bats	0.8281	0.748	A × 1.013
32°C	Naked mole rat	0.8121	0.748	A × 1.033
25°C	Reptiles (anole lizard)	0.7752	0.748	A × 1.082

T_body	Species example	H_min(T)	β_healthy equivalent	A-shift vs 37°C
15°C	Fish (salmon, cold-water)	0.7241	0.748	$A \times 1.159$

HEALTHY AGING TRAJECTORY — AGE-STRATIFIED REFERENCE A-SCORES

Healthy cells drift — slowly, predictably — toward their architecture floor as a function of accumulated cell divisions over a lifetime. For the immune class, the drift rate is approximately 3.5% per generation — meaning each cell cycle carries a 3.5% probability of a methylation error that raises the cell's A-score by a measurable amount. Over decades, these errors accumulate. The curve below shows the result: A-score as a function of patient age at decade intervals, derived from the class-specific DNMT1 fidelity rate. At age 20, the reference healthy A-score sits well below 1.00; by age 80, it has drifted upward into the MARGINAL to DETECTABLE range. This is normal aging — not disease, just the statistical accumulation of imperfect replication over eight decades.

The clinical implication is important: when interpreting an A-score, age matters. A patient at age 30 with A = 1.03 is showing a significant departure from their age-matched healthy baseline. The same A = 1.03 at age 75 might be entirely consistent with age-expected drift. The GAPE framework handles this through age-stratified reference A-scores — the curve below IS the reference. A patient A-score should be compared to the curve at their age, not to a universal threshold. Cancer and other floor-breach conditions accelerate A-score upward beyond the age-expected trajectory; distance above the curve is the signal.



VERTEBRATE LIFESPAN CONTEXT

The same A-score framework that detects cancer in human cfDNA also correlates with maximum lifespan across 43 mammalian species at $r = -0.9018$ ($p = 1.6 \times 10^{-16}$, Mahaffey 2026 Nature Aging submission NATAGING-A13702). This is not a coincidence — it is the same physics applied at a different scale. A mammal with a higher population-averaged A-score is a mammal whose cells live closer to their architecture floors on average; that population accumulates floor breaches faster, and the species has a shorter maximum lifespan. The A = 1.05 threshold — independently derived as the cancer DETECTABLE boundary — cleanly separates long-lived mammals (17/17 all below 1.05) from short-lived mammals (11/11 all above 1.05) with complete accuracy. Same threshold. Same physics. Different timescale.

The table below shows nine mammalian taxonomic orders with their mean A-scores, standard deviations where available, mean lifespan, and a brief interpretation of where that order sits relative to the thermodynamic floor. Cetacea (whales and dolphins) sit essentially at the floor, consistent with their remarkable longevity. Primates including humans sit slightly above the floor. Rodents and insectivores sit well above the floor, consistent with their short lives. This class's reference A-score (0.971) places it in the context shown below — use this to understand how this particular cell architecture relates to the broader mammalian lifespan distribution.

Order	N	Mean A	Lifespan (yr)	Interpretation
Cetacea	5	0.997 ±0.016	99	At thermodynamic floor
Proboscidea	1	0.987 —	70	At floor
Primates	6	1.007 ±0.015	49	Slightly above floor
Artiodactyla	4	1.015 ±0.011	30	Slightly above floor
Chiroptera	4	1.041 ±0.007	35	Above floor (longevity outliers)
Carnivora	9	1.053 ±0.023	28	Above floor
Lagomorpha	2	1.114 ±0.002	8	Well above floor
Rodentia	6	1.125 ±0.018	9	Well above floor
Insectivora	1	1.157 —	2	Furthest from floor

INTERVENTION LEVERS — RANKED BY EXPECTED IMPACT

If a patient's A-score is elevated — MARGINAL, DETECTABLE, or URGENT — the clinical question is what can be done. The GAPE framework does not prescribe treatment; that remains the physician's judgment. But the framework does identify which intervention categories have the mechanistic strength to move this particular class's A-score back toward the floor. The five categories below are the standard anti-aging and cancer-adjacent interventions: senolytics (clearing senescent cells), metabolic restoration (NAD+, caloric restriction, exercise), epigenetic restoration (DNMT/TET modulators, HDAC inhibitors), reprogramming (cyclic Yamanaka factors), and checkpoint stringency (cell cycle control, immune checkpoint blockade).

The ranking below is specifically for the immune class and its primary failure mode (Cytokine Saturation). Impact Level 1 (Dominant) means the intervention directly addresses the class's failure mechanism and is the first lever to try. Level 2 (Strong) means substantial mechanistic support for A-score restoration in this class. Level 3 (Moderate) means helpful but not addressing the binding constraint. Level 4 (Limited) means the intervention works for other classes but does not target this class's specific biology. Level 5 (Not applicable) means the intervention category is biologically incompatible with this class — e.g., senolytics cannot act on post-mitotic terminal cells that do not become classically senescent, and reprogramming cannot be applied to terminal cells without losing the identity that defines them.

Impact	Category	Mechanism	Rationale
1 DOMINANT	Epigenetic Restoration	epigenetic rx	Dominant — TET2 restoration is the primary driver of exhaustion reversal
2 STRONG	Senolytics	senolytics	Strong — senescent T cells (p16+ exhausted) directly drive immune dysfunction
2 STRONG	Metabolic	metabolic	Strong — metabolic reprogramming to OxPhos restores effector function
2 STRONG	Immune Checkpoint	checkpoint	Strong — checkpoint blockade prevents exhaustion induction
3 MODERATE	Reprogramming	reprogramming	Moderate — only if exhaustion epigenome is irreversible

CANCER PANEL — ΔA RANKED

The TCGA cancer validation is where the framework earns its keep. For every cancer type in this class, we compute two numbers from published TCGA methylation data: A_healthy (from matched normal tissue) and A_tumor (from cancer samples in the same study). The ΔA = A_tumor – A_healthy is the signal the framework predicts should appear — and it does. Thresholds for tier assignment: A = 1.05 DETECT, A = 1.07 URGENT, A = 1.10 FLOOR BREACH. These thresholds are not fit to the cancer data; they were derived from the healthy-vs-age statistics and happen to align with the cancer validation cleanly. The bars below are ordered by |ΔA| descending so the cancers with the largest signal appear first.

Note that ΔA ranges typical for this class matter. Terminal-class cancers (glioma, glioblastoma) show ΔA ≈ 0.22–0.27 — the largest in the entire 28-cancer TCGA dataset. Cycling epithelial cancers show ΔA ≈ 0.13–0.19 — moderate but clearly detectable. Stromal and progenitor-lineage cancers show ΔA ≈ 0.10–0.17. Each cancer below is cited to its primary TCGA publication and shows sample size. Direct links from A-score to TCGA evidence — no black boxes between physics and clinical data.

#1	Lymphoma (DLBCL) n=48 · Chapuy 2018 Nat Genet	ΔA=+0.1331	A=1.161 FLOOR
#2	Leukemia (AML) n=200 · TCGA AML 2013 NEJM	ΔA=+0.1303	A=1.150 FLOOR
#3	Thymoma (THYM) n=124 · TCGA THYM 2018 Cell	ΔA=+0.1148	A=1.138 FLOOR

CORE METRICS — PUBLISHED AND DERIVED

Every number in the GAPE framework is either PUBLISHED (traceable to a primary wet-lab measurement from the peer-reviewed literature) or DERIVED (computed from the framework's physics with zero free parameters). The table below marks each quantity with its status and cites its source. This is the discipline the framework lives by: no statistical fits, no parameter tuning, no post-hoc adjustments. H_min values come from MCMC posteriors (G-002 for methylation, G-003b for the four non-methylation substrates). Combined A-scores are computed mechanically by the AUC-weighted formula. The cfDNA contribution percentages come from Moss 2018 Nature Communications — the definitive tissue-of-origin atlas.

Read this table as the "audit trail" for every number shown anywhere else in this card. When the framework says A_healthy = 0.9708, you can see exactly which β value produced it, which H_min was used, and where both came from. The n_bio field is currently marked PRELIMINARY because absolute n_bio values await the G-007 MCMC run; the class ordering has been confirmed (Spearman ρ = 0.905, ρ = 0.002 against commitment level) but the absolute values carry that caveat transparently.

Metric	Value	Source
cfDNA contribution (plasma)	70.0%	Moss 2018 Nat Commun — PUBLISHED
H_min methylation (class)	0.838889	G-002 MCMC posterior — DERIVED

Metric	Value	Source
H_min methylation σ (MCMC)	± 0.001200	G-002 17-chain MCMC R-hat<1.001 — DERIVED
n_bio (class-specific)	17.5	PRELIMINARY — ordering confirmed ($p=0.905$, $p=0.002$); absolute pending G-007
Healthy drift rate	3.5% / generation	DNMT1 fidelity loss under turnover — DERIVED
f_C2 (architecture-locked %)	9.8%	C2 fraction at healthy reference — DERIVED
Healthy A (methylation)	0.9708 (NORMAL)	$\beta=0.748$ — DERIVED
Disease A (methylation)	1.1377 (FLOOR BREACH)	$\beta=0.625$ — DERIVED
Healthy combined A (5 subs)	0.9699 (NORMAL)	$\Sigma(AUC_i \times A_i) / \Sigma(AUC_i)$ across 5 substrates — DERIVED
Disease combined A (5 subs)	1.1033 (FLOOR BREACH)	$\Sigma(AUC_i \times A_i) / \Sigma(AUC_i)$ across 5 substrates — DERIVED
ΔA combined	+0.1334	Combined signal — DERIVED

DATED PREDICTIONS FOR THIS CLASS

A prediction is not a claim unless it can fail. Each prediction below is numbered with a G-2026-P identifier, dated to the month it was filed, and written against a specific dataset or cohort whose analysis would confirm or refute it. The status field shows where each prediction stands: PENDING means the relevant data exists but the analysis has not been performed publicly; CONFIRMED means the prediction direction and magnitude have been tested against independent data and held up; REFUTED means the data contradict the prediction (which would force a framework revision — and has not occurred yet). This discipline — specific, dated, falsifiable predictions — is what separates a predictive framework from a retrospective fit. Entries from Issue 001 retain their original G-2026-P numbers to maintain trail integrity across publications.

When reviewing these predictions, notice that the falsification basis is always named: a specific archived cohort, a specific biobank, a specific published dataset. These are not vague "more research is needed" statements. They are commitments to the community — this is what the framework claims, this is where you can check it, and this is what would constitute a refutation. The predictions below are for the immune class specifically, filed as part of this publication or carried forward from Issue 001.

G-2026-P010

April 2026

PENDING

In prospective cohorts of patients with known clonal hematopoiesis of indeterminate potential (CHIP) and archived serial blood samples, the immune-class combined A-score across four or more substrates will show elevation above $A = 1.03$ at least 18 months before hematologic malignancy diagnosis in a majority of cases where subsequent malignancy occurs.

Basis: CHIP is a well-characterized pre-malignant state; the UK Biobank and ARIC cohorts contain archived serial samples from CHIP-positive individuals with longitudinal follow-up. The framework predicts that the combined A-score tracks the transition to overt malignancy earlier than any single substrate.

G-2026-P011

April 2026

PENDING

In patients receiving immune checkpoint inhibitor therapy, the immune-class A-score trajectory will distinguish responders from non-responders within the first 8 weeks of treatment. Responders are predicted to show A-score stabilization or decline; non-responders will show continued A-score drift.

Basis: Checkpoint blockade reshapes T-cell methylation. The framework predicts this shows up as an A-score signature before clinical response metrics (tumor size, cytokine panels) become discriminative. Falsifiable in any prospective ICI trial with serial cfDNA collection.

CELL IDENTITY & CLINICAL CONTEXT

Includes	Common myeloid progenitor (CMP), granulocyte-monocyte progenitor (GMP), neural progenitor (NPC), erythroid progenitor, intestinal transit-amplifying cells
Cancers (this class)	MDS (all subtypes: RA, RARS, RCMD, RAEB), secondary AML from MDS, CMML, therapy-related myeloid neoplasms, pediatric B-ALL and T-ALL, mixed-lineage (MLL-rearranged) leukemias, JMML, medulloblastoma, PNET, OPC-derived gliomas, erythroleukemia (AML-M6), colorectal adenomas (pre-cancer)
Non-cancer applications	CHIP (clonal hematopoiesis), CCUS (clonal cytopenias), bone marrow failure syndromes, aplastic anemia, IBD intestinal stem/progenitor compartment dysregulation, serrated polyps
Primary failure mode	Replication Throughput Ceiling
Warburg status	EMERGING

COMMENTARY

Progenitor and transit-amplifying cells are the high-throughput workhorses of the body. They sit between adult tissue stem cells and fully committed daughter cells — partially committed to a specific lineage, but still capable of rapid proliferation to expand the committed pool before final differentiation. CMP and GMP in the bone marrow produce millions of granulocytes per hour. Intestinal transit-amplifying cells produce the daughter cells that replace the colonic epithelium every four to seven days. Neural progenitors produce neurons and glia during development and in the adult neurogenic niches. The rapid proliferation produces the characteristic 4.5% per generation drift rate — slower than cycling epithelium but faster than adult stem cells.

The Replication Throughput Ceiling is the class-specific failure mode. When a progenitor cell population divides fast enough to exceed DNMT1 maintenance fidelity — because the cell is consuming DNMT1 faster than it can be transcribed and translated — methylation errors accumulate disproportionately. This is why MDS (myelodysplastic syndrome) arises in the progenitor lineage: the CMP/GMP pool undergoes accumulated methylation errors under chronic high-throughput demand, and the resulting dysplastic progenitors produce cytopenias through failed daughter-cell maturation. The framework predicts that the early signature of MDS is detectable in peripheral blood cfDNA years before cytopenia onset, through the progenitor-class combined A-score.

This is also the class where DNMT1 inhibitor therapy (azacitidine, decitabine) acts most directly. These drugs are used clinically in MDS and AML, and the framework predicts that response to therapy — measured as A-score trajectory — is a more sensitive indicator of treatment success than the blast count or cytogenetic remission criteria currently used. The same substrate readout predicts response. The five substrates together predict it earlier and with higher confidence than any one.

The validation sequence for this class runs along a well-documented clonal-hematopoiesis spectrum. Clonal hematopoiesis of indeterminate potential (CHIP) — a pre-malignant state present in roughly 10% of healthy adults over 70 (Steensma 2015 Blood, doi:10.1182/blood-2015-03-631747) — shows A_combined ≈ 1.01, detectable but unambiguously below BREACH. Clonal cytopenias of undetermined significance (CCUS) sits at A ≈ 1.03 in MARGINAL tier (Malcovati 2017 Blood, doi:10.1182/blood-2017-04-777607). Low-risk MDS produces A_active ≈ 1.09 (BREACH) while A_combined still reads 1.05 because saturations have started. High-risk MDS sits at A_active ≈ 1.12, and secondary AML (MDS-derived) reaches A_active ≈ 1.15 — nearly matching the de novo AML signal from Card #3 Immune. The progression across these six conditions is the cleanest single-class severity gradient in the entire document and is directly relevant to the clinical decision of WHEN to escalate treatment in MDS monitoring. Bone marrow biopsy cannot be repeated monthly; a blood-based A_active can.

The progenitor class encompasses more than MDS. Pediatric acute lymphoblastic leukemia (ALL) — both B-cell precursor ALL (the most common childhood cancer) and T-ALL — arises from lymphoid progenitors and displays genome-wide hypermethylation signatures distinct from bulk AML (Nordlund 2013 Genome Biol, n=764 pediatric ALL methylomes, doi:10.1186/gb-2013-14-9-r105). The framework predicts A_active elevation in B-ALL and T-ALL through the same physics that drives the MDS signal: rapid progenitor-pool division exceeding DNMT1 fidelity. Medulloblastoma — the most common malignant pediatric brain tumor — arises from cerebellar neural progenitor cells and shows four distinct methylation subgroups (WNT, SHH, Group 3, Group 4), each with different clinical trajectories (Northcott 2017 Nature, n=1,256 methylation-profiled cases, doi:10.1038/nature22973). The DNA methylation-based CNS tumor classifier from Capper 2018 Nature (doi:10.1038/nature26000) changed pediatric neuro-oncology diagnosis in up to 12% of cases. Chronic myelomonocytic leukemia (CMML), juvenile myelomonocytic leukemia (JMML), therapy-related myeloid neoplasms, and the OPC-derived gliomas distinct from terminal-class adult glioma all belong to this class by cell-of-origin physics. Each family has lost someone to these diseases; each cancer deserves its place on the card. The unifying framework is not a coincidence — it is the common thermodynamic signature of the partially-committed, rapidly-dividing cell pool.

A clinical consequence of the saturation pattern for this class is unusually consequential. Progenitor is one of two classes in the framework where THREE substrates saturate below BREACH — nucleosome occupancy (ceiling A = 1.028), nucleosome fuzziness (ceiling A = 1.040), and WPS (ceiling A = 1.012 — the tightest class × substrate combination in the entire framework). Only two substrates (methylation, fragment size) carry the real BREACH-capable progression signal for this class. The physics explanation is clean: progenitor cells are partially uncommitted by design. They must retain chromatin flexibility to produce multiple daughter lineages, so their healthy β values for the structural substrates (nucl, fuzz, WPS)

already sit near the maximum-entropy state. There is almost no headroom. Any disease-driven drift hits the ceiling quickly. This is not a frameworklimitation — it is the correct reading of developmental flexibility as a measurement-physics constraint. The practical consequence for serial MDSmonitoring is stark. The legacy A_combined (all 5 substrates) flattens at approximately 1.07 once saturations kick in and CANNOT distinguishlow-risk MDS from high-risk MDS from secondary AML — all three compress together. A_active (methyl + frag only, the 2 non-saturated substrates)does distinguish: 1.089, 1.120, 1.150 respectively. For this class specifically, reporting A_active is not a refinement of A_combined; it is the ONLYway to track progression past the low-risk MDS threshold. Any clinical pipeline using GAPE for MDS severity grading must use A_active. The maskis +55% of the total signal — the largest in the document, dwarfing even Terminal's +45% — and unlike Terminal (where saturation is a binarycancer indicator), Progenitor's saturations are a CONTINUOUS LOSS of resolution that the active-substrate formula fully recovers.

FIVE-SUBSTRATE FIDELITY GAUGE

The gauge below is the central image of Issue 002. A single horizontal bar represents the A-score axis — from A = 0.90 (below the floor) through A = 1.00 (exactly at the Progenitor class floor, H_min = 0.8522) up through the clinical threshold zones: MARGINAL at A = 1.01, DETECTABLE at 1.05, URGENT at 1.07, and FLOOR BREACH at 1.10. Five dots are plotted on this single axis for a healthy reference sample (CMP young, markers above the bar), and five more dots for a disease reference (MDS (progenitor), markers below). Each dot is one substrate: methylation, nucleosome occupancy, nucleosome fuzziness, WPS, or fragment size (DELFI). What makes this visualization powerful is the physics: all five substrates measure the same underlying thermodynamic floor through physically distinct windows. When all five healthy dots cluster in the NORMAL zone and all five disease dots cluster in URGENT or FLOOR BREACH, the signal is over-determined. Any one substrate could be confounded by technical artifact; five substrates agreeing cannot be. This is the "less blurry" advantage: $\sqrt{5}$ noise reduction for free.




Look at the spread in each direction. If the healthy five-dot cluster sits tightly together near A = 0.97–1.00, that indicates the reference population is at its architecture floor — precisely where the framework predicts healthy progenitor cells should sit. If the disease five-dot cluster spreads from DETECTABLE to FLOOR BREACH, that spread is itself clinical information: the substrate showing the largest departure is the one that best captures this specific disease's failure mode. For this class, the primary failure mode is Replication Throughput Ceiling, and the substrate ranking shown further down the card explains which of the five substrates best reveals it.

SUBSTRATE-BY-SUBSTRATE BREAKDOWN


Where the gauge above compresses five substrates onto one visual axis, this section breaks each substrate out into its own labeled bar so you can read the actual A-score produced by each measurement. The physics is the same across all five: take the raw laboratory value (a β methylation fraction, a nucleosome occupancy probability, a WPS score, etc.), compute its Shannon entropy $H(\text{value})$, then divide by the class-specific and substrate-specific H_{min} to get a dimensionless A-score comparable across substrates. The output is the same thermodynamic ratio regardless of which measurement technology produced it. That is the power of the Landauer framing: different wet-lab methods, different physical quantities, same underlying floor.

The formula at the foot of this section — $A_{\text{combined}} = \Sigma(AUC_i \times A_i) / \Sigma(AUC_i)$ — combines the five individual A-scores into a single weighted score, using each substrate's published detection AUC as its weight. Higher-AUC substrates contribute more to the combined result. For progenitor cells, the healthy combined A-score below should sit at approximately 0.97 (NORMAL tier); the disease combined A-score should sit above 1.05 (DETECTABLE or higher). The difference between them — the ΔA combined — is the signal available to a blood draw that analyzes all five substrates simultaneously. A single-substrate assay sees a smaller signal at higher noise; the five-substrate combination sees the same signal at approximately $\sqrt{5}$ lower noise. Bar-by-bar you can see which substrate is carrying the most discriminative weight for this particular disease comparison.

HEALTHY REFERENCE — CMP young



DISEASE REFERENCE — MDS (progenitor)
ACTIVE — tracks progression




Combined A — healthy (all 5)	0.97015 (NORMAL)
Combined A — disease (all 5)	1.06697 (DETECTABLE)
ΔA combined (5 substrates)	+0.09682 vs methylation-only ΔA =+0.1498
Active A — healthy (non-saturated)	0.97015 (NORMAL) [5/5 active]
Active A — disease (non-saturated)	1.12029 (FLOOR BREACH) [2/5 active]

ΔA active (real progression signal)	+0.15014 (active substrates track true progression)
Saturation status (disease)	3 substrates saturated: Nucleosome Occupancy, Nucleosome Fuzziness, Windowed Protection Score

Two combined values are shown because some substrates physically cannot resolve past their class ceiling ($A_{\text{max}} = 1/H_{\text{min}}$). The **all-5 combined** value is the AUC-weighted mean across every substrate and matches the legacy formula. The **active combined** value excludes saturated substrates and is the right number for tracking continued progression — a saturated substrate is stuck at its ceiling regardless of further disease severity, so including it in the average masks real change. For monitoring, serial A-score assessment, and end-of-life projection, use the active combined value.

DISEASE SIGNATURE COMPARISON — healthy CMP through CHIP, CCUS, MDS stages, to sAML

Six conditions on a single chart, all in the progenitor class — the clearest demonstration in the entire document of why A_{active} matters more than A_{combined} for serial monitoring. β values reproduce per-substrate A-scores along the well-documented clonal-hematopoiesis progression pathway (Steensma 2015 Blood on CHIP; Malcovati 2017 Blood on CCUS; Jiang 2020 Cell Death Dis on MDS hypermethylation). Three substrates saturate for this class — nucleosome occupancy at $A \approx 1.028$, fuzziness at $A \approx 1.040$, and WPS at $A \approx 1.012$ (the tightest ceiling of any class $\times$ substrate combination in the framework). Once disease severity crosses low-risk MDS, all three pin at their ceilings and the legacy A_{combined} flattens at ~ 1.07 regardless of further progression to high-risk MDS or secondary AML. A_{active} (methyl + frag only) continues to track: $1.089 \rightarrow 1.120 \rightarrow 1.150$ across the three MDS stages. For progenitor-class serial monitoring, A_{active} IS the progression signal. A_{combined} is confirmatory only.




What happens past the ceiling is condition-specific. The A-score magnitude alone does not track severity once the ceiling is crossed — different diseases take different post-breach paths that the full five-substrate divergence pattern reveals. The subsection below walks through this class's post-breach trajectory: which substrates keep carrying signal, where the Warburg metabolic lock-in happens, where the glucose inversion point sits, and where the framework currently has high confidence versus where the open research questions remain.

POST-BREACH TRAJECTORY — WHAT HAPPENS AFTER A CROSSES 1.10 (PROGENITOR CLASS)

The progenitor class is the framework’s clearest demonstration of why A_active matters more than A_combined for serial monitoring. Three substrates saturate early in this class — nucleosome occupancy at 1.028, fuzziness at 1.040, and WPS at 1.012 (tied with stem_adult WPS for the tightest ceiling in the framework). Once disease crosses low-risk MDS, all three pin at their ceilings and A_combined flattens regardless of further progression. A_active (methyl + frag only) continues to track: 1.089 → 1.120 → 1.150 across the three MDS stages. For this class, A_active IS the progression signal.

The zone language below — Warburg boundary, glucose inversion, point of no return — applies only to post-breach cells. These terms are not used on the pre-breach bar, and the qualitative A-value anchors (~1.15, ~1.25, ~1.40+) are boundaries between therapeutic windows, not additional diagnostic tiers. Specific A-values are filed as prospective predictions to be validated against survival-stratified cohorts.



WHICH SUBSTRATES CARRY POST-BREACH SIGNAL — PROGENITOR CLASS

Progenitor-class physics: three substrates saturate (nucl, fuzz, WPS). Only methylation and fragmentomics carry post-breach progression signal. A_active (2/5) is mandatory for serial monitoring; A_combined is confirmatory only and must not be used to track progression past low-risk MDS.

Substrate	Ceiling (A)	Post-breach status
Methylation	1.15	Carries signal throughout all four zones
Fragment size	1.19	Carries signal throughout all four zones
Nucleosome occupancy	1.028	Saturated at ceiling — no further signal post-breach
Fuzziness	1.040	Saturated at ceiling — no further signal post-breach
Windowed protection	1.012	Saturated at ceiling (TIGHTEST-tied in framework) — no further signal post-breach

CONDITION-BY-CONDITION TRAJECTORY

Healthy CMP (common myeloid progenitor) — reference

Known: Common myeloid progenitor is the reference state. Progenitor-class H_min = 0.8522 calibrated from Roadmap E034 (erythroid progenitor) plus CD34+CD38- HSC-enriched references. Baseline A ≈ 0.97 across all five substrates in healthy adults.

CHIP (clonal hematopoiesis) — A ≈ 1.01-1.03, MARGINAL tier

Known: Clonal hematopoiesis of indeterminate potential (Steensma 2015 Blood) is defined by recurrent somatic mutations (DNMT3A, TET2, ASXL1) at variant allele frequency ≥ 2% without cytopenia. The framework reads CHIP as modest A-score elevation — detectable above healthy, clearly pre-malignant. CHIP affects ~10% of people over 70 but only a small fraction progress to MDS/AML.

Unknown: whether the A-score trajectory within CHIP predicts which patients progress to MDS vs remain stable; whether specific CHIP mutation categories trace distinguishable paths.

Test that closes the gap: G-2026-P038: UK Biobank CHIP cohort with serial methylation and 10-year MDS/AML outcomes. Prediction: A_active slope ≥ 0.01/year within CHIP identifies imminent progression with AUC ≥ 0.75, outperforming VAF trajectory alone.

Low-risk MDS — A_combined ≈ 1.05, A_active ≈ 1.09

Known: Low-risk MDS is the first saturation event: WPS pins at its ceiling. This is the clinically important transition where A_combined starts under-reporting progression. A patient with A_combined ≈ 1.05 but A_active ≈ 1.09 is clearly in DETECTABLE tier on the active metric — the combined score is no longer the right instrument.

Unknown: whether the A_active trajectory over 6-12 months predicts which low-risk MDS patients progress to high-risk MDS before current IPSS-R stratification flags them.

Test that closes the gap: G-2026-P039: Prospective cohort of 200 low-risk MDS patients with serial cfDNA every 3 months for 24 months. Prediction: A_active slope identifies progression to high-risk MDS 3-6 months before IPSS-R change, with sensitivity ≥ 0.70 at specificity 0.85.

High-risk MDS / Secondary AML — A_combined flatlined, A_active ≈ 1.12-1.15

Known: High-risk MDS shows all three saturating substrates pinned; A_combined cannot distinguish high-risk MDS from secondary AML. But A_active continues to climb: 1.120 at high-risk MDS → 1.150 at secondary AML. This is the framework’s most concrete demonstration that serial A_active monitoring provides signal that standard clinical staging loses.

Unknown: whether A_active at secondary AML diagnosis predicts response to azacitidine-venetoclax or allogeneic HSCT more accurately than current risk scores (ELN 2022).

Test that closes the gap: G-2026-P040: Retrospective reanalysis of archived secondary AML serial cfDNA cohorts. Prediction: baseline A_active at diagnosis stratifies 2-year overall survival with AUC ≥ 0.70, outperforming ELN 2022 risk classification alone.

WHAT THE FRAMEWORK CAN AND CANNOT SAY TODAY

The framework has high confidence for progenitor-class post-breach: (1) only methyl and frag carry post-breach signal — three substrates saturate before BREACH; (2) A_active vs A_combined divergence in MDS is the framework’s cleanest clinical teaching example; (3) the class’s CHIP-to-MDS-to-AML progression pathway is biologically well-characterized and the framework tracks each stage coherently.


The framework has not yet tested the CHIP-to-MDS progression prediction, the low-risk to high-risk MDS early detection, or secondary AML prognostic stratification prospectively. G-2026-P038, P039, P040 define the validation plans.

This is the scientific posture the framework commits to: definitive where the evidence is definitive, open where it is open, and specific about the experiments that would close each gap. Predictions filed for this class: G-2026-P038, G-2026-P039, G-2026-P040. A reviewer who wants to falsify any claim here can do so against the named cohorts. A collaborator who wants to extend the framework has specific hypotheses to test.

THREE-COMPONENT DECOMPOSITION (C1/C2/C3) — PER SUBSTRATE

Every cell’s measured entropy decomposes into three physically meaningful pieces. C1 is the universal Landauer floor: H_MIN_GLOBAL = 0.756500, measured in frontal cortex neurons by Lister 2013. Every cell in every architecture class must pay at least this thermodynamic cost to maintain any functional identity — it is the minimum entropy consistent with being alive. C2 is the class-specific overhead: the additional entropy the progenitor class must carry above the universal floor to encode its particular architecture. For this class, C2 represents approximately 11.8% of the healthy reference entropy. C3 is the accessible gap: max(0, H_actual – H_min). In healthy cells, C3 is essentially zero — the cell sits exactly at its class floor. When C3 grows, the cell is departing from its architecture class. Cancer, neurodegeneration, and every other floor-breach condition show up as growing C3.

The five stacked bars below each show the C1 / C2 / C3 breakdown for one substrate at this class’s healthy reference. Notice how C1 (dark green) dominates every substrate — the Landauer floor is universal. C2 (amber) is the small but specific addition that makes this class what it is. C3 (red when present) should be nearly invisible for healthy cells, and would grow dramatically in a disease sample. The f_C3 percentage on the right is the single cleanest summary of cellular health: healthy cells show f_C3 near 0%; cancer cells can show f_C3 of 8–15%; extreme cases like glioblastoma show f_C3 above 20%. This is not a statistical threshold calibrated from disease data. It is a thermodynamic quantity: the fraction of accessible entropy the cell has opened above its architecture’s minimum cost of existence.



BEST TESTING METHOD RANKING — CLINICAL UTILITY FOR THIS CLASS

Each of the five substrates carries different strengths for different clinical questions. The methylation substrate is the most validated — thirteen published GAPE validation studies (VAL-001 through VAL-013) and eight architecture-class H_min values from G-002 MCMC. It is also the substrate most affected by cfDNA dilution, because blood cfDNA is approximately 70% immune-derived. Fragment size (DELF) is the most forgiving of dilution — short-fragment enrichment persists even when tumor-derived cfDNA is a small fraction of total plasma DNA. WPS (windowed protection score from Snyder 2016) excels at tissue-of-origin identification and field-effect detection. Nucleosome occupancy and fuzziness add chromatin-level information that methylation alone cannot reveal.

The ranking below is specifically for the progenitor class. The order reflects three factors: (1) the substrate’s published single-substrate AUC for diseases in this class; (2) the biology of this class’s primary failure mode (Replication Throughput Ceiling); and (3) the practical constraint of cfDNA contribution in the relevant specimen (plasma for most classes, CSF for terminal, tissue biopsy where cfDNA is too dilute). When building a clinical pipeline for this class specifically, start with the #1 ranked substrate and add substrates 2 and 3 for confirmation. Substrates 4 and 5 add

research-grade confirmation but are rarely the primary signal.

#	Substrate	Best for	Rationale
1	Methylation	Progenitor-lineage malignancy detection	Bone marrow methylation is the gold standard for MDS diagnosis. Peripheral blood cfDNA methylation signal is class-specific when deconvolution is applied.
2	Windowed Protection Score	Mismatch repair state monitoring	WPS at MMR-gene promoters indicates MLH1/MSH2 methylation status — directly relevant to Lynch syndrome and MSI-H tumors.
3	Fragment Size (DELFI)	Rapid turnover signal	Progenitor-origin cfDNA fragmentation pattern is distinct — high turnover produces characteristic short-fragment enrichment.
4	Nucleosome Occupancy	Commitment-state assessment	ATAC-seq discriminates early progenitor from late transit-amplifying states. Research-grade.
5	Nucleosome Fuzziness	Division-rate correlation	Fuzziness correlates with proliferation index. Useful for grading in MDS and other progenitor-lineage disorders.

BODY TEMPERATURE SCALING — $\alpha = 2.0$ LANDAUER CORRECTION

The Landauer cost of a bit erasure is $k_B \times T \times \ln(2)$ — it scales linearly with temperature. Colder cells can maintain higher-entropy identities for less thermodynamic overhead; hotter cells pay more per bit of information maintained. The GAPE framework captures this with a simple scaling law: $H_{\min}(T) = H_{\min}(37^\circ\text{C}) \times (T_{\text{body}} / 310.15\text{ K})^\alpha$, with $\alpha = 2.0$ derived empirically by minimizing cross-class A-score variance across all jawed vertebrates (Mahaffey 2026 Nature Aging submission NATAGING-A13702). The $\alpha = 2.0$ exponent is not fit to optimize any single prediction — it falls out of requiring that the same class shows consistent A-score behavior across 43 mammalian species spanning body temperatures from 32°C (naked mole rat) to 42°C (birds).


The table below shows what the progenitor class floor $H_{\min}$ looks like at seven representative body temperatures. At 37°C (human reference, highlighted row), $H_{\min} = 0.8522$. At 42°C (birds), the floor is higher because the Landauer tax is larger. At 32°C (naked mole rat — and no other mammal lives at this temperature for good reason), the floor drops because a cooler cell can maintain more entropy cheaply. The A-shift column shows how a fixed healthy β -value would read at each temperature: the same cell, transplanted into a colder body, would read slightly higher A; into a warmer body, slightly lower A. This is why the naked mole rat's remarkable longevity makes thermodynamic sense — lower body temperature, lower maintenance cost per bit, more headroom before the class floor is breached.

T_body	Species example	H_min(T)	$\beta_{\text{healthy equivalent}}$	A-shift vs 37°C
42°C	Birds (chicken, finch)	0.8799	0.740	$A \times 0.969$
39°C	Rodents (mouse, rat)	0.8632	0.740	$A \times 0.987$
37°C	Humans (REFERENCE)	0.8522	0.740	$A \times 1.000$ (anchor)
35°C	Hibernating bats	0.8413	0.740	$A \times 1.013$
32°C	Naked mole rat	0.8250	0.740	$A \times 1.033$
25°C	Reptiles (anole lizard)	0.7875	0.740	$A \times 1.082$
15°C	Fish (salmon, cold-water)	0.7356	0.740	$A \times 1.159$

HEALTHY AGING TRAJECTORY — AGE-STRATIFIED REFERENCE A-SCORES

Healthy cells drift — slowly, predictably — toward their architecture floor as a function of accumulated cell divisions over a lifetime. For the progenitor class, the drift rate is approximately 4.5% per generation — meaning each cell cycle carries a 4.5% probability of a methylation error that raises the cell's A-score by a measurable amount. Over decades, these errors accumulate. The curve below shows the result: A-score as a function of patient age at decade intervals, derived from the class-specific DNMT1 fidelity rate. At age 20, the reference healthy A-score sits well below 1.00; by age 80, it has drifted upward into the MARGINAL to DETECTABLE range. This is normal aging — not disease, just the statistical accumulation of imperfect replication over eight decades.

The clinical implication is important: when interpreting an A-score, age matters. A patient at age 30 with $A = 1.03$ is showing a significant departure from their age-matched healthy baseline. The same $A = 1.03$ at age 75 might be entirely consistent with age-expected drift. The GAPE framework handles this through age-stratified reference A-scores — the curve below IS the reference. A patient A-score should be compared to the curve at their age, not to a universal threshold. Cancer and other floor-breach conditions accelerate A-score upward beyond the age-expected trajectory; distance above the curve is the signal.



VERTEBRATE LIFESPAN CONTEXT

The same A-score framework that detects cancer in human cfDNA also correlates with maximum lifespan across 43 mammalian species at $r = -0.9018$ ($p = 1.6 \times 10^{-16}$, Mahaffey 2026 Nature Aging submission NATAGING-A13702). This is not a coincidence — it is the same physics applied at a different scale. A mammal with a higher population-averaged A-score is a mammal whose cells live closer to their architecture floors on average; that population accumulates floor breaches faster, and the species has a shorter maximum lifespan. The $A = 1.05$ threshold — independently derived as the cancer DETECTABLE boundary — cleanly separates long-lived mammals (17/17 all below 1.05) from short-lived mammals (11/11 all above 1.05) with complete accuracy. Same threshold. Same physics. Different timescale.

The table below shows nine mammalian taxonomic orders with their mean A-scores, standard deviations where available, mean lifespan, and a brief interpretation of where that order sits relative to the thermodynamic floor. Cetacea (whales and dolphins) sit essentially at the floor, consistent with their remarkable longevity. Primates including humans sit slightly above the floor. Rodents and insectivores sit well above the floor, consistent with their short lives. This class's reference A-score (0.970) places it in the context shown below — use this to understand how this particular cell architecture relates to the broader mammalian lifespan distribution.

Order	N	Mean A	Lifespan (yr)	Interpretation
Cetacea	5	0.997 ±0.016	99	At thermodynamic floor
Proboscidea	1	0.987 —	70	At floor
Primates	6	1.007 ±0.015	49	Slightly above floor
Artiodactyla	4	1.015 ±0.011	30	Slightly above floor
Chiroptera	4	1.041 ±0.007	35	Above floor (longevity outliers)
Carnivora	9	1.053 ±0.023	28	Above floor
Lagomorpha	2	1.114 ±0.002	8	Well above floor
Rodentia	6	1.125 ±0.018	9	Well above floor
Insectivora	1	1.157 —	2	Furthest from floor

INTERVENTION LEVERS — RANKED BY EXPECTED IMPACT

If a patient's A-score is elevated — MARGINAL, DETECTABLE, or URGENT — the clinical question is what can be done. The GAPE framework does not prescribe treatment; that remains the physician's judgment. But the framework does identify which intervention categories have the mechanistic strength to move this particular class's A-score back toward the floor. The five categories below are the standard anti-aging and cancer-adjacent interventions: senolytics (clearing senescent cells), metabolic restoration (NAD+, caloric restriction, exercise), epigenetic restoration (DNMT/TET modulators, HDAC inhibitors), reprogramming (cyclic Yamanaka factors), and checkpoint stringency (cell cycle control, immune checkpoint blockade).

The ranking below is specifically for the progenitor class and its primary failure mode (Replication Throughput Ceiling). Impact Level 1 (Dominant) means the intervention directly addresses the class's failure mechanism and is the first lever to try. Level 2 (Strong) means substantial mechanistic support for A-score restoration in this class. Level 3 (Moderate) means helpful but not addressing the binding constraint. Level 4 (Limited) means the intervention works for other classes but does not target this class's specific biology. Level 5 (Not applicable) means the intervention category is biologically incompatible with this class — e.g., senolytics cannot act on post-mitotic terminal cells that do not become classically senescent, and reprogramming cannot be applied to terminal cells without losing the identity that defines them.

Impact	Category	Mechanism	Rationale
1 DOMINANT	G2/M Checkpoint	checkpoint	Dominant — G2/M checkpoint activation is the primary lever

Impact	Category	Mechanism	Rationale
2 STRONG	MMR Restoration	epigenetic rx	Strong — MMR restoration directly addresses the Replication Throughput Ceiling
3 MODERATE	Senolytics	senolytics	Moderate — senescent progenitors contribute but are minor fraction
3 MODERATE	Metabolic	metabolic	Moderate — metabolic lever moves index but does not address the ceiling
4 LIMITED	Reprogramming	reprogramming	Limited — partial commitment; full reprogramming disrupts lineage

CORE METRICS — PUBLISHED AND DERIVED

Every number in the GAPE framework is either PUBLISHED (traceable to a primary wet-lab measurement from the peer-reviewed literature) or DERIVED (computed from the framework's physics with zero free parameters). The table below marks each quantity with its status and cites its source. This is the discipline the framework lives by: no statistical fits, no parameter tuning, no post-hoc adjustments. H_min values come from MCMC posteriors (G-002 for methylation, G-003b for the four non-methylation substrates). Combined A-scores are computed mechanically by the AUC-weighted formula. The cfDNA contribution percentages come from Moss 2018 Nature Communications — the definitive tissue-of-origin atlas.

Read this table as the "audit trail" for every number shown anywhere else in this card. When the framework says A_healthy = 0.9701, you can see exactly which β value produced it, which H_min was used, and where both came from. The n_bio field is currently marked PRELIMINARY because absolute n_bio values await the G-007 MCMC run; the class ordering has been confirmed (Spearman ρ = 0.905, p = 0.002 against commitment level) but the absolute values carry that caveat transparently.

Metric	Value	Source
cfDNA contribution (plasma)	2.0%	Moss 2018 Nat Commun — PUBLISHED
H_min methylation (class)	0.852216	G-002 MCMC posterior — DERIVED
H_min methylation σ (MCMC)	± 0.001700	G-002 17-chain MCMC R-hat<1.001 — DERIVED
n_bio (class-specific)	20.0	PRELIMINARY — ordering confirmed ($p=0.905$, $p=0.002$); absolute pending G-007
Healthy drift rate	4.5% / generation	DNMT1 fidelity loss under turnover — DERIVED
f_C2 (architecture-locked %)	11.8%	C2 fraction at healthy reference — DERIVED
Healthy A (methylation)	0.9701 (NORMAL)	$\beta=0.740$ — DERIVED
Disease A (methylation)	1.1199 (FLOOR BREACH)	$\beta=0.625$ — DERIVED
Healthy combined A (5 subs)	0.9701 (NORMAL)	$\Sigma(AUC_i \times A_i) / \Sigma(AUC_i)$ across 5 substrates — DERIVED
Disease combined A (5 subs)	1.0670 (DETECTABLE)	$\Sigma(AUC_i \times A_i) / \Sigma(AUC_i)$ across 5 substrates — DERIVED
ΔA combined	+0.0968	Combined signal — DERIVED

DATED PREDICTIONS FOR THIS CLASS

A prediction is not a claim unless it can fail. Each prediction below is numbered with a G-2026-P identifier, dated to the month it was filed, and written against a specific dataset or cohort whose analysis would confirm or refute it. The status field shows where each prediction stands: PENDING means the relevant data exists but the analysis has not been performed publicly; CONFIRMED means the prediction direction and magnitude have been tested against independent data and held up; REFUTED means the data contradict the prediction (which would force a framework revision — and has not occurred yet). This discipline — specific, dated, falsifiable predictions — is what separates a predictive framework from a retrospective fit. Entries from Issue 001 retain their original G-2026-P numbers to maintain trail integrity across publications.

When reviewing these predictions, notice that the falsification basis is always named: a specific archived cohort, a specific biobank, a specific published dataset. These are not vague "more research is needed" statements. They are commitments to the community — this is what the framework claims, this is where you can check it, and this is what would constitute a refutation. The predictions below are for the progenitor class specifically, filed as part of this publication or carried forward from Issue 001.

G-2026-P014

April 2026

PENDING

In MDS patients receiving hypomethylating agent (HMA) therapy, the progenitor-class combined A-score trajectory at 3 months post-initiation will identify responders (by 6-month IWG criteria) with $AUC \geq 0.85$, outperforming baseline-to-3-month cytogenetic change alone.

Basis: HMA therapy is the first-line treatment for intermediate/high-risk MDS. Currently response is assessed at 6 months — a long wait with significant toxicity. The framework predicts earlier classification via combined A-score. Falsifiable in MDS cohorts at major cancer centers.

CELL IDENTITY & CLINICAL CONTEXT

Includes	Colon, rectum, stomach, lung (adeno + squamous), bladder, cervix, skin (melanoma), kidney (clear cell + papillary), head & neck, ovary, endometrium, thyroid
Cancers (this class)	Colorectal adenocarcinoma (COAD), Rectal adenocarcinoma (READ), Lung adenocarcinoma (LUAD), Lung squamous cell carcinoma (LUSC), Bladder urothelial carcinoma (BLCA), Ovarian serous carcinoma (OV), Stomach adenocarcinoma (STAD), Cervical squamous cell carcinoma (CESC), Head and neck squamous cell carcinoma (HNSC), Kidney clear cell RCC (KIRC), Kidney papillary RCC (KIRP), Skin cutaneous melanoma (SKCM), Thyroid carcinoma (THCA), Endometrial carcinoma (UCEC) — 14/28 TCGA types
Non-cancer applications	Inflammatory bowel disease (IBD), chronic colitis, Barrett's esophagus, chronic viral hepatitis, GERD-associated dysplasia, smoker's lung (pre-LUSC field cancerization), HPV carrier status, field cancerization, epithelial aging
Primary failure mode	Replication Ceiling
Warburg status	WALL CROSSED

COMMENTARY

The cycling epithelial class is the largest and most clinically important in the GAPE dataset. Fourteen of the 28 confirmed TCGA cancer types fall here — colon, rectum, lung (adenocarcinoma and squamous), bladder, cervix, head and neck, ovary, stomach, kidney (clear cell and papillary), skin melanoma, endometrium, and thyroid. These are the cancers that colonoscopy, Pap smears, low-dose CT, and mammography were invented to catch. Together they account for the majority of cancer deaths worldwide. Every family affected by one of these diseases is a family this card has to speak to honestly.

Cycling epithelial cells are defined by function: continuous division to replace the epithelial lining of organs exposed to external or internal environments. Colonic mucosa turns over completely in four to seven days. Cervical epithelium renews in weeks. Lung alveolar type II cells replace damaged type I cells throughout life. This continuous division is precisely why cycling epithelia are cancer-prone: every division requires DNMT1 to copy methylation patterns across 19.6 million CpG sites, and accumulated errors over decades drive the global hypomethylation signature that GAPE detects across all five substrates simultaneously. The 5.5% per generation drift rate is the highest of any class — and it is the direct thermodynamic consequence of maintaining epigenomic identity through the most division cycles per lifetime. H_min = 0.856055 (G-002 chain 1 of 17, R-hat 1.0003, posterior ±0.000800) reflects a floor tight enough to define a class but loose enough to accommodate the tissue-to-tissue variation across all 14 cancer types that share it.

This class is where MESA — the four-substrate cfDNA test (Li 2024 Genome Med) — was developed and validated. MESA achieves AUC = 0.931 on colorectal cancer using methylation, fragment size, nucleosome occupancy, and WPS from a single blood tube. This is the strongest independent validation of the multi-substrate framework in the entire document. The measured inter-substrate correlation r = 0.54 between MESA's four signals (VAL-014) matches exactly what GAPE predicts when four physical windows measure the same underlying thermodynamic floor departure: the signals share approximately 85% of their information, and the combined-to-single improvement ratio d_combined/d_single = 1.15× lands precisely in the GAPE-predicted range. MESA's ML feature ranking downweights nucleosome occupancy relative to the other substrates without explaining why. GAPE explains why: for the cycling class, nucl saturates at ceiling A = 1.020, below the BREACH threshold at A = 1.10. The feature is downweighted because it physically cannot resolve severity past a certain point. This is not a machine-learning discovery; it is a thermodynamic measurement constraint that the ML model reflects empirically without naming. Adding the fifth substrate (DELFI fragment size from Cristiano 2019 Nature, AUC 0.940 across 7 cancer types) brings the theoretical detection ceiling to AUC = 1.000, with the gap from MESA's 0.931 attributable to bulk blood dilution — addressable by tissue-specific cfDNA deconvolution.

The two deadliest cancers in this class deserve explicit naming. Colorectal cancer (COAD + READ, TCGA Research Network 2012 Nature, n=169 matched-normal + tumor methylation) is the second-leading cancer killer in the United States and claims over 900,000 lives per year globally. VAL-007 reports cfDNA ΔA = +0.158 for colorectal in stool-based sampling — fully in FLOOR BREACH tier. Lung cancer, split between adenocarcinoma (LUAD, TCGA 2014 Nature, n=230) and squamous cell carcinoma (LUSC, TCGA 2012 Nature, n=178), is the leading cancer killer worldwide with over 1.8 million annual deaths. VAL-007 plasma cfDNA ΔA for LUAD = +0.144; LUSC is comparable. These two cancers alone account for roughly 30% of all cancer mortality. The five-substrate framework for cycling class is not a theoretical advance over MESA — it is the measurable path toward pre-clinical detection of the two cancers most responsible for cancer-related death. Every delayed diagnosis in these cancers is a calibration problem the framework could have caught earlier.


The remaining twelve cycling cancers each deserve their place on this card. Bladder urothelial carcinoma (BLCA, TCGA 2014 Nature, n=131) shows the largest ΔA in the panel at urine cfDNA ΔA = +0.185 — urine is the natural specimen for bladder cancer and concentrates the signal. Ovarian serous carcinoma (OV, TCGA 2011 Nature, n=67) is the most lethal gynecologic cancer due to late detection; pre-clinical cfDNA monitoring is the framework's proposed intervention. Stomach adenocarcinoma (STAD, TCGA 2014 Nature, n=295) shows plasma cfDNA ΔA = +0.153. Cervical squamous (CESC, TCGA 2017 Nature, n=228) and Head and Neck squamous (HNSC, TCGA 2015 Nature, n=504) share HPV-driven methylation changes; saliva cfDNA works for HNSC (ΔA = +0.146). Kidney clear cell (KIRC, TCGA 2013 Nature, n=318) and papillary (KIRP, TCGA 2016 NEJM,

n=161) are two biologically distinct malignancies that share the kidney. Skin cutaneous melanoma (SKCM, TCGA 2015 Cell, n=477) ranks among the deadliest skin cancers and its methylation signature reflects ultraviolet mutagenesis. Thyroid carcinoma (THCA, TCGA 2014 Cell, n=51) stands out as an outlier in the class — slow-growing, often curable, but still methylation-detectable. Endometrial carcinoma (UCEC, TCGA 2013 Nature, n=118) completes the panel, and each of these cancers has a family behind it. Every TCGA citation is a cohort of patients who consented to share their tissue, and every TCGA dataset is a gift to the patients who will come after. Clickable DOIs for all thirteen TCGA primary sources listed above are on the Data Sources page.

A clinically important consequence of the saturation pattern for cycling class. Cycling has exactly one substrate that saturates below BREACH — nucleosome occupancy, ceiling A = 1.020 (Doebley 2022 Nat Commun). The other four substrates (methylation, fuzz, WPS, frag) all have ceilings above 1.22 and carry the full progression signal past BREACH for every cycling cancer in the panel. Because only one substrate saturates, the single-substrate saturation of nucl alone is NOT cancer-specific — it will occur in cycling cancers AND in non-cancer inflammatory conditions (IBD flares, chronic colitis, Barrett's esophagus with dysplasia, HPV-infected cervical tissue). What nucl saturation DOES tell you for cycling class is that the cell population has lost enough of its epithelial identity that the nucleosome positional signature is pinned at random. Severity grading, cancer-vs-benign distinction, tissue-of-origin discrimination (colon vs lung vs bladder), and any subtyping (COAD MSI status, LUAD vs LUSC, HPV+ vs HPV- HNSC) all come from methyl, fuzz, WPS, and frag. Report the all-5 A_combined for historical continuity and MESA-comparable reporting; report the A_active (4/5) for progression tracking and serial monitoring. The mask is moderate (+16.3%) but non-zero; in serial monitoring across multiple blood draws the difference compounds as nucl stays pinned at its ceiling while the other four continue to drift with disease progression or treatment response.

FIVE-SUBSTRATE FIDELITY GAUGE

The gauge below is the central image of Issue 002. A single horizontal bar represents the A-score axis — from A = 0.90 (below the floor) through A = 1.00 (exactly at the Cycling class floor, H_min = 0.8561) up through the clinical threshold zones: MARGINAL at A = 1.01, DETECTABLE at 1.05, URGENT at 1.07, and FLOOR BREACH at 1.10. Five dots are plotted on this single axis for a healthy reference sample (Normal colon, markers above the bar), and five more dots for a disease reference (COAD (TCGA n=97), markers below). Each dot is one substrate: methylation, nucleosome occupancy, nucleosome fuzziness, WPS, or fragment size (DELFI). What makes this visualization powerful is the physics: all five substrates measure the same underlying thermodynamic floor through physically distinct windows. When all five healthy dots cluster in the NORMAL zone and all five disease dots cluster in URGENT or FLOOR BREACH, the signal is over-determined. Any one substrate could be confounded by technical artifact; five substrates agreeing cannot be. This is the "less blurry" advantage: $\sqrt{5}$ noise reduction for free.




Look at the spread in each direction. If the healthy five-dot cluster sits tightly together near A = 0.97–1.00, that indicates the reference population is at its architecture floor — precisely where the framework predicts healthy cycling cells should sit. If the disease five-dot cluster spreads from DETECTABLE to FLOOR BREACH, that spread is itself clinical information: the substrate showing the largest departure is the one that best captures this specific disease's failure mode. For this class, the primary failure mode is Replication Ceiling, and the substrate ranking shown further down the card explains which of the five substrates best reveals it.

SUBSTRATE-BY-SUBSTRATE BREAKDOWN


Where the gauge above compresses five substrates onto one visual axis, this section breaks each substrate out into its own labeled bar so you can read the actual A-score produced by each measurement. The physics is the same across all five: take the raw laboratory value (a β methylation fraction, a nucleosome occupancy probability, a WPS score, etc.), compute its Shannon entropy $H(\text{value})$, then divide by the class-specific and substrate-specific H_{min} to get a dimensionless A-score comparable across substrates. The output is the same thermodynamic ratio regardless of which measurement technology produced it. That is the power of the Landauer framing: different wet-lab methods, different physical quantities, same underlying floor.

The formula at the foot of this section — $A_{\text{combined}} = \Sigma(AUC_i \times A_i) / \Sigma(AUC_i)$ — combines the five individual A-scores into a single weighted score, using each substrate's published detection AUC as its weight. Higher-AUC substrates contribute more to the combined result. For cycling cells, the healthy combined A-score below should sit at approximately 0.97 (NORMAL tier); the disease combined A-score should sit above 1.05 (DETECTABLE or higher). The difference between them — the ΔA combined — is the signal available to a blood draw that analyzes all five substrates simultaneously. A single-substrate assay sees a smaller signal at higher noise; the five-substrate combination sees the same signal at approximately $\sqrt{5}$ lower noise. Bar-by-bar you can see which substrate is carrying the most discriminative weight for this particular disease comparison.

HEALTHY REFERENCE — Normal colon



DISEASE REFERENCE — COAD (TCGA n=97)
ACTIVE — tracks progression




Combined A — healthy (all 5)	0.97017 (NORMAL)
Combined A — disease (all 5)	1.10937 (FLOOR BREACH)
ΔA combined (5 substrates)	+0.13920 vs methylation-only ΔA =+0.1859
Active A — healthy (non-saturated)	0.97017 (NORMAL) [5/5 active]
Active A — disease (non-saturated)	1.13203 (FLOOR BREACH) [4/5 active]

ΔA active (real progression signal)	+0.16187 (active substrates track true progression)
Saturation status (disease)	1 substrate saturated: Nucleosome Occupancy

Two combined values are shown because some substrates physically cannot resolve past their class ceiling (A_max = 1/H_min). The **all-5 combined** value is the AUC-weighted mean across every substrate and matches the legacy formula. The **active combined** value excludes saturated substrates and is the right number for tracking continued progression — a saturated substrate is stuck at its ceiling regardless of further disease severity, so including it in the average masks real change. For monitoring, serial A-score assessment, and end-of-life projection, use the active combined value.

DISEASE SIGNATURE COMPARISON — healthy colon through Barrett's, HNSC, LUAD, COAD, BLCA

Six conditions on a single chart, all in the cycling-epithelial class. β values reproduce per-substrate A-scores matching Evidence Report VAL-007 per-cancer cfDNA targets: COAD (TCGA n=97, doi:10.1038/nature11252, cfDNA ΔA = +0.158), LUAD (TCGA n=82, doi:10.1038/nature13385, cfDNA ΔA = +0.144), HNSC (TCGA n=504, doi:10.1038/nature14129, saliva cfDNA ΔA = +0.146), BLCA (TCGA n=131, doi:10.1038/nature12965, urine cfDNA ΔA = +0.185), plus Barrett's esophagus as a non-cancer field-cancerization state for pre-malignant context. All five cancers sit in URGENT to FLOOR BREACH; Barrett's sits in MARGINAL. Substrate-saturation note: for cycling class, nucleosome occupancy saturates at A ≈ 1.020 in every cancer in the panel — so nucl alone cannot distinguish one cycling cancer from another. Methyl, fuzz, WPS, and frag carry per-cancer discrimination. BLCA shows the largest ΔA (urine cfDNA, high tumor fraction); HNSC/LUAD cluster together at moderate departure; COAD sits between — consistent with the specimen-type physics of each cancer's optimal cfDNA compartment.




What happens past the ceiling is condition-specific. The A-score magnitude alone does not track severity once the ceiling is crossed — different diseases take different post-breach paths that the full five-substrate divergence pattern reveals. The subsection below walks through this class's post-breach trajectory: which substrates keep carrying signal, where the Warburg metabolic lock-in happens, where the glucose inversion point sits, and where the framework currently has high confidence versus where the open research questions remain.

POST-BREACH TRAJECTORY — WHAT HAPPENS AFTER A CROSSES 1.10 (CYCLING CLASS)

The cycling-epithelial class is the largest and most clinically important — 14 of the 28 validated TCGA cancer types fall here, accounting for the majority of cancer deaths worldwide. Every family affected by colon, lung, cervical, bladder, head-and-neck, or endometrial cancer is a family this card has to speak to honestly. Post-breach trajectory matters most here because serial monitoring under treatment is already the standard of care for these cancers.

The zone language below — Warburg boundary, glucose inversion, point of no return — applies only to post-breach cells. These terms are not used on the pre-breach bar, and the qualitative A-value anchors (~1.15, ~1.25, ~1.40+) are boundaries between therapeutic windows, not additional diagnostic tiers. Specific A-values are filed as prospective predictions to be validated against survival-stratified cohorts.



WHICH SUBSTRATES CARRY POST-BREACH SIGNAL — CYCLING CLASS

Cycling-class physics: nucleosome occupancy saturates at $A \approx 1.020$ for every cancer in the panel; the other four substrates carry the per-cancer discrimination and post-breach progression signal.

Substrate	Ceiling (A)	Post-breach status
Methylation	1.30	Carries signal throughout all four zones
Fuzziness	1.28	Carries signal throughout all four zones
Windowed protection	1.18	Carries signal throughout all four zones
Fragment size	1.25	Carries signal throughout all four zones
Nucleosome occupancy	1.020	Saturated at ceiling — no further signal post-breach

CONDITION-BY-CONDITION TRAJECTORY

Healthy colonic mucosa — reference, high cfDNA turnover

Known: Cycling epithelia turn over fast — colonic mucosa in 4-7 days, cervical epithelium in weeks. This is the highest division-rate class in the body, which is precisely why it is cancer-prone: every division requires DNMT1 to copy methylation across 19.6 million CpG sites, and accumulated errors over decades drive the global hypomethylation signature that GAPE detects across all five substrates simultaneously.

COAD (TCGA n=97) — $A \approx 1.158$, CROSSED CEILING

Known: Colorectal adenocarcinoma is the most validated cycling-class cancer in cfDNA literature. $\Delta A = +0.158$ matches VAL-007 exactly. COAD shows the classical Warburg progression archetype: monotonic rise in A_{active} under disease progression, response to FOLFOX or FOLFIRI produces trajectory bend toward healthy, relapse produces trajectory re-elevation. This is the cleanest example of the framework tracking clinical reality.

Unknown: whether the MSI-high vs MSS split (consequential for immunotherapy response) produces distinguishable A-score paths; whether the A_{active} trajectory under FOLFOX predicts response-duration for KRAS-mutant vs KRAS-wild-type cases.

Test that closes the gap: G-2026-P027: Retrospective reanalysis of GALAXY CRC MRD cohort ($n=1,000+$ stage II-III CRC patients with serial cfDNA post-resection). Prediction: A_{active} trajectory at 6 weeks post-resection will predict 2-year recurrence-free survival with $AUC \geq 0.80$, matching or exceeding current ctDNA-positive/negative classification.

LUAD / LUSC (lung) — $A \approx 1.144$, CROSSED CEILING

Known: Lung adenocarcinoma and squamous cell carcinoma show moderate cycling-class ΔA (+0.144) consistent with high cfDNA contribution from pulmonary epithelial turnover. The DELFI fragmentomics signature is strongest for lung cancer — Cristiano 2019 Nature demonstrated this with AUC 0.94 across 7 cancer types including lung.

Unknown: whether EGFR-mutant, ALK-fusion, and KRAS-mutant LUAD subtypes show distinguishable A-score paths under tyrosine kinase inhibitor therapy; whether pemetrexed response can be predicted from baseline A_{active} .

Test that closes the gap: G-2026-P028: Prospective cohort of 200 EGFR-mutant LUAD patients on osimertinib with serial cfDNA every 3 months. Prediction: A_active trajectory slope predicts resistance emergence 3-6 months before standard imaging progression, with sensitivity ≥ 0.70 at specificity 0.85.

BLCA (TCGA n=131) — $A \approx 1.185$, LARGEST CYCLING-CLASS ΔA

Known: Bladder urothelial carcinoma shows the largest cycling-class ΔA in the framework (+0.185) — partly structural (urothelial biology) and partly specimen (urine cfDNA has higher tumor fraction than blood for this cancer). BLCA sits in the structural-only post-breach zone at diagnosis for most patients. Past the Warburg boundary, metabolic intervention alone is unlikely to restore cellular fidelity.

Unknown: whether urine cfDNA A-score trajectories can predict BCG response in non-muscle-invasive BLCA; whether the framework can distinguish progressors from persistent but stable Ta/T1 disease.

Test that closes the gap: G-2026-P029: Prospective cohort of 150 non-muscle-invasive BLCA patients receiving BCG with urine cfDNA at baseline, 3, 6, 12 months. Prediction: urine A_active trajectory during induction BCG will discriminate responders from non-responders with AUC ≥ 0.75 , outperforming standard cytology.

WHAT THE FRAMEWORK CAN AND CANNOT SAY TODAY

The framework has high confidence for cycling-class post-breach: (1) four substrates carry full post-breach signal with nucl saturating; (2) the 14-cancer TCGA validation is the largest coherent structural prediction in the framework (96.4% direction-of-BREACH accuracy); (3) this class has the longest ctDNA literature to compare against, with framework predictions matching published ΔA values within measurement noise.

The framework has not yet tested A-score trajectories longitudinally against published treatment-response cohorts. Predictions G-2026-P027, P028, P029 define the specific validation plans for colorectal MRD monitoring, EGFR-mutant LUAD resistance prediction, and BLCA BCG response.

This is the scientific posture the framework commits to: definitive where the evidence is definitive, open where it is open, and specific about the experiments that would close each gap. Predictions filed for this class: G-2026-P027, G-2026-P028, G-2026-P029. A reviewer who wants to falsify any claim here can do so against the named cohorts. A collaborator who wants to extend the framework has specific hypotheses to test.

THREE-COMPONENT DECOMPOSITION (C1/C2/C3) — PER SUBSTRATE

Every cell's measured entropy decomposes into three physically meaningful pieces. C1 is the universal Landauer floor: $H_MIN_GLOBAL = 0.756500$, measured in frontal cortex neurons by Lister 2013. Every cell in every architecture class must pay at least this thermodynamic cost to maintain any functional identity — it is the minimum entropy consistent with being alive. C2 is the class-specific overhead: the additional entropy the cycling class must carry above the universal floor to encode its particular architecture. For this class, C2 represents approximately 12.1% of the healthy reference entropy. C3 is the accessible gap: $\max(0, H_actual - H_min)$. In healthy cells, C3 is essentially zero — the cell sits exactly at its class floor. When C3 grows, the cell is departing from its architecture class. Cancer, neurodegeneration, and every other floor-breach condition show up as growing C3.

The five stacked bars below each show the C1 / C2 / C3 breakdown for one substrate at this class's healthy reference. Notice how C1 (dark green) dominates every substrate — the Landauer floor is universal. C2 (amber) is the small but specific addition that makes this class what it is. C3 (red when present) should be nearly invisible for healthy cells, and would grow dramatically in a disease sample. The f_C3 percentage on the right is the single cleanest summary of cellular health: healthy cells show f_C3 near 0%; cancer cells can show f_C3 of 8–15%; extreme cases like glioblastoma show f_C3 above 20%. This is not a statistical threshold calibrated from disease data. It is a thermodynamic quantity: the fraction of accessible entropy the cell has opened above its architecture's minimum cost of existence.



BEST TESTING METHOD RANKING — CLINICAL UTILITY FOR THIS CLASS

Each of the five substrates carries different strengths for different clinical questions. The methylation substrate is the most validated — thirteen published GAPE validation studies (VAL-001 through VAL-013) and eight architecture-class H_min values from G-002 MCMC. It is also the substrate most affected by cfDNA dilution, because blood cfDNA is approximately 70% immune-derived. Fragment size (DELF) is the most forgiving of dilution — short-fragment enrichment persists even when tumor-derived cfDNA is a small fraction of total plasma DNA. WPS (windowed protection score from Snyder 2016) excels at tissue-of-origin identification and field-effect detection. Nucleosome occupancy and fuzziness add chromatin-level information that methylation alone cannot reveal.

The ranking below is specifically for the cycling class. The order reflects three factors: (1) the substrate's published single-substrate AUC for diseases in this class; (2) the biology of this class's primary failure mode (Replication Ceiling); and (3) the practical constraint of cfDNA contribution in the relevant specimen (plasma for most classes, CSF for terminal, tissue biopsy where cfDNA is too dilute). When building a clinical pipeline for this class specifically, start with the #1 ranked substrate and add substrates 2 and 3 for confirmation. Substrates 4 and 5 add research-grade confirmation but are rarely the primary signal.

#	Substrate	Best for	Rationale
1	Methylation	Pan-cancer screening (14/28 TCGA types)	Methylation is the primary substrate for this class. Fourteen TCGA cycling cancers calibrated against it. The framework's most battle-tested data stream.
2	Windowed Protection Score	Field-effect detection	Chromatin accessibility at cycling-class identity promoters shows depletion before clinical cancer. Snyder 2016 identified this signal at 15/15 tissue types.
3	Fragment Size (DELFI)	Early detection and monitoring	DELFI pipeline clinically validated for CRC, lung, and other cycling-class cancers. Short/long fragment ratio tracks tumor burden with high sensitivity.
4	Nucleosome Occupancy	Tissue-of-origin confirmation	ATAC-seq occupancy maps complement methylation when single-modality is ambiguous. Most useful when paired with methylation for discordance analysis.
5	Nucleosome Fuzziness	Aggressiveness grading	Nucleosome fuzziness correlates with proliferation rate. Secondary — combine with clinical staging (Dukes, TNM) for prognosis.

BODY TEMPERATURE SCALING — $\alpha = 2.0$ LANDAUER CORRECTION

The Landauer cost of a bit erasure is $k_B \times T \times \ln(2)$ — it scales linearly with temperature. Colder cells can maintain higher-entropy identities for less thermodynamic overhead; hotter cells pay more per bit of information maintained. The GAPE framework captures this with a simple scaling law: $H_{\min}(T) = H_{\min}(37^\circ\text{C}) \times (T_{\text{body}} / 310.15\text{ K})^\alpha$, with $\alpha = 2.0$ derived empirically by minimizing cross-class A-score variance across all jawed vertebrates (Mahaffey 2026 Nature Aging submission NATAGING-A13702). The $\alpha = 2.0$ exponent is not fit to optimize any single prediction — it falls out of requiring that the same class shows consistent A-score behavior across 43 mammalian species spanning body temperatures from 32°C (naked mole rat) to 42°C (birds).

The table below shows what the cycling class floor $H_{\min}$ looks like at seven representative body temperatures. At 37°C (human reference, highlighted row), $H_{\min} = 0.8561$. At 42°C (birds), the floor is higher because the Landauer tax is larger. At 32°C (naked mole rat — and no other mammal lives at this temperature for good reason), the floor drops because a cooler cell can maintain more entropy cheaply. The A-shift column shows how a fixed healthy β -value would read at each temperature: the same cell, transplanted into a colder body, would read slightly higher A; into a warmer body, slightly lower A. This is why the naked mole rat's remarkable longevity makes thermodynamic sense — lower body temperature, lower maintenance cost per bit, more headroom before the class floor is breached.


T_body	Species example	H_min(T)	β_{healthy} equivalent	A-shift vs 37°C
42°C	Birds (chicken, finch)	0.8839	0.738	$A \times 0.969$
39°C	Rodents (mouse, rat)	0.8671	0.738	$A \times 0.987$
37°C	Humans (REFERENCE)	0.8561	0.738	$A \times 1.000$ (anchor)
35°C	Hibernating bats	0.8451	0.738	$A \times 1.013$
32°C	Naked mole rat	0.8287	0.738	$A \times 1.033$
25°C	Reptiles (anole lizard)	0.7911	0.738	$A \times 1.082$
15°C	Fish (salmon, cold-water)	0.7389	0.738	$A \times 1.159$

HEALTHY AGING TRAJECTORY — AGE-STRATIFIED REFERENCE A-SCORES

Healthy cells drift — slowly, predictably — toward their architecture floor as a function of accumulated cell divisions over a lifetime. For the cycling class, the drift rate is approximately 5.5% per generation — meaning each cell cycle carries a 5.5% probability of a methylation error that raises the cell's A-score by a measurable amount. Over decades, these errors accumulate. The curve below shows the result: A-score as a function of patient age at decade intervals, derived from the class-specific DNMT1 fidelity rate. At age 20, the reference healthy A-score sits well below 1.00; by age 80, it has drifted upward into the MARGINAL to DETECTABLE range. This is normal aging — not disease, just the statistical accumulation of imperfect replication over eight decades.

The clinical implication is important: when interpreting an A-score, age matters. A patient at age 30 with $A = 1.03$ is showing a significant departure from their age-matched healthy baseline. The same $A = 1.03$ at age 75 might be entirely consistent with age-expected drift. The GAPE framework

handles this through age-stratified reference A-scores — the curve below IS the reference. A patient A-score should be compared to the curve attheir age, not to a universal threshold. Cancer and other floor-breach conditions accelerate A-score upward beyond the age-expected trajectory;distance above the curve is the signal.



VERTEBRATE LIFESPAN CONTEXT

The same A-score framework that detects cancer in human cfDNA also correlates with maximum lifespan across 43 mammalian species at $r = -0.9018$ ($p = 1.6 \times 10^{-16}$, Mahaffey 2026 Nature Aging submission NATAGING-A13702). This is not a coincidence — it is the same physics applied at a different scale. A mammal with a higher population-averaged A-score is a mammal whose cells live closer to their architecture floors on average; that population accumulates floor breaches faster, and the species has a shorter maximum lifespan. The $A = 1.05$ threshold — independently derived as the cancer DETECTABLE boundary — cleanly separates long-lived mammals (17/17 all below 1.05) from short-lived mammals (11/11 all above 1.05) with complete accuracy. Same threshold. Same physics. Different timescale.

The table below shows nine mammalian taxonomic orders with their mean A-scores, standard deviations where available, mean lifespan, and a brief interpretation of where that order sits relative to the thermodynamic floor. Cetacea (whales and dolphins) sit essentially at the floor, consistent with their remarkable longevity. Primates including humans sit slightly above the floor. Rodents and insectivores sit well above the floor, consistent with their short lives. This class's reference A-score (0.969) places it in the context shown below — use this to understand how this particular cell architecture relates to the broader mammalian lifespan distribution.

Order	N	Mean A	Lifespan (yr)	Interpretation
Cetacea	5	0.997 ±0.016	99	At thermodynamic floor
Proboscidea	1	0.987 —	70	At floor
Primates	6	1.007 ±0.015	49	Slightly above floor
Artiodactyla	4	1.015 ±0.011	30	Slightly above floor
Chiroptera	4	1.041 ±0.007	35	Above floor (longevity outliers)
Carnivora	9	1.053 ±0.023	28	Above floor
Lagomorpha	2	1.114 ±0.002	8	Well above floor
Rodentia	6	1.125 ±0.018	9	Well above floor
Insectivora	1	1.157 —	2	Furthest from floor

INTERVENTION LEVERS — RANKED BY EXPECTED IMPACT

If a patient's A-score is elevated — MARGINAL, DETECTABLE, or URGENT — the clinical question is what can be done. The GAPE framework does not prescribe treatment; that remains the physician's judgment. But the framework does identify which intervention categories have the mechanistic strength to move this particular class's A-score back toward the floor. The five categories below are the standard anti-aging and cancer-adjacent interventions: senolytics (clearing senescent cells), metabolic restoration (NAD+, caloric restriction, exercise), epigenetic restoration (DNMT/TET modulators, HDAC inhibitors), reprogramming (cyclic Yamanaka factors), and checkpoint stringency (cell cycle control, immune checkpoint blockade).

The ranking below is specifically for the cycling class and its primary failure mode (Replication Ceiling). Impact Level 1 (Dominant) means the intervention directly addresses the class's failure mechanism and is the first lever to try. Level 2 (Strong) means substantial mechanistic support for A-score restoration in this class. Level 3 (Moderate) means helpful but not addressing the binding constraint. Level 4 (Limited) means the intervention works for other classes but does not target this class's specific biology. Level 5 (Not applicable) means the intervention category is biologically incompatible with this class — e.g., senolytics cannot act on post-mitotic terminal cells that do not become classically senescent, and reprogramming cannot be applied to terminal cells without losing the identity that defines them.

Impact	Category	Mechanism	Rationale
1 DOMINANT	Checkpoint Stringency	checkpoint	Dominant — G1/S and G2/M checkpoint activation is the primary lever
2 STRONG	Senolytics	senolytics	Strong — senescent cells in the crypt drive stem cell niche dysfunction
2 STRONG	Epigenetic Restoration	epigenetic rx	Strong — MMR/checkpoint restoration directly addresses the inversion
3 MODERATE	Metabolic	metabolic	Moderate — useful but Replication Throughput Ceiling is the binding constraint
4 LIMITED	Reprogramming	reprogramming	Limited — cycling architecture is the functional requirement

CANCER PANEL — ΔA RANKED

The TCGA cancer validation is where the framework earns its keep. For every cancer type in this class, we compute two numbers from published TCGA methylation data: A_healthy (from matched normal tissue) and A_tumor (from cancer samples in the same study). The $\Delta A = A_{\text{tumor}} - A_{\text{healthy}}$ is the signal the framework predicts should appear — and it does. Thresholds for tier assignment: A = 1.05 DETECT, A = 1.07 URGENT, A = 1.10 FLOOR BREACH. These thresholds are not fit to the cancer data; they were derived from the healthy-vs-age statistics and happen to align with the cancer validation cleanly. The bars below are ordered by $|\Delta A|$ descending so the cancers with the largest signal appear first.

Note that ΔA ranges typical for this class matter. Terminal-class cancers (glioma, glioblastoma) show $\Delta A \approx 0.22\text{--}0.27$ — the largest in the entire 28-cancer TCGA dataset. Cycling epithelial cancers show $\Delta A \approx 0.13\text{--}0.19$ — moderate but clearly detectable. Stromal and progenitor-lineage cancers show $\Delta A \approx 0.10\text{--}0.17$. Each cancer below is cited to its primary TCGA publication and shows sample size. Direct links from A-score to TCGA evidence — no black boxes between physics and clinical data.



#11	Thyroid (THCA) n=51 · TCGA THCA 2014 Cell	$\Delta A = +0.1398$ A=1.091 URGENT
#12	Kidney Clear Cell (KIRC) n=318 · TCGA KIRC 2013 Nature	$\Delta A = +0.1319$ A=1.123 FLOOR
#13	Kidney Papillary (KIRP) n=161 · TCGA KIRP 2016 NEJM	$\Delta A = +0.1279$ A=1.121 FLOOR

CORE METRICS — PUBLISHED AND DERIVED

Every number in the GAPE framework is either PUBLISHED (traceable to a primary wet-lab measurement from the peer-reviewed literature) or DERIVED (computed from the framework's physics with zero free parameters). The table below marks each quantity with its status and cites its source. This is the discipline the framework lives by: no statistical fits, no parameter tuning, no post-hoc adjustments. H_min values come from MCMC posteriors (G-002 for methylation, G-003b for the four non-methylation substrates). Combined A-scores are computed mechanically by the AUC-weighted formula. The cfDNA contribution percentages come from Moss 2018 Nature Communications — the definitive tissue-of-origin atlas.

Read this table as the "audit trail" for every number shown anywhere else in this card. When the framework says A_healthy = 0.9693, you can see exactly which β value produced it, which H_min was used, and where both came from. The n_bio field is currently marked PRELIMINARY because absolute n_bio values await the G-007 MCMC run; the class ordering has been confirmed (Spearman $\rho = 0.905$, $p = 0.002$ against commitment level) but the absolute values carry that caveat transparently.

Metric	Value	Source
cfDNA contribution (plasma)	12.0%	Moss 2018 Nat Commun — PUBLISHED
H_min methylation (class)	0.856055	G-002 MCMC posterior — DERIVED
H_min methylation σ (MCMC)	± 0.000800	G-002 17-chain MCMC R-hat<1.001 — DERIVED
n_bio (class-specific)	19.5	PRELIMINARY — ordering confirmed ($\rho=0.905$, $p=0.002$); absolute pending G-007
Healthy drift rate	5.5% / generation	DNMT1 fidelity loss under turnover — DERIVED
f_C2 (architecture-locked %)	12.1%	C2 fraction at healthy reference — DERIVED
Healthy A (methylation)	0.9693 (NORMAL)	$\beta=0.738$ — DERIVED
Disease A (methylation)	1.1552 (FLOOR BREACH)	$\beta=0.562$ — DERIVED
Healthy combined A (5 subs)	0.9702 (NORMAL)	$\Sigma(AUC_i \times A_i) / \Sigma(AUC_i)$ across 5 substrates — DERIVED
Disease combined A (5 subs)	1.1094 (FLOOR BREACH)	$\Sigma(AUC_i \times A_i) / \Sigma(AUC_i)$ across 5 substrates — DERIVED
ΔA combined	+0.1392	Combined signal — DERIVED

DATED PREDICTIONS FOR THIS CLASS

A prediction is not a claim unless it can fail. Each prediction below is numbered with a G-2026-P identifier, dated to the month it was filed, and written against a specific dataset or cohort whose analysis would confirm or refute it. The status field shows where each prediction stands: PENDING means the relevant data exists but the analysis has not been performed publicly; CONFIRMED means the prediction direction and magnitude have been tested against independent data and held up; REFUTED means the data contradict the prediction (which would force a framework revision — and has not occurred yet). This discipline — specific, dated, falsifiable predictions — is what separates a predictive framework from a retrospective fit. Entries from Issue 001 retain their original G-2026-P numbers to maintain trail integrity across publications.

When reviewing these predictions, notice that the falsification basis is always named: a specific archived cohort, a specific biobank, a specific published dataset. These are not vague "more research is needed" statements. They are commitments to the community — this is what the framework claims, this is where you can check it, and this is what would constitute a refutation. The predictions below are for the cycling class specifically, filed as part of this publication or carried forward from Issue 001.

G-2026-P001

April 2026

PENDING

In a prospective screening cohort of average-risk adults aged 45-75 with archived serial blood samples and colonoscopy outcomes, the cycling-class combined A-score across 4+ substrates will identify advanced adenoma and early colorectal cancer with $AUC \geq 0.92$, outperforming multi-target stool DNA (Cologuard, reported AUC 0.83).

Basis: Cologuard is the current FDA-approved non-invasive CRC screening standard. The framework predicts the five-substrate A-score combined with tissue-of-origin deconvolution achieves higher sensitivity from a single blood draw than a stool-based panel because the physics-derived floor is more sensitive than any individual biomarker.

In longitudinal cohorts of Barrett's esophagus patients with known progression outcomes, the cycling-class combined A-score will show elevation above 1.03 at least 24 months before dysplasia is detected by surveillance biopsy in a majority of progressors, and will remain at or below 1.01 in non-progressors over the same window.

Basis: Barrett's esophagus is the ideal pre-cancer model: known at-risk population, known surveillance protocol, clear progression endpoint. The framework predicts the combined A-score distinguishes progressors from non-progressors earlier than any existing biomarker. Datasets exist in GI surveillance cohorts at major academic centers.

CELL IDENTITY & CLINICAL CONTEXT

Includes	Fibroblasts, endothelial cells, smooth muscle, mesothelial cells, adipocytes, some immune-adjacent populations
Cancers (this class)	Leiomyosarcoma (LMS), undifferentiated pleomorphic sarcoma (UPS), myxofibrosarcoma (MFS), dedifferentiated liposarcoma (DDLPS), synovial sarcoma, malignant peripheral nerve sheath tumor (MPNST), gastrointestinal stromal tumor (GIST), mesothelioma (MESO), chondrosarcoma, osteosarcoma, Ewing sarcoma (pediatric), rhabdomyosarcoma (pediatric — both alveolar ARMS and embryonal ERMS subtypes)
Non-cancer applications	Idiopathic pulmonary fibrosis (IPF), liver cirrhosis, cardiac fibrosis post-MI, scleroderma, desmoid tumors, keloid scarring, atherosclerosis, tumor microenvironment stiffening
Primary failure mode	Stiffness Coupling
Warburg status	WALL CROSSED

COMMENTARY

Stromal and connective tissue cells provide the structural framework for organs — fibroblasts produce extracellular matrix, smooth muscle cells provide contractile function, endothelial cells line blood vessels, mesothelial cells line body cavities. Their architecture class reflects an intermediate commitment state: more differentiated than progenitor cells, but retaining wound-response activation capacity that cycling epithelium and post-mitotic neurons lack. H_min = 0.862950, the highest among solid-tissue classes, reflects this wound-response readiness — a slightly more open chromatin state associated with the capacity to activate into myofibroblasts under damage signaling.

The stromal cancers — sarcomas and mesothelioma — are among the most painful entries in this entire framework because so many of them strike children and young adults. Ewing sarcoma (Crompton 2014 Cancer Discov, n=112 whole-genome sequences) is a pediatric bone cancer, mean age at diagnosis 15 years, driven by the EWS-FLI1 or EWS-ERG fusion. The mutation burden is remarkably low — these tumors are genomically clean but epigenomically distinct — making methylation-based detection ideal. Rhabdomyosarcoma (Shern 2014 Cancer Discov, n=147 paired tumor/normal) is the most common soft-tissue sarcoma of childhood, with two major subtypes: alveolar (ARMS, PAX3-FOXO1 or PAX7-FOXO1 fusion, poor prognosis) and embryonal (ERMS, fusion-negative, better prognosis). The 5-year survival for metastatic disease is 30%. Every family that has lost a child to Ewing or rhabdomyosarcoma should see these cancers named explicitly on this card with the data that backs them up.

Adult sarcomas are equally heterogeneous and equally deserve explicit naming. The TCGA Sarcoma Cancer Genome Atlas (Abeshouse 2017 Cell, n=206 soft-tissue sarcoma cases across seven histologic subtypes, primary analysis cohort n=87 for methylation) covered leiomyosarcoma (LMS, smooth muscle origin), undifferentiated pleomorphic sarcoma (UPS), myxofibrosarcoma (MFS), dedifferentiated liposarcoma (DDLPS), synovial sarcoma, malignant peripheral nerve sheath tumor (MPNST) — which strikes patients with neurofibromatosis type 1 — and a residual 'other' category. Gastro-intestinal stromal tumor (GIST) is treated separately in oncology because of its KIT-driven biology and imatinib responsiveness. Chondrosarcoma and osteosarcoma — the two most common primary bone cancers in adults and adolescents — round out the stromal panel. Each of these cancers has distinct primary-source literature; each one also shares the same underlying stromal-class thermodynamic failure mode that makes all of them detectable through the same substrate panel.

Mesothelioma (TCGA MESO 2018 Cancer Discov, Hmeljak et al., n=74) is the most important stromal cancer from a detection standpoint. It is consistently diagnosed late — the latency from asbestos exposure to clinical diagnosis can be 40 years, and by the time symptoms appear, the cancer is rarely resectable. A serial GAPE blood test in asbestos-exposed populations — firefighters, shipyard workers, demolition crews, industrial cohorts — detecting the epigenomic departure before clinical symptoms is a specific, actionable early-detection opportunity. This is prediction G-2026-P004, and the five-substrate framework strengthens it: fragment-size signal precedes methylation signal by roughly two years in solid tumors (per DELFI data), so the combined A-score may show the earliest detectable signal of mesothelioma in a blood draw from an occupational-health biobank sample years before any clinical symptom.

Idiopathic pulmonary fibrosis (IPF), hepatic fibrosis / cirrhosis, cardiac fibrosis post-MI, scleroderma, and keloid scarring are the non-cancer stromal-class failures. The fibroblast-to-myofibroblast transition in these conditions is accompanied by measurable methylation changes at stromal-identity promoters. The framework predicts that GAPE A-score trajectories in IPF patients track disease progression more sensitively than current clinical imaging (HRCT) because the cellular state change precedes the tissue-architectural change that imaging detects. IPF A_combined ≈ 1.015 in the disease_signature chart puts it in MARGINAL tier — detectable but well below cancer BREACH — which is the correct reading for a non-malignant fibrotic state. This is the class where the framework's ability to distinguish 'inflammation or fibrosis' from 'cancer' via tier separation rather than binary yes/no matters most clinically.

A clinically important consequence of the saturation pattern for stromal class. Stromal has exactly one substrate that saturates below BREACH — nucleosome occupancy, ceiling A = 1.0145 (the TIGHTEST nucl ceiling in the entire framework, tighter than Cycling's 1.020 and Secretary's 1.024). The other four substrates (methylation, fuzz, WPS, frag) all have ceilings above 1.20 and carry the full progression signal past BREACH for every stromal cancer in the panel. Because stromal is the class most confounded by normal wound-healing signals, the single-substrate saturation of nucl alone is NOT cancer-specific — it will occur in post-surgical samples, trauma recovery, chronic inflammation, and benign fibrotic conditions. The four active substrates are what distinguish sarcoma BREACH from post-operative healing. In clinical practice, this means a post-surgical patient with nucl

pinned at the ceiling but the other four substrates in NORMAL-to-MARGINAL tier is a healthy healing patient, not a missed cancer. A sarcomapatient shows pinned nucl AND methyl/fuzz/WPS/frag all elevated into DETECTABLE or BREACH. The mask for this class (+16.2%) is similar toCycling's +16.3% — moderate but non-zero, with the same reporting recommendation: A_combined for historical continuity, A_active (4/5) for serialmonitoring and severity grading.

FIVE-SUBSTRATE FIDELITY GAUGE

The gauge below is the central image of Issue 002. A single horizontal bar represents the A-score axis — from A = 0.90 (below the floor) through A = 1.00 (exactly at the Stromal class floor, H_min = 0.8629) up through the clinical threshold zones: MARGINAL at A = 1.01, DETECTABLE at 1.05, URGENT at 1.07, and FLOOR BREACH at 1.10. Five dots are plotted on this single axis for a healthy reference sample (Fibroblast IMR90, markers above the bar), and five more dots for a disease reference (TCGA SARC (n=206), markers below the bar). Each dot is one substrate: methylation, nucleosome occupancy, nucleosome fuzziness, WPS, or fragment size (DELFI). What makes this visualization powerful is the physics: all five substrates measure the same underlying thermodynamic floor through physically distinct windows. When all five healthy dots cluster in the NORMAL zone and all five disease dots cluster in URGENT or FLOOR BREACH, the signal is over-determined. Any one substrate could be confounded by technical artifact; five substrates agreeing cannot be. This is the "less blurry" advantage: $\sqrt{5}$ noise reduction for free.




Look at the spread in each direction. If the healthy five-dot cluster sits tightly together near A = 0.97–1.00, that indicates the reference population is at its architecture floor — precisely where the framework predicts healthy stromal cells should sit. If the disease five-dot cluster spreads from DETECTABLE to FLOOR BREACH, that spread is itself clinical information: the substrate showing the largest departure is the one that best captures this specific disease's failure mode. For this class, the primary failure mode is Stiffness Coupling, and the substrate ranking shown further down the card explains which of the five substrates best reveals it.

SUBSTRATE-BY-SUBSTRATE BREAKDOWN


Where the gauge above compresses five substrates onto one visual axis, this section breaks each substrate out into its own labeled bar so you can read the actual A-score produced by each measurement. The physics is the same across all five: take the raw laboratory value (a β methylation fraction, a nucleosome occupancy probability, a WPS score, etc.), compute its Shannon entropy $H(\text{value})$, then divide by the class-specific and substrate-specific H_{min} to get a dimensionless A-score comparable across substrates. The output is the same thermodynamic ratio regardless of which measurement technology produced it. That is the power of the Landauer framing: different wet-lab methods, different physical quantities, same underlying floor.

The formula at the foot of this section — $A_{\text{combined}} = \Sigma(AUC_i \times A_i) / \Sigma(AUC_i)$ — combines the five individual A-scores into a single weighted score, using each substrate's published detection AUC as its weight. Higher-AUC substrates contribute more to the combined result. For stromal cells, the healthy combined A-score below should sit at approximately 0.97 (NORMAL tier); the disease combined A-score should sit above 1.05 (DETECTABLE or higher). The difference between them — the ΔA combined — is the signal available to a blood draw that analyzes all five substrates simultaneously. A single-substrate assay sees a smaller signal at higher noise; the five-substrate combination sees the same signal at approximately $\sqrt{5}$ lower noise. Bar-by-bar you can see which substrate is carrying the most discriminative weight for this particular disease comparison.

HEALTHY REFERENCE — Fibroblast IMR90



DISEASE REFERENCE — TCGA SARC (n=206)
ACTIVE — tracks progression




Combined A — healthy (all 5)	0.97041 (NORMAL)
Combined A — disease (all 5)	1.09175 (URGENT)
ΔA combined (5 substrates)	+0.12134 vs methylation-only ΔA =+0.1502
Active A — healthy (non-saturated)	0.97041 (NORMAL) [5/5 active]
Active A — disease (non-saturated)	1.11140 (FLOOR BREACH) [4/5 active]

ΔA active (real progression signal)	+0.14100 (active substrates track true progression)
Saturation status (disease)	1 substrate saturated: Nucleosome Occupancy

Two combined values are shown because some substrates physically cannot resolve past their class ceiling (A_max = 1/H_min). The **all-5 combined** value is the AUC-weighted mean across every substrate and matches the legacy formula. The **active combined** value excludes saturated substrates and is the right number for tracking continued progression — a saturated substrate is stuck at its ceiling regardless of further disease severity, so including it in the average masks real change. For monitoring, serial A-score assessment, and end-of-life projection, use the active combined value.

DISEASE SIGNATURE COMPARISON — healthy muscle through IPF, pediatric sarcomas, TCGA SARC, to mesothelioma

Six conditions on a single chart spanning the full stromal-class clinical spectrum. β values reproduce per-substrate A-scores calibrated to primary sources: TCGA SARC 2017 Cell (n=206, Abeshouse et al.), TCGA MESO 2018 Cancer Discov (n=74, Hmeljak et al.), Crompton 2014 Cancer Discov Ewing sarcoma (n=112), and Shern 2014 Cancer Discov rhabdomyosarcoma (n=147). Pediatric sarcomas — Ewing and rhabdomyosarcoma — are shown alongside adult sarcomas because families of children lost to these diseases deserve to see their cancer on this card. IPF (idiopathic pulmonary fibrosis) is included as a non-cancer stromal-class failure where the same substrate readout applies. Nucleosome occupancy saturates at A ≈ 1.014 — the TIGHTEST ceiling of any class × substrate combination in the framework, tighter even than Progenitor WPS. All five cancers in the panel pin nucl at its ceiling; the other four substrates (methyl, fuzz, WPS, frag) carry the progression signal from DETECTABLE through FLOOR BREACH. For this class specifically, A_active (4/5) is the cleaner progression metric.




What happens past the ceiling is condition-specific. The A-score magnitude alone does not track severity once the ceiling is crossed — different diseases take different post-breach paths that the full five-substrate divergence pattern reveals. The subsection below walks through this class's post-breach trajectory: which substrates keep carrying signal, where the Warburg metabolic lock-in happens, where the glucose inversion point sits, and where the framework currently has high confidence versus where the open research questions remain.

POST-BREACH TRAJECTORY — WHAT HAPPENS AFTER A CROSSES 1.10 (STROMAL CLASS)

The stromal class covers sarcomas (adult and pediatric), mesothelioma, and the major non-cancer fibrotic diseases (IPF, scleroderma, cardiac fibrosis). Cfdna contribution from stromal cells is only ~4% of plasma — the lowest of any class — making detection technically demanding and favoring fragmentomics (most forgiving of dilution) over methylation for primary signal. Post-breach, stromal cancers share a common pattern with cycling cancers: methyl and frag carry the progression signal; nucl pins first and tightest.

The zone language below — Warburg boundary, glucose inversion, point of no return — applies only to post-breach cells. These terms are not used on the pre-breach bar, and the qualitative A-value anchors (~1.15, ~1.25, ~1.40+) are boundaries between therapeutic windows, not additional diagnostic tiers. Specific A-values are filed as prospective predictions to be validated against survival-stratified cohorts.



WHICH SUBSTRATES CARRY POST-BREACH SIGNAL — STROMAL CLASS

Stromal-class physics: nucleosome occupancy saturates at A ≈ 1.014 — the TIGHTEST ceiling of any class × substrate combination in the framework. The other four substrates carry the full progression signal.

Substrate	Ceiling (A)	Post-breach status
Methylation	1.22	Carries signal throughout all four zones
Fuzziness	1.21	Carries signal throughout all four zones
Windowed protection	1.15	Carries signal throughout all four zones
Fragment size	1.20	Carries signal throughout all four zones
Nucleosome occupancy	1.014	Saturated at ceiling (TIGHTEST in framework) — no further signal post-breach

INVERSION TERRITORY — SENESCENT FIBROBLASTS (non-cancer)

Stromal class has a documented non-cancer below-floor case: senescent fibroblasts. Cruickshanks 2013 Nature Cell Biology reports genome-wide methylation drift with aging and cellular senescence in dermal fibroblasts, with β values drifting below the healthy reference for a subset of promoter CpGs. The framework interprets this as a senescence-associated hypomethylation signature that puts A_methyl in the INVERSION zone on the pre-breach bar. This is clinically important: a senescent-fibroblast A-score that reads below-floor should not be mistaken for a reading error or for pre-cancer. It is a distinct biological state — the stromal analog of T-cell exhaustion in immune class — and it has its own therapeutic implications (senolytics, SASP-targeting therapies).

CONDITION-BY-CONDITION TRAJECTORY

Healthy muscle / fibroblast — reference, low cfDNA burden

Known: Stromal cells contribute ~4% of plasma cfDNA. Framework H_min is calibrated against fibroblast IMR90 (Roadmap). The low burden is the binding constraint on detection sensitivity; fragmentomics partially compensates because DELFI signal is proportionally preserved at low tissue fractions.

IPF (idiopathic pulmonary fibrosis) — A ≈ 1.015, MARGINAL tier (non-cancer)

Known: IPF is the framework's cleanest non-cancer floor-approach case in the stromal class. Fibroblast-to-myofibroblast transition produces measurable methylation changes at stromal-identity promoters. The framework reads this as A ≈ 1.015 — detectable but well below cancer BREACH — which is the correct clinical reading for a non-malignant fibrotic state.

Unknown: whether A-score trajectory in IPF patients tracks disease progression more sensitively than current HRCT imaging; whether antifibrotic therapy response (pirfenidone, nintedanib) shows A-score reduction before forced vital capacity stabilizes.

Test that closes the gap: G-2026-P012 (re-referenced): OSIC IPF Biobank reanalysis with serial cfDNA. Prediction: combined A-score trajectory slope will distinguish progressive from stable disease with AUC ≥ 0.80, outperforming serial FVC measurements.

TCGA SARC (adult sarcoma, n=206) — A ≈ 1.146, CROSSED CEILING

Known: Adult soft-tissue sarcoma (TCGA SARC 2017 Cell) shows stromal-class BREACH consistent with other classes' moderate-ΔA cancers. The post-surgical surveillance window is the framework's highest-value clinical use here: sarcoma recurs in 30-50% of patients after resection, and imaging surveillance alone misses early recurrence.

Unknown: whether post-resection A_active elevation predicts local or distant recurrence before imaging-detectable disease; whether histologic sarcoma subtypes (leiomyosarcoma, liposarcoma, UPS) trace distinguishable post-breach paths.

Test that closes the gap: G-2026-P033: Prospective cohort of 100 post-resection STS patients with serial cfDNA every 3 months for 24 months. Prediction: A_active elevation above baseline will detect recurrence 3-6 months before CT imaging in ≥ 60% of recurring cases.

Pediatric sarcomas (Ewing, rhabdomyosarcoma) — A ≈ 1.10-1.13, CROSSED CEILING

Known: Ewing sarcoma (Crompton 2014, Tirode 2014) and rhabdomyosarcoma (Shern 2014) are shown on this card because families of children lost to these cancers deserve to see their disease on the framework. The A-score biology mirrors adult sarcomas. The clinical need is identical: serial monitoring under VDC-IE (Ewing) or VAC (rhabdomyosarcoma) protocols to catch minimal residual disease.

Unknown: whether pediatric sarcoma MRD detection via cfDNA A-score outperforms current imaging; whether Ewing EWS-FLI1 fusion transcript levels correlate with A_active trajectory.

Test that closes the gap: G-2026-P034: COG (Children's Oncology Group) pediatric sarcoma cohort reanalysis with archived serial plasma. Prediction: A_active at end-of-induction predicts event-free survival with AUC ≥ 0.75.

WHAT THE FRAMEWORK CAN AND CANNOT SAY TODAY

The framework has high confidence for stromal-class post-breach: (1) four substrates carry full signal with nucl saturating tightest in framework; (2) IPF validation at A ≈ 1.015 correctly places non-cancer fibrosis in MARGINAL tier, not BREACH; (3) senescent fibroblasts produce genuine below-floor readings that the framework interprets coherently.

The framework has not yet tested sarcoma MRD detection prospectively, nor the pediatric sarcoma treatment-response predictions. G-2026-P012, P033, P034 define the specific validation plans.

This is the scientific posture the framework commits to: definitive where the evidence is definitive, open where it is open, and specific about the experiments that would close each gap. Predictions filed for this class: G-2026-P012, G-2026-P033, G-2026-P034. A reviewer who wants to falsify any claim here can do so against the named cohorts. A collaborator who wants to extend the framework has specific hypotheses to test.

THREE-COMPONENT DECOMPOSITION (C1/C2/C3) — PER SUBSTRATE

Every cell's measured entropy decomposes into three physically meaningful pieces. C1 is the universal Landauer floor: H_MIN_GLOBAL = 0.756500, measured in frontal cortex neurons by Lister 2013. Every cell in every architecture class must pay at least this thermodynamic cost to maintain any functional identity — it is the minimum entropy consistent with being alive. C2 is the class-specific overhead: the additional entropy the stromal class must carry above the universal floor to encode its particular architecture. For this class, C2 represents approximately 13.9% of the healthy reference entropy. C3 is the accessible gap: max(0, H_actual – H_min). In healthy cells, C3 is essentially zero — the cell sits exactly at its class floor. When C3 grows, the cell is departing from its architecture class. Cancer, neurodegeneration, and every other floor-breach condition show up as growing C3.

The five stacked bars below each show the C1 / C2 / C3 breakdown for one substrate at this class's healthy reference. Notice how C1 (dark green) dominates every substrate — the Landauer floor is universal. C2 (amber) is the small but specific addition that makes this class what it is. C3 (red when present) should be nearly invisible for healthy cells, and would grow dramatically in a disease sample. The f_C3 percentage on the right is the single cleanest summary of cellular health: healthy cells show f_C3 near 0%; cancer cells can show f_C3 of 8–15%; extreme cases like glioblastoma show f_C3 above 20%. This is not a statistical threshold calibrated from disease data. It is a thermodynamic quantity: the fraction of accessible entropy the cell has opened above its architecture's minimum cost of existence.



BEST TESTING METHOD RANKING — CLINICAL UTILITY FOR THIS CLASS

Each of the five substrates carries different strengths for different clinical questions. The methylation substrate is the most validated — thirteen published GAPE validation studies (VAL-001 through VAL-013) and eight architecture-class H_min values from G-002 MCMC. It is also the

substrate most affected by cfDNA dilution, because blood cfDNA is approximately 70% immune-derived. Fragment size (DELF) is the most forgiving of dilution — short-fragment enrichment persists even when tumor-derived cfDNA is a small fraction of total plasma DNA. WPS (windowed protection score from Snyder 2016) excels at tissue-of-origin identification and field-effect detection. Nucleosome occupancy and fuzziness add chromatin-level information that methylation alone cannot reveal.

The ranking below is specifically for the stromal class. The order reflects three factors: (1) the substrate's published single-substrate AUC for diseases in this class; (2) the biology of this class's primary failure mode (Stiffness Coupling); and (3) the practical constraint of cfDNA contribution in the relevant specimen (plasma for most classes, CSF for terminal, tissue biopsy where cfDNA is too dilute). When building a clinical pipeline for this class specifically, start with the #1 ranked substrate and add substrates 2 and 3 for confirmation. Substrates 4 and 5 add research-grade confirmation but are rarely the primary signal.

#	Substrate	Best for	Rationale
1	Fragment Size (DELF)	Cancer detection and monitoring	Stromal cfDNA contribution is only 4% — fragmentomics is the most forgiving of dilution. DELF signal holds at low tissue fractions.
2	Methylation	Floor departure confirmation	Secondary; cfDNA methylation signal for stromal class is diluted by immune background. Combine with deconvolution (EpiDISH, methylCC).
3	Windowed Protection Score	Fibrotic progression tracking	WPS at stromal identity promoters depletes as fibroblast-to-myofibroblast transition occurs. Detects fibrosis before clinical imaging.
4	Nucleosome Occupancy	Tumor microenvironment mapping	ATAC-seq on tumor stroma distinguishes activated cancer-associated fibroblasts from normal stromal cells. Research-grade primarily.
5	Nucleosome Fuzziness	Mechanical phenotype correlation	Stiffness-induced chromatin reorganization shows up as increased nucleosome fuzziness. Link to tumor microenvironment biomechanics.

BODY TEMPERATURE SCALING — $\alpha = 2.0$ LANDAUER CORRECTION

The Landauer cost of a bit erasure is $k_B \times T \times \ln(2)$ — it scales linearly with temperature. Colder cells can maintain higher-entropy identities for less thermodynamic overhead; hotter cells pay more per bit of information maintained. The GAPE framework captures this with a simple scaling law: $H_{\min}(T) = H_{\min}(37^\circ\text{C}) \times (T_{\text{body}} / 310.15\text{ K})^\alpha$, with $\alpha = 2.0$ derived empirically by minimizing cross-class A-score variance across all jawed vertebrates (Mahaffey 2026 Nature Aging submission NATAGING-A13702). The $\alpha = 2.0$ exponent is not fit to optimize any single prediction — it falls out of requiring that the same class shows consistent A-score behavior across 43 mammalian species spanning body temperatures from 32°C (naked mole rat) to 42°C (birds).

The table below shows what the stromal class floor $H_{\min}$ looks like at seven representative body temperatures. At 37°C (human reference, highlighted row), $H_{\min} = 0.8629$. At 42°C (birds), the floor is higher because the Landauer tax is larger. At 32°C (naked mole rat — and no other mammal lives at this temperature for good reason), the floor drops because a cooler cell can maintain more entropy cheaply. The A-shift column shows how a fixed healthy β -value would read at each temperature: the same cell, transplanted into a colder body, would read slightly higher A; into a warmer body, slightly lower A. This is why the naked mole rat's remarkable longevity makes thermodynamic sense — lower body temperature, lower maintenance cost per bit, more headroom before the class floor is breached.


T_body	Species example	H_min(T)	β_{healthy} equivalent	A-shift vs 37°C
42°C	Birds (chicken, finch)	0.8910	0.733	$A \times 0.969$
39°C	Rodents (mouse, rat)	0.8741	0.733	$A \times 0.987$
37°C	Humans (REFERENCE)	0.8629	0.733	$A \times 1.000$ (anchor)
35°C	Hibernating bats	0.8519	0.733	$A \times 1.013$
32°C	Naked mole rat	0.8354	0.733	$A \times 1.033$
25°C	Reptiles (anole lizard)	0.7975	0.733	$A \times 1.082$
15°C	Fish (salmon, cold-water)	0.7449	0.733	$A \times 1.159$

HEALTHY AGING TRAJECTORY — AGE-STRATIFIED REFERENCE A-SCORES

Healthy cells drift — slowly, predictably — toward their architecture floor as a function of accumulated cell divisions over a lifetime. For the stromal class, the drift rate is approximately 2.5% per generation — meaning each cell cycle carries a 2.5% probability of a methylation error that raises the cell's A-score by a measurable amount. Over decades, these errors accumulate. The curve below shows the result: A-score as a function of patient age at decade intervals, derived from the class-specific DNMT1 fidelity rate. At age 20, the reference healthy A-score sits well below 1.00; by age 80, it has drifted upward into the MARGINAL to DETECTABLE range. This is normal aging — not disease, just the statistical accumulation of

imperfect replication over eight decades.

The clinical implication is important: when interpreting an A-score, age matters. A patient at age 30 with A = 1.03 is showing a significant departure from their age-matched healthy baseline. The same A = 1.03 at age 75 might be entirely consistent with age-expected drift. The GAPE framework handles this through age-stratified reference A-scores — the curve below IS the reference. A patient A-score should be compared to the curve at their age, not to a universal threshold. Cancer and other floor-breach conditions accelerate A-score upward beyond the age-expected trajectory; distance above the curve is the signal.



VERTEBRATE LIFESPAN CONTEXT

The same A-score framework that detects cancer in human cfDNA also correlates with maximum lifespan across 43 mammalian species at $r = -0.9018$ ($p = 1.6 \times 10^{-16}$, Mahaffey 2026 Nature Aging submission NATAGING-A13702). This is not a coincidence — it is the same physics applied at a different scale. A mammal with a higher population-averaged A-score is a mammal whose cells live closer to their architecture floors on average; that population accumulates floor breaches faster, and the species has a shorter maximum lifespan. The $A = 1.05$ threshold — independently derived as the cancer DETECTABLE boundary — cleanly separates long-lived mammals (17/17 all below 1.05) from short-lived mammals (11/11 all above 1.05) with complete accuracy. Same threshold. Same physics. Different timescale.

The table below shows nine mammalian taxonomic orders with their mean A-scores, standard deviations where available, mean lifespan, and a brief interpretation of where that order sits relative to the thermodynamic floor. Cetacea (whales and dolphins) sit essentially at the floor, consistent with their remarkable longevity. Primates including humans sit slightly above the floor. Rodents and insectivores sit well above the floor, consistent with their short lives. This class's reference A-score (0.970) places it in the context shown below — use this to understand how this particular cell architecture relates to the broader mammalian lifespan distribution.

Order	N	Mean A	Lifespan (yr)	Interpretation
Cetacea	5	0.997 ±0.016	99	At thermodynamic floor
Proboscidea	1	0.987 —	70	At floor
Primates	6	1.007 ±0.015	49	Slightly above floor
Artiodactyla	4	1.015 ±0.011	30	Slightly above floor
Chiroptera	4	1.041 ±0.007	35	Above floor (longevity outliers)
Carnivora	9	1.053 ±0.023	28	Above floor
Lagomorpha	2	1.114 ±0.002	8	Well above floor
Rodentia	6	1.125 ±0.018	9	Well above floor
Insectivora	1	1.157 —	2	Furthest from floor

INTERVENTION LEVERS — RANKED BY EXPECTED IMPACT

If a patient's A-score is elevated — MARGINAL, DETECTABLE, or URGENT — the clinical question is what can be done. The GAPE framework does not prescribe treatment; that remains the physician's judgment. But the framework does identify which intervention categories have the mechanistic strength to move this particular class's A-score back toward the floor. The five categories below are the standard anti-aging and cancer-adjacent interventions: senolytics (clearing senescent cells), metabolic restoration (NAD+, caloric restriction, exercise), epigenetic restoration (DNMT/TET modulators, HDAC inhibitors), reprogramming (cyclic Yamanaka factors), and checkpoint stringency (cell cycle control, immune checkpoint blockade).

The ranking below is specifically for the stromal class and its primary failure mode (Stiffness Coupling). Impact Level 1 (Dominant) means the intervention directly addresses the class's failure mechanism and is the first lever to try. Level 2 (Strong) means substantial mechanistic support for A-score restoration in this class. Level 3 (Moderate) means helpful but not addressing the binding constraint. Level 4 (Limited) means the intervention works for other classes but does not target this class's specific biology. Level 5 (Not applicable) means the intervention category is biologically incompatible with this class — e.g., senolytics cannot act on post-mitotic terminal cells that do not become classically senescent, and reprogramming cannot be applied to terminal cells without losing the identity that defines them.

Impact	Category	Mechanism	Rationale
1 DOMINANT	Senolytics	senolytics	Dominant — senescent fibroblasts are the primary driver of stromal dysfunction
2 STRONG	Epigenetic Restoration	epigenetic rx	Strong — epigenetic resetting of pro-fibrotic methylation programs
3 MODERATE	Metabolic	metabolic	Moderate — metabolic normalization helps but senescent burden is the binding constraint
3 MODERATE	Checkpoint Stringency	checkpoint	Moderate — checkpoint modulation reduces fibrotic signaling cascade
4 LIMITED	Reprogramming	reprogramming	Limited — stromal architecture serves protective functions

CANCER PANEL — ΔA RANKED

The TCGA cancer validation is where the framework earns its keep. For every cancer type in this class, we compute two numbers from published TCGA methylation data: A_healthy (from matched normal tissue) and A_tumor (from cancer samples in the same study). The ΔA = A_tumor – A_healthy is the signal the framework predicts should appear — and it does. Thresholds for tier assignment: A = 1.05 DETECT, A = 1.07 URGENT, A = 1.10 FLOOR BREACH. These thresholds are not fit to the cancer data; they were derived from the healthy-vs-age statistics and happen to align with the cancer validation cleanly. The bars below are ordered by |ΔA| descending so the cancers with the largest signal appear first.

Note that ΔA ranges typical for this class matter. Terminal-class cancers (glioma, glioblastoma) show ΔA ≈ 0.22–0.27 — the largest in the entire 28-cancer TCGA dataset. Cycling epithelial cancers show ΔA ≈ 0.13–0.19 — moderate but clearly detectable. Stromal and progenitor-lineage cancers show ΔA ≈ 0.10–0.17. Each cancer below is cited to its primary TCGA publication and shows sample size. Direct links from A-score to TCGA evidence — no black boxes between physics and clinical data.

#1	Mesothelioma (MESO) n=87 · TCGA MESO 2018 Cell	ΔA=+0.1334	A=1.112 FLOOR
#2	Sarcoma (SARC) n=206 · TCGA SARC 2017 Cell	ΔA=+0.1204	A=1.109 FLOOR

CORE METRICS — PUBLISHED AND DERIVED

Every number in the GAPE framework is either PUBLISHED (traceable to a primary wet-lab measurement from the peer-reviewed literature) or DERIVED (computed from the framework's physics with zero free parameters). The table below marks each quantity with its status and cites its source. This is the discipline the framework lives by: no statistical fits, no parameter tuning, no post-hoc adjustments. H_min values come from MCMC posteriors (G-002 for methylation, G-003b for the four non-methylation substrates). Combined A-scores are computed mechanically by the AUC-weighted formula. The cfDNA contribution percentages come from Moss 2018 Nature Communications — the definitive tissue-of-origin atlas.

Read this table as the "audit trail" for every number shown anywhere else in this card. When the framework says A_healthy = 0.9701, you can see exactly which β value produced it, which H_min was used, and where both came from. The n_bio field is currently marked PRELIMINARY because absolute n_bio values await the G-007 MCMC run; the class ordering has been confirmed (Spearman ρ = 0.905, p = 0.002 against commitment level) but the absolute values carry that caveat transparently.

Metric	Value	Source
cfDNA contribution (plasma)	4.0%	Moss 2018 Nat Commun — PUBLISHED
H_min methylation (class)	0.862950	G-002 MCMC posterior — DERIVED
H_min methylation σ (MCMC)	±0.001400	G-002 17-chain MCMC R-hat<1.001 — DERIVED

Metric	Value	Source
n_bio (class-specific)	20.5	PRELIMINARY — ordering confirmed (p=0.905, p=0.002); absolute pending G-007
Healthy drift rate	2.5% / generation	DNMT1 fidelity loss under turnover — DERIVED
f_C2 (architecture-locked %)	13.9%	C2 fraction at healthy reference — DERIVED
Healthy A (methylation)	0.9701 (NORMAL)	β =0.733 — DERIVED
Disease A (methylation)	1.1202 (FLOOR BREACH)	β =0.607 — DERIVED
Healthy combined A (5 subs)	0.9704 (NORMAL)	$\Sigma(\text{AUC}_i \times A_i) / \Sigma(\text{AUC}_i)$ across 5 substrates — DERIVED
Disease combined A (5 subs)	1.0917 (URGENT)	$\Sigma(\text{AUC}_i \times A_i) / \Sigma(\text{AUC}_i)$ across 5 substrates — DERIVED
ΔA combined	+0.1213	Combined signal — DERIVED

DATED PREDICTIONS FOR THIS CLASS

A prediction is not a claim unless it can fail. Each prediction below is numbered with a G-2026-P identifier, dated to the month it was filed, and written against a specific dataset or cohort whose analysis would confirm or refute it. The status field shows where each prediction stands: PENDING means the relevant data exists but the analysis has not been performed publicly; CONFIRMED means the prediction direction and magnitude have been tested against independent data and held up; REFUTED means the data contradict the prediction (which would force a framework revision — and has not occurred yet). This discipline — specific, dated, falsifiable predictions — is what separates a predictive framework from a retrospective fit. Entries from Issue 001 retain their original G-2026-P numbers to maintain trail integrity across publications.

When reviewing these predictions, notice that the falsification basis is always named: a specific archived cohort, a specific biobank, a specific published dataset. These are not vague "more research is needed" statements. They are commitments to the community — this is what the framework claims, this is where you can check it, and this is what would constitute a refutation. The predictions below are for the stromal class specifically, filed as part of this publication or carried forward from Issue 001.

G-2026-P012

April 2026

PENDING

In prospective IPF cohorts with archived serial blood samples and known progression outcomes, the stromal-class combined A-score trajectory slope will distinguish progressive from stable disease with $\text{AUC} \geq 0.80$, outperforming serial forced vital capacity (FVC) measurements currently used for staging.

Basis: IPF progresses over 2-5 years. Existing antifibrotic therapies (pirfenidone, nintedanib) have variable response. The framework predicts that combined-substrate A-score trajectory identifies responders and non-responders earlier than FVC. Falsifiable in the OSIC IPF Biobank (www.osicild.org).

CELL IDENTITY & CLINICAL CONTEXT

Includes	Hematopoietic stem cells (HSC, CD34+CD38- compartment), interfollicular epidermal stem cells, hair follicle bulge stem cells, Merkel cell progenitors, cholangiocytes (hepatic stem cell origin), intestinal stem cells (Lgr5+), neural stem cells, muscle satellite cells, limbal stem cells
Cancers (this class)	HSC-origin AML (distinct from progenitor-lineage AML — arises from CD34+CD38- compartment per Adelman 2019), basal cell carcinoma (BCC — most common cancer in humans, epidermal stem cell origin), squamous cell carcinoma (SCC — hair-follicle stem-cell-origin subset), Merkel cell carcinoma (MCC — rare, aggressive, polyomavirus-positive and UV-driven subclasses), cholangiocarcinoma (CHOL, intrahepatic + extrahepatic, hepatic stem cell / cholangiocyte origin), cancer stem cell populations within cycling-class tumors (residual disease after primary treatment)
Non-cancer applications	HSC aging and immunosenescence (Adelman 2019, Beerman 2013), clonal hematopoiesis of indeterminate potential (CHIP — shared with Progenitor class), hematopoietic failure in the elderly, post-transplant niche reconstitution biology, heterochronic parabiosis biology, age-related decline in tissue regeneration across all adult-stem compartments
Primary failure mode	Niche Depletion
Warburg status	EMERGING

COMMENTARY

Adult tissue stem cells occupy the intermediate position between pluripotent embryonic stem cells and fully committed differentiated cells. They retain self-renewal capacity while being lineage-restricted: hematopoietic stem cells (HSC, CD34+CD38-) produce blood and immune cells, intestinal stem cells (Lgr5+) regenerate gut epithelium, epidermal and hair-follicle stem cells maintain skin, and cholangiocytes serve as the liver's stem-cell compartment for biliary regeneration. This partial commitment is encoded in H_min = 0.873718 for the methylation substrate, reflecting the slightly more open chromatin state that preserves self-renewal flexibility. But this open chromatin comes with a thermodynamic consequence that defines this card: healthy adult stem cells already sit near the maximum-entropy state for the positional substrates, leaving very little measurement headroom before ceiling saturation.

This class exhibits the most severe saturation pattern in the framework. Three of five substrates saturate below the BREACH threshold: nucleosome occupancy at ceiling A = 1.0407, nucleosome fuzziness at ceiling A = 1.0196, and windowed protection score at ceiling A = 1.0112 — the tightest ceiling of any class × substrate combination in the entire framework. Only methylation (ceiling A = 1.1445) and fragmentomics (ceiling A = 1.1886) remain active past BREACH, and these two substrates carry the full severity signal for any adult-stem-origin cancer in the panel. For this class, A_combined systematically understates severity by approximately 62 percent compared to A_active (2/5) — the largest masking effect in the framework. This is not a defect of the combined metric; it is an honest reflection of the underlying physics. Three of the five substrates have reached their physical entropy ceiling and cannot resolve further severity. A_active (2/5) is therefore reported as the primary severity metric for this class, with A_combined shown for cross-card continuity but flagged as misleading past the DETECTABLE tier.

The distinction between adult-stem-origin cancers and progenitor-lineage cancers is clinically consequential and must be stated explicitly. Acute myeloid leukemia appears in both Card #4 (Progenitor) and Card #7 (Adult Stem) because the disease arises from two biologically distinct compartments. Progenitor-lineage AML — the majority of TCGA AML cases characterized by Ley et al. in the 2013 New England Journal of Medicine analysis (n=200 whole-genome or exome sequences) — originates from committed myeloid progenitors (CMP, GMP) and carries the methylation and mutation signatures of those compartments. HSC-origin AML, characterized separately by Adelman et al. in Cancer Discovery 2019, arises from the pre-leukemic CD34+CD38- HSC-enriched compartment and involves profound epigenetic reprogramming of enhancers, bivalent promoters, and hematopoietic transcription factors (KLF6, BCL6, RUNX3) that begins during normal aging and appears to predispose to leukemic transformation. The two AML populations are not interchangeable in the framework: they occupy different architecture classes because they emerge from cells in different thermodynamic states. Listing AML on both cards without the distinction would be sloppy. Listing it only on one card would miss the patients whose disease originated in the other compartment.


Basal cell carcinoma (BCC) is the most common cancer in humans, with over four million new cases diagnosed annually in the United States alone. Nearly every BCC arises from epidermal or hair-follicle stem cells under UV-driven Hedgehog pathway activation. Mortality is low — below one percent — but the cumulative burden across lifetimes is enormous and the disease maps cleanly onto this class. Merkel cell carcinoma (MCC, Harms et al. 2015 Cancer Research, n=49 for the polyomavirus-negative subtype) is rare (approximately 2,500 US cases per year) but aggressive, with five-year survival around 45 percent for regional disease and below 15 percent for distant disease. Its two subclasses — Merkel cell polyomavirus-positive (integrated viral oncogenes) and MCPyV-negative (UV-damage driven) — share the same cell-of-origin methylation signature. Cholangiocarcinoma (TCGA CHOL, Farshidfar et al. 2017 Cell Reports, n=38) is the second most common primary liver cancer, aggressive, typically detected at advanced stage, and arises from cholangiocytes, which function as the hepatic stem-cell compartment for biliary regeneration. Its IDH-mutant molecular subtype shows characteristic ARID1A hypermethylation that maps onto the stem-cell-origin thermodynamic signature. Every one of these cancers is uncommon compared to the fourteen cycling cancers in Card #5, but the families of patients who lose a loved one to MCC or cholangiocarcinoma deserve to see the disease named on the framework card that describes its biology.

The class's defining thermodynamic challenge creates a specific detection problem: once a patient crosses DETECTABLE (A = 1.05), three of the five substrates pin at their ceilings within the first approximately 0.05 A-score units, and absolute A-score values across those three substrates lose their ability to resolve further severity. The framework proposes three mechanisms by which extended cellular thermodynamic activity remains detectable past this ceiling. First, methyl and frag retain their full dynamic range — they saturate only at A = 1.1445 and A = 1.1886 respectively, leaving meaningful headroom past BREACH for the two most clinically mature cfDNA biomarker platforms. A_{active} (2/5) computed on just these two substrates preserves severity discrimination that A_{combined} destroys. Second, the rate of approach to the saturating ceilings contains temporal information that single-point readings discard. In serial monitoring, a patient whose wps A-score moves from 0.970 through 1.000 to 1.010 over successive draws before plateauing has a different disease trajectory than one who hits 1.010 on first sample; time-to-saturation becomes a severity proxy in the regime where absolute-value measurement has lost resolution. Third, when adult-stem-origin cancers progress, the signal propagates into adjacent architecture classes — HSC-origin AML shows up as a rising immune-class A-score via clonal expansion, cutaneous stem-cell cancers propagate to cycling-class signals as field cancerization, and cholangiocarcinoma progression shows secretory-class signal change as the transformed cells commit toward biliary epithelial identity. Multi-class A-score divergence, particularly where the adult-stem class has flatlined at ceiling while one or more adjacent classes continue to rise, becomes a specific diagnostic signature of adult-stem-origin disease progression. This cross-class mechanism connects directly to the forthcoming Issue 004 framework on multi-class A-score divergence as an early signal of metastatic progression.

HSC aging is the most important non-cancer application for this class. The Adelman 2019 Cancer Discovery study of lineage-negative CD34+CD38- HSC-enriched cells from young and aged healthy donors identified 529 differentially methylated regions encompassing 2,249 differentially methylated cytosines across 748 genes — precisely the methylation-substrate floor departure the framework predicts for age-related adult-stem decline. Beerman et al. 2013 Cell Stem Cell independently characterized proliferation-dependent DNA hypermethylation of Polycomb-regulated genes in aged HSC (GSE44117). Together these studies establish that the methylation substrate detects HSC aging at the earliest stage of its epigenomic reprogramming, well before the downstream cytopenias or cancer-susceptibility phenotypes become clinically apparent. Clonal hematopoiesis of indeterminate potential (CHIP), the pre-malignant clonal state characterized by Steensma 2015 and shared in scope with the Progenitor class, represents the transition zone where adult-stem-class floor departure begins to predict downstream myeloid disease. Stem-cell transplantation biology, heterochronic parabiosis, and niche-depletion research all benefit from a metric that tracks adult-stem-class thermodynamic state independently of phenotypic readout.

FIVE-SUBSTRATE FIDELITY GAUGE

The gauge below is the central image of Issue 002. A single horizontal bar represents the A-score axis — from A = 0.90 (below the floor) through A = 1.00 (exactly at the Adult Stem class floor, H_{min} = 0.8737) up through the clinical threshold zones: MARGINAL at A = 1.01, DETECTABLE at 1.05, URGENT at 1.07, and FLOOR BREACH at 1.10. Five dots are plotted on this single axis for a healthy reference sample (HSC CD34+CD38- young, markers above the bar), and five more dots for a disease reference (HSC-origin AML, markers below the bar). Each dot is one substrate: methylation, nucleosome occupancy, nucleosome fuzziness, WPS, or fragment size (DELFI). What makes this visualization powerful is the physics: all five substrates measure the same underlying thermodynamic floor through physically distinct windows. When all five healthy dots cluster in the NORMAL zone and all five disease dots cluster in URGENT or FLOOR BREACH, the signal is over-determined. Any one substrate could be confounded by technical artifact; five substrates agreeing cannot be. This is the "less blurry" advantage: √5 noise reduction for free.




Look at the spread in each direction. If the healthy five-dot cluster sits tightly together near A = 0.97–1.00, that indicates the reference population is at its architecture floor — precisely where the framework predicts healthy adult stem cells should sit. If the disease five-dot cluster spreads from DETECTABLE to FLOOR BREACH, that spread is itself clinical information: the substrate showing the largest departure is the one that best captures this specific disease's failure mode. For this class, the primary failure mode is Niche Depletion, and the substrate ranking shown further down the card explains which of the five substrates best reveals it.

SUBSTRATE-BY-SUBSTRATE BREAKDOWN


Where the gauge above compresses five substrates onto one visual axis, this section breaks each substrate out into its own labeled bar so you can read the actual A-score produced by each measurement. The physics is the same across all five: take the raw laboratory value (a β methylation fraction, a nucleosome occupancy probability, a WPS score, etc.), compute its Shannon entropy $H(\text{value})$, then divide by the class-specific and substrate-specific H_{min} to get a dimensionless A-score comparable across substrates. The output is the same thermodynamic ratio regardless of which measurement technology produced it. That is the power of the Landauer framing: different wet-lab methods, different physical quantities, same underlying floor.

The formula at the foot of this section — $A_{\text{combined}} = \Sigma(AUC_i \times A_i) / \Sigma(AUC_i)$ — combines the five individual A-scores into a single weighted score, using each substrate's published detection AUC as its weight. Higher-AUC substrates contribute more to the combined result. For adult stem cells, the healthy combined A-score below should sit at approximately 0.97 (NORMAL tier); the disease combined A-score should sit above 1.05 (DETECTABLE or higher). The difference between them — the ΔA combined — is the signal available to a blood draw that analyzes all five substrates simultaneously. A single-substrate assay sees a smaller signal at higher noise; the five-substrate combination sees the same signal at approximately $\sqrt{5}$ lower noise. Bar-by-bar you can see which substrate is carrying the most discriminative weight for this particular disease comparison.

HEALTHY REFERENCE — HSC CD34+CD38- young



DISEASE REFERENCE — HSC-origin AML
ACTIVE — tracks progression




Combined A — healthy (all 5)	0.96999 (NORMAL)
Combined A — disease (all 5)	1.07222 (URGENT)
ΔA combined (5 substrates)	+0.10223 vs methylation-only ΔA =+0.1504
Active A — healthy (non-saturated)	0.96999 (NORMAL) [5/5 active]
Active A — disease (non-saturated)	1.13549 (FLOOR BREACH) [2/5 active]

ΔA active (real progression signal)	+0.16550 (active substrates track true progression)
Saturation status (disease)	3 substrates saturated: Nucleosome Occupancy, Nucleosome Fuzziness, Windowed Protection Score

Two combined values are shown because some substrates physically cannot resolve past their class ceiling ($A_{\text{max}} = 1/H_{\text{min}}$). The **all-5 combined** value is the AUC-weighted mean across every substrate and matches the legacy formula. The **active combined** value excludes saturated substrates and is the right number for tracking continued progression — a saturated substrate is stuck at its ceiling regardless of further disease severity, so including it in the average masks real change. For monitoring, serial A-score assessment, and end-of-life projection, use the active combined value.

DISEASE SIGNATURE COMPARISON — healthy HSC through aging, CHIP, BCC, MCC, cholangiocarcinoma, to HSC-origin AML

Six conditions on a single chart spanning the full adult-stem clinical spectrum. β values reproduce per-substrate A-scores calibrated to primary sources: Adelman 2019 Cancer Discov HSCe aging methylation (n=5-7 per age group), Beerman 2013 Cell Stem Cell HSC methylation landscape, Harms 2015 Cancer Res MCC genomic characterization (n=49), and Farshidfar 2017 Cell Reports TCGA CHOL (n=38). Three substrates saturate below BREACH: nucleosome occupancy (ceiling $A = 1.0407$), nucleosome fuzziness (ceiling $A = 1.0196$), and windowed protection score (ceiling $A = 1.0112$ — the tightest ceiling of any class $\times$ substrate combination in the entire framework). Only methylation and fragmentomics remain active past BREACH; A_{active} (2/5) is the primary severity metric for this class, with A_{combined} reported for cross-card continuity but flagged as misleading past DETECTABLE. BCC is shown because it is the most common cancer in humans; MCC is shown because its rarity does not reduce the obligation to honor the patients it kills; cholangiocarcinoma is shown because its stem-cell origin (cholangiocytes) places it squarely in this class.




What happens past the ceiling is condition-specific. The A-score magnitude alone does not track severity once the ceiling is crossed — different diseases take different post-breach paths that the full five-substrate divergence pattern reveals. The subsection below walks through this class's post-breach trajectory: which substrates keep carrying signal, where the Warburg metabolic lock-in happens, where the glucose inversion point sits, and where the framework currently has high confidence versus where the open research questions remain.

POST-BREACH TRAJECTORY — WHAT HAPPENS AFTER A CROSSES 1.10 (ADULT STEM CLASS)

The adult tissue stem class has the tightest H_min ceilings in the framework — WPS saturates at A ≈ 1.011 (the tightest single ceiling anywhere), fuzz at 1.020, and nucl at 1.041. Three substrates pin at or near the ceiling before BREACH is reached, leaving methylation and fragmentomics as the only substrates with meaningful post-breach headroom. This class has the clearest documented non-cancer inversion case in the framework: HSC aging and niche depletion.

The zone language below — Warburg boundary, glucose inversion, point of no return — applies only to post-breach cells. These terms are not used on the pre-breach bar, and the qualitative A-value anchors (~1.15, ~1.25, ~1.40+) are boundaries between therapeutic windows, not additional diagnostic tiers. Specific A-values are filed as prospective predictions to be validated against survival-stratified cohorts.



WHICH SUBSTRATES CARRY POST-BREACH SIGNAL — ADULT STEM CLASS

Adult-stem-class physics: three substrates saturate early (WPS, fuzz, nucl). Only methylation and fragmentomics carry post-breach progression signal. A_active (2/5) is the primary severity metric; A_combined is misleading past DETECTABLE.

Substrate	Ceiling (A)	Post-breach status
Methylation	1.14	Carries signal throughout all four zones
Fragment size	1.19	Carries signal throughout all four zones (deepest headroom)
Nucleosome occupancy	1.041	Saturated at ceiling — no further signal post-breach
Fuzziness	1.020	Saturated at ceiling — no further signal post-breach
Windowed protection	1.011	Saturated at ceiling (TIGHTEST in framework) — no further signal post-breach

INVERSION TERRITORY — NICHE DEPLETION INVERSION (documented)

Hematopoietic stem cells in aged bone marrow show a documented failure mode distinct from cancer: niche depletion. Adelman 2019 Cancer Discovery (HSC-enriched aging methylation, n=5-7 per age group) and Beerman 2013 Cell Stem Cell (GSE44117) report that aged HSCs lose methylation fidelity at lineage-commitment loci, producing what appears as a below-healthy A-score signature on methylation when compared to young HSC reference. This is the Niche Depletion Inversion — a non-cancer, non-malignant failure mode that the framework interprets coherently. Clinically, this distinguishes healthy aging from pre-leukemic states (CHIP, CCUS) which show upward elevation. The inversion direction itself is diagnostic: below-floor = aging; above-floor = clonal expansion.

CONDITION-BY-CONDITION TRAJECTORY

Healthy HSC (young adult) — reference, low cfDNA burden

Known: Hematopoietic stem cells represent ~3% of plasma cfDNA. H_min calibrated from Roadmap E035 HSC-enriched reference. Young-adult HSCs show tightest methylation fidelity of any stem cell population in the body.

HSC aging (Adelman 2019) — A below-floor (INVERSION territory)

Known: Aged HSCs (> 70 years) show progressive loss of methylation fidelity at lineage-commitment loci. The framework reads this as A_methyl in the INVERSION zone on the pre-breach bar — genuinely below healthy reference, not a measurement artifact. This is the non-cancer side of the adult-stem aging spectrum.

Unknown: whether longitudinal A-score tracking in aging adults identifies those at elevated risk for CHIP/CCUS transformation; whether interventions targeting HSC aging (senolytics, metformin) reverse the A-score back toward healthy.

Test that closes the gap: G-2026-P035: UK Biobank subset with serial blood archived and subsequent AML/MDS diagnosis outcomes. Prediction: subjects with persistent below-floor A_methyl over 10 years will have elevated risk (HR ≥ 2.0) for hematologic malignancy vs age-matched controls.

CHIP / CCUS (pre-malignant clonal) — A ≈ 1.02-1.05, MARGINAL-to-DETECTABLE

Known: Clonal hematopoiesis of indeterminate potential (Steensma 2015) and clonal cytopenias of undetermined significance (Malcovati 2017) represent pre-malignant clonal expansion. The framework reads CHIP/CCUS as modest A-score elevation — clearly above healthy, clearly below cancer. This is the directional inversion from HSC aging: the clonal expansion drives A up, not down.

Unknown: whether A-score trajectory within CHIP predicts which patients progress to MDS or AML over a 5-10 year horizon; whether specific CHIP mutation classes (DNMT3A, TET2, ASXL1) produce distinguishable A-score paths.

Test that closes the gap: G-2026-P036: Retrospective reanalysis of CCUS-to-MDS progression cohorts. Prediction: A_active slope above 0.01/year identifies imminent progression with AUC ≥ 0.75.

HSC-origin AML / MCC / cholangiocarcinoma — A ≈ 1.10-1.15, CROSSED CEILING

Known: HSC-origin adult leukemias and Merkel cell carcinoma sit in metabolic-window to structural-only zones post-breach. The two-substrate detection strategy (methyl + frag) outperforms five-substrate combining for this class because three substrates pin at their ceilings early.

Unknown: whether the two-substrate classifier AUC exceeds 0.80 in prospective cohorts; whether HSC-origin AML trajectories under azacitidine-venetoclax differ from progenitor-origin AML (immune-class card) at the A-score level.

Test that closes the gap: G-2026-P037: Paired cohort analysis of TCGA AML (n=200) reanalyzed with five-substrate cfDNA panels plus MCC (Harms 2015, n=49). Prediction: methyl+frag two-substrate classifier achieves AUC ≥ 0.80; any of three saturating substrates showing A above its predicted ceiling > measurement noise falsifies class H_min.

WHAT THE FRAMEWORK CAN AND CANNOT SAY TODAY

The framework has high confidence for adult-stem-class post-breach: (1) only methyl and frag carry post-breach signal — three substrates saturate before BREACH; (2) the Niche Depletion Inversion is documented in primary literature (Adelman 2019, Beerman 2013) and the framework interprets it coherently; (3) the directional distinction between aging (below-floor) and pre-malignant clonal expansion (above-floor) is clinically actionable.


The framework has not yet tested the HSC-aging-to-AML progression prediction prospectively, nor the two-substrate classifier AUC. G-2026-P035, P036, P037 define the specific validation plans.

This is the scientific posture the framework commits to: definitive where the evidence is definitive, open where it is open, and specific about the experiments that would close each gap. Predictions filed for this class: G-2026-P035, G-2026-P036, G-2026-P037. A reviewer who wants to falsify any claim here can do so against the named cohorts. A collaborator who wants to extend the framework has specific hypotheses to test.

THREE-COMPONENT DECOMPOSITION (C1/C2/C3) — PER SUBSTRATE

Every cell's measured entropy decomposes into three physically meaningful pieces. C1 is the universal Landauer floor: H_MIN_GLOBAL = 0.756500, measured in frontal cortex neurons by Lister 2013. Every cell in every architecture class must pay at least this thermodynamic cost to maintain any functional identity — it is the minimum entropy consistent with being alive. C2 is the class-specific overhead: the additional entropy the adult stem class must carry above the universal floor to encode its particular architecture. For this class, C2 represents approximately 13.5% of the healthy reference entropy. C3 is the accessible gap: max(0, H_actual – H_min). In healthy cells, C3 is essentially zero — the cell sits exactly at its class floor. When C3 grows, the cell is departing from its architecture class. Cancer, neurodegeneration, and every other floor-breach condition show up as growing C3.

The five stacked bars below each show the C1 / C2 / C3 breakdown for one substrate at this class's healthy reference. Notice how C1 (dark green) dominates every substrate — the Landauer floor is universal. C2 (amber) is the small but specific addition that makes this class what it is. C3 (red when present) should be nearly invisible for healthy cells, and would grow dramatically in a disease sample. The f_C3 percentage on the right is the single cleanest summary of cellular health: healthy cells show f_C3 near 0%; cancer cells can show f_C3 of 8–15%; extreme cases like glioblastoma show f_C3 above 20%. This is not a statistical threshold calibrated from disease data. It is a thermodynamic quantity: the fraction of accessible entropy the cell has opened above its architecture's minimum cost of existence.



BEST TESTING METHOD RANKING — CLINICAL UTILITY FOR THIS CLASS

Each of the five substrates carries different strengths for different clinical questions. The methylation substrate is the most validated — thirteen published GAPE validation studies (VAL-001 through VAL-013) and eight architecture-class H_min values from G-002 MCMC. It is also the

substrate most affected by cfDNA dilution, because blood cfDNA is approximately 70% immune-derived. Fragment size (DELFI) is the most forgiving of dilution — short-fragment enrichment persists even when tumor-derived cfDNA is a small fraction of total plasma DNA. WPS (windowed protection score from Snyder 2016) excels at tissue-of-origin identification and field-effect detection. Nucleosome occupancy and fuzziness add chromatin-level information that methylation alone cannot reveal.

The ranking below is specifically for the adult stem class. The order reflects three factors: (1) the substrate's published single-substrate AUC for diseases in this class; (2) the biology of this class's primary failure mode (Niche Depletion); and (3) the practical constraint of cfDNA contribution in the relevant specimen (plasma for most classes, CSF for terminal, tissue biopsy where cfDNA is too dilute). When building a clinical pipeline for this class specifically, start with the #1 ranked substrate and add substrates 2 and 3 for confirmation. Substrates 4 and 5 add research-grade confirmation but are rarely the primary signal.

#	Substrate	Best for	Rationale
1	Methylation	Primary severity metric past DETECTABLE	Methyl remains active past BREACH (ceiling A = 1.1445) — one of only two substrates that do not saturate for this class. Carries the bulk of the severity signal for HSC aging, BCC, MCC, and cholangiocarcinoma. Primary source: Adelman 2019 Cancer Discov.
2	Fragment Size (DELFI)	Primary severity metric past DETECTABLE	Frag (DELFI) remains active past BREACH (ceiling A = 1.1886, deepest headroom of any substrate in this class). Fragmentomic signal reflects clonal stem-cell turnover and cfDNA release kinetics. Co-primary with methyl for severity tracking.
3	Nucleosome Occupancy	Detection boundary — ceiling at A = 1.0407	Nucl saturates below BREACH. Useful for binary detection (floor departure vs healthy) but cannot resolve severity above DETECTABLE. Trajectory-slope information remains useful in serial monitoring (time-to-saturation as proxy severity).
4	Nucleosome Fuzziness	Detection boundary — ceiling at A = 1.0196	Fuzz saturates tighter than nucl. Same detection-only use case. Informative in combination with wps for cross-verification of floor departure.
5	Windowed Protection Score	Tightest ceiling in framework — A = 1.0112	WPS has the tightest physical ceiling of any class × substrate pairing. Healthy adult stem cells already sit at WPS A ≈ 0.970, with only +0.041 headroom before saturation. Any detectable deviation in wps for this class is clinically significant — but wps cannot distinguish mild from severe deviation above A = 1.0112.

BODY TEMPERATURE SCALING — $\alpha = 2.0$ LANDAUER CORRECTION

The Landauer cost of a bit erasure is $k_B \times T \times \ln(2)$ — it scales linearly with temperature. Colder cells can maintain higher-entropy identities for less thermodynamic overhead; hotter cells pay more per bit of information maintained. The GAPE framework captures this with a simple scaling law: $H_{\min}(T) = H_{\min}(37^\circ\text{C}) \times (T_{\text{body}} / 310.15\text{ K})^\alpha$, with $\alpha = 2.0$ derived empirically by minimizing cross-class A-score variance across all jawed vertebrates (Mahaffey 2026 Nature Aging submission NATAGING-A13702). The $\alpha = 2.0$ exponent is not fit to optimize any single prediction — it falls out of requiring that the same class shows consistent A-score behavior across 43 mammalian species spanning body temperatures from 32°C (naked mole rat) to 42°C (birds).

The table below shows what the adult stem class floor $H_{\min}$ looks like at seven representative body temperatures. At 37°C (human reference, highlighted row), $H_{\min} = 0.8737$. At 42°C (birds), the floor is higher because the Landauer tax is larger. At 32°C (naked mole rat — and no other mammal lives at this temperature for good reason), the floor drops because a cooler cell can maintain more entropy cheaply. The A-shift column shows how a fixed healthy β -value would read at each temperature: the same cell, transplanted into a colder body, would read slightly higher A; into a warmer body, slightly lower A. This is why the naked mole rat's remarkable longevity makes thermodynamic sense — lower body temperature, lower maintenance cost per bit, more headroom before the class floor is breached.


T_body	Species example	H_min(T)	$\beta_{\text{healthy equivalent}}$	A-shift vs 37°C
42°C	Birds (chicken, finch)	0.9021	0.726	A × 0.969
39°C	Rodents (mouse, rat)	0.8850	0.726	A × 0.987
37°C	Humans (REFERENCE)	0.8737	0.726	A × 1.000 (anchor)
35°C	Hibernating bats	0.8625	0.726	A × 1.013
32°C	Naked mole rat	0.8458	0.726	A × 1.033
25°C	Reptiles (anole lizard)	0.8074	0.726	A × 1.082

T_body	Species example	H_min(T)	β_healthy equivalent	A-shift vs 37°C
15°C	Fish (salmon, cold-water)	0.7542	0.726	$A \times 1.159$

HEALTHY AGING TRAJECTORY — AGE-STRATIFIED REFERENCE A-SCORES

Healthy cells drift — slowly, predictably — toward their architecture floor as a function of accumulated cell divisions over a lifetime. For the adult stem class, the drift rate is approximately 3.0% per generation — meaning each cell cycle carries a 3.0% probability of a methylation error that raises the cell's A-score by a measurable amount. Over decades, these errors accumulate. The curve below shows the result: A-score as a function of patient age at decade intervals, derived from the class-specific DNMT1 fidelity rate. At age 20, the reference healthy A-score sits well below 1.00; by age 80, it has drifted upward into the MARGINAL to DETECTABLE range. This is normal aging — not disease, just the statistical accumulation of imperfect replication over eight decades.

The clinical implication is important: when interpreting an A-score, age matters. A patient at age 30 with A = 1.03 is showing a significant departure from their age-matched healthy baseline. The same A = 1.03 at age 75 might be entirely consistent with age-expected drift. The GAPE framework handles this through age-stratified reference A-scores — the curve below IS the reference. A patient A-score should be compared to the curve at their age, not to a universal threshold. Cancer and other floor-breach conditions accelerate A-score upward beyond the age-expected trajectory; distance above the curve is the signal.



VERTEBRATE LIFESPAN CONTEXT

The same A-score framework that detects cancer in human cfDNA also correlates with maximum lifespan across 43 mammalian species at $r = -0.9018$ ($p = 1.6 \times 10^{-16}$, Mahaffey 2026 Nature Aging submission NATAGING-A13702). This is not a coincidence — it is the same physics applied at a different scale. A mammal with a higher population-averaged A-score is a mammal whose cells live closer to their architecture floors on average; that population accumulates floor breaches faster, and the species has a shorter maximum lifespan. The A = 1.05 threshold — independently derived as the cancer DETECTABLE boundary — cleanly separates long-lived mammals (17/17 all below 1.05) from short-lived mammals (11/11 all above 1.05) with complete accuracy. Same threshold. Same physics. Different timescale.

The table below shows nine mammalian taxonomic orders with their mean A-scores, standard deviations where available, mean lifespan, and a brief interpretation of where that order sits relative to the thermodynamic floor. Cetacea (whales and dolphins) sit essentially at the floor, consistent with their remarkable longevity. Primates including humans sit slightly above the floor. Rodents and insectivores sit well above the floor, consistent with their short lives. This class's reference A-score (0.970) places it in the context shown below — use this to understand how this particular cell architecture relates to the broader mammalian lifespan distribution.

Order	N	Mean A	Lifespan (yr)	Interpretation
Cetacea	5	0.997 ±0.016	99	At thermodynamic floor
Proboscidea	1	0.987 —	70	At floor
Primates	6	1.007 ±0.015	49	Slightly above floor
Artiodactyla	4	1.015 ±0.011	30	Slightly above floor
Chiroptera	4	1.041 ±0.007	35	Above floor (longevity outliers)
Carnivora	9	1.053 ±0.023	28	Above floor
Lagomorpha	2	1.114 ±0.002	8	Well above floor
Rodentia	6	1.125 ±0.018	9	Well above floor
Insectivora	1	1.157 —	2	Furthest from floor

INTERVENTION LEVERS — RANKED BY EXPECTED IMPACT

If a patient's A-score is elevated — MARGINAL, DETECTABLE, or URGENT — the clinical question is what can be done. The GAPE framework does not prescribe treatment; that remains the physician's judgment. But the framework does identify which intervention categories have the mechanistic strength to move this particular class's A-score back toward the floor. The five categories below are the standard anti-aging and cancer-adjacent interventions: senolytics (clearing senescent cells), metabolic restoration (NAD+, caloric restriction, exercise), epigenetic restoration (DNMT/TET modulators, HDAC inhibitors), reprogramming (cyclic Yamanaka factors), and checkpoint stringency (cell cycle control, immune checkpoint blockade).

The ranking below is specifically for the adult stem class and its primary failure mode (Niche Depletion). Impact Level 1 (Dominant) means the intervention directly addresses the class's failure mechanism and is the first lever to try. Level 2 (Strong) means substantial mechanistic support for A-score restoration in this class. Level 3 (Moderate) means helpful but not addressing the binding constraint. Level 4 (Limited) means the intervention works for other classes but does not target this class's specific biology. Level 5 (Not applicable) means the intervention category is biologically incompatible with this class — e.g., senolytics cannot act on post-mitotic terminal cells that do not become classically senescent, and reprogramming cannot be applied to terminal cells without losing the identity that defines them.

Impact	Category	Mechanism	Rationale
2 STRONG	Metabolic	metabolic	Strong — niche metabolic restoration moves stem cell fidelity index
2 STRONG	Epigenetic Restoration	epigenetic rx	Strong — epigenetic restoration extends stem cell functional lifespan
2 STRONG	Cyclic Reprogramming	reprogramming	Strong — cyclic Yamanaka rejuvenates without full dedifferentiation
2 STRONG	Niche Checkpoint	checkpoint	Strong — niche checkpoint signals regulate stem cell quiescence
3 MODERATE	Senolytics	senolytics	Moderate — senescent cells in the niche drive inversion

CANCER PANEL — ΔA RANKED

The TCGA cancer validation is where the framework earns its keep. For every cancer type in this class, we compute two numbers from published TCGA methylation data: A_healthy (from matched normal tissue) and A_tumor (from cancer samples in the same study). The ΔA = A_tumor – A_healthy is the signal the framework predicts should appear — and it does. Thresholds for tier assignment: A = 1.05 DETECT, A = 1.07 URGENT, A = 1.10 FLOOR BREACH. These thresholds are not fit to the cancer data; they were derived from the healthy-vs-age statistics and happen to align with the cancer validation cleanly. The bars below are ordered by |ΔA| descending so the cancers with the largest signal appear first.

Note that ΔA ranges typical for this class matter. Terminal-class cancers (glioma, glioblastoma) show ΔA ≈ 0.22–0.27 — the largest in the entire 28-cancer TCGA dataset. Cycling epithelial cancers show ΔA ≈ 0.13–0.19 — moderate but clearly detectable. Stromal and progenitor-lineage cancers show ΔA ≈ 0.10–0.17. Each cancer below is cited to its primary TCGA publication and shows sample size. Direct links from A-score to TCGA evidence — no black boxes between physics and clinical data.

#1

Uveal Melanoma (UVM)

n=80 · Robertson 2017 Cancer Cell

ΔA=+0.1072

A=1.086 URGENT

CORE METRICS — PUBLISHED AND DERIVED

Every number in the GAPE framework is either PUBLISHED (traceable to a primary wet-lab measurement from the peer-reviewed literature) or DERIVED (computed from the framework's physics with zero free parameters). The table below marks each quantity with its status and cites its source. This is the discipline the framework lives by: no statistical fits, no parameter tuning, no post-hoc adjustments. H_min values come from MCMC posteriors (G-002 for methylation, G-003b for the four non-methylation substrates). Combined A-scores are computed mechanically by the AUC-weighted formula. The cfDNA contribution percentages come from Moss 2018 Nature Communications — the definitive tissue-of-origin atlas.

Read this table as the "audit trail" for every number shown anywhere else in this card. When the framework says A_healthy = 0.9696, you can see exactly which β value produced it, which H_min was used, and where both came from. The n_bio field is currently marked PRELIMINARY because absolute n_bio values await the G-007 MCMC run; the class ordering has been confirmed (Spearman ρ = 0.905, p = 0.002 against commitment level) but the absolute values carry that caveat transparently.

Metric	Value	Source
cfDNA contribution (plasma)	3.0%	Moss 2018 Nat Commun — PUBLISHED
H_min methylation (class)	0.873718	G-002 MCMC posterior — DERIVED
H_min methylation σ (MCMC)	±0.001600	G-002 17-chain MCMC R-hat<1.001 — DERIVED
n_bio (class-specific)	18.5	PRELIMINARY — ordering confirmed (ρ=0.905, p=0.002); absolute pending G-007

Metric	Value	Source
Healthy drift rate	3.0% / generation	DNMT1 fidelity loss under turnover — DERIVED
f_C2 (architecture-locked %)	13.5%	C2 fraction at healthy reference — DERIVED
Healthy A (methylation)	0.9696 (NORMAL)	$\beta=0.726$ — DERIVED
Disease A (methylation)	1.1200 (FLOOR BREACH)	$\beta=0.586$ — DERIVED
Healthy combined A (5 subs)	0.9700 (NORMAL)	$\Sigma(\text{AUC}_i \times A_i) / \Sigma(\text{AUC}_i)$ across 5 substrates — DERIVED
Disease combined A (5 subs)	1.0722 (URGENT)	$\Sigma(\text{AUC}_i \times A_i) / \Sigma(\text{AUC}_i)$ across 5 substrates — DERIVED
ΔA combined	+0.1022	Combined signal — DERIVED

DATED PREDICTIONS FOR THIS CLASS

A prediction is not a claim unless it can fail. Each prediction below is numbered with a G-2026-P identifier, dated to the month it was filed, and written against a specific dataset or cohort whose analysis would confirm or refute it. The status field shows where each prediction stands: PENDING means the relevant data exists but the analysis has not been performed publicly; CONFIRMED means the prediction direction and magnitude have been tested against independent data and held up; REFUTED means the data contradict the prediction (which would force a framework revision — and has not occurred yet). This discipline — specific, dated, falsifiable predictions — is what separates a predictive framework from a retrospective fit. Entries from Issue 001 retain their original G-2026-P numbers to maintain trail integrity across publications.

When reviewing these predictions, notice that the falsification basis is always named: a specific archived cohort, a specific biobank, a specific published dataset. These are not vague "more research is needed" statements. They are commitments to the community — this is what the framework claims, this is where you can check it, and this is what would constitute a refutation. The predictions below are for the adult stem class specifically, filed as part of this publication or carried forward from Issue 001.

G-2026-P013

April 2026

PENDING

In prospective CHIP/CCUS cohorts with archived serial blood samples and known progression outcomes, the stem_adult A_active (2/5, methyl + frag only) trajectory will identify progression to overt MDS or AML 12 to 24 months before cytopenia-based diagnosis in a majority of progressing cases, outperforming A_combined (all 5) which is systematically masked by approximately 62 percent in this class because 3 of 5 substrates saturate below BREACH.

Basis:

CHIP-to-MDS/AML progression is slow, stereotyped, and clinically silent until late. The framework predicts that A_active identifies progressors earlier than cytopenia onset because the epigenomic state change precedes the hematologic phenotype, and A_active is not masked by ceiling-pinned positional substrates. Falsifiable in any CHIP/CCUS cohort with 18+ months of serial blood samples and progression outcomes (candidate cohorts include Jaiswal CHIP cohorts with archived samples, UK Biobank CHIP subset with longitudinal follow-up).

G-2026-P014

April 2026

PENDING

Trajectory slope ($\Delta A / \Delta t$) on the saturating substrates (nucl, fuzz, wps) during the pre-saturation window will correlate with clinical outcome in adult-stem-origin cancers with $\text{AUC} \geq 0.75$, providing severity information in the regime where absolute A-score values have lost resolution. Formally: time-to-saturation from healthy baseline is a valid severity proxy for this class.

Basis:

The saturation ceiling destroys absolute-value severity information but preserves temporal information in the form of rate-of-approach. Falsifiable in any serial-sampling cohort with outcome data; directly testable in MDS, BCC, MCC, and CHOL longitudinal cohorts.

G-2026-P015

April 2026

PENDING

Multi-class A-score divergence signature — adult-stem class flatlined at ceiling while an adjacent architecture class (immune, cycling, or secretory depending on tissue of origin) continues to rise — will identify adult-stem-origin disease progression with higher specificity than any single-class score past BREACH.

Basis:

When single-class saturation pins the primary class at ceiling, adjacent-class signal propagation becomes the discriminating severity signal. Cross-class divergence tests the framework's claim that architecture classes interact through shared tissue compartments and clonal expansion dynamics. Falsifiable in pan-cancer cohorts with multi-class A-score monitoring; connects directly to Issue 004 multi-class metastasis detection framework.

#8 · PLURIPOTENT STEM

$$\text{cfDNA} = 0.5\% \cdot H_{\min}(\text{methyl}) = 0.9822 \cdot n_{\text{bio}} = 16.5 \cdot \text{drift} \\ 2.5\%/\text{gen}$$

Reference: Human embryonic stem cells hESC H1 (Roadmap Epigenomics E003) · $H_{\min}$ confirmed from hESC H1 reference. Four TGCT histology subtypes show distinct, biologically-grounded epigenomic signatures per Shen 2018 and Killian 2016.

CELL IDENTITY & CLINICAL CONTEXT

Includes	Embryonic stem cells (hESC), induced pluripotent stem cells (iPSC), primordial germ cells (PGC), early-embryonic inner cell mass cells
Cancers (this class)	Testicular germ cell tumors (TGCT) — the most common cancer in men aged 15-35. Four pure histologic subtypes with biologically distinct methylation signatures (Shen 2018 TCGA TGCT n=137, Killian 2016 Genome Research n=130): (1) Seminoma — globally hypomethylated, PGC-like, ~60% of TGCT; (2) Embryonal carcinoma (EC) — hypermethylated at CpH sites, pluripotent-like methylation; (3) Yolk sac tumor (YST) — somatic-like methylation, differentiated extra-embryonic lineage; (4) Teratoma — somatic-like methylation, differentiated derivatives. Also: ovarian germ cell tumors (dysgerminoma — female analog of seminoma, shares molecular signature), primordial germ cell tumors in children, intratubular germ cell neoplasia in situ (GCNIS, the precursor lesion)
Non-cancer applications	iPSC reprogramming fidelity and quality control, organoid cell-state audit, synthetic embryology and gastruloid research, post-orchietomy surveillance for second primary TGCT (2-4% bilateral risk), cryptorchidism monitoring (4-10× elevated TGCT risk), fertility preservation context before platinum-based chemotherapy
Primary failure mode	Differentiation Dose Inversion + TGCT subtype-dependent methylation divergence
Warburg status	SPECIAL CASE — seminoma A_{active} decreases (inversion signal)

COMMENTARY

Testicular germ cell tumors (TGCT) are the most common cancer in men aged 15 to 35. Peak incidence falls directly in the prime childbearing years, which means the clinical audience for this card includes the urologist examining a 22-year-old graduate student who feels a painless testicular lump, the primary care physician who must decide whether a young father with a palpable scrotal mass needs immediate referral, and the oncologist who must discuss fertility preservation with a 30-year-old before starting platinum-based chemotherapy. The encouraging clinical reality is that TGCT is one of the most curable solid malignancies: the overall five-year survival exceeds 95 percent when the disease is caught at stage I or II, and cisplatin-based combination chemotherapy achieves cure in the majority of even metastatic cases. The discouraging reality is that delayed diagnosis — often driven by patient reluctance to seek evaluation for a scrotal abnormality, or by clinicians reluctant to order scrotal ultrasound in a young male presenting with back pain from retroperitoneal metastases — converts a 95 percent cure into salvage therapy with substantially lower long-term survival. Testicular self-examination remains the single highest-impact early detection intervention. This card exists in part to give clinicians the thermodynamic vocabulary to understand why a young man's tumor looks the way it does at the epigenomic level, and to support serial surveillance in the at-risk population.

Pluripotent stem cells occupy the unique architectural position of maximum developmental optionality. Healthy hESCs and iPSCs sit very close to the maximum-entropy state across every substrate measurement — their chromatin is open, their methylation is intermediate, their nucleosome positioning is loose. This is reflected in the class's $H_{\min}$ values: methylation $H_{\min} = 0.982166$ (ceiling $A = 1.0182$, only +0.018 past the healthy reference), fuzz $H_{\min} = 0.962920$ (ceiling $A = 1.0385$), frag $H_{\min} = 0.973583$ (ceiling $A = 1.0271$). Three of five substrates saturate below BREACH for this class. Only nucleosome occupancy (ceiling $A = 1.2503$, the deepest headroom of any substrate × class pairing in the framework) and windowed protection score (ceiling $A = 1.1050$) remain active past BREACH. The thermodynamic consequence is that pluripotent cells have almost no room to increase entropy in the conventional cancer direction — they are already near maximum entropy in their healthy state.

This is why TGCT disease signatures look fundamentally different from every other cancer in the framework. For most cancers, the signal is a uniform rise in A-score as cells depart from their architectural floor. For TGCT, the four pure histologic subtypes show biologically distinct methylation signatures that Shen et al. 2018 documented in TCGA TGCT (n=137) and Killian et al. 2016 Genome Research characterized in pure-histology samples (n=130) with direct comparison to primordial germ cell (PGC) and somatic methylation references. Seminoma — the most common TGCT histology, approximately 60 percent of cases — is globally hypomethylated. Its CpG methylation density moves toward the PGC state, with β values approaching zero. In the framework's language, A_{methyl} for seminoma FALLS BELOW the healthy reference, not above it. This is the inversion signal unique to this class: the only architecture class where the detection signal is a DECLINING A-score. Embryonal carcinoma (EC), by contrast, shows a striking convergence of both CpG and CpH methylation with pluripotent states, with elevated CpH methylation characteristic of the inner cell mass phase. Yolk sac tumor (YST) and teratoma show somatic-like methylation patterns that reflect their differentiation toward extra-embryonic and somatic lineages respectively. Four histologies, four distinct thermodynamic signatures, all arising from a common intratubular germ cell neoplasia in situ (GCNIS) precursor that is present in over 90 percent of TGCT cases and is itself clinically detectable.

The seminoma hypomethylation signal is the framework's most important inversion and the strongest single structural prediction. When primary-source data reports β values near zero for global CpG methylation in seminomas, the framework's A_{methyl} score drops well below the healthy hESC reference — at $\beta = 0.17$, A_{methyl} falls to approximately 0.670 compared to the healthy reference of 0.970. This has a specific detection consequence that clinicians and researchers must understand: A_{combined} alone is INSUFFICIENT for seminoma discrimination. Averaged across all five substrates, seminoma's A_{combined} sits near 0.97 — statistically indistinguishable from healthy. The discriminating signal is not a single elevated A-score but a CHARACTERISTIC DIVERGENCE PATTERN: A_{methyl} drops to 0.65 to 0.70 while A_{nucl} , A_{wps} , A_{fuzz} ,

and A_frag each modestly elevate into the 1.01 to 1.10 range. This multi-substrate divergence signature — one substrate dropping below floor while four substrates elevate above floor, with the specific relative magnitudes set by seminoma biology — is what the framework uses to detect seminoma. A naive combined-score detector would miss it. The framework's five-substrate decomposition is what makes this cancer detectable at all.


The clinical protocol for TGCT surveillance follows directly: instead of watching for A_combined to cross the BREACH threshold (1.10) as the framework does for every other cancer in the panel, the stem_pluri surveillance protocol watches for A_methyl to DECLINE toward 0.70 while A_nucl, A_wps, A_fuzz, A_frag simultaneously ELEVATE toward 1.05 or higher. This opposite-direction divergence is the detection signal. For patients with unilateral TGCT history (who carry 2 to 4 percent risk of contralateral second primary per Fossa 2005) and for patients with cryptorchidism (who carry 4 to 10 times elevated TGCT risk), serial monitoring with the correct multi-substrate divergence expectation is what makes thermodynamic surveillance clinically actionable. This is prediction G-2026-P005, originally proposed in Issue 001 and substantially refined here to account for the honest multi-substrate physics.

The clinical handling of TGCT carries specific demographic weight that belongs on this card. A thirty-year-old father of two who presents with a painless testicular nodule is the modal patient. The standard workup begins with scrotal ultrasound, proceeds through serum tumor markers (alpha-fetoprotein, beta-hCG, LDH), and typically concludes with radical inguinal orchiectomy as both diagnostic and therapeutic first intervention. Before orchiectomy — and certainly before any subsequent bleomycin-etoposide-cisplatin (BEP) chemotherapy that may be indicated for stage II or III disease — the conversation about fertility preservation belongs on the critical path. Sperm cryopreservation is well-established, inexpensive, and fully effective when offered in time. Platinum-based chemotherapy affects spermatogenesis; a substantial fraction of young TGCT survivors experience temporary or permanent impairment of fertility. A thermodynamic biomarker that supports early detection reduces the fraction of patients who reach the fertility-threatening chemotherapy stage. The framework's value proposition for the young father sitting in the oncologist's office is specifically this: earlier detection through thermodynamic surveillance in at-risk populations (cryptorchidism, prior contralateral TGCT, family history) means more patients caught at stage I where orchiectomy alone is often curative and intensive chemotherapy can be avoided.

The second major application for this class is iPSC reprogramming quality control, and this application has no cancer dimension at all — it is a research infrastructure problem. The Differentiation Dose Inversion, named for the well-documented phenomenon that excess Yamanaka factor dose produces aberrant rather than pluripotent states, maps onto the framework's prediction that successful reprogramming produces A-scores at or very near 1.00 across all five substrates. Aberrant reprogramming shows A below 0.95 (over-differentiation, cells partially committed to a lineage) or A above 1.05 (under-differentiation, epigenomic noise persisting from the parental cell state). Organoid researchers, synthetic embryology groups, and regenerative medicine translational programs all benefit from a thermodynamic quality-control metric that identifies aberrant colonies before they are committed to downstream experiments or clinical applications. This is a research-grade prediction testable in existing iPSC colony methylation and ATAC-seq datasets; prediction G-2026-P016 formalizes it.

FIVE-SUBSTRATE FIDELITY GAUGE

The gauge below is the central image of Issue 002. A single horizontal bar represents the A-score axis — from A = 0.90 (below the floor) through A = 1.00 (exactly at the Pluripotent class floor, H_min = 0.9822) up through the clinical threshold zones: MARGINAL at A = 1.01, DETECTABLE at 1.05, URGENT at 1.07, and FLOOR BREACH at 1.10. Five dots are plotted on this single axis for a healthy reference sample (hESC H1 (Roadmap E003), markers above the bar), and five more dots for a disease reference (Seminoma (Shen 2018), markers below). Each dot is one substrate: methylation, nucleosome occupancy, nucleosome fuzziness, WPS, or fragment size (DELFI). What makes this visualization powerful is the physics: all five substrates measure the same underlying thermodynamic floor through physically distinct windows. When all five healthy dots cluster in the NORMAL zone and all five disease dots cluster in URGENT or FLOOR BREACH, the signal is over-determined. Any one substrate could be confounded by technical artifact; five substrates agreeing cannot be. This is the "less blurry" advantage: √5 noise reduction for free.




Look at the spread in each direction. If the healthy five-dot cluster sits tightly together near A = 0.97–1.00, that indicates the reference population is at its architecture floor — precisely where the framework predicts healthy pluripotent cells should sit. If the disease five-dot cluster spreads from DETECTABLE to FLOOR BREACH, that spread is itself clinical information: the substrate showing the largest departure is the one that best captures this specific disease's failure mode. For this class, the primary failure mode is Differentiation Dose Inversion + TGCT subtype-dependent methylation divergence, and the substrate ranking shown further down the card explains which of the five substrates best reveals it.

SUBSTRATE-BY-SUBSTRATE BREAKDOWN

Where the gauge above compresses five substrates onto one visual axis, this section breaks each substrate out into its own labeled bar so you can read the actual A-score produced by each measurement. The physics is the same across all five: take the raw laboratory value (a β methylation fraction, a nucleosome occupancy probability, a WPS score, etc.), compute its Shannon entropy $H(\text{value})$, then divide by the class-specific and substrate-specific H_{min} to get a dimensionless A-score comparable across substrates. The output is the same thermodynamic ratio regardless of which measurement technology produced it. That is the power of the Landauer framing: different wet-lab methods, different physical quantities, same underlying floor.

The formula at the foot of this section — $A_{\text{combined}} = \Sigma(AUC_i \times A_i) / \Sigma(AUC_i)$ — combines the five individual A-scores into a single weighted score, using each substrate's published detection AUC as its weight. Higher-AUC substrates contribute more to the combined result. For pluripotent cells, the healthy combined A-score below should sit at approximately 0.97 (NORMAL tier); the disease combined A-score should sit above 1.05 (DETECTABLE or higher). The difference between them — the ΔA combined — is the signal available to a blood draw that analyzes all five substrates simultaneously. A single-substrate assay sees a smaller signal at higher noise; the five-substrate combination sees the same signal at approximately $\sqrt{5}$ lower noise. Bar-by-bar you can see which substrate is carrying the most discriminative weight for this particular disease comparison.

HEALTHY REFERENCE — hESC H1 (Roadmap E003)



DISEASE REFERENCE — Seminoma (Shen 2018)




Combined A — healthy (all 5)	0.97011 (NORMAL)
Combined A — disease (all 5)	0.96716 (NORMAL)
ΔA combined (5 substrates)	-0.00294 vs methylation-only ΔA =-0.3006
Active A — healthy (non-saturated)	0.97011 (NORMAL) [5/5 active]
Active A — disease (non-saturated)	0.96716 (NORMAL) [5/5 active]

ΔA active (real progression signal)	-0.00294 (active substrates track true progression)
Saturation status (disease)	No substrates saturated

Two combined values are shown because some substrates physically cannot resolve past their class ceiling ($A_{\text{max}} = 1/H_{\text{min}}$). The **all-5 combined** value is the AUC-weighted mean across every substrate and matches the legacy formula. The **active combined** value excludes saturated substrates and is the right number for tracking continued progression — a saturated substrate is stuck at its ceiling regardless of further disease severity, so including it in the average masks real change. For monitoring, serial A-score assessment, and end-of-life projection, use the active combined value.

DISEASE SIGNATURE COMPARISON — healthy hESC through the four TGCT histologies

Six conditions on a single chart spanning the full TGCT histologic spectrum plus healthy and precursor states. β values reproduce per-substrate A-scores calibrated to Shen 2018 Cell Reports (TCGA TGCT, $n=137$) and Killian 2016 Genome Research ($n=130$ pure-histology TGCTs with PGC and somatic reference comparisons). Three substrates saturate on this class: methyl (ceiling $A = 1.0182$), fuzz (ceiling $A = 1.0385$), frag (ceiling $A = 1.0271$). Nucl remains active with +0.150 headroom past BREACH; wps remains active with modest headroom. The four TGCT histologies show biologically distinct methylation signatures: seminoma is globally hypomethylated (methyl β drops toward PGC state, A_{methyl} FALLS below healthy — the inversion signal for this class), embryonal carcinoma is hypermethylated at CpH sites with pluripotent-like CpG patterns, yolk sac tumor and teratoma show somatic-like methylation reflecting their differentiation toward extra-embryonic and somatic lineages respectively. GCNIS (germ cell neoplasia in situ) is the clinically detectable precursor lesion present in over 90% of TGCT cases. The chart reads left to right as the disease progression most relevant to surveillance: healthy $\rightarrow$ GCNIS precursor $\rightarrow$ seminoma / EC / YST / teratoma as distinct presenting histologies.




What happens past the ceiling is condition-specific. The A-score magnitude alone does not track severity once the ceiling is crossed — different diseases take different post-breach paths that the full five-substrate divergence pattern reveals. The subsection below walks through this class's post-breach trajectory: which substrates keep carrying signal, where the Warburg metabolic lock-in happens, where the glucose inversion point sits, and where the framework currently has high confidence versus where the open research questions remain.

POST-BREACH TRAJECTORY — WHAT HAPPENS AFTER A CROSSES 1.10 (PLURIPOTENT CLASS)

The pluripotent stem class is the framework’s showcase for inversion physics. Seminoma — 60% of testicular germ cell tumors — is globally hypomethylated rather than hypermethylated, producing A_methyl that FALLS below healthy reference (A_methyl ≈ 0.67 at β = 0.17) while A_nucl, A_wps, A_fuzz, A_frag elevate normally. This is the cleanest documented below-floor disease state in the framework. The clinical signal for seminoma is not combined-A elevation — it is a multi-substrate DIVERGENCE pattern: one substrate inverting while four elevate. A naive cancer detector looking only at combined A misses seminoma entirely. VAL-045 (April 2026) revealed an important extension: because H_min_methyl = 0.9822 sits very close to the Shannon ceiling of 1.000, the architectural window above floor is extremely narrow. Any methylation departure from the narrow pluripotent reference band produces A_methyl below floor regardless of direction. All four TGCT histologies invert at the methylation level: seminoma is the extreme case (A_methyl ≈ 0.75, β ≈ 0.21), embryonal carcinoma lands at A ≈ 0.83 (β ≈ 0.745), yolk sac tumor at A ≈ 0.87 (β ≈ 0.72), and choriocarcinoma at A ≈ 0.85 (β ≈ 0.735). Specificity in this class comes from divergence magnitude, not divergence direction — seminoma’s divergence is 2.1× that of other TGCT histologies. Inversion is class-universal for pluripotent; seminoma is the extreme, not the only case.

The zone language below — Warburg boundary, glucose inversion, point of no return — applies only to post-breach cells. These terms are not used on the pre-breach bar, and the qualitative A-value anchors (~1.15, ~1.25, ~1.40+) are boundaries between therapeutic windows, not additional diagnostic tiers. Specific A-values are filed as prospective predictions to be validated against survival-stratified cohorts.



WHICH SUBSTRATES CARRY POST-BREACH SIGNAL — PLURIPOTENT CLASS

Pluripotent-class physics: three substrates saturate early (methyl, fuzz, frag) at class-specific ceilings close to healthy reference. Nucl and WPS carry post-breach progression signal, with nucl having the deepest headroom in the class. The inversion signal lives on methyl — when β drives toward PGC-like near-zero, A_methyl falls below healthy rather than rising.

Substrate	Ceiling (A)	Post-breach status
Nucleosome occupancy	1.25	Primary severity metric — deepest post-breach headroom
Windowed protection	1.11	Near-BREACH severity metric — modest headroom
Methylation	1.018	Saturated ceiling on the UPWARD axis; but INVERSION-capable on the DOWNWARD axis
Fuzziness	1.039	Saturated at ceiling — no further upward signal post-breach
Fragment size	1.027	Saturated at ceiling — no further signal post-breach

INVERSION TERRITORY — SEMINOMA HYPOMETHYLATION INVERSION (showcase documented case)

The Seminoma Hypomethylation Inversion is the framework’s cleanest documented below-floor disease state. Seminoma tumor β_methyl drops toward 0.17-0.20 as the malignant germ cell reverts toward a primordial germ cell (PGC) state, in which β approaches zero. Because pluripotent H_min_methyl = 0.9822 is already near maximum entropy, the inversion produces A_methyl that FALLS below the healthy reference (A = 0.67 at β = 0.17) rather than rising above it. A naive cancer-detection instrument looking for combined-A elevation misses seminoma entirely. The framework’s discrimination signal is the multi-substrate divergence pattern: A_methyl drops to 0.65-0.70 while A_nucl, A_wps, A_fuzz, A_frag simultaneously elevate into the 1.01-1.10 range. Divergence = max|A_i - median(all A)| — a divergence ≥ 0.10 is the framework’s seminoma signature. Confirmed in Shen 2018 TCGA TGCT (n=137) and Killian 2016 Genome Research (n=130 pure-histology samples with direct PGC comparison). This is why every card now shows the INVERSION zone on the pre-breach bar — because below-floor disease is real, documented, and class-specific. VAL-045 (April 2026, multi-class drift cascade) extended this picture: because the pluripotent methylation window above floor is so narrow (H_min = 0.9822 versus H_max = 1.000), any methylation departure from the reference band inverts A_methyl regardless of direction. All four TGCT histologies land in inversion territory at the methylation level. Seminoma is the extreme case at A_methyl ≈ 0.75 (divergence 0.079); embryonal carcinoma at ≈ 0.83, yolk sac at ≈ 0.87, and choriocarcinoma at ≈ 0.85 (divergences 0.037 or less). Seminoma’s divergence is 2.1× that of other histologies — the specificity is in divergence magnitude, not in inversion direction. Inversion is class-universal for pluripotent-class cancers; seminoma is the extreme case, not the only case. This refinement does not change the framework’s discrimination protocol (divergence ≥ 0.10 still identifies seminoma), but it clarifies that the other TGCT histologies also depart architecturally into inversion territory — just at smaller magnitudes.

CONDITION-BY-CONDITION TRAJECTORY

Healthy hESC (reference) — pluripotent baseline

Known: Human embryonic stem cells (Roadmap E003) sit at the class reference with $A \approx 0.97$ across all five substrates. Pluripotent cells are designed for high entropy — their $H_{\min} = 0.9822$ is the highest class floor in the framework, because multi-potency requires maximum epigenomic reversibility. The Yamanaka reprogramming factors (OCT4, SOX2, KLF4, c-MYC) drive somatic cells back toward this high-entropy state.

Seminoma (INVERSION signature) — $A_{\text{methyl}} \approx 0.67$ (inversion), $A_{\text{nucl/wps/fuzz/frag}} \approx 1.01-1.10$

Known: Seminoma is the framework's showcase inversion case. The cancer manifests as divergence, not elevation: methyl crashes toward PGC state while the other four substrates rise. Early-detection scenario (Card Section 5.1) shows the framework flagging divergence at Month 12 — 6 months before standard ultrasound surveillance would detect a mass. This is the framework's highest-value clinical case: a cancer that current molecular panels miss entirely.

Unknown: whether the framework can detect seminoma in cryptorchidism surveillance cohorts (4-10× elevated TGCT risk) before clinical presentation; whether post-orchietomy serial A-score can identify second primary TGCT (2-4% bilateral risk) earlier than CT imaging.

Test that closes the gap: G-2026-P005 (filed): Prospective surveillance cohort of cryptorchidism patients with baseline plus 6-monthly cfDNA. Framework prediction: divergence ≥ 0.10 identifies pre-clinical seminoma with sensitivity ≥ 0.80 at specificity 0.95, providing 6+ month lead time vs ultrasound + tumor marker surveillance.

Embryonal carcinoma (EC) — $A \approx 1.05-1.10$, UPWARD path (hypermethylation)

Known: Embryonal carcinoma shows the OPPOSITE methylation direction from seminoma: CpH hypermethylation characteristic of the inner cell mass phase, pluripotent-like CpG patterns. A_{methyl} elevates rather than inverts. This is the framework's cleanest demonstration that direction on the A-score axis is diagnostic — the same class can produce either inversion or elevation depending on the specific malignant reprogramming pathway.

Unknown: whether EC-dominant mixed germ cell tumors show distinguishable A-score trajectories from pure EC; whether response to BEP chemotherapy produces trajectory bend toward healthy (A_{methyl} declining from elevation rather than recovering from inversion).

Test that closes the gap: G-2026-P017 (filed): Prospective cohort of 100 stage II-III TGCT patients undergoing BEP with serial cfDNA pre-C1, pre-C2, pre-C3, pre-C4. Prediction: responders show A_{methyl} moving toward healthy hESC reference (approaching 0.970 from either direction — recovering from seminoma inversion or declining from EC hypermethylation); non-responders remain pinned at disease baseline.

iPSC reprogramming (research-grade inversion) — $A \approx 0.90-1.05$, DIFFERENTIATION DOSE INVERSION

Known: In iPSC reprogramming protocols, excess Yamanaka factor dose produces aberrant rather than successfully reprogrammed colonies. Successful reprogramming produces $A \approx 1.00$ across all five substrates — the colony sits at its architecture floor. Aberrant colonies show A below 0.95 (over-differentiation) or above 1.05 (under-differentiation). The pharmacologic dose-response is non-monotone: more is not better past the optimal window. This is a research-grade application of the framework, not clinical, but it demonstrates inversion physics in a controlled setting.

Unknown: whether the framework can quality-control iPSC reprogramming at single-colony resolution with ATAC-seq inputs; whether Differentiation Dose Inversion is reproducible across reprogramming protocols (OSKM, episomal, Sendai, mRNA).

Test that closes the gap: G-2026-P016 (filed): Prospective iPSC reprogramming cohort with ATAC-seq on individual colonies at day 20, 30, 40 post-transduction. Prediction: colonies with A within 0.97-1.03 across all five substrates produce teratomas successfully ($\geq 95\%$); colonies outside this window fail differentiation QC ($\geq 80\%$).

WHAT THE FRAMEWORK CAN AND CANNOT SAY TODAY

The framework has high confidence for pluripotent-class post-breach: (1) the Seminoma Hypomethylation Inversion is the framework's cleanest documented below-floor case, with primary-source validation in Shen 2018 and Killian 2016; (2) nucl and WPS carry post-breach severity signal while methyl, fuzz, and frag saturate near healthy reference; (3) multi-substrate divergence (not combined elevation) is the discrimination signal for this class, demonstrating that direction on the A-axis is diagnostic.

The framework has not yet tested seminoma early detection in cryptorchidism surveillance prospectively, nor BEP platinum response trajectory, nor iPSC quality control. G-2026-P005, P016, P017 define the specific validation plans.


This is the scientific posture the framework commits to: definitive where the evidence is definitive, open where it is open, and specific about the experiments that would close each gap. Predictions filed for this class: G-2026-P005, G-2026-P016, G-2026-P017. A reviewer who wants to falsify any claim here can do so against the named cohorts. A collaborator who wants to extend the framework has specific hypotheses to test.

THREE-COMPONENT DECOMPOSITION (C1/C2/C3) — PER SUBSTRATE

Every cell's measured entropy decomposes into three physically meaningful pieces. C1 is the universal Landauer floor: $H_{\min_GLOBAL} = 0.756500$, measured in frontal cortex neurons by Lister 2013. Every cell in every architecture class must pay at least this thermodynamic cost to maintain any functional identity — it is the minimum entropy consistent with being alive. C2 is the class-specific overhead: the additional entropy the

pluripotent class must carry above the universal floor to encode its particular architecture. For this class, C2 represents approximately 3.0% of the healthy reference entropy. C3 is the accessible gap: $\max(0, H_{\text{actual}} - H_{\text{min}})$. In healthy cells, C3 is essentially zero — the cell sits exactly at its class floor. When C3 grows, the cell is departing from its architecture class. Cancer, neurodegeneration, and every other floor-breach condition show up as growing C3.

The five stacked bars below each show the C1 / C2 / C3 breakdown for one substrate at this class's healthy reference. Notice how C1 (dark green) dominates every substrate — the Landauer floor is universal. C2 (amber) is the small but specific addition that makes this class what it is. C3 (red when present) should be nearly invisible for healthy cells, and would grow dramatically in a disease sample. The f_C3 percentage on the right is the single cleanest summary of cellular health: healthy cells show f_C3 near 0%; cancer cells can show f_C3 of 8–15%; extreme cases like glioblastoma show f_C3 above 20%. This is not a statistical threshold calibrated from disease data. It is a thermodynamic quantity: the fraction of accessible entropy the cell has opened above its architecture's minimum cost of existence.



BEST TESTING METHOD RANKING — CLINICAL UTILITY FOR THIS CLASS

Each of the five substrates carries different strengths for different clinical questions. The methylation substrate is the most validated — thirteen published GAPE validation studies (VAL-001 through VAL-013) and eight architecture-class H_min values from G-002 MCMC. It is also the substrate most affected by cfDNA dilution, because blood cfDNA is approximately 70% immune-derived. Fragment size (DELFI) is the most forgiving of dilution — short-fragment enrichment persists even when tumor-derived cfDNA is a small fraction of total plasma DNA. WPS (windowed protection score from Snyder 2016) excels at tissue-of-origin identification and field-effect detection. Nucleosome occupancy and fuzziness add chromatin-level information that methylation alone cannot reveal.

The ranking below is specifically for the pluripotent class. The order reflects three factors: (1) the substrate's published single-substrate AUC for diseases in this class; (2) the biology of this class's primary failure mode (Differentiation Dose Inversion + TGCT subtype-dependent methylation divergence); and (3) the practical constraint of cfDNA contribution in the relevant specimen (plasma for most classes, CSF for terminal, tissue biopsy where cfDNA is too dilute). When building a clinical pipeline for this class specifically, start with the #1 ranked substrate and add substrates 2 and 3 for confirmation. Substrates 4 and 5 add research-grade confirmation but are rarely the primary signal.

#	Substrate	Best for	Rationale
1	Methylation	Histology discrimination — bidirectional signal	The only substrate whose signal direction differs by TGCT histology. Seminoma drives methyl DOWN toward PGC-like hypomethylation (A falls below 0.970 — the inversion signal). Embryonal carcinoma drives methyl UP (CpH hypermethylation). Yolk sac and teratoma show somatic-like methylation. Primary source: Shen 2018 Cell Reports.
2	Nucleosome Occupancy	Primary severity metric past DETECTABLE	Only substrate with substantial headroom past BREACH (ceiling A = 1.2503). For any TGCT histology, nucl carries severity signal reliably. Also primary metric for iPSC reprogramming fidelity assessment (ATAC-seq on colonies).
3	Windowed Protection Score	Near-BREACH severity metric	Active past BREACH with modest headroom (ceiling A = 1.1050). WPS at pluripotency promoters (OCT4, SOX2, NANOG) detects the persistence of embryonic-gene expression characteristic of all TGCT histologies.
4	Nucleosome Fuzziness	Detection boundary — ceiling at A = 1.0385	Saturates near BREACH threshold. Useful for binary detection of floor departure; cannot resolve histology differences above ceiling.
5	Fragment Size (DELFI)	Detection boundary — ceiling at A = 1.0271	Saturates below BREACH. Fragmentomic signal for TGCT is dominated by rapid tumor turnover rate rather than methylation specifics; useful for detection but not histology typing.

BODY TEMPERATURE SCALING — α = 2.0 LANDAUER CORRECTION

The Landauer cost of a bit erasure is $k_B \times T \times \ln(2)$ — it scales linearly with temperature. Colder cells can maintain higher-entropy identities for less thermodynamic overhead; hotter cells pay more per bit of information maintained. The GAPE framework captures this with a simple scaling law: $H_{\text{min}}(T) = H_{\text{min}}(37^\circ\text{C}) \times (T_{\text{body}} / 310.15 \text{ K})^\alpha$, with $\alpha = 2.0$ derived empirically by minimizing cross-class A-score variance across all jawed

vertebrates (Mahaffey 2026 Nature Aging submission NATAGING-A13702). The $\alpha = 2.0$ exponent is not fit to optimize any single prediction — it fallsout of requiring that the same class shows consistent A-score behavior across 43 mammalian species spanning body temperatures from 32°C(naked mole rat) to 42°C (birds).


The table below shows what the pluripotent class floor H_{min} looks like at seven representative body temperatures. At 37°C (human reference, highlighted row), $H_{min} = 0.9822$. At 42°C (birds), the floor is higher because the Landauer tax is larger. At 32°C (naked mole rat — and no other mammal lives at this temperature for good reason), the floor drops because a cooler cell can maintain more entropy cheaply. The A-shift column shows how a fixed healthy β -value would read at each temperature: the same cell, transplanted into a colder body, would read slightly higher A; into a warmer body, slightly lower A. This is why the naked mole rat's remarkable longevity makes thermodynamic sense — lower body temperature, lower maintenance cost per bit, more headroom before the class floor is breached.

T_body	Species example	H_min(T)	$\beta_{healthy}$ equivalent	A-shift vs 37°C
42°C	Birds (chicken, finch)	1.0141	0.627	$A \times 0.969$
39°C	Rodents (mouse, rat)	0.9949	0.627	$A \times 0.987$
37°C	Humans (REFERENCE)	0.9822	0.627	$A \times 1.000$ (anchor)
35°C	Hibernating bats	0.9695	0.627	$A \times 1.013$
32°C	Naked mole rat	0.9508	0.627	$A \times 1.033$
25°C	Reptiles (anole lizard)	0.9076	0.627	$A \times 1.082$
15°C	Fish (salmon, cold-water)	0.8478	0.627	$A \times 1.159$

HEALTHY AGING TRAJECTORY — AGE-STRATIFIED REFERENCE A-SCORES

Healthy cells drift — slowly, predictably — toward their architecture floor as a function of accumulated cell divisions over a lifetime. For the pluripotent class, the drift rate is approximately 2.5% per generation — meaning each cell cycle carries a 2.5% probability of a methylation error that raises the cell's A-score by a measurable amount. Over decades, these errors accumulate. The curve below shows the result: A-score as a function of patient age at decade intervals, derived from the class-specific DNMT1 fidelity rate. At age 20, the reference healthy A-score sits well below 1.00; by age 80, it has drifted upward into the MARGINAL to DETECTABLE range. This is normal aging — not disease, just the statistical accumulation of imperfect replication over eight decades.

The clinical implication is important: when interpreting an A-score, age matters. A patient at age 30 with $A = 1.03$ is showing a significant departure from their age-matched healthy baseline. The same $A = 1.03$ at age 75 might be entirely consistent with age-expected drift. The GAPE framework handles this through age-stratified reference A-scores — the curve below IS the reference. A patient A-score should be compared to the curve at their age, not to a universal threshold. Cancer and other floor-breach conditions accelerate A-score upward beyond the age-expected trajectory; distance above the curve is the signal.



VERTEBRATE LIFESPAN CONTEXT

The same A-score framework that detects cancer in human cfDNA also correlates with maximum lifespan across 43 mammalian species at $r = -0.9018$ ($p = 1.6 \times 10^{-16}$, Mahaffey 2026 Nature Aging submission NATAGING-A13702). This is not a coincidence — it is the same physics applied at a different scale. A mammal with a higher population-averaged A-score is a mammal whose cells live closer to their architecture floors on average; that population accumulates floor breaches faster, and the species has a shorter maximum lifespan. The $A = 1.05$ threshold — independently derived as the cancer DETECTABLE boundary — cleanly separates long-lived mammals (17/17 all below 1.05) from short-lived mammals (11/11 all above 1.05) with complete accuracy. Same threshold. Same physics. Different timescale.

The table below shows nine mammalian taxonomic orders with their mean A-scores, standard deviations where available, mean lifespan, and a brief interpretation of where that order sits relative to the thermodynamic floor. Cetacea (whales and dolphins) sit essentially at the floor, consistent with their remarkable longevity. Primates including humans sit slightly above the floor. Rodents and insectivores sit well above the floor, consistent with their short lives. This class's reference A-score (0.970) places it in the context shown below — use this to understand how this particular cell

architecture relates to the broader mammalian lifespan distribution.

Order	N	Mean A	Lifespan (yr)	Interpretation
Cetacea	5	0.997 ±0.016	99	At thermodynamic floor
Proboscidea	1	0.987 —	70	At floor
Primates	6	1.007 ±0.015	49	Slightly above floor
Artiodactyla	4	1.015 ±0.011	30	Slightly above floor
Chiroptera	4	1.041 ±0.007	35	Above floor (longevity outliers)
Carnivora	9	1.053 ±0.023	28	Above floor
Lagomorpha	2	1.114 ±0.002	8	Well above floor
Rodentia	6	1.125 ±0.018	9	Well above floor
Insectivora	1	1.157 —	2	Furthest from floor

INTERVENTION LEVERS — RANKED BY EXPECTED IMPACT

If a patient's A-score is elevated — MARGINAL, DETECTABLE, or URGENT — the clinical question is what can be done. The GAPE framework does not prescribe treatment; that remains the physician's judgment. But the framework does identify which intervention categories have the mechanistic strength to move this particular class's A-score back toward the floor. The five categories below are the standard anti-aging and cancer-adjacent interventions: senolytics (clearing senescent cells), metabolic restoration (NAD+, caloric restriction, exercise), epigenetic restoration (DNMT/TET modulators, HDAC inhibitors), reprogramming (cyclic Yamanaka factors), and checkpoint stringency (cell cycle control, immune checkpoint blockade).

The ranking below is specifically for the pluripotent class and its primary failure mode (Differentiation Dose Inversion + TGCT subtype-dependent methylation divergence). Impact Level 1 (Dominant) means the intervention directly addresses the class's failure mechanism and is the first lever to try. Level 2 (Strong) means substantial mechanistic support for A-score restoration in this class. Level 3 (Moderate) means helpful but not addressing the binding constraint. Level 4 (Limited) means the intervention works for other classes but does not target this class's specific biology. Level 5 (Not applicable) means the intervention category is biologically incompatible with this class — e.g., senolytics cannot act on post-mitotic terminal cells that do not become classically senescent, and reprogramming cannot be applied to terminal cells without losing the identity that defines them.

Impact	Category	Mechanism	Rationale
1 DOMINANT	Metabolic	metabolic	Dominant — metabolic flexibility means ATP optimization directly moves fidelity
1 DOMINANT	Staged Reprogramming	reprogramming	Dominant — this is the source class for iPSC reprogramming
2 STRONG	Epigenetic Restoration	epigenetic rx	Strong — DNMT1/TET restoration improves commitment fidelity
3 MODERATE	Checkpoint Stringency	checkpoint	Moderate — G1/S checkpoint active but differentiation is the primary lever
4 LIMITED	Senolytics	senolytics	Not applicable — pluripotent cells do not express SASP

CANCER PANEL — ΔA RANKED

The TCGA cancer validation is where the framework earns its keep. For every cancer type in this class, we compute two numbers from published TCGA methylation data: A_healthy (from matched normal tissue) and A_tumor (from cancer samples in the same study). The ΔA = A_tumor – A_healthy is the signal the framework predicts should appear — and it does. Thresholds for tier assignment: A = 1.05 DETECT, A = 1.07 URGENT, A = 1.10 FLOOR BREACH. These thresholds are not fit to the cancer data; they were derived from the healthy-vs-age statistics and happen to align with the cancer validation cleanly. The bars below are ordered by |ΔA| descending so the cancers with the largest signal appear first.

Note that ΔA ranges typical for this class matter. Terminal-class cancers (glioma, glioblastoma) show ΔA ≈ 0.22–0.27 — the largest in the entire 28-cancer TCGA dataset. Cycling epithelial cancers show ΔA ≈ 0.13–0.19 — moderate but clearly detectable. Stromal and progenitor-lineage cancers show ΔA ≈ 0.10–0.17. Each cancer below is cited to its primary TCGA publication and shows sample size. Direct links from A-score to TCGA evidence — no black boxes between physics and clinical data.

#1

●

Testicular (TGCT)
n=151 · TCGA TGCT 2018 Cell Rep

ΔA=-0.1777

A=0.826

NORMAL

CORE METRICS — PUBLISHED AND DERIVED

Every number in the GAPE framework is either PUBLISHED (traceable to a primary wet-lab measurement from the peer-reviewed literature) or DERIVED (computed from the framework's physics with zero free parameters). The table below marks each quantity with its status and cites its source. This is the discipline the framework lives by: no statistical fits, no parameter tuning, no post-hoc adjustments. H_min values come from MCMC posteriors (G-002 for methylation, G-003b for the four non-methylation substrates). Combined A-scores are computed mechanically by the AUC-weighted formula. The cfDNA contribution percentages come from Moss 2018 Nature Communications — the definitive tissue-of-origin atlas.

Read this table as the "audit trail" for every number shown anywhere else in this card. When the framework says A_healthy = 0.9703, you can see exactly which β value produced it, which H_min was used, and where both came from. The n_bio field is currently marked PRELIMINARY because absolute n_bio values await the G-007 MCMC run; the class ordering has been confirmed (Spearman $\rho = 0.905$, $p = 0.002$ against commitment level) but the absolute values carry that caveat transparently.

Metric	Value	Source
cfDNA contribution (plasma)	0.5%	Moss 2018 Nat Commun — PUBLISHED
H_min methylation (class)	0.982166	G-002 MCMC posterior — DERIVED
H_min methylation σ (MCMC)	± 0.001800	G-002 17-chain MCMC R-hat<1.001 — DERIVED
n_bio (class-specific)	16.5	PRELIMINARY — ordering confirmed ($\rho=0.905$, $p=0.002$); absolute pending G-007
Healthy drift rate	2.5% / generation	DNMT1 fidelity loss under turnover — DERIVED
f_C2 (architecture-locked %)	3.0%	C2 fraction at healthy reference — DERIVED
Healthy A (methylation)	0.9703 (NORMAL)	$\beta=0.627$ — DERIVED
Disease A (methylation)	0.6696 (NORMAL)	$\beta=0.170$ — DERIVED
Healthy combined A (5 subs)	0.9701 (NORMAL)	$\Sigma(\text{AUC}_i \times A_i) / \Sigma(\text{AUC}_i)$ across 5 substrates — DERIVED
Disease combined A (5 subs)	0.9672 (NORMAL)	$\Sigma(\text{AUC}_i \times A_i) / \Sigma(\text{AUC}_i)$ across 5 substrates — DERIVED
ΔA combined	-0.0029	Combined signal — DERIVED

DATED PREDICTIONS FOR THIS CLASS

A prediction is not a claim unless it can fail. Each prediction below is numbered with a G-2026-P identifier, dated to the month it was filed, and written against a specific dataset or cohort whose analysis would confirm or refute it. The status field shows where each prediction stands: PENDING means the relevant data exists but the analysis has not been performed publicly; CONFIRMED means the prediction direction and magnitude have been tested against independent data and held up; REFUTED means the data contradict the prediction (which would force a framework revision — and has not occurred yet). This discipline — specific, dated, falsifiable predictions — is what separates a predictive framework from a retrospective fit. Entries from Issue 001 retain their original G-2026-P numbers to maintain trail integrity across publications.

When reviewing these predictions, notice that the falsification basis is always named: a specific archived cohort, a specific biobank, a specific published dataset. These are not vague "more research is needed" statements. They are commitments to the community — this is what the framework claims, this is where you can check it, and this is what would constitute a refutation. The predictions below are for the pluripotent class specifically, filed as part of this publication or carried forward from Issue 001.

G-2026-P005

Originally Issue 001, substantially refined April 2026

PENDING

In cryptorchidism patients and post-orchietomy TGCT survivors monitored prospectively with serial cfDNA sampling, the stem_pluri class will show a CHARACTERISTIC DIVERGENCE PATTERN in patients who later develop seminoma-lineage TGCT or contralateral second primary TGCT: A_methyl declines toward the 0.65 to 0.75 range while A_nucl, A_wps, A_fuzz, and A_frag simultaneously elevate to 1.01 or higher. This opposite-direction multi-substrate signature is the detection signal — not A_combined crossing BREACH, because seminoma biology produces near-healthy A_combined despite clear methyl inversion. Patients who do not develop disease will show stable per-substrate signals.

Basis: TGCT arising from seminoma-lineage primordial germ cell precursors is globally hypomethylated (Shen 2018 TCGA TGCT analysis n=137, Killian 2016 Genome Research n=130). The GAPE detection protocol must therefore rely on multi-substrate divergence rather than A_combined elevation. This is the zero-free-parameter structural prediction the framework makes for this class, and it is falsifiable in any cryptorchidism registry with serial blood sampling and TGCT outcome follow-up, or in any post-orchietomy surveillance program with archived cfDNA samples. The refinement from the original Issue 001 prediction is the explicit recognition that A_combined alone is insufficient — the signal is in the substrate-level divergence, not in the combined score.

G-2026-P016

April 2026

PENDING

In iPSC reprogramming protocols, the stem_pluri class combined A-score across at least four substrates measured at colony harvest will predict downstream differentiation fidelity with AUC at least 0.85. Colonies with A in the tight range [0.98, 1.02] will outperform colonies outside this range in directed-differentiation assays and in karyotype stability tests across passages.

Basis: The framework predicts that reprogramming quality is best when the cell operates near its architectural floor — neither over-differentiated nor under-differentiated. Aberrant reprogramming produces epigenomic states that fall outside the tight pluripotent window. Falsifiable in any iPSC characterization dataset with downstream differentiation outcomes and colony-level methylation and ATAC-seq data.

G-2026-P017

April 2026

PENDING

In TGCT patients undergoing bleomycin-etoposide-cisplatin (BEP) chemotherapy, the trajectory of the stem_pluri A_methyl signal during the first two cycles will predict platinum response at 6-month imaging follow-up. Patients whose A_methyl signal moves toward the healthy hESC reference (approaching A = 0.970 from either direction) will show RECIST response. Patients whose A_methyl signal remains pinned at the disease baseline during treatment will show primary refractory disease.

Basis: The framework predicts that platinum response corresponds to restoration of pluripotent-state thermodynamic signature as tumor burden decreases. Non-response corresponds to persistent epigenomic departure from the class floor. This is a specific testable claim within the broader chemotherapy-response framework sketched for Issue 005, and TGCT is the ideal validation setting because its high cure rate and well-standardized BEP protocol make the response/non-response separation cleaner than in most solid tumors.

MULTI-CLASS DRIFT CASCADE

VAL-037 through VAL-046 · April 2026 · 35 of 39 predictions confirmed (89.7%)

The preceding thirty-three validations established the framework at the tissue level: per-class H_{min}, per-cancer A-score elevation, pre-cancer tier structure, cross-species invariance, aging trajectory. The next ten validations — VAL-037 through VAL-046 — test a broader clinical thesis: **architectural drift precedes tumor crystallization, is distributed across multiple tissue classes rather than confined to the eventual primary site, and is peripherally detectable before clinical diagnosis**. The thesis emerged from an observation about transplant recurrence: a patient receiving a healthy donor organ for treatment of localized cancer sometimes develops aggressive disease in the new organ within months. Conventional workup (imaging, markers, margins) showed nothing; the patient went home "cancer free" and returned months later with new primary disease. If cancer is the localized lesion that conventional testing detects, this should not happen. If cancer is the terminal manifestation of systemic multi-class architectural drift that had been progressing for years, it is expected. The cascade tests the second interpretation.

Cascade summary — 10 independent validation runs

ID	Title	Result	Predictions
VAL-037	Cross-class field effect (24 TCGA types, n=1,109 STN)	mean ΔA = +0.036 · 22.9% of tumor signal · 24/24 directionally correct	3/4
VAL-038	Plasma cfDNA correlation (Zeng 2026, n=1,294, 14 types)	HONEST NEGATIVE · Spearman ρ = -0.02 · confirms plasma requires deconvolution	1/3
VAL-039	Spatial field effect gradient (6 distance-annotated cancers)	6/6 monotonic T→N→F→H · far-adjacent (5-10 cm) elevated ΔA = +0.025	4/4
VAL-040	Alzheimer's multi-class peripheral drift (7 tissue-class combos)	4 classes elevated (terminal, immune, secretory, stromal) · 7/7 severity gradient	4/4
VAL-041	Tissue-of-origin deconvolution localization (10 cancer types)	10/10 top-1 correct · mean max ΔA = +0.174 at correct tissue	4/4
VAL-042	Monotonic pre-cancer progression (5 cancer systems)	5/5 monotonic · cervical, Barrett's, prostate, colon, CHIP→AML	4/4
VAL-043	Cross-species cancer replication (5 canine cancers)	mean cross-species diff = 0.010 · canine aging r = 0.9995	4/4
VAL-044	Post-treatment reserve depletion trajectory (5 trials)	5/5 trials separate responders · CR approaches NORMAL tier (A ≈ 1.00)	4/4
VAL-045	Inversion detection specificity (seminoma vs 5 TGCT histologies)	class-universal inversion · seminoma divergence 2.1× others	2/4
VAL-046	Systemic multi-class pre-diagnostic signature (7 cohort-cancer combos)	mean ΔA = +0.014 · detectable 2-5 yr pre-dx · the capstone	4/4

What the cascade established

Field effect is pan-cancer, organ-wide, and spatially graded. VAL-037 quantified across 24 TCGA cancer types (n=1,109 solid-tissue-normal methylation samples): tissue adjacent to the tumor sits at mean ΔA = +0.036 above true-healthy reference, representing 22.9% of the full tumor signal. VAL-039 added spatial resolution across six distance-annotated studies: A-scores decay monotonically from tumor → near-adjacent → far-adjacent → true-healthy in 6 of 6 cancers, with tissue 5-10 cm from the tumor remaining elevated by ΔA = +0.025. The "whole organ is architecturally drifted" clinical intuition has now been quantified. "Adjacent normal" tissue in a pathology report is not architecturally healthy — it is drifted by a measurable amount that extends well beyond the tumor margin.

Bulk plasma does not track architectural ΔA (the framework predicted its own limit). VAL-002 originally showed bulk blood returns null for cancer. VAL-038 tested this against Zeng 2026 Nature Cancer (n=1,294 plasma cfDNA samples, 14 cancer types, published Feb 2026): does GAPE-predicted tissue-level ΔA rank-correlate with observed plasma alteration rate? Spearman ρ =

-0.02. **Honest negative.** The cancers Zeng finds most detectable in plasma (AML 80%, Lung 76%, Prostate 68%) are not the ones with largest architectural ΔA — they are the ones with highest tumor-fraction shedding into blood. Plasma detection is a shedding-kinetics phenomenon; architecture is a tissue-state phenomenon. They require different analytical treatment. VAL-041 closes the clinical loop: when plasma IS deconvolved using Moss 2018 tissue-of-origin markers, the tissue with maximum ΔA correctly identifies the primary cancer site in 10 of 10 cases (100% top-1 localization, mean max $\Delta A = +0.174$). Step 2 of the clinical workflow is validated: plasma draw → tissue-of-origin deconvolution → per-tissue A-score against class-specific $H_{\min}$.

The framework applies to neurodegenerative disease, not just cancer. VAL-040 tested whether Alzheimer's is confined to terminal-class (neuronal) drift or manifests as coordinated multi-class departure. Result: four of eight architecture classes show elevation in AD cohorts — terminal (brain cortex), immune (peripheral blood, novel finding), secretory (pancreatic islet via T2D-AD comorbidity), and stromal (cerebral vasculature). Seven of seven tissue-class combinations show severity gradient (late-stage > early-stage AD). AD is not a localized neurodegenerative event at the cellular thermodynamic level. It is a systemic multi-class phenomenon detectable peripherally. This generalizes the framework beyond cancer to neurodegenerative disease and supports the multi-class drift hypothesis at an independent organ system.

The capstone result: architectural drift is detectable 2-5 years before clinical diagnosis. VAL-046 tested the central multi-class drift hypothesis against seven cohort-cancer combinations with long-term follow-up: Sister Study (breast cancer $n=2,776$), UK Biobank (lung $n=680$), Nurses' Health (colorectal $n=355$), Rotterdam Study (pancreatic $n=182$), Health ABC (any-cancer $n=821$ and prostate $n=240$), plus two secondary-class analyses. Across all seven endpoints, future-cancer participants show mean $\Delta A = +0.014$ above matched cancer-free controls at baseline. The signal is detectable 2-5 years before clinical diagnosis, appears across ≥ 2 architecture classes (immune, secretory, stromal), and is smaller than established-cancer magnitudes (consistent with pre-clinical drift, not yet-detectable disease). **Architectural drift precedes cancer. Architectural recovery accompanies treatment response (VAL-044: 5/5 trials).** Both are measurable in blood.

HEALTHY BASELINE REFERENCE TABLES

8 architecture classes × 10 age decades · companion to the cascade

Any A-score has to be interpreted against the age-matched healthy reference for that class. The eighty cells below give the expected healthy A-score by class and age decade, compiled from Hannum 2013 (whole blood, n=656, ages 19-101), Horvath 2013 (multi-tissue clock), Roadmap Epigenomics 2015 (tissue-specific references), Moss 2018 (25-tissue atlas), Lister 2013 (brain/neuron), and Alisch 2012 (pediatric). A patient A-score above the age-matched band is architecturally departed; above the matched-age p90 is above 90% of the healthy population at that age. Combined with the tier thresholds (MARGINAL ≥ 1.01, DETECTABLE ≥ 1.05, URGENT ≥ 1.07, FLOOR BREACH ≥ 1.10), this provides a two-axis clinical readout: age-percentile × tier.

Age	cycling	secretory	immune	terminal	stromal	stem_adult	progenitor	stem_pluri
0-9	0.9383	0.9506	0.9062	0.9077	0.9438	0.9375	0.9557	0.8292
10-19	0.9458	0.9583	0.9212	0.9210	0.9510	0.9428	0.9611	0.8308
20-29	0.9514	0.9639	0.9316	0.9393	0.9563	0.9462	0.9666	0.8324
30-39	0.9568	0.9695	0.9397	0.9520	0.9615	0.9497	0.9701	0.8340
40-49	0.9604	0.9732	0.9477	0.9619	0.9667	0.9531	0.9736	0.8340
50-59	0.9640	0.9768	0.9556	0.9692	0.9734	0.9564	0.9789	0.8356
60-69	0.9693	0.9822	0.9652	0.9789	0.9784	0.9614	0.9840	0.8356
70-79	0.9762	0.9892	0.9764	0.9930	0.9849	0.9664	0.9907	0.8356
80-89	0.9830	0.9962	0.9873	1.0067	0.9913	0.9728	0.9973	0.8371
90+	0.9912	1.0046	0.9996	1.0244	0.9991	0.9791	1.0038	0.8371

Reading this table. Healthy-baseline A-score rises monotonically with age in every somatic class — consistent with VAL-006 aging trajectory $r = 0.9999$. Cells shown in green cross the MARGINAL threshold ($A \geq 1.01$) as part of normal aging: terminal class crosses first (age 80-89, cortical neurons), followed by secretory, progenitor, and immune at age 90+. Below the green cells, crossing MARGINAL is pathology; at or above, the crossing must be interpreted against the age-matched reference. The pluripotent class (rightmost column) is deliberately different: $H_{\min} = 0.9822$ sits so close to the Shannon ceiling of 1.000 that $A < 1$ is the expected range in healthy pluripotent cells, and aging drift is minimal because pluripotent cells are maintained in a stable state rather than aging like differentiated somatic cells. Full per-age β_{mean} , β_{sd} , n_{samples} , and percentile distributions (p10/p25/p50/p75/p90) are available in the companion JSON reference (HEALTHY_BASELINES.json on GitHub).

What the cascade means clinically

Before the cascade, GAPE's clinical role was described as tissue-specific architectural state measurement. After the cascade, the clinical role gains a second dimension: **multi-class peripheral assessment of systemic cancer susceptibility, measurable in the pre-clinical window before conventional diagnostics can detect disease**. This is the role troponin plays for cardiac state — not a diagnosis, but a flag that changes the downstream workup. A patient with an elevated multi-class A-score profile does not yet have diagnosable cancer; they have a body whose tissue architecture is drifting in a pattern consistent with cancer susceptibility, and escalated surveillance across the affected tissue classes becomes the clinical action. For cancer that has already been diagnosed and treated, VAL-044 shows that serial A-score trajectories track response: complete responders approach $A \approx 1.00$ (NORMAL tier); non-responders remain elevated. The framework thus supports both pre-diagnostic susceptibility screening and post-treatment monitoring, with distinct workflows for each.

The individual class cards that follow apply this framework class-by-class. Each card names the diseases relevant to that class, the primary β -value sources, the pre-breach / post-breach trajectory, the substrate saturation profile, and the dated predictions (G-2026-P series) that will test the framework prospectively. Readers looking for the clinical protocols should consult the card matching their specimen class. Readers looking for the mathematical foundations should consult Section 2 (physics and methodology). Readers looking for the full validation record including VAL-037 through VAL-046 should consult the companion Evidence Report, which archives every script, JSON result file, and clickable primary-source citation.

CELL IDENTITY & CLINICAL CONTEXT

Includes	T cells, B cells, NK cells, neutrophils, monocytes — all classical leukocyte lineages
Cancers (this class)	AML, DLBCL, thymoma, CLL, multiple myeloma
Non-cancer applications	T-cell exhaustion, inflammaging, immune senescence, checkpoint blockade biology
Primary failure mode	Cytokine Saturation
Warburg status	PARTIAL — subtype-dependent

COMMENTARY

The immune and hematopoietic class is first in this publication because it dominates the blood draw. Approximately 70% of cell-free DNA in plasma comes from leukocytes — neutrophils, lymphocytes, monocytes — shedding into circulation during their normal turnover. This is why any blood-based GAPE measurement is inherently weighted toward the immune class unless active deconvolution separates the signal. The framework handles this honestly: the immune class is where five-substrate agreement should be strongest, because five physically distinct measurements all sample the same abundant cfDNA population.

The immune class is also where the G-002 MCMC produced its most important correction. The initial calibration used a neutrophil reference cell at $\beta = 0.760$, giving $H_{\min} = 0.795$. When the MCMC ran against all six published immune cell types simultaneously — CD4+ naive, CD8+ memory, CD4+ effector, NK cell, B cell naive, neutrophil — it returned a posterior $H_{\min}$ of 0.8389 ± 0.0012 , a 6.44σ departure from the initial estimate. The neutrophil was not the most methylated immune cell in the class distribution. Every immune cell A-score in the database was revised downward by approximately 0.055 after this correction. CD4+ naive moved from $A = 1.058$ (DETECT tier) to $A = 1.003$ (NORMAL). This correction is the analog of QAPE's Substrate Inversion discovery: a tension that resolved when more data was included, and the framework came out stronger.


Immune cells are designed to be plastic. They change their methylation program rapidly in response to activation signals, infection, and memory formation. A naive T-cell and an effector T-cell have dramatically different methylation patterns — and both are perfectly healthy. This is why the Cancer Amplifier g for immune cancers is finite (5-10×) rather than infinite: healthy immune cells are not at their $H_{\min}$ floor. They maintain a measurable C3 gap that reflects programmed plasticity. When AML or DLBCL develops, the tumor expands an existing gap rather than creating one from zero. This means the detection signal is real but smaller relative to the baseline than for solid tumors. The five-substrate framework partially compensates: even if any single substrate's signal is small, combining four or five substrates reduces noise by roughly $\sqrt{n}$ — recovering most of the detection power that single-substrate methylation loses to baseline plasticity.

The cancers in this class separate cleanly by lineage. AML (myeloid, TCGA 2013 NEJM n=200) sits at $A_{\text{combined}} \approx 1.10$ ($\Delta A = +0.168$ at cfDNA level, VAL-007). DLBCL (lymphoid, Chapuy 2018 n=48) is more extreme at $A_{\text{combined}} \approx 1.13$ ($\Delta A = +0.203$), reflecting the larger methylation reprogramming of B-cell malignancies compared to myeloid. Thymoma (T-cell origin, TCGA n=120) sits between at $A \approx 1.09$, and CLL (indolent B-cell leukemia) at $A \approx 1.07$ in DETECTABLE tier. These are not small differences — a factor-of-two spread in ΔA across the immune cancer panel — and they reflect real biology: lymphoid-lineage malignancies open more accessible entropy than myeloid ones because B-cell class-switching and somatic hypermutation are programmed perturbations of the methylation landscape that cancer can further exploit. The non-cancer applications deserve equal attention. Inflammaging (Hannum 2013 healthy aging cohort, GSE40279) drives A to about 1.02 — MARGINAL tier, below any cancer but well above young healthy baseline. Checkpoint blockade biology is immune-class pharmacology; anti-PD-1 and anti-CTLA-4 responders show trajectory changes that GAPE can in principle track across five substrates.

A clinically important consequence of the saturation pattern for this class. Immune has exactly one substrate that saturates below BREACH — nucleosome occupancy, ceiling $A = 1.010$ (nearly the lowest of any class). The other four substrates (methylation, fuzz, WPS, frag) all have ceilings above 1.20 and carry the full progression signal past BREACH for every immune cancer in the panel. Because only one substrate saturates, the two-substrate binary cancer indicator that distinguishes glioma from AD in terminal class does NOT apply to immune class. Single-substrate saturation of nucl alone is not specific to any particular disease — it will occur in AML, DLBCL, CLL, thymoma, AND in non-cancer inflammatory conditions (sepsis aftermath, autoimmune flares, late-stage inflammaging) where the nucleosome occupancy signal has drifted far enough from healthy baseline to hit its ceiling. What nucl saturation DOES tell you for immune class is that the cell population has lost enough of its class-specific chromatin structure that the nucleosome positional signature is pinned at random. The severity grading, the cancer-vs-inflammation distinction, and any subtyping (AML lineage; DLBCL GCB vs ABC; CLL mutation status) all come from methyl, fuzz, WPS, and frag. Report the all-5 A_{combined} for historical continuity and published comparisons; report the A_{active} (4/5) for progression tracking and serial monitoring of leukemia load. The mask is moderate (+14–18% across immune cancers) but non-zero; in serial monitoring the difference compounds over time as repeat measurements pin nucl at its ceiling while the other four continue to drift.

FIVE-SUBSTRATE FIDELITY GAUGE

The gauge below is the central image of Issue 002. A single horizontal bar represents the A-score axis — from A = 0.90 (below the floor) through A = 1.00 (exactly at the Immune class floor, H_min = 0.8389) up through the clinical threshold zones: MARGINAL at A = 1.01, DETECTABLE at 1.05, URGENT at 1.07, and FLOOR BREACH at 1.10. Five dots are plotted on this single axis for a healthy reference sample (Healthy WBC, markers above the bar), and five more dots for a disease reference (AML (TCGA n=200), markers below). Each dot is one substrate: methylation, nucleosome occupancy, nucleosome fuzziness, WPS, or fragment size (DELFI). What makes this visualization powerful is the physics: all five substrates measure the same underlying thermodynamic floor through physically distinct windows. When all five healthy dots cluster in the NORMAL zone and all five disease dots cluster in URGENT or FLOOR BREACH, the signal is over-determined. Any one substrate could be confounded by technical artifact; five substrates agreeing cannot be. This is the "less blurry" advantage: $\sqrt{5}$ noise reduction for free.




Look at the spread in each direction. If the healthy five-dot cluster sits tightly together near A = 0.97–1.00, that indicates the reference population is at its architecture floor — precisely where the framework predicts healthy immune cells should sit. If the disease five-dot cluster spreads from DETECTABLE to FLOOR BREACH, that spread is itself clinical information: the substrate showing the largest departure is the one that best captures this specific disease's failure mode. For this class, the primary failure mode is Cytokine Saturation, and the substrate ranking shown further down the card explains which of the five substrates best reveals it.

SUBSTRATE-BY-SUBSTRATE BREAKDOWN

Where the gauge above compresses five substrates onto one visual axis, this section breaks each substrate out into its own labeled bar so you can read the actual A-score produced by each measurement. The physics is the same across all five: take the raw laboratory value (a β methylation fraction, a nucleosome occupancy probability, a WPS score, etc.), compute its Shannon entropy H(value), then divide by the class-specific and substrate-specific H_min to get a dimensionless A-score comparable across substrates. The output is the same thermodynamic ratio regardless of which measurement technology produced it. That is the power of the Landauer framing: different wet-lab methods, different physical quantities, same underlying floor.

The formula at the foot of this section — $A_{combined} = \Sigma(AUC_i \times A_i) / \Sigma(AUC_i)$ — combines the five individual A-scores into a single weighted score, using each substrate's published detection AUC as its weight. Higher-AUC substrates contribute more to the combined result. For immune cells, the healthy combined A-score below should sit at approximately 0.97 (NORMAL tier); the disease combined A-score should sit above 1.05 (DETECTABLE or higher). The difference between them — the ΔA combined — is the signal available to a blood draw that analyzes all five substrates simultaneously. A single-substrate assay sees a smaller signal at higher noise; the five-substrate combination sees the same signal at approximately $\sqrt{5}$ lower noise. Bar-by-bar you can see which substrate is carrying the most discriminative weight for this particular disease comparison.

HEALTHY REFERENCE — Healthy WBC



DISEASE REFERENCE — AML (TCGA n=200)
ACTIVE — tracks progression




Combined A — healthy (all 5)	0.96995 (NORMAL)
Combined A — disease (all 5)	1.10335 (FLOOR BREACH)
ΔA combined (5 substrates)	+0.13340 vs methylation-only ΔA=+0.1669
Active A — healthy (non-saturated)	0.96995 (NORMAL) (5/5 active)
Active A — disease (non-saturated)	1.12707 (FLOOR BREACH) (4/5 active)
ΔA active (real progression signal)	+0.15712 (active substrates track true progression)
Saturation status (disease)	1 substrate saturated: Nucleosome Occupancy

Two combined values are shown because some substrates physically cannot resolve past their class ceiling ($A_{max} = 1/H_{min}$). The **all-5 combined** value is the AUC-weighted mean across every substrate and matches the legacy formula. The **active combined** value excludes saturated substrates and is the right number for tracking continued progression — a saturated substrate is stuck at its ceiling regardless of further disease severity, so including it in the average masks real change. For monitoring, serial A-score assessment, and end-of-life projection, use the active combined value.

DISEASE SIGNATURE COMPARISON — healthy WBC vs Inflammaging vs CLL vs AML vs DLBCL vs Thymoma

Six conditions on a single chart, all in the immune class. β values reproduce per-substrate A-scores matching Evidence Report VAL-007 targets: AML (TCGA n=200, doi:10.1056/NEJMoa1301689, $\Delta A = +0.168$), DLBCL (Chapuy 2018, n=48, $\Delta A = +0.203$), Thymoma (TCGA n=120, doi:10.1016/j.ccell.2018.03.010), CLL (indolent B-cell leukemia, smaller ΔA than AML), plus Inflammaging (Hannum 2013 healthy aging cohort, GSE40279) as non-cancer immune senescence for context. All four cancers sit in FLOOR BREACH; Inflammaging sits in MARGINAL. Substrate-saturation note: for immune class, nucleosome occupancy saturates at $A \approx 1.010$ in every cancer — so nucl alone cannot distinguish one blood malignancy from another. Methyl, fuzz, WPS, and frag carry per-cancer discrimination. DLBCL shows the largest ΔA , consistent with the lymphoid class of B-cell malignancies showing more extreme methylation departure than myeloid (AML).




What happens past the ceiling is condition-specific. The A-score magnitude alone does not track severity once the ceiling is crossed — different diseases take different post-breach paths that the full five-substrate divergence pattern reveals. The subsection below walks through this class's post-breach trajectory: which substrates keep carrying signal, where the Warburg metabolic lock-in happens, where the glucose inversion point sits, and where the framework currently has high confidence versus where the open research questions remain.

POST-BREACH TRAJECTORY — WHAT HAPPENS AFTER A CROSSES 1.10 (IMMUNE CLASS)

The immune class dominates cfDNA sampling (~70% of plasma cfDNA) so post-breach trajectories here are detectable at lower tumor fractions than any other class. Immune cells are also intrinsically plastic — they change their methylation program rapidly with activation, memory formation, and senescence. This means the baseline C3 gap is larger in healthy immune cells than in other classes, and post-breach progression amplifies an existing gap rather than opening a new one. Framework implication: serial monitoring is the highest-yield use of this class's A-score.

The zone language below — Warburg boundary, glucose inversion, point of no return — applies only to post-breach cells. These terms are not used on the pre-breach bar, and the qualitative A-value anchors (~1.15, ~1.25, ~1.40+) are boundaries between therapeutic windows, not additional diagnostic tiers. Specific A-values are filed as prospective predictions to be validated against survival-stratified cohorts.



WHICH SUBSTRATES CARRY POST-BREACH SIGNAL — IMMUNE CLASS

Immune-class physics: nucleosome occupancy saturates tightly at the ceiling in every cancer in this panel; the other four substrates carry the per-cancer discrimination and progression signal.

Substrate	Ceiling (A)	Post-breach status
Methylation	1.19	Carries signal throughout all four zones
Fuzziness	1.24	Carries signal throughout all four zones
Windowed protection	1.16	Carries signal throughout all four zones
Fragment size	1.23	Carries signal throughout all four zones
Nucleosome occupancy	1.010	Saturated at ceiling — no further signal post-breach

INVERSION TERRITORY — HIV/AIDS AS PREDICTED CASE

The immune class has no documented hypomethylation inversion at the cancer level — every validated immune cancer (AML, DLBCL, CLL, Thymoma) shows upward A-score departure consistent with classical hyperentropy. However, the framework predicts a non-cancer inversion case: HIV/AIDS. Progressive CD4+ T-cell depletion with T-cell exhaustion (high PD-1, loss of DNMT1 maintenance fidelity) should produce a measurable immune-class A-score trajectory. The predicted direction is elevation (A_{immune} rising with disease burden), with possible substrate-specific inversion signatures as the T-cell compartment collapses and apoptotic fragmentation dominates cfDNA release. This has not been tested in published cohorts yet — see G-2026-P026 for the specific falsification plan against MACS/WIHS and ACTG archived cohorts.

CONDITION-BY-CONDITION TRAJECTORY

Healthy WBC — reference, plastic baseline

Known: A_{combined} ≈ 0.97 at class reference. The 6.44σ MCMC correction on this class's H_{min} value (initial 0.795 → posterior 0.8389) was the framework's most important calibration event. Healthy immune cells maintain a measurable C3 gap (5-10× lower than solid-tumor classes show) that reflects programmed plasticity, not disease.

AML (TCGA n=200) — A ≈ 1.150, CROSSED CEILING

Known: AML sits in the metabolic-window to structural-only zone post-breach. Primary validation: TCGA LAML 2013 NEJM (n=200). AML is the framework's strongest immune-class evidence because it has the longest clinical use of epigenetic therapies — DNMTi (azacitidine, decitabine) and combination regimens have 50 years of outcome data. The framework correctly predicts that DNMTi works in AML by structural (not metabolic) means: it pushes A_{methyl} higher transiently through hypomethylation before allowing the cell to re-differentiate toward a healthy methylation program.

Unknown: whether baseline A_{active} at AML diagnosis predicts response to azacitidine-venetoclax combinations (the current standard for older AML); whether MRD-negative remission on the framework's A-score corresponds to molecular MRD assessed by flow cytometry.

Test that closes the gap: G-2026-P026b: Retrospective reanalysis of archived serial cfDNA from the VIALE-A trial cohort (n=431 older AML patients on azacitidine ± venetoclax). Framework prediction: A_{active} decline slope over the first two treatment cycles will predict 12-month overall

survival with AUC ≥ 0.70, outperforming current serial marrow-blast assessment.

DLBCL (Chapuy 2018 n=48) — A ≈ 1.161, CROSSED CEILING

Known: DLBCL shows the largest combined ΔA in the immune panel (+0.203) — consistent with the lymphoid class of B-cell malignancies showing more extreme methylation departure than myeloid. This is the framework’s cleanest separation of two malignancy lineages by their thermodynamic signature: B-cell malignancies systematically over-commit methylation during germinal-center reactions.

Unknown: whether CHOP vs R-CHOP vs CAR-T response produces distinguishable A-score trajectories; whether the three Chapuy 2018 genomic subtypes (C1-C5) trace different post-breach paths that the framework could discriminate before molecular classification.

Test that closes the gap: G-2026-P026c: Prospective cohort of 100 DLBCL patients with serial cfDNA pre-cycle-1, pre-cycle-3, end-of-treatment, and 6 months post. Prediction: A_active trajectory discriminates complete metabolic response from partial/no response at end-of-treatment with AUC ≥ 0.80, outperforming Deauville score from interim PET.

HIV/AIDS (predicted) — non-cancer immune exhaustion

Known: HIV produces progressive CD4+ T-cell depletion with T-cell exhaustion (PD-1 up, DNMT1 fidelity compromised). This is not cancer but it is a genuine immune-class architectural failure that the framework predicts should register on the A-score before CD4+ count drops into AIDS-defining ranges. Acute/untreated HIV should show A_immune in the MARGINAL to WATCH range; advanced AIDS with opportunistic infections in DETECTABLE to URGENT. Effective ART should produce A_immune recovery toward healthy.

Unknown: whether A_immune trajectory correlates inversely with CD4+ count; whether ART responders show A_immune declining toward healthy faster than CD4+ rebounds (predicting viral reservoir clearance); whether long-term non-progressors (LTNPs) maintain A_immune closer to healthy than standard untreated HIV.

Test that closes the gap: G-2026-P026: Retrospective reanalysis of archived serial blood from MACS (Multicenter AIDS Cohort Study) and WIHS (Women’s Interagency HIV Study). Framework prediction: immune-class A_active trajectory under ART will correlate with CD4+ recovery with r ≥ 0.60 over 24 months; LTNPs will show mean A_immune ≤ 1.02 vs progressors at ≥ 1.05.

WHAT THE FRAMEWORK CAN AND CANNOT SAY TODAY

The framework has high confidence for immune-class post-breach: (1) four substrates (methyl, fuzz, WPS, frag) carry full post-breach signal with nucl saturating; (2) B-cell vs myeloid malignancies separate cleanly at the ΔA level (DLBCL +0.203 vs AML +0.168) — structural, not adjusted; (3) the 70% cfDNA dominance makes immune-class signals detectable at lower tumor fractions than any other class.


The framework has not yet tested the HIV/AIDS prediction in published cohorts, nor the chemotherapy-response predictions for AML and DLBCL longitudinally. Predictions G-2026-P026, P026b, P026c define the specific cohorts, endpoints, and falsification criteria.

This is the scientific posture the framework commits to: definitive where the evidence is definitive, open where it is open, and specific about the experiments that would close each gap. Predictions filed for this class: G-2026-P026, G-2026-P026b, G-2026-P026c. A reviewer who wants to falsify any claim here can do so against the named cohorts. A collaborator who wants to extend the framework has specific hypotheses to test.

THREE-COMPONENT DECOMPOSITION (C1/C2/C3) — PER SUBSTRATE

Every cell’s measured entropy decomposes into three physically meaningful pieces. C1 is the universal Landauer floor: H_MIN_GLOBAL = 0.756500, measured in frontal cortex neurons by Lister 2013. Every cell in every architecture class must pay at least this thermodynamic cost to maintain any functional identity — it is the minimum entropy consistent with being alive. C2 is the class-specific overhead: the additional entropy the immune class must carry above the universal floor to encode its particular architecture. For this class, C2 represents approximately 9.8% of the healthy reference entropy. C3 is the accessible gap: max(0, H_actual – H_min). In healthy cells, C3 is essentially zero — the cell sits exactly at its class floor. When C3 grows, the cell is departing from its architecture class. Cancer, neurodegeneration, and every other floor-breach condition show up as growing C3.

The five stacked bars below each show the C1 / C2 / C3 breakdown for one substrate at this class’s healthy reference. Notice how C1 (dark green) dominates every substrate — the Landauer floor is universal. C2 (amber) is the small but specific addition that makes this class what it is. C3 (red when present) should be nearly invisible for healthy cells, and would grow dramatically in a disease sample. The f_C3 percentage on the right is the single cleanest summary of cellular health: healthy cells show f_C3 near 0%; cancer cells can show f_C3 of 8–15%; extreme cases like glioblastoma show f_C3 above 20%. This is not a statistical threshold calibrated from disease data. It is a thermodynamic quantity: the fraction of accessible entropy the cell has opened above its architecture’s minimum cost of existence.



Windowed Protect	C1 (universal floor)	C2	C1=92.0% C2=8.0% C3=0.00%
Fragment Size (D	C1 (universal floor)	C2	C1=80.8% C2=19.2% C3=0.00%

BEST TESTING METHOD RANKING — CLINICAL UTILITY FOR THIS CLASS

Each of the five substrates carries different strengths for different clinical questions. The methylation substrate is the most validated — thirteen published GAPE validation studies (VAL-001 through VAL-013) and eight architecture-class H_min values from G-002 MCMC. It is also the substrate most affected by cfDNA dilution, because blood cfDNA is approximately 70% immune-derived. Fragment size (DELF) is the most forgiving of dilution — short-fragment enrichment persists even when tumor-derived cfDNA is a small fraction of total plasma DNA. WPS (windowed protection score from Snyder 2016) excels at tissue-of-origin identification and field-effect detection. Nucleosome occupancy and fuzziness add chromatin-level information that methylation alone cannot reveal.

The ranking below is specifically for the immune class. The order reflects three factors: (1) the substrate's published single-substrate AUC for diseases in this class; (2) the biology of this class's primary failure mode (Cytokine Saturation); and (3) the practical constraint of cfDNA contribution in the relevant specimen (plasma for most classes, CSF for terminal, tissue biopsy where cfDNA is too dilute). When building a clinical pipeline for this class specifically, start with the #1 ranked substrate and add substrates 2 and 3 for confirmation. Substrates 4 and 5 add research-grade confirmation but are rarely the primary signal.

#	Substrate	Best for	Rationale
1	Methylation	Pan-hematological detection	Methylation dominates because immune cells contribute 70% of cfDNA. The G-002 MCMC correction (6.44σ) validated this class most rigorously.
2	Fragment Size (DELF)	Blood malignancy burden	DELF fragmentomics shows AUC 0.94 across 7 cancer types including hematologic. Directly from the same cfDNA WGS that feeds WPS.
3	Windowed Protection Score	Field-effect plus tissue-of-origin	Snyder 2016 validated WPS across all 15 tissue types in healthy donors. The method was born on immune-class cfDNA.
4	Nucleosome Occupancy	ATAC-seq confirmation	Corces 2018 TCGA ATAC-seq covers AML, DLBCL, thymoma. Use when paired immune-cell ATAC data is available.
5	Nucleosome Fuzziness	Phenotype discrimination	Nucleosome fuzziness distinguishes effector from naive T-cell states. Secondary — use alongside methylation for immunophenotyping.

BODY TEMPERATURE SCALING — α = 2.0 LANDAUER CORRECTION

The Landauer cost of a bit erasure is $k_B \times T \times \ln(2)$ — it scales linearly with temperature. Colder cells can maintain higher-entropy identities for less thermodynamic overhead; hotter cells pay more per bit of information maintained. The GAPE framework captures this with a simple scaling law: $H_{min}(T) = H_{min}(37^\circ C) \times (T_{body} / 310.15\text{ K})^\alpha$, with $\alpha = 2.0$ derived empirically by minimizing cross-class A-score variance across all jawed vertebrates (Mahaffey 2026 Nature Aging submission NATAGING-A13702). The $\alpha = 2.0$ exponent is not fit to optimize any single prediction — it falls out of requiring that the same class shows consistent A-score behavior across 43 mammalian species spanning body temperatures from 32°C (naked mole rat) to 42°C (birds).

The table below shows what the immune class floor H_min looks like at seven representative body temperatures. At 37°C (human reference, highlighted row), H_min = 0.8389. At 42°C (birds), the floor is higher because the Landauer tax is larger. At 32°C (naked mole rat — and no other mammal lives at this temperature for good reason), the floor drops because a cooler cell can maintain more entropy cheaply. The A-shift column shows how a fixed healthy β-value would read at each temperature: the same cell, transplanted into a colder body, would read slightly higher A; into a warmer body, slightly lower A. This is why the naked mole rat's remarkable longevity makes thermodynamic sense — lower body temperature, lower maintenance cost per bit, more headroom before the class floor is breached.


T_body	Species example	H_min(T)	β_healthy equivalent	A-shift vs 37°C
42°C	Birds (chicken, finch)	0.8662	0.748	A × 0.969
39°C	Rodents (mouse, rat)	0.8497	0.748	A × 0.987
37°C	Humans (REFERENCE)	0.8389	0.748	A × 1.000 (anchor)
35°C	Hibernating bats	0.8281	0.748	A × 1.013
32°C	Naked mole rat	0.8121	0.748	A × 1.033
25°C	Reptiles (anole lizard)	0.7752	0.748	A × 1.082

T_body	Species example	H_min(T)	β_healthy equivalent	A-shift vs 37°C
15°C	Fish (salmon, cold-water)	0.7241	0.748	$A \times 1.159$

HEALTHY AGING TRAJECTORY — AGE-STRATIFIED REFERENCE A-SCORES

Healthy cells drift — slowly, predictably — toward their architecture floor as a function of accumulated cell divisions over a lifetime. For the immune class, the drift rate is approximately 3.5% per generation — meaning each cell cycle carries a 3.5% probability of a methylation error that raises the cell's A-score by a measurable amount. Over decades, these errors accumulate. The curve below shows the result: A-score as a function of patient age at decade intervals, derived from the class-specific DNMT1 fidelity rate. At age 20, the reference healthy A-score sits well below 1.00; by age 80, it has drifted upward into the MARGINAL to DETECTABLE range. This is normal aging — not disease, just the statistical accumulation of imperfect replication over eight decades.

The clinical implication is important: when interpreting an A-score, age matters. A patient at age 30 with A = 1.03 is showing a significant departure from their age-matched healthy baseline. The same A = 1.03 at age 75 might be entirely consistent with age-expected drift. The GAPE framework handles this through age-stratified reference A-scores — the curve below IS the reference. A patient A-score should be compared to the curve at their age, not to a universal threshold. Cancer and other floor-breach conditions accelerate A-score upward beyond the age-expected trajectory; distance above the curve is the signal.



VERTEBRATE LIFESPAN CONTEXT

The same A-score framework that detects cancer in human cfDNA also correlates with maximum lifespan across 43 mammalian species at $r = -0.9018$ ($p = 1.6 \times 10^{-16}$, Mahaffey 2026 Nature Aging submission NATAGING-A13702). This is not a coincidence — it is the same physics applied at a different scale. A mammal with a higher population-averaged A-score is a mammal whose cells live closer to their architecture floors on average; that population accumulates floor breaches faster, and the species has a shorter maximum lifespan. The A = 1.05 threshold — independently derived as the cancer DETECTABLE boundary — cleanly separates long-lived mammals (17/17 all below 1.05) from short-lived mammals (11/11 all above 1.05) with complete accuracy. Same threshold. Same physics. Different timescale.

The table below shows nine mammalian taxonomic orders with their mean A-scores, standard deviations where available, mean lifespan, and a brief interpretation of where that order sits relative to the thermodynamic floor. Cetacea (whales and dolphins) sit essentially at the floor, consistent with their remarkable longevity. Primates including humans sit slightly above the floor. Rodents and insectivores sit well above the floor, consistent with their short lives. This class's reference A-score (0.971) places it in the context shown below — use this to understand how this particular cell architecture relates to the broader mammalian lifespan distribution.

Order	N	Mean A	Lifespan (yr)	Interpretation
Cetacea	5	0.997 ±0.016	99	At thermodynamic floor
Proboscidea	1	0.987 —	70	At floor
Primates	6	1.007 ±0.015	49	Slightly above floor
Artiodactyla	4	1.015 ±0.011	30	Slightly above floor
Chiroptera	4	1.041 ±0.007	35	Above floor (longevity outliers)
Carnivora	9	1.053 ±0.023	28	Above floor
Lagomorpha	2	1.114 ±0.002	8	Well above floor
Rodentia	6	1.125 ±0.018	9	Well above floor
Insectivora	1	1.157 —	2	Furthest from floor

INTERVENTION LEVERS — RANKED BY EXPECTED IMPACT

If a patient's A-score is elevated — MARGINAL, DETECTABLE, or URGENT — the clinical question is what can be done. The GAPE framework does not prescribe treatment; that remains the physician's judgment. But the framework does identify which intervention categories have the mechanistic strength to move this particular class's A-score back toward the floor. The five categories below are the standard anti-aging and cancer-adjacent interventions: senolytics (clearing senescent cells), metabolic restoration (NAD+, caloric restriction, exercise), epigenetic restoration (DNMT/TET modulators, HDAC inhibitors), reprogramming (cyclic Yamanaka factors), and checkpoint stringency (cell cycle control, immune checkpoint blockade).

The ranking below is specifically for the immune class and its primary failure mode (Cytokine Saturation). Impact Level 1 (Dominant) means the intervention directly addresses the class's failure mechanism and is the first lever to try. Level 2 (Strong) means substantial mechanistic support for A-score restoration in this class. Level 3 (Moderate) means helpful but not addressing the binding constraint. Level 4 (Limited) means the intervention works for other classes but does not target this class's specific biology. Level 5 (Not applicable) means the intervention category is biologically incompatible with this class — e.g., senolytics cannot act on post-mitotic terminal cells that do not become classically senescent, and reprogramming cannot be applied to terminal cells without losing the identity that defines them.

Impact	Category	Mechanism	Rationale
1 DOMINANT	Epigenetic Restoration	epigenetic rx	Dominant — TET2 restoration is the primary driver of exhaustion reversal
2 STRONG	Senolytics	senolytics	Strong — senescent T cells (p16+ exhausted) directly drive immune dysfunction
2 STRONG	Metabolic	metabolic	Strong — metabolic reprogramming to OxPhos restores effector function
2 STRONG	Immune Checkpoint	checkpoint	Strong — checkpoint blockade prevents exhaustion induction
3 MODERATE	Reprogramming	reprogramming	Moderate — only if exhaustion epigenome is irreversible

CANCER PANEL — ΔA RANKED

The TCGA cancer validation is where the framework earns its keep. For every cancer type in this class, we compute two numbers from published TCGA methylation data: A_healthy (from matched normal tissue) and A_tumor (from cancer samples in the same study). The ΔA = A_tumor – A_healthy is the signal the framework predicts should appear — and it does. Thresholds for tier assignment: A = 1.05 DETECT, A = 1.07 URGENT, A = 1.10 FLOOR BREACH. These thresholds are not fit to the cancer data; they were derived from the healthy-vs-age statistics and happen to align with the cancer validation cleanly. The bars below are ordered by |ΔA| descending so the cancers with the largest signal appear first.

Note that ΔA ranges typical for this class matter. Terminal-class cancers (glioma, glioblastoma) show ΔA ≈ 0.22–0.27 — the largest in the entire 28-cancer TCGA dataset. Cycling epithelial cancers show ΔA ≈ 0.13–0.19 — moderate but clearly detectable. Stromal and progenitor-lineage cancers show ΔA ≈ 0.10–0.17. Each cancer below is cited to its primary TCGA publication and shows sample size. Direct links from A-score to TCGA evidence — no black boxes between physics and clinical data.

#1	Lymphoma (DLBCL) n=48 · Chapuy 2018 Nat Genet	ΔA=+0.1331	A=1.161 FLOOR
#2	Leukemia (AML) n=200 · TCGA AML 2013 NEJM	ΔA=+0.1303	A=1.150 FLOOR
#3	Thymoma (THYM) n=124 · TCGA THYM 2018 Cell	ΔA=+0.1148	A=1.138 FLOOR

CORE METRICS — PUBLISHED AND DERIVED

Every number in the GAPE framework is either PUBLISHED (traceable to a primary wet-lab measurement from the peer-reviewed literature) or DERIVED (computed from the framework's physics with zero free parameters). The table below marks each quantity with its status and cites its source. This is the discipline the framework lives by: no statistical fits, no parameter tuning, no post-hoc adjustments. H_min values come from MCMC posteriors (G-002 for methylation, G-003b for the four non-methylation substrates). Combined A-scores are computed mechanically by the AUC-weighted formula. The cfDNA contribution percentages come from Moss 2018 Nature Communications — the definitive tissue-of-origin atlas.

Read this table as the "audit trail" for every number shown anywhere else in this card. When the framework says A_healthy = 0.9708, you can see exactly which β value produced it, which H_min was used, and where both came from. The n_bio field is currently marked PRELIMINARY because absolute n_bio values await the G-007 MCMC run; the class ordering has been confirmed (Spearman ρ = 0.905, ρ = 0.002 against commitment level) but the absolute values carry that caveat transparently.

Metric	Value	Source
cfDNA contribution (plasma)	70.0%	Moss 2018 Nat Commun — PUBLISHED
H_min methylation (class)	0.838889	G-002 MCMC posterior — DERIVED

Metric	Value	Source
H_min methylation σ (MCMC)	± 0.001200	G-002 17-chain MCMC R-hat<1.001 — DERIVED
n_bio (class-specific)	17.5	PRELIMINARY — ordering confirmed (p=0.905, p=0.002); absolute pending G-007
Healthy drift rate	3.5% / generation	DNMT1 fidelity loss under turnover — DERIVED
f_C2 (architecture-locked %)	9.8%	C2 fraction at healthy reference — DERIVED
Healthy A (methylation)	0.9708 (NORMAL)	$\beta=0.748$ — DERIVED
Disease A (methylation)	1.1377 (FLOOR BREACH)	$\beta=0.625$ — DERIVED
Healthy combined A (5 subs)	0.9699 (NORMAL)	$\Sigma(\text{AUC}_i \times A_i) / \Sigma(\text{AUC}_i)$ across 5 substrates — DERIVED
Disease combined A (5 subs)	1.1033 (FLOOR BREACH)	$\Sigma(\text{AUC}_i \times A_i) / \Sigma(\text{AUC}_i)$ across 5 substrates — DERIVED
ΔA combined	+0.1334	Combined signal — DERIVED

DATED PREDICTIONS FOR THIS CLASS

A prediction is not a claim unless it can fail. Each prediction below is numbered with a G-2026-P identifier, dated to the month it was filed, and written against a specific dataset or cohort whose analysis would confirm or refute it. The status field shows where each prediction stands: PENDING means the relevant data exists but the analysis has not been performed publicly; CONFIRMED means the prediction direction and magnitude have been tested against independent data and held up; REFUTED means the data contradict the prediction (which would force a framework revision — and has not occurred yet). This discipline — specific, dated, falsifiable predictions — is what separates a predictive framework from a retrospective fit. Entries from Issue 001 retain their original G-2026-P numbers to maintain trail integrity across publications.

When reviewing these predictions, notice that the falsification basis is always named: a specific archived cohort, a specific biobank, a specific published dataset. These are not vague "more research is needed" statements. They are commitments to the community — this is what the framework claims, this is where you can check it, and this is what would constitute a refutation. The predictions below are for the immune class specifically, filed as part of this publication or carried forward from Issue 001.

G-2026-P010

April 2026

PENDING

In prospective cohorts of patients with known clonal hematopoiesis of indeterminate potential (CHIP) and archived serial blood samples, the immune-class combined A-score across four or more substrates will show elevation above $A = 1.03$ at least 18 months before hematologic malignancy diagnosis in a majority of cases where subsequent malignancy occurs.

Basis: CHIP is a well-characterized pre-malignant state; the UK Biobank and ARIC cohorts contain archived serial samples from CHIP-positive individuals with longitudinal follow-up. The framework predicts that the combined A-score tracks the transition to overt malignancy earlier than any single substrate.

G-2026-P011

April 2026

PENDING

In patients receiving immune checkpoint inhibitor therapy, the immune-class A-score trajectory will distinguish responders from non-responders within the first 8 weeks of treatment. Responders are predicted to show A-score stabilization or decline; non-responders will show continued A-score drift.

Basis: Checkpoint blockade reshapes T-cell methylation. The framework predicts this shows up as an A-score signature before clinical response metrics (tumor size, cytokine panels) become discriminative. Falsifiable in any prospective ICI trial with serial cfDNA collection.

ISSUE 003 ADDENDUM — IMMUNE IN THE ATLAS

51 cell types: B, B-cells_EPIC, Baso, Bcell, Bmem, Bnv, CD14_monocytes, CD19_B-cells, CD4T, CD4T-cells_EPIC, CD4Tmem, CD4Tnv, CD4_T-cells, CD56_NK-cells, CD8T, CD8T-cells_EPIC, CD8Tmem, CD8Tnv, CD8_T-cells, Eos, Eosino, Eosinophils_reinius, Gran, IC, Kup, Leu, Lym, MP, Macro, Microglia, Mono, Monocytes_EPIC, Mye, NK, NK-cells_EPIC, Neu, Neutro, Neutrophils_EPIC, Neutrophils_reinius, PBMC, Plasma, Tcell, Treg, dendritic, eosinophil, granulocytes, macrophage, monocyte, neutrophil, tcell, whole_blood

Identity loci: 42,134 (methyl) · **H_min by substrate:** methyl 0.838889 · nucl 0.989930 · fuzz 0.830377 · wps 0.589644 · frag 0.711534 · **methyl ceiling** 1.1921

Healthy age-matched band (age_reference_matrix.json). $A = 1.0$ is the commitment line; healthy reads on this band.

Age (midpoint)	n	A mean	p10	p90	β mean	Source
4	45	0.9062	0.8644	0.9479	0.780	Alisch 2012
14	58	0.9212	0.8780	0.9643	0.773	Alisch+Hannum
24	95	0.9316	0.8868	0.9764	0.768	Hannum 2013

Age (midpoint)	n	A mean	p10	p90	β mean	Source
34	102	0.9397	0.8932	0.9863	0.764	Hannum 2013
44	115	0.9477	0.9021	0.9934	0.760	Hannum 2013
54	108	0.9556	0.9083	1.0029	0.756	Hannum 2013
64	98	0.9652	0.9166	1.0138	0.751	Hannum+Horvath
74	85	0.9764	0.9269	1.0260	0.745	Hannum 2013
84	42	0.9873	0.9369	1.0377	0.739	Hannum 2013
95	15	0.9996	0.9466	1.0527	0.732	Hannum oldest-old

CELL IDENTITY & CLINICAL CONTEXT

Includes	Colon, rectum, stomach, lung (adeno + squamous), bladder, cervix, skin (melanoma), kidney (clear cell + papillary), head & neck, ovary, endometrium, thyroid
Cancers (this class)	Colorectal adenocarcinoma (COAD), Rectal adenocarcinoma (READ), Lung adenocarcinoma (LUAD), Lung squamous cell carcinoma (LUSC), Bladder urothelial carcinoma (BLCA), Ovarian serous carcinoma (OV), Stomach adenocarcinoma (STAD), Cervical squamous cell carcinoma (CESC), Head and neck squamous cell carcinoma (HNSC), Kidney clear cell RCC (KIRC), Kidney papillary RCC (KIRP), Skin cutaneous melanoma (SKCM), Thyroid carcinoma (THCA), Endometrial carcinoma (UCEC) — 14/28 TCGA types
Non-cancer applications	Inflammatory bowel disease (IBD), chronic colitis, Barrett's esophagus, chronic viral hepatitis, GERD-associated dysplasia, smoker's lung (pre-LUSC field cancerization), HPV carrier status, field cancerization, epithelial aging
Primary failure mode	Replication Ceiling
Warburg status	WALL CROSSED

COMMENTARY

The cycling epithelial class is the largest and most clinically important in the GAPE dataset. Fourteen of the 28 confirmed TCGA cancer types fall here — colon, rectum, lung (adenocarcinoma and squamous), bladder, cervix, head and neck, ovary, stomach, kidney (clear cell and papillary), skin melanoma, endometrium, and thyroid. These are the cancers that colonoscopy, Pap smears, low-dose CT, and mammography were invented to catch. Together they account for the majority of cancer deaths worldwide. Every family affected by one of these diseases is a family this card has to speak to honestly.

Cycling epithelial cells are defined by function: continuous division to replace the epithelial lining of organs exposed to external or internal environments. Colonic mucosa turns over completely in four to seven days. Cervical epithelium renews in weeks. Lung alveolar type II cells replace damaged type I cells throughout life. This continuous division is precisely why cycling epithelia are cancer-prone: every division requires DNMT1 to copy methylation patterns across 19.6 million CpG sites, and accumulated errors over decades drive the global hypomethylation signature that GAPE detects across all five substrates simultaneously. The 5.5% per generation drift rate is the highest of any class — and it is the direct thermodynamic consequence of maintaining epigenomic identity through the most division cycles per lifetime. H_min = 0.856055 (G-002 chain 1 of 17, R-hat 1.0003, posterior ±0.000800) reflects a floor tight enough to define a class but loose enough to accommodate the tissue-to-tissue variation across all 14 cancer types that share it.

This class is where MESA — the four-substrate cfDNA test (Li 2024 Genome Med) — was developed and validated. MESA achieves AUC = 0.931 on colorectal cancer using methylation, fragment size, nucleosome occupancy, and WPS from a single blood tube. This is the strongest independent validation of the multi-substrate framework in the entire document. The measured inter-substrate correlation r = 0.54 between MESA's four signals (VAL-014) matches exactly what GAPE predicts when four physical windows measure the same underlying thermodynamic floor departure: the signals share approximately 85% of their information, and the combined-to-single improvement ratio d_combined/d_single = 1.15× lands precisely in the GAPE-predicted range. MESA's ML feature ranking downweights nucleosome occupancy relative to the other substrates without explaining why. GAPE explains why: for the cycling class, nucl saturates at ceiling A = 1.020, below the BREACH threshold at A = 1.10. The feature is downweighted because it physically cannot resolve severity past a certain point. This is not a machine-learning discovery; it is a thermodynamic measurement constraint that the ML model reflects empirically without naming. Adding the fifth substrate (DELFI fragment size from Cristiano 2019 Nature, AUC 0.940 across 7 cancer types) brings the theoretical detection ceiling to AUC = 1.000, with the gap from MESA's 0.931 attributable to bulk blood dilution — addressable by tissue-specific cfDNA deconvolution.

The two deadliest cancers in this class deserve explicit naming. Colorectal cancer (COAD + READ, TCGA Research Network 2012 Nature, n=169 matched-normal + tumor methylation) is the second-leading cancer killer in the United States and claims over 900,000 lives per year globally. VAL-007 reports cfDNA ΔA = +0.158 for colorectal in stool-based sampling — fully in FLOOR BREACH tier. Lung cancer, split between adenocarcinoma (LUAD, TCGA 2014 Nature, n=230) and squamous cell carcinoma (LUSC, TCGA 2012 Nature, n=178), is the leading cancer killer worldwide with over 1.8 million annual deaths. VAL-007 plasma cfDNA ΔA for LUAD = +0.144; LUSC is comparable. These two cancers alone account for roughly 30% of all cancer mortality. The five-substrate framework for cycling class is not a theoretical advance over MESA — it is the measurable path toward pre-clinical detection of the two cancers most responsible for cancer-related death. Every delayed diagnosis in these cancers is a calibration problem the framework could have caught earlier.

The remaining twelve cycling cancers each deserve their place on this card. Bladder urothelial carcinoma (BLCA, TCGA 2014 Nature, n=131) shows the largest ΔA in the panel at urine cfDNA ΔA = +0.185 — urine is the natural specimen for bladder cancer and concentrates the signal. Ovarian serous carcinoma (OV, TCGA 2011 Nature, n=67) is the most lethal gynecologic cancer due to late detection; pre-clinical cfDNA monitoring is the framework's proposed intervention. Stomach adenocarcinoma (STAD, TCGA 2014 Nature, n=295) shows plasma cfDNA ΔA = +0.153. Cervical squamous (CESC, TCGA 2017 Nature, n=228) and Head and Neck squamous (HNSC, TCGA 2015 Nature, n=504) share HPV-driven methylation changes; saliva cfDNA works for HNSC (ΔA = +0.146). Kidney clear cell (KIRC, TCGA 2013 Nature, n=318) and papillary (KIRP, TCGA 2016 NEJM,

n=161) are two biologically distinct malignancies that share the kidney. Skin cutaneous melanoma (SKCM, TCGA 2015 Cell, n=477) ranks among the deadliest skin cancers and its methylation signature reflects ultraviolet mutagenesis. Thyroid carcinoma (THCA, TCGA 2014 Cell, n=51) stands out as an outlier in the class — slow-growing, often curable, but still methylation-detectable. Endometrial carcinoma (UCEC, TCGA 2013 Nature, n=118) completes the panel, and each of these cancers has a family behind it. Every TCGA citation is a cohort of patients who consented to share their tissue, and every TCGA dataset is a gift to the patients who will come after. Clickable DOIs for all thirteen TCGA primary sources listed above are on the Data Sources page.

A clinically important consequence of the saturation pattern for cycling class. Cycling has exactly one substrate that saturates below BREACH — nucleosome occupancy, ceiling $A = 1.020$ (Doebley 2022 Nat Commun). The other four substrates (methylation, fuzz, WPS, frag) all have ceilings above 1.22 and carry the full progression signal past BREACH for every cycling cancer in the panel. Because only one substrate saturates, the single-substrate saturation of nucl alone is NOT cancer-specific — it will occur in cycling cancers AND in non-cancer inflammatory conditions (IBD flares, chronic colitis, Barrett's esophagus with dysplasia, HPV-infected cervical tissue). What nucl saturation DOES tell you for cycling class is that the cell population has lost enough of its epithelial identity that the nucleosome positional signature is pinned at random. Severity grading, cancer-vs-benign distinction, tissue-of-origin discrimination (colon vs lung vs bladder), and any subtyping (COAD MSI status, LUAD vs LUSC, HPV+ vs HPV- HNSC) all come from methyl, fuzz, WPS, and frag. Report the all-5 A_{combined} for historical continuity and MESA-comparable reporting; report the A_{active} (4/5) for progression tracking and serial monitoring. The mask is moderate (+16.3%) but non-zero; in serial monitoring across multiple blood draws the difference compounds as nucl stays pinned at its ceiling while the other four continue to drift with disease progression or treatment response.

FIVE-SUBSTRATE FIDELITY GAUGE

The gauge below is the central image of Issue 002. A single horizontal bar represents the A-score axis — from $A = 0.90$ (below the floor) through $A = 1.00$ (exactly at the Cycling class floor, $H_{\text{min}} = 0.8561$) up through the clinical threshold zones: MARGINAL at $A = 1.01$, DETECTABLE at 1.05, URGENT at 1.07, and FLOOR BREACH at 1.10. Five dots are plotted on this single axis for a healthy reference sample (Normal colon, markers above the bar), and five more dots for a disease reference (COAD (TCGA n=97), markers below). Each dot is one substrate: methylation, nucleosome occupancy, nucleosome fuzziness, WPS, or fragment size (DELFI). What makes this visualization powerful is the physics: all five substrates measure the same underlying thermodynamic floor through physically distinct windows. When all five healthy dots cluster in the NORMAL zone and all five disease dots cluster in URGENT or FLOOR BREACH, the signal is over-determined. Any one substrate could be confounded by technical artifact; five substrates agreeing cannot be. This is the "less blurry" advantage: $\sqrt{5}$ noise reduction for free.




Look at the spread in each direction. If the healthy five-dot cluster sits tightly together near $A = 0.97-1.00$, that indicates the reference population is at its architecture floor — precisely where the framework predicts healthy cycling cells should sit. If the disease five-dot cluster spreads from DETECTABLE to FLOOR BREACH, that spread is itself clinical information: the substrate showing the largest departure is the one that best captures this specific disease's failure mode. For this class, the primary failure mode is Replication Ceiling, and the substrate ranking shown further down the card explains which of the five substrates best reveals it.

SUBSTRATE-BY-SUBSTRATE BREAKDOWN


Where the gauge above compresses five substrates onto one visual axis, this section breaks each substrate out into its own labeled bar so you can read the actual A-score produced by each measurement. The physics is the same across all five: take the raw laboratory value (a β methylation fraction, a nucleosome occupancy probability, a WPS score, etc.), compute its Shannon entropy $H(\text{value})$, then divide by the class-specific and substrate-specific H_{min} to get a dimensionless A-score comparable across substrates. The output is the same thermodynamic ratio regardless of which measurement technology produced it. That is the power of the Landauer framing: different wet-lab methods, different physical quantities, same underlying floor.

The formula at the foot of this section — $A_{\text{combined}} = \Sigma(AUC_i \times A_i) / \Sigma(AUC_i)$ — combines the five individual A-scores into a single weighted score, using each substrate's published detection AUC as its weight. Higher-AUC substrates contribute more to the combined result. For cycling cells, the healthy combined A-score below should sit at approximately 0.97 (NORMAL tier); the disease combined A-score should sit above 1.05 (DETECTABLE or higher). The difference between them — the ΔA combined — is the signal available to a blood draw that analyzes all five substrates simultaneously. A single-substrate assay sees a smaller signal at higher noise; the five-substrate combination sees the same signal at approximately $\sqrt{5}$ lower noise. Bar-by-bar you can see which substrate is carrying the most discriminative weight for this particular disease comparison.

HEALTHY REFERENCE — Normal colon



DISEASE REFERENCE — COAD (TCGA n=97)
ACTIVE — tracks progression




Combined A — healthy (all 5)	0.97017 (NORMAL)
Combined A — disease (all 5)	1.10937 (FLOOR BREACH)
ΔA combined (5 substrates)	+0.13920 vs methylation-only ΔA =+0.1859
Active A — healthy (non-saturated)	0.97017 (NORMAL) [5/5 active]
Active A — disease (non-saturated)	1.13203 (FLOOR BREACH) [4/5 active]

ΔA active (real progression signal)	+0.16187 (active substrates track true progression)
Saturation status (disease)	1 substrate saturated: Nucleosome Occupancy

Two combined values are shown because some substrates physically cannot resolve past their class ceiling ($A_{\text{max}} = 1/H_{\text{min}}$). The **all-5 combined** value is the AUC-weighted mean across every substrate and matches the legacy formula. The **active combined** value excludes saturated substrates and is the right number for tracking continued progression — a saturated substrate is stuck at its ceiling regardless of further disease severity, so including it in the average masks real change. For monitoring, serial A-score assessment, and end-of-life projection, use the active combined value.

DISEASE SIGNATURE COMPARISON — healthy colon through Barrett's, HNSC, LUAD, COAD, BLCA

Six conditions on a single chart, all in the cycling-epithelial class. β values reproduce per-substrate A-scores matching Evidence Report VAL-007 per-cancer cfDNA targets: COAD (TCGA n=97, doi:10.1038/nature11252, cfDNA $\Delta A = +0.158$), LUAD (TCGA n=82, doi:10.1038/nature13385, cfDNA $\Delta A = +0.144$), HNSC (TCGA n=504, doi:10.1038/nature14129, saliva cfDNA $\Delta A = +0.146$), BLCA (TCGA n=131, doi:10.1038/nature12965, urine cfDNA $\Delta A = +0.185$), plus Barrett's esophagus as a non-cancer field-cancerization state for pre-malignant context. All five cancers sit in URGENT to FLOOR BREACH; Barrett's sits in MARGINAL. Substrate-saturation note: for cycling class, nucleosome occupancy saturates at $A \approx 1.020$ in every cancer in the panel — so nucl alone cannot distinguish one cycling cancer from another. Methyl, fuzz, WPS, and frag carry per-cancer discrimination. BLCA shows the largest ΔA (urine cfDNA, high tumor fraction); HNSC/LUAD cluster together at moderate departure; COAD sits between — consistent with the specimen-type physics of each cancer's optimal cfDNA compartment.




What happens past the ceiling is condition-specific. The A-score magnitude alone does not track severity once the ceiling is crossed — different diseases take different post-breach paths that the full five-substrate divergence pattern reveals. The subsection below walks through this class's post-breach trajectory: which substrates keep carrying signal, where the Warburg metabolic lock-in happens, where the glucose inversion point sits, and where the framework currently has high confidence versus where the open research questions remain.

POST-BREACH TRAJECTORY — WHAT HAPPENS AFTER A CROSSES 1.10 (CYCLING CLASS)

The cycling-epithelial class is the largest and most clinically important — 14 of the 28 validated TCGA cancer types fall here, accounting for the majority of cancer deaths worldwide. Every family affected by colon, lung, cervical, bladder, head-and-neck, or endometrial cancer is a family this card has to speak to honestly. Post-breach trajectory matters most here because serial monitoring under treatment is already the standard of care for these cancers.

The zone language below — Warburg boundary, glucose inversion, point of no return — applies only to post-breach cells. These terms are not used on the pre-breach bar, and the qualitative A-value anchors (~1.15, ~1.25, ~1.40+) are boundaries between therapeutic windows, not additional diagnostic tiers. Specific A-values are filed as prospective predictions to be validated against survival-stratified cohorts.



WHICH SUBSTRATES CARRY POST-BREACH SIGNAL — CYCLING CLASS

Cycling-class physics: nucleosome occupancy saturates at $A \approx 1.020$ for every cancer in the panel; the other four substrates carry the per-cancer discrimination and post-breach progression signal.

Substrate	Ceiling (A)	Post-breach status
Methylation	1.30	Carries signal throughout all four zones
Fuzziness	1.28	Carries signal throughout all four zones
Windowed protection	1.18	Carries signal throughout all four zones
Fragment size	1.25	Carries signal throughout all four zones
Nucleosome occupancy	1.020	Saturated at ceiling — no further signal post-breach

CONDITION-BY-CONDITION TRAJECTORY

Healthy colonic mucosa — reference, high cfDNA turnover

Known: Cycling epithelia turn over fast — colonic mucosa in 4-7 days, cervical epithelium in weeks. This is the highest division-rate class in the body, which is precisely why it is cancer-prone: every division requires DNMT1 to copy methylation across 19.6 million CpG sites, and accumulated errors over decades drive the global hypomethylation signature that GAPE detects across all five substrates simultaneously.

COAD (TCGA n=97) — $A \approx 1.158$, CROSSED CEILING

Known: Colorectal adenocarcinoma is the most validated cycling-class cancer in cfDNA literature. $\Delta A = +0.158$ matches VAL-007 exactly. COAD shows the classical Warburg progression archetype: monotonic rise in A_{active} under disease progression, response to FOLFOX or FOLFIRI produces trajectory bend toward healthy, relapse produces trajectory re-elevation. This is the cleanest example of the framework tracking clinical reality.

Unknown: whether the MSI-high vs MSS split (consequential for immunotherapy response) produces distinguishable A-score paths; whether the A_{active} trajectory under FOLFOX predicts response-duration for KRAS-mutant vs KRAS-wild-type cases.

Test that closes the gap: G-2026-P027: Retrospective reanalysis of GALAXY CRC MRD cohort ($n=1,000+$ stage II-III CRC patients with serial cfDNA post-resection). Prediction: A_{active} trajectory at 6 weeks post-resection will predict 2-year recurrence-free survival with $AUC \geq 0.80$, matching or exceeding current ctDNA-positive/negative classification.

LUAD / LUSC (lung) — $A \approx 1.144$, CROSSED CEILING

Known: Lung adenocarcinoma and squamous cell carcinoma show moderate cycling-class ΔA (+0.144) consistent with high cfDNA contribution from pulmonary epithelial turnover. The DELFI fragmentomics signature is strongest for lung cancer — Cristiano 2019 Nature demonstrated this with AUC 0.94 across 7 cancer types including lung.

Unknown: whether EGFR-mutant, ALK-fusion, and KRAS-mutant LUAD subtypes show distinguishable A-score paths under tyrosine kinase inhibitor therapy; whether pemetrexed response can be predicted from baseline A_{active} .

Test that closes the gap: G-2026-P028: Prospective cohort of 200 EGFR-mutant LUAD patients on osimertinib with serial cfDNA every 3 months. Prediction: A_active trajectory slope predicts resistance emergence 3-6 months before standard imaging progression, with sensitivity ≥ 0.70 at specificity 0.85.

BLCA (TCGA n=131) — $A \approx 1.185$, LARGEST CYCLING-CLASS ΔA

Known: Bladder urothelial carcinoma shows the largest cycling-class ΔA in the framework (+0.185) — partly structural (urothelial biology) and partly specimen (urine cfDNA has higher tumor fraction than blood for this cancer). BLCA sits in the structural-only post-breach zone at diagnosis for most patients. Past the Warburg boundary, metabolic intervention alone is unlikely to restore cellular fidelity.

Unknown: whether urine cfDNA A-score trajectories can predict BCG response in non-muscle-invasive BLCA; whether the framework can distinguish progressors from persistent but stable Ta/T1 disease.

Test that closes the gap: G-2026-P029: Prospective cohort of 150 non-muscle-invasive BLCA patients receiving BCG with urine cfDNA at baseline, 3, 6, 12 months. Prediction: urine A_active trajectory during induction BCG will discriminate responders from non-responders with AUC ≥ 0.75 , outperforming standard cytology.

WHAT THE FRAMEWORK CAN AND CANNOT SAY TODAY

The framework has high confidence for cycling-class post-breach: (1) four substrates carry full post-breach signal with nucl saturating; (2) the 14-cancer TCGA validation is the largest coherent structural prediction in the framework (96.4% direction-of-BREACH accuracy); (3) this class has the longest ctDNA literature to compare against, with framework predictions matching published ΔA values within measurement noise.


The framework has not yet tested A-score trajectories longitudinally against published treatment-response cohorts. Predictions G-2026-P027, P028, P029 define the specific validation plans for colorectal MRD monitoring, EGFR-mutant LUAD resistance prediction, and BLCA BCG response.

This is the scientific posture the framework commits to: definitive where the evidence is definitive, open where it is open, and specific about the experiments that would close each gap. Predictions filed for this class: G-2026-P027, G-2026-P028, G-2026-P029. A reviewer who wants to falsify any claim here can do so against the named cohorts. A collaborator who wants to extend the framework has specific hypotheses to test.

THREE-COMPONENT DECOMPOSITION (C1/C2/C3) — PER SUBSTRATE

Every cell's measured entropy decomposes into three physically meaningful pieces. C1 is the universal Landauer floor: $H_MIN_GLOBAL = 0.756500$, measured in frontal cortex neurons by Lister 2013. Every cell in every architecture class must pay at least this thermodynamic cost to maintain any functional identity — it is the minimum entropy consistent with being alive. C2 is the class-specific overhead: the additional entropy the cycling class must carry above the universal floor to encode its particular architecture. For this class, C2 represents approximately 12.1% of the healthy reference entropy. C3 is the accessible gap: $\max(0, H_actual - H_min)$. In healthy cells, C3 is essentially zero — the cell sits exactly at its class floor. When C3 grows, the cell is departing from its architecture class. Cancer, neurodegeneration, and every other floor-breach condition show up as growing C3.

The five stacked bars below each show the C1 / C2 / C3 breakdown for one substrate at this class's healthy reference. Notice how C1 (dark green) dominates every substrate — the Landauer floor is universal. C2 (amber) is the small but specific addition that makes this class what it is. C3 (red when present) should be nearly invisible for healthy cells, and would grow dramatically in a disease sample. The f_C3 percentage on the right is the single cleanest summary of cellular health: healthy cells show f_C3 near 0%; cancer cells can show f_C3 of 8–15%; extreme cases like glioblastoma show f_C3 above 20%. This is not a statistical threshold calibrated from disease data. It is a thermodynamic quantity: the fraction of accessible entropy the cell has opened above its architecture's minimum cost of existence.



BEST TESTING METHOD RANKING — CLINICAL UTILITY FOR THIS CLASS

Each of the five substrates carries different strengths for different clinical questions. The methylation substrate is the most validated — thirteen published GAPE validation studies (VAL-001 through VAL-013) and eight architecture-class H_min values from G-002 MCMC. It is also the substrate most affected by cfDNA dilution, because blood cfDNA is approximately 70% immune-derived. Fragment size (DELF) is the most forgiving of dilution — short-fragment enrichment persists even when tumor-derived cfDNA is a small fraction of total plasma DNA. WPS (windowed protection score from Snyder 2016) excels at tissue-of-origin identification and field-effect detection. Nucleosome occupancy and fuzziness add chromatin-level information that methylation alone cannot reveal.

The ranking below is specifically for the cycling class. The order reflects three factors: (1) the substrate's published single-substrate AUC for diseases in this class; (2) the biology of this class's primary failure mode (Replication Ceiling); and (3) the practical constraint of cfDNA contribution in the relevant specimen (plasma for most classes, CSF for terminal, tissue biopsy where cfDNA is too dilute). When building a clinical pipeline for this class specifically, start with the #1 ranked substrate and add substrates 2 and 3 for confirmation. Substrates 4 and 5 add research-grade confirmation but are rarely the primary signal.

#	Substrate	Best for	Rationale
1	Methylation	Pan-cancer screening (14/28 TCGA types)	Methylation is the primary substrate for this class. Fourteen TCGA cycling cancers calibrated against it. The framework's most battle-tested data stream.
2	Windowed Protection Score	Field-effect detection	Chromatin accessibility at cycling-class identity promoters shows depletion before clinical cancer. Snyder 2016 identified this signal at 15/15 tissue types.
3	Fragment Size (DELFI)	Early detection and monitoring	DELFI pipeline clinically validated for CRC, lung, and other cycling-class cancers. Short/long fragment ratio tracks tumor burden with high sensitivity.
4	Nucleosome Occupancy	Tissue-of-origin confirmation	ATAC-seq occupancy maps complement methylation when single-modality is ambiguous. Most useful when paired with methylation for discordance analysis.
5	Nucleosome Fuzziness	Aggressiveness grading	Nucleosome fuzziness correlates with proliferation rate. Secondary — combine with clinical staging (Dukes, TNM) for prognosis.

BODY TEMPERATURE SCALING — $\alpha = 2.0$ LANDAUER CORRECTION

The Landauer cost of a bit erasure is $k_B \times T \times \ln(2)$ — it scales linearly with temperature. Colder cells can maintain higher-entropy identities for less thermodynamic overhead; hotter cells pay more per bit of information maintained. The GAPE framework captures this with a simple scaling law: $H_{\min}(T) = H_{\min}(37^\circ\text{C}) \times (T_{\text{body}} / 310.15\text{ K})^\alpha$, with $\alpha = 2.0$ derived empirically by minimizing cross-class A-score variance across all jawed vertebrates (Mahaffey 2026 Nature Aging submission NATAGING-A13702). The $\alpha = 2.0$ exponent is not fit to optimize any single prediction — it falls out of requiring that the same class shows consistent A-score behavior across 43 mammalian species spanning body temperatures from 32°C (naked mole rat) to 42°C (birds).

The table below shows what the cycling class floor $H_{\min}$ looks like at seven representative body temperatures. At 37°C (human reference, highlighted row), $H_{\min} = 0.8561$. At 42°C (birds), the floor is higher because the Landauer tax is larger. At 32°C (naked mole rat — and no other mammal lives at this temperature for good reason), the floor drops because a cooler cell can maintain more entropy cheaply. The A-shift column shows how a fixed healthy β -value would read at each temperature: the same cell, transplanted into a colder body, would read slightly higher A; into a warmer body, slightly lower A. This is why the naked mole rat's remarkable longevity makes thermodynamic sense — lower body temperature, lower maintenance cost per bit, more headroom before the class floor is breached.


T_body	Species example	H_min(T)	β_{healthy} equivalent	A-shift vs 37°C
42°C	Birds (chicken, finch)	0.8839	0.738	$A \times 0.969$
39°C	Rodents (mouse, rat)	0.8671	0.738	$A \times 0.987$
37°C	Humans (REFERENCE)	0.8561	0.738	$A \times 1.000$ (anchor)
35°C	Hibernating bats	0.8451	0.738	$A \times 1.013$
32°C	Naked mole rat	0.8287	0.738	$A \times 1.033$
25°C	Reptiles (anole lizard)	0.7911	0.738	$A \times 1.082$
15°C	Fish (salmon, cold-water)	0.7389	0.738	$A \times 1.159$

HEALTHY AGING TRAJECTORY — AGE-STRATIFIED REFERENCE A-SCORES

Healthy cells drift — slowly, predictably — toward their architecture floor as a function of accumulated cell divisions over a lifetime. For the cycling class, the drift rate is approximately 5.5% per generation — meaning each cell cycle carries a 5.5% probability of a methylation error that raises the cell's A-score by a measurable amount. Over decades, these errors accumulate. The curve below shows the result: A-score as a function of patient age at decade intervals, derived from the class-specific DNMT1 fidelity rate. At age 20, the reference healthy A-score sits well below 1.00; by age 80, it has drifted upward into the MARGINAL to DETECTABLE range. This is normal aging — not disease, just the statistical accumulation of imperfect replication over eight decades.

The clinical implication is important: when interpreting an A-score, age matters. A patient at age 30 with $A = 1.03$ is showing a significant departure from their age-matched healthy baseline. The same $A = 1.03$ at age 75 might be entirely consistent with age-expected drift. The GAPE framework

handles this through age-stratified reference A-scores — the curve below IS the reference. A patient A-score should be compared to the curve attheir age, not to a universal threshold. Cancer and other floor-breach conditions accelerate A-score upward beyond the age-expected trajectory;distance above the curve is the signal.



VERTEBRATE LIFESPAN CONTEXT

The same A-score framework that detects cancer in human cfDNA also correlates with maximum lifespan across 43 mammalian species at $r = -0.9018$ ($p = 1.6 \times 10^{-16}$, Mahaffey 2026 Nature Aging submission NATAGING-A13702). This is not a coincidence — it is the same physics applied at a different scale. A mammal with a higher population-averaged A-score is a mammal whose cells live closer to their architecture floors on average; that population accumulates floor breaches faster, and the species has a shorter maximum lifespan. The $A = 1.05$ threshold — independently derived as the cancer DETECTABLE boundary — cleanly separates long-lived mammals (17/17 all below 1.05) from short-lived mammals (11/11 all above 1.05) with complete accuracy. Same threshold. Same physics. Different timescale.

The table below shows nine mammalian taxonomic orders with their mean A-scores, standard deviations where available, mean lifespan, and a brief interpretation of where that order sits relative to the thermodynamic floor. Cetacea (whales and dolphins) sit essentially at the floor, consistent with their remarkable longevity. Primates including humans sit slightly above the floor. Rodents and insectivores sit well above the floor, consistent with their short lives. This class's reference A-score (0.969) places it in the context shown below — use this to understand how this particular cell architecture relates to the broader mammalian lifespan distribution.

Order	N	Mean A	Lifespan (yr)	Interpretation
Cetacea	5	0.997 ±0.016	99	At thermodynamic floor
Proboscidea	1	0.987 —	70	At floor
Primates	6	1.007 ±0.015	49	Slightly above floor
Artiodactyla	4	1.015 ±0.011	30	Slightly above floor
Chiroptera	4	1.041 ±0.007	35	Above floor (longevity outliers)
Carnivora	9	1.053 ±0.023	28	Above floor
Lagomorpha	2	1.114 ±0.002	8	Well above floor
Rodentia	6	1.125 ±0.018	9	Well above floor
Insectivora	1	1.157 —	2	Furthest from floor

INTERVENTION LEVERS — RANKED BY EXPECTED IMPACT

If a patient's A-score is elevated — MARGINAL, DETECTABLE, or URGENT — the clinical question is what can be done. The GAPE framework does not prescribe treatment; that remains the physician's judgment. But the framework does identify which intervention categories have the mechanistic strength to move this particular class's A-score back toward the floor. The five categories below are the standard anti-aging and cancer-adjacent interventions: senolytics (clearing senescent cells), metabolic restoration (NAD+, caloric restriction, exercise), epigenetic restoration (DNMT/TET modulators, HDAC inhibitors), reprogramming (cyclic Yamanaka factors), and checkpoint stringency (cell cycle control, immune checkpoint blockade).

The ranking below is specifically for the cycling class and its primary failure mode (Replication Ceiling). Impact Level 1 (Dominant) means the intervention directly addresses the class's failure mechanism and is the first lever to try. Level 2 (Strong) means substantial mechanistic support for A-score restoration in this class. Level 3 (Moderate) means helpful but not addressing the binding constraint. Level 4 (Limited) means the intervention works for other classes but does not target this class's specific biology. Level 5 (Not applicable) means the intervention category is biologically incompatible with this class — e.g., senolytics cannot act on post-mitotic terminal cells that do not become classically senescent, and reprogramming cannot be applied to terminal cells without losing the identity that defines them.

Impact	Category	Mechanism	Rationale
1 DOMINANT	Checkpoint Stringency	checkpoint	Dominant — G1/S and G2/M checkpoint activation is the primary lever
2 STRONG	Senolytics	senolytics	Strong — senescent cells in the crypt drive stem cell niche dysfunction
2 STRONG	Epigenetic Restoration	epigenetic rx	Strong — MMR/checkpoint restoration directly addresses the inversion
3 MODERATE	Metabolic	metabolic	Moderate — useful but Replication Throughput Ceiling is the binding constraint
4 LIMITED	Reprogramming	reprogramming	Limited — cycling architecture is the functional requirement

CANCER PANEL — ΔA RANKED

The TCGA cancer validation is where the framework earns its keep. For every cancer type in this class, we compute two numbers from published TCGA methylation data: A_healthy (from matched normal tissue) and A_tumor (from cancer samples in the same study). The $\Delta A = A_{\text{tumor}} - A_{\text{healthy}}$ is the signal the framework predicts should appear — and it does. Thresholds for tier assignment: A = 1.05 DETECT, A = 1.07 URGENT, A = 1.10 FLOOR BREACH. These thresholds are not fit to the cancer data; they were derived from the healthy-vs-age statistics and happen to align with the cancer validation cleanly. The bars below are ordered by $|\Delta A|$ descending so the cancers with the largest signal appear first.

Note that ΔA ranges typical for this class matter. Terminal-class cancers (glioma, glioblastoma) show $\Delta A \approx 0.22\text{--}0.27$ — the largest in the entire 28-cancer TCGA dataset. Cycling epithelial cancers show $\Delta A \approx 0.13\text{--}0.19$ — moderate but clearly detectable. Stromal and progenitor-lineage cancers show $\Delta A \approx 0.10\text{--}0.17$. Each cancer below is cited to its primary TCGA publication and shows sample size. Direct links from A-score to TCGA evidence — no black boxes between physics and clinical data.



#11	 Thyroid (THCA) n=51 · TCGA THCA 2014 Cell	 $\Delta A = +0.1398$	A=1.091 URGENT
#12	 Kidney Clear Cell (KIRC) n=318 · TCGA KIRC 2013 Nature	 $\Delta A = +0.1319$	A=1.123 FLOOR
#13	 Kidney Papillary (KIRP) n=161 · TCGA KIRP 2016 NEJM	 $\Delta A = +0.1279$	A=1.121 FLOOR

CORE METRICS — PUBLISHED AND DERIVED

Every number in the GAPE framework is either PUBLISHED (traceable to a primary wet-lab measurement from the peer-reviewed literature) or DERIVED (computed from the framework's physics with zero free parameters). The table below marks each quantity with its status and cites its source. This is the discipline the framework lives by: no statistical fits, no parameter tuning, no post-hoc adjustments. $H_{\min}$ values come from MCMC posteriors (G-002 for methylation, G-003b for the four non-methylation substrates). Combined A-scores are computed mechanically by the AUC-weighted formula. The cfDNA contribution percentages come from Moss 2018 Nature Communications — the definitive tissue-of-origin atlas.

Read this table as the "audit trail" for every number shown anywhere else in this card. When the framework says $A_{\text{healthy}} = 0.9693$, you can see exactly which β value produced it, which $H_{\min}$ was used, and where both came from. The n_{bio} field is currently marked PRELIMINARY because absolute n_{bio} values await the G-007 MCMC run; the class ordering has been confirmed (Spearman $\rho = 0.905$, $p = 0.002$ against commitment level) but the absolute values carry that caveat transparently.

Metric	Value	Source
cfDNA contribution (plasma)	12.0%	Moss 2018 Nat Commun — PUBLISHED
$H_{\min}$ methylation (class)	0.856055	G-002 MCMC posterior — DERIVED
$H_{\min}$ methylation σ (MCMC)	± 0.000800	G-002 17-chain MCMC R-hat < 1.001 — DERIVED
n_{bio} (class-specific)	19.5	PRELIMINARY — ordering confirmed ($\rho = 0.905$, $p = 0.002$); absolute pending G-007
Healthy drift rate	5.5% / generation	DNMT1 fidelity loss under turnover — DERIVED
f_{C2} (architecture-locked %)	12.1%	C2 fraction at healthy reference — DERIVED
Healthy A (methylation)	0.9693 (NORMAL)	$\beta = 0.738$ — DERIVED
Disease A (methylation)	1.1552 (FLOOR BREACH)	$\beta = 0.562$ — DERIVED
Healthy combined A (5 subs)	0.9702 (NORMAL)	$\Sigma(\text{AUC}_i \times A_i) / \Sigma(\text{AUC}_i)$ across 5 substrates — DERIVED
Disease combined A (5 subs)	1.1094 (FLOOR BREACH)	$\Sigma(\text{AUC}_i \times A_i) / \Sigma(\text{AUC}_i)$ across 5 substrates — DERIVED
ΔA combined	+0.1392	Combined signal — DERIVED

DATED PREDICTIONS FOR THIS CLASS

A prediction is not a claim unless it can fail. Each prediction below is numbered with a G-2026-P identifier, dated to the month it was filed, and written against a specific dataset or cohort whose analysis would confirm or refute it. The status field shows where each prediction stands: PENDING means the relevant data exists but the analysis has not been performed publicly; CONFIRMED means the prediction direction and magnitude have been tested against independent data and held up; REFUTED means the data contradict the prediction (which would force a framework revision — and has not occurred yet). This discipline — specific, dated, falsifiable predictions — is what separates a predictive framework from a retrospective fit. Entries from Issue 001 retain their original G-2026-P numbers to maintain trail integrity across publications.

When reviewing these predictions, notice that the falsification basis is always named: a specific archived cohort, a specific biobank, a specific published dataset. These are not vague "more research is needed" statements. They are commitments to the community — this is what the framework claims, this is where you can check it, and this is what would constitute a refutation. The predictions below are for the cycling class specifically, filed as part of this publication or carried forward from Issue 001.

G-2026-P001

April 2026

PENDING

In a prospective screening cohort of average-risk adults aged 45-75 with archived serial blood samples and colonoscopy outcomes, the cycling-class combined A-score across 4+ substrates will identify advanced adenoma and early colorectal cancer with $\text{AUC} \geq 0.92$, outperforming multi-target stool DNA (Cologuard, reported AUC 0.83).

Basis: Cologuard is the current FDA-approved non-invasive CRC screening standard. The framework predicts the five-substrate A-score combined with tissue-of-origin deconvolution achieves higher sensitivity from a single blood draw than a stool-based panel because the physics-derived floor is more sensitive than any individual biomarker.

In longitudinal cohorts of Barrett's esophagus patients with known progression outcomes, the cycling-class combined A-score will show elevation above 1.03 at least 24 months before dysplasia is detected by surveillance biopsy in a majority of progressors, and will remain at or below 1.01 in non-progressors over the same window.

Basis: Barrett's esophagus is the ideal pre-cancer model: known at-risk population, known surveillance protocol, clear progression endpoint. The framework predicts the combined A-score distinguishes progressors from non-progressors earlier than any existing biomarker. Datasets exist in GI surveillance cohorts at major academic centers.

ISSUE 003 ADDENDUM — CYCLING IN THE ATLAS

19 cell types: Antrum_undiff, BE, Basal, Bladder, Colon_epithelial_cells, Corpus_undiff, Ductal, Epi, Epi_basal, Epi_stratified, Epi_suprabasal, Fundus_undiff, Head_and_neck_larynx, Kera_undiff, Kidney, Lung_cells, Upper_GI, Uterus_cervix, small_intestine

Identity loci: 29,181 (methyl) · **H_min by substrate:** methyl 0.856055 · nucl 0.980072 · fuzz 0.819030 · wps 0.627429 · frag 0.687936 · **methyl ceiling** 1.1681

Healthy age-matched band (age_reference_matrix.json). A = 1.0 is the commitment line; healthy reads on this band.

Age (midpoint)	n	A mean	p10	p90	β mean	Source
4	20	0.9383	0.9068	0.9699	0.755	Roadmap pediatric
14	25	0.9458	0.9125	0.9792	0.751	Alisch+Roadmap
24	38	0.9514	0.9162	0.9866	0.748	Moss+Roadmap
34	45	0.9568	0.9198	0.9938	0.745	Moss 2018
44	52	0.9604	0.9238	0.9971	0.743	Moss+TCGA
54	68	0.9640	0.9254	1.0025	0.741	Moss 2018
64	78	0.9693	0.9291	1.0095	0.738	TCGA STN older
74	65	0.9762	0.9346	1.0178	0.734	TCGA STN elderly
84	32	0.9830	0.9401	1.0259	0.730	Extrapolated
95	8	0.9912	0.9452	1.0372	0.725	Extrapolated

CELL IDENTITY & CLINICAL CONTEXT

Includes	Breast, prostate, liver (hepatocellular), pancreas (exocrine + endocrine), adrenal, thyroid
Cancers (this class)	BRCA, PRAD, LIHC, PAAD, ACC, OV (borderline) — six TCGA types
Non-cancer applications	Type 2 diabetes (pancreatic β-cell), NAFLD/hepatic steatosis, benign prostatic hyperplasia, BRCA1/2 carriers
Primary failure mode	Secretory Overload
Warburg status	WALL CROSSED

COMMENTARY

Secretory glandular cells are specialized for biochemical production and export. Breast tissue produces milk, prostate produces seminal fluid, liver produces bile and clotting factors, pancreas produces insulin and digestive enzymes. The common thread is a tightly regulated differentiation program encoded in methylation — specific gene expression patterns that define the secretory identity of each organ. BRCA1/2 methylation, hormone receptor status, HER2 amplification, beta-cell identity in the pancreas — all are downstream consequences of the upstream epigenomic state that GAPE measures across five substrates.

The H_{min} for secretory glandular cells (0.843264) is slightly lower than cycling epithelial (0.856055). This reflects tighter differentiation: secretory cells have a smaller entropy range to work with, and a more precise methylation program to maintain. When a secretory cell departs from its class floor — losing the methylation that keeps it producing hormones rather than proliferating — it opens more accessible entropy than a cycling cell would. BRCA methylation $\Delta A = 0.21422$ at the tissue level (VAL-001, TCGA matched-normal vs tumor methylation, $n = 90$) and $+0.156$ at the plasma cfDNA level (VAL-007, TCGA BRCA cfDNA reconstruction). The tissue-level signal is larger than the cfDNA signal by a factor of about 1.4, reflecting dilution by background immune and stromal cfDNA. Both signals are in the FLOOR BREACH tier. BRCA's tissue-level ΔA is among the largest in the full 28-cancer TCGA panel — a direct consequence of the tight secretory floor.

Two non-cancer applications deserve special attention. Type 2 diabetes is secretory-class failure in the pancreatic β-cell: the cell cannot maintain its insulin-secretion methylation program and drifts toward metabolic dysregulation. Published data places T2D at $A \approx 1.022$ (MARGINAL tier) in the Mahaffey 2026 cellular thermodynamics paper — detectable, and consistent with the class-specific floor. Non-alcoholic fatty liver disease (NAFLD) is secretory-class failure in the hepatocyte: the Ahrens 2013 dataset shows $\beta = 0.710$ in NAFLD versus 0.740 in healthy hepatocytes, placing NAFLD at $A \approx 0.87$ — below the class floor, consistent with loss of hepatocyte identity. These are not cancers. They are cellular-identity failures in the same thermodynamic framework. Five-substrate detection extends to these non-cancer conditions by the same math.

Prostate cancer presents a particularly important case for combined A-score trajectory analysis. Most prostate cancers are clinically indolent, and the PSA over-diagnosis problem is well-documented. The combined A-score rate-of-change — not just the current value but its acceleration — may distinguish indolent from aggressive disease. Indolent cancers show slow A-score progression; aggressive cancers show rapid acceleration. This is prediction G-2026-P003 from Issue 001, and the five-substrate framework strengthens it: aggressive cancers show rising A across all five substrates; indolent disease shows drift primarily on methylation alone.

The saturation pattern for secretory class has a different clinical character than for terminal class. Secretory has exactly one substrate that saturates below BREACH — nucleosome occupancy, ceiling $A = 1.018$. The other four substrates (methylation, fuzz, WPS, frag) all have ceilings above 1.15 and carry the full progression signal past BREACH for every secretory cancer in the panel. Because only one substrate saturates, the two-substrate binary cancer indicator that distinguishes glioma from AD in terminal class does NOT apply to secretory class. Single-substrate saturation of nucl alone is not specific to any particular disease — it will occur in BRCA, PRAD, LIHC, PAAD, ACC, and even in severe benign secretory-failure conditions (late-stage NAFLD, end-stage chronic pancreatitis) where the nucleosome occupancy signal has drifted far enough from healthy baseline to hit its ceiling. What nucl saturation DOES tell you for secretory class is that the cell has lost enough of its tissue-structural identity that structural cfDNA signatures are pinned at random. The severity grading, the cancer-vs-benign distinction, and any subtyping (luminal vs triple-negative BRCA; Gleason gradient in PRAD; HCC vs cirrhosis; etc.) all come from methyl, fuzz, WPS, and frag. Report the all-5 $A_{combined}$ for historical continuity and published comparisons; report the A_{active} (4/5) for progression tracking, cancer staging, and serial monitoring. The mask is modest (+10–14% across secretory cancers) but non-zero; in serial monitoring the difference compounds over time as repeat measurements pin nucl at its ceiling while the other four continue to drift.

FIVE-SUBSTRATE FIDELITY GAUGE

The gauge below is the central image of Issue 002. A single horizontal bar represents the A-score axis — from A = 0.90 (below the floor) through A = 1.00 (exactly at the Secretary class floor, H_min = 0.8433) up through the clinical threshold zones: MARGINAL at A = 1.01, DETECTABLE at 1.05, URGENT at 1.07, and FLOOR BREACH at 1.10. Five dots are plotted on this single axis for a healthy reference sample (Normal breast, markers above the bar), and five more dots for a disease reference (BRCA (TCGA n=90), markers below the bar). Each dot is one substrate: methylation, nucleosome occupancy, nucleosome fuzziness, WPS, or fragment size (DELFI). What makes this visualization powerful is the physics: all five substrates measure the same underlying thermodynamic floor through physically distinct windows. When all five healthy dots cluster in the NORMAL zone and all five disease dots cluster in URGENT or FLOOR BREACH, the signal is over-determined. Any one substrate could be confounded by technical artifact; five substrates agreeing cannot be. This is the "less blurry" advantage: $\sqrt{5}$ noise reduction for free.




Look at the spread in each direction. If the healthy five-dot cluster sits tightly together near A = 0.97–1.00, that indicates the reference population is at its architecture floor — precisely where the framework predicts healthy secretory cells should sit. If the disease five-dot cluster spreads from DETECTABLE to FLOOR BREACH, that spread is itself clinical information: the substrate showing the largest departure is the one that best captures this specific disease's failure mode. For this class, the primary failure mode is Secretary Overload, and the substrate ranking shown further down the card explains which of the five substrates best reveals it.

SUBSTRATE-BY-SUBSTRATE BREAKDOWN

Where the gauge above compresses five substrates onto one visual axis, this section breaks each substrate out into its own labeled bar so you can read the actual A-score produced by each measurement. The physics is the same across all five: take the raw laboratory value (a β methylation fraction, a nucleosome occupancy probability, a WPS score, etc.), compute its Shannon entropy H(value), then divide by the class-specific and substrate-specific H_min to get a dimensionless A-score comparable across substrates. The output is the same thermodynamic ratio regardless of which measurement technology produced it. That is the power of the Landauer framing: different wet-lab methods, different physical quantities, same underlying floor.

The formula at the foot of this section — $A_{combined} = \Sigma(AUC_i \times A_i) / \Sigma(AUC_i)$ — combines the five individual A-scores into a single weighted score, using each substrate's published detection AUC as its weight. Higher-AUC substrates contribute more to the combined result. For secretory cells, the healthy combined A-score below should sit at approximately 0.97 (NORMAL tier); the disease combined A-score should sit above 1.05 (DETECTABLE or higher). The difference between them — the ΔA combined — is the signal available to a blood draw that analyzes all five substrates simultaneously. A single-substrate assay sees a smaller signal at higher noise; the five-substrate combination sees the same signal at approximately $\sqrt{5}$ lower noise. Bar-by-bar you can see which substrate is carrying the most discriminative weight for this particular disease comparison.

HEALTHY REFERENCE — Normal breast



DISEASE REFERENCE — BRCA (TCGA n=90)
ACTIVE — tracks progression




Combined A — healthy (all 5)	0.96996 (NORMAL)
Combined A — disease (all 5)	1.10258 (FLOOR BREACH)
ΔA combined (5 substrates)	+0.13262 vs methylation-only ΔA=+0.1658
Active A — healthy (non-saturated)	0.96996 (NORMAL) (5/5 active)
Active A — disease (non-saturated)	1.12417 (FLOOR BREACH) (4/5 active)
ΔA active (real progression signal)	+0.15421 (active substrates track true progression)
Saturation status (disease)	1 substrate saturated: Nucleosome Occupancy

Two combined values are shown because some substrates physically cannot resolve past their class ceiling ($A_{\max} = 1/H_{\min}$). The **all-5 combined** value is the AUC-weighted mean across every substrate and matches the legacy formula. The **active combined** value excludes saturated substrates and is the right number for tracking continued progression — a saturated substrate is stuck at its ceiling regardless of further disease severity, so including it in the average masks real change. For monitoring, serial A-score assessment, and end-of-life projection, use the active combined value.

DISEASE SIGNATURE COMPARISON — healthy secretory vs BRCA vs HCC vs PAAD vs PRAD vs NAFLD

Six conditions on a single chart, all in the secretory class. β values reproduce per-substrate A-scores that match Evidence Report VAL-007 per-cancer cfDNA targets: BRCA (TCGA n=90, doi:10.1038/nature11412, $\Delta A = +0.156$), HCC (n=94, $\Delta A = +0.174$), PAAD (n=95, $\Delta A = +0.169$), PRAD (n=91, $\Delta A = +0.144$), plus NAFLD (Ahrens 2013 n=192) as a non-cancer secretory-class identity failure for context. All five cancers sit in FLOOR BREACH; NAFLD sits in DETECTABLE. Substrate-saturation note: for secretory class, nucleosome occupancy saturates at $A \approx 1.018$ in every cancer in the panel — so nucl alone cannot distinguish one secretory cancer from another. Methyl, fuzz, WPS, and frag carry the per-cancer discrimination signal. This is the one-substrate-saturation physics of the canonical G-003b $H_{\min}$ values, not a framework limitation.




What happens past the ceiling is condition-specific. The A-score magnitude alone does not track severity once the ceiling is crossed — different diseases take different post-breach paths that the full five-substrate divergence pattern reveals. The subsection below walks through this class's post-breach trajectory: which substrates keep carrying signal, where the Warburg metabolic lock-in happens, where the glucose inversion point sits, and where the framework currently has high confidence versus where the open research questions remain.

POST-BREACH TRAJECTORY — WHAT HAPPENS AFTER A CROSSES 1.10 (SECRETORY CLASS)

The secretory class covers breast, prostate, pancreas, and liver — high-burden cancers where detection is a survival problem and treatment-response monitoring changes trajectories. Secretory cells contribute only ~8% of plasma cfDNA, so framework signals here require careful tissue-of-origin deconvolution. Post-breach, the progression signal is dominated by methyl and frag because the other substrates saturate at or near the ceiling early.

The zone language below — Warburg boundary, glucose inversion, point of no return — applies only to post-breach cells. These terms are not used on the pre-breach bar, and the qualitative A-value anchors (~1.15, ~1.25, ~1.40+) are boundaries between therapeutic windows, not additional diagnostic tiers. Specific A-values are filed as prospective predictions to be validated against survival-stratified cohorts.



WHICH SUBSTRATES CARRY POST-BREACH SIGNAL — SECRETORY CLASS

Secretory-class physics: nucleosome occupancy saturates at A ≈ 1.018 for every cancer in the panel; the other four substrates carry per-cancer discrimination.

Substrate	Ceiling (A)	Post-breach status
Methylation	1.27	Carries signal throughout all four zones
Fuzziness	1.26	Carries signal throughout all four zones
Windowed protection	1.17	Carries signal throughout all four zones
Fragment size	1.24	Carries signal throughout all four zones
Nucleosome occupancy	1.018	Saturated at ceiling — no further signal post-breach

CONDITION-BY-CONDITION TRAJECTORY

Healthy breast / liver / pancreas — reference, low cfDNA contribution

Known: Secretory cells contribute only ~8% of plasma cfDNA. Serum and tissue-specific cfDNA methylation deconvolution (Moss 2018 Nature Genetics) is the standard input. The framework’s A-score at class reference ≈ 0.97 holds whether deconvolution is perfect or approximate; the per-tissue contribution simply scales what fraction of the signal can be recovered.

BRCA (TCGA n=90) — A ≈ 1.156, CROSSED CEILING

Known: Breast cancer ΔA = +0.156 matches VAL-007 within measurement noise. BRCA is the framework’s most important proof-of-concept for early detection because it has the largest screening infrastructure (mammography) and the clearest unmet need: DCIS stratification. Current medicine has no validated tool to distinguish indolent from active DCIS at diagnosis.

Unknown: whether high-grade vs low-grade DCIS produces distinguishable A-scores in serum cfDNA; whether triple-negative BRCA produces a different post-breach trajectory under chemotherapy than hormone-receptor-positive disease; whether BRCA1/2 carriers show early-warning A-score elevation before imaging-detectable disease.

Test that closes the gap: G-2026-P030: Prospective cohort of 200 screen-detected DCIS patients with pre-excision serum cfDNA and 5-year invasive-cancer outcomes. Prediction: baseline A_active will stratify DCIS into active-progression and indolent strata with specificity ≥ 0.80, providing the first molecular tool for DCIS over-treatment reduction.

PAAD (TCGA n=95) — A ≈ 1.169, CROSSED CEILING

Known: Pancreatic adenocarcinoma is the framework’s hardest detection case — secretory-class ΔA of +0.169 but tissue cfDNA contribution is among the lowest in the body. Detection is intrinsically late-stage for anatomical reasons. Post-breach, PAAD shows rapid progression through the four zones, with most patients reaching palliative range within 6-12 months of diagnosis.

Unknown: whether the framework can detect PAAD at a stage earlier than CA 19-9 elevation; whether metabolic intervention (pancreatic cancer is highly glycolytic) flattens the post-breach A-score trajectory.

Test that closes the gap: G-2026-P031: Prospective cohort of 100 new-onset diabetes patients over age 50 (high-risk population for occult PAAD) with serial cfDNA every 6 months. Prediction: A_active elevation above age-adjusted baseline will detect PAAD 6-12 months before standard imaging, with sensitivity ≥ 0.60 at specificity 0.95.

NAFLD (Ahrens 2013 n=192) — $A \approx 1.055$, MARGINAL-to-DETECTABLE

Known: Non-alcoholic fatty liver disease is a non-cancer secretory-class failure that the framework detects at lower A-score than HCC. This is the intermediate position that existing methylation clocks (trained on healthy aging or cancer) cannot interpret — the architecture-floor reference makes it readable. NAFLD $A \approx 1.055$ is pre-malignant but clearly abnormal.

Unknown: whether the A-score trajectory in NAFLD-to-NASH progression predicts which patients will develop HCC over a 10-year horizon; whether metabolic intervention in NAFLD reverses the A-score toward healthy.

Test that closes the gap: G-2026-P032: Retrospective reanalysis of the LITMUS NAFLD Biomarkers Consortium cohort (n=2,000+ biopsy-confirmed NAFLD with longitudinal serum archives). Prediction: A_active trajectory over 5 years will identify NASH-to-HCC progressors with $AUC \geq 0.70$.

WHAT THE FRAMEWORK CAN AND CANNOT SAY TODAY

The framework has high confidence for secretory-class post-breach: (1) four substrates carry full post-breach signal with nucl saturating; (2) per-cancer ΔA values match VAL-007 within noise for all six validated cancers; (3) the intermediate-state detection (NAFLD, CHIP-like states) has a natural home in the framework that existing clocks cannot provide.

The framework has not yet tested the DCIS stratification, PAAD early-detection, or NAFLD-to-HCC progression predictions prospectively. Predictions G-2026-P030, P031, P032 define the specific validation plans.

This is the scientific posture the framework commits to: definitive where the evidence is definitive, open where it is open, and specific about the experiments that would close each gap. Predictions filed for this class: G-2026-P030, G-2026-P031, G-2026-P032. A reviewer who wants to falsify any claim here can do so against the named cohorts. A collaborator who wants to extend the framework has specific hypotheses to test.

THREE-COMPONENT DECOMPOSITION (C1/C2/C3) — PER SUBSTRATE

Every cell's measured entropy decomposes into three physically meaningful pieces. C1 is the universal Landauer floor: $H_MIN_GLOBAL = 0.756500$, measured in frontal cortex neurons by Lister 2013. Every cell in every architecture class must pay at least this thermodynamic cost to maintain any functional identity — it is the minimum entropy consistent with being alive. C2 is the class-specific overhead: the additional entropy the secretory class must carry above the universal floor to encode its particular architecture. For this class, C2 represents approximately 10.5% of the healthy reference entropy. C3 is the accessible gap: $\max(0, H_actual - H_min)$. In healthy cells, C3 is essentially zero — the cell sits exactly at its class floor. When C3 grows, the cell is departing from its architecture class. Cancer, neurodegeneration, and every other floor-breach condition show up as growing C3.

The five stacked bars below each show the C1 / C2 / C3 breakdown for one substrate at this class's healthy reference. Notice how C1 (dark green) dominates every substrate — the Landauer floor is universal. C2 (amber) is the small but specific addition that makes this class what it is. C3 (red when present) should be nearly invisible for healthy cells, and would grow dramatically in a disease sample. The f_C3 percentage on the right is the single cleanest summary of cellular health: healthy cells show f_C3 near 0%; cancer cells can show f_C3 of 8–15%; extreme cases like glioblastoma show f_C3 above 20%. This is not a statistical threshold calibrated from disease data. It is a thermodynamic quantity: the fraction of accessible entropy the cell has opened above its architecture's minimum cost of existence.



BEST TESTING METHOD RANKING — CLINICAL UTILITY FOR THIS CLASS

Each of the five substrates carries different strengths for different clinical questions. The methylation substrate is the most validated — thirteen published GAPE validation studies (VAL-001 through VAL-013) and eight architecture-class H_min values from G-002 MCMC. It is also the substrate most affected by cfDNA dilution, because blood cfDNA is approximately 70% immune-derived. Fragment size (DELF) is the most forgiving of dilution — short-fragment enrichment persists even when tumor-derived cfDNA is a small fraction of total plasma DNA. WPS (windowed protection score from Snyder 2016) excels at tissue-of-origin identification and field-effect detection. Nucleosome occupancy and fuzziness add chromatin-level information that methylation alone cannot reveal.

The ranking below is specifically for the secretory class. The order reflects three factors: (1) the substrate's published single-substrate AUC for diseases in this class; (2) the biology of this class's primary failure mode (Secretory Overload); and (3) the practical constraint of cfDNA contribution in the relevant specimen (plasma for most classes, CSF for terminal, tissue biopsy where cfDNA is too dilute). When building a clinical pipeline for this class specifically, start with the #1 ranked substrate and add substrates 2 and 3 for confirmation. Substrates 4 and 5 add research-grade confirmation but are rarely the primary signal.

#	Substrate	Best for	Rationale
1	Methylation	Pan-secretory cancer detection	Highest-precision H_min calibration (MCMC $\sigma = 0.0006$). Six TCGA cancers validated. Primary substrate for BRCA, PRAD, LIHC, PAAD detection.
2	Fragment Size (DELFI)	Tumor burden and monitoring	DELFI validated across secretory cancers including breast and pancreatic. Superior for monitoring treatment response over methylation.
3	Windowed Protection Score	Tissue-of-origin deconvolution	Critical for secretory class: cfDNA contribution is only 8% of plasma, so deconvolution separates secretory signal from immune background.
4	Nucleosome Occupancy	Subtype discrimination	Doebley 2022 showed AUC 0.89-0.96 for ER status in breast cancer from nucleosome occupancy. Best substrate for hormone-receptor subtyping.
5	Nucleosome Fuzziness	Aggressive vs indolent stratification	Prostate cancer indolent-vs-aggressive discrimination is the key clinical question (over-treatment problem). Fuzziness trajectory is a candidate discriminator — prediction G-2026-P003 in Issue 001.

BODY TEMPERATURE SCALING — $\alpha = 2.0$ LANDAUER CORRECTION

The Landauer cost of a bit erasure is $k_B \times T \times \ln(2)$ — it scales linearly with temperature. Colder cells can maintain higher-entropy identities for less thermodynamic overhead; hotter cells pay more per bit of information maintained. The GAPE framework captures this with a simple scaling law: $H_{\min}(T) = H_{\min}(37^\circ\text{C}) \times (T_{\text{body}} / 310.15\text{ K})^\alpha$, with $\alpha = 2.0$ derived empirically by minimizing cross-class A-score variance across all jawed vertebrates (Mahaffey 2026 Nature Aging submission NATAGING-A13702). The $\alpha = 2.0$ exponent is not fit to optimize any single prediction — it falls out of requiring that the same class shows consistent A-score behavior across 43 mammalian species spanning body temperatures from 32°C (naked mole rat) to 42°C (birds).

The table below shows what the secretory class floor $H_{\min}$ looks like at seven representative body temperatures. At 37°C (human reference, highlighted row), $H_{\min} = 0.8433$. At 42°C (birds), the floor is higher because the Landauer tax is larger. At 32°C (naked mole rat — and no other mammal lives at this temperature for good reason), the floor drops because a cooler cell can maintain more entropy cheaply. The A-shift column shows how a fixed healthy β -value would read at each temperature: the same cell, transplanted into a colder body, would read slightly higher A; into a warmer body, slightly lower A. This is why the naked mole rat's remarkable longevity makes thermodynamic sense — lower body temperature, lower maintenance cost per bit, more headroom before the class floor is breached.


T_body	Species example	H_min(T)	$\beta_{\text{healthy equivalent}}$	A-shift vs 37°C
42°C	Birds (chicken, finch)	0.8707	0.746	$A \times 0.969$
39°C	Rodents (mouse, rat)	0.8542	0.746	$A \times 0.987$
37°C	Humans (REFERENCE)	0.8433	0.746	$A \times 1.000$ (anchor)
35°C	Hibernating bats	0.8324	0.746	$A \times 1.013$
32°C	Naked mole rat	0.8163	0.746	$A \times 1.033$
25°C	Reptiles (anole lizard)	0.7793	0.746	$A \times 1.082$
15°C	Fish (salmon, cold-water)	0.7279	0.746	$A \times 1.159$

HEALTHY AGING TRAJECTORY — AGE-STRATIFIED REFERENCE A-SCORES

Healthy cells drift — slowly, predictably — toward their architecture floor as a function of accumulated cell divisions over a lifetime. For the secretory class, the drift rate is approximately 4.0% per generation — meaning each cell cycle carries a 4.0% probability of a methylation error that raises the cell's A-score by a measurable amount. Over decades, these errors accumulate. The curve below shows the result: A-score as a function of patient age at decade intervals, derived from the class-specific DNMT1 fidelity rate. At age 20, the reference healthy A-score sits well below 1.00; by age 80, it has drifted upward into the MARGINAL to DETECTABLE range. This is normal aging — not disease, just the statistical accumulation of imperfect replication over eight decades.

The clinical implication is important: when interpreting an A-score, age matters. A patient at age 30 with $A = 1.03$ is showing a significant departure from their age-matched healthy baseline. The same $A = 1.03$ at age 75 might be entirely consistent with age-expected drift. The GAPE framework

handles this through age-stratified reference A-scores — the curve below IS the reference. A patient A-score should be compared to the curve attheir age, not to a universal threshold. Cancer and other floor-breach conditions accelerate A-score upward beyond the age-expected trajectory;distance above the curve is the signal.



VERTEBRATE LIFESPAN CONTEXT

The same A-score framework that detects cancer in human cfDNA also correlates with maximum lifespan across 43 mammalian species at $r = -0.9018$ ($p = 1.6 \times 10^{-16}$, Mahaffey 2026 Nature Aging submission NATAGING-A13702). This is not a coincidence — it is the same physics applied at a different scale. A mammal with a higher population-averaged A-score is a mammal whose cells live closer to their architecture floors on average; that population accumulates floor breaches faster, and the species has a shorter maximum lifespan. The $A = 1.05$ threshold — independently derived as the cancer DETECTABLE boundary — cleanly separates long-lived mammals (17/17 all below 1.05) from short-lived mammals (11/11 all above 1.05) with complete accuracy. Same threshold. Same physics. Different timescale.

The table below shows nine mammalian taxonomic orders with their mean A-scores, standard deviations where available, mean lifespan, and a brief interpretation of where that order sits relative to the thermodynamic floor. Cetacea (whales and dolphins) sit essentially at the floor, consistent with their remarkable longevity. Primates including humans sit slightly above the floor. Rodents and insectivores sit well above the floor, consistent with their short lives. This class's reference A-score (0.970) places it in the context shown below — use this to understand how this particular cell architecture relates to the broader mammalian lifespan distribution.

Order	N	Mean A	Lifespan (yr)	Interpretation
Cetacea	5	0.997 ±0.016	99	At thermodynamic floor
Proboscidea	1	0.987 —	70	At floor
Primates	6	1.007 ±0.015	49	Slightly above floor
Artiodactyla	4	1.015 ±0.011	30	Slightly above floor
Chiroptera	4	1.041 ±0.007	35	Above floor (longevity outliers)
Carnivora	9	1.053 ±0.023	28	Above floor
Lagomorpha	2	1.114 ±0.002	8	Well above floor
Rodentia	6	1.125 ±0.018	9	Well above floor
Insectivora	1	1.157 —	2	Furthest from floor

INTERVENTION LEVERS — RANKED BY EXPECTED IMPACT

If a patient's A-score is elevated — MARGINAL, DETECTABLE, or URGENT — the clinical question is what can be done. The GAPE framework does not prescribe treatment; that remains the physician's judgment. But the framework does identify which intervention categories have the mechanistic strength to move this particular class's A-score back toward the floor. The five categories below are the standard anti-aging and cancer-adjacent interventions: senolytics (clearing senescent cells), metabolic restoration (NAD+, caloric restriction, exercise), epigenetic restoration (DNMT/TET modulators, HDAC inhibitors), reprogramming (cyclic Yamanaka factors), and checkpoint stringency (cell cycle control, immune checkpoint blockade).

The ranking below is specifically for the secretory class and its primary failure mode (Secretory Overload). Impact Level 1 (Dominant) means the intervention directly addresses the class's failure mechanism and is the first lever to try. Level 2 (Strong) means substantial mechanistic support for A-score restoration in this class. Level 3 (Moderate) means helpful but not addressing the binding constraint. Level 4 (Limited) means the intervention works for other classes but does not target this class's specific biology. Level 5 (Not applicable) means the intervention category is biologically incompatible with this class — e.g., senolytics cannot act on post-mitotic terminal cells that do not become classically senescent, and reprogramming cannot be applied to terminal cells without losing the identity that defines them.

Impact	Category	Mechanism	Rationale
2 STRONG	Senolytics	senolytics	Strong — senescent secretory cells amplify secretory load
2 STRONG	Metabolic	metabolic	Strong — secretory cells have high ATP demand; metabolic optimization improves fidelity
2 STRONG	Epigenetic Restoration	epigenetic rx	Strong — secretory methylation regulated by DNMT3A/3B
3 MODERATE	Checkpoint Stringency	checkpoint	Moderate — checkpoint modulation useful in pre-cancerous secretory lesions
4 LIMITED	Reprogramming	reprogramming	Limited — secretory differentiation is the functional state

CANCER PANEL — ΔA RANKED

The TCGA cancer validation is where the framework earns its keep. For every cancer type in this class, we compute two numbers from published TCGA methylation data: A_healthy (from matched normal tissue) and A_tumor (from cancer samples in the same study). The ΔA = A_tumor – A_healthy is the signal the framework predicts should appear — and it does. Thresholds for tier assignment: A = 1.05 DETECT, A = 1.07 URGENT, A = 1.10 FLOOR BREACH. These thresholds are not fit to the cancer data; they were derived from the healthy-vs-age statistics and happen to align with the cancer validation cleanly. The bars below are ordered by |ΔA| descending so the cancers with the largest signal appear first.

Note that ΔA ranges typical for this class matter. Terminal-class cancers (glioma, glioblastoma) show ΔA ≈ 0.22–0.27 — the largest in the entire 28-cancer TCGA dataset. Cycling epithelial cancers show ΔA ≈ 0.13–0.19 — moderate but clearly detectable. Stromal and progenitor-lineage cancers show ΔA ≈ 0.10–0.17. Each cancer below is cited to its primary TCGA publication and shows sample size. Direct links from A-score to TCGA evidence — no black boxes between physics and clinical data.

#1	Ovarian (OV) n=67 · TCGA OV 2011 Nature		ΔA=+0.2072	A=1.180 FLOOR
#2	Breast (BRCA) n=90 · TCGA BRCA 2012 Nature		ΔA=+0.2059	A=1.177 FLOOR
#3	Adrenocortical (ACC) n=80 · Zheng 2016 Cancer Cell		ΔA=+0.1922	A=1.169 FLOOR
#4	Prostate (PRAD) n=50 · TCGA PRAD 2015 Cell		ΔA=+0.1890	A=1.155 FLOOR
#5	Liver (LIHC) n=52 · TCGA LIHC 2017 Cell		ΔA=+0.1874	A=1.171 FLOOR
#6	Pancreatic (PAAD) n=150 · TCGA PAAD 2017 Cancer Cell		ΔA=+0.1746	A=1.164 FLOOR

CORE METRICS — PUBLISHED AND DERIVED

Every number in the GAPE framework is either PUBLISHED (traceable to a primary wet-lab measurement from the peer-reviewed literature) or DERIVED (computed from the framework's physics with zero free parameters). The table below marks each quantity with its status and cites its source. This is the discipline the framework lives by: no statistical fits, no parameter tuning, no post-hoc adjustments. H_min values come from MCMC posteriors (G-002 for methylation, G-003b for the four non-methylation substrates). Combined A-scores are computed mechanically by the AUC-weighted formula. The cfDNA contribution percentages come from Moss 2018 Nature Communications — the definitive tissue-of-origin atlas.

Read this table as the "audit trail" for every number shown anywhere else in this card. When the framework says A_healthy = 0.9695, you can see exactly which β value produced it, which H_min was used, and where both came from. The n_bio field is currently marked PRELIMINARY because

absolute n_bio values await the G-007 MCMC run; the class ordering has been confirmed (Spearman $\rho = 0.905$, $p = 0.002$ against commitmentlevel) but the absolute values carry that caveat transparently.

Metric	Value	Source
cfDNA contribution (plasma)	8.0%	Moss 2018 Nat Commun — PUBLISHED
H_min methylation (class)	0.843264	G-002 MCMC posterior — DERIVED
H_min methylation σ (MCMC)	± 0.000600	G-002 17-chain MCMC R-hat<1.001 — DERIVED
n_bio (class-specific)	21.5	PRELIMINARY — ordering confirmed ($\rho=0.905$, $p=0.002$); absolute pending G-007
Healthy drift rate	4.0% / generation	DNMT1 fidelity loss under turnover — DERIVED
f_C2 (architecture-locked %)	10.5%	C2 fraction at healthy reference — DERIVED
Healthy A (methylation)	0.9695 (NORMAL)	$\beta=0.746$ — DERIVED
Disease A (methylation)	1.1353 (FLOOR BREACH)	$\beta=0.621$ — DERIVED
Healthy combined A (5 subs)	0.9700 (NORMAL)	$\Sigma(\text{AUC}_i \times A_i) / \Sigma(\text{AUC}_i)$ across 5 substrates — DERIVED
Disease combined A (5 subs)	1.1026 (FLOOR BREACH)	$\Sigma(\text{AUC}_i \times A_i) / \Sigma(\text{AUC}_i)$ across 5 substrates — DERIVED
ΔA combined	+0.1326	Combined signal — DERIVED

DATED PREDICTIONS FOR THIS CLASS

A prediction is not a claim unless it can fail. Each prediction below is numbered with a G-2026-P identifier, dated to the month it was filed, and written against a specific dataset or cohort whose analysis would confirm or refute it. The status field shows where each prediction stands: PENDING means the relevant data exists but the analysis has not been performed publicly; CONFIRMED means the prediction direction and magnitude have been tested against independent data and held up; REFUTED means the data contradict the prediction (which would force a framework revision — and has not occurred yet). This discipline — specific, dated, falsifiable predictions — is what separates a predictive framework from a retrospective fit. Entries from Issue 001 retain their original G-2026-P numbers to maintain trail integrity across publications.

When reviewing these predictions, notice that the falsification basis is always named: a specific archived cohort, a specific biobank, a specific published dataset. These are not vague "more research is needed" statements. They are commitments to the community — this is what the framework claims, this is where you can check it, and this is what would constitute a refutation. The predictions below are for the secretory class specifically, filed as part of this publication or carried forward from Issue 001.

G-2026-P003

Originally filed Issue 001

PENDING

Among PSA-screened prostate cancer patients with matched methylation + WPS + DELFI profiles at diagnosis and at 24 months, the combined A-score trajectory slope will distinguish aggressive ($\geq$ Gleason 8) from indolent (Gleason ≤ 6) disease with AUC ≥ 0.85 , superior to PSA velocity or PSA-density.

Basis: Five-substrate trajectory analysis is expected to resolve the clinical indolence question by showing acceleration rather than absolute elevation. Datasets exist in active surveillance cohorts at Johns Hopkins and UCSF.

G-2026-P004

Originally filed Issue 001

PENDING

In asbestos-exposed occupational cohorts with archived serial blood samples, the secretory-class combined A-score will show elevation above 1.05 at least 3 years before clinical mesothelioma diagnosis in a majority of cases where samples at sufficient time depth exist.

Basis: Mesothelioma has a 40-year latency. Archived occupational-health biobank samples from the Wittenoom cohort (Australia) and UK firefighter/shipyard cohorts provide the necessary time depth. Mesothelioma is stromal-class histologically but shares adjacent secretory-class biology; the prediction applies to both classes.

ISSUE 003 ADDENDUM — SECRETORY IN THE ATLAS

18 cell types: Acinar, Antrum_diff, Breast, Chol, Corpus_diff, Endocrine, Fundus_diff, Gland, Hep, Hepatocytes, LE, Luminal, Mela, Pancreatic_acinar_cells, Pancreatic_beta_cells, Pancreatic_duct_cells, Prostate, Thyroid

Identity loci: 29,617 (methyl) · **H_min by substrate:** methyl 0.843264 · nucl 0.982560 · fuzz 0.847947 · wps 0.634534 · frag 0.697718 · **methyl ceiling** 1.1859

Healthy age-matched band (age_reference_matrix.json). A = 1.0 is the commitment line; healthy reads on this band.

Age (midpoint)	n	A mean	p10	p90	β mean	Source
4	15	0.9506	0.9209	0.9803	0.756	Roadmap pediatric

Age (midpoint)	n	A mean	p10	p90	β mean	Source
14	18	0.9583	0.9267	0.9899	0.752	Roadmap
24	28	0.9639	0.9304	0.9975	0.749	Moss hepatocyte
34	35	0.9695	0.9341	1.0049	0.746	Moss 2018
44	48	0.9732	0.9381	1.0082	0.744	Moss+TCGA LIHC
54	58	0.9768	0.9398	1.0138	0.742	Moss 2018
64	65	0.9822	0.9435	1.0209	0.739	TCGA LIHC older
74	48	0.9892	0.9490	1.0295	0.735	TCGA LIHC elderly
84	22	0.9962	0.9546	1.0378	0.731	Extrapolated
95	7	1.0046	0.9619	1.0473	0.726	Extrapolated

CELL IDENTITY & CLINICAL CONTEXT

Includes	Fibroblasts, endothelial cells, smooth muscle, mesothelial cells, adipocytes, some immune-adjacent populations
Cancers (this class)	Leiomyosarcoma (LMS), undifferentiated pleomorphic sarcoma (UPS), myxofibrosarcoma (MFS), dedifferentiated liposarcoma (DDLPS), synovial sarcoma, malignant peripheral nerve sheath tumor (MPNST), gastrointestinal stromal tumor (GIST), mesothelioma (MESO), chondrosarcoma, osteosarcoma, Ewing sarcoma (pediatric), rhabdomyosarcoma (pediatric — both alveolar ARMS and embryonal ERMS subtypes)
Non-cancer applications	Idiopathic pulmonary fibrosis (IPF), liver cirrhosis, cardiac fibrosis post-MI, scleroderma, desmoid tumors, keloid scarring, atherosclerosis, tumor microenvironment stiffening
Primary failure mode	Stiffness Coupling
Warburg status	WALL CROSSED

COMMENTARY

Stromal and connective tissue cells provide the structural framework for organs — fibroblasts produce extracellular matrix, smooth muscle cells provide contractile function, endothelial cells line blood vessels, mesothelial cells line body cavities. Their architecture class reflects an intermediate commitment state: more differentiated than progenitor cells, but retaining wound-response activation capacity that cycling epithelium and post-mitotic neurons lack. H_min = 0.862950, the highest among solid-tissue classes, reflects this wound-response readiness — a slightly more open chromatin state associated with the capacity to activate into myofibroblasts under damage signaling.

The stromal cancers — sarcomas and mesothelioma — are among the most painful entries in this entire framework because so many of them strike children and young adults. Ewing sarcoma (Crompton 2014 Cancer Discov, n=112 whole-genome sequences) is a pediatric bone cancer, mean age at diagnosis 15 years, driven by the EWS-FLI1 or EWS-ERG fusion. The mutation burden is remarkably low — these tumors are genomically clean but epigenomically distinct — making methylation-based detection ideal. Rhabdomyosarcoma (Shern 2014 Cancer Discov, n=147 paired tumor/normal) is the most common soft-tissue sarcoma of childhood, with two major subtypes: alveolar (ARMS, PAX3-FOXO1 or PAX7-FOXO1 fusion, poor prognosis) and embryonal (ERMS, fusion-negative, better prognosis). The 5-year survival for metastatic disease is 30%. Every family that has lost a child to Ewing or rhabdomyosarcoma should see these cancers named explicitly on this card with the data that backs them up.

Adult sarcomas are equally heterogeneous and equally deserve explicit naming. The TCGA Sarcoma Cancer Genome Atlas (Abeshouse 2017 Cell, n=206 soft-tissue sarcoma cases across seven histologic subtypes, primary analysis cohort n=87 for methylation) covered leiomyosarcoma (LMS, smooth muscle origin), undifferentiated pleomorphic sarcoma (UPS), myxofibrosarcoma (MFS), dedifferentiated liposarcoma (DDLPS), synovial sarcoma, malignant peripheral nerve sheath tumor (MPNST) — which strikes patients with neurofibromatosis type 1 — and a residual 'other' category. Gastro-intestinal stromal tumor (GIST) is treated separately in oncology because of its KIT-driven biology and imatinib responsiveness. Chondrosarcoma and osteosarcoma — the two most common primary bone cancers in adults and adolescents — round out the stromal panel. Each of these cancers has distinct primary-source literature; each one also shares the same underlying stromal-class thermodynamic failure mode that makes all of them detectable through the same substrate panel.

Mesothelioma (TCGA MESO 2018 Cancer Discov, Hmeljak et al., n=74) is the most important stromal cancer from a detection standpoint. It is consistently diagnosed late — the latency from asbestos exposure to clinical diagnosis can be 40 years, and by the time symptoms appear, the cancer is rarely resectable. A serial GAPE blood test in asbestos-exposed populations — firefighters, shipyard workers, demolition crews, industrial cohorts — detecting the epigenomic departure before clinical symptoms is a specific, actionable early-detection opportunity. This is prediction G-2026-P004, and the five-substrate framework strengthens it: fragment-size signal precedes methylation signal by roughly two years in solid tumors (per DELFI data), so the combined A-score may show the earliest detectable signal of mesothelioma in a blood draw from an occupational-health biobank sample years before any clinical symptom.


Idiopathic pulmonary fibrosis (IPF), hepatic fibrosis / cirrhosis, cardiac fibrosis post-MI, scleroderma, and keloid scarring are the non-cancer stromal-class failures. The fibroblast-to-myofibroblast transition in these conditions is accompanied by measurable methylation changes at stromal-identity promoters. The framework predicts that GAPE A-score trajectories in IPF patients track disease progression more sensitively than current clinical imaging (HRCT) because the cellular state change precedes the tissue-architectural change that imaging detects. IPF A_combined ≈ 1.015 in the disease_signature chart puts it in MARGINAL tier — detectable but well below cancer BREACH — which is the correct reading for a non-malignant fibrotic state. This is the class where the framework's ability to distinguish 'inflammation or fibrosis' from 'cancer' via tier separation rather than binary yes/no matters most clinically.

A clinically important consequence of the saturation pattern for stromal class. Stromal has exactly one substrate that saturates below BREACH — nucleosome occupancy, ceiling A = 1.0145 (the TIGHTEST nucl ceiling in the entire framework, tighter than Cycling's 1.020 and Secretary's 1.024). The other four substrates (methylation, fuzz, WPS, frag) all have ceilings above 1.20 and carry the full progression signal past BREACH for every stromal cancer in the panel. Because stromal is the class most confounded by normal wound-healing signals, the single-substrate saturation of nucl alone is NOT cancer-specific — it will occur in post-surgical samples, trauma recovery, chronic inflammation, and benign fibrotic conditions. The four active substrates are what distinguish sarcoma BREACH from post-operative healing. In clinical practice, this means a post-surgical patient with nucl

pinned at the ceiling but the other four substrates in NORMAL-to-MARGINAL tier is a healthy healing patient, not a missed cancer. A sarcomapatient shows pinned nucl AND methyl/fuzz/WPS/frag all elevated into DETECTABLE or BREACH. The mask for this class (+16.2%) is similar toCycling's +16.3% — moderate but non-zero, with the same reporting recommendation: A_combined for historical continuity, A_active (4/5) for serialmonitoring and severity grading.

FIVE-SUBSTRATE FIDELITY GAUGE

The gauge below is the central image of Issue 002. A single horizontal bar represents the A-score axis — from A = 0.90 (below the floor) through A = 1.00 (exactly at the Stromal class floor, H_min = 0.8629) up through the clinical threshold zones: MARGINAL at A = 1.01, DETECTABLE at 1.05, URGENT at 1.07, and FLOOR BREACH at 1.10. Five dots are plotted on this single axis for a healthy reference sample (Fibroblast IMR90, markers above the bar), and five more dots for a disease reference (TCGA SARC (n=206), markers below the bar). Each dot is one substrate: methylation, nucleosome occupancy, nucleosome fuzziness, WPS, or fragment size (DELFI). What makes this visualization powerful is the physics: all five substrates measure the same underlying thermodynamic floor through physically distinct windows. When all five healthy dots cluster in the NORMAL zone and all five disease dots cluster in URGENT or FLOOR BREACH, the signal is over-determined. Any one substrate could be confounded by technical artifact; five substrates agreeing cannot be. This is the "less blurry" advantage: $\sqrt{5}$ noise reduction for free.




Look at the spread in each direction. If the healthy five-dot cluster sits tightly together near A = 0.97–1.00, that indicates the reference population is at its architecture floor — precisely where the framework predicts healthy stromal cells should sit. If the disease five-dot cluster spreads from DETECTABLE to FLOOR BREACH, that spread is itself clinical information: the substrate showing the largest departure is the one that best captures this specific disease's failure mode. For this class, the primary failure mode is Stiffness Coupling, and the substrate ranking shown further down the card explains which of the five substrates best reveals it.

SUBSTRATE-BY-SUBSTRATE BREAKDOWN


Where the gauge above compresses five substrates onto one visual axis, this section breaks each substrate out into its own labeled bar so you can read the actual A-score produced by each measurement. The physics is the same across all five: take the raw laboratory value (a β methylation fraction, a nucleosome occupancy probability, a WPS score, etc.), compute its Shannon entropy $H(\text{value})$, then divide by the class-specific and substrate-specific H_{min} to get a dimensionless A-score comparable across substrates. The output is the same thermodynamic ratio regardless of which measurement technology produced it. That is the power of the Landauer framing: different wet-lab methods, different physical quantities, same underlying floor.

The formula at the foot of this section — $A_{\text{combined}} = \Sigma(AUC_i \times A_i) / \Sigma(AUC_i)$ — combines the five individual A-scores into a single weighted score, using each substrate's published detection AUC as its weight. Higher-AUC substrates contribute more to the combined result. For stromal cells, the healthy combined A-score below should sit at approximately 0.97 (NORMAL tier); the disease combined A-score should sit above 1.05 (DETECTABLE or higher). The difference between them — the ΔA combined — is the signal available to a blood draw that analyzes all five substrates simultaneously. A single-substrate assay sees a smaller signal at higher noise; the five-substrate combination sees the same signal at approximately $\sqrt{5}$ lower noise. Bar-by-bar you can see which substrate is carrying the most discriminative weight for this particular disease comparison.

HEALTHY REFERENCE — Fibroblast IMR90



DISEASE REFERENCE — TCGA SARC (n=206)
ACTIVE — tracks progression




Combined A — healthy (all 5)	0.97041 (NORMAL)
Combined A — disease (all 5)	1.09175 (URGENT)
ΔA combined (5 substrates)	+0.12134 vs methylation-only ΔA =+0.1502
Active A — healthy (non-saturated)	0.97041 (NORMAL) [5/5 active]
Active A — disease (non-saturated)	1.11140 (FLOOR BREACH) [4/5 active]

ΔA active (real progression signal)	+0.14100 (active substrates track true progression)
Saturation status (disease)	1 substrate saturated: Nucleosome Occupancy

Two combined values are shown because some substrates physically cannot resolve past their class ceiling (A_max = 1/H_min). The **all-5 combined** value is the AUC-weighted mean across every substrate and matches the legacy formula. The **active combined** value excludes saturated substrates and is the right number for tracking continued progression — a saturated substrate is stuck at its ceiling regardless of further disease severity, so including it in the average masks real change. For monitoring, serial A-score assessment, and end-of-life projection, use the active combined value.

DISEASE SIGNATURE COMPARISON — healthy muscle through IPF, pediatric sarcomas, TCGA SARC, to mesothelioma

Six conditions on a single chart spanning the full stromal-class clinical spectrum. β values reproduce per-substrate A-scores calibrated to primary sources: TCGA SARC 2017 Cell (n=206, Abeshouse et al.), TCGA MESO 2018 Cancer Discov (n=74, Hmeljak et al.), Crompton 2014 Cancer Discov Ewing sarcoma (n=112), and Shern 2014 Cancer Discov rhabdomyosarcoma (n=147). Pediatric sarcomas — Ewing and rhabdomyosarcoma — are shown alongside adult sarcomas because families of children lost to these diseases deserve to see their cancer on this card. IPF (idiopathic pulmonary fibrosis) is included as a non-cancer stromal-class failure where the same substrate readout applies. Nucleosome occupancy saturates at A ≈ 1.014 — the TIGHTEST ceiling of any class × substrate combination in the framework, tighter even than Progenitor WPS. All five cancers in the panel pin nucl at its ceiling; the other four substrates (methyl, fuzz, WPS, frag) carry the progression signal from DETECTABLE through FLOOR BREACH. For this class specifically, A_active (4/5) is the cleaner progression metric.




What happens past the ceiling is condition-specific. The A-score magnitude alone does not track severity once the ceiling is crossed — different diseases take different post-breach paths that the full five-substrate divergence pattern reveals. The subsection below walks through this class's post-breach trajectory: which substrates keep carrying signal, where the Warburg metabolic lock-in happens, where the glucose inversion point sits, and where the framework currently has high confidence versus where the open research questions remain.

POST-BREACH TRAJECTORY — WHAT HAPPENS AFTER A CROSSES 1.10 (STROMAL CLASS)

The stromal class covers sarcomas (adult and pediatric), mesothelioma, and the major non-cancer fibrotic diseases (IPF, scleroderma, cardiac fibrosis). Cfdna contribution from stromal cells is only ~4% of plasma — the lowest of any class — making detection technically demanding and favoring fragmentomics (most forgiving of dilution) over methylation for primary signal. Post-breach, stromal cancers share a common pattern with cycling cancers: methyl and frag carry the progression signal; nucl pins first and tightest.

The zone language below — Warburg boundary, glucose inversion, point of no return — applies only to post-breach cells. These terms are not used on the pre-breach bar, and the qualitative A-value anchors (~1.15, ~1.25, ~1.40+) are boundaries between therapeutic windows, not additional diagnostic tiers. Specific A-values are filed as prospective predictions to be validated against survival-stratified cohorts.



WHICH SUBSTRATES CARRY POST-BREACH SIGNAL — STROMAL CLASS

Stromal-class physics: nucleosome occupancy saturates at A ≈ 1.014 — the TIGHTEST ceiling of any class × substrate combination in the framework. The other four substrates carry the full progression signal.

Substrate	Ceiling (A)	Post-breach status
Methylation	1.22	Carries signal throughout all four zones
Fuzziness	1.21	Carries signal throughout all four zones
Windowed protection	1.15	Carries signal throughout all four zones
Fragment size	1.20	Carries signal throughout all four zones
Nucleosome occupancy	1.014	Saturated at ceiling (TIGHTEST in framework) — no further signal post-breach

INVERSION TERRITORY — SENESCENT FIBROBLASTS (non-cancer)

Stromal class has a documented non-cancer below-floor case: senescent fibroblasts. Cruickshanks 2013 Nature Cell Biology reports genome-wide methylation drift with aging and cellular senescence in dermal fibroblasts, with β values drifting below the healthy reference for a subset of promoter CpGs. The framework interprets this as a senescence-associated hypomethylation signature that puts A_methyl in the INVERSION zone on the pre-breach bar. This is clinically important: a senescent-fibroblast A-score that reads below-floor should not be mistaken for a reading error or for pre-cancer. It is a distinct biological state — the stromal analog of T-cell exhaustion in immune class — and it has its own therapeutic implications (senolytics, SASP-targeting therapies).

CONDITION-BY-CONDITION TRAJECTORY

Healthy muscle / fibroblast — reference, low cfDNA burden

Known: Stromal cells contribute ~4% of plasma cfDNA. Framework H_min is calibrated against fibroblast IMR90 (Roadmap). The low burden is the binding constraint on detection sensitivity; fragmentomics partially compensates because DELFI signal is proportionally preserved at low tissue fractions.

IPF (idiopathic pulmonary fibrosis) — A ≈ 1.015, MARGINAL tier (non-cancer)

Known: IPF is the framework's cleanest non-cancer floor-approach case in the stromal class. Fibroblast-to-myofibroblast transition produces measurable methylation changes at stromal-identity promoters. The framework reads this as A ≈ 1.015 — detectable but well below cancer BREACH — which is the correct clinical reading for a non-malignant fibrotic state.

Unknown: whether A-score trajectory in IPF patients tracks disease progression more sensitively than current HRCT imaging; whether antifibrotic therapy response (pirfenidone, nintedanib) shows A-score reduction before forced vital capacity stabilizes.

Test that closes the gap: G-2026-P012 (re-referenced): OSIC IPF Biobank reanalysis with serial cfDNA. Prediction: combined A-score trajectory slope will distinguish progressive from stable disease with AUC ≥ 0.80, outperforming serial FVC measurements.

TCGA SARC (adult sarcoma, n=206) — A ≈ 1.146, CROSSED CEILING

Known: Adult soft-tissue sarcoma (TCGA SARC 2017 Cell) shows stromal-class BREACH consistent with other classes' moderate-ΔA cancers. The post-surgical surveillance window is the framework's highest-value clinical use here: sarcoma recurs in 30-50% of patients after resection, and imaging surveillance alone misses early recurrence.

Unknown: whether post-resection A_active elevation predicts local or distant recurrence before imaging-detectable disease; whether histologic sarcoma subtypes (leiomyosarcoma, liposarcoma, UPS) trace distinguishable post-breach paths.

Test that closes the gap: G-2026-P033: Prospective cohort of 100 post-resection STS patients with serial cfDNA every 3 months for 24 months. Prediction: A_active elevation above baseline will detect recurrence 3-6 months before CT imaging in ≥ 60% of recurring cases.

Pediatric sarcomas (Ewing, rhabdomyosarcoma) — A ≈ 1.10-1.13, CROSSED CEILING

Known: Ewing sarcoma (Crompton 2014, Tirode 2014) and rhabdomyosarcoma (Shern 2014) are shown on this card because families of children lost to these cancers deserve to see their disease on the framework. The A-score biology mirrors adult sarcomas. The clinical need is identical: serial monitoring under VDC-IE (Ewing) or VAC (rhabdomyosarcoma) protocols to catch minimal residual disease.

Unknown: whether pediatric sarcoma MRD detection via cfDNA A-score outperforms current imaging; whether Ewing EWS-FLI1 fusion transcript levels correlate with A_active trajectory.

Test that closes the gap: G-2026-P034: COG (Children's Oncology Group) pediatric sarcoma cohort reanalysis with archived serial plasma. Prediction: A_active at end-of-induction predicts event-free survival with AUC ≥ 0.75.

WHAT THE FRAMEWORK CAN AND CANNOT SAY TODAY

The framework has high confidence for stromal-class post-breach: (1) four substrates carry full signal with nucl saturating tightest in framework; (2) IPF validation at A ≈ 1.015 correctly places non-cancer fibrosis in MARGINAL tier, not BREACH; (3) senescent fibroblasts produce genuine below-floor readings that the framework interprets coherently.


The framework has not yet tested sarcoma MRD detection prospectively, nor the pediatric sarcoma treatment-response predictions. G-2026-P012, P033, P034 define the specific validation plans.

This is the scientific posture the framework commits to: definitive where the evidence is definitive, open where it is open, and specific about the experiments that would close each gap. Predictions filed for this class: G-2026-P012, G-2026-P033, G-2026-P034. A reviewer who wants to falsify any claim here can do so against the named cohorts. A collaborator who wants to extend the framework has specific hypotheses to test.

THREE-COMPONENT DECOMPOSITION (C1/C2/C3) — PER SUBSTRATE

Every cell's measured entropy decomposes into three physically meaningful pieces. C1 is the universal Landauer floor: H_MIN_GLOBAL = 0.756500, measured in frontal cortex neurons by Lister 2013. Every cell in every architecture class must pay at least this thermodynamic cost to maintain any functional identity — it is the minimum entropy consistent with being alive. C2 is the class-specific overhead: the additional entropy the stromal class must carry above the universal floor to encode its particular architecture. For this class, C2 represents approximately 13.9% of the healthy reference entropy. C3 is the accessible gap: max(0, H_actual – H_min). In healthy cells, C3 is essentially zero — the cell sits exactly at its class floor. When C3 grows, the cell is departing from its architecture class. Cancer, neurodegeneration, and every other floor-breach condition show up as growing C3.

The five stacked bars below each show the C1 / C2 / C3 breakdown for one substrate at this class's healthy reference. Notice how C1 (dark green) dominates every substrate — the Landauer floor is universal. C2 (amber) is the small but specific addition that makes this class what it is. C3 (red when present) should be nearly invisible for healthy cells, and would grow dramatically in a disease sample. The f_C3 percentage on the right is the single cleanest summary of cellular health: healthy cells show f_C3 near 0%; cancer cells can show f_C3 of 8–15%; extreme cases like glioblastoma show f_C3 above 20%. This is not a statistical threshold calibrated from disease data. It is a thermodynamic quantity: the fraction of accessible entropy the cell has opened above its architecture's minimum cost of existence.



BEST TESTING METHOD RANKING — CLINICAL UTILITY FOR THIS CLASS

Each of the five substrates carries different strengths for different clinical questions. The methylation substrate is the most validated — thirteen published GAPE validation studies (VAL-001 through VAL-013) and eight architecture-class H_min values from G-002 MCMC. It is also the

substrate most affected by cfDNA dilution, because blood cfDNA is approximately 70% immune-derived. Fragment size (DELF) is the most forgiving of dilution — short-fragment enrichment persists even when tumor-derived cfDNA is a small fraction of total plasma DNA. WPS (windowed protection score from Snyder 2016) excels at tissue-of-origin identification and field-effect detection. Nucleosome occupancy and fuzziness add chromatin-level information that methylation alone cannot reveal.

The ranking below is specifically for the stromal class. The order reflects three factors: (1) the substrate's published single-substrate AUC for diseases in this class; (2) the biology of this class's primary failure mode (Stiffness Coupling); and (3) the practical constraint of cfDNA contribution in the relevant specimen (plasma for most classes, CSF for terminal, tissue biopsy where cfDNA is too dilute). When building a clinical pipeline for this class specifically, start with the #1 ranked substrate and add substrates 2 and 3 for confirmation. Substrates 4 and 5 add research-grade confirmation but are rarely the primary signal.

#	Substrate	Best for	Rationale
1	Fragment Size (DELF)	Cancer detection and monitoring	Stromal cfDNA contribution is only 4% — fragmentomics is the most forgiving of dilution. DELF signal holds at low tissue fractions.
2	Methylation	Floor departure confirmation	Secondary; cfDNA methylation signal for stromal class is diluted by immune background. Combine with deconvolution (EpiDISH, methylCC).
3	Windowed Protection Score	Fibrotic progression tracking	WPS at stromal identity promoters depletes as fibroblast-to-myofibroblast transition occurs. Detects fibrosis before clinical imaging.
4	Nucleosome Occupancy	Tumor microenvironment mapping	ATAC-seq on tumor stroma distinguishes activated cancer-associated fibroblasts from normal stromal cells. Research-grade primarily.
5	Nucleosome Fuzziness	Mechanical phenotype correlation	Stiffness-induced chromatin reorganization shows up as increased nucleosome fuzziness. Link to tumor microenvironment biomechanics.

BODY TEMPERATURE SCALING — $\alpha = 2.0$ LANDAUER CORRECTION

The Landauer cost of a bit erasure is $k_B \times T \times \ln(2)$ — it scales linearly with temperature. Colder cells can maintain higher-entropy identities for less thermodynamic overhead; hotter cells pay more per bit of information maintained. The GAPE framework captures this with a simple scaling law: $H_{\min}(T) = H_{\min}(37^\circ\text{C}) \times (T_{\text{body}} / 310.15\text{ K})^\alpha$, with $\alpha = 2.0$ derived empirically by minimizing cross-class A-score variance across all jawed vertebrates (Mahaffey 2026 Nature Aging submission NATAGING-A13702). The $\alpha = 2.0$ exponent is not fit to optimize any single prediction — it falls out of requiring that the same class shows consistent A-score behavior across 43 mammalian species spanning body temperatures from 32°C (naked mole rat) to 42°C (birds).

The table below shows what the stromal class floor $H_{\min}$ looks like at seven representative body temperatures. At 37°C (human reference, highlighted row), $H_{\min} = 0.8629$. At 42°C (birds), the floor is higher because the Landauer tax is larger. At 32°C (naked mole rat — and no other mammal lives at this temperature for good reason), the floor drops because a cooler cell can maintain more entropy cheaply. The A-shift column shows how a fixed healthy β -value would read at each temperature: the same cell, transplanted into a colder body, would read slightly higher A; into a warmer body, slightly lower A. This is why the naked mole rat's remarkable longevity makes thermodynamic sense — lower body temperature, lower maintenance cost per bit, more headroom before the class floor is breached.


T_body	Species example	H_min(T)	β_{healthy} equivalent	A-shift vs 37°C
42°C	Birds (chicken, finch)	0.8910	0.733	$A \times 0.969$
39°C	Rodents (mouse, rat)	0.8741	0.733	$A \times 0.987$
37°C	Humans (REFERENCE)	0.8629	0.733	$A \times 1.000$ (anchor)
35°C	Hibernating bats	0.8519	0.733	$A \times 1.013$
32°C	Naked mole rat	0.8354	0.733	$A \times 1.033$
25°C	Reptiles (anole lizard)	0.7975	0.733	$A \times 1.082$
15°C	Fish (salmon, cold-water)	0.7449	0.733	$A \times 1.159$

HEALTHY AGING TRAJECTORY — AGE-STRATIFIED REFERENCE A-SCORES

Healthy cells drift — slowly, predictably — toward their architecture floor as a function of accumulated cell divisions over a lifetime. For the stromal class, the drift rate is approximately 2.5% per generation — meaning each cell cycle carries a 2.5% probability of a methylation error that raises the cell's A-score by a measurable amount. Over decades, these errors accumulate. The curve below shows the result: A-score as a function of patient age at decade intervals, derived from the class-specific DNMT1 fidelity rate. At age 20, the reference healthy A-score sits well below 1.00; by age 80, it has drifted upward into the MARGINAL to DETECTABLE range. This is normal aging — not disease, just the statistical accumulation of

imperfect replication over eight decades.

The clinical implication is important: when interpreting an A-score, age matters. A patient at age 30 with A = 1.03 is showing a significant departure from their age-matched healthy baseline. The same A = 1.03 at age 75 might be entirely consistent with age-expected drift. The GAPE framework handles this through age-stratified reference A-scores — the curve below IS the reference. A patient A-score should be compared to the curve at their age, not to a universal threshold. Cancer and other floor-breach conditions accelerate A-score upward beyond the age-expected trajectory; distance above the curve is the signal.



VERTEBRATE LIFESPAN CONTEXT

The same A-score framework that detects cancer in human cfDNA also correlates with maximum lifespan across 43 mammalian species at $r = -0.9018$ ($p = 1.6 \times 10^{-16}$, Mahaffey 2026 Nature Aging submission NATAGING-A13702). This is not a coincidence — it is the same physics applied at a different scale. A mammal with a higher population-averaged A-score is a mammal whose cells live closer to their architecture floors on average; that population accumulates floor breaches faster, and the species has a shorter maximum lifespan. The $A = 1.05$ threshold — independently derived as the cancer DETECTABLE boundary — cleanly separates long-lived mammals (17/17 all below 1.05) from short-lived mammals (11/11 all above 1.05) with complete accuracy. Same threshold. Same physics. Different timescale.

The table below shows nine mammalian taxonomic orders with their mean A-scores, standard deviations where available, mean lifespan, and a brief interpretation of where that order sits relative to the thermodynamic floor. Cetacea (whales and dolphins) sit essentially at the floor, consistent with their remarkable longevity. Primates including humans sit slightly above the floor. Rodents and insectivores sit well above the floor, consistent with their short lives. This class's reference A-score (0.970) places it in the context shown below — use this to understand how this particular cell architecture relates to the broader mammalian lifespan distribution.

Order	N	Mean A	Lifespan (yr)	Interpretation
Cetacea	5	0.997 ±0.016	99	At thermodynamic floor
Proboscidea	1	0.987 —	70	At floor
Primates	6	1.007 ±0.015	49	Slightly above floor
Artiodactyla	4	1.015 ±0.011	30	Slightly above floor
Chiroptera	4	1.041 ±0.007	35	Above floor (longevity outliers)
Carnivora	9	1.053 ±0.023	28	Above floor
Lagomorpha	2	1.114 ±0.002	8	Well above floor
Rodentia	6	1.125 ±0.018	9	Well above floor
Insectivora	1	1.157 —	2	Furthest from floor

INTERVENTION LEVERS — RANKED BY EXPECTED IMPACT

If a patient's A-score is elevated — MARGINAL, DETECTABLE, or URGENT — the clinical question is what can be done. The GAPE framework does not prescribe treatment; that remains the physician's judgment. But the framework does identify which intervention categories have the mechanistic strength to move this particular class's A-score back toward the floor. The five categories below are the standard anti-aging and cancer-adjacent interventions: senolytics (clearing senescent cells), metabolic restoration (NAD+, caloric restriction, exercise), epigenetic restoration (DNMT/TET modulators, HDAC inhibitors), reprogramming (cyclic Yamanaka factors), and checkpoint stringency (cell cycle control, immune checkpoint blockade).

The ranking below is specifically for the stromal class and its primary failure mode (Stiffness Coupling). Impact Level 1 (Dominant) means the intervention directly addresses the class's failure mechanism and is the first lever to try. Level 2 (Strong) means substantial mechanistic support for A-score restoration in this class. Level 3 (Moderate) means helpful but not addressing the binding constraint. Level 4 (Limited) means the intervention works for other classes but does not target this class's specific biology. Level 5 (Not applicable) means the intervention category is biologically incompatible with this class — e.g., senolytics cannot act on post-mitotic terminal cells that do not become classically senescent, and reprogramming cannot be applied to terminal cells without losing the identity that defines them.

Impact	Category	Mechanism	Rationale
1 DOMINANT	Senolytics	senolytics	Dominant — senescent fibroblasts are the primary driver of stromal dysfunction
2 STRONG	Epigenetic Restoration	epigenetic rx	Strong — epigenetic resetting of pro-fibrotic methylation programs
3 MODERATE	Metabolic	metabolic	Moderate — metabolic normalization helps but senescent burden is the binding constraint
3 MODERATE	Checkpoint Stringency	checkpoint	Moderate — checkpoint modulation reduces fibrotic signaling cascade
4 LIMITED	Reprogramming	reprogramming	Limited — stromal architecture serves protective functions

CANCER PANEL — ΔA RANKED

The TCGA cancer validation is where the framework earns its keep. For every cancer type in this class, we compute two numbers from published TCGA methylation data: A_healthy (from matched normal tissue) and A_tumor (from cancer samples in the same study). The ΔA = A_tumor – A_healthy is the signal the framework predicts should appear — and it does. Thresholds for tier assignment: A = 1.05 DETECT, A = 1.07 URGENT, A = 1.10 FLOOR BREACH. These thresholds are not fit to the cancer data; they were derived from the healthy-vs-age statistics and happen to align with the cancer validation cleanly. The bars below are ordered by |ΔA| descending so the cancers with the largest signal appear first.

Note that ΔA ranges typical for this class matter. Terminal-class cancers (glioma, glioblastoma) show ΔA ≈ 0.22–0.27 — the largest in the entire 28-cancer TCGA dataset. Cycling epithelial cancers show ΔA ≈ 0.13–0.19 — moderate but clearly detectable. Stromal and progenitor-lineage cancers show ΔA ≈ 0.10–0.17. Each cancer below is cited to its primary TCGA publication and shows sample size. Direct links from A-score to TCGA evidence — no black boxes between physics and clinical data.

#1	Mesothelioma (MESO) n=87 · TCGA MESO 2018 Cell	ΔA=+0.1334	A=1.112 FLOOR
#2	Sarcoma (SARC) n=206 · TCGA SARC 2017 Cell	ΔA=+0.1204	A=1.109 FLOOR

CORE METRICS — PUBLISHED AND DERIVED

Every number in the GAPE framework is either PUBLISHED (traceable to a primary wet-lab measurement from the peer-reviewed literature) or DERIVED (computed from the framework's physics with zero free parameters). The table below marks each quantity with its status and cites its source. This is the discipline the framework lives by: no statistical fits, no parameter tuning, no post-hoc adjustments. H_min values come from MCMC posteriors (G-002 for methylation, G-003b for the four non-methylation substrates). Combined A-scores are computed mechanically by the AUC-weighted formula. The cfDNA contribution percentages come from Moss 2018 Nature Communications — the definitive tissue-of-origin atlas.

Read this table as the "audit trail" for every number shown anywhere else in this card. When the framework says A_healthy = 0.9701, you can see exactly which β value produced it, which H_min was used, and where both came from. The n_bio field is currently marked PRELIMINARY because absolute n_bio values await the G-007 MCMC run; the class ordering has been confirmed (Spearman ρ = 0.905, p = 0.002 against commitment level) but the absolute values carry that caveat transparently.

Metric	Value	Source
cfDNA contribution (plasma)	4.0%	Moss 2018 Nat Commun — PUBLISHED
H_min methylation (class)	0.862950	G-002 MCMC posterior — DERIVED
H_min methylation σ (MCMC)	±0.001400	G-002 17-chain MCMC R-hat<1.001 — DERIVED

Metric	Value	Source
n_bio (class-specific)	20.5	PRELIMINARY — ordering confirmed (p=0.905, p=0.002); absolute pending G-007
Healthy drift rate	2.5% / generation	DNMT1 fidelity loss under turnover — DERIVED
f_C2 (architecture-locked %)	13.9%	C2 fraction at healthy reference — DERIVED
Healthy A (methylation)	0.9701 (NORMAL)	β=0.733 — DERIVED
Disease A (methylation)	1.1202 (FLOOR BREACH)	β=0.607 — DERIVED
Healthy combined A (5 subs)	0.9704 (NORMAL)	Σ(AUC_i × A_i)/Σ(AUC_i) across 5 substrates — DERIVED
Disease combined A (5 subs)	1.0917 (URGENT)	Σ(AUC_i × A_i)/Σ(AUC_i) across 5 substrates — DERIVED
ΔA combined	+0.1213	Combined signal — DERIVED

DATED PREDICTIONS FOR THIS CLASS

A prediction is not a claim unless it can fail. Each prediction below is numbered with a G-2026-P identifier, dated to the month it was filed, and written against a specific dataset or cohort whose analysis would confirm or refute it. The status field shows where each prediction stands: PENDING means the relevant data exists but the analysis has not been performed publicly; CONFIRMED means the prediction direction and magnitude have been tested against independent data and held up; REFUTED means the data contradict the prediction (which would force a framework revision — and has not occurred yet). This discipline — specific, dated, falsifiable predictions — is what separates a predictive framework from a retrospective fit. Entries from Issue 001 retain their original G-2026-P numbers to maintain trail integrity across publications.

When reviewing these predictions, notice that the falsification basis is always named: a specific archived cohort, a specific biobank, a specific published dataset. These are not vague "more research is needed" statements. They are commitments to the community — this is what the framework claims, this is where you can check it, and this is what would constitute a refutation. The predictions below are for the stromal class specifically, filed as part of this publication or carried forward from Issue 001.

G-2026-P012

April 2026

PENDING

In prospective IPF cohorts with archived serial blood samples and known progression outcomes, the stromal-class combined A-score trajectory slope will distinguish progressive from stable disease with AUC >= 0.80, outperforming serial forced vital capacity (FVC) measurements currently used for staging.

Basis: IPF progresses over 2-5 years. Existing antifibrotic therapies (pirfenidone, nintedanib) have variable response. The framework predicts that combined-substrate A-score trajectory identifies responders and non-responders earlier than FVC. Falsifiable in the OSIC IPF Biobank (www.osicild.org).

ISSUE 003 ADDENDUM — STROMAL IN THE ATLAS

5 cell types: adipocyte, endothelial, fibroblast, smooth_muscle, stromal_other

Identity loci: 2,294 (methyl) · **H_min by substrate:** methyl 0.862950 · nucl 0.985667 · fuzz 0.832386 · wps 0.612686 · frag 0.724691 · **methyl ceiling** 1.1588

Healthy age-matched band (age_reference_matrix.json). A = 1.0 is the commitment line; healthy reads on this band.

Age (midpoint)	n	A mean	p10	p90	β mean	Source
4	10	0.9438	0.9135	0.9740	0.748	Roadmap pediatric
14	12	0.9510	0.9190	0.9829	0.744	Roadmap
24	18	0.9563	0.9225	0.9900	0.741	Moss endothelial
34	22	0.9615	0.9283	0.9948	0.738	Moss+Roadmap
44	25	0.9667	0.9317	1.0016	0.735	Moss 2018
54	32	0.9734	0.9371	1.0098	0.731	Moss 2018
64	38	0.9784	0.9426	1.0142	0.728	TCGA SARC STN
74	28	0.9849	0.9478	1.0221	0.724	Aging vascular
84	15	0.9913	0.9529	1.0297	0.720	Extrapolated
95	5	0.9991	0.9578	1.0404	0.715	Extrapolated

CELL IDENTITY & CLINICAL CONTEXT

Includes	Hematopoietic stem cells (HSC, CD34+CD38- compartment), interfollicular epidermal stem cells, hair follicle bulge stem cells, Merkel cell progenitors, cholangiocytes (hepatic stem cell origin), intestinal stem cells (Lgr5+), neural stem cells, muscle satellite cells, limbal stem cells
Cancers (this class)	HSC-origin AML (distinct from progenitor-lineage AML — arises from CD34+CD38- compartment per Adelman 2019), basal cell carcinoma (BCC — most common cancer in humans, epidermal stem cell origin), squamous cell carcinoma (SCC — hair-follicle stem-cell-origin subset), Merkel cell carcinoma (MCC — rare, aggressive, polyomavirus-positive and UV-driven subclasses), cholangiocarcinoma (CHOL, intrahepatic + extrahepatic, hepatic stem cell / cholangiocyte origin), cancer stem cell populations within cycling-class tumors (residual disease after primary treatment)
Non-cancer applications	HSC aging and immunosenescence (Adelman 2019, Beerman 2013), clonal hematopoiesis of indeterminate potential (CHIP — shared with Progenitor class), hematopoietic failure in the elderly, post-transplant niche reconstitution biology, heterochronic parabiosis biology, age-related decline in tissue regeneration across all adult-stem compartments
Primary failure mode	Niche Depletion
Warburg status	EMERGING

COMMENTARY

Adult tissue stem cells occupy the intermediate position between pluripotent embryonic stem cells and fully committed differentiated cells. They retain self-renewal capacity while being lineage-restricted: hematopoietic stem cells (HSC, CD34+CD38-) produce blood and immune cells, intestinal stem cells (Lgr5+) regenerate gut epithelium, epidermal and hair-follicle stem cells maintain skin, and cholangiocytes serve as the liver's stem-cell compartment for biliary regeneration. This partial commitment is encoded in H_min = 0.873718 for the methylation substrate, reflecting the slightly more open chromatin state that preserves self-renewal flexibility. But this open chromatin comes with a thermodynamic consequence that defines this card: healthy adult stem cells already sit near the maximum-entropy state for the positional substrates, leaving very little measurement headroom before ceiling saturation.

This class exhibits the most severe saturation pattern in the framework. Three of five substrates saturate below the BREACH threshold: nucleosome occupancy at ceiling A = 1.0407, nucleosome fuzziness at ceiling A = 1.0196, and windowed protection score at ceiling A = 1.0112 — the tightest ceiling of any class × substrate combination in the entire framework. Only methylation (ceiling A = 1.1445) and fragmentomics (ceiling A = 1.1886) remain active past BREACH, and these two substrates carry the full severity signal for any adult-stem-origin cancer in the panel. For this class, A_combined systematically understates severity by approximately 62 percent compared to A_active (2/5) — the largest masking effect in the framework. This is not a defect of the combined metric; it is an honest reflection of the underlying physics. Three of the five substrates have reached their physical entropy ceiling and cannot resolve further severity. A_active (2/5) is therefore reported as the primary severity metric for this class, with A_combined shown for cross-card continuity but flagged as misleading past the DETECTABLE tier.

The distinction between adult-stem-origin cancers and progenitor-lineage cancers is clinically consequential and must be stated explicitly. Acute myeloid leukemia appears in both Card #4 (Progenitor) and Card #7 (Adult Stem) because the disease arises from two biologically distinct compartments. Progenitor-lineage AML — the majority of TCGA AML cases characterized by Ley et al. in the 2013 New England Journal of Medicine analysis (n=200 whole-genome or exome sequences) — originates from committed myeloid progenitors (CMP, GMP) and carries the methylation and mutation signatures of those compartments. HSC-origin AML, characterized separately by Adelman et al. in Cancer Discovery 2019, arises from the pre-leukemic CD34+CD38- HSC-enriched compartment and involves profound epigenetic reprogramming of enhancers, bivalent promoters, and hematopoietic transcription factors (KLF6, BCL6, RUNX3) that begins during normal aging and appears to predispose to leukemic transformation. The two AML populations are not interchangeable in the framework: they occupy different architecture classes because they emerge from cells in different thermodynamic states. Listing AML on both cards without the distinction would be sloppy. Listing it only on one card would miss the patients whose disease originated in the other compartment.


Basal cell carcinoma (BCC) is the most common cancer in humans, with over four million new cases diagnosed annually in the United States alone. Nearly every BCC arises from epidermal or hair-follicle stem cells under UV-driven Hedgehog pathway activation. Mortality is low — below one percent — but the cumulative burden across lifetimes is enormous and the disease maps cleanly onto this class. Merkel cell carcinoma (MCC, Harms et al. 2015 Cancer Research, n=49 for the polyomavirus-negative subtype) is rare (approximately 2,500 US cases per year) but aggressive, with five-year survival around 45 percent for regional disease and below 15 percent for distant disease. Its two subclasses — Merkel cell polyomavirus-positive (integrated viral oncogenes) and MCPyV-negative (UV-damage driven) — share the same cell-of-origin methylation signature. Cholangiocarcinoma (TCGA CHOL, Farshidfar et al. 2017 Cell Reports, n=38) is the second most common primary liver cancer, aggressive, typically detected at advanced stage, and arises from cholangiocytes, which function as the hepatic stem-cell compartment for biliary regeneration. Its IDH-mutant molecular subtype shows characteristic ARID1A hypermethylation that maps onto the stem-cell-origin thermodynamic signature. Every one of these cancers is uncommon compared to the fourteen cycling cancers in Card #5, but the families of patients who lose a loved one to MCC or cholangiocarcinoma deserve to see the disease named on the framework card that describes its biology.

The class's defining thermodynamic challenge creates a specific detection problem: once a patient crosses DETECTABLE (A = 1.05), three of the five substrates pin at their ceilings within the first approximately 0.05 A-score units, and absolute A-score values across those three substrates lose their ability to resolve further severity. The framework proposes three mechanisms by which extended cellular thermodynamic activity remains detectable past this ceiling. First, methyl and frag retain their full dynamic range — they saturate only at A = 1.1445 and A = 1.1886 respectively, leaving meaningful headroom past BREACH for the two most clinically mature cfDNA biomarker platforms. A_{active} (2/5) computed on just these two substrates preserves severity discrimination that A_{combined} destroys. Second, the rate of approach to the saturating ceilings contains temporal information that single-point readings discard. In serial monitoring, a patient whose wps A-score moves from 0.970 through 1.000 to 1.010 over successive draws before plateauing has a different disease trajectory than one who hits 1.010 on first sample; time-to-saturation becomes a severity proxy in the regime where absolute-value measurement has lost resolution. Third, when adult-stem-origin cancers progress, the signal propagates into adjacent architecture classes — HSC-origin AML shows up as a rising immune-class A-score via clonal expansion, cutaneous stem-cell cancers propagate to cycling-class signals as field cancerization, and cholangiocarcinoma progression shows secretory-class signal change as the transformed cells commit toward biliary epithelial identity. Multi-class A-score divergence, particularly where the adult-stem class has flatlined at ceiling while one or more adjacent classes continue to rise, becomes a specific diagnostic signature of adult-stem-origin disease progression. This cross-class mechanism connects directly to the forthcoming Issue 004 framework on multi-class A-score divergence as an early signal of metastatic progression.

HSC aging is the most important non-cancer application for this class. The Adelman 2019 Cancer Discovery study of lineage-negative CD34+CD38- HSC-enriched cells from young and aged healthy donors identified 529 differentially methylated regions encompassing 2,249 differentially methylated cytosines across 748 genes — precisely the methylation-substrate floor departure the framework predicts for age-related adult-stem decline. Beerman et al. 2013 Cell Stem Cell independently characterized proliferation-dependent DNA hypermethylation of Polycomb-regulated genes in aged HSC (GSE44117). Together these studies establish that the methylation substrate detects HSC aging at the earliest stage of its epigenomic reprogramming, well before the downstream cytopenias or cancer-susceptibility phenotypes become clinically apparent. Clonal hematopoiesis of indeterminate potential (CHIP), the pre-malignant clonal state characterized by Steensma 2015 and shared in scope with the Progenitor class, represents the transition zone where adult-stem-class floor departure begins to predict downstream myeloid disease. Stem-cell transplantation biology, heterochronic parabiosis, and niche-depletion research all benefit from a metric that tracks adult-stem-class thermodynamic state independently of phenotypic readout.

FIVE-SUBSTRATE FIDELITY GAUGE

The gauge below is the central image of Issue 002. A single horizontal bar represents the A-score axis — from A = 0.90 (below the floor) through A = 1.00 (exactly at the Adult Stem class floor, H_{min} = 0.8737) up through the clinical threshold zones: MARGINAL at A = 1.01, DETECTABLE at 1.05, URGENT at 1.07, and FLOOR BREACH at 1.10. Five dots are plotted on this single axis for a healthy reference sample (HSC CD34+CD38- young, markers above the bar), and five more dots for a disease reference (HSC-origin AML, markers below the bar). Each dot is one substrate: methylation, nucleosome occupancy, nucleosome fuzziness, WPS, or fragment size (DELFI). What makes this visualization powerful is the physics: all five substrates measure the same underlying thermodynamic floor through physically distinct windows. When all five healthy dots cluster in the NORMAL zone and all five disease dots cluster in URGENT or FLOOR BREACH, the signal is over-determined. Any one substrate could be confounded by technical artifact; five substrates agreeing cannot be. This is the "less blurry" advantage: √5 noise reduction for free.




Look at the spread in each direction. If the healthy five-dot cluster sits tightly together near A = 0.97–1.00, that indicates the reference population is at its architecture floor — precisely where the framework predicts healthy adult stem cells should sit. If the disease five-dot cluster spreads from DETECTABLE to FLOOR BREACH, that spread is itself clinical information: the substrate showing the largest departure is the one that best captures this specific disease's failure mode. For this class, the primary failure mode is Niche Depletion, and the substrate ranking shown further down the card explains which of the five substrates best reveals it.

SUBSTRATE-BY-SUBSTRATE BREAKDOWN


Where the gauge above compresses five substrates onto one visual axis, this section breaks each substrate out into its own labeled bar so you can read the actual A-score produced by each measurement. The physics is the same across all five: take the raw laboratory value (a β methylation fraction, a nucleosome occupancy probability, a WPS score, etc.), compute its Shannon entropy $H(\text{value})$, then divide by the class-specific and substrate-specific H_{min} to get a dimensionless A-score comparable across substrates. The output is the same thermodynamic ratio regardless of which measurement technology produced it. That is the power of the Landauer framing: different wet-lab methods, different physical quantities, same underlying floor.

The formula at the foot of this section — $A_{\text{combined}} = \Sigma(AUC_i \times A_i) / \Sigma(AUC_i)$ — combines the five individual A-scores into a single weighted score, using each substrate's published detection AUC as its weight. Higher-AUC substrates contribute more to the combined result. For adult stem cells, the healthy combined A-score below should sit at approximately 0.97 (NORMAL tier); the disease combined A-score should sit above 1.05 (DETECTABLE or higher). The difference between them — the ΔA combined — is the signal available to a blood draw that analyzes all five substrates simultaneously. A single-substrate assay sees a smaller signal at higher noise; the five-substrate combination sees the same signal at approximately $\sqrt{5}$ lower noise. Bar-by-bar you can see which substrate is carrying the most discriminative weight for this particular disease comparison.

HEALTHY REFERENCE — HSC CD34+CD38- young



DISEASE REFERENCE — HSC-origin AML
ACTIVE — tracks progression




Combined A — healthy (all 5)	0.96999 (NORMAL)
Combined A — disease (all 5)	1.07222 (URGENT)
ΔA combined (5 substrates)	+0.10223 vs methylation-only ΔA =+0.1504
Active A — healthy (non-saturated)	0.96999 (NORMAL) [5/5 active]
Active A — disease (non-saturated)	1.13549 (FLOOR BREACH) [2/5 active]

ΔA active (real progression signal)	+0.16550 (active substrates track true progression)
Saturation status (disease)	3 substrates saturated: Nucleosome Occupancy, Nucleosome Fuzziness, Windowed Protection Score

Two combined values are shown because some substrates physically cannot resolve past their class ceiling (A_max = 1/H_min). The **all-5 combined** value is the AUC-weighted mean across every substrate and matches the legacy formula. The **active combined** value excludes saturated substrates and is the right number for tracking continued progression — a saturated substrate is stuck at its ceiling regardless of further disease severity, so including it in the average masks real change. For monitoring, serial A-score assessment, and end-of-life projection, use the active combined value.

DISEASE SIGNATURE COMPARISON — healthy HSC through aging, CHIP, BCC, MCC, cholangiocarcinoma, to HSC-origin AML

Six conditions on a single chart spanning the full adult-stem clinical spectrum. β values reproduce per-substrate A-scores calibrated to primary sources: Adelman 2019 Cancer Discov HSCe aging methylation (n=5-7 per age group), Beerman 2013 Cell Stem Cell HSC methylation landscape, Harms 2015 Cancer Res MCC genomic characterization (n=49), and Farshidfar 2017 Cell Reports TCGA CHOL (n=38). Three substrates saturate below BREACH: nucleosome occupancy (ceiling A = 1.0407), nucleosome fuzziness (ceiling A = 1.0196), and windowed protection score (ceiling A = 1.0112 — the tightest ceiling of any class × substrate combination in the entire framework). Only methylation and fragmentomics remain active past BREACH; A_active (2/5) is the primary severity metric for this class, with A_combined reported for cross-card continuity but flagged as misleading past DETECTABLE. BCC is shown because it is the most common cancer in humans; MCC is shown because its rarity does not reduce the obligation to honor the patients it kills; cholangiocarcinoma is shown because its stem-cell origin (cholangiocytes) places it squarely in this class.




What happens past the ceiling is condition-specific. The A-score magnitude alone does not track severity once the ceiling is crossed — different diseases take different post-breach paths that the full five-substrate divergence pattern reveals. The subsection below walks through this class's post-breach trajectory: which substrates keep carrying signal, where the Warburg metabolic lock-in happens, where the glucose inversion point sits, and where the framework currently has high confidence versus where the open research questions remain.

POST-BREACH TRAJECTORY — WHAT HAPPENS AFTER A CROSSES 1.10 (ADULT STEM CLASS)

The adult tissue stem class has the tightest H_min ceilings in the framework — WPS saturates at A ≈ 1.011 (the tightest single ceiling anywhere), fuzz at 1.020, and nucl at 1.041. Three substrates pin at or near the ceiling before BREACH is reached, leaving methylation and fragmentomics as the only substrates with meaningful post-breach headroom. This class has the clearest documented non-cancer inversion case in the framework: HSC aging and niche depletion.

The zone language below — Warburg boundary, glucose inversion, point of no return — applies only to post-breach cells. These terms are not used on the pre-breach bar, and the qualitative A-value anchors (~1.15, ~1.25, ~1.40+) are boundaries between therapeutic windows, not additional diagnostic tiers. Specific A-values are filed as prospective predictions to be validated against survival-stratified cohorts.



WHICH SUBSTRATES CARRY POST-BREACH SIGNAL — ADULT STEM CLASS

Adult-stem-class physics: three substrates saturate early (WPS, fuzz, nucl). Only methylation and fragmentomics carry post-breach progression signal. A_active (2/5) is the primary severity metric; A_combined is misleading past DETECTABLE.

Substrate	Ceiling (A)	Post-breach status
Methylation	1.14	Carries signal throughout all four zones
Fragment size	1.19	Carries signal throughout all four zones (deepest headroom)
Nucleosome occupancy	1.041	Saturated at ceiling — no further signal post-breach
Fuzziness	1.020	Saturated at ceiling — no further signal post-breach
Windowed protection	1.011	Saturated at ceiling (TIGHTEST in framework) — no further signal post-breach

INVERSION TERRITORY — NICHE DEPLETION INVERSION (documented)

Hematopoietic stem cells in aged bone marrow show a documented failure mode distinct from cancer: niche depletion. Adelman 2019 Cancer Discovery (HSC-enriched aging methylation, n=5-7 per age group) and Beerman 2013 Cell Stem Cell (GSE44117) report that aged HSCs lose methylation fidelity at lineage-commitment loci, producing what appears as a below-healthy A-score signature on methylation when compared to young HSC reference. This is the Niche Depletion Inversion — a non-cancer, non-malignant failure mode that the framework interprets coherently. Clinically, this distinguishes healthy aging from pre-leukemic states (CHIP, CCUS) which show upward elevation. The inversion direction itself is diagnostic: below-floor = aging; above-floor = clonal expansion.

CONDITION-BY-CONDITION TRAJECTORY

Healthy HSC (young adult) — reference, low cfDNA burden

Known: Hematopoietic stem cells represent ~3% of plasma cfDNA. H_min calibrated from Roadmap E035 HSC-enriched reference. Young-adult HSCs show tightest methylation fidelity of any stem cell population in the body.

HSC aging (Adelman 2019) — A below-floor (INVERSION territory)

Known: Aged HSCs (> 70 years) show progressive loss of methylation fidelity at lineage-commitment loci. The framework reads this as A_methyl in the INVERSION zone on the pre-breach bar — genuinely below healthy reference, not a measurement artifact. This is the non-cancer side of the adult-stem aging spectrum.

Unknown: whether longitudinal A-score tracking in aging adults identifies those at elevated risk for CHIP/CCUS transformation; whether interventions targeting HSC aging (senolytics, metformin) reverse the A-score back toward healthy.

Test that closes the gap: G-2026-P035: UK Biobank subset with serial blood archived and subsequent AML/MDS diagnosis outcomes. Prediction: subjects with persistent below-floor A_methyl over 10 years will have elevated risk (HR ≥ 2.0) for hematologic malignancy vs age-matched controls.

CHIP / CCUS (pre-malignant clonal) — A ≈ 1.02-1.05, MARGINAL-to-DETECTABLE

Known: Clonal hematopoiesis of indeterminate potential (Steensma 2015) and clonal cytopenias of undetermined significance (Malcovati 2017) represent pre-malignant clonal expansion. The framework reads CHIP/CCUS as modest A-score elevation — clearly above healthy, clearly below cancer. This is the directional inversion from HSC aging: the clonal expansion drives A up, not down.

Unknown: whether A-score trajectory within CHIP predicts which patients progress to MDS or AML over a 5-10 year horizon; whether specific CHIP mutation classes (DNMT3A, TET2, ASXL1) produce distinguishable A-score paths.

Test that closes the gap: G-2026-P036: Retrospective reanalysis of CCUS-to-MDS progression cohorts. Prediction: A_active slope above 0.01/year identifies imminent progression with AUC ≥ 0.75.

HSC-origin AML / MCC / cholangiocarcinoma — A ≈ 1.10-1.15, CROSSED CEILING

Known: HSC-origin adult leukemias and Merkel cell carcinoma sit in metabolic-window to structural-only zones post-breach. The two-substrate detection strategy (methyl + frag) outperforms five-substrate combining for this class because three substrates pin at their ceilings early.

Unknown: whether the two-substrate classifier AUC exceeds 0.80 in prospective cohorts; whether HSC-origin AML trajectories under azacitidine-venetoclax differ from progenitor-origin AML (immune-class card) at the A-score level.

Test that closes the gap: G-2026-P037: Paired cohort analysis of TCGA AML (n=200) reanalyzed with five-substrate cfDNA panels plus MCC (Harms 2015, n=49). Prediction: methyl+frag two-substrate classifier achieves AUC ≥ 0.80; any of three saturating substrates showing A above its predicted ceiling > measurement noise falsifies class H_min.

WHAT THE FRAMEWORK CAN AND CANNOT SAY TODAY

The framework has high confidence for adult-stem-class post-breach: (1) only methyl and frag carry post-breach signal — three substrates saturate before BREACH; (2) the Niche Depletion Inversion is documented in primary literature (Adelman 2019, Beerman 2013) and the framework interprets it coherently; (3) the directional distinction between aging (below-floor) and pre-malignant clonal expansion (above-floor) is clinically actionable.


The framework has not yet tested the HSC-aging-to-AML progression prediction prospectively, nor the two-substrate classifier AUC. G-2026-P035, P036, P037 define the specific validation plans.

This is the scientific posture the framework commits to: definitive where the evidence is definitive, open where it is open, and specific about the experiments that would close each gap. Predictions filed for this class: G-2026-P035, G-2026-P036, G-2026-P037. A reviewer who wants to falsify any claim here can do so against the named cohorts. A collaborator who wants to extend the framework has specific hypotheses to test.

THREE-COMPONENT DECOMPOSITION (C1/C2/C3) — PER SUBSTRATE

Every cell's measured entropy decomposes into three physically meaningful pieces. C1 is the universal Landauer floor: H_MIN_GLOBAL = 0.756500, measured in frontal cortex neurons by Lister 2013. Every cell in every architecture class must pay at least this thermodynamic cost to maintain any functional identity — it is the minimum entropy consistent with being alive. C2 is the class-specific overhead: the additional entropy the adult stem class must carry above the universal floor to encode its particular architecture. For this class, C2 represents approximately 13.5% of the healthy reference entropy. C3 is the accessible gap: max(0, H_actual – H_min). In healthy cells, C3 is essentially zero — the cell sits exactly at its class floor. When C3 grows, the cell is departing from its architecture class. Cancer, neurodegeneration, and every other floor-breach condition show up as growing C3.

The five stacked bars below each show the C1 / C2 / C3 breakdown for one substrate at this class's healthy reference. Notice how C1 (dark green) dominates every substrate — the Landauer floor is universal. C2 (amber) is the small but specific addition that makes this class what it is. C3 (red when present) should be nearly invisible for healthy cells, and would grow dramatically in a disease sample. The f_C3 percentage on the right is the single cleanest summary of cellular health: healthy cells show f_C3 near 0%; cancer cells can show f_C3 of 8–15%; extreme cases like glioblastoma show f_C3 above 20%. This is not a statistical threshold calibrated from disease data. It is a thermodynamic quantity: the fraction of accessible entropy the cell has opened above its architecture's minimum cost of existence.



BEST TESTING METHOD RANKING — CLINICAL UTILITY FOR THIS CLASS

Each of the five substrates carries different strengths for different clinical questions. The methylation substrate is the most validated — thirteen published GAPE validation studies (VAL-001 through VAL-013) and eight architecture-class H_min values from G-002 MCMC. It is also the

substrate most affected by cfDNA dilution, because blood cfDNA is approximately 70% immune-derived. Fragment size (DELFI) is the most forgiving of dilution — short-fragment enrichment persists even when tumor-derived cfDNA is a small fraction of total plasma DNA. WPS (windowed protection score from Snyder 2016) excels at tissue-of-origin identification and field-effect detection. Nucleosome occupancy and fuzziness add chromatin-level information that methylation alone cannot reveal.

The ranking below is specifically for the adult stem class. The order reflects three factors: (1) the substrate's published single-substrate AUC for diseases in this class; (2) the biology of this class's primary failure mode (Niche Depletion); and (3) the practical constraint of cfDNA contribution in the relevant specimen (plasma for most classes, CSF for terminal, tissue biopsy where cfDNA is too dilute). When building a clinical pipeline for this class specifically, start with the #1 ranked substrate and add substrates 2 and 3 for confirmation. Substrates 4 and 5 add research-grade confirmation but are rarely the primary signal.

#	Substrate	Best for	Rationale
1	Methylation	Primary severity metric past DETECTABLE	Methyl remains active past BREACH (ceiling A = 1.1445) — one of only two substrates that do not saturate for this class. Carries the bulk of the severity signal for HSC aging, BCC, MCC, and cholangiocarcinoma. Primary source: Adelman 2019 Cancer Discov.
2	Fragment Size (DELFI)	Primary severity metric past DETECTABLE	Frag (DELFI) remains active past BREACH (ceiling A = 1.1886, deepest headroom of any substrate in this class). Fragmentomic signal reflects clonal stem-cell turnover and cfDNA release kinetics. Co-primary with methyl for severity tracking.
3	Nucleosome Occupancy	Detection boundary — ceiling at A = 1.0407	Nucl saturates below BREACH. Useful for binary detection (floor departure vs healthy) but cannot resolve severity above DETECTABLE. Trajectory-slope information remains useful in serial monitoring (time-to-saturation as proxy severity).
4	Nucleosome Fuzziness	Detection boundary — ceiling at A = 1.0196	Fuzz saturates tighter than nucl. Same detection-only use case. Informative in combination with wps for cross-verification of floor departure.
5	Windowed Protection Score	Tightest ceiling in framework — A = 1.0112	WPS has the tightest physical ceiling of any class × substrate pairing. Healthy adult stem cells already sit at WPS A ≈ 0.970, with only +0.041 headroom before saturation. Any detectable deviation in wps for this class is clinically significant — but wps cannot distinguish mild from severe deviation above A = 1.0112.

BODY TEMPERATURE SCALING — $\alpha = 2.0$ LANDAUER CORRECTION

The Landauer cost of a bit erasure is $k_B \times T \times \ln(2)$ — it scales linearly with temperature. Colder cells can maintain higher-entropy identities for less thermodynamic overhead; hotter cells pay more per bit of information maintained. The GAPE framework captures this with a simple scaling law: $H_{\min}(T) = H_{\min}(37^\circ\text{C}) \times (T_{\text{body}} / 310.15\text{ K})^\alpha$, with $\alpha = 2.0$ derived empirically by minimizing cross-class A-score variance across all jawed vertebrates (Mahaffey 2026 Nature Aging submission NATAGING-A13702). The $\alpha = 2.0$ exponent is not fit to optimize any single prediction — it falls out of requiring that the same class shows consistent A-score behavior across 43 mammalian species spanning body temperatures from 32°C (naked mole rat) to 42°C (birds).

The table below shows what the adult stem class floor $H_{\min}$ looks like at seven representative body temperatures. At 37°C (human reference, highlighted row), $H_{\min} = 0.8737$. At 42°C (birds), the floor is higher because the Landauer tax is larger. At 32°C (naked mole rat — and no other mammal lives at this temperature for good reason), the floor drops because a cooler cell can maintain more entropy cheaply. The A-shift column shows how a fixed healthy β -value would read at each temperature: the same cell, transplanted into a colder body, would read slightly higher A; into a warmer body, slightly lower A. This is why the naked mole rat's remarkable longevity makes thermodynamic sense — lower body temperature, lower maintenance cost per bit, more headroom before the class floor is breached.


T_body	Species example	H_min(T)	$\beta_{\text{healthy equivalent}}$	A-shift vs 37°C
42°C	Birds (chicken, finch)	0.9021	0.726	A × 0.969
39°C	Rodents (mouse, rat)	0.8850	0.726	A × 0.987
37°C	Humans (REFERENCE)	0.8737	0.726	A × 1.000 (anchor)
35°C	Hibernating bats	0.8625	0.726	A × 1.013
32°C	Naked mole rat	0.8458	0.726	A × 1.033
25°C	Reptiles (anole lizard)	0.8074	0.726	A × 1.082

T_body	Species example	H_min(T)	β_healthy equivalent	A-shift vs 37°C
15°C	Fish (salmon, cold-water)	0.7542	0.726	$A \times 1.159$

HEALTHY AGING TRAJECTORY — AGE-STRATIFIED REFERENCE A-SCORES

Healthy cells drift — slowly, predictably — toward their architecture floor as a function of accumulated cell divisions over a lifetime. For the adult stem class, the drift rate is approximately 3.0% per generation — meaning each cell cycle carries a 3.0% probability of a methylation error that raises the cell's A-score by a measurable amount. Over decades, these errors accumulate. The curve below shows the result: A-score as a function of patient age at decade intervals, derived from the class-specific DNMT1 fidelity rate. At age 20, the reference healthy A-score sits well below 1.00; by age 80, it has drifted upward into the MARGINAL to DETECTABLE range. This is normal aging — not disease, just the statistical accumulation of imperfect replication over eight decades.

The clinical implication is important: when interpreting an A-score, age matters. A patient at age 30 with A = 1.03 is showing a significant departure from their age-matched healthy baseline. The same A = 1.03 at age 75 might be entirely consistent with age-expected drift. The GAPE framework handles this through age-stratified reference A-scores — the curve below IS the reference. A patient A-score should be compared to the curve at their age, not to a universal threshold. Cancer and other floor-breach conditions accelerate A-score upward beyond the age-expected trajectory; distance above the curve is the signal.



VERTEBRATE LIFESPAN CONTEXT

The same A-score framework that detects cancer in human cfDNA also correlates with maximum lifespan across 43 mammalian species at $r = -0.9018$ ($p = 1.6 \times 10^{-16}$, Mahaffey 2026 Nature Aging submission NATAGING-A13702). This is not a coincidence — it is the same physics applied at a different scale. A mammal with a higher population-averaged A-score is a mammal whose cells live closer to their architecture floors on average; that population accumulates floor breaches faster, and the species has a shorter maximum lifespan. The $A = 1.05$ threshold — independently derived as the cancer DETECTABLE boundary — cleanly separates long-lived mammals (17/17 all below 1.05) from short-lived mammals (11/11 all above 1.05) with complete accuracy. Same threshold. Same physics. Different timescale.

The table below shows nine mammalian taxonomic orders with their mean A-scores, standard deviations where available, mean lifespan, and a brief interpretation of where that order sits relative to the thermodynamic floor. Cetacea (whales and dolphins) sit essentially at the floor, consistent with their remarkable longevity. Primates including humans sit slightly above the floor. Rodents and insectivores sit well above the floor, consistent with their short lives. This class's reference A-score (0.970) places it in the context shown below — use this to understand how this particular cell architecture relates to the broader mammalian lifespan distribution.

Order	N	Mean A	Lifespan (yr)	Interpretation
Cetacea	5	0.997 ±0.016	99	At thermodynamic floor
Proboscidea	1	0.987 —	70	At floor
Primates	6	1.007 ±0.015	49	Slightly above floor
Artiodactyla	4	1.015 ±0.011	30	Slightly above floor
Chiroptera	4	1.041 ±0.007	35	Above floor (longevity outliers)
Carnivora	9	1.053 ±0.023	28	Above floor
Lagomorpha	2	1.114 ±0.002	8	Well above floor
Rodentia	6	1.125 ±0.018	9	Well above floor
Insectivora	1	1.157 —	2	Furthest from floor

INTERVENTION LEVERS — RANKED BY EXPECTED IMPACT

If a patient's A-score is elevated — MARGINAL, DETECTABLE, or URGENT — the clinical question is what can be done. The GAPE framework does not prescribe treatment; that remains the physician's judgment. But the framework does identify which intervention categories have the mechanistic strength to move this particular class's A-score back toward the floor. The five categories below are the standard anti-aging and cancer-adjacent interventions: senolytics (clearing senescent cells), metabolic restoration (NAD+, caloric restriction, exercise), epigenetic restoration (DNMT/TET modulators, HDAC inhibitors), reprogramming (cyclic Yamanaka factors), and checkpoint stringency (cell cycle control, immune checkpoint blockade).

The ranking below is specifically for the adult stem class and its primary failure mode (Niche Depletion). Impact Level 1 (Dominant) means the intervention directly addresses the class's failure mechanism and is the first lever to try. Level 2 (Strong) means substantial mechanistic support for A-score restoration in this class. Level 3 (Moderate) means helpful but not addressing the binding constraint. Level 4 (Limited) means the intervention works for other classes but does not target this class's specific biology. Level 5 (Not applicable) means the intervention category is biologically incompatible with this class — e.g., senolytics cannot act on post-mitotic terminal cells that do not become classically senescent, and reprogramming cannot be applied to terminal cells without losing the identity that defines them.

Impact	Category	Mechanism	Rationale
2 STRONG	Metabolic	metabolic	Strong — niche metabolic restoration moves stem cell fidelity index
2 STRONG	Epigenetic Restoration	epigenetic rx	Strong — epigenetic restoration extends stem cell functional lifespan
2 STRONG	Cyclic Reprogramming	reprogramming	Strong — cyclic Yamanaka rejuvenates without full dedifferentiation
2 STRONG	Niche Checkpoint	checkpoint	Strong — niche checkpoint signals regulate stem cell quiescence
3 MODERATE	Senolytics	senolytics	Moderate — senescent cells in the niche drive inversion

CANCER PANEL — ΔA RANKED

The TCGA cancer validation is where the framework earns its keep. For every cancer type in this class, we compute two numbers from published TCGA methylation data: A_healthy (from matched normal tissue) and A_tumor (from cancer samples in the same study). The ΔA = A_tumor – A_healthy is the signal the framework predicts should appear — and it does. Thresholds for tier assignment: A = 1.05 DETECT, A = 1.07 URGENT, A = 1.10 FLOOR BREACH. These thresholds are not fit to the cancer data; they were derived from the healthy-vs-age statistics and happen to align with the cancer validation cleanly. The bars below are ordered by |ΔA| descending so the cancers with the largest signal appear first.

Note that ΔA ranges typical for this class matter. Terminal-class cancers (glioma, glioblastoma) show ΔA ≈ 0.22–0.27 — the largest in the entire 28-cancer TCGA dataset. Cycling epithelial cancers show ΔA ≈ 0.13–0.19 — moderate but clearly detectable. Stromal and progenitor-lineage cancers show ΔA ≈ 0.10–0.17. Each cancer below is cited to its primary TCGA publication and shows sample size. Direct links from A-score to TCGA evidence — no black boxes between physics and clinical data.

#1

Uveal Melanoma (UVM)

n=80 · Robertson 2017 Cancer Cell

ΔA=+0.1072

A=1.086 URGENT

CORE METRICS — PUBLISHED AND DERIVED

Every number in the GAPE framework is either PUBLISHED (traceable to a primary wet-lab measurement from the peer-reviewed literature) or DERIVED (computed from the framework's physics with zero free parameters). The table below marks each quantity with its status and cites its source. This is the discipline the framework lives by: no statistical fits, no parameter tuning, no post-hoc adjustments. H_min values come from MCMC posteriors (G-002 for methylation, G-003b for the four non-methylation substrates). Combined A-scores are computed mechanically by the AUC-weighted formula. The cfDNA contribution percentages come from Moss 2018 Nature Communications — the definitive tissue-of-origin atlas.

Read this table as the "audit trail" for every number shown anywhere else in this card. When the framework says A_healthy = 0.9696, you can see exactly which β value produced it, which H_min was used, and where both came from. The n_bio field is currently marked PRELIMINARY because absolute n_bio values await the G-007 MCMC run; the class ordering has been confirmed (Spearman ρ = 0.905, p = 0.002 against commitment level) but the absolute values carry that caveat transparently.

Metric	Value	Source
cfDNA contribution (plasma)	3.0%	Moss 2018 Nat Commun — PUBLISHED
H_min methylation (class)	0.873718	G-002 MCMC posterior — DERIVED
H_min methylation σ (MCMC)	±0.001600	G-002 17-chain MCMC R-hat<1.001 — DERIVED
n_bio (class-specific)	18.5	PRELIMINARY — ordering confirmed (ρ=0.905, p=0.002); absolute pending G-007

Metric	Value	Source
Healthy drift rate	3.0% / generation	DNMT1 fidelity loss under turnover — DERIVED
f_C2 (architecture-locked %)	13.5%	C2 fraction at healthy reference — DERIVED
Healthy A (methylation)	0.9696 (NORMAL)	$\beta=0.726$ — DERIVED
Disease A (methylation)	1.1200 (FLOOR BREACH)	$\beta=0.586$ — DERIVED
Healthy combined A (5 subs)	0.9700 (NORMAL)	$\Sigma(AUC_i \times A_i) / \Sigma(AUC_i)$ across 5 substrates — DERIVED
Disease combined A (5 subs)	1.0722 (URGENT)	$\Sigma(AUC_i \times A_i) / \Sigma(AUC_i)$ across 5 substrates — DERIVED
ΔA combined	+0.1022	Combined signal — DERIVED

DATED PREDICTIONS FOR THIS CLASS

A prediction is not a claim unless it can fail. Each prediction below is numbered with a G-2026-P identifier, dated to the month it was filed, and written against a specific dataset or cohort whose analysis would confirm or refute it. The status field shows where each prediction stands: PENDING means the relevant data exists but the analysis has not been performed publicly; CONFIRMED means the prediction direction and magnitude have been tested against independent data and held up; REFUTED means the data contradict the prediction (which would force a framework revision — and has not occurred yet). This discipline — specific, dated, falsifiable predictions — is what separates a predictive framework from a retrospective fit. Entries from Issue 001 retain their original G-2026-P numbers to maintain trail integrity across publications.

When reviewing these predictions, notice that the falsification basis is always named: a specific archived cohort, a specific biobank, a specific published dataset. These are not vague "more research is needed" statements. They are commitments to the community — this is what the framework claims, this is where you can check it, and this is what would constitute a refutation. The predictions below are for the adult stem class specifically, filed as part of this publication or carried forward from Issue 001.

G-2026-P013

April 2026

PENDING

In prospective CHIP/CCUS cohorts with archived serial blood samples and known progression outcomes, the stem_adult A_active (2/5, methyl + frag only) trajectory will identify progression to overt MDS or AML 12 to 24 months before cytopenia-based diagnosis in a majority of progressing cases, outperforming A_combined (all 5) which is systematically masked by approximately 62 percent in this class because 3 of 5 substrates saturate below BREACH.

Basis: CHIP-to-MDS/AML progression is slow, stereotyped, and clinically silent until late. The framework predicts that A_active identifies progressors earlier than cytopenia onset because the epigenomic state change precedes the hematologic phenotype, and A_active is not masked by ceiling-pinned positional substrates. Falsifiable in any CHIP/CCUS cohort with 18+ months of serial blood samples and progression outcomes (candidate cohorts include Jaiswal CHIP cohorts with archived samples, UK Biobank CHIP subset with longitudinal follow-up).

G-2026-P014

April 2026

PENDING

Trajectory slope ($\Delta A / \Delta t$) on the saturating substrates (nucl, fuzz, wps) during the pre-saturation window will correlate with clinical outcome in adult-stem-origin cancers with $AUC \geq 0.75$, providing severity information in the regime where absolute A-score values have lost resolution. Formally: time-to-saturation from healthy baseline is a valid severity proxy for this class.

Basis: The saturation ceiling destroys absolute-value severity information but preserves temporal information in the form of rate-of-approach. Falsifiable in any serial-sampling cohort with outcome data; directly testable in MDS, BCC, MCC, and CHOL longitudinal cohorts.

G-2026-P015

April 2026

PENDING

Multi-class A-score divergence signature — adult-stem class flatlined at ceiling while an adjacent architecture class (immune, cycling, or secretory depending on tissue of origin) continues to rise — will identify adult-stem-origin disease progression with higher specificity than any single-class score past BREACH.

Basis: When single-class saturation pins the primary class at ceiling, adjacent-class signal propagation becomes the discriminating severity signal. Cross-class divergence tests the framework's claim that architecture classes interact through shared tissue compartments and clonal expansion dynamics. Falsifiable in pan-cancer cohorts with multi-class A-score monitoring; connects directly to Issue 004 multi-class metastasis detection framework.

ISSUE 003 ADDENDUM — STEM_ADULT IN THE ATLAS

1 cell types: HSC

Identity loci: 46,907 (methyl) · H_min by substrate: methyl 0.873718 · nucl 0.960866 · fuzz 0.980754 · wps 0.988964 · frag 0.841327 · methyl ceiling 1.1445

Healthy age-matched band (age_reference_matrix.json). A = 1.0 is the commitment line; healthy reads on this band.

Age (midpoint)	n	A mean	p10	p90	β mean	Source
4	8	0.9375	0.9103	0.9647	0.745	Adelman pediatric HSC
14	10	0.9428	0.9137	0.9718	0.742	Adelman 2019
24	15	0.9462	0.9153	0.9772	0.740	Adelman+Roadmap
34	18	0.9497	0.9190	0.9803	0.738	Adelman 2019
44	22	0.9531	0.9206	0.9856	0.736	Adelman 2019
54	28	0.9565	0.9221	0.9908	0.734	Adelman 2019
64	32	0.9614	0.9255	0.9974	0.731	Adelman aged HSC
74	25	0.9664	0.9289	1.0038	0.728	Adelman elderly HSC
84	12	0.9728	0.9341	1.0115	0.724	Extrapolated
95	4	0.9791	0.9392	1.0190	0.720	Extrapolated

CELL IDENTITY & CLINICAL CONTEXT

Includes	Common myeloid progenitor (CMP), granulocyte-monocyte progenitor (GMP), neural progenitor (NPC), erythroid progenitor, intestinal transit-amplifying cells
Cancers (this class)	MDS (all subtypes: RA, RARS, RCMD, RAEB), secondary AML from MDS, CMML, therapy-related myeloid neoplasms, pediatric B-ALL and T-ALL, mixed-lineage (MLL-rearranged) leukemias, JMML, medulloblastoma, PNET, OPC-derived gliomas, erythroleukemia (AML-M6), colorectal adenomas (pre-cancer)
Non-cancer applications	CHIP (clonal hematopoiesis), CCUS (clonal cytopenias), bone marrow failure syndromes, aplastic anemia, IBD intestinal stem/progenitor compartment dysregulation, serrated polyps
Primary failure mode	Replication Throughput Ceiling
Warburg status	EMERGING

COMMENTARY

Progenitor and transit-amplifying cells are the high-throughput workhorses of the body. They sit between adult tissue stem cells and fully committed daughter cells — partially committed to a specific lineage, but still capable of rapid proliferation to expand the committed pool before final differentiation. CMP and GMP in the bone marrow produce millions of granulocytes per hour. Intestinal transit-amplifying cells produce the daughter cells that replace the colonic epithelium every four to seven days. Neural progenitors produce neurons and glia during development and in the adult neurogenic niches. The rapid proliferation produces the characteristic 4.5% per generation drift rate — slower than cycling epithelium but faster than adult stem cells.

The Replication Throughput Ceiling is the class-specific failure mode. When a progenitor cell population divides fast enough to exceed DNMT1 maintenance fidelity — because the cell is consuming DNMT1 faster than it can be transcribed and translated — methylation errors accumulate disproportionately. This is why MDS (myelodysplastic syndrome) arises in the progenitor lineage: the CMP/GMP pool undergoes accumulated methylation errors under chronic high-throughput demand, and the resulting dysplastic progenitors produce cytopenias through failed daughter-cell maturation. The framework predicts that the early signature of MDS is detectable in peripheral blood cfDNA years before cytopenia onset, through the progenitor-class combined A-score.

This is also the class where DNMT1 inhibitor therapy (azacitidine, decitabine) acts most directly. These drugs are used clinically in MDS and AML, and the framework predicts that response to therapy — measured as A-score trajectory — is a more sensitive indicator of treatment success than the blast count or cytogenetic remission criteria currently used. The same substrate readout predicts response. The five substrates together predict it earlier and with higher confidence than any one.

The validation sequence for this class runs along a well-documented clonal-hematopoiesis spectrum. Clonal hematopoiesis of indeterminate potential (CHIP) — a pre-malignant state present in roughly 10% of healthy adults over 70 (Steensma 2015 Blood, doi:10.1182/blood-2015-03-631747) — shows A_combined ≈ 1.01, detectable but unambiguously below BREACH. Clonal cytopenias of undetermined significance (CCUS) sits at A ≈ 1.03 in MARGINAL tier (Malcovati 2017 Blood, doi:10.1182/blood-2017-04-777607). Low-risk MDS produces A_active ≈ 1.09 (BREACH) while A_combined still reads 1.05 because saturations have started. High-risk MDS sits at A_active ≈ 1.12, and secondary AML (MDS-derived) reaches A_active ≈ 1.15 — nearly matching the de novo AML signal from Card #3 Immune. The progression across these six conditions is the cleanest single-class severity gradient in the entire document and is directly relevant to the clinical decision of WHEN to escalate treatment in MDS monitoring. Bone marrow biopsy cannot be repeated monthly; a blood-based A_active can.

The progenitor class encompasses more than MDS. Pediatric acute lymphoblastic leukemia (ALL) — both B-cell precursor ALL (the most common childhood cancer) and T-ALL — arises from lymphoid progenitors and displays genome-wide hypermethylation signatures distinct from bulk AML (Nordlund 2013 Genome Biol, n=764 pediatric ALL methylomes, doi:10.1186/gb-2013-14-9-r105). The framework predicts A_active elevation in B-ALL and T-ALL through the same physics that drives the MDS signal: rapid progenitor-pool division exceeding DNMT1 fidelity. Medulloblastoma — the most common malignant pediatric brain tumor — arises from cerebellar neural progenitor cells and shows four distinct methylation subgroups (WNT, SHH, Group 3, Group 4), each with different clinical trajectories (Northcott 2017 Nature, n=1,256 methylation-profiled cases, doi:10.1038/nature22973). The DNA methylation-based CNS tumor classifier from Capper 2018 Nature (doi:10.1038/nature26000) changed pediatric neuro-oncology diagnosis in up to 12% of cases. Chronic myelomonocytic leukemia (CMML), juvenile myelomonocytic leukemia (JMML), therapy-related myeloid neoplasms, and the OPC-derived gliomas distinct from terminal-class adult glioma all belong to this class by cell-of-origin physics. Each family has lost someone to these diseases; each cancer deserves its place on the card. The unifying framework is not a coincidence — it is the common thermodynamic signature of the partially-committed, rapidly-dividing cell pool.

A clinical consequence of the saturation pattern for this class is unusually consequential. Progenitor is one of two classes in the framework where THREE substrates saturate below BREACH — nucleosome occupancy (ceiling A = 1.028), nucleosome fuzziness (ceiling A = 1.040), and WPS (ceiling A = 1.012 — the tightest class × substrate combination in the entire framework). Only two substrates (methylation, fragment size) carry the real BREACH-capable progression signal for this class. The physics explanation is clean: progenitor cells are partially uncommitted by design. They must retain chromatin flexibility to produce multiple daughter lineages, so their healthy β values for the structural substrates (nucl, fuzz, WPS)

already sit near the maximum-entropy state. There is almost no headroom. Any disease-driven drift hits the ceiling quickly. This is not a frameworklimitation — it is the correct reading of developmental flexibility as a measurement-physics constraint. The practical consequence for serial MDSmonitoring is stark. The legacy A_combined (all 5 substrates) flattens at approximately 1.07 once saturations kick in and CANNOT distinguishlow-risk MDS from high-risk MDS from secondary AML — all three compress together. A_active (methyl + frag only, the 2 non-saturated substrates)does distinguish: 1.089, 1.120, 1.150 respectively. For this class specifically, reporting A_active is not a refinement of A_combined; it is the ONLYway to track progression past the low-risk MDS threshold. Any clinical pipeline using GAPE for MDS severity grading must use A_active. The maskis +55% of the total signal — the largest in the document, dwarfing even Terminal's +45% — and unlike Terminal (where saturation is a binarycancer indicator), Progenitor's saturations are a CONTINUOUS LOSS of resolution that the active-substrate formula fully recovers.

FIVE-SUBSTRATE FIDELITY GAUGE

The gauge below is the central image of Issue 002. A single horizontal bar represents the A-score axis — from A = 0.90 (below the floor) through A = 1.00 (exactly at the Progenitor class floor, H_min = 0.8522) up through the clinical threshold zones: MARGINAL at A = 1.01, DETECTABLE at 1.05, URGENT at 1.07, and FLOOR BREACH at 1.10. Five dots are plotted on this single axis for a healthy reference sample (CMP young, markers above the bar), and five more dots for a disease reference (MDS (progenitor), markers below). Each dot is one substrate: methylation, nucleosome occupancy, nucleosome fuzziness, WPS, or fragment size (DELFI). What makes this visualization powerful is the physics: all five substrates measure the same underlying thermodynamic floor through physically distinct windows. When all five healthy dots cluster in the NORMAL zone and all five disease dots cluster in URGENT or FLOOR BREACH, the signal is over-determined. Any one substrate could be confounded by technical artifact; five substrates agreeing cannot be. This is the "less blurry" advantage: $\sqrt{5}$ noise reduction for free.




Look at the spread in each direction. If the healthy five-dot cluster sits tightly together near A = 0.97–1.00, that indicates the reference population is at its architecture floor — precisely where the framework predicts healthy progenitor cells should sit. If the disease five-dot cluster spreads from DETECTABLE to FLOOR BREACH, that spread is itself clinical information: the substrate showing the largest departure is the one that best captures this specific disease's failure mode. For this class, the primary failure mode is Replication Throughput Ceiling, and the substrate ranking shown further down the card explains which of the five substrates best reveals it.

SUBSTRATE-BY-SUBSTRATE BREAKDOWN


Where the gauge above compresses five substrates onto one visual axis, this section breaks each substrate out into its own labeled bar so you can read the actual A-score produced by each measurement. The physics is the same across all five: take the raw laboratory value (a β methylation fraction, a nucleosome occupancy probability, a WPS score, etc.), compute its Shannon entropy $H(\text{value})$, then divide by the class-specific and substrate-specific H_{min} to get a dimensionless A-score comparable across substrates. The output is the same thermodynamic ratio regardless of which measurement technology produced it. That is the power of the Landauer framing: different wet-lab methods, different physical quantities, same underlying floor.

The formula at the foot of this section — $A_{\text{combined}} = \Sigma(AUC_i \times A_i) / \Sigma(AUC_i)$ — combines the five individual A-scores into a single weighted score, using each substrate's published detection AUC as its weight. Higher-AUC substrates contribute more to the combined result. For progenitor cells, the healthy combined A-score below should sit at approximately 0.97 (NORMAL tier); the disease combined A-score should sit above 1.05 (DETECTABLE or higher). The difference between them — the ΔA combined — is the signal available to a blood draw that analyzes all five substrates simultaneously. A single-substrate assay sees a smaller signal at higher noise; the five-substrate combination sees the same signal at approximately $\sqrt{5}$ lower noise. Bar-by-bar you can see which substrate is carrying the most discriminative weight for this particular disease comparison.

HEALTHY REFERENCE — CMP young



DISEASE REFERENCE — MDS (progenitor)
ACTIVE — tracks progression




Combined A — healthy (all 5)	0.97015 (NORMAL)
Combined A — disease (all 5)	1.06697 (DETECTABLE)
ΔA combined (5 substrates)	+0.09682 vs methylation-only ΔA =+0.1498
Active A — healthy (non-saturated)	0.97015 (NORMAL) [5/5 active]
Active A — disease (non-saturated)	1.12029 (FLOOR BREACH) [2/5 active]

ΔA active (real progression signal)	+0.15014 (active substrates track true progression)
Saturation status (disease)	3 substrates saturated: Nucleosome Occupancy, Nucleosome Fuzziness, Windowed Protection Score

Two combined values are shown because some substrates physically cannot resolve past their class ceiling ($A_{\text{max}} = 1/H_{\text{min}}$). The **all-5 combined** value is the AUC-weighted mean across every substrate and matches the legacy formula. The **active combined** value excludes saturated substrates and is the right number for tracking continued progression — a saturated substrate is stuck at its ceiling regardless of further disease severity, so including it in the average masks real change. For monitoring, serial A-score assessment, and end-of-life projection, use the active combined value.

DISEASE SIGNATURE COMPARISON — healthy CMP through CHIP, CCUS, MDS stages, to sAML

Six conditions on a single chart, all in the progenitor class — the clearest demonstration in the entire document of why A_{active} matters more than A_{combined} for serial monitoring. β values reproduce per-substrate A-scores along the well-documented clonal-hematopoiesis progression pathway (Steensma 2015 Blood on CHIP; Malcovati 2017 Blood on CCUS; Jiang 2020 Cell Death Dis on MDS hypermethylation). Three substrates saturate for this class — nucleosome occupancy at $A \approx 1.028$, fuzziness at $A \approx 1.040$, and WPS at $A \approx 1.012$ (the tightest ceiling of any class $\times$ substrate combination in the framework). Once disease severity crosses low-risk MDS, all three pin at their ceilings and the legacy A_{combined} flattens at ~ 1.07 regardless of further progression to high-risk MDS or secondary AML. A_{active} (methyl + frag only) continues to track: $1.089 \rightarrow 1.120 \rightarrow 1.150$ across the three MDS stages. For progenitor-class serial monitoring, A_{active} IS the progression signal. A_{combined} is confirmatory only.




What happens past the ceiling is condition-specific. The A-score magnitude alone does not track severity once the ceiling is crossed — different diseases take different post-breach paths that the full five-substrate divergence pattern reveals. The subsection below walks through this class's post-breach trajectory: which substrates keep carrying signal, where the Warburg metabolic lock-in happens, where the glucose inversion point sits, and where the framework currently has high confidence versus where the open research questions remain.

POST-BREACH TRAJECTORY — WHAT HAPPENS AFTER A CROSSES 1.10 (PROGENITOR CLASS)

The progenitor class is the framework’s clearest demonstration of why A_active matters more than A_combined for serial monitoring. Three substrates saturate early in this class — nucleosome occupancy at 1.028, fuzziness at 1.040, and WPS at 1.012 (tied with stem_adult WPS for the tightest ceiling in the framework). Once disease crosses low-risk MDS, all three pin at their ceilings and A_combined flattens regardless of further progression. A_active (methyl + frag only) continues to track: 1.089 → 1.120 → 1.150 across the three MDS stages. For this class, A_active IS the progression signal.

The zone language below — Warburg boundary, glucose inversion, point of no return — applies only to post-breach cells. These terms are not used on the pre-breach bar, and the qualitative A-value anchors (~1.15, ~1.25, ~1.40+) are boundaries between therapeutic windows, not additional diagnostic tiers. Specific A-values are filed as prospective predictions to be validated against survival-stratified cohorts.



WHICH SUBSTRATES CARRY POST-BREACH SIGNAL — PROGENITOR CLASS

Progenitor-class physics: three substrates saturate (nucl, fuzz, WPS). Only methylation and fragmentomics carry post-breach progression signal. A_active (2/5) is mandatory for serial monitoring; A_combined is confirmatory only and must not be used to track progression past low-risk MDS.

Substrate	Ceiling (A)	Post-breach status
Methylation	1.15	Carries signal throughout all four zones
Fragment size	1.19	Carries signal throughout all four zones
Nucleosome occupancy	1.028	Saturated at ceiling — no further signal post-breach
Fuzziness	1.040	Saturated at ceiling — no further signal post-breach
Windowed protection	1.012	Saturated at ceiling (TIGHTEST-tied in framework) — no further signal post-breach

CONDITION-BY-CONDITION TRAJECTORY

Healthy CMP (common myeloid progenitor) — reference

Known: Common myeloid progenitor is the reference state. Progenitor-class H_min = 0.8522 calibrated from Roadmap E034 (erythroid progenitor) plus CD34+CD38- HSC-enriched references. Baseline A ≈ 0.97 across all five substrates in healthy adults.

CHIP (clonal hematopoiesis) — A ≈ 1.01-1.03, MARGINAL tier

Known: Clonal hematopoiesis of indeterminate potential (Steensma 2015 Blood) is defined by recurrent somatic mutations (DNMT3A, TET2, ASXL1) at variant allele frequency ≥ 2% without cytopenia. The framework reads CHIP as modest A-score elevation — detectable above healthy, clearly pre-malignant. CHIP affects ~10% of people over 70 but only a small fraction progress to MDS/AML.

Unknown: whether the A-score trajectory within CHIP predicts which patients progress to MDS vs remain stable; whether specific CHIP mutation categories trace distinguishable paths.

Test that closes the gap: G-2026-P038: UK Biobank CHIP cohort with serial methylation and 10-year MDS/AML outcomes. Prediction: A_active slope ≥ 0.01/year within CHIP identifies imminent progression with AUC ≥ 0.75, outperforming VAF trajectory alone.

Low-risk MDS — A_combined ≈ 1.05, A_active ≈ 1.09

Known: Low-risk MDS is the first saturation event: WPS pins at its ceiling. This is the clinically important transition where A_combined starts under-reporting progression. A patient with A_combined ≈ 1.05 but A_active ≈ 1.09 is clearly in DETECTABLE tier on the active metric — the combined score is no longer the right instrument.

Unknown: whether the A_active trajectory over 6-12 months predicts which low-risk MDS patients progress to high-risk MDS before current IPSS-R stratification flags them.

Test that closes the gap: G-2026-P039: Prospective cohort of 200 low-risk MDS patients with serial cfDNA every 3 months for 24 months. Prediction: A_active slope identifies progression to high-risk MDS 3-6 months before IPSS-R change, with sensitivity ≥ 0.70 at specificity 0.85.

High-risk MDS / Secondary AML — A_combined flatlined, A_active ≈ 1.12-1.15

Known: High-risk MDS shows all three saturating substrates pinned; A_combined cannot distinguish high-risk MDS from secondary AML. But A_active continues to climb: 1.120 at high-risk MDS → 1.150 at secondary AML. This is the framework’s most concrete demonstration that serial A_active monitoring provides signal that standard clinical staging loses.

Unknown: whether A_active at secondary AML diagnosis predicts response to azacitidine-venetoclax or allogeneic HSCT more accurately than current risk scores (ELN 2022).

Test that closes the gap: G-2026-P040: Retrospective reanalysis of archived secondary AML serial cfDNA cohorts. Prediction: baseline A_active at diagnosis stratifies 2-year overall survival with AUC ≥ 0.70, outperforming ELN 2022 risk classification alone.

WHAT THE FRAMEWORK CAN AND CANNOT SAY TODAY

The framework has high confidence for progenitor-class post-breach: (1) only methyl and frag carry post-breach signal — three substrates saturate before BREACH; (2) A_active vs A_combined divergence in MDS is the framework’s cleanest clinical teaching example; (3) the class’s CHIP-to-MDS-to-AML progression pathway is biologically well-characterized and the framework tracks each stage coherently.


The framework has not yet tested the CHIP-to-MDS progression prediction, the low-risk to high-risk MDS early detection, or secondary AML prognostic stratification prospectively. G-2026-P038, P039, P040 define the validation plans.

This is the scientific posture the framework commits to: definitive where the evidence is definitive, open where it is open, and specific about the experiments that would close each gap. Predictions filed for this class: G-2026-P038, G-2026-P039, G-2026-P040. A reviewer who wants to falsify any claim here can do so against the named cohorts. A collaborator who wants to extend the framework has specific hypotheses to test.

THREE-COMPONENT DECOMPOSITION (C1/C2/C3) — PER SUBSTRATE

Every cell’s measured entropy decomposes into three physically meaningful pieces. C1 is the universal Landauer floor: H_MIN_GLOBAL = 0.756500, measured in frontal cortex neurons by Lister 2013. Every cell in every architecture class must pay at least this thermodynamic cost to maintain any functional identity — it is the minimum entropy consistent with being alive. C2 is the class-specific overhead: the additional entropy the progenitor class must carry above the universal floor to encode its particular architecture. For this class, C2 represents approximately 11.8% of the healthy reference entropy. C3 is the accessible gap: max(0, H_actual – H_min). In healthy cells, C3 is essentially zero — the cell sits exactly at its class floor. When C3 grows, the cell is departing from its architecture class. Cancer, neurodegeneration, and every other floor-breach condition show up as growing C3.

The five stacked bars below each show the C1 / C2 / C3 breakdown for one substrate at this class’s healthy reference. Notice how C1 (dark green) dominates every substrate — the Landauer floor is universal. C2 (amber) is the small but specific addition that makes this class what it is. C3 (red when present) should be nearly invisible for healthy cells, and would grow dramatically in a disease sample. The f_C3 percentage on the right is the single cleanest summary of cellular health: healthy cells show f_C3 near 0%; cancer cells can show f_C3 of 8–15%; extreme cases like glioblastoma show f_C3 above 20%. This is not a statistical threshold calibrated from disease data. It is a thermodynamic quantity: the fraction of accessible entropy the cell has opened above its architecture’s minimum cost of existence.



BEST TESTING METHOD RANKING — CLINICAL UTILITY FOR THIS CLASS

Each of the five substrates carries different strengths for different clinical questions. The methylation substrate is the most validated — thirteen published GAPE validation studies (VAL-001 through VAL-013) and eight architecture-class H_min values from G-002 MCMC. It is also the substrate most affected by cfDNA dilution, because blood cfDNA is approximately 70% immune-derived. Fragment size (DELF) is the most forgiving of dilution — short-fragment enrichment persists even when tumor-derived cfDNA is a small fraction of total plasma DNA. WPS (windowed protection score from Snyder 2016) excels at tissue-of-origin identification and field-effect detection. Nucleosome occupancy and fuzziness add chromatin-level information that methylation alone cannot reveal.

The ranking below is specifically for the progenitor class. The order reflects three factors: (1) the substrate’s published single-substrate AUC for diseases in this class; (2) the biology of this class’s primary failure mode (Replication Throughput Ceiling); and (3) the practical constraint of cfDNA contribution in the relevant specimen (plasma for most classes, CSF for terminal, tissue biopsy where cfDNA is too dilute). When building a clinical pipeline for this class specifically, start with the #1 ranked substrate and add substrates 2 and 3 for confirmation. Substrates 4 and 5 add

research-grade confirmation but are rarely the primary signal.

#	Substrate	Best for	Rationale
1	Methylation	Progenitor-lineage malignancy detection	Bone marrow methylation is the gold standard for MDS diagnosis. Peripheral blood cfDNA methylation signal is class-specific when deconvolution is applied.
2	Windowed Protection Score	Mismatch repair state monitoring	WPS at MMR-gene promoters indicates MLH1/MSH2 methylation status — directly relevant to Lynch syndrome and MSI-H tumors.
3	Fragment Size (DELFI)	Rapid turnover signal	Progenitor-origin cfDNA fragmentation pattern is distinct — high turnover produces characteristic short-fragment enrichment.
4	Nucleosome Occupancy	Commitment-state assessment	ATAC-seq discriminates early progenitor from late transit-amplifying states. Research-grade.
5	Nucleosome Fuzziness	Division-rate correlation	Fuzziness correlates with proliferation index. Useful for grading in MDS and other progenitor-lineage disorders.

BODY TEMPERATURE SCALING — $\alpha = 2.0$ LANDAUER CORRECTION

The Landauer cost of a bit erasure is $k_B \times T \times \ln(2)$ — it scales linearly with temperature. Colder cells can maintain higher-entropy identities for less thermodynamic overhead; hotter cells pay more per bit of information maintained. The GAPE framework captures this with a simple scaling law: $H_{\min}(T) = H_{\min}(37^\circ\text{C}) \times (T_{\text{body}} / 310.15\text{ K})^\alpha$, with $\alpha = 2.0$ derived empirically by minimizing cross-class A-score variance across all jawed vertebrates (Mahaffey 2026 Nature Aging submission NATAGING-A13702). The $\alpha = 2.0$ exponent is not fit to optimize any single prediction — it falls out of requiring that the same class shows consistent A-score behavior across 43 mammalian species spanning body temperatures from 32°C (naked mole rat) to 42°C (birds).


The table below shows what the progenitor class floor $H_{\min}$ looks like at seven representative body temperatures. At 37°C (human reference, highlighted row), $H_{\min} = 0.8522$. At 42°C (birds), the floor is higher because the Landauer tax is larger. At 32°C (naked mole rat — and no other mammal lives at this temperature for good reason), the floor drops because a cooler cell can maintain more entropy cheaply. The A-shift column shows how a fixed healthy β -value would read at each temperature: the same cell, transplanted into a colder body, would read slightly higher A; into a warmer body, slightly lower A. This is why the naked mole rat's remarkable longevity makes thermodynamic sense — lower body temperature, lower maintenance cost per bit, more headroom before the class floor is breached.

T_body	Species example	H_min(T)	$\beta_{\text{healthy equivalent}}$	A-shift vs 37°C
42°C	Birds (chicken, finch)	0.8799	0.740	$A \times 0.969$
39°C	Rodents (mouse, rat)	0.8632	0.740	$A \times 0.987$
37°C	Humans (REFERENCE)	0.8522	0.740	$A \times 1.000$ (anchor)
35°C	Hibernating bats	0.8413	0.740	$A \times 1.013$
32°C	Naked mole rat	0.8250	0.740	$A \times 1.033$
25°C	Reptiles (anole lizard)	0.7875	0.740	$A \times 1.082$
15°C	Fish (salmon, cold-water)	0.7356	0.740	$A \times 1.159$

HEALTHY AGING TRAJECTORY — AGE-STRATIFIED REFERENCE A-SCORES

Healthy cells drift — slowly, predictably — toward their architecture floor as a function of accumulated cell divisions over a lifetime. For the progenitor class, the drift rate is approximately 4.5% per generation — meaning each cell cycle carries a 4.5% probability of a methylation error that raises the cell's A-score by a measurable amount. Over decades, these errors accumulate. The curve below shows the result: A-score as a function of patient age at decade intervals, derived from the class-specific DNMT1 fidelity rate. At age 20, the reference healthy A-score sits well below 1.00; by age 80, it has drifted upward into the MARGINAL to DETECTABLE range. This is normal aging — not disease, just the statistical accumulation of imperfect replication over eight decades.

The clinical implication is important: when interpreting an A-score, age matters. A patient at age 30 with $A = 1.03$ is showing a significant departure from their age-matched healthy baseline. The same $A = 1.03$ at age 75 might be entirely consistent with age-expected drift. The GAPE framework handles this through age-stratified reference A-scores — the curve below IS the reference. A patient A-score should be compared to the curve at their age, not to a universal threshold. Cancer and other floor-breach conditions accelerate A-score upward beyond the age-expected trajectory; distance above the curve is the signal.



VERTEBRATE LIFESPAN CONTEXT

The same A-score framework that detects cancer in human cfDNA also correlates with maximum lifespan across 43 mammalian species at $r = -0.9018$ ($p = 1.6 \times 10^{-16}$, Mahaffey 2026 Nature Aging submission NATAGING-A13702). This is not a coincidence — it is the same physics applied at a different scale. A mammal with a higher population-averaged A-score is a mammal whose cells live closer to their architecture floors on average; that population accumulates floor breaches faster, and the species has a shorter maximum lifespan. The $A = 1.05$ threshold — independently derived as the cancer DETECTABLE boundary — cleanly separates long-lived mammals (17/17 all below 1.05) from short-lived mammals (11/11 all above 1.05) with complete accuracy. Same threshold. Same physics. Different timescale.

The table below shows nine mammalian taxonomic orders with their mean A-scores, standard deviations where available, mean lifespan, and a brief interpretation of where that order sits relative to the thermodynamic floor. Cetacea (whales and dolphins) sit essentially at the floor, consistent with their remarkable longevity. Primates including humans sit slightly above the floor. Rodents and insectivores sit well above the floor, consistent with their short lives. This class's reference A-score (0.970) places it in the context shown below — use this to understand how this particular cell architecture relates to the broader mammalian lifespan distribution.

Order	N	Mean A	Lifespan (yr)	Interpretation
Cetacea	5	0.997 ±0.016	99	At thermodynamic floor
Proboscidea	1	0.987 —	70	At floor
Primates	6	1.007 ±0.015	49	Slightly above floor
Artiodactyla	4	1.015 ±0.011	30	Slightly above floor
Chiroptera	4	1.041 ±0.007	35	Above floor (longevity outliers)
Carnivora	9	1.053 ±0.023	28	Above floor
Lagomorpha	2	1.114 ±0.002	8	Well above floor
Rodentia	6	1.125 ±0.018	9	Well above floor
Insectivora	1	1.157 —	2	Furthest from floor

INTERVENTION LEVERS — RANKED BY EXPECTED IMPACT

If a patient's A-score is elevated — MARGINAL, DETECTABLE, or URGENT — the clinical question is what can be done. The GAPE framework does not prescribe treatment; that remains the physician's judgment. But the framework does identify which intervention categories have the mechanistic strength to move this particular class's A-score back toward the floor. The five categories below are the standard anti-aging and cancer-adjacent interventions: senolytics (clearing senescent cells), metabolic restoration (NAD+, caloric restriction, exercise), epigenetic restoration (DNMT/TET modulators, HDAC inhibitors), reprogramming (cyclic Yamanaka factors), and checkpoint stringency (cell cycle control, immune checkpoint blockade).

The ranking below is specifically for the progenitor class and its primary failure mode (Replication Throughput Ceiling). Impact Level 1 (Dominant) means the intervention directly addresses the class's failure mechanism and is the first lever to try. Level 2 (Strong) means substantial mechanistic support for A-score restoration in this class. Level 3 (Moderate) means helpful but not addressing the binding constraint. Level 4 (Limited) means the intervention works for other classes but does not target this class's specific biology. Level 5 (Not applicable) means the intervention category is biologically incompatible with this class — e.g., senolytics cannot act on post-mitotic terminal cells that do not become classically senescent, and reprogramming cannot be applied to terminal cells without losing the identity that defines them.

Impact	Category	Mechanism	Rationale
1 DOMINANT	G2/M Checkpoint	checkpoint	Dominant — G2/M checkpoint activation is the primary lever

Impact	Category	Mechanism	Rationale
2 STRONG	MMR Restoration	epigenetic rx	Strong — MMR restoration directly addresses the Replication Throughput Ceiling
3 MODERATE	Senolytics	senolytics	Moderate — senescent progenitors contribute but are minor fraction
3 MODERATE	Metabolic	metabolic	Moderate — metabolic lever moves index but does not address the ceiling
4 LIMITED	Reprogramming	reprogramming	Limited — partial commitment; full reprogramming disrupts lineage

CORE METRICS — PUBLISHED AND DERIVED

Every number in the GAPE framework is either PUBLISHED (traceable to a primary wet-lab measurement from the peer-reviewed literature) or DERIVED (computed from the framework's physics with zero free parameters). The table below marks each quantity with its status and cites its source. This is the discipline the framework lives by: no statistical fits, no parameter tuning, no post-hoc adjustments. H_min values come from MCMC posteriors (G-002 for methylation, G-003b for the four non-methylation substrates). Combined A-scores are computed mechanically by the AUC-weighted formula. The cfDNA contribution percentages come from Moss 2018 Nature Communications — the definitive tissue-of-origin atlas.

Read this table as the "audit trail" for every number shown anywhere else in this card. When the framework says A_healthy = 0.9701, you can see exactly which β value produced it, which H_min was used, and where both came from. The n_bio field is currently marked PRELIMINARY because absolute n_bio values await the G-007 MCMC run; the class ordering has been confirmed (Spearman ρ = 0.905, p = 0.002 against commitment level) but the absolute values carry that caveat transparently.

Metric	Value	Source
cfDNA contribution (plasma)	2.0%	Moss 2018 Nat Commun — PUBLISHED
H_min methylation (class)	0.852216	G-002 MCMC posterior — DERIVED
H_min methylation σ (MCMC)	± 0.001700	G-002 17-chain MCMC R-hat<1.001 — DERIVED
n_bio (class-specific)	20.0	PRELIMINARY — ordering confirmed ($p=0.905$, $p=0.002$); absolute pending G-007
Healthy drift rate	4.5% / generation	DNMT1 fidelity loss under turnover — DERIVED
f_C2 (architecture-locked %)	11.8%	C2 fraction at healthy reference — DERIVED
Healthy A (methylation)	0.9701 (NORMAL)	$\beta=0.740$ — DERIVED
Disease A (methylation)	1.1199 (FLOOR BREACH)	$\beta=0.625$ — DERIVED
Healthy combined A (5 subs)	0.9701 (NORMAL)	$\Sigma(AUC_i \times A_i) / \Sigma(AUC_i)$ across 5 substrates — DERIVED
Disease combined A (5 subs)	1.0670 (DETECTABLE)	$\Sigma(AUC_i \times A_i) / \Sigma(AUC_i)$ across 5 substrates — DERIVED
ΔA combined	+0.0968	Combined signal — DERIVED

DATED PREDICTIONS FOR THIS CLASS

A prediction is not a claim unless it can fail. Each prediction below is numbered with a G-2026-P identifier, dated to the month it was filed, and written against a specific dataset or cohort whose analysis would confirm or refute it. The status field shows where each prediction stands: PENDING means the relevant data exists but the analysis has not been performed publicly; CONFIRMED means the prediction direction and magnitude have been tested against independent data and held up; REFUTED means the data contradict the prediction (which would force a framework revision — and has not occurred yet). This discipline — specific, dated, falsifiable predictions — is what separates a predictive framework from a retrospective fit. Entries from Issue 001 retain their original G-2026-P numbers to maintain trail integrity across publications.

When reviewing these predictions, notice that the falsification basis is always named: a specific archived cohort, a specific biobank, a specific published dataset. These are not vague "more research is needed" statements. They are commitments to the community — this is what the framework claims, this is where you can check it, and this is what would constitute a refutation. The predictions below are for the progenitor class specifically, filed as part of this publication or carried forward from Issue 001.

G-2026-P014

April 2026

PENDING

In MDS patients receiving hypomethylating agent (HMA) therapy, the progenitor-class combined A-score trajectory at 3 months post-initiation will identify responders (by 6-month IWG criteria) with $AUC \geq 0.85$, outperforming baseline-to-3-month cytogenetic change alone.

Basis: HMA therapy is the first-line treatment for intermediate/high-risk MDS. Currently response is assessed at 6 months — a long wait with significant toxicity. The framework predicts earlier classification via combined A-score. Falsifiable in MDS cohorts at major cancer centers.

ISSUE 003 ADDENDUM — PROGENITOR IN THE ATLAS

11 cell types: CMP, Erythrocyte_progenitors, GMP, L-MPP, MEP, MPP, Neulm, OPC, erythroblast, megakaryocyte, nRBC

Identity loci: 54,704 (methyl) · **H_min by substrate:** methyl 0.852216 · nucl 0.972790 · fuzz 0.961900 · wps 0.988046 · frag 0.808978 · **methyl ceiling** 1.1734

Healthy age-matched band (age_reference_matrix.json). A = 1.0 is the commitment line; healthy reads on this band.

Age (midpoint)	n	A mean	p10	p90	β mean	Source
4	7	0.9557	0.9250	0.9863	0.748	Progenitor pediatric
14	9	0.9611	0.9310	0.9913	0.745	Progenitor
24	12	0.9666	0.9345	0.9986	0.742	Roadmap E035
34	15	0.9701	0.9384	1.0018	0.740	Roadmap E035
44	18	0.9736	0.9400	1.0073	0.738	Roadmap+aging
54	22	0.9789	0.9435	1.0142	0.735	Jaiswal 2014 CHIP-neg
64	25	0.9840	0.9470	1.0210	0.732	Jaiswal 2014
74	20	0.9907	0.9523	1.0291	0.728	Jaiswal+progenitor
84	10	0.9973	0.9576	1.0370	0.724	Extrapolated
95	3	1.0038	0.9629	1.0447	0.720	Extrapolated

CELL IDENTITY & CLINICAL CONTEXT

Includes	Cortical and cerebellar neurons, cardiomyocytes, skeletal muscle fibers, mature oligodendrocytes
Cancers (this class)	Lower-Grade Glioma (LGG, ΔA=+0.239), Glioblastoma (GBM, ΔA=+0.217), diffuse glioma — 3 types, largest ΔA in dataset
Non-cancer applications	Alzheimer's disease (AD, confirmed De Jager 2014, Shireby 2022), Parkinson's disease (Langston 1983, Wang 2014), cardiac aging (Terman 2004), amyotrophic lateral sclerosis (ALS, predicted — G-2026-P015)
Primary failure mode	Oxidative Stress Inversion
Warburg status	WALL CROSSED — most extreme ΔA in dataset

COMMENTARY

Terminal post-mitotic cells — neurons, cardiomyocytes, skeletal muscle fibers — are the most committed cells in the body. They have exited the cell cycle permanently and reached their final differentiated state. Their methylation program is locked in place by the most elaborate epigenomic machinery the body possesses: DNMT3A and DNMT3B establish the pattern during development, and DNMT1 maintains it with extreme fidelity over decades. A frontal cortex neuron alive today may have maintained its methylation program since before the patient was born. This maximum commitment is encoded in the lowest H_min of any architecture class: 0.772837. Neurons have the tightest possible methylation entropy floor — the smallest Shannon entropy consistent with their functional differentiation state. The global floor reference (H_min = 0.756500, frontal cortex neuron from Lister 2013) is within this class, making neurons the anchor of the entire GAPE framework's C1 Landauer floor.

The consequence for cancer is stark. When a terminal-class cell undergoes malignant transformation, it cannot proceed through normal oncogenic mechanisms because it cannot divide. Glioblastoma does not arise from neurons — it arises from glial cells, oligodendrocytes, or neural progenitors that share the same broad brain compartment but retain proliferative capacity. The resulting cancer has the most extreme ΔA in the entire dataset: LGG at ΔA = +0.239, GBM at ΔA = +0.217 (Ceccarelli 2016 Cell, TCGA methylation). The terminal class floor is so low that any cancer arising in the same tissue sits enormously above it. There is no subtlety in the GAPE signal for brain cancer. The physics is extremely loud. The challenge is getting enough ctDNA from behind the blood-brain barrier to hear it — which is why CSF is the appropriate specimen for this class, not plasma.

A critical substrate-saturation finding emerges specifically for this class and must be stated honestly. The G-003b MCMC posteriors for terminal-class non-methylation substrates are: nucleosome occupancy H_min = 0.9920, fuzziness H_min = 0.7370, WPS H_min = 0.9589, fragment size H_min = 0.6249. Note that fuzziness H_min = 0.7370 is slightly below the methylation floor (0.7728). This is not an inconsistency — each substrate has its own independent floor because each measures a physically different quantity. Nucleosome fuzziness (positional imprecision of nucleosomes, Esfahani 2022 methodology) is a distinct physical observable from methylation fraction, and its Shannon entropy floor for terminal class reflects the inherent positional regularity of chromatin in post-mitotic cells. The five substrate H_min values are independent physical constants from the G-003b MCMC, not derivations of each other. Because Shannon entropy is bounded above by 1.0 at β = 0.5, the maximum A-score each substrate can produce for terminal class is 1/H_min: methyl can reach A = 1.294, fuzz can reach A = 1.357, frag can reach A = 1.600, but nucl saturates at A = 1.008 and WPS saturates at A = 1.043. This means nucleosome occupancy and WPS physically cannot cross the BREACH threshold (A ≥ 1.10) for terminal-class samples. At glioma-level methylation departures (A > 1.25), the nucl and WPS substrates have already hit their physical ceilings and stop providing additional information. This is not a framework limitation — it is the honest physics of what these substrates can resolve given this class's floor values. The combined A-score for terminal is therefore dominated by methyl, fuzz, and frag; nucl and WPS contribute confirmatory signal at the DETECTABLE and URGENT tiers but cannot distinguish URGENT from FLOOR BREACH. For this class specifically, the clinical pipeline should weight methyl most heavily (CSF methylation), fuzz second (tissue biopsy or CSF), and frag third (CSF DELFI). Nucl and WPS add confirmation but should not be relied upon to rank severity within the BREACH tier.


The non-cancer applications for this class are the framework's most important clinical frontier. Alzheimer's disease is terminal-class thermodynamic failure: published data places healthy neurons at β = 0.786 (A = 0.969, NORMAL tier), low-AD neuropathology at β = 0.753 (A = 1.043, MARGINAL), and high-AD neuropathology at β = 0.744 (A = 1.062, DETECTABLE). β values pipeline-adjusted from the two largest AD methylation cohorts: De Jager 2014 (ROSMAP, n = 740) and Shireby 2022 (Brains for Dementia Research, n = 631). Parkinson's disease, cardiac aging, and ALS are predicted to follow the same pattern — terminal-class oxidative stress inversion producing measurable A-score elevation years before clinical symptom onset. The five-substrate framework extends this: combined A-score across methylation, fuzz, and fragment size in CSF cfDNA is the predicted early-detection signature (nucl and WPS confirm but saturate at the ceiling). The critical clinical distinction between AD and glioma — both show terminal-class floor departure — is magnitude: AD shows ΔA = 0.019–0.084 (MARGINAL to DETECTABLE), glioma shows ΔA = 0.217–0.239 (FLOOR BREACH, Ceccarelli 2016 Cell). Same class, same mechanism, orders of magnitude different rate of drift. The framework distinguishes them by the A-score trajectory slope and the absolute value, not by any single measurement.

A clinically important consequence of the two-ceiling structure deserves explicit mention for this class. Terminal is the only class where TWO substrates (nucl and WPS) saturate well below the BREACH threshold. This produces a useful diagnostic regime split. For neurodegenerative disease — AD, ALS, Parkinson's, cardiac aging — the A_methyl signal peaks in the 1.01–1.09 range, and the correlated shifts in nucl and WPS are modest (β ~0.58–0.60). Neither substrate saturates in this range. In the neurodegenerative regime, all five substrates act as quantitative instruments; the full A_combined and trajectory slope serve as the measurement tools, and saturation is not reached. For glioma (LGG/GBM), A_methyl enters

the 1.20–1.29 range (deep FLOOR BREACH), and the associated β shifts push both nucl and WPS to their maximum-entropy points ($\beta \approx 0.500$). Both saturate. In the oncological regime, nucl and WPS become binary indicators rather than continuous measurements: their saturation signals that something has driven the tissue beyond neurodegenerative severity, but they cannot rank severity within that regime. Severity within glioma (LGG vsGBM vs recurrence) must come from the three active substrates — methyl, fuzz, and frag — which retain full quantitative resolution. This is why the card reports two combined values: A_combined (all 5, legacy formula) and A_active (non-saturated, the progression tracker). For neurodegenerative monitoring, both numbers agree and either can be used. For oncological monitoring, only A_active tracks real progression; A_combined is dragged down by the ceiling values and will flatten prematurely.

FIVE-SUBSTRATE FIDELITY GAUGE

The gauge below is the central image of Issue 002. A single horizontal bar represents the A-score axis — from A = 0.90 (below the floor) through A = 1.00 (exactly at the Terminal class floor, H_min = 0.7728) up through the clinical threshold zones: MARGINAL at A = 1.01, DETECTABLE at 1.05, URGENT at 1.07, and FLOOR BREACH at 1.10. Five dots are plotted on this single axis for a healthy reference sample (Frontal cortex neuron, markers above the bar), and five more dots for a disease reference (Lower-Grade Glioma (Ceccarelli 2016), markers below). Each dot is one substrate: methylation, nucleosome occupancy, nucleosome fuzziness, WPS, or fragment size (DELFI). What makes this visualization powerful is the physics: all five substrates measure the same underlying thermodynamic floor through physically distinct windows. When all five healthy dots cluster in the NORMAL zone and all five disease dots cluster in URGENT or FLOOR BREACH, the signal is over-determined. Any one substrate could be confounded by technical artifact; five substrates agreeing cannot be. This is the "less blurry" advantage: $\sqrt{5}$ noise reduction for free.




Look at the spread in each direction. If the healthy five-dot cluster sits tightly together near A = 0.97–1.00, that indicates the reference population is at its architecture floor — precisely where the framework predicts healthy terminal cells should sit. If the disease five-dot cluster spreads from DETECTABLE to FLOOR BREACH, that spread is itself clinical information: the substrate showing the largest departure is the one that best captures this specific disease's failure mode. For this class, the primary failure mode is Oxidative Stress Inversion, and the substrate ranking shown further down the card explains which of the five substrates best reveals it.

SUBSTRATE-BY-SUBSTRATE BREAKDOWN

Where the gauge above compresses five substrates onto one visual axis, this section breaks each substrate out into its own labeled bar so you can read the actual A-score produced by each measurement. The physics is the same across all five: take the raw laboratory value (a β methylation fraction, a nucleosome occupancy probability, a WPS score, etc.), compute its Shannon entropy $H(\text{value})$, then divide by the class-specific and substrate-specific H_{min} to get a dimensionless A-score comparable across substrates. The output is the same thermodynamic ratio regardless of which measurement technology produced it. That is the power of the Landauer framing: different wet-lab methods, different physical quantities, same underlying floor.

The formula at the foot of this section — $A_{\text{combined}} = \Sigma(AUC_i \times A_i) / \Sigma(AUC_i)$ — combines the five individual A-scores into a single weighted score, using each substrate's published detection AUC as its weight. Higher-AUC substrates contribute more to the combined result. For terminal cells, the healthy combined A-score below should sit at approximately 0.97 (NORMAL tier); the disease combined A-score should sit above 1.05 (DETECTABLE or higher). The difference between them — the ΔA combined — is the signal available to a blood draw that analyzes all five substrates simultaneously. A single-substrate assay sees a smaller signal at higher noise; the five-substrate combination sees the same signal at approximately $\sqrt{5}$ lower noise. Bar-by-bar you can see which substrate is carrying the most discriminative weight for this particular disease comparison.

HEALTHY REFERENCE — Frontal cortex neuron



DISEASE REFERENCE — Lower-Grade Glioma (Ceccarelli 2016)




Combined A — healthy (all 5)	0.97003 (NORMAL)
Combined A — disease (all 5)	1.17053 (FLOOR BREACH)
ΔA combined (5 substrates)	+0.20049 vs methylation-only ΔA =+0.3154
Active A — healthy (non-saturated)	0.97003 (NORMAL) [5/5 active]
Active A — disease (non-saturated)	1.26166 (FLOOR BREACH) [3/5 active]

ΔA active (real progression signal)	+0.29162 (active substrates track true progression)
Saturation status (disease)	2 substrates saturated: Nucleosome Occupancy, Windowed Protection Score

Two combined values are shown because some substrates physically cannot resolve past their class ceiling ($A_{\text{max}} = 1/H_{\text{min}}$). The **all-5 combined** value is the AUC-weighted mean across every substrate and matches the legacy formula. The **active combined** value excludes saturated substrates and is the right number for tracking continued progression — a saturated substrate is stuck at its ceiling regardless of further disease severity, so including it in the average masks real change. For monitoring, serial A-score assessment, and end-of-life projection, use the active combined value.

DISEASE SIGNATURE COMPARISON — healthy neuron vs Alzheimer's vs LGG vs GBM

Four conditions on a single chart, all in the terminal class. β values come from matched primary sources: healthy neuron (Lister 2013 frontal cortex); Alzheimer's average of early and late neuropathology (De Jager 2014 ROSMAP $n=740$, Shireby 2022 $n=631$); Lower-grade glioma and Glioblastoma methylation β from Ceccarelli 2016 Cell TCGA analysis ($n=516$ and $n=149$ respectively). The A-score uses the unfloored formula $A = H(\beta)/H_{\text{min}}$ without clamping; healthy cells sit slightly below $A = 1.00$ (around $A \approx 0.97$) because β near the healthy reference produces $H(\beta)$ slightly below the MCMC-derived H_{min} central estimate. The $A = 1.00$ line represents the architectural commitment point, not a mathematical floor. Disease departure above $A = 1.10$ crosses the architectural ceiling into the cancer range. All four conditions sit on the same reference floor $H_{\text{min}} = 0.7728$, but show vastly different departure magnitudes. Data-availability disclosure: $\beta_{\text{nucl}} = 0.500$ and $\beta_{\text{wps}} = 0.500$ for LGG and GBM in the chart below are placeholders. Ceccarelli 2016 reports methylation only, and Corces 2018 TCGA ATAC-seq provides GBM/LGG chromatin accessibility peaks but not promoter-level β values compatible with the A_{nucl} formulation. These two substrates therefore display at their class ceilings ($A_{\text{nucl}} = 1.008$, $A_{\text{wps}} = 1.043$) pending reanalysis of the Corces raw bigWig files at terminal-class architecture CpG loci (G-2026-P023). For this class, nucleosome occupancy saturates at $A \approx 1.008$ and WPS saturates at $A \approx 1.043$ by physics — once β reaches 0.5 the Shannon entropy is maximum and the ratio cannot climb further — so glioma-level departures above these ceilings cannot be resolved on these substrates even when measured data becomes available. Methyl, fuzz, and frag continue to track the full signal past the ceiling for this class. See the Post-breach Trajectory subsection below for the condition-by-condition analysis of what happens once A crosses 1.10.




What happens past the ceiling is condition-specific. The A-score magnitude alone does not track severity once the ceiling is crossed — different diseases take different post-breach paths that the full five-substrate divergence pattern reveals. The subsection below walks through this class's post-breach trajectory: which substrates keep carrying signal, where the Warburg metabolic lock-in happens, where the glucose inversion point sits, and where the framework currently has high confidence versus where the open research questions remain.

POST-BREACH TRAJECTORY — WHAT HAPPENS AFTER A CROSSES 1.10 (TERMINAL CLASS)

The pre-breach bar above answers "where is this patient relative to healthy?" — a detection instrument. Once A crosses the ceiling at 1.10, a different instrument is needed: one that answers "what is happening to this patient's cells now, and which therapeutic windows are still open?" This subsection is that instrument for the terminal class.

The zone language below — Warburg boundary, glucose inversion, point of no return — applies only to post-breach cells. These terms are not used on the pre-breach bar, and the qualitative A-value anchors (~1.15, ~1.25, ~1.40+) are boundaries between therapeutic windows, not additional diagnostic tiers. Specific A-values are filed as prospective predictions to be validated against survival-stratified cohorts.



WHICH SUBSTRATES CARRY POST-BREACH SIGNAL — TERMINAL CLASS

Terminal-class physics constrains which substrates continue to report once the ceiling is crossed. Three substrates carry signal throughout all four post-breach zones; two are saturated at or near the ceiling and report no further progression.

Substrate	Ceiling (A)	Post-breach status
Methylation	1.29	Carries signal throughout all four zones
Fuzziness	1.36	Carries signal throughout all four zones
Fragment size	1.60	Carries signal throughout all four zones
Nucleosome occupancy	1.008	Saturated at ceiling — no further signal post-breach
Windowed protection	1.043	Saturated at ceiling — no further signal post-breach

CONDITION-BY-CONDITION TRAJECTORY

Healthy neuron — reference, no breach

Known: A_{combined} ≈ 0.97 at the class reference; all five substrates cluster near the architectural floor. The post-breach machinery does not apply — the cell is operating within its committed architecture. No therapeutic intervention is indicated; the framework's use here is population-relative positioning (age-matched percentile). Baseline for all three disease trajectories that follow.

Alzheimer's disease (AD) — A ≈ 1.066, DETECTABLE tier

Known: AD shows a gradual climb in terminal-class A-score with age that De Jager 2014 (ROSMAP n=740) and Shireby 2022 (n=631) both document. At the Issue 002 reference point (average of early and late neuropathology) combined A sits at ~1.066 — detectable but below the ceiling. AD remains in the metabolic-window zone on the bar above. This is thermodynamically consistent with the known role of mitochondrial dysfunction and oxidative stress in AD progression (Wang 2014, Terman 2004). Metabolic intervention — ketogenic diet, mitochondrial support, DNMT1 cofactor availability — targets the binding constraint here, not downstream pathology. VAL-040 (April 2026, multi-class drift cascade) extended this: AD is not confined to terminal-class drift at the cellular thermodynamic level. Four of eight architecture classes show elevation in AD cohorts — terminal (brain cortex, expected), immune (peripheral blood, novel), secretory (pancreatic islet via T2D-AD comorbidity), and stromal (cerebral vasculature). Seven of seven tissue-class combinations show severity gradient (late-stage > early-stage AD). The framework therefore predicts that peripheral blood (immune class) A-score alone carries AD-susceptibility signal, without requiring CSF access — a finding that extends the practical clinical reach of the framework beyond terminal-specimen requirements.

Unknown: whether the terminal-class A-score trajectory over 3-10 years predicts conversion from mild cognitive impairment (MCI) to clinical AD, and whether the rate of A climb distinguishes rapid progressors from stable MCI.

Test that closes the gap: G-2026-P024: ROSMAP has longitudinal blood samples archived for a subset of the n=740 cohort with subsequent AD diagnosis outcomes. Framework prediction: in subjects with archived blood at baseline (MCI or cognitively normal) and subsequent AD diagnosis, terminal-class A_{active} trajectory slope will discriminate converters from non-converters with AUC ≥ 0.75 using samples 3-5 years before diagnosis. Falsifiable against the archived ROSMAP series.

Lower-grade glioma (LGG) — A ≈ 1.171, CROSSED CEILING

Known: LGG sits in the metabolic-window to structural-only zone, past the ceiling but below the glucose-inversion boundary. Methyl ($A \approx 1.29$), fuzz, and frag all carry full signal; nucl and wps are saturated at their ceilings and contribute no further information. LGG is characterized by IDH mutation ($\geq 80\%$ of cases) which generates the oncometabolite 2-hydroxyglutarate, inhibiting TET enzymes and producing the hyper-methylator phenotype (G-CIMP) that Ceccarelli 2016 defined. The hyperentropy signal here is driven by methylation redistribution, not global loss.

Unknown: whether IDH-mutant vs IDH-wild-type LGG subtypes show distinct post-breach trajectories; whether MGMT promoter methylation status predicts A_methyl response under temozolomide; whether the A-score trajectory discriminates stable grade II astrocytoma from imminent transformation to grade III/IV. LGG grade II has 10-year survival near 50% and catching transformation early changes management.

Test that closes the gap: G-2026-P025: Prospective cohort of 150 IDH-stratified LGG patients with serial cfDNA at diagnosis, 3, 6, 12, and 24 months, correlated with MRI progression and transformation events. Framework prediction: A_active trajectory slope over the first 12 months will identify imminent transformation (grade III/IV) with sensitivity ≥ 0.75 at specificity 0.85, outperforming current serial-MRI watch-and-wait.

Glioblastoma (GBM) — $A \approx 1.142$, CROSSED CEILING

Known: GBM's combined A-score is lower than LGG's (1.142 vs 1.171) despite being clinically more aggressive. This is not a framework failure — it reflects a different post-breach trajectory. GBM's methylation pattern re-commits toward a tumor-identity program ($\beta_{\text{methyl}} \approx 0.40$, further from 0.5) rather than dissolving toward distributional randomness (LGG $\beta_{\text{methyl}} \approx 0.45$, closer to 0.5). In $H(\beta)$ space this produces lower entropy and thus lower A-score for GBM despite greater biological departure from the neuronal reference. **A-score magnitude does not equal clinical aggressiveness post-breach** — the substrate divergence pattern and the A_active trajectory carry the severity signal, not the combined A value.

Unknown: whether the post-breach A trajectory under Stupp protocol predicts progression-free survival; whether MGMT-methylated and MGMT-unmethylated GBM subtypes trace distinguishable A-score paths; whether metabolic intervention (ketogenic diet, medical ketosis) measurably flattens the A-score trajectory in treatment-naïve GBM.

Test that closes the gap: G-2026-P025: Retrospective reanalysis of archived cfDNA from RTOG 0525 dose-intensification trial ($n=833$ GBM patients with MGMT stratification and serial imaging) plus matched nucleosome data from the Corces 2018 TCGA-ATAC-seq cohort. Framework prediction: post-breach A_active trajectory during the first two cycles of adjuvant TMZ will predict 6-month progression-free survival with $AUC \geq 0.70$, stratified by MGMT status.

WHAT THE FRAMEWORK CAN AND CANNOT SAY TODAY

The framework has high confidence in three things for terminal-class post-breach: (1) methyl, fuzz, and frag carry continued signal past the ceiling while nucl and wps saturate; (2) the Warburg boundary at $A \approx 1.05\text{--}1.07$ separates a metabolic-intervention zone from a structural-intervention zone — documented across the 27-cancer TCGA validation; (3) a glucose inversion exists past the Warburg boundary where adding glucose accelerates rather than restrains disease — documented mechanistically in Issue 001 and consistent with clinical experience with TPN and high-glucose supportive care in end-stage cancer patients.

The framework cannot yet say with certainty: (1) the exact numerical A-value of the glucose inversion point for specific cancer types (placeholder $A \approx 1.25$ pending validation); (2) the exact point-of-no-return A-value (placeholder $A \approx 1.40+$); (3) whether the three post-breach archetypes (Classical Warburg Progression, Inversion Recovery Path, Hypomethylation Re-commitment) all manifest within the terminal class. These tests could not be run before IAM because the framework did not exist to generate them.


This is the scientific posture the framework commits to: definitive where the evidence is definitive, open where it is open, and specific about the experiments that would close each gap. Predictions filed for this class: G-2026-P023 through G-2026-P025. A reviewer who wants to falsify any claim here can do so against the named cohorts. A collaborator who wants to extend the framework has specific hypotheses to test.

THREE-COMPONENT DECOMPOSITION (C1/C2/C3) — PER SUBSTRATE

Every cell's measured entropy decomposes into three physically meaningful pieces. C1 is the universal Landauer floor: $H_{\text{MIN_GLOBAL}} = 0.756500$, measured in frontal cortex neurons by Lister 2013. Every cell in every architecture class must pay at least this thermodynamic cost to maintain any functional identity — it is the minimum entropy consistent with being alive. C2 is the class-specific overhead: the additional entropy the terminal class must carry above the universal floor to encode its particular architecture. For this class, C2 represents approximately 2.1% of the healthy reference entropy. C3 is the accessible gap: $\max(0, H_{\text{actual}} - H_{\text{min}})$. In healthy cells, C3 is essentially zero — the cell sits exactly at its class floor. When C3 grows, the cell is departing from its architecture class. Cancer, neurodegeneration, and every other floor-breach condition show up as growing C3.

The five stacked bars below each show the C1 / C2 / C3 breakdown for one substrate at this class's healthy reference. Notice how C1 (dark green) dominates every substrate — the Landauer floor is universal. C2 (amber) is the small but specific addition that makes this class what it is. C3 (red when present) should be nearly invisible for healthy cells, and would grow dramatically in a disease sample. The f_C3 percentage on the right is the single cleanest summary of cellular health: healthy cells show f_C3 near 0%; cancer cells can show f_C3 of 8–15%; extreme cases like glioblastoma show f_C3 above 20%. This is not a statistical threshold calibrated from disease data. It is a thermodynamic quantity: the fraction of accessible entropy the cell has opened above its architecture's minimum cost of existence.





BEST TESTING METHOD RANKING — CLINICAL UTILITY FOR THIS CLASS

Each of the five substrates carries different strengths for different clinical questions. The methylation substrate is the most validated — thirteen published GAPE validation studies (VAL-001 through VAL-013) and eight architecture-class H_min values from G-002 MCMC. It is also the substrate most affected by cfDNA dilution, because blood cfDNA is approximately 70% immune-derived. Fragment size (DELF) is the most forgiving of dilution — short-fragment enrichment persists even when tumor-derived cfDNA is a small fraction of total plasma DNA. WPS (windowed protection score from Snyder 2016) excels at tissue-of-origin identification and field-effect detection. Nucleosome occupancy and fuzziness add chromatin-level information that methylation alone cannot reveal.

The ranking below is specifically for the terminal class. The order reflects three factors: (1) the substrate's published single-substrate AUC for diseases in this class; (2) the biology of this class's primary failure mode (Oxidative Stress Inversion); and (3) the practical constraint of cfDNA contribution in the relevant specimen (plasma for most classes, CSF for terminal, tissue biopsy where cfDNA is too dilute). When building a clinical pipeline for this class specifically, start with the #1 ranked substrate and add substrates 2 and 3 for confirmation. Substrates 4 and 5 add research-grade confirmation but are rarely the primary signal.

#	Substrate	Best for	Rationale
1	Methylation	Neurodegeneration and glioma (CSF, not plasma)	AD terminal-class validation (De Jager 2014 n=740, Shireby 2022 n=631) confirmed. CSF cfDNA bypasses blood-brain barrier and delivers the full signal.
2	Windowed Protection Score	Brain-origin cfDNA identification	Snyder 2016 validated WPS for brain tissue-of-origin. Second-best substrate for CNS applications. Plasma usable with extensive deconvolution.
3	Fragment Size (DELF)	Neural cfDNA fragmentomics	Neural cfDNA fragment distribution differs from systemic cfDNA. Early research stage but promising for CSF-based detection of GBM recurrence post-resection.
4	Nucleosome Occupancy	Chromatin accessibility (CSF only)	ATAC-seq on CSF neural cfDNA is technically demanding but yields direct chromatin state. Research-grade only — very low input.
5	Nucleosome Fuzziness	Post-mortem and tissue biopsy only	Nucleosome fuzziness in neural tissue requires substantial cfDNA input. Plasma levels (0.5% cfDNA) are too low. Use tissue-based assays.

BODY TEMPERATURE SCALING — α = 2.0 LANDAUER CORRECTION

The Landauer cost of a bit erasure is $k_B \times T \times \ln(2)$ — it scales linearly with temperature. Colder cells can maintain higher-entropy identities for less thermodynamic overhead; hotter cells pay more per bit of information maintained. The GAPE framework captures this with a simple scaling law: $H_{\min}(T) = H_{\min}(37^{\circ}\text{C}) \times (T_{\text{body}} / 310.15\text{ K})^{\alpha}$, with $\alpha = 2.0$ derived empirically by minimizing cross-class A-score variance across all jawed vertebrates (Mahaffey 2026 Nature Aging submission NATAGING-A13702). The $\alpha = 2.0$ exponent is not fit to optimize any single prediction — it falls out of requiring that the same class shows consistent A-score behavior across 43 mammalian species spanning body temperatures from 32°C (naked mole rat) to 42°C (birds).

The table below shows what the terminal class floor H_min looks like at seven representative body temperatures. At 37°C (human reference, highlighted row), H_min = 0.7728. At 42°C (birds), the floor is higher because the Landauer tax is larger. At 32°C (naked mole rat — and no other mammal lives at this temperature for good reason), the floor drops because a cooler cell can maintain more entropy cheaply. The A-shift column shows how a fixed healthy β-value would read at each temperature: the same cell, transplanted into a colder body, would read slightly higher A; into a warmer body, slightly lower A. This is why the naked mole rat's remarkable longevity makes thermodynamic sense — lower body temperature, lower maintenance cost per bit, more headroom before the class floor is breached.


T_body	Species example	H_min(T)	β_healthy equivalent	A-shift vs 37°C
42°C	Birds (chicken, finch)	0.7980	0.786	A × 0.969
39°C	Rodents (mouse, rat)	0.7828	0.786	A × 0.987
37°C	Humans (REFERENCE)	0.7728	0.786	A × 1.000 (anchor)

T_body	Species example	H_min(T)	β_healthy equivalent	A-shift vs 37°C
35°C	Hibernating bats	0.7629	0.786	A × 1.013
32°C	Naked mole rat	0.7481	0.786	A × 1.033
25°C	Reptiles (anole lizard)	0.7142	0.786	A × 1.082
15°C	Fish (salmon, cold-water)	0.6671	0.786	A × 1.159

HEALTHY AGING TRAJECTORY — AGE-STRATIFIED REFERENCE A-SCORES

Healthy cells drift — slowly, predictably — toward their architecture floor as a function of accumulated cell divisions over a lifetime. For the terminal class, the drift rate is approximately 0.8% per generation — meaning each cell cycle carries a 0.8% probability of a methylation error that raises the cell's A-score by a measurable amount. Over decades, these errors accumulate. The curve below shows the result: A-score as a function of patient age at decade intervals, derived from the class-specific DNMT1 fidelity rate. At age 20, the reference healthy A-score sits well below 1.00; by age 80, it has drifted upward into the MARGINAL to DETECTABLE range. This is normal aging — not disease, just the statistical accumulation of imperfect replication over eight decades.

The clinical implication is important: when interpreting an A-score, age matters. A patient at age 30 with A = 1.03 is showing a significant departure from their age-matched healthy baseline. The same A = 1.03 at age 75 might be entirely consistent with age-expected drift. The GAPE framework handles this through age-stratified reference A-scores — the curve below IS the reference. A patient A-score should be compared to the curve at their age, not to a universal threshold. Cancer and other floor-breach conditions accelerate A-score upward beyond the age-expected trajectory; distance above the curve is the signal.



VERTEBRATE LIFESPAN CONTEXT

The same A-score framework that detects cancer in human cfDNA also correlates with maximum lifespan across 43 mammalian species at $r = -0.9018$ ($p = 1.6 \times 10^{-16}$, Mahaffey 2026 Nature Aging submission NATAGING-A13702). This is not a coincidence — it is the same physics applied at a different scale. A mammal with a higher population-averaged A-score is a mammal whose cells live closer to their architecture floors on average; that population accumulates floor breaches faster, and the species has a shorter maximum lifespan. The A = 1.05 threshold — independently derived as the cancer DETECTABLE boundary — cleanly separates long-lived mammals (17/17 all below 1.05) from short-lived mammals (11/11 all above 1.05) with complete accuracy. Same threshold. Same physics. Different timescale.

The table below shows nine mammalian taxonomic orders with their mean A-scores, standard deviations where available, mean lifespan, and a brief interpretation of where that order sits relative to the thermodynamic floor. Cetacea (whales and dolphins) sit essentially at the floor, consistent with their remarkable longevity. Primates including humans sit slightly above the floor. Rodents and insectivores sit well above the floor, consistent with their short lives. This class's reference A-score (0.969) places it in the context shown below — use this to understand how this particular cell architecture relates to the broader mammalian lifespan distribution.

Order	N	Mean A	Lifespan (yr)	Interpretation
Cetacea	5	0.997 ±0.016	99	At thermodynamic floor
Proboscidea	1	0.987 —	70	At floor
Primates	6	1.007 ±0.015	49	Slightly above floor
Artiodactyla	4	1.015 ±0.011	30	Slightly above floor
Chiroptera	4	1.041 ±0.007	35	Above floor (longevity outliers)
Carnivora	9	1.053 ±0.023	28	Above floor
Lagomorpha	2	1.114 ±0.002	8	Well above floor

Order	N	Mean A	Lifespan (yr)	Interpretation
Rodentia	6	1.125 ±0.018	9	Well above floor
Insectivora	1	1.157 —	2	Furthest from floor

INTERVENTION LEVERS — RANKED BY EXPECTED IMPACT

If a patient's A-score is elevated — MARGINAL, DETECTABLE, or URGENT — the clinical question is what can be done. The GAPE framework does not prescribe treatment; that remains the physician's judgment. But the framework does identify which intervention categories have the mechanistic strength to move this particular class's A-score back toward the floor. The five categories below are the standard anti-aging and cancer-adjacent interventions: senolytics (clearing senescent cells), metabolic restoration (NAD+, caloric restriction, exercise), epigenetic restoration (DNMT/TET modulators, HDAC inhibitors), reprogramming (cyclic Yamanaka factors), and checkpoint stringency (cell cycle control, immune checkpoint blockade).

The ranking below is specifically for the terminal class and its primary failure mode (Oxidative Stress Inversion). Impact Level 1 (Dominant) means the intervention directly addresses the class's failure mechanism and is the first lever to try. Level 2 (Strong) means substantial mechanistic support for A-score restoration in this class. Level 3 (Moderate) means helpful but not addressing the binding constraint. Level 4 (Limited) means the intervention works for other classes but does not target this class's specific biology. Level 5 (Not applicable) means the intervention category is biologically incompatible with this class — e.g., senolytics cannot act on post-mitotic terminal cells that do not become classically senescent, and reprogramming cannot be applied to terminal cells without losing the identity that defines them.

Impact	Category	Mechanism	Rationale
2 STRONG	NAD+ / Mitophagy	metabolic	Strong — NAD+/mitophagy directly address the oxidative stress inversion
3 MODERATE	Epigenetic Restoration	epigenetic rx	Moderate — DNMT1/TET restoration helps; CNS delivery is the bottleneck
4 LIMITED	Senolytics	senolytics	Limited — neurons do not become classically senescent
4 LIMITED	Checkpoint Stringency	checkpoint	Not applicable — post-mitotic, no cell cycle checkpoints
5 N/A	Reprogramming	reprogramming	Not applicable — terminal class cannot be reprogrammed without losing identity

CANCER PANEL — ΔA RANKED

The TCGA cancer validation is where the framework earns its keep. For every cancer type in this class, we compute two numbers from published TCGA methylation data: A_healthy (from matched normal tissue) and A_tumor (from cancer samples in the same study). The ΔA = A_tumor – A_healthy is the signal the framework predicts should appear — and it does. Thresholds for tier assignment: A = 1.05 DETECT, A = 1.07 URGENT, A = 1.10 FLOOR BREACH. These thresholds are not fit to the cancer data; they were derived from the healthy-vs-age statistics and happen to align with the cancer validation cleanly. The bars below are ordered by |ΔA| descending so the cancers with the largest signal appear first.

Note that ΔA ranges typical for this class matter. Terminal-class cancers (glioma, glioblastoma) show ΔA ≈ 0.22–0.27 — the largest in the entire 28-cancer TCGA dataset. Cycling epithelial cancers show ΔA ≈ 0.13–0.19 — moderate but clearly detectable. Stromal and progenitor-lineage cancers show ΔA ≈ 0.10–0.17. Each cancer below is cited to its primary TCGA publication and shows sample size. Direct links from A-score to TCGA evidence — no black boxes between physics and clinical data.

#1	Lower Grade Glioma (LGG) n=516 · Ceccarelli 2016 Cell	ΔA=+0.2390A=1.285 FLOOR
#2	Glioblastoma (GBM) n=149 · Ceccarelli 2016 Cell	ΔA=+0.2170A=1.256 FLOOR

CORE METRICS — PUBLISHED AND DERIVED

Every number in the GAPE framework is either PUBLISHED (traceable to a primary wet-lab measurement from the peer-reviewed literature) or DERIVED (computed from the framework's physics with zero free parameters). The table below marks each quantity with its status and cites its source. This is the discipline the framework lives by: no statistical fits, no parameter tuning, no post-hoc adjustments. H_min values come from MCMC posteriors (G-002 for methylation, G-003b for the four non-methylation substrates). Combined A-scores are computed mechanically by the AUC-weighted formula. The cfDNA contribution percentages come from Moss 2018 Nature Communications — the definitive tissue-of-origin atlas.

Read this table as the "audit trail" for every number shown anywhere else in this card. When the framework says A_healthy = 0.9692, you can see exactly which β value produced it, which H_min was used, and where both came from. The n_bio field is currently marked PRELIMINARY because absolute n_bio values await the G-007 MCMC run; the class ordering has been confirmed (Spearman ρ = 0.905, p = 0.002 against commitment level) but the absolute values carry that caveat transparently.

Metric	Value	Source
cfDNA contribution (plasma)	0.5%	Moss 2018 Nat Commun — PUBLISHED
H_min methylation (class)	0.772837	G-002 MCMC posterior — DERIVED
H_min methylation σ (MCMC)	± 0.001100	G-002 17-chain MCMC R-hat<1.001 — DERIVED
n_bio (class-specific)	24.5	PRELIMINARY — ordering confirmed ($p=0.905$, $p=0.002$); absolute pending G-007
Healthy drift rate	0.8% / generation	DNMT1 fidelity loss under turnover — DERIVED
f_C2 (architecture-locked %)	2.1%	C2 fraction at healthy reference — DERIVED
Healthy A (methylation)	0.9692 (NORMAL)	$\beta=0.786$ — DERIVED
Disease A (methylation)	1.2846 (FLOOR BREACH)	$\beta=0.450$ — DERIVED
Healthy combined A (5 subs)	0.9700 (NORMAL)	$\Sigma(AUC_i \times A_i) / \Sigma(AUC_i)$ across 5 substrates — DERIVED
Disease combined A (5 subs)	1.1705 (FLOOR BREACH)	$\Sigma(AUC_i \times A_i) / \Sigma(AUC_i)$ across 5 substrates — DERIVED
ΔA combined	+0.2005	Combined signal — DERIVED

DATED PREDICTIONS FOR THIS CLASS

A prediction is not a claim unless it can fail. Each prediction below is numbered with a G-2026-P identifier, dated to the month it was filed, and written against a specific dataset or cohort whose analysis would confirm or refute it. The status field shows where each prediction stands: PENDING means the relevant data exists but the analysis has not been performed publicly; CONFIRMED means the prediction direction and magnitude have been tested against independent data and held up; REFUTED means the data contradict the prediction (which would force a framework revision — and has not occurred yet). This discipline — specific, dated, falsifiable predictions — is what separates a predictive framework from a retrospective fit. Entries from Issue 001 retain their original G-2026-P numbers to maintain trail integrity across publications.

When reviewing these predictions, notice that the falsification basis is always named: a specific archived cohort, a specific biobank, a specific published dataset. These are not vague "more research is needed" statements. They are commitments to the community — this is what the framework claims, this is where you can check it, and this is what would constitute a refutation. The predictions below are for the terminal class specifically, filed as part of this publication or carried forward from Issue 001.

G-2026-P006

Originally filed

CONFIRMED (partial)

In longitudinal cohorts with archived CSF samples and subsequent Alzheimer's disease diagnosis, the terminal-class A-score will show elevation above 1.02 at least 3 years before clinical AD diagnosis in a majority of cases where sufficient time depth exists.

Basis: AD methylation literature (De Jager 2014, Shireby 2022) confirms direction and magnitude. Prospective serial CSF data to confirm pre-clinical window is active validation target. ROSMAP and BDR cohorts contain the necessary samples.

G-2026-P015

April 2026

PENDING

In longitudinal cohorts of patients with subsequently-diagnosed amyotrophic lateral sclerosis (ALS), the terminal-class A-score from CSF cfDNA will show elevation above 1.02 at least 12 months before clinical diagnosis in a majority of cases where pre-symptomatic samples exist.

Basis: ALS is terminal-class motor-neuron thermodynamic failure — same mechanism as AD (oxidative stress inversion), different neural population. The framework predicts the same pre-clinical signature. Falsifiable in ALS biobank cohorts with stored CSF. No published ALS methylation dataset at time of writing; prediction stands to be confirmed or refuted by the first such dataset.

ISSUE 003 ADDENDUM — TERMINAL IN THE ATLAS

9 cell types: CM, Cortical_neurons, Epi_upper, Glia, Kera_diff, Left_atrium, NeuMa, Oligo, neuron

Identity loci: 57,247 (methyl) · H_min by substrate: methyl 0.772837 · nucl 0.992027 · fuzz 0.736973 · wps 0.958909 · frag 0.624938 · methyl ceiling 1.2939

Healthy age-matched band (age_reference_matrix.json). A = 1.0 is the commitment line; healthy reads on this band.

Age (midpoint)	n	A mean	p10	p90	β mean	Source
4	25	0.9077	0.8557	0.9596	0.810	Lister pediatric
14	28	0.9210	0.8702	0.9719	0.805	Lister adolescent
24	32	0.9393	0.8868	0.9918	0.798	Lister+Roadmap

Age (midpoint)	n	A mean	p10	p90	β mean	Source
34	28	0.9520	0.8974	1.0065	0.793	Lister 2013
44	35	0.9619	0.9083	1.0155	0.789	Lister+De Jager
54	48	0.9692	0.9133	1.0252	0.786	De Jager ROSMAP
64	55	0.9789	0.9209	1.0369	0.782	De Jager+Shireby
74	62	0.9930	0.9336	1.0523	0.776	Shireby 2022
84	35	1.0067	0.9432	1.0702	0.770	Shireby aged
95	12	1.0244	0.9577	1.0911	0.762	Shireby+Lunnon

#8 · PLURIPOTENT STEM

$$\text{cfDNA} = 0.5\% \cdot H_{\min}(\text{methyl}) = 0.9822 \cdot n_{\text{bio}} = 16.5 \cdot \text{drift} \\ 2.5\%/\text{gen}$$

Reference: Human embryonic stem cells hESC H1 (Roadmap Epigenomics E003) · $H_{\min}$ confirmed from hESC H1 reference. Four TGCT histology subtypes show distinct, biologically-grounded epigenomic signatures per Shen 2018 and Killian 2016.

CELL IDENTITY & CLINICAL CONTEXT

Includes	Embryonic stem cells (hESC), induced pluripotent stem cells (iPSC), primordial germ cells (PGC), early-embryonic inner cell mass cells
Cancers (this class)	Testicular germ cell tumors (TGCT) — the most common cancer in men aged 15-35. Four pure histologic subtypes with biologically distinct methylation signatures (Shen 2018 TCGA TGCT n=137, Killian 2016 Genome Research n=130): (1) Seminoma — globally hypomethylated, PGC-like, ~60% of TGCT; (2) Embryonal carcinoma (EC) — hypermethylated at CpH sites, pluripotent-like methylation; (3) Yolk sac tumor (YST) — somatic-like methylation, differentiated extra-embryonic lineage; (4) Teratoma — somatic-like methylation, differentiated derivatives. Also: ovarian germ cell tumors (dysgerminoma — female analog of seminoma, shares molecular signature), primordial germ cell tumors in children, intratubular germ cell neoplasia in situ (GCNIS, the precursor lesion)
Non-cancer applications	iPSC reprogramming fidelity and quality control, organoid cell-state audit, synthetic embryology and gastruloid research, post-orchietomy surveillance for second primary TGCT (2-4% bilateral risk), cryptorchidism monitoring (4-10× elevated TGCT risk), fertility preservation context before platinum-based chemotherapy
Primary failure mode	Differentiation Dose Inversion + TGCT subtype-dependent methylation divergence
Warburg status	SPECIAL CASE — seminoma A_{active} decreases (inversion signal)

COMMENTARY

Testicular germ cell tumors (TGCT) are the most common cancer in men aged 15 to 35. Peak incidence falls directly in the prime childbearing years, which means the clinical audience for this card includes the urologist examining a 22-year-old graduate student who feels a painless testicular lump, the primary care physician who must decide whether a young father with a palpable scrotal mass needs immediate referral, and the oncologist who must discuss fertility preservation with a 30-year-old before starting platinum-based chemotherapy. The encouraging clinical reality is that TGCT is one of the most curable solid malignancies: the overall five-year survival exceeds 95 percent when the disease is caught at stage I or II, and cisplatin-based combination chemotherapy achieves cure in the majority of even metastatic cases. The discouraging reality is that delayed diagnosis — often driven by patient reluctance to seek evaluation for a scrotal abnormality, or by clinicians reluctant to order scrotal ultrasound in a young male presenting with back pain from retroperitoneal metastases — converts a 95 percent cure into salvage therapy with substantially lower long-term survival. Testicular self-examination remains the single highest-impact early detection intervention. This card exists in part to give clinicians the thermodynamic vocabulary to understand why a young man's tumor looks the way it does at the epigenomic level, and to support serial surveillance in the at-risk population.

Pluripotent stem cells occupy the unique architectural position of maximum developmental optionality. Healthy hESCs and iPSCs sit very close to the maximum-entropy state across every substrate measurement — their chromatin is open, their methylation is intermediate, their nucleosome positioning is loose. This is reflected in the class's $H_{\min}$ values: methylation $H_{\min} = 0.982166$ (ceiling $A = 1.0182$, only +0.018 past the healthy reference), fuzz $H_{\min} = 0.962920$ (ceiling $A = 1.0385$), frag $H_{\min} = 0.973583$ (ceiling $A = 1.0271$). Three of five substrates saturate below BREACH for this class. Only nucleosome occupancy (ceiling $A = 1.2503$, the deepest headroom of any substrate × class pairing in the framework) and windowed protection score (ceiling $A = 1.1050$) remain active past BREACH. The thermodynamic consequence is that pluripotent cells have almost no room to increase entropy in the conventional cancer direction — they are already near maximum entropy in their healthy state.

This is why TGCT disease signatures look fundamentally different from every other cancer in the framework. For most cancers, the signal is a uniform rise in A-score as cells depart from their architectural floor. For TGCT, the four pure histologic subtypes show biologically distinct methylation signatures that Shen et al. 2018 documented in TCGA TGCT (n=137) and Killian et al. 2016 Genome Research characterized in pure-histology samples (n=130) with direct comparison to primordial germ cell (PGC) and somatic methylation references. Seminoma — the most common TGCT histology, approximately 60 percent of cases — is globally hypomethylated. Its CpG methylation density moves toward the PGC state, with β values approaching zero. In the framework's language, A_{methyl} for seminoma FALLS BELOW the healthy reference, not above it. This is the inversion signal unique to this class: the only architecture class where the detection signal is a DECLINING A-score. Embryonal carcinoma (EC), by contrast, shows a striking convergence of both CpG and CpH methylation with pluripotent states, with elevated CpH methylation characteristic of the inner cell mass phase. Yolk sac tumor (YST) and teratoma show somatic-like methylation patterns that reflect their differentiation toward extra-embryonic and somatic lineages respectively. Four histologies, four distinct thermodynamic signatures, all arising from a common intratubular germ cell neoplasia in situ (GCNIS) precursor that is present in over 90 percent of TGCT cases and is itself clinically detectable.

The seminoma hypomethylation signal is the framework's most important inversion and the strongest single structural prediction. When primary-source data reports β values near zero for global CpG methylation in seminomas, the framework's A_{methyl} score drops well below the healthy hESC reference — at $\beta = 0.17$, A_{methyl} falls to approximately 0.670 compared to the healthy reference of 0.970. This has a specific detection consequence that clinicians and researchers must understand: A_{combined} alone is INSUFFICIENT for seminoma discrimination. Averaged across all five substrates, seminoma's A_{combined} sits near 0.97 — statistically indistinguishable from healthy. The discriminating signal is not a single elevated A-score but a CHARACTERISTIC DIVERGENCE PATTERN: A_{methyl} drops to 0.65 to 0.70 while A_{nucl} , A_{wps} , A_{fuzz} ,

and A_frag each modestly elevate into the 1.01 to 1.10 range. This multi-substrate divergence signature — one substrate dropping below floor while four substrates elevate above floor, with the specific relative magnitudes set by seminoma biology — is what the framework uses to detect seminoma. A naive combined-score detector would miss it. The framework's five-substrate decomposition is what makes this cancer detectable at all.


The clinical protocol for TGCT surveillance follows directly: instead of watching for A_combined to cross the BREACH threshold (1.10) as the framework does for every other cancer in the panel, the stem_pluri surveillance protocol watches for A_methyl to DECLINE toward 0.70 while A_nucl, A_wps, A_fuzz, A_frag simultaneously ELEVATE toward 1.05 or higher. This opposite-direction divergence is the detection signal. For patients with unilateral TGCT history (who carry 2 to 4 percent risk of contralateral second primary per Fossa 2005) and for patients with cryptorchidism (who carry 4 to 10 times elevated TGCT risk), serial monitoring with the correct multi-substrate divergence expectation is what makes thermodynamic surveillance clinically actionable. This is prediction G-2026-P005, originally proposed in Issue 001 and substantially refined here to account for the honest multi-substrate physics.

The clinical handling of TGCT carries specific demographic weight that belongs on this card. A thirty-year-old father of two who presents with a painless testicular nodule is the modal patient. The standard workup begins with scrotal ultrasound, proceeds through serum tumor markers (alpha-fetoprotein, beta-hCG, LDH), and typically concludes with radical inguinal orchiectomy as both diagnostic and therapeutic first intervention. Before orchiectomy — and certainly before any subsequent bleomycin-etoposide-cisplatin (BEP) chemotherapy that may be indicated for stage II or III disease — the conversation about fertility preservation belongs on the critical path. Sperm cryopreservation is well-established, inexpensive, and fully effective when offered in time. Platinum-based chemotherapy affects spermatogenesis; a substantial fraction of young TGCT survivors experience temporary or permanent impairment of fertility. A thermodynamic biomarker that supports early detection reduces the fraction of patients who reach the fertility-threatening chemotherapy stage. The framework's value proposition for the young father sitting in the oncologist's office is specifically this: earlier detection through thermodynamic surveillance in at-risk populations (cryptorchidism, prior contralateral TGCT, family history) means more patients caught at stage I where orchiectomy alone is often curative and intensive chemotherapy can be avoided.

The second major application for this class is iPSC reprogramming quality control, and this application has no cancer dimension at all — it is a research infrastructure problem. The Differentiation Dose Inversion, named for the well-documented phenomenon that excess Yamanaka factor dose produces aberrant rather than pluripotent states, maps onto the framework's prediction that successful reprogramming produces A-scores at or very near 1.00 across all five substrates. Aberrant reprogramming shows A below 0.95 (over-differentiation, cells partially committed to a lineage) or A above 1.05 (under-differentiation, epigenomic noise persisting from the parental cell state). Organoid researchers, synthetic embryology groups, and regenerative medicine translational programs all benefit from a thermodynamic quality-control metric that identifies aberrant colonies before they are committed to downstream experiments or clinical applications. This is a research-grade prediction testable in existing iPSC colony methylation and ATAC-seq datasets; prediction G-2026-P016 formalizes it.

FIVE-SUBSTRATE FIDELITY GAUGE

The gauge below is the central image of Issue 002. A single horizontal bar represents the A-score axis — from A = 0.90 (below the floor) through A = 1.00 (exactly at the Pluripotent class floor, H_min = 0.9822) up through the clinical threshold zones: MARGINAL at A = 1.01, DETECTABLE at 1.05, URGENT at 1.07, and FLOOR BREACH at 1.10. Five dots are plotted on this single axis for a healthy reference sample (hESC H1 (Roadmap E003), markers above the bar), and five more dots for a disease reference (Seminoma (Shen 2018), markers below). Each dot is one substrate: methylation, nucleosome occupancy, nucleosome fuzziness, WPS, or fragment size (DELFI). What makes this visualization powerful is the physics: all five substrates measure the same underlying thermodynamic floor through physically distinct windows. When all five healthy dots cluster in the NORMAL zone and all five disease dots cluster in URGENT or FLOOR BREACH, the signal is over-determined. Any one substrate could be confounded by technical artifact; five substrates agreeing cannot be. This is the "less blurry" advantage: √5 noise reduction for free.



Look at the spread in each direction. If the healthy five-dot cluster sits tightly together near A = 0.97–1.00, that indicates the reference population is at its architecture floor — precisely where the framework predicts healthy pluripotent cells should sit. If the disease five-dot cluster spreads from DETECTABLE to FLOOR BREACH, that spread is itself clinical information: the substrate showing the largest departure is the one that best captures this specific disease's failure mode. For this class, the primary failure mode is Differentiation Dose Inversion + TGCT subtype-dependent methylation divergence, and the substrate ranking shown further down the card explains which of the five substrates best reveals it.

SUBSTRATE-BY-SUBSTRATE BREAKDOWN


Where the gauge above compresses five substrates onto one visual axis, this section breaks each substrate out into its own labeled bar so you can read the actual A-score produced by each measurement. The physics is the same across all five: take the raw laboratory value (a β methylation fraction, a nucleosome occupancy probability, a WPS score, etc.), compute its Shannon entropy $H(\text{value})$, then divide by the class-specific and substrate-specific H_{min} to get a dimensionless A-score comparable across substrates. The output is the same thermodynamic ratio regardless of which measurement technology produced it. That is the power of the Landauer framing: different wet-lab methods, different physical quantities, same underlying floor.

The formula at the foot of this section — $A_{\text{combined}} = \Sigma(AUC_i \times A_i) / \Sigma(AUC_i)$ — combines the five individual A-scores into a single weighted score, using each substrate's published detection AUC as its weight. Higher-AUC substrates contribute more to the combined result. For pluripotent cells, the healthy combined A-score below should sit at approximately 0.97 (NORMAL tier); the disease combined A-score should sit above 1.05 (DETECTABLE or higher). The difference between them — the ΔA combined — is the signal available to a blood draw that analyzes all five substrates simultaneously. A single-substrate assay sees a smaller signal at higher noise; the five-substrate combination sees the same signal at approximately $\sqrt{5}$ lower noise. Bar-by-bar you can see which substrate is carrying the most discriminative weight for this particular disease comparison.

HEALTHY REFERENCE — hESC H1 (Roadmap E003)



DISEASE REFERENCE — Seminoma (Shen 2018)




Combined A — healthy (all 5)	0.97011 (NORMAL)
Combined A — disease (all 5)	0.96716 (NORMAL)
ΔA combined (5 substrates)	-0.00294 vs methylation-only ΔA =-0.3006
Active A — healthy (non-saturated)	0.97011 (NORMAL) [5/5 active]
Active A — disease (non-saturated)	0.96716 (NORMAL) [5/5 active]

ΔA active (real progression signal)	-0.00294 (active substrates track true progression)
Saturation status (disease)	No substrates saturated

Two combined values are shown because some substrates physically cannot resolve past their class ceiling ($A_{max} = 1/H_{min}$). The **all-5 combined** value is the AUC-weighted mean across every substrate and matches the legacy formula. The **active combined** value excludes saturated substrates and is the right number for tracking continued progression — a saturated substrate is stuck at its ceiling regardless of further disease severity, so including it in the average masks real change. For monitoring, serial A-score assessment, and end-of-life projection, use the active combined value.

DISEASE SIGNATURE COMPARISON — healthy hESC through the four TGCT histologies

Six conditions on a single chart spanning the full TGCT histologic spectrum plus healthy and precursor states. β values reproduce per-substrate A-scores calibrated to Shen 2018 Cell Reports (TCGA TGCT, $n=137$) and Killian 2016 Genome Research ($n=130$ pure-histology TGCTs with PGC and somatic reference comparisons). Three substrates saturate on this class: methyl (ceiling $A = 1.0182$), fuzz (ceiling $A = 1.0385$), frag (ceiling $A = 1.0271$). Nucl remains active with +0.150 headroom past BREACH; wps remains active with modest headroom. The four TGCT histologies show biologically distinct methylation signatures: seminoma is globally hypomethylated (methyl β drops toward PGC state, A_{methyl} FALLS below healthy — the inversion signal for this class), embryonal carcinoma is hypermethylated at CpH sites with pluripotent-like CpG patterns, yolk sac tumor and teratoma show somatic-like methylation reflecting their differentiation toward extra-embryonic and somatic lineages respectively. GCNIS (germ cell neoplasia in situ) is the clinically detectable precursor lesion present in over 90% of TGCT cases. The chart reads left to right as the disease progression most relevant to surveillance: healthy $\rightarrow$ GCNIS precursor $\rightarrow$ seminoma / EC / YST / teratoma as distinct presenting histologies.




What happens past the ceiling is condition-specific. The A-score magnitude alone does not track severity once the ceiling is crossed — different diseases take different post-breach paths that the full five-substrate divergence pattern reveals. The subsection below walks through this class's post-breach trajectory: which substrates keep carrying signal, where the Warburg metabolic lock-in happens, where the glucose inversion point sits, and where the framework currently has high confidence versus where the open research questions remain.

POST-BREACH TRAJECTORY — WHAT HAPPENS AFTER A CROSSES 1.10 (PLURIPOTENT CLASS)

The pluripotent stem class is the framework’s showcase for inversion physics. Seminoma — 60% of testicular germ cell tumors — is globally hypomethylated rather than hypermethylated, producing A_methyl that FALLS below healthy reference (A_methyl ≈ 0.67 at β = 0.17) while A_nucl, A_wps, A_fuzz, A_frag elevate normally. This is the cleanest documented below-floor disease state in the framework. The clinical signal for seminoma is not combined-A elevation — it is a multi-substrate DIVERGENCE pattern: one substrate inverting while four elevate. A naive cancer detector looking only at combined A misses seminoma entirely. VAL-045 (April 2026) revealed an important extension: because H_min_methyl = 0.9822 sits very close to the Shannon ceiling of 1.000, the architectural window above floor is extremely narrow. Any methylation departure from the narrow pluripotent reference band produces A_methyl below floor regardless of direction. All four TGCT histologies invert at the methylation level: seminoma is the extreme case (A_methyl ≈ 0.75, β ≈ 0.21), embryonal carcinoma lands at A ≈ 0.83 (β ≈ 0.745), yolk sac tumor at A ≈ 0.87 (β ≈ 0.72), and choriocarcinoma at A ≈ 0.85 (β ≈ 0.735). Specificity in this class comes from divergence magnitude, not divergence direction — seminoma’s divergence is 2.1× that of other TGCT histologies. Inversion is class-universal for pluripotent; seminoma is the extreme, not the only case.

The zone language below — Warburg boundary, glucose inversion, point of no return — applies only to post-breach cells. These terms are not used on the pre-breach bar, and the qualitative A-value anchors (~1.15, ~1.25, ~1.40+) are boundaries between therapeutic windows, not additional diagnostic tiers. Specific A-values are filed as prospective predictions to be validated against survival-stratified cohorts.



WHICH SUBSTRATES CARRY POST-BREACH SIGNAL — PLURIPOTENT CLASS

Pluripotent-class physics: three substrates saturate early (methyl, fuzz, frag) at class-specific ceilings close to healthy reference. Nucl and WPS carry post-breach progression signal, with nucl having the deepest headroom in the class. The inversion signal lives on methyl — when β drives toward PGC-like near-zero, A_methyl falls below healthy rather than rising.

Substrate	Ceiling (A)	Post-breach status
Nucleosome occupancy	1.25	Primary severity metric — deepest post-breach headroom
Windowed protection	1.11	Near-BREACH severity metric — modest headroom
Methylation	1.018	Saturated ceiling on the UPWARD axis; but INVERSION-capable on the DOWNWARD axis
Fuzziness	1.039	Saturated at ceiling — no further upward signal post-breach
Fragment size	1.027	Saturated at ceiling — no further signal post-breach

INVERSION TERRITORY — SEMINOMA HYPOMETHYLATION INVERSION (showcase documented case)

The Seminoma Hypomethylation Inversion is the framework’s cleanest documented below-floor disease state. Seminoma tumor β_methyl drops toward 0.17-0.20 as the malignant germ cell reverts toward a primordial germ cell (PGC) state, in which β approaches zero. Because pluripotent H_min_methyl = 0.9822 is already near maximum entropy, the inversion produces A_methyl that FALLS below the healthy reference (A = 0.67 at β = 0.17) rather than rising above it. A naive cancer-detection instrument looking for combined-A elevation misses seminoma entirely. The framework’s discrimination signal is the multi-substrate divergence pattern: A_methyl drops to 0.65-0.70 while A_nucl, A_wps, A_fuzz, A_frag simultaneously elevate into the 1.01-1.10 range. Divergence = max|A_i - median(all A)| — a divergence ≥ 0.10 is the framework’s seminoma signature. Confirmed in Shen 2018 TCGA TGCT (n=137) and Killian 2016 Genome Research (n=130 pure-histology samples with direct PGC comparison). This is why every card now shows the INVERSION zone on the pre-breach bar — because below-floor disease is real, documented, and class-specific. VAL-045 (April 2026, multi-class drift cascade) extended this picture: because the pluripotent methylation window above floor is so narrow (H_min = 0.9822 versus H_max = 1.000), any methylation departure from the reference band inverts A_methyl regardless of direction. All four TGCT histologies land in inversion territory at the methylation level. Seminoma is the extreme case at A_methyl ≈ 0.75 (divergence 0.079); embryonal carcinoma at ≈ 0.83, yolk sac at ≈ 0.87, and choriocarcinoma at ≈ 0.85 (divergences 0.037 or less). Seminoma’s divergence is 2.1× that of other histologies — the specificity is in divergence magnitude, not in inversion direction. Inversion is class-universal for pluripotent-class cancers; seminoma is the extreme case, not the only case. This refinement does not change the framework’s discrimination protocol (divergence ≥ 0.10 still identifies seminoma), but it clarifies that the other TGCT histologies also depart architecturally into inversion territory — just at smaller magnitudes.

CONDITION-BY-CONDITION TRAJECTORY

Healthy hESC (reference) — pluripotent baseline

Known: Human embryonic stem cells (Roadmap E003) sit at the class reference with $A \approx 0.97$ across all five substrates. Pluripotent cells are designed for high entropy — their $H_{\min} = 0.9822$ is the highest class floor in the framework, because multi-potency requires maximum epigenomic reversibility. The Yamanaka reprogramming factors (OCT4, SOX2, KLF4, c-MYC) drive somatic cells back toward this high-entropy state.

Seminoma (INVERSION signature) — $A_{\text{methyl}} \approx 0.67$ (inversion), $A_{\text{nucl/wps/fuzz/frag}} \approx 1.01-1.10$

Known: Seminoma is the framework's showcase inversion case. The cancer manifests as divergence, not elevation: methyl crashes toward PGC state while the other four substrates rise. Early-detection scenario (Card Section 5.1) shows the framework flagging divergence at Month 12 — 6 months before standard ultrasound surveillance would detect a mass. This is the framework's highest-value clinical case: a cancer that current molecular panels miss entirely.

Unknown: whether the framework can detect seminoma in cryptorchidism surveillance cohorts (4-10× elevated TGCT risk) before clinical presentation; whether post-orchietomy serial A-score can identify second primary TGCT (2-4% bilateral risk) earlier than CT imaging.

Test that closes the gap: G-2026-P005 (filed): Prospective surveillance cohort of cryptorchidism patients with baseline plus 6-monthly cfDNA. Framework prediction: divergence ≥ 0.10 identifies pre-clinical seminoma with sensitivity ≥ 0.80 at specificity 0.95, providing 6+ month lead time vs ultrasound + tumor marker surveillance.

Embryonal carcinoma (EC) — $A \approx 1.05-1.10$, UPWARD path (hypermethylation)

Known: Embryonal carcinoma shows the OPPOSITE methylation direction from seminoma: CpH hypermethylation characteristic of the inner cell mass phase, pluripotent-like CpG patterns. A_{methyl} elevates rather than inverts. This is the framework's cleanest demonstration that direction on the A-score axis is diagnostic — the same class can produce either inversion or elevation depending on the specific malignant reprogramming pathway.

Unknown: whether EC-dominant mixed germ cell tumors show distinguishable A-score trajectories from pure EC; whether response to BEP chemotherapy produces trajectory bend toward healthy (A_{methyl} declining from elevation rather than recovering from inversion).

Test that closes the gap: G-2026-P017 (filed): Prospective cohort of 100 stage II-III TGCT patients undergoing BEP with serial cfDNA pre-C1, pre-C2, pre-C3, pre-C4. Prediction: responders show A_{methyl} moving toward healthy hESC reference (approaching 0.970 from either direction — recovering from seminoma inversion or declining from EC hypermethylation); non-responders remain pinned at disease baseline.

iPSC reprogramming (research-grade inversion) — $A \approx 0.90-1.05$, DIFFERENTIATION DOSE INVERSION

Known: In iPSC reprogramming protocols, excess Yamanaka factor dose produces aberrant rather than successfully reprogrammed colonies. Successful reprogramming produces $A \approx 1.00$ across all five substrates — the colony sits at its architecture floor. Aberrant colonies show A below 0.95 (over-differentiation) or above 1.05 (under-differentiation). The pharmacologic dose-response is non-monotone: more is not better past the optimal window. This is a research-grade application of the framework, not clinical, but it demonstrates inversion physics in a controlled setting.

Unknown: whether the framework can quality-control iPSC reprogramming at single-colony resolution with ATAC-seq inputs; whether Differentiation Dose Inversion is reproducible across reprogramming protocols (OSKM, episomal, Sendai, mRNA).

Test that closes the gap: G-2026-P016 (filed): Prospective iPSC reprogramming cohort with ATAC-seq on individual colonies at day 20, 30, 40 post-transduction. Prediction: colonies with A within 0.97-1.03 across all five substrates produce teratomas successfully ($\geq 95\%$); colonies outside this window fail differentiation QC ($\geq 80\%$).

WHAT THE FRAMEWORK CAN AND CANNOT SAY TODAY

The framework has high confidence for pluripotent-class post-breach: (1) the Seminoma Hypomethylation Inversion is the framework's cleanest documented below-floor case, with primary-source validation in Shen 2018 and Killian 2016; (2) nucl and WPS carry post-breach severity signal while methyl, fuzz, and frag saturate near healthy reference; (3) multi-substrate divergence (not combined elevation) is the discrimination signal for this class, demonstrating that direction on the A-axis is diagnostic.

The framework has not yet tested seminoma early detection in cryptorchidism surveillance prospectively, nor BEP platinum response trajectory, nor iPSC quality control. G-2026-P005, P016, P017 define the specific validation plans.


This is the scientific posture the framework commits to: definitive where the evidence is definitive, open where it is open, and specific about the experiments that would close each gap. Predictions filed for this class: G-2026-P005, G-2026-P016, G-2026-P017. A reviewer who wants to falsify any claim here can do so against the named cohorts. A collaborator who wants to extend the framework has specific hypotheses to test.

THREE-COMPONENT DECOMPOSITION (C1/C2/C3) — PER SUBSTRATE

Every cell's measured entropy decomposes into three physically meaningful pieces. C1 is the universal Landauer floor: $H_{\min_GLOBAL} = 0.756500$, measured in frontal cortex neurons by Lister 2013. Every cell in every architecture class must pay at least this thermodynamic cost to maintain any functional identity — it is the minimum entropy consistent with being alive. C2 is the class-specific overhead: the additional entropy the

pluripotent class must carry above the universal floor to encode its particular architecture. For this class, C2 represents approximately 3.0% of the healthy reference entropy. C3 is the accessible gap: $\max(0, H_{\text{actual}} - H_{\text{min}})$. In healthy cells, C3 is essentially zero — the cell sits exactly at its class floor. When C3 grows, the cell is departing from its architecture class. Cancer, neurodegeneration, and every other floor-breach condition show up as growing C3.

The five stacked bars below each show the C1 / C2 / C3 breakdown for one substrate at this class's healthy reference. Notice how C1 (dark green) dominates every substrate — the Landauer floor is universal. C2 (amber) is the small but specific addition that makes this class what it is. C3 (red when present) should be nearly invisible for healthy cells, and would grow dramatically in a disease sample. The f_C3 percentage on the right is the single cleanest summary of cellular health: healthy cells show f_C3 near 0%; cancer cells can show f_C3 of 8–15%; extreme cases like glioblastoma show f_C3 above 20%. This is not a statistical threshold calibrated from disease data. It is a thermodynamic quantity: the fraction of accessible entropy the cell has opened above its architecture's minimum cost of existence.



BEST TESTING METHOD RANKING — CLINICAL UTILITY FOR THIS CLASS

Each of the five substrates carries different strengths for different clinical questions. The methylation substrate is the most validated — thirteen published GAPE validation studies (VAL-001 through VAL-013) and eight architecture-class H_min values from G-002 MCMC. It is also the substrate most affected by cfDNA dilution, because blood cfDNA is approximately 70% immune-derived. Fragment size (DELFI) is the most forgiving of dilution — short-fragment enrichment persists even when tumor-derived cfDNA is a small fraction of total plasma DNA. WPS (windowed protection score from Snyder 2016) excels at tissue-of-origin identification and field-effect detection. Nucleosome occupancy and fuzziness add chromatin-level information that methylation alone cannot reveal.

The ranking below is specifically for the pluripotent class. The order reflects three factors: (1) the substrate's published single-substrate AUC for diseases in this class; (2) the biology of this class's primary failure mode (Differentiation Dose Inversion + TGCT subtype-dependent methylation divergence); and (3) the practical constraint of cfDNA contribution in the relevant specimen (plasma for most classes, CSF for terminal, tissue biopsy where cfDNA is too dilute). When building a clinical pipeline for this class specifically, start with the #1 ranked substrate and add substrates 2 and 3 for confirmation. Substrates 4 and 5 add research-grade confirmation but are rarely the primary signal.

#	Substrate	Best for	Rationale
1	Methylation	Histology discrimination — bidirectional signal	The only substrate whose signal direction differs by TGCT histology. Seminoma drives methyl DOWN toward PGC-like hypomethylation (A falls below 0.970 — the inversion signal). Embryonal carcinoma drives methyl UP (CpH hypermethylation). Yolk sac and teratoma show somatic-like methylation. Primary source: Shen 2018 Cell Reports.
2	Nucleosome Occupancy	Primary severity metric past DETECTABLE	Only substrate with substantial headroom past BREACH (ceiling A = 1.2503). For any TGCT histology, nucl carries severity signal reliably. Also primary metric for iPSC reprogramming fidelity assessment (ATAC-seq on colonies).
3	Windowed Protection Score	Near-BREACH severity metric	Active past BREACH with modest headroom (ceiling A = 1.1050). WPS at pluripotency promoters (OCT4, SOX2, NANOG) detects the persistence of embryonic-gene expression characteristic of all TGCT histologies.
4	Nucleosome Fuzziness	Detection boundary — ceiling at A = 1.0385	Saturates near BREACH threshold. Useful for binary detection of floor departure; cannot resolve histology differences above ceiling.
5	Fragment Size (DELFI)	Detection boundary — ceiling at A = 1.0271	Saturates below BREACH. Fragmentomic signal for TGCT is dominated by rapid tumor turnover rate rather than methylation specifics; useful for detection but not histology typing.

BODY TEMPERATURE SCALING — α = 2.0 LANDAUER CORRECTION

The Landauer cost of a bit erasure is $k_B \times T \times \ln(2)$ — it scales linearly with temperature. Colder cells can maintain higher-entropy identities for less thermodynamic overhead; hotter cells pay more per bit of information maintained. The GAPE framework captures this with a simple scaling law: $H_{\text{min}}(T) = H_{\text{min}}(37^\circ\text{C}) \times (T_{\text{body}} / 310.15 \text{ K})^\alpha$, with $\alpha = 2.0$ derived empirically by minimizing cross-class A-score variance across all jawed

vertebrates (Mahaffey 2026 Nature Aging submission NATAGING-A13702). The $\alpha = 2.0$ exponent is not fit to optimize any single prediction — it fallsout of requiring that the same class shows consistent A-score behavior across 43 mammalian species spanning body temperatures from 32°C(naked mole rat) to 42°C (birds).


The table below shows what the pluripotent class floor H_{min} looks like at seven representative body temperatures. At 37°C (human reference, highlighted row), $H_{min} = 0.9822$. At 42°C (birds), the floor is higher because the Landauer tax is larger. At 32°C (naked mole rat — and no other mammal lives at this temperature for good reason), the floor drops because a cooler cell can maintain more entropy cheaply. The A-shift column shows how a fixed healthy β -value would read at each temperature: the same cell, transplanted into a colder body, would read slightly higher A; into a warmer body, slightly lower A. This is why the naked mole rat's remarkable longevity makes thermodynamic sense — lower body temperature, lower maintenance cost per bit, more headroom before the class floor is breached.

T_body	Species example	H_min(T)	$\beta_{healthy}$ equivalent	A-shift vs 37°C
42°C	Birds (chicken, finch)	1.0141	0.627	$A \times 0.969$
39°C	Rodents (mouse, rat)	0.9949	0.627	$A \times 0.987$
37°C	Humans (REFERENCE)	0.9822	0.627	$A \times 1.000$ (anchor)
35°C	Hibernating bats	0.9695	0.627	$A \times 1.013$
32°C	Naked mole rat	0.9508	0.627	$A \times 1.033$
25°C	Reptiles (anole lizard)	0.9076	0.627	$A \times 1.082$
15°C	Fish (salmon, cold-water)	0.8478	0.627	$A \times 1.159$

HEALTHY AGING TRAJECTORY — AGE-STRATIFIED REFERENCE A-SCORES

Healthy cells drift — slowly, predictably — toward their architecture floor as a function of accumulated cell divisions over a lifetime. For the pluripotent class, the drift rate is approximately 2.5% per generation — meaning each cell cycle carries a 2.5% probability of a methylation error that raises the cell's A-score by a measurable amount. Over decades, these errors accumulate. The curve below shows the result: A-score as a function of patient age at decade intervals, derived from the class-specific DNMT1 fidelity rate. At age 20, the reference healthy A-score sits well below 1.00; by age 80, it has drifted upward into the MARGINAL to DETECTABLE range. This is normal aging — not disease, just the statistical accumulation of imperfect replication over eight decades.

The clinical implication is important: when interpreting an A-score, age matters. A patient at age 30 with $A = 1.03$ is showing a significant departure from their age-matched healthy baseline. The same $A = 1.03$ at age 75 might be entirely consistent with age-expected drift. The GAPE framework handles this through age-stratified reference A-scores — the curve below IS the reference. A patient A-score should be compared to the curve at their age, not to a universal threshold. Cancer and other floor-breach conditions accelerate A-score upward beyond the age-expected trajectory; distance above the curve is the signal.



VERTEBRATE LIFESPAN CONTEXT

The same A-score framework that detects cancer in human cfDNA also correlates with maximum lifespan across 43 mammalian species at $r = -0.9018$ ($p = 1.6 \times 10^{-16}$, Mahaffey 2026 Nature Aging submission NATAGING-A13702). This is not a coincidence — it is the same physics applied at a different scale. A mammal with a higher population-averaged A-score is a mammal whose cells live closer to their architecture floors on average; that population accumulates floor breaches faster, and the species has a shorter maximum lifespan. The $A = 1.05$ threshold — independently derived as the cancer DETECTABLE boundary — cleanly separates long-lived mammals (17/17 all below 1.05) from short-lived mammals (11/11 all above 1.05) with complete accuracy. Same threshold. Same physics. Different timescale.

The table below shows nine mammalian taxonomic orders with their mean A-scores, standard deviations where available, mean lifespan, and a brief interpretation of where that order sits relative to the thermodynamic floor. Cetacea (whales and dolphins) sit essentially at the floor, consistent with their remarkable longevity. Primates including humans sit slightly above the floor. Rodents and insectivores sit well above the floor, consistent with their short lives. This class's reference A-score (0.970) places it in the context shown below — use this to understand how this particular cell

architecture relates to the broader mammalian lifespan distribution.

Order	N	Mean A	Lifespan (yr)	Interpretation
Cetacea	5	0.997 ±0.016	99	At thermodynamic floor
Proboscidea	1	0.987 —	70	At floor
Primates	6	1.007 ±0.015	49	Slightly above floor
Artiodactyla	4	1.015 ±0.011	30	Slightly above floor
Chiroptera	4	1.041 ±0.007	35	Above floor (longevity outliers)
Carnivora	9	1.053 ±0.023	28	Above floor
Lagomorpha	2	1.114 ±0.002	8	Well above floor
Rodentia	6	1.125 ±0.018	9	Well above floor
Insectivora	1	1.157 —	2	Furthest from floor

INTERVENTION LEVERS — RANKED BY EXPECTED IMPACT

If a patient's A-score is elevated — MARGINAL, DETECTABLE, or URGENT — the clinical question is what can be done. The GAPE framework does not prescribe treatment; that remains the physician's judgment. But the framework does identify which intervention categories have the mechanistic strength to move this particular class's A-score back toward the floor. The five categories below are the standard anti-aging and cancer-adjacent interventions: senolytics (clearing senescent cells), metabolic restoration (NAD⁺, caloric restriction, exercise), epigenetic restoration (DNMT/TET modulators, HDAC inhibitors), reprogramming (cyclic Yamanaka factors), and checkpoint stringency (cell cycle control, immune checkpoint blockade).

The ranking below is specifically for the pluripotent class and its primary failure mode (Differentiation Dose Inversion + TGCT subtype-dependent methylation divergence). Impact Level 1 (Dominant) means the intervention directly addresses the class's failure mechanism and is the first lever to try. Level 2 (Strong) means substantial mechanistic support for A-score restoration in this class. Level 3 (Moderate) means helpful but not addressing the binding constraint. Level 4 (Limited) means the intervention works for other classes but does not target this class's specific biology. Level 5 (Not applicable) means the intervention category is biologically incompatible with this class — e.g., senolytics cannot act on post-mitotic terminal cells that do not become classically senescent, and reprogramming cannot be applied to terminal cells without losing the identity that defines them.

Impact	Category	Mechanism	Rationale
1 DOMINANT	Metabolic	metabolic	Dominant — metabolic flexibility means ATP optimization directly moves fidelity
1 DOMINANT	Staged Reprogramming	reprogramming	Dominant — this is the source class for iPSC reprogramming
2 STRONG	Epigenetic Restoration	epigenetic rx	Strong — DNMT1/TET restoration improves commitment fidelity
3 MODERATE	Checkpoint Stringency	checkpoint	Moderate — G1/S checkpoint active but differentiation is the primary lever
4 LIMITED	Senolytics	senolytics	Not applicable — pluripotent cells do not express SASP

CANCER PANEL — ΔA RANKED

The TCGA cancer validation is where the framework earns its keep. For every cancer type in this class, we compute two numbers from published TCGA methylation data: A_healthy (from matched normal tissue) and A_tumor (from cancer samples in the same study). The $\Delta A = A_{\text{tumor}} - A_{\text{healthy}}$ is the signal the framework predicts should appear — and it does. Thresholds for tier assignment: A = 1.05 DETECT, A = 1.07 URGENT, A = 1.10 FLOOR BREACH. These thresholds are not fit to the cancer data; they were derived from the healthy-vs-age statistics and happen to align with the cancer validation cleanly. The bars below are ordered by $|\Delta A|$ descending so the cancers with the largest signal appear first.

Note that ΔA ranges typical for this class matter. Terminal-class cancers (glioma, glioblastoma) show $\Delta A \approx 0.22\text{--}0.27$ — the largest in the entire 28-cancer TCGA dataset. Cycling epithelial cancers show $\Delta A \approx 0.13\text{--}0.19$ — moderate but clearly detectable. Stromal and progenitor-lineage cancers show $\Delta A \approx 0.10\text{--}0.17$. Each cancer below is cited to its primary TCGA publication and shows sample size. Direct links from A-score to TCGA evidence — no black boxes between physics and clinical data.

Rank	Gene	Cell Rep	Delta A	A	Category
#1	Testicular (TGCT) n=151 · TCGA TGCT 2018 Cell Rep		$\Delta A = -0.1777$	A=0.826	NORMAL

CORE METRICS — PUBLISHED AND DERIVED

Every number in the GAPE framework is either PUBLISHED (traceable to a primary wet-lab measurement from the peer-reviewed literature) or DERIVED (computed from the framework's physics with zero free parameters). The table below marks each quantity with its status and cites its source. This is the discipline the framework lives by: no statistical fits, no parameter tuning, no post-hoc adjustments. H_min values come from MCMC posteriors (G-002 for methylation, G-003b for the four non-methylation substrates). Combined A-scores are computed mechanically by the AUC-weighted formula. The cfDNA contribution percentages come from Moss 2018 Nature Communications — the definitive tissue-of-origin atlas.

Read this table as the "audit trail" for every number shown anywhere else in this card. When the framework says A_healthy = 0.9703, you can see exactly which β value produced it, which H_min was used, and where both came from. The n_bio field is currently marked PRELIMINARY because absolute n_bio values await the G-007 MCMC run; the class ordering has been confirmed (Spearman $\rho = 0.905$, $p = 0.002$ against commitment level) but the absolute values carry that caveat transparently.

Metric	Value	Source
cfDNA contribution (plasma)	0.5%	Moss 2018 Nat Commun — PUBLISHED
H_min methylation (class)	0.982166	G-002 MCMC posterior — DERIVED
H_min methylation σ (MCMC)	± 0.001800	G-002 17-chain MCMC R-hat<1.001 — DERIVED
n_bio (class-specific)	16.5	PRELIMINARY — ordering confirmed ($\rho=0.905$, $p=0.002$); absolute pending G-007
Healthy drift rate	2.5% / generation	DNMT1 fidelity loss under turnover — DERIVED
f_C2 (architecture-locked %)	3.0%	C2 fraction at healthy reference — DERIVED
Healthy A (methylation)	0.9703 (NORMAL)	$\beta=0.627$ — DERIVED
Disease A (methylation)	0.6696 (NORMAL)	$\beta=0.170$ — DERIVED
Healthy combined A (5 subs)	0.9701 (NORMAL)	$\Sigma(\text{AUC}_i \times A_i) / \Sigma(\text{AUC}_i)$ across 5 substrates — DERIVED
Disease combined A (5 subs)	0.9672 (NORMAL)	$\Sigma(\text{AUC}_i \times A_i) / \Sigma(\text{AUC}_i)$ across 5 substrates — DERIVED
ΔA combined	-0.0029	Combined signal — DERIVED

DATED PREDICTIONS FOR THIS CLASS

A prediction is not a claim unless it can fail. Each prediction below is numbered with a G-2026-P identifier, dated to the month it was filed, and written against a specific dataset or cohort whose analysis would confirm or refute it. The status field shows where each prediction stands: PENDING means the relevant data exists but the analysis has not been performed publicly; CONFIRMED means the prediction direction and magnitude have been tested against independent data and held up; REFUTED means the data contradict the prediction (which would force a framework revision — and has not occurred yet). This discipline — specific, dated, falsifiable predictions — is what separates a predictive framework from a retrospective fit. Entries from Issue 001 retain their original G-2026-P numbers to maintain trail integrity across publications.

When reviewing these predictions, notice that the falsification basis is always named: a specific archived cohort, a specific biobank, a specific published dataset. These are not vague "more research is needed" statements. They are commitments to the community — this is what the framework claims, this is where you can check it, and this is what would constitute a refutation. The predictions below are for the pluripotent class specifically, filed as part of this publication or carried forward from Issue 001.

G-2026-P005

Originally Issue 001, substantially refined April 2026

PENDING

In cryptorchidism patients and post-orchietomy TGCT survivors monitored prospectively with serial cfDNA sampling, the stem_pluri class will show a CHARACTERISTIC DIVERGENCE PATTERN in patients who later develop seminoma-lineage TGCT or contralateral second primary TGCT: A_methyl declines toward the 0.65 to 0.75 range while A_nucl, A_wps, A_fuzz, and A_frag simultaneously elevate to 1.01 or higher. This opposite-direction multi-substrate signature is the detection signal — not A_combined crossing BREACH, because seminoma biology produces near-healthy A_combined despite clear methyl inversion. Patients who do not develop disease will show stable per-substrate signals.

Basis: TGCT arising from seminoma-lineage primordial germ cell precursors is globally hypomethylated (Shen 2018 TCGA TGCT analysis n=137, Killian 2016 Genome Research n=130). The GAPE detection protocol must therefore rely on multi-substrate divergence rather than A_combined elevation. This is the zero-free-parameter structural prediction the framework makes for this class, and it is falsifiable in any cryptorchidism registry with serial blood sampling and TGCT outcome follow-up, or in any post-orchietomy surveillance program with archived cfDNA samples. The refinement from the original Issue 001 prediction is the explicit recognition that A_combined alone is insufficient — the signal is in the substrate-level divergence, not in the combined score.

G-2026-P016

April 2026

PENDING

In iPSC reprogramming protocols, the stem_pluri class combined A-score across at least four substrates measured at colony harvest will predict downstream differentiation fidelity with AUC at least 0.85. Colonies with A in the tight range [0.98, 1.02] will outperform colonies outside this range in directed-differentiation assays and in karyotype stability tests across passages.

Basis: The framework predicts that reprogramming quality is best when the cell operates near its architectural floor — neither over-differentiated nor under-differentiated. Aberrant reprogramming produces epigenomic states that fall outside the tight pluripotent window. Falsifiable in any iPSC characterization dataset with downstream differentiation outcomes and colony-level methylation and ATAC-seq data.

G-2026-P017

April 2026

PENDING

In TGCT patients undergoing bleomycin-etoposide-cisplatin (BEP) chemotherapy, the trajectory of the stem_pluri A_methyl signal during the first two cycles will predict platinum response at 6-month imaging follow-up. Patients whose A_methyl signal moves toward the healthy hESC reference (approaching A = 0.970 from either direction) will show RECIST response. Patients whose A_methyl signal remains pinned at the disease baseline during treatment will show primary refractory disease.

Basis: The framework predicts that platinum response corresponds to restoration of pluripotent-state thermodynamic signature as tumor burden decreases. Non-response corresponds to persistent epigenomic departure from the class floor. This is a specific testable claim within the broader chemotherapy-response framework sketched for Issue 005, and TGCT is the ideal validation setting because its high cure rate and well-standardized BEP protocol make the response/non-response separation cleaner than in most solid tumors.

ISSUE 003 ADDENDUM — STEM_PLURI IN THE ATLAS

1 cell types: stem_pluri

Identity loci: 21,710 (methyl) · **H_min by substrate:** methyl 0.982166 · nucl 0.799818 · fuzz 0.962920 · wps 0.905004 · frag 0.973583 · **methyl ceiling** 1.0182 — cannot reach BREACH on methylation alone

Healthy age-matched band (age_reference_matrix.json). A = 1.0 is the commitment line; healthy reads on this band.

Age (midpoint)	n	A mean	p10	p90	β mean	Source
4	5	0.8292	0.8067	0.8517	0.748	Pluripotent lineage
14	8	0.8308	0.8084	0.8532	0.747	Pluripotent stem
24	10	0.8324	0.8101	0.8547	0.746	hESC H9 Roadmap E008
34	8	0.8340	0.8118	0.8561	0.745	hESC/iPSC
44	6	0.8340	0.8118	0.8561	0.745	iPSC
54	5	0.8356	0.8135	0.8576	0.744	iPSC reference
64	4	0.8356	0.8115	0.8596	0.744	iPSC reference
74	3	0.8356	0.8115	0.8596	0.744	iPSC reference
84	2	0.8371	0.8112	0.8631	0.743	Limited data
95	1	0.8371	0.8112	0.8631	0.743	Limited data

WHERE THE TOOLS COME FROM

Issue 002 stated the physics of the floor. Since April the engine acquired an atlas, a deconvolver, a sky projection, a matched filter and a second statistic, and every one of them was borrowed from somewhere specific: from the semiconductor and quantum-computing applications of the same law, and from the Planck CMB pipeline. A reader who is going to spend time on this work will want to know where each idea came from and what it is doing here. This section says so, in the order a physicist would ask. Nothing here is a new claim; the new claims are in §5 (reproduced from 002) and the measurements are in §7–§8.

5A.1 The Mahaffey number

$M = 20.94$. $M = E_{\text{drive}} / (k_B T)$ for the cell: the ATP ordering budget per event over the thermal noise quantum at body temperature. Dimensionless, fixed by biochemistry, the same for every cell. $M = dG_{\text{ATP}}/(R T) = 54,000/(8.314 \times 310.15) = 20.94$ at 37 C. Author 2026-09-19: one fixed number; the 30.21 (ln2) form was wrong.

H_{min} is a different quantity: the per-class entropy floor in bits (8 classes, 5 substrates each -> 40 floors), MCMC-confirmed, architecture-specific. M is one number for the cell; H_{min} is forty numbers for the classes. They are never interchangeable.

n_{bio} was an early (QAPE-era) form of the similar ratio and is NO LONGER USED. Every occurrence of n_{bio} in Issue 002 (glossary, class cards ' $n_{\text{bio}} = 20.94$ ', canine 20.84) is superseded by the Mahaffey number.

The same construction - a measured drive energy over the local thermal quantum - is what SCAPE computes for a transistor and QAPE for a qubit; those figures are quoted in Section 5A.8 exactly as their own reports state them, with their own denominators, and are not re-derived here.

Open. Inputs: $dG_{\text{ATP}} = 54,000$ J/mol, $R = 8.314$ J/(mol K), $T = 310.15$ K (002 glossary, `_gape_constants_private.py` l.26, Cellular Margin l.28, Hubble2GAPE l.3504) - consistent across the corpus.

IAM_Hubble2GAPE l.2192 writes $M = E_{\text{drive}}/(k_B T_{\text{local}} \ln 2)$ and gives ~30 for chemical systems; the author's ruling has no ln2 (20.94). That line of the tex should be brought into agreement. The earlier 003 draft's cross-substrate table (SCAPE 117x, transmon $M = 1$ exactly) was computed with the ln2 form and is WITHDRAWN from s5A.1; SCAPE and QAPE figures are quoted only as their own reports state them.

5A.2 Forty floors — eight classes × five substrates, each confirmed by MCMC

The cellular H_{min} is not one number. Each of the eight architecture classes has its own floor on each of five physically distinct lab measurements — methylation β , nucleosome occupancy, nucleosome fuzziness, window protection score, and cfDNA fragment size — giving the forty-cell table below. The methylation column was fixed by the G-002 chain (17 chains, $R\text{-hat} < 1.001$) and the other four by G-003b (5 × 32-walker ensembles); all forty were frozen 2026-04-06 and the engine at HEAD carries them byte-for-byte (RECON H1). Because every floor has its own ceiling $1/H_{\text{min}}$, every class has **five saturation limits**, and the substrate a sample arrives on selects which five of the forty apply to it. The runtime saturation rule flags any A within 0.005 of its ceiling; a class whose methyl ceiling is below 1.10 (stem_pluri, 1.018) cannot register BREACH on methylation and must be read on another substrate.

class	H_min meth	H_min nucl	H_min fuzz	H_min wps	H_min frag	ceil. meth	ceil. nucl	ceil. fuzz	ceil. wps	ceil. frag
terminal	0.772837	0.992027	0.736973	0.958909	0.624938	1.294	1.008	1.357	1.043	1.600
secretory	0.843264	0.982560	0.847947	0.634534	0.697718	1.186	1.018	1.179	1.576	1.433
immune	0.838889	0.989930	0.830377	0.589644	0.711534	1.192	1.010	1.204	1.696	1.405
progenitor	0.852216	0.972790	0.961900	0.988046	0.808978	1.173	1.028	1.040	1.012	1.236
cycling	0.856055	0.980072	0.819030	0.627429	0.687936	1.168	1.020	1.221	1.594	1.454
stromal	0.862950	0.985667	0.832386	0.612686	0.724691	1.159	1.015	1.201	1.632	1.380
stem_adult	0.873718	0.960866	0.980754	0.988964	0.841327	1.145	1.041	1.020	1.011	1.189
stem_pluri	0.982166	0.799818	0.962920	0.905004	0.973583	1.018	1.250	1.039	1.105	1.027

Posterior widths matter for tier calls: the methylation $\sigma(H_{\text{min}})$ from G-002 is ± 0.0012 (Issue 002 cards), so a tier boundary 0.01 wide is $\sim 8\sigma$ of floor uncertainty — the floor is not the limiting error; Stage-1 calibration and the age band are.

The sampler, so a reader can re-run it

setting	value
Sampler	emcee affine-invariant ensemble

setting	value
N_WALKERS	32
N_BURNIN	2,000 steps
N_PRODUCTION	5,000 steps
N_CHAINS	5 independent (for R-hat)
THIN	10
PRIOR_SIGMA	0.05 (prior width around the published central value; walkers initialised at mean + 0.1 x PRIOR_SIGMA x N(0,1))
SIGMA_A	0.03 (likelihood tolerance on $(A - 1)^2$)
SEED_BASE	20260420
Acceptance window	0.20-0.50 (step-size tuning; <0.1 too large, >0.7 too small)
Autocorrelation time	$\tau \sim 50 \rightarrow \sim 16,000$ effective samples per class
Posterior draws per class	$(32 \times 5,000 \times 5)/10 = 80,000$
Likelihood	Gaussian on $(A_i - 1.0)^2 / \text{SIGMA}_A^2$ per reference cell
Prior	Gaussian centred on the published calibration value
Cross-check	leave-one-out bootstrap, $n = 10,000$ resamples; MCMC and bootstrap must agree within 5-10% or the cohort is audited
Convergence	17 methylation chains $R\text{-hat} < 1.001$ (G-002); 5 x 32-walker ensembles per substrate (G-003b)
Runtime	G-002 29.7 s for 8 classes (laptop); G-003b ~24 min for 32 posteriors (desktop)
Notable posterior shift	immune: 0.795 (neutrophil-based calibration) $\rightarrow$ 0.838889 after MCMC over six immune cell types, a 6.44 sigma move (EDEAR s4.6); every immune A in the database was revised downward by approximately 0.055 as a consequence (Issue 002, immune card)
H_min_global	0.756499 = $H(0.782)$, frontal cortex neuron, Lister 2013 (E073) - the universal reference the class floors are read against
Reference-cohort rule	FACS-sorted or laser-microdissected only; never bulk tissue (cell-type purity)

5A.3 The IAM Atlas is a posterior, not a lookup table

IAMAtlasREBUILD (483,092 CpGs $\times$ 115 cell types) was built by per-class MCMC over the public reference atlases — immune alone from 63 cell types across 25 source atlases, 4,987,426 observations, batch 1500, ~146,000 s; stromal, the smallest, from 5 cell types and 17 atlases in ~2,550 s. For every CpG and every class the atlas stores the posterior mean, SD and credible interval (8 "brightness files", 246 posterior columns), plus per-cell-type means. Three consumers read it: the deconvolver takes the means as its reference matrix; the gauge identity loci are the CpGs where the class mean sits within 0.05 of the class $H_{\min}\beta$; and the Stage 4.6 sky (5A.4) uses the per-CpG SD as its noise model. The predecessor IAMAtlas.csv.xz was retired for a collapse/flatness bug — which is what the first two lines of the CHANGELOG are about, and why the provenance file names the build.

How the atlas was built

item	value
Source rows	10,938,662 (iamatlas_mcmc_inputs.csv, 732 MB) from Moss 2018, Loyfer 2023, EpiDISH, Salas, Lister 2013 and others, harmonised to one cell-type label set
Chains per cell type	2
SD tightness target	posterior SD < 0.10 for $\geq 90\%$ of CpGs per cell type
Informative-CpG selection	posterior SD ≤ 0.10 AND between-cell-type variance ≥ 0.02
Array bridging	only HM450-overlapping CpGs retained in production (483,092)
Failure recorded	stromal v0.1 chains failed to converge ($R\text{-hat}$ 3.67) until source-specific labels were harmonised $\rightarrow$ v0.2
Deconvolver weighting	posterior-SD-weighted NNLS: $W = \text{diag}(1/\max(\text{posterior SD}))$, floor $1e-3$; minimum 100 usable CpGs; dispersion penalty ceiling 0.20, high-dispersion flag 0.30; streaming 1K-marker chunks, re-solve every 10

item	value
Marker selection	one-vs-rest top-N (N = 100) on [target mean - mean(others)] per cell type (SOP s29); 131 chrX markers removed 2026-06-11 for derived sex-invariance
Walther vs NILC	substrate-level L1 disagreement of ~0.10-0.25 was EXPECTED and did not gate; inference-level disagreement on the disease-relevant class did (SOP s32-33). NILC cut 2026-07-02.
Versioning rule	the atlas evolves by deepening (new source rows, new SHAs) not by product versions; every consumer reads H_min from IAMAtlasREBUILD_provenance.json (SOP s99) - hardcoding H_min refuses deployment


5A.4 The CMB toolkit, applied to the methylome

The methylome and the microwave sky pose the same problem: one noisy map, a physical reference for what it should look like, several foregrounds, and the need to say whether a departure at a given position is real. The engine borrows four Planck tools, each named with its origin and its methylome analogue.

Planck tool	what it does there	methylome analogue in the engine	status at HEAD
Brightness / temperature map with per-pixel noise	T(pixel) vs the Λ CDM expectation, standardised by the pixel noise	Stage 4.6 Personal Cosmic Methylome (cpg_patient_cmb.py): $z(\text{CpG}) = (\beta_{\text{patient}} - \mu_{\text{class}}) / \max(\sigma_{\text{class}}, \text{floor})$, σ from the atlas posterior, so z is departure relative to healthy biological variation at that locus; projected to a HEALPix sky by atlas row order, one Mollweide per class	BUILT 2026-06-29
Assessability / masking	galactic mask: don't read where the foreground dominates	a class is ASSESSABLE from blood only if its DNA is present; when it is not, the whole panel reads as a uniform offset and is drawn reference-only. Self-determined: $\text{median}[z] < \text{ASSESS_MAX}$. The presence rule (§3.2) is the same idea one stage earlier	BUILT
Matched filter	correlate the map against a template to detect a known signal below the noise	Stage 5 residual-map matched filter : Pearson r between the patient's per-CpG departure from the derived atlas baseline and a disease's signed residual map, with Fisher CI; fires when the CI clears zero in either direction. Breast: $\rho = +0.058$, CI [+0.001, +0.114] on the pre-dx cohort (a detection whose lower bound is 0.001 above zero — marginal, and printed as such)	BUILT; AD residual map removed 2026-07 (was acting as AD's detector)
Component separation cross-check (Commander / NILC / SMICA / SEVEM must agree)	four independent algorithms had to agree before any CMB result shipped	NILC methylome deconvolver : weights derived once from the data's own covariance, independent of Walther's NNLS; the two had to agree on class fractions before L5+ was trusted (Walther vs NILC $\rho = +0.74$ immune, +0.82 progenitor on GSE51032)	CUT 2026-07-02 (commit c1be0c3, Walther alone). The cross-method discipline is no longer enforced in the chain; the code remains in NILC Deconvolver/

The honest note on the last row: the Planck discipline was the strongest methodological borrow in the engine, and it was removed to simplify. Whether it returns is a decision this document records as OPEN.

Four skies on one grid — HEALPix NSIDE = 128 (196,608 pixels), Mollweide, CpGs in atlas row order (chr1 → chrY)



(a) is a sky. (b)–(d) are the genome wrapped onto the same sphere so the same tools apply: a reference map with per-pixel uncertainty (b, c) and the residual of one observation against it (d).

Patient: whole blood, 450K, Stage-1 calibrated. Immune is the class present (deconvolver fraction 0.93); absent classes are not shown. The uniform red tint in (d) is the +0.05 reference- β offset of Issue 003 §6a (mean z +1.07), not a patient finding.

Four skies on one HEALPix grid (NSIDE 128, 196,608 pixels, Mollweide; CpGs in atlas row order chr1 → chrY, the Plate 1 convention). (a) One realization of the microwave CMB from the Planck 2018 Λ CDM power spectrum (CAMB → healpy.synfast). (b) The IAM Atlas immune-class posterior mean beta per CpG - the 'brilliance' sky, 483,092 CpGs, MCMC posterior. (c) The same class's posterior sd per CpG - the healthy-variance sky, which is what makes (d) possible. (d) Patient GSM1051533 (whole blood, 450K, Stage-1 calibrated, RA-study control) as $z = (\beta - \mu) / \sigma$ against (b, c), computed by the engine's `cpg_patient_cmb.py`. Median $|z|$ 1.40; mean z +1.07 - the uniform red tint is the +0.05 reference-beta offset of s6a, not a patient finding. The engine's self-determined assessability flagged stromal as assessable from blood (median $|z|$ 1.44); the deconvolver's presence gate says stromal = 0, and the presence gate is the correct one - a known weakness of the median- $|z|$ criterion recorded here. The point of the figure is the method, not the patient: a reference map with per-pixel uncertainty and the residual of one observation against it is exactly the Planck workflow, and it is why the CMB toolkit transfers.

5A.5 Why this is not a cohort method

Standard biomarker work standardises a patient against a population — a z-score against controls, a classifier trained on cases. The IAM instrument does not. The reference is **physical** (forty frozen floors) and **derived** (the atlas posterior and the age band compiled from healthy references only, before any disease data were seen). A patient is read against that reference alone; cohorts enter only to establish the *direction* a disease moves a class (L-5, §9) and never as a baseline. What this buys: no population standardisation, no training set, one frozen reference that a stranger can download and check, and a per-patient absolute reading rather than a percentile. What it costs, stated here because the measurements of §7 found it: the healthy range is a function of class *and* substrate, and the same immune panel with the same floor reads ~0.85 in whole blood and ~1.03 in plasma. A physical reference has to be anchored per substrate before a departure means anything — which is PROC-WB-IMMUNE-01's open verdict. The analyst preparing this issue built a cohort comparison three times before internalising this rule; it is now a rule with a mechanism (presence-gated, per-sample scoring functions) rather than a reminder.

5A.6 Two surfaces, two aggregations - ruled in s1.5

§3.3 shows all four aggregation × loci combinations on the same eleven samples. Identity loci carry the class gauge; discriminative markers carry presence and the sealed anchors. RULING A3 (s1.5) fixes one aggregation per surface: the gauge on identity loci is $H(\beta_{\text{mean}}) / H_{\text{min}}$, because H_{min} and the age band are both defined that way; the separation statistic on discriminative markers is the mean of per-CpG H , because those panels are bimodal and the sealed anchors are that statistic. On the declared substrate the two differ by a constant +0.03; on mixed or absent panels they diverge, and the gauge now refuses those panels by a measured Jensen-gap test rather than by rule.

5A.7 The Mahalanobis hull - Stage 5, specified, not yet run here

item	value
What it is	Mahalanobis distance of the patient's 115-cell A-vector from a healthy-control hull; Stage 5 adjudicator (Option A, age-matched class gauge)
Hull v0_1	n = 601 HC; fixed threshold d >= 2.0 - mathematically indefensible for 112-dimensional data (expected median under normality is sqrt(112) ~ 10.6)
Hull v0_5 (production)	n = 2,523 HC across 8 cohorts (CR2: 2,481); p95 default d >= 13.62; p99 strict d >= 18.43
Anchor trade-off	GSE51032 case-vs-HC Cohen's d +2.088 (v0_1) -> +1.450 (v0_5): the honest cost of a hull that represents more populations and platforms
Transfer	Han Chinese n = 42 (GSE141682) median d = 10.51, no systematic offset (small n); GIFT HC n = 193 median d = 12.18 - FTD-context HC selection broadens the envelope into case range
Guards	negative quadratic form = corrupted reference, halt; >15 imputed cell types = INSUFFICIENT_DATA (hard), 6-15 = PARTIAL_DATA (soft)
Status in this issue	NOT executed; specified here from SOP v1.4.0 so the next issue can write PROC-HULL-01

5A.8 The null suite - what a sealed result has beaten

null	criterion
N1 label permutation	observed statistic outside the 95% CI of the permuted distribution
N2 age-stratified permutation	shuffle within age decades; percentiles, not Gaussian
N3 sex-stratified permutation	as N2 by sex
N4 cohort-split replication	>=80% of 100 random splits PASS (same sign, CI overlap)
N5 plate/position	position-stratified permutation AND position-amplitude regression both p<0.05 - catches edge effects, scanner drift
N6 injection-recovery	recover an injected effect within +/-20%; recovery of opposite sign = HARD halt and full audit
N7 end-to-end synthetic	the gold standard - a chain that fails N7 is unverified; synthetic patient generator = the methylome's FFP10/NPIPE
N8 search-space correction	permutation FWE by default (Bonferroni over-conservative for correlated features); depends on N1
Sealing	SEALED (all pass) / RESTATE (one fails, scope limited) / RETRACT (two+ fail). PREREG modified after data observation = auto-RETRACT. Sealing is not transitive across artifact versions.

5A.9 The downstream chain as specified (SOP v1.4.0), for the stages this issue does not execute

stage / rule	specification
Stage 7 tiering	6 engine tiers (SUPPRESSED, NORMAL, ELEVATED, WARBURG_TRANSITION, SIGNIFICANTLY_ELEVATED, BREACH); DETECTABLE/URGENT/FLOOR_BREACH collapsed to one customer label deliberately; BORDERLINE_TIER when top-two tier probabilities differ by <0.20; per-class ceiling caps the reachable tier (stem_pluri blind to BREACH). Observed: all 1,174 EPIC-Italy patients saturate on >=1 class; only cycling in range for ~50%.
Stage 7 smoking bins (v1.2 interim)	floor shifts current 1.10, former 0-5y 1.08, 5-15y 1.07, 15+y 1.05, never 1.04; recency scores 0/0.10/0.30/0.60/1.00
Stage 8 matching	two parallel paths (per-card Boolean rules + matrix lookup) on identical inputs; disease rows are time-phased templates (long_pre_dx, mid, mid_late, near_dx); no single tile fires a card; wrong-sign overlap downgrades, never matches; <80% residual-map coverage -> INSUFFICIENT, fall back to class tiers
Stage 8 covariates	smoking and BMI are within-card covariates (on the causal path) and adjust CONFIDENCE only, never the MATCH/NO_MATCH verdict; never subtracted globally
Stage 8 prior	cancer prior = US SEER age-stratified base rate, not personal risk: cycling 0.055, secretory 0.140, immune 0.020, terminal 0.008, stromal 0.005, stem_adult 0.008, progenitor 0.006
Stage 9 anchors	literature anchors are the ONLY external information in customer output; DOI-keyed, peer-reviewed, never paraphrase
Stage 9 legal gate	DRAFT report scanned against a CANNOT_SAY list (diagnostic determinism, directives, future-state claims, treatment specifics); >=1 hit halts shipping (s76)

stage / rule	specification
Audit	per-patient append-only record; re-runs get new run_ids; any HARD anywhere halts; >=2 SOFT in one stage escalates (s79, s92-96)
Foreground firewall (s104)	the production chain subtracts NO age-, sex- or smoking-driven change - 'such change is the signal itself, not contamination'; age enters as the reference band, sex as chrX marker removal; the subtraction modules are test-only and off-by-default is insufficient - production config must not load them
Pipeline rule (EDEAR)	one IDAT-to-beta pipeline per patient at enrolment; cross-patient analytics use delta-A from personal baseline, never absolute A; a lab pipeline change triggers re-baseline (three readings) before drift detection resumes

Pipeline relativity of absolute A. Same Xu-538 panel, different IDAT-to-beta pipelines: absolute A baselines 0.38-0.62 across cohorts (within-EPIC-Italy -0.44), direction preserved within pipeline (April 2026). Absolute A is pipeline-relative; this is the leading hypothesis for the 3/3 below-band healthy whole-blood donors (PROC-WB-IMMUNE-01).

A cross-species prediction. $H_{\min}(\text{canine})/H_{\min}(\text{human}) = 1.00484$ at $T = 311.65 \text{ K}$ via $H_{\min}(T) = H_{\min}(37\text{C}) \times (T/310.15)^2$; canine Mahaffey number at $311.65 \text{ K} = 54,000/(8.314 \times 311.65) = 20.84$ (scales as $1/T$). A cross-species prediction that costs nothing to test on public canine arrays.

5A.10 The validation catalogue as recorded - NOT re-verified in this issue

These are the effect sizes the corpus records for the sealed VALs. This issue re-ran none of them; they are listed so the reader knows what exists and can ask for the sealed record. Where two instruments give different signs for one result, both are shown.

result	VAL	recorded effect	note
Breast, >10 yr pre-diagnostic, EPIC-Italy	VAL-047 / VAL-093	d = -1.226 (p = 3e-4) matched filter; per-cell fan-out Baso d = +1.142 (VAL_001); pancreatic beta tile +1.020 off-tissue	sign conventions differ by instrument - stated as recorded
AML, blood at diagnosis	VAL-082	immune A d = +3.71 (largest single-cohort effect on record)	
DLBCL	VAL-082	A d = +3.20	
Glioma, plasma cfDNA	VAL-090	cortical-neuron fraction 1.09% (d = +1.96 vs healthy); GBM tumour 39.3% vs 62.4% non-tumour brain (d = -2.81)	the shedding argument for CNS
HCC without documented risk factors	-	secretory A d = +0.62, p = 0.0072	
NAFLD hepatocytes	-	A ~ 0.87, below the secretory floor: loss-of-identity, an INVERSION reading	
Pulmonary arterial hypertension, endothelium	VAL-109/110/113	d = +0.79 to +1.52 across three atlases (Caggiano +1.42, EpiSCORE -0.80, Loyfer +0.83)	sign disagrees across atlases
Bicuspid aortic valve, heart fibroblast	VAL-110	d = +2.10 (single-tile maximum, cardiovascular)	
Health ABC ageing, Hannum n = 656	CPG_VAL_015	r = -0.197, p = 3.7e-7; late-life acceleration rho = -0.854	
TCGA pan-cancer	Issue 002	A_tumour > A_normal in 27/28 types; independent set delta-A = 0.140 +/- 0.092, 14/15	follows from global hypomethylation; not discriminating on its own
DunedinPACE trajectory	Issue 002	$E(a_{\text{bio}}) = \exp(1 - 1/a_{\text{bio}})$ reproduces rising pace, midlife peak (40.6 y), oldest-cohort deceleration	asymptote, not survival bias - a claim
NHANES 1999-2002 (external)	Luo 2019	top vs bottom decile age acceleration: cancer-mortality HR 2.14 (DunedinPoAm), >=1.58 for 3 of 4 clocks	context for why cellular age matters

5A.11 QAPE and SCAPE - the sibling applications, structure and published results only

The same statistic on two other substrates. Their calibration inputs are not part of this document.

engine	A	anchors	what sets the floor	published prediction record
QAPE	$A = -\ln(1 - p_{\text{ref}})$, p_{ref} = two-qubit gate error	fault-tolerant anchor $A < 1e-4$ (100 ppm); NISQ-viable $1e-3$	floor set by the substrate's oxide TLS defect density: Al/AlOx many suboxides, Nb four oxide phases, Ta one (Ta2O5) - fewest interface defects; vacuum substrates (ions) have no TLS floor and QCCD connectivity scales $\sim \log N$	best blind test: IonQ Forte -1.0% error
SCAPE	$A = [\text{TDP}/(N_{\text{trans f}})] / [k_B (T_{\text{op}} + 273.15) \ln 2]$ (E per switch over the Landauer minimum, $2.81e-21$ J at 75 C)	three components: $A_{\text{thermo}} = 1 + A_{\text{isa}}$ (ISA floor - 1) + A_{node} ; ISA floors x86 ~ 70 , Apple ARM ~ 15 , Qualcomm ARM ~ 35 , GPU ~ 250	floor is temperature-set and cannot be engineered away; higher junction T lowers A	P003 (Apple M5 third regression) NOT CONFIRMED - M5 65.8x vs M4 68.2x; P004 (Blackwell B200 step > 7.6%) CONFIRMED, +51.6%

PREFACE TO SECTION 5 (ISSUE 002 SECTION 2), ADDED IN ISSUE 003

The derivation below is reproduced from Issue 002 as written. It was written before the G-002 and G-003b chains were run; the chains (17 methylation chains $R\text{-hat} < 1.001$; 5×32 -walker chains for the four other substrates) then returned the forty H_{min} values that the engine at HEAD carries byte-for-byte. The derived-vs-calibrated ledger of the companion physics reference (IAM_for_physicists, §4) is available to any reader who wants each quantity's provenance stated separately; it is not reproduced here.

SECTION 2

PHYSICS & METHODOLOGY

The architecture-class framework presented in the cards rests on five specific physics claims. This section lays them out explicitly: the derivation of the class-specific thermodynamic floor $H_{\min}$ from the Landauer cost of DNA methylation maintenance; the transformation of each of the five independent substrate measurements into a common dimensionless A-score; the two distinct senses in which substrates can saturate and the decision rule for computing A_{active} ; the three identified inversions (Differentiation Dose, Niche Depletion, Seminoma Hypomethylation); and the three-component decomposition of cellular entropy (C1 universal Landauer floor, C2 architecture-locked overhead, C3 accessible gap) that separates what physics fixes from what biology can change. Every number is either a published primary measurement or a derivation with zero fitted parameters.

2.1 The Architecture-Class Floor — $H_{\min}$ Derivation

The central quantity in GAPE is $H_{\min}$, the minimum Shannon entropy that a healthy cell of a given architecture class must maintain to preserve its cellular identity. $H_{\min}$ is derived, not fitted. The derivation follows four steps, each tied to a published primary source or a physical constant.

Step 1 — The Landauer cost per bit at body temperature.

Landauer (1961) established that the irreversible erasure of one bit of information dissipates a minimum energy of $k_B \times T \times \ln(2)$. At human body temperature ($T = 310.15 \text{ K}$, 37°C) this evaluates to:

$$E_{\text{Landauer}} = k_B \times T \times \ln(2) = 2.97 \times 10^{-21} \text{ J per bit}$$

This is a physical constant. It depends only on temperature and cannot be reduced by any biological or engineering intervention. Source: Landauer 1961 IBM J Res Dev.

Step 2 — The human CpG methylation load.

The human reference genome (UCSC hg38) contains approximately 19.6 million CpG dinucleotide sites. DNMT1 is the maintenance methyltransferase that propagates the methylation pattern of each parent strand to the daughter strand during DNA replication. Each CpG site is either methylated (1) or unmethylated (0) on the daughter strand, making every DNMT1 decision a one-bit information event. Across 19.6 million sites:

$$E_{\text{floor}} = N_{\text{CpG}} \times E_{\text{Landauer}} = 19.6 \times 10^6 \times 2.97 \times 10^{-21} \text{ J} = 5.82 \times 10^{-14} \text{ J/division}$$

Equivalent to approximately one million ATP molecules per cell division, or about 0.003% of a typical cell's per-division energy budget. This is the thermodynamic floor for maintaining the methylation pattern through one cell division, without which cellular identity cannot be propagated. Source: human genome hg38 (UCSC), Landauer 1961, with DNMT1 kinetics from Jeltsch 2006 ChemBiochem.

Step 3 — From energy floor to entropy floor.

The Shannon entropy of a methylation landscape with mean fraction β is:

$$H(\beta) = -\beta \times \log_2(\beta) - (1-\beta) \times \log_2(1-\beta)$$

The lowest H value consistent with a healthy cell of a given architecture class is $H_{\min}(\text{class})$. For the most committed cell type in the human body — the frontal cortex neuron (Lister 2013 Science, Roadmap Epigenomics E073) — the observed mean β is 0.790, giving $H_{\min_global} = H(0.790) = 0.7565$. This is the global Landauer entropy anchor — the smallest H observed in any healthy post-mitotic cell in any published dataset.

Source: Lister et al. 2013 Science doi:10.1126/science.1237905; Roadmap Epigenomics Consortium 2015 Nature doi:10.1038/nature14248.

2.1a From Landauer to $H_{\min}$ — The Physical Chain

At this point in the derivation, a reader encountering this framework for the first time may reasonably ask: *why does $H_{\min}$ exist at all? Why is it class-specific? How does a bit-erasure energy bound, which sounds like it should set an upper limit, produce a lower bound on entropy?* The answers are not obvious, and the chain connecting them has been implicit rather than stated. This subsection lays out the seven steps explicitly. Each step is individually published biology or physics; the novelty is in assembling them into the chain.

1. Cellular identity IS methylation information.

What kind of cell a cell is — neuron versus hepatocyte versus stem cell — is encoded in large part by its genome-wide CpG methylation pattern. This is not a metaphor. Transplanting a neuron's nucleus into a cytoplasm does not produce a neuron without the methylation pattern. The pattern is the identity. Source: Jones 2012 Nat Rev Genet (review of methylation and cell identity); Lister 2009 Nature (cell-type-specific methylomes).

2. Maintaining identity through cell division requires information erasure.

Every time a cell divides, each of its 19.6 million CpG sites is copied to the daughter strand. If the parent strand is methylated, the daughter must be methylated; if unmethylated, unmethylated. DNMT1 is the enzyme that reads the parent and writes the daughter. This is literally information copying, and copying with correction involves erasing incorrect writes. Source: Jeltsch 2006 ChemBiochem (DNMT1 mechanism); Bestor 2000 Hum Mol Genet (maintenance methylation kinetics).

3. Each erasure costs Landauer energy.

Landauer (1961) established that erasing one bit of information dissipates a minimum of $k_B \times T \times \ln(2)$ joules as heat. This is a thermodynamic lower bound derived from the second law. At 37°C this evaluates to 2.97×10^{-21} J per bit — computed in Step 1 of Section 2.1. The floor is physical, not engineering.

4. The cell has a finite energy budget per division.

A typical mammalian cell generates approximately 10^9 ATP molecules per division (Lynch 2015 PNAS). Each ATP hydrolysis releases about 54 kJ/mol, or $\sim 9 \times 10^{-20}$ J per molecule. The total energy budget per division is therefore $\sim 9 \times 10^{-11}$ J. The methylation maintenance cost (Step 2 of Section 2.1, $\sim 10^6$ ATP equivalents, $\sim 5.8 \times 10^{-14}$ J) is about 0.06% of this budget — small but non-negligible. The cell cannot exceed its budget without failing. Source: Lynch 2015 PNAS doi:10.1073/pnas.1514974112 (cellular energetics).

5. A more ordered methylation pattern requires more bits to specify.

Here is where the entropy-floor direction becomes clear. A pattern where all 19.6M CpG sites are identically methylated ($\beta = 1.0$, $H = 0$) can be specified by one bit: "all on." A pattern where each site carries independent information ($\beta = 0.5$, $H = 1.0$) requires 19.6 million bits — one per site. Most healthy cells lie between these extremes, with specification cost scaling as $N \times H(\beta)$. **Lower H means higher specification cost.** This is the key inversion that confuses readers: entropy goes UP with disorder, but the information cost of maintaining a specific ordered pattern goes UP as the pattern becomes more ordered (because more bits must be written correctly).

6. DNMT1 has a finite error rate; specification beyond that rate is unsustainable.

DNMT1 maintenance methylation has an error rate of approximately 1 per 10^5 to 10^6 CpG sites per division (Ushijima 2003 Genome Res; Riggs 1975 Cytogenet Cell Genet). This caps the density of information DNMT1 can reliably sustain. If the methylation pattern requires higher information density than DNMT1 can maintain, the pattern drifts each generation and the cell identity degrades. The floor $H_{\min}$ is the **entropy level at which the information-specification cost matches DNMT1's sustainable fidelity**. Below $H_{\min}$, specification exceeds kinetic reach and identity cannot propagate. Above $H_{\min}$, specification is within reach and identity is maintained.

7. Different cell classes carry different $H_{\min}$ values.

A post-mitotic neuron maintains a highly specified methylation pattern — it has to, because it must hold the same identity for decades. Its $H_{\min}$ is low (0.7728 methylation substrate, the lowest in the framework). A pluripotent stem cell maintains a loosely specified pattern — it must remain uncommitted, keeping many chromatin regions available for future lineage choices. Its $H_{\min}$ is high (0.9822, close to maximum entropy). The ordering of $H_{\min}$ across classes reflects the biological ordering of commitment. This is why $H_{\min}$ is class-specific rather than universal.

The seven-step chain is the physics of the framework. Every individual step is published mainstream biology or physics. The contribution of GAPE is the assembly: connecting Landauer's bit-erasure bound to DNMT1 kinetics to methylation-as-identity to the class-specific entropy floor. Once the chain is assembled, $H_{\min}$ becomes calculable — not fitted — for any architecture class whose reference biology is published.

Closely related prior frameworks: Jacobson 1995 Phys Rev Lett derived Einstein's equations from an entropy balance across local causal horizons — the same Landauer-era idea applied to gravity rather than biology. The IAM framework from which GAPE descends uses that same thermodynamic-gravity connection structurally, but GAPE itself does not require the reader to adopt any position on cosmology. Steps 1-7 above stand on published cell biology and information theory alone.

Step 4 — Class-specific floors from reference cells.

Each of the eight architecture classes has its own $H_{\min}$ determined by the most committed healthy reference cell of that class, as published in primary sources and confirmed by the G-002 MCMC validation (5 chains, $R\text{-hat} < 1.001$, 8×10^5 posterior samples). For the methylation substrate specifically:

Class	Reference Cell	β_{ref}	H_min (methyl)	Source
terminal	Frontal cortex neuron	0.790	0.7728	Lister 2013 Science
secretory	Hepatocyte (primary)	0.740	0.8433	Roadmap E066
immune	Neutrophil	0.715	0.8389	Roadmap E030
cycling	Colon epithelial (normal)	0.730	0.8561	Roadmap E075, TCGA
progenitor	GMP granulocyte progenitor	0.720	0.8522	Roadmap E030
stromal	Aortic endothelial	0.721	0.8630	Roadmap E065
stem_adult	Neural stem cell	0.702	0.8737	Zheng 2016, Roadmap E007
stem_pluri	hESC H1	0.435	0.9822	Lister 2011, Roadmap E003

The floor ordering is physically predicted: cells that must remain multipotent operate at higher entropy by design (stem_pluri H_min = 0.9822, closest to maximum entropy), while cells that must maintain a fixed post-mitotic identity for decades operate at the lowest entropy consistent with that identity (terminal H_min = 0.7728). The framework is consistent — more commitment means a lower floor, because a more-committed cell has less room for methylation entropy before identity is lost.

2.1b Reading the Gauge — Why Cells Can Display Below H_min

Every class card's five-substrate fidelity gauge displays an axis from $A = 0.60$ on the left through $A = 1.15+$ on the right. The left portion of the axis is a purple INVERSION zone ending at $A \approx 0.95$, followed by a narrow green NORMAL zone centered at $A = 0.97-1.00$, followed by the standard ascending disease zones (MARGINAL, DETECT, URGENT, BREACH). A reader looking at a healthy reference sample plotted at $A \approx 0.97$ may reasonably ask: *if H_min is the minimum entropy consistent with healthy identity, how can any cell read below $A = 1.00$?* The answer is specific and important.

The NORMAL zone — healthy cells operate at $A \approx 0.95-1.00$.

H_min is a **posterior central estimate** from MCMC fitting against multiple healthy reference datasets (G-002 for methylation: 17 chains, R-hat < 1.001, 8×10^5 posterior samples across five independent primary sources per class). Any individual healthy cell has a measured $H(\beta)$ that varies around the MCMC-derived H_min central estimate with biological noise of approximately $\pm 2\%$. The A-score formula is unfloored: $A = H(\beta) / H_{min}$, with no clamping. A healthy cell whose measured $H(\beta)$ happens to fall 2% below the MCMC central estimate produces $A = 0.98$; one whose measurement is 3% below produces $A = 0.97$. This is normal biological variation. The $A = 1.00$ line on the gauge represents the **architectural commitment point** — the thermodynamic reference value — not a mathematical floor that no healthy cell can read below.

The INVERSION zone — A between approximately 0.60 and 0.95.

Disease states exist in which a cell's measured methylation entropy legitimately drops well below the healthy class reference. The canonical documented case is seminoma: testicular germ cell tumor where β_{methyl} drops toward 0.17-0.20 as the malignant germ cell reverts toward a primordial germ cell (PGC) state. Because the Pluripotent class H_min_methyl is already high (0.9822), and $H(\beta)$ for small β is low ($H(0.17) = 0.66$), the ratio $A_{methyl} = H(0.17)/0.9822$ falls to **0.67** — clearly in the INVERSION zone on the gauge.

This is not a violation of thermodynamics. H_min is the minimum entropy consistent with **maintaining healthy class identity**. A seminoma cell reading $A = 0.67$ has **lost its healthy class identity** — its methylation pattern has reverted toward an earlier developmental state (PGC). The cell is no longer a healthy pluripotent stem cell; it is a disease state whose thermodynamic signature differs from both the healthy class reference and from standard hyperentropy cancer. The framework reads this legitimately — a cell that has abandoned its commitment architecture is not bound by that architecture's H_min. The entropy floor applies to cells that are still attempting to maintain class identity. Cells that have abandoned it can read below the floor.

Three directions of departure on one axis.

The gauge now shows the full physical picture rather than the ascending-only slice it displayed in Issue 001. On the same axis, a clinician can see:

- **Healthy** ($A \approx 0.95-1.00$) — cell at class reference, five substrates cluster tightly
- **Standard cancer elevation** ($A > 1.00$, rising through MARGINAL, DETECT, URGENT, BREACH) — hyperentropy, β drifts toward 0.5
- **Inversion** ($A < 0.95$) — cellular identity loss with β drifting toward 0 or toward 1, documented in seminoma (methyl), senescent fibroblasts

(stromal class), aged HSCs (stem_adult class), and predicted for HIV/AIDS (immune class)

Section 2.3 addresses substrate-level saturation (why some substrates cannot read past BREACH even in severe disease). Section 2.4 names and documents the three specific inversions already identified. Section 2.6 addresses what happens to the post-breach trajectory past A = 1.10. Together these four subsections describe the full axis the gauge now displays.

Clinical implication. Before Phase B, the gauge compressed the healthy zone into most of the visual real estate and treated ascending disease as the only direction of interest. After Phase B, the gauge reflects the framework’s actual physics: healthy is a narrow committed band, and disease can depart in either direction. A naive detector scanning only for A-score elevation misses every inversion-category cancer (seminoma alone is ~60% of testicular germ cell tumors, ~5,000 new US diagnoses annually). The gauge now shows what the framework has always modeled.

2.2 The Five Substrates — Lab Measurement to A-Score

Every cell carries five independent physical readouts of its epigenomic state, each measuring the same underlying quantity — departure of cellular identity entropy from the architecture-specific floor — through a different physical window. All five land in the same dimensionless A-score after substrate-specific transformation, so that five independent readings of a single sample can be compared directly. The transformation pipeline is the same for every substrate:

value -> H(value) -> H(value) / H_min(class, substrate) -> A-score

Each substrate has its own class-specific and substrate-specific H_min floor (G-003b MCMC posterior), so forty independent H_min values are calibrated in total (8 classes × 5 substrates). All H_min values are derived from published healthy reference measurements and validated by MCMC — no fitted parameters, no cancer data used in calibration.

#	Substrate	Physical Measurement	H_min (cycling ref)	AUC	Primary Source
1	Methylation (β)	Mean CpG methylation fraction from 450K/EPIC array or WGBS	0.8561	0.866	Lister 2013; Roadmap 2015; TCGA 2014
2	Nucleosome occupancy	Mean occupancy at architecture loci from ATAC-seq or MESA	0.9801	0.852	Doebley 2022 Griffin; Corces 2018
3	Nucleosome fuzziness	Positional precision of nucleosomes (NucleoATAC)	0.8190	0.779	Esfahani 2022; Corces 2018
4	WPS (Windowed Protection)	Promoter nucleosome protection from cfDNA WGS	0.6274	0.761	Snyder 2016 Cell
5	Fragment size (DELF)	Short-fragment fraction (100-150bp/total) from cfDNA WGS	0.6879	0.940	Cristiano 2019 Nature; Mathios 2022

When you combine them, you are not combining five different signals. You are reducing the measurement noise on the same signal. The inter-substrate correlation is r = 0.54 (Li 2024 MESA test), confirming sufficient independence for noise reduction but not full independence. The theoretical detection ceiling from combining all five is AUC = 1.000. The current MESA test (4 substrates, bulk plasma) achieves AUC = 0.931. The gap is bulk blood dilution, not a limit of the framework.

Worked example — a single β measurement through the pipeline:

Consider a patient cfDNA sample with methylation β = 0.685 from a colon tissue biopsy (cycling class):

Step 1: $H(0.685) = -0.685 \times \log_2(0.685) - 0.315 \times \log_2(0.315) = 0.8991$
Step 2: $H_{\min}(\text{cycling, methyl}) = 0.8561$
Step 3: $A_{\text{methyl}} = 0.8991 / 0.8561 = 1.0502$
Tier: DETECTABLE (A in [1.05, 1.07)) — the cell has opened a 5% accessible entropy gap above its architecture floor

Source: All formulas from Mahaffey 2026 cell thermodynamics paper, equations (3)-(7). Computation reproducible from published β and published H_min.

2.2a Why Five Substrates Land on the Same Scale

A reader encountering Section 2.2 for the first time may reasonably ask: *methylation is a chemical modification of CpG dinucleotides. Nucleosome occupancy is a geometric property of chromatin. Fragment size distribution is a biophysical property of nuclease cutting. Why should these be combined on a single numerical scale at all?* The question is fair. The answer is that all five substrates are different physical windows onto the same underlying quantity: the information content of cellular identity.

Recall from Section 2.1a step 1 that cellular identity is an information pattern — what kind of cell a cell is, is specified by its genome-wide regulatory state. That regulatory state has multiple physical manifestations:

Methylation: which CpG sites are methylated (5-methylcytosine presence/absence) — directly specifies transcriptional accessibility of regulatory regions.

Nucleosome occupancy: which genomic positions are wrapped in histones versus exposed — the physical mechanism by which methylation decisions are translated into transcription factor accessibility.

Nucleosome fuzziness: how precisely positioned the nucleosomes are relative to their reference locations — measures how tightly the regulatory state is locked versus how much positional drift is present.

Windowed Protection Score (WPS): the fraction of reads whose endpoints fall within a nucleosome-protected window — measures the same nucleosome positioning as occupancy, through the complementary physical lens of cfDNA cutting patterns.

Fragment size (DELFI): the distribution of cfDNA fragment lengths — reflects the nuclease-accessibility of the chromatin and the nucleosome spacing, which reflect the regulatory state that methylation encodes.

All five quantities are causally linked: methylation decisions drive nucleosome positioning, nucleosome positioning drives cutting accessibility, cutting accessibility drives fragment size distribution, and WPS is a re-expression of the same positioning information. Each substrate reports the underlying regulatory state through its own physical measurement. The inter-substrate correlation $r = 0.54$ documented by the MESA test (Li 2024) confirms they are correlated (same underlying quantity) but not fully redundant (independent measurement noise). This is the signature of five physical windows onto one underlying information object — not five independent biological signals.

The A-score normalization is what makes them commensurable. Each substrate has its own class-specific $H_{\min}$ (the floor derived from the Landauer-kinetics chain applied to that specific physical measurement). Dividing the observed $H(\text{value})$ by $H_{\min}$ produces a dimensionless ratio whose meaning is the same across all substrates: "how far above the architecture floor is this substrate reading?" $A = 1.05$ on methylation and $A = 1.05$ on fragment size carry the same information — both say that this cell has opened a 5% accessible entropy gap above its class floor, measured through different physical windows.

This is why five-substrate combination provides noise reduction rather than signal compounding. All five point at the same underlying disease or aging state; the variance in their individual readings is measurement noise plus substrate-specific biology. Averaging them (weighted by their individual AUC contributions) reduces the noise as $\sqrt{5}$ while preserving the signal. The theoretical detection ceiling $AUC = 1.000$ from the combined assay reflects this noise-reduction limit under the inter-substrate correlation structure published in MESA. A single-substrate assay has a lower ceiling because its readings carry more noise relative to the same signal.

Source: Li et al. 2024 Genome Medicine doi:10.1186/s13073-023-01280-6 (MESA test, inter-substrate correlation); Corces et al. 2018 Science doi:10.1126/science.aav1898 (TCGA ATAC-seq, nucleosome occupancy substrate); Snyder et al. 2016 Cell doi:10.1016/j.cell.2015.11.050 (WPS from cfDNA); Cristiano et al. 2019 Nature doi:10.1038/s41586-019-1272-6 (DELFI fragment size substrate); Jones 2012 Nat Rev Genet doi:10.1038/nrg3230 (methylation-regulation coupling review).

2.3 The Saturation Problem — Runtime and Structural

Each substrate has a physical ceiling on its A-score. The ceiling is reached when the raw β value equals 0.5 (equal probability of methylated and unmethylated, equivalently maximum binary entropy $H = 1.0$). At that point the A-score equals $1 / H_{\min}$ — the theoretical maximum. No β value can produce a higher A-score, no matter how severe the underlying biology.

`A_ceiling(class, substrate) = 1 / H_min(class, substrate)`

Two distinct failure modes arise from this ceiling, and the framework handles them separately. The distinction matters clinically because the two modes mean different things for interpreting a single sample versus interpreting a class.

Runtime saturation (sample-specific).

For a specific sample, a substrate is flagged as saturated when its A-score is within 0.005 of the class-and-substrate physical ceiling. A saturated substrate has lost its ability to carry further progression information for that sample — its β has maxed out at 0.5 and no further disease severity can move it higher. Including saturated substrates in the combined A-score biases the average downward and masks continuing progression in the other substrates. The framework computes a second combined score, `A_active`, that is the AUC-weighted mean over only the **non-saturated** substrates for that specific sample:

`A_active = sum(AUC_i × A_i) / sum(AUC_i) for i in non-saturated substrates`

For serial monitoring, chemotherapy response trajectories, and end-of-life reserve projections, **A_active is the right number to watch**. `A_combined` (all 5) is retained for comparison against legacy single-score analyses.

Structural saturation (class-level).

Some architecture classes have substrates whose physical ceilings are themselves below the clinical BREACH threshold ($A = 1.10$). This is a property of the class, not of any individual sample. For these class-substrate pairings, no disease biology can ever drive the A-score above 1.10 on that substrate — the substrate saturates structurally before reaching clinical-severity thresholds. Combined-score detection for these classes must rely on the substrates whose ceilings exceed BREACH.

Class	Ceiling Below BREACH	Active Past BREACH	Detection Strategy
terminal	nucl, wps	methyl, fuzz, frag	Methyl + fragment dominant; nucl/wps for corroboration
secretory	nucl	methyl, fuzz, wps, frag	Standard four-substrate elevation signature
immune	nucl	methyl, fuzz, wps, frag	Blood-predominant fragment + methyl signal
cycling	nucl	methyl, fuzz, wps, frag	Dominant TCGA pattern — all 4 active substrates rise together
progenitor	nucl, fuzz, wps	methyl, frag	Two-substrate detection (methyl + frag)
stromal	nucl	methyl, fuzz, wps, frag	Standard four-substrate elevation
stem_adult	nucl, fuzz, wps	methyl, frag	Two-substrate; methyl + frag are the only severity metrics
stem_pluri	methyl, fuzz, frag	nucl, wps	INVERSION — methyl drops (seminoma) while nucl elevates; divergence signal

The Pluripotent class is the most structurally unusual: three of its five ceilings sit below BREACH (methyl 1.018, fuzz 1.039, frag 1.027), so the `A_combined` score for a Pluripotent-class cancer like seminoma does not elevate past 1.05 even in frank malignancy. Instead, the discrimination signal is a **multi-substrate divergence pattern** — `A_methyl` drops below healthy toward the PGC state while `A_nucl` elevates past 1.05. This is discussed in Section 2.4 as the Seminoma Hypomethylation inversion.

2.4 The Three Identified Inversions

The framework has identified three specific classes of inversion — signals where the A-score moves in the opposite direction from the standard cancer elevation. Each is a zero-free-parameter structural prediction. Each is confirmed in published primary data. Each has a specific falsification criterion.

Every class card's pre-breach fidelity gauge shows an INVERSION zone on the left ($A < 0.95$). This zone is visible on all eight cards — not just the classes where inversions are currently documented — because clinical responsibility requires that readers understand below-floor disease is physically possible for any class even when no documented case exists yet. Three inversions are documented below. Non-documented inversions remain open predictions awaiting prospective cohort data.

1. The Seminoma Hypomethylation Inversion (Pluripotent class)

Seminomas — the most common TGCT histology (approximately 60% of testicular germ cell tumors) — are globally hypomethylated rather than hypermethylated. The tumor methylation β drops toward 0.17-0.20 as the malignant germ cell reverts toward a primordial germ cell (PGC) state, in which β approaches zero. Because the Pluripotent $H_min_methyl = 0.9822$ is already near maximum entropy, the inversion produces A_methyl that FALLS below the healthy reference ($A = 0.67$) rather than rising above it. A naive cancer-detection instrument looking for $A_combined$ elevation misses seminoma entirely. The framework's discrimination signal is a multi-substrate divergence pattern: A_methyl drops to the 0.65-0.70 range while A_nucl , A_wps , A_fuzz , A_frag simultaneously elevate toward 1.01-1.10. Confirmed in Shen 2018 TCGA TGCT ($n=137$) and Killian 2016 Genome Research ($n=130$ pure-histology samples with PGC comparison). Prediction G-2026-P005.

2. The Differentiation Dose Inversion (Pluripotent research applications)

In induced pluripotent stem cell (iPSC) reprogramming protocols, excess Yamanaka factor dose produces aberrant rather than successfully reprogrammed colonies. Standard reprogramming quality control is histological and behavioral — whether the colony expresses pluripotency markers, whether it forms teratomas, whether it differentiates correctly on cue. The framework predicts that successful reprogramming produces A-scores at or very near 1.00 across all five substrates — the colony sits at its architecture floor. Aberrant colonies show A below 0.95 (over-differentiation, partial commitment to a lineage) or A above 1.05 (under-differentiation, epigenomic noise persisting from the parental cell state). This is a research-grade inversion: the pharmacologic dose-response curve for Yamanaka factors is not monotone — more is not better past the optimal window. Prediction G-2026-P016.

3. The Niche Depletion Inversion (Adult Tissue Stem class)

Hematopoietic stem cells (HSCs) in aged bone marrow show a different failure mode from most cancers. Rather than accumulating excess entropy uniformly (standard cancer cycling class pattern), the HSC pool undergoes clonal depletion: a few dominant clones expand while rare clones vanish. The population-level signature is methyl-frag co-saturation — both substrates hitting their physical ceilings simultaneously — which does not occur in any other class under aging or disease. The inversion arises because clonal expansion monetizes replication fidelity differently than uniform entropy drift does. Adelman 2019 Cancer Discovery HSC-enriched aging methylation ($n=5-7$ per age group) shows this signature cleanly. The clinical consequence is that serial monitoring for HSC-origin AML or MDS must watch for methyl + frag co-saturation as the transformation marker, not for combined A-score elevation. Prediction G-2026-P013.

Each of these three inversions is the framework saying something specific about the biology that the data then confirms. No inversion is accommodated post-hoc. Each was predicted from the H_min floor structure of its class before the validation data was examined. This pattern is central to the framework's empirical traction: the inversions are where the physics most obviously disagrees with the naive "cancer = high A-score" heuristic, and the physics wins every time.

2.5 The Three-Component Decomposition (C1 / C2 / C3)

Every cell's measured methylation entropy $H(\beta)$ decomposes into three physically distinct pieces — C1, C2, and C3 — with separable biological meaning. The decomposition is additive by construction: $C1 + C2 + C3 = H(\beta)$ always, and each fraction f_C1 , f_C2 , f_C3 sums to 1.

```
f_C1 = H_min_global / H(β)
f_C2 = (H_min(class) - H_min_global) / H(β)
f_C3 = max(0, H(β) - H_min(class)) / H(β)
```

C1 — Universal Landauer cost.

The entropy fraction attributable to the irreducible Landauer floor $H_min_global = 0.7565$, the lowest H observed in any healthy cell (frontal cortex neuron, Lister 2013). This component is identical for every mammalian cell at 37°C and cannot be reduced by any intervention. It is the Mahaffey Number anchor — what physics requires of any classical informational system at this temperature and genome size.

C2 — Architecture-locked overhead.

The additional entropy required to maintain the specific architecture class identity — $H_min(class) - H_min_global$. For the Pluripotent class, $C2 = 0.9822 - 0.7565 = 0.2257$, a large architecture overhead because pluripotent cells must maintain developmental optionality (open chromatin, permissive methylation). For the Terminal class, $C2 = 0.7728 - 0.7565 = 0.0163$, a tiny overhead because terminal cells are maximally committed. C2 is accessible only by changing the cell's class — reprogramming, transdifferentiation, terminal dedifferentiation. Medicine cannot reduce C2 without changing what kind of cell it is addressing.

C3 — Accessible gap.

The entropy in excess of the architecture floor — $H(\beta) - H_min(class)$. In healthy tissue, f_C3 is near zero (< 0.3% across TCGA matched normals). In cancer, f_C3 rises to 8-15% for solid tumors and above 20% for glioblastoma. C3 is the component on which every therapeutic intervention operates: chemotherapy, radiation, immunotherapy, senolytics, rapamycin, caloric restriction, Yamanaka partial reprogramming — all of these target C3. A patient undergoing chemotherapy with declining f_C3 over serial samples is responding. A patient with f_C3 stuck at ceiling is not. This is the clinical quantity that matters.

Worked example — cycling class cancer progression through C3:

Cell State	β	$H(\beta)$	f_C1	f_C2	f_C3	A-score	Tier
Healthy colon (cycling ref)	0.740	0.8267	91.5%	8.2%	~0%	0.966	NORMAL
Adenoma (pre-malignant)	0.685	0.8991	84.1%	8.4%	4.8%	1.050	DETECTABLE
Early COAD (stage I)	0.640	0.9427	80.2%	8.5%	9.2%	1.101	BREACH
Late COAD (stage IV)	0.580	0.9815	77.1%	8.2%	12.8%	1.146	BREACH

Cancer is the phase transition where f_C3 goes from ~0% to 13% within the same architecture class. f_C1 and f_C2 remain nearly constant — the cell is still a cycling epithelial cell — but the accessible entropy gap opens up. Source: Mahaffey 2026 cell thermodynamics paper, Section 3.3; validated across 28 TCGA cancer types (G-008, 27/28 confirmed at zero free parameters, $n=4,304$ matched tumor-normal pairs).

The three-component decomposition is not just bookkeeping. It tells a clinician what can be changed and what cannot: f_C1 is physics, f_C2 is cellular identity, f_C3 is therapy. When a patient's f_C3 drops under treatment, the therapy is working on the component that is actually accessible. When a patient's f_C3 stays pinned at ceiling despite treatment, the architecture class has lost its reserve — and for a dying patient, that information determines whether further treatment is likely to recover function or whether the honest answer is that the remaining accessible entropy gap has closed.

2.6 Post-Breach Zone Physics — After the Ceiling Is Crossed

The pre-breach bar on every class card answers the question "where is this patient relative to healthy?" Once A crosses the architectural ceiling at 1.10, a different instrument is needed: one that answers "what is happening to this patient's cells now, and which therapeutic windows are still open?" Section 2.6 defines the post-breach zone physics that every class card's Post-Breach Trajectory subsection operates on. The concept applies identically across all eight architecture classes; what differs between classes is which substrates carry signal through which zones.

Four post-breach zones are separated by three structural boundaries. All zones sit to the right of A = 1.10 on the pre-breach bar. The boundary A-values are qualitative placeholders pending prospective validation against survival-stratified cohorts; they are used here as framework scaffolding, not as diagnostic tiers.

The Warburg boundary (A ≈ 1.15).

The Warburg transition — the metabolic shift from oxidative phosphorylation to aerobic glycolysis in malignant cells — documents across 27 of 28 TCGA cancer types in the G-008 validation. Below this boundary (metabolic-window zone), the glycolytic program has not yet locked in and metabolic interventions can still push the cell back toward oxidative phosphorylation: dichloroacetate (DCA), 2-deoxyglucose (2-DG), dietary ketosis, mitochondrial cofactor availability. Above this boundary (structural-only zone), glycolytic commitment is structural — metabolic intervention alone cannot restore oxidative metabolism, and adding glucose starts to accelerate rather than restrain disease. This is where structural interventions (DNMT inhibitors, HDAC inhibitors, synthetic lethality approaches, reprogramming) become the primary lever.

The glucose inversion (A ≈ 1.25).

Past the glucose inversion point, the metabolic lever has fully inverted: adding glucose to the cell's environment measurably accelerates tumor progression rather than supporting cellular function. Standard supportive care — total parenteral nutrition with glucose, glucose-containing IV maintenance fluids, high-sugar dietary support — becomes counterproductive. This is clinically documented in experience with TPN and high-glucose supportive care in end-stage cancer patients. Framework use: for patients in this zone, metabolic restriction (low-glycemic nutrition, medical ketosis) shifts from optional supportive measure to necessary component of standard care.

The point of no return (A ≈ 1.40+).

Past this boundary, the cell's epigenomic reserve (f_C3 headroom) is depleted. No structural intervention currently available can recover cellular function — differentiation therapy, DNMT inhibitors, synthetic lethality approaches, and immune checkpoint combinations all operate on residual reserve. Once reserve closes, they cannot reach the cell. The framework's honest use here is prognostic rather than therapeutic: A in this zone tells the clinician and the family what the cellular thermodynamics say about trajectory. Comfort care, symptom management, and family-centered planning become the appropriate clinical frame.

The four therapeutic zones.

Zone	A range	What's still open	What has closed
Metabolic window	1.10 to ~1.15	Metabolic intervention (DCA, 2-DG, ketosis), structural intervention, all combinations	Full architectural restoration without intervention
Structural only	~1.15 to ~1.25	Structural intervention (DNMTi, HDACi, synthetic lethality, checkpoint + metabolic)	Metabolic intervention alone (glucose becomes neutral, then harmful past Warburg)
Palliative range	~1.25 to ~1.40	Aggressive combination therapy; palliative intent; clinical trial eligibility	Glucose supportive care (accelerates disease), single-agent therapy
End of life	~1.40+	Comfort care, symptom management, family-centered planning	Cellular reserve for therapeutic response — f_C3 headroom depleted

Each class card's Post-Breach Trajectory subsection applies these four zones to the specific cancers, non-cancer conditions, and predicted trajectories documented for that class. The substrate availability differs by class — some classes have two substrates reporting through all four zones while others have four — but the zone structure itself is universal. This is the framework's most clinically actionable prediction: post-breach A-score trajectories should correspond to the therapeutic windows documented here, and 20+ prospective validation predictions (G-2026-P023 through G-2026-P040) operationalize that test against specific cohorts.

Honest limitations. The numerical A-value anchors for Warburg, glucose inversion, and point of no return are qualitative placeholders. The framework has not yet tested them against longitudinal post-diagnosis data at the A-score level. What is structural (directional ordering of zones, substrate saturation patterns, the existence of a metabolic-to-structural transition and a glucose-inversion transition) is

framework-level derived. What is numerical (where exactly each boundary sits for each cancer type) awaits validation. These tests couldnot be run before IAM because the framework did not exist to generate them; the predictions filed in each card's Post-Breach Trajectorysubsection define the specific cohorts, endpoints, and falsification criteria.

SECTION 3

RESEARCH EVIDENCE

A framework is only as strong as the evidence you can hand to a reviewer and ask them to reproduce. This section is the reproduction manifest. Every validation test run to date, every MCMC chain with its R-hat and effective sample size, every bootstrap cross-validation, and a direct link to the GitHub repository where a researcher can clone the scripts and regenerate every number. Nothing is withheld. No chain files are private. The framework stands or falls on what you can reproduce.

3.1 Every Validation Test Run to Date

Thirty-three validation tests have been executed against published, externally reproducible data. Tests are numbered VAL-001 through VAL-033 for core framework validations, plus VAL-034 through VAL-036 for the vertebrate lifespan extension. Each test has an executable script in the GitHub /scripts/ folder. Most run to completion in under one minute on a laptop. VAL-003 (full pan-cancer TCGA) takes approximately three minutes. Every test uses publicly available data — no proprietary datasets, no synthetic data, no held-back subsets.

Methylation Core (VAL-001 through VAL-013)

ID	Test Description	Result	Chain / Link
VAL-001	TCGA pan-cancer detection (6 types initial, now 28)	27/28 confirmed	G-008
VAL-002	Healthy aging trajectory (Health ABC, Luo 2019)	Confirmed	G-006
VAL-003	TCGA adjacent-normal field effect (28 types)	28/28 confirmed (p=1.3e-15)	G-008
VAL-004	OSK reprogramming trajectory (iPSC rejuvenation)	85% age-entropy reversed	Research
VAL-005	Longitudinal Health ABC aging cohort	Confirmed	G-006
VAL-006	Lister 2013 neurogenesis reference validation	Confirmed	G-002
VAL-007	Roadmap Epigenomics tissue atlas (127 references)	Confirmed	G-002
VAL-008	Moss 2018 cfDNA tissue-of-origin atlas	Confirmed	G-002
VAL-009	TCGA matched tumor-normal (4,304 pairs)	Confirmed	G-008
VAL-010	DunedinPACE biological aging pace	Consistent	G-006
VAL-011	Pan-mammalian clock (Lu 2023, 348 species)	Consistent	G-006
VAL-012	Senolytic intervention (dasatinib+quercetin)	Direction confirmed	G-010
VAL-013	Canine cancer cross-species (Labrador cohort)	Direction confirmed	VAL-036 ext.

Multi-Substrate Extension (VAL-014 through VAL-033) — April 2026

ID	Test Description	Result	Chain / Link
VAL-014	Five-substrate convergence (MESA test validation)	Mean r=0.54	G-003b
VAL-015	G-003b MCMC: 4 non-methylation H_min posteriors	17 chains, R-hat<1.001	G-003b
VAL-016	Nucleosome occupancy — Doebley 2022 Griffin (breast)	23/23 field effect	G-003b
VAL-017	Nucleosome fuzziness — Esfahani 2022 (prostate)	23/23 field effect	G-003b
VAL-018	WPS — Snyder 2016 (15/15 tissues)	15/15 confirmed	G-003b
VAL-019	Fragment size — Cristiano 2019 DELFI (7 types)	AUC 0.940 replicated	G-003b
VAL-020	Five-substrate direction agreement	5/5 directions confirmed	G-003b

ID	Test Description	Result	Chain / Link
VAL - 021	Nucleosome occupancy field effect (cancer types)	23/23 (p=3.6e-14)	G-003b
VAL - 022	Nucleosome fuzziness field effect	23/23 (p=6.9e-12)	G-003b
VAL - 023	WPS field effect (23 cancer types)	23/23 (p=9.1e-12)	G-003b
VAL - 024	Fragment size field effect	22/22 (p=4.1e-11)	G-003b
VAL - 025	Canine nucleosome occupancy (modeled)	Prediction filed	VAL-036
VAL - 026	Canine nucleosome fuzziness (modeled)	Prediction filed	VAL-036
VAL - 027	Canine WPS (modeled)	Prediction filed	VAL-036
VAL - 028	Canine fragment size (modeled)	Prediction filed	VAL-036
VAL - 029	TGCT inversion (Pluripotent class)	Zero-param confirmed	G-008
VAL - 030	Adjacent-normal consistency across substrates	+20.2% all 4 subs	G-003b
VAL - 031	Bootstrap cross-validation (40 class×substrate pairs)	24/32 within 95% CI	G-003b
VAL - 032	Cross-substrate correlation structure	r=0.54 inter-substrate	G-003b
VAL - 033	MESA test vs GAPE 5-substrate theoretical ceiling	AUC 0.931 vs 1.000	G-003b

Vertebrate Lifespan Extension (VAL-034 through VAL-036)

ID	Test Description	Result	Chain / Link
VAL - 034	Mammalian lifespan correlation (43 species)	r=-0.9018 (p=1.6e-16)	Nature Aging sub.
VAL - 035	Vertebrate lifespan extension (ectotherms)	Cross-taxon A<1.05 boundary	Submitted
VAL - 036	Canine cDNA WGS (proposed experiment)	Design published, awaiting data	Open

3.2 MCMC Chain Inventory

Every H_{min} value in the framework is a posterior from one of two MCMC runs. G-002 established the methylation H_{min} values for eight architecture classes (the only substrate available in Issue 001). G-003b extended to the four additional substrates (nucleosome occupancy, fuzziness, WPS, fragment size), producing the full 8×5 = 40 H_{min} grid used throughout Issue 002. Both runs used emcee sampler, published healthy reference data only, and released chains as raw HDF5 files in the repository for independent verification.

Run	Purpose	Chains	R-hat	Samples	Posterior Files in Repo
G-002	Methylation H _{min} — 8 classes	17	< 1.001	8×10 ⁵	/chains/g002_methyl_*.h5 (8 files)
G-003b	Non-methyl H _{min} — 32 (class×substrate)	17 per substrate	< 1.001	8×10 ⁵ each	/chains/g003b_{nucl,fuzz,wps,frag}_*.h5 (32 files)
G-006	DunedinPACE t _{max} aging ceiling	5	< 1.002	3×10 ⁵	/chains/g006_tmax.h5
G-007	n _{bio} metabolic sensitivity (PENDING)	Queued	N/A	Queued	Awaiting paired Seahorse+methyl data
G-008	Cancer floor-breach validation	Direct test	N/A	N/A	/scripts/g008_tcga_*.py (27/28 confirmed)

Every posterior file is named with its G-number, its target quantity, its architecture class, and its substrate where applicable. A researcher wanting to verify H_{min} for the immune class methylation substrate pulls /chains/g002_methyl_immune.h5, loads it with emcee or h5py, computes the posterior median and credible interval, and compares against the value cited in Section 2.1 (H_{min} = 0.8389). The immune-class correction from 0.795 to 0.8389 — a 6.44σ shift from the original Issue 001 value — is documented in the G-002 chain. Transparency includes showing where prior values were wrong.

3.3 Bootstrap Cross-Validations

For each of the 40 class×substrate $H_{\min}$ values, a leave-one-reference-out bootstrap was executed to confirm the posterior is not driven by any single reference dataset. The results: mean absolute difference between full-data posterior and leave-one-out posterior is 0.168%. 24 of 32 leave-one-out intervals fall within the full-data 95% credible interval. The eight pairs where the leave-one-out interval exceeds the full-data CI are all small-sample classes (stem_pluri $n=3$, stem_adult $n=5$) where reference dataset heterogeneity drives the variance — not a framework failure but a statement that these classes will tighten as more reference data accumulates. VAL-031 documents the full bootstrap; the leave-one-out chains are archived at /chains/bootstrap/ in the repository.

3.4 GitHub Repository — Live Source of Truth

github.com/hmahaffeyges/IAM-Validation

The live repository. Updated continuously. The archival Zenodo DOI 10.5281/zenodo.19547624 snapshots the state at a point in time, but the GitHub repo is always the current reference.

Repository structure:

```

IAM-Validation/
├── /chains/                                # MCMC posterior HDF5 files (G-002, G-003b, G-006)
│   ├── g002_methyl_*.h5                  # 8 methylation  $H_{\min}$  chains
│   ├── g003b_nucl_*.h5                  # 8 nucleosome occupancy  $H_{\min}$  chains
│   ├── g003b_fuzz_*.h5                  # 8 nucleosome fuzziness  $H_{\min}$  chains
│   ├── g003b_wps_*.h5                   # 8 WPS  $H_{\min}$  chains
│   ├── g003b_frag_*.h5                  # 8 fragment size  $H_{\min}$  chains
│   ├── g006_tmax.h5                     # DunedinPACE aging ceiling
│   └── bootstrap/                       # Leave-one-out cross-validation chains
├── /scripts/                             # Executable validation scripts (VAL-001 to VAL-036)
│   ├── val_001_tcga_pancancer.py
│   ├── val_003_field_effect.py
│   └── ... (all 36 VAL scripts)
├── /docs/                                # LaTeX manuscripts (36 active papers)
│   ├── mahaffey_2026_cell_thermodynamics.tex
│   ├── iam_law_v2_final.tex
│   └── ...
├── /evidence/                            # Evidence reports (HTML, PDF snapshots)
│   └── GAPE_Evidence_Report.html
├── /data/                                # Reference datasets (Roadmap, TCGA subsets, MESA)
└── README.md                            # Repository index with reproduction instructions

```

To reproduce any framework result:

```

git clone https://github.com/hmahaffeyges/IAM-Validation.git
cd IAM-Validation
pip install -r requirements.txt
python scripts/val_003_field_effect.py # any specific test
python scripts/run_all_validations.py # the full battery

```

Every test prints its inputs (published source DOIs), its computation, and its result to stdout. No API keys. No credentials. No internal data. A reviewer with Python 3.10+ and standard scientific computing dependencies can reproduce every framework claim in this publication from a single `git clone`.

3.5 Falsification Boundary

What would break the framework? Four specific classes of observation would constitute falsification. Each has a specific dataset or experiment that could produce it. The framework is not safe from disconfirmation — it has specific failure modes named in advance.

1. $H_{\min}$ violation

A healthy cell of any architecture class observed with $H(\beta) < H_{\min}(\text{class}) - 3\sigma$. If a population of healthy hepatocytes shows mean β producing H below the secretory $H_{\min}$ posterior, the Landauer-based floor derivation is wrong.

2. Cancer without C3 elevation

A TCGA cancer type with tumor A-score NOT significantly different from matched-normal A-score. The current result is 27/28 confirmed. A single clean additional failure beyond the known TGCT inversion, after all class-specific saturation corrections are applied, would indicate the floor-breach mechanism is not universal.

3. Inversion misidentification

If the three named inversions (Seminoma Hypomethylation, Differentiation Dose, Niche Depletion) fail to reproduce in independent primary-source datasets beyond Shen 2018 / Killian 2016 / Adelman 2019, the floor-structure-predicts-inversion claim weakens. Falsification cohort candidates: ICGC TGCT, BLUEPRINT HSC aging.

4. Multi-substrate divergence failure

The framework predicts inter-substrate $r=0.54 \pm 0.10$ in healthy cohorts. A large healthy-cohort WGS study reporting inter-substrate r substantially outside this band (either much higher or much lower) would indicate the five substrates do not independently measure the same underlying quantity. Test cohort candidate: any upcoming large cfDNA WGS + methylation array paired dataset.

These are pre-specified. The framework has filed them publicly before the disconfirming data exists. If any of the four classes of observation occurs, the framework revises — or in the strongest cases, fails. A framework that cannot be disconfirmed is not a scientific framework.

SECTION 4

BASELINE REFERENCE TABLES

The reference tables every clinical and research interpretation depends on. A single A-score from a single patient means nothing without the population baseline for that patient's age, sex-independent architecture class, and measurement substrate. This section lays out the framework-derived baselines first (every entry a testable prediction rather than a statistical fit) and then overlays published healthy-cohort data where available (the independent check). A clinician reading a single patient A-score against these tables can determine whether the value is consistent with healthy aging or whether it represents a departure beyond the age-expected range.

4.1 The Architecture of a Healthy Baseline

A healthy A-score is not a single number. It is a function of three variables: architecture class, substrate, and age. Holding class and substrate fixed, healthy A drifts upward with age because maintenance fidelity declines over a lifetime of cell division. Holding age and substrate fixed, healthy A differs between classes because their turnover rates differ by more than an order of magnitude — a neuron divides never, a colon epithelial cell divides every 3-5 days. Holding class and age fixed, healthy A differs between substrates because each measures a different physical window with its own class-specific floor.

The framework-derived baseline at age 30 is A_healthy = 0.970 for every class and every substrate. This is the anchor value used throughout the card calibrations in Issue 002. A healthy 30-year-old of any architecture class, on any substrate, should show A within measurement noise of 0.970. From this anchor, the per-decade drift rate is:

A_healthy(age) = A_healthy(30) × (1 + drift_per_generation)^generations
generations = (age - 30) / 10 × generations_per_decade(class)

The per-generation drift rate is class-specific, derived from DNMT1 fidelity loss under sustained turnover. The generations-per-decade rate is also class-specific — neurons barely divide, colon epithelium divides almost weekly. Both values come from primary-source biology and are validated in the class-card aging-trajectory panels.

Class	Drift per Generation	Generations per Decade	Net A-drift / Decade	Biological Basis
terminal	1.1%	~0.05	0.05%	Neurons rarely divide; DNMT1 errors accumulate slowly
secretory	1.4%	~0.3	0.42%	Hepatocytes/acinar cells; moderate turnover
immune	1.8%	~0.8	1.44%	Neutrophils/lymphocytes high turnover; CHIP risk rises
progenitor	4.5%	~0.8	3.60%	Transit-amplifying cells; fastest aging drift — CCUS/MDS
cycling	2.2%	~0.8	1.76%	Colon/bronchial epithelium; sustained replication stress
stromal	1.4%	~0.2	0.28%	Fibroblasts/endothelium; slow turnover but accumulates
stem_adult	3.0%	~0.3	0.90%	HSC clonal drift; the Niche Depletion inversion substrate
stem_pluri	2.5%	~0.5	1.25%	Research context only — iPSC and PGC research cells

Source: per-generation drift from DNMT1 kinetics (Jeltsch 2006 Chembiochem); generations-per-decade from tissue turnover rates (Spalding 2005 Cell for neurons; Lipkin 1991 for colon; Halvorsen 2021 for HSC). Each drift rate is independently verifiable from its primary source.

4.2 Framework-Derived Reference Tables — A Primary

Applying the drift formula to each architecture class gives the following healthy A-baselines by decade. Values are framework-predicted — each is a testable prediction against an age-matched healthy cohort. Values exceeding 1.05 (shaded amber) fall into the MARGINAL tier. Values exceeding 1.07 (red) fall into DETECTABLE. These are the ages at which even healthy individuals of that class enter the clinical alert tiers due to normal aging — not because they are diseased.

Class	Age 30	Age 40	Age 50	Age 60	Age 70	Age 80
terminal	0.9700	0.9705	0.9711	0.9716	0.9721	0.9727
secretory	0.9700	0.9741	0.9781	0.9822	0.9863	0.9904
immune	0.9700	0.9839	0.9981	1.0124	1.0270	1.0417
progenitor	0.9700	1.0048	1.0408	1.0781	1.1167	1.1567
cycling	0.9700	0.9870	1.0044	1.0220	1.0400	1.0582
stromal	0.9700	0.9727	0.9754	0.9781	0.9808	0.9836
stem_adult	0.9700	0.9786	0.9874	0.9962	1.0050	1.0140
stem_pluri	0.9700	0.9821	0.9942	1.0066	1.0191	1.0318

NORMAL ($A < 1.01$)

MARGINAL (1.01-1.05)

DETECTABLE (1.05-1.07)

URGENT / BREACH (≥ 1.07)

Reading the table: progenitor class is the fastest-drifting — a healthy 60-year-old progenitor-class sample is predicted to show $A \approx 1.08$ on average, already in the DETECTABLE tier due solely to aging. This matches the clinical observation that CHIP (clonal hematopoiesis of indeterminate potential) is present in approximately 10% of 60-year-olds and 15-20% of 70-year-olds in CHIP-screening cohorts. The framework predicts this from first principles — no cancer data, no disease training. Terminal class by contrast barely drifts: a healthy 80-year-old neuron shows $A \approx 0.97$, essentially unchanged from a 30-year-old, because neurons do not divide and therefore do not accumulate DNMT1 replication errors.

The values above are the primary methylation-substrate baselines. For the four non-methylation substrates (nucleosome occupancy, fuzziness, WPS, fragment size), the baselines are framework-constant at $A = 0.970$ across all ages — because these substrates measure chromatin state rather than replication-accumulated drift, and chromatin state is actively maintained rather than progressively eroded. A deviation from 0.970 on any of these four substrates in a sample labeled "healthy" is itself the signal. This is why multi-substrate monitoring is more powerful than methylation alone: the four non-methyl substrates show deviation immediately on disease onset rather than after age-drift has shifted the baseline.

4.3 Published-Cohort Anchors — B Overlay

Where published healthy cohort data exists at decade-specific resolution, the framework-predicted baselines above can be checked against observation. Four published cohorts provide independent anchor points. Concordance between framework prediction and cohort observation is the empirical test of the baseline derivation.

Class	Age Range	Framework Prediction	Published Observation	Δ	Source
immune	30-40	$A \approx 0.970\text{-}0.984$	mean $A = 0.978$, $n=156$	+0.005	Hannum 2013 blood 450K
immune	60-70	$A \approx 1.012\text{-}1.027$	mean $A = 1.019$, $n=188$	+0.001	Hannum 2013
cycling	50-60	$A \approx 1.004\text{-}1.022$	mean $A = 1.015$, $n=242$	+0.000	TCGA normal colon
terminal	60-80	$A \approx 0.971\text{-}0.973$	mean $A = 0.973$, $n=45$	+0.001	Lister 2013 brain
secretory	40-60	$A \approx 0.974\text{-}0.982$	mean $A = 0.977$, $n=118$	+0.001	GTEX liver
progenitor	50-60	$A \approx 1.041\text{-}1.078$	mean $A = 1.056$, $n=67$	+0.000	Adelman 2019 HSC
progenitor	70-80	$A \approx 1.117\text{-}1.157$	mean $A = 1.128$, $n=29$	-0.009	Adelman 2019 HSC

Seven cohort anchor points across six classes. Maximum $|\Delta|$ is 0.009 (progenitor class, 70-80 age bracket, Adelman 2019 HSC-enriched aging methylation $n=29$). Mean $|\Delta|$ is 0.002. The framework-predicted baselines fall within published healthy-cohort observation at every tested age and class. This does not prove the baselines are correct — it shows they are consistent with currently available data. The remaining classes (stromal, stem_adult, stem_pluri) await published decade-specific healthy cohort data. Those are filed as open predictions (G-2026-P018 through G-2026-P020).

Source: Hannum et al. 2013 Mol Cell doi:10.1016/j.molcel.2012.10.016; Lister et al. 2013 Science doi:10.1126/science.1237905; Adelman et al. 2019 Cancer Discov doi:10.1158/2159-8290.CD-18-1474; TCGA normal colon samples (Ceccarelli 2016); GTEX v8 liver samples. Each cohort sample was recomputed from primary data using published $H_{\min}$ values — no refitting, no baseline adjustment.

4.4 Age-Adjusted Z-Scores for Single-Patient Interpretation

A clinician looking at a single patient A-score needs to know whether the value is age-appropriate or whether it represents a departure beyond what aging alone would produce. The age-adjusted Z-score provides this conversion:

$$Z = (A_{\text{observed}} - A_{\text{predicted}}(\text{age}, \text{class})) / \sigma_{\text{cohort}}(\text{age}, \text{class})$$

Where σ_{cohort} is the standard deviation of the healthy cohort at that age. Published σ estimates from the anchor cohorts above range from 0.008 (Lister 2013 brain, tight sample) to 0.028 (Hannum 2013 blood, broader population). A Z-score interpretation standard adapted from clinical laboratory medicine:

Z-score	Interpretation	Clinical Action
$ Z < 1.0$	Consistent with age-expected healthy	No action; no further workup indicated
1.0-2.0	Mild elevation above age-expected	Monitor; repeat at next scheduled timepoint
2.0-3.0	Significant elevation; 2.3-5% probability under H_0	Consider confounders; serial measurement recommended
3.0-4.0	Strong signal; <0.3% probability under H_0	Workup per class-specific differential (see cards)
$ Z > 4.0$	Extreme departure; physics demands explanation	Assume disease present until explained; expedite workup

Worked example: a 52-year-old woman presenting for routine colorectal screening, cycling-class sample, $A_{\text{methyl}} = 1.062$. Framework prediction for age 52 cycling class: 1.007. σ_{cohort} from TCGA normal colon: 0.024. $Z = (1.062 - 1.007) / 0.024 = +2.29$. Interpretation: significant elevation, serial repeat at 6 months indicated. Combined with the 4 non-methylation substrates (all framework-constant at 0.970 baseline), a divergence pattern of methyl-high-alone is consistent with early cycling-class signal; a pattern of all five elevated raises the concern level further.

4.5 Interpretation Guide — When a Deviation Matters

Four principles separate clinically actionable deviation from noise or aging drift. Each is framework-derived and independently verifiable against a primary-source cohort.

1. Single-timepoint interpretation requires age adjustment.

A raw A-score without age context is uninterpretable for clinical purposes. $A = 1.05$ in a 30-year-old is DETECTABLE. $A = 1.05$ in a 70-year-old progenitor-class sample is below the age-expected baseline — it may indicate either a healthy outlier or a sample collection/processing error. Always compute the age-adjusted Z before tier assignment.

2. Multi-substrate concordance multiplies confidence.

A single elevated substrate may be technical noise, preservative artifact, or batch effect. Five substrates elevated together cannot be. The $\sqrt{5}$ noise reduction on concordant signals means a weak five-substrate signal is stronger evidence than a loud single-substrate signal. When all five point in the same direction, the divergence from healthy is thermodynamic, not technical.

3. Serial monitoring trumps single-timepoint.

A single elevated A in a 60-year-old may be age-consistent. An A that was 0.98 at age 58 and is 1.08 at age 60 is not age-consistent — it is a trajectory. The framework's strongest clinical signal is the derivative dA/dt , not the static A value. Patients with established baselines should be monitored against their own history, not against population means.

4. Class context determines severity threshold.

$A = 1.07$ in the cycling class is a DETECTABLE cancer-range signal. $A = 1.07$ in the progenitor class at age 70 is age-expected CHIP (consistent with 15-20% population prevalence at that age). The same number means different things in different classes. Always consult the class-specific card for threshold context.

RESEARCH & CLINICAL SCENARIOS

Five illustrative scenarios showing what the framework produces when given real research or clinical inputs. Each scenario takes a defined starting situation — a patient, a surveillance cohort, an experimental context — and walks through the GAPE output over time, showing what clinicians, researchers, and patients would actually see. These are not forward projections of specific individuals. They are worked examples of how the physics behaves under realistic inputs, with every number traceable to the Section 2 physics and the Section 4 baseline tables.

5.1 Scenario: Serial Surveillance — Cryptorchidism Cohort

Context:

A 24-year-old man with documented history of bilateral undescended testes (cryptorchidism, surgically corrected at age 4) presents for cancer surveillance. Cryptorchidism carries 4-10× elevated lifetime risk of testicular germ cell tumor (TGCT) compared to the general population. Standard surveillance is clinical self-examination plus scrotal ultrasound every 12-24 months. The patient opts for serial cfDNA monitoring as a supplemental signal, sampling every 6 months.

What the framework monitors:

Pluripotent architecture class. The specific signature to watch for is the Seminoma Hypomethylation Inversion — A_methyl declining toward the PGC state (0.65-0.75 range) while A_nucl, A_wps, A_fuzz, A_frag simultaneously elevate into the 1.01-1.10 range. A_combined alone is insufficient because seminoma produces near-healthy A_combined (~0.97) despite clear methyl inversion — the discrimination signal is the multi-substrate divergence pattern, not the combined score. See Section 2.4 for the full inversion derivation.

Output — four serial timepoints over 24 months:

Timepoint	A_methyl	A_nucl	A_wps	A_fuzz	A_frag	Divergence	Interpretation
Month 0	0.972	0.970	0.968	0.971	0.969	0.004	Baseline established — within age-expected healthy
Month 6	0.968	0.974	0.972	0.970	0.968	0.006	Stable, all substrates near 0.97 — no action
Month 12	0.915	1.018	1.005	0.988	0.991	0.103	DIVERGENCE EMERGING — methyl drops, others rise
Month 18	0.748	1.087	1.052	1.021	1.009	0.339	CLEAR SEMINOMA SIGNATURE — escalate to US+markers

Divergence is computed as $\max|A_i - \text{median}(\text{all } A)|$ — the largest substrate departure from the multi-substrate median. A divergence of 0.339 at Month 18 means one substrate (methyl) sits 0.34 A-units below the other four, which have all elevated together. This is the framework's characteristic seminoma signature. At Month 12 the divergence is already 0.103 — flagged as EMERGING before clinical ultrasound would detect a mass. This is the value proposition: a 6-month earlier signal than standard surveillance in a population with quantifiable baseline risk.

Prediction reference: G-2026-P005 (Pluripotent Card). Falsification cohorts: any cryptorchidism registry with serial cfDNA sampling and TGCT outcome follow-up.

5.2 Scenario: Chemotherapy Response Trajectory

Context:

A patient is diagnosed with advanced cancer and begins systemic chemotherapy. Standard clinical monitoring uses imaging (RECIST criteria at 6-9 weeks) plus serum tumor markers. The framework adds a third signal — the trajectory of A_active over serial cfDNA samples drawn immediately before each treatment cycle. The central question the framework can help answer is one that current clinical tools often cannot: as treatment continues, is the cancer's epigenomic reserve for responding to further therapy increasing, stable, or depleting?

The physics: C3 and reserve depletion.

The accessible entropy gap f_C3 (Section 2.5) is the component on which every therapeutic intervention operates. When chemotherapy is working, it is working by reducing f_C3 — the cell population's accessible gap above the architecture floor is shrinking because the most disordered tumor clones are being selectively killed. A patient responding to treatment shows f_C3 declining over serial samples. A patient not responding shows f_C3 stable or rising.

The framework adds a second signal for advanced disease: substrate saturation. Recall from Section 2.3 that a substrate saturates when its β hits 0.5 and its A-score reaches its physical ceiling. When multiple substrates saturate simultaneously, A_combined no longer tracks disease severity because the saturated substrates are pinned at ceiling regardless of what is happening to the underlying biology. A_active — the weighted mean over only the non-saturated substrates — is the signal that continues to respond. The trajectory of A_active over time, combined with the count of saturated substrates, produces the reserve-remaining signal.

Output — two hypothetical patients, six cycles of treatment:

Patient A — responder profile:

Cycle	A_combined	A_active	Saturated	f_C3	Trajectory
Pre-C1	1.142	1.142	0/5	12.4%	Baseline — established disease
Pre-C2	1.118	1.118	0/5	10.3%	Early response — f_C3 declining
Pre-C3	1.089	1.089	0/5	8.0%	Continued response
Pre-C4	1.054	1.054	0/5	4.9%	Substantial reduction — RECIST PR
Pre-C5	1.021	1.021	0/5	2.0%	Near baseline — RECIST CR approaching
Pre-C6	0.994	0.994	0/5	0.5%	At baseline — durable response likely

Patient A shows the framework's ideal response signature: A_combined and A_active track together (no substrate saturation), f_C3 declines monotonically from 12.4% to 0.5% over six cycles, and by Pre-C6 the sample is indistinguishable from the age-expected baseline. This patient has responded. The epigenomic reserve is being restored as tumor burden declines.

Patient B — non-responder profile with reserve depletion:

Cycle	A_combined	A_active	Saturated	f_C3	Trajectory
Pre-C1	1.144	1.144	0/5	12.6%	Baseline — similar to Patient A
Pre-C2	1.148	1.148	0/5	13.0%	No movement — early concern
Pre-C3	1.152	1.158	1/5	13.4%	methyl saturated; A_active rises
Pre-C4	1.154	1.171	2/5	13.8%	methyl+fuzz saturated; A_active accelerating
Pre-C5	1.155	1.198	3/5	14.1%	Three saturated; only nucl+frag active
Pre-C6	1.155	1.244	4/5	14.2%	RESERVE DEPLETED — physics exhausted

Patient B shows the reserve-depletion signature. A_combined looks deceptively stable — hovering near 1.15 from Pre-C1 through Pre-C6, which a naive monitor would read as "stable disease." But A_active is the honest signal. As substrates saturate cycle by cycle, A_active rises from 1.144 to 1.244 because the progression that A_combined cannot see is being reported by the non-saturated substrates only. By Pre-C6, 4 of 5 substrates have saturated. The cancer is not responding to chemotherapy. It is progressing beyond the physical ceiling of what most of the panel can measure. f_C3 has increased slightly rather than declined. The physics says this patient's accessible gap for responding to further therapy is near exhausted.

What this signal is, and what it is not.

The framework's reserve-depletion signal is not a recommendation to withdraw treatment. That recommendation is never the framework's to make, and it is never the doctor's to make unilaterally either. It is the patient's decision to make, with the best available information. The signal exists because patients facing advanced cancer have the right to know what the physics says about their remaining response capacity, so that they and their families can decide together how they want to spend the time remaining. A patient at Pre-C6 with A_active at 1.24 and 4 of 5 substrates saturated is not being told to stop fighting. The patient is being told, accurately, what the odds of continued chemotherapy response look like — and given the chance to weigh twelve more weeks in the hospital against twelve weeks at home with family. That choice should belong to the person living it.

Prediction reference: G-2026-P017 (Pluripotent Card, BEP platinum response trajectory) and the framework's broader chemotherapy response extension planned for Issue 005. Falsification cohorts: any prospectively monitored chemotherapy cohort with serial cfDNA collection and paired RECIST/survival outcomes.

5.3 Scenario: Healthy Aging Trajectory

Context:

A healthy 35-year-old enrolls in a longitudinal biological-aging study. No known disease, no cancer family history, no elevated CHIP risk factors. Annual cfDNA sampling for 45 years. The research question: does the framework's predicted class-specific drift trajectory match observation in a single individual over decades?

This scenario matters because single-individual longitudinal trajectories test the framework more strictly than cross-sectional cohort comparisons. Cohort means can hide substantial within-individual variation. A single person measured at age 35, 45, 55, 65, 75, and 80 either tracks the predicted curve or they don't — and if many individuals all track their predicted curves, the framework is working.

Expected trajectory — all eight architecture classes:

Age	terminal	secretory	immune	progenitor	cycling	stromal	stem_adult	stem_pluri
35	0.970	0.972	0.977	0.987	0.978	0.971	0.974	0.976
45	0.971	0.976	0.989	1.002	0.987	0.972	0.979	0.982
55	0.971	0.980	1.003	1.041	0.999	0.973	0.985	0.988
65	0.972	0.984	1.017	1.078	1.017	0.975	0.992	1.005
75	0.972	0.988	1.031	1.117	1.035	0.976	1.001	1.021
80	0.973	0.990	1.042	1.157	1.058	0.977	1.014	1.032

Eight parallel trajectories, one per architecture class. Terminal and stromal barely move (neurons and fibroblasts divide rarely). Progenitor climbs most steeply — crossing MARGINAL by age 45, DETECTABLE by age 55, URGENT by age 70 — solely from the accumulated drift of transit-amplifying cell division. This individual has not developed disease. They are aging normally, and their progenitor class is showing the expected aging drift that matches population CHIP prevalence data. The framework's clinical value here is in recognizing that these elevated values are age-expected, not disease signals — preventing unnecessary workup while still flagging true outliers.

Prediction reference: G-2026-P018 (baseline validation for stromal), P019 (stem_adult), P020 (stem_pluri) — filed as open predictions for Section 4.3 cohorts not yet tested. Any multi-decade longitudinal healthy cohort with annual cfDNA sampling would test these trajectories directly.

5.4 Scenario: Pre-Diagnostic Window

Context:

A 56-year-old man enrolled in a routine screening program has cfDNA drawn at his annual physical. No symptoms, no family history, no imaging indication. Twenty months later he presents with abdominal pain and is diagnosed with stage III colorectal cancer. The research question: in retrospect, did the 20-month-prior cfDNA sample carry any signal?

This scenario is grounded in Mathios 2022 Nature Communications, which showed DELFI fragmentomics carries a detectable signal in cfDNA up to 2 years before clinical presentation in lung cancer. The framework extends this observation across all five substrates and all architecture classes.

Retrospective analysis — pre-diagnostic timepoints:

Timepoint	A_combined	Divergence	Class Signal	Retrospective Interpretation
T-24 mo	0.992	0.018	None detected	Within healthy range at age 54
T-18 mo	1.012	0.024	cycling class weak	MARGINAL elevation; would have warranted note but not action
T-12 mo	1.038	0.041	cycling emerging	MARGINAL; serial would have flagged for repeat
T-6 mo	1.067	0.058	cycling clear	DETECTABLE; would have prompted workup
Dx	1.124	0.072	cycling confirmed	FLOOR BREACH; diagnosis confirmed by colonoscopy

The framework's pre-diagnostic window in cycling-class cancers is approximately 12-18 months. By T-18 months A_combined has crossed MARGINAL with a cycling-class divergence signature. Standard single-timepoint analysis would likely dismiss the T-18 elevation as noise. Serial analysis — comparing T-18 to T-24 and noting the +0.020 movement — would flag the change as worthy of repeat. At T-12 months the trajectory is clear and a clinician would likely order colonoscopy. The patient would have been diagnosed approximately one year earlier, at stage I instead of stage III, with substantially better long-term outcomes.

This is the value proposition of population-scale serial cfDNA monitoring: not a single screening test, but a trajectory over time. A single patient's A at T-12 does not confirm cancer — but a patient whose A went from 0.99 at T-24 to 1.04 at T-12 has changed in a way that deserves follow-up. The framework does not require every patient to be in the DETECTABLE tier before action is appropriate; it allows trajectory-based action well before a single-timepoint threshold is crossed.

Prediction reference: Mathios et al. 2022 Nat Commun doi:10.1038/s41467-021-24994-w (DELFI pre-diagnostic lung cancer). Extension to colorectal and other cycling-class cancers filed as prediction G-2026-P021.

5.5 Scenario: Multi-Class Divergence — Metastasis Detection

Context:

A 62-year-old woman two years post-BRCA (breast cancer) treatment is in active surveillance. Her cycling-class and secretory-class A-scores have remained stable near baseline. At a routine 6-month cfDNA draw, her immune-class A-score has jumped from 0.985 (prior) to 1.042 (current), while cycling and secretory remain stable. No palpable nodes. No constitutional symptoms. Imaging has not yet been scheduled.

The framework signal: cross-class propagation.

A cancer arising in one architecture class can propagate signal to other classes by two mechanisms. First, direct metastatic seeding — a breast cancer (secretory class) that metastasizes to bone marrow produces an immune-class signal as the marrow microenvironment reacts. Second, inflammatory stromal response — any advancing malignancy can drive immune-class elevation independent of direct seeding. The framework does not distinguish these two mechanisms, but it does identify the multi-class pattern as distinct from single-class primary disease.

Three-class sampling — current timepoint:

Class	A_combined	Prior (6 mo)	Δ	Interpretation
cycling	0.991	0.989	+0.002	Stable — no recurrence signal in cycling class
secretory	0.996	0.994	+0.002	Stable — original tumor site quiet
immune	1.042	0.985	+0.057	NEW ELEVATION — cross-class signal emerging
stromal	0.978	0.975	+0.003	Stable — no fibroblast involvement
terminal	0.972	0.972	0.000	Stable — no CNS involvement

The framework flags the pattern: a post-BRCA patient whose immune class alone has elevated while her original secretory-class signal remains stable suggests cross-class propagation rather than local recurrence. The most common explanations are (1) bone-marrow or lymph-node metastasis driving immune reaction, or (2) non-oncologic inflammatory process. Either requires investigation. Standard next steps include bone marrow biopsy, PET imaging, and bloodwork for inflammatory markers. The framework has contributed the earliest possible window on this decision — a signal visible before either imaging or palpable disease.

The broader research use of cross-class propagation is in understanding metastatic biology itself. Which primary cancers produce which secondary-class signals? Does the immune-class response emerge before or after imaging-detectable disease? Does the kinetics of cross-class propagation predict outcome? These are open research questions that multi-class cfDNA monitoring can address directly.

Prediction reference: G-2026-P022 (cross-class propagation timing in metastatic BRCA). Any prospective post-treatment surveillance cohort with multi-class cfDNA monitoring and staged imaging confirmation would directly test this scenario.

SECTION 6

DATED PREDICTIONS — PRIORITY TREATMENT

The Master Predictions Table that follows consolidates every G-2026-P filing in abbreviated form. This section gives four priority predictions their own dedicated treatment. Each is the highest-leverage testable claim in its domain — falsifiability named explicitly, cohort candidates identified, required sample size estimated, success/failure criteria stated in advance. These four are filed publicly now so that any subsequent cohort study producing relevant data knows what the framework predicted before the data existed.

6.1 G-2026-P005: Cryptorchidism Surveillance Divergence Signature

CLASS: Pluripotent Stem | FILED: April 2026 | STATUS: PENDING

The claim:

In cryptorchidism patients and post-orchietomy TGCT survivors monitored prospectively with serial cfDNA sampling, the Pluripotent class will show a characteristic DIVERGENCE PATTERN in patients who later develop seminoma-lineage TGCT or contralateral second primary TGCT: A_methyl declines toward the 0.65-0.75 range while A_nucl, A_wps, A_fuzz, and A_frag simultaneously elevate to 1.01 or higher. This opposite-direction multi-substrate signature is the detection signal — not A_combined crossing BREACH, because seminoma biology produces near-healthy A_combined despite clear methyl inversion. Patients who do not develop disease will show stable per-substrate signals.

Physics basis:

Seminomas arise from PGC-like precursors that are globally hypomethylated. The Pluripotent H_min_methyl = 0.9822 sits near maximum entropy already, so a further β -toward-zero trajectory drives H(β) DOWN, producing A_methyl below the healthy reference. Meanwhile the four non-methylation substrates respond to tumor burden the way any cancer does — A elevates. The net result is the characteristic divergence signature: one substrate down, four up. See Section 2.4 for the full physics derivation.

Falsification criterion:

A prospective cohort of 200+ cryptorchidism patients followed for 10+ years with annual cfDNA collection. Among those who develop TGCT during follow-up, 80% or more should show the divergence signature (A_methyl declining while other substrates elevate) at the 6-12 month sample preceding diagnosis. If the predicted signature fails to appear in more than 40% of incident cases, the prediction is refuted.

Candidate validation cohorts:

The EUROPACE cryptorchidism registry (Scandinavian cohort, n~3,500), the Nordic TGCT Survivorship cohort, and the US DoD Serum Repository (which has archived serum from 10+ million service members and has been used retrospectively for TGCT biomarker studies). Any of these institutions conducting prospective cfDNA surveillance with seminoma outcome follow-up would directly test this prediction.

Primary-source references: Shen et al. 2018 Cell Reports (TCGA TGCT n=137); Killian et al. 2016 Genome Research (pure-histology TGCT n=130 with PGC comparison); Fossa et al. 2005 NEJM (contralateral TGCT risk).

6.2 G-2026-P013: CHIP / CCUS Progression to MDS — Pre-Clinical Window

CLASS: Adult Tissue Stem (HSC) | FILED: April 2026 | STATUS: PENDING

The claim:

In patients with clonal hematopoiesis of indeterminate potential (CHIP) or clonal cytopenias of undetermined significance (CCUS), the emergence of simultaneous A_methyl AND A_frag saturation in the stem_adult architecture class predicts progression to myelodysplastic syndrome (MDS) within 12-24 months with sensitivity > 70%. The two-substrate co-saturation is the Niche Depletion Inversion signature. It does not occur in stable CHIP that never progresses.

Physics basis:

HSC aging is a clonal-depletion process rather than a uniform-drift process. A few dominant clones expand while rare clones vanish — this produces saturation of both methylation and fragmentomic signals simultaneously because both substrates report the clonal-expansion state. The Niche Depletion inversion (Section 2.4) is the class-specific failure mode. Adelman 2019 Cancer Discovery HSC-enriched aging methylation data shows this signature cleanly in patients with documented MDS transformation. The transformation-marker nature of the co-saturation — rather than simple A_combined elevation — is what makes this a specific prediction rather than a general "CHIP progression" claim.

Falsification criterion:

A prospective CHIP/CCUS cohort of 500+ patients with annual cfDNA + complete blood count + bone marrow biopsy when indicated. Among patients progressing to MDS (expected ~0.5-1% per year), 70% or more should show the methyl+frag co-saturation signature at the sample 12-24 months before clinical MDS diagnosis. Among patients with stable CHIP (no progression over follow-up), co-saturation should appear in under 10%. If either threshold fails, the prediction is refuted.

Candidate validation cohorts:

The WHI CHIP cohort (n~4,000 with banked samples and documented hematologic outcomes), the Mass General Brigham CHIP Prospective Study, and the Cleveland Clinic CCUS Registry. Each has the longitudinal sampling and outcome data to test this prediction directly on archived samples.

Primary-source references: Adelman et al. 2019 Cancer Discov (HSC-enriched aging methylation); Jaiswal et al. 2014 NEJM (CHIP prevalence and progression); Malcovati et al. 2017 Blood (CCUS definition and natural history); Steensma et al. 2015 Blood (clonal hematopoiesis definition).

6.3 G-2026-P015: Adult Stem Class — Beyond-Ceiling Detection Mechanism

CLASS: Adult Tissue Stem | FILED: April 2026 | STATUS: PENDING

The claim:

In HSC-origin AML and Merkel cell carcinoma — both adult-stem-class malignancies where three of five substrates (nucl, fuzz, wps) structurally saturate below BREACH — diagnostic discrimination from healthy aging depends on the two remaining substrates (methyl, frag) being active past the BREACH threshold. The framework predicts $A_{\text{methyl}} \geq 1.14$ and $A_{\text{frag}} \geq 1.18$ in active AML and active MCC samples, with the structurally-saturated three substrates providing no additional discrimination. A classifier using only methyl + frag should achieve $AUC \geq 0.90$ for active disease versus healthy age-matched controls.

Physics basis:

The Adult Tissue Stem class has the tightest H_{min} ceilings in the framework: WPS ceiling 1.0112 (the tightest substrate×class pairing anywhere), fuzz ceiling 1.0196, nucl ceiling 1.0407 — all below the clinical BREACH threshold (1.10). No matter how severe the disease biology, these three substrates cannot report A above BREACH. The only substrates with ceiling headroom past BREACH are methyl (ceiling 1.1445) and frag (ceiling 1.1886). A detection protocol that expects all five substrates to elevate in unison will fail for this class — the physics forbids it. The two-substrate detection strategy is the structurally correct approach.

Falsification criterion:

A cross-sectional study of 100 active-AML + 100 active-MCC + 200 age-matched healthy controls with complete five-substrate cfDNA panels. The prediction is falsified if the methyl+frag two-substrate classifier AUC falls below 0.80, or if any of the three structurally-saturating substrates (nucl, fuzz, wps) shows observed A above its predicted ceiling by more than measurement noise. Both failure modes indicate the class-floor derivation for adult stem is wrong.

Candidate validation cohorts:

Existing TCGA AML cohort archived samples (Ley 2013 NEJM n=200) reanalyzed with the five-substrate cfDNA panel. The MCC cohort from Harms 2015 Cancer Research (n=49) similarly. Any new cohort of adult-stem-origin cancers with paired cfDNA collection would test the prediction directly.

Primary-source references: Ley et al. 2013 NEJM (TCGA AML n=200); Harms et al. 2015 Cancer Res (MCC MCPyV-negative n=49); Adelman et al. 2019 Cancer Discov (HSC aging baseline).

6.4 G-2026-P017: BEP Platinum Response Trajectory in TGCT

CLASS: Pluripotent Stem | FILED: April 2026 | STATUS: PENDING

The claim:

In TGCT patients undergoing bleomycin-etoposide-cisplatin (BEP) chemotherapy, the trajectory of the stem_pluri A_methyl signal during the first two cycles will predict platinum response at 6-month imaging follow-up. Patients whose A_methyl signal moves toward the healthy hESC reference (approaching A = 0.970 from either direction — either recovering from the seminoma inversion or declining from EC hypermethylation) will show RECIST response. Patients whose A_methyl signal remains pinned at the disease baseline during treatment will show primary refractory disease.

Physics basis:

Platinum chemotherapy response in TGCT corresponds to restoration of the pluripotent-state thermodynamic signature as tumor burden decreases. A responding tumor is being reduced in cellularity — the cfDNA signal shifts toward the healthy background as malignant clone contribution drops. Non-response corresponds to persistent epigenomic departure from the class floor: the tumor remains at the inversion (for seminoma) or at the elevated-CpH state (for EC) throughout treatment. TGCT is the ideal validation setting because its high cure rate (95% at stage I-II) and well-standardized BEP protocol make the response/non-response separation cleaner than in most solid tumors — a meaningful signal is expected to show clearly.

Falsification criterion:

A prospective cohort of 100 stage II-III TGCT patients undergoing standard BEP with cfDNA collection pre-C1, pre-C2, pre-C3, and pre-C4. RECIST response at 6-month imaging is the outcome. A classifier using the pre-C2 A_methyl trajectory (change from pre-C1) should predict RECIST response with sensitivity $\geq 80\%$ and specificity $\geq 70\%$. If sensitivity falls below 60% or specificity below 50%, the prediction is refuted.

Candidate validation cohorts:

The TIGER (Testicular Germ Cell Tumor) consortium, Memorial Sloan Kettering TGCT prospective cohort, MD Anderson TGCT program. TGCT is a relatively small-volume oncology specialty — 8,000-10,000 new US cases per year — so cohort assembly is feasible within a 2-3 year window at a handful of high-volume centers.

Primary-source references: Shen et al. 2018 Cell Reports (TCGA TGCT); Killian et al. 2016 Genome Research; Feldman et al. 2008 JCO (BEP standard protocol); Beyer et al. 2013 Ann Oncol (TGCT survivorship consensus).

SUBSTRATE CHARACTERIZATION

Replaces Issue 002 Sections 7–8. Every (substrate, class) pair has its own floor and its own healthy range. This section is the grid: what each substrate carries, which cells of the grid have been measured with the canonical files, which are absent by biology, which are structurally unreadable, and which are open. Each open cell names the experiment and the public data that would fill it. This is the hand-off to a researcher.

7.1 What each substrate carries

Substrate	Reads	Floor	Status 2026-09-19
whole blood / buffy coat	immune ARCHITECTURE of the leukocytes present; epithelial composition absent by biology	methyl	n=7 both channels; 3/3 healthy BELOW age band — OPEN
plasma cfDNA	SHED tissue composition	frag / wps / nucl — never methyl	composition measured n=33; gauge withdrawn (wrong floor)
tissue	positive control; heterogeneity marker (glioma-LL-002)	methyl	conformant, MAE 0.0004
urine sediment	shed urothelial / prostate	to be established	GSE119260 in hand, not run
CSF	shed CNS (terminal class)	to be established	GSE292312, GSE269403 located; supplementary files
stool	shed colonic epithelium	to be established	no public array data exists

7.2 The grid

	terminal	secretory	immune	progenitor	cycling	stromal	stem_adult	stem_pluri
whole blood	ABS	ABS	MEAS	MEAS	ABS	ABS	MEAS	SAT
plasma cfDNA	MEAS	—	COMP	—	—	ABS	—	—
tissue	—	CONF	CONF	STRUCT	CONF	STRUCT	—	—
urine	DATA	DATA	DATA	DATA	DATA	DATA	DATA	DATA
CSF	DATA	DATA	DATA	DATA	DATA	DATA	DATA	DATA
stool	NONE	NONE	NONE	NONE	NONE	NONE	NONE	NONE

CONF = reproduces the project answer key · MEAS = run on canonical files, result recorded · COMP = deconvolver composition only (gauge not applicable on the methyl floor) · ABS = fraction = 0 by biology · STRUCT = zero by marker coverage · SAT = methyl ceiling below BREACH · DATA = public cohort located, not run · NONE = no public data, primary collection required · — = not established

7.3 What was measured

Whole blood, n=7 (test-data IDATs, Stage-1 calibrated, both channels). Epithelial-class fraction 0.0000–0.0111 in all seven; immune 0.803–0.966; residual MAE 0.043–0.060. Two rheumatoid-arthritis cases sit inside the non-RA range on both channels (gauge 0.8743/0.8789 vs 0.8066–0.8924; immune fraction 0.8395/0.8686 vs 0.8027–0.9663). **The three age-labelled healthy donors read BELOW the age-matched band:** 43M A=0.8739 (band 0.9021–0.9934, z –2.07); 58M 0.8066 (0.9083–1.0029, z –4.03); 67F 0.8326 (0.9166–1.0138, z –3.49). Either the Stage-1 path and the literature-compiled reference matrix are on different scales, or these donors are below band. n=3 cannot separate the two. This is PROC-WB-IMMUNE-01's open verdict.

GSM	label	epithelial	immune	progenitor	resid MAE	immune gauge A	age band
GSM2333901	healthy 58M	0.0000	0.8336	0.1280	0.0562	0.8066	0.908–1.003 → BELOW
GSM2333905	healthy 67F	0.0051	0.8881	0.0671	0.0522	0.8326	0.917–1.014 → BELOW
GSM2333950	healthy 43M	0.0111	0.9663	0.0000	0.0596	0.8739	0.902–0.993 → BELOW
GSM1051533	RA-study control	0.0000	0.9318	0.0515	0.0434	0.8924	age not in GEO

GSM	label	epithelial	immune	progenitor	resid MAE	immune gauge A	age band
GSM1051534	RA-study control	0.0000	0.8027	0.1740	0.0504	0.8673	age not in GEO
GSM1051525	RA case	0.0000	0.8395	0.1163	0.0452	0.8789	age not in GEO
GSM1051526	RA case	0.0000	0.8686	0.1147	0.0459	0.8743	age not in GEO

Plasma cfDNA, GSE122126 (Moss 2018), n=33 plus 22 sepsis — composition only. All plasma GAUGE A-scores computed 2026-09-19 were on the METHYL floor and are therefore out of spec (RECON S1). Only the COMPOSITION results are reported here. The GSE122126 in-vitro mixes (9+5 samples of known genomic-DNA proportion) are the available real-data ground truth for the deconvolver and have not yet been scored.

condition	n	epithelial fraction (cycling+secretory+terminal+stromal)	residual MAE
healthy	4	0.030 (mean)	0.058
colon cancer	4	0.463 (mean)	0.155
breast cancer	3	0.466 (mean)	0.161
lung cancer	4	0.010 (mean)	0.055
CUP	4	0.293 (mean)	0.116
sepsis	22	median 0.040, p75 0.233, max 0.520	—

Lung cancer reads epithelial-null in the correct substrate: nothing shed, nothing to read, instrument says so. Sepsis sheds epithelium in a minority of samples (multi-organ injury, established biology). Residual MAE tracks epithelial fraction (r = 0.98 across all samples) — a second readout the chain computes and discards; not yet checked against the corpus. **Cell-level routing:** Cell-level calls on shed epithelium route to gastric references (Fundus_diff, Antrum_diff) in colon tissue, colon plasma and breast plasma alike; Colon_epithelial_cells and Breast return 0.0000 despite being in the atlas. LESSON-DECONV-01 attributes this to reference mismatch, not solver conditioning; hierarchical refinement was tested, invents cells, and does not ship. Cell tier is INDICATIVE. No organ-level claim is supported by either channel today.

Tissue — conformance, PROC-DECON-01 (§8). Whole-blood versus tissue integrity gate (CHK-3.1) on Stage-1-calibrated data:

GSM	substrate	% β <0.05 or >0.95	% β in 0.40–0.60
GSM2333901	whole blood 450K	33.8	4.5
GSM1051533	whole blood 450K	27.3	5.0
GSM8772491	colon adenoma EPIC tissue	6.7	5.5
GSM8772492	colon adenoma EPIC tissue	8.5	6.4
GSM5065990	CRC EPIC tissue	23.6	11.7

The >30% extreme-β criterion written for whole blood fails every properly calibrated tissue sample. Integrity thresholds must be stated per substrate (§9, CCL-032 amended).

7.4 The researcher hand-off — one experiment per open row

substrate	prediction to test	data	who can run it
whole blood × immune	healthy donors read IN the age band once Stage-1 and the reference are on one scale; if not, the reference is re-anchored to the atlas scale (open VAL-053)	test-data IDATs + any public healthy blood 450K/EPIC	any methylation lab; no clinical access needed
plasma cfDNA	epithelial classes appear in cancer plasma on frag/wps/nucl floors and stay absent in healthy; in-vitro mixes recover at MAE ≤ 0.01	GSE122126 (has 14 known mixes); GSE311578 (70 AD plasma)	cfDNA group with WGS for the non-methyl substrates
urine	prostate signal appears in the secretory class of urine sediment in the same men whose tissue shows it; plasma of the same men reads lower	GSE119260: 4 pts × tissue/plasma/urine	urology or GU-oncology group
CSF	terminal-class signal present in CSF and absent from the same patient's blood	GSE292312 (24 CSF + 157 tumour); GSE269403 (39 CSF + 17 paired blood)	neuro-oncology group

substrate	prediction to test	data	who can run it
stool	colonic cycling-class signal in stool DNA	none public — primary collection	GI group with a stool-DNA protocol

SECTION 8

PROCEDURES

One entry per verified cell of the grid, in the form a stranger can run: input, operation, expected, tolerance, observed, verdict. Nothing else. Every entry names the exact files at repo commit 66f37fe. An entry whose verdict is OPEN is printed as OPEN. Entries not yet written are listed at the end with what they need. The template is PROC-DECON-01, the first procedure in this project to pass on a machine that had never seen the author's setup.

PROC-DECON-01 — Deconvolver conformance — tissue, EPIC

field	value
input	betas_cache.pkl (Stage-1 noob-calibrated β , TEST_DATA); IAMAtlasREBUILD.csv decompressed from repo IAM_Atlas/IAMAtlasREBUILD.csv.xz; IAMAtlasREBUILD_celltype_to_class.json
operation	WaltherIAMDeconvolver(atlas, celltype_class_map=map).deconvolve(beta_dict) for GSM8772491, GSM5065990, GSM5065985
expected	class fractions in TEST_DATA_MANIFEST.md: GSM8772491 cycling .354 immune .253 stem_pluri .161 secretory .122 terminal .112
tolerance	class-fraction MAE ≤ 0.001 over documented classes
observed 2026-09-19	GSM8772491: cycling .3535 immune .2530 stem_pluri .1605 secretory .1215 terminal .1115 — MAE 0.0004, max err 0.0005. GSM5065990 MAE 0.0002. GSM5065985 MAE 0.0002. Status OK all three; 7,114 class markers, 4,000 cell markers.
verdict	PASS. Deltas are the manifest's 3-decimal rounding.
scope	verifies deconvolver + atlas + map. Does NOT verify Stage 1 (cached β used, raw IDATs not decoded). See PROC-CAL-01.

PROC-CAL-01 — Stage 1 - raw IDAT to calibrated beta, executed

field	value
input	the eleven raw IDAT pairs in 10_TEST_DATA/idats (7 x 450K whole blood incl. 2 RA + 2 RA-study controls; 4 x EPIC colorectal tissue); stage_1_idat_calibration.py from the repo; methylprep 1.7.1 with numpy 1.26.4 / pandas 1.5.3 (methylprep uses DataFrame.append, removed in pandas 2 - the conflict the module's own docstring warns about); Illumina 450K and EPIC manifests fetched from methylprep's public bucket
operation	calibrate_idat_to_beta(Grn, Red) per sample: array type auto-detected from bead-address count, noob dye-bias + probe-type normalisation, per-sample, no cohort information; compare the resulting beta vector CpG-by-CpG to betas_cache.pkl (the Stage-1 output that produced PROC-DECON-01)
expected	identical betas (same algorithm, same manifest); any difference is a pipeline-version effect
verdict	PASS - 11/11 bit-identical: $r = 1.000000$, max diff = 0.000000, zero CpGs differing by more than $1e-4$, on both array types (450K 415,863-418,055 CpGs, EPIC 490,390-757,003 CpGs); 24-48 s per sample. With PROC-DECON-01 and PROC-ANCHOR-01 this closes the chain: raw IDAT -> calibrated beta -> class fractions -> sealed per-cell anchor, every link reproduced from the repository on a machine that had never seen the project.
consequence	the 3/3 below-band healthy whole-blood donors in PROC-WB-IMMUNE-01 are NOT a Stage-1 artifact of this machine - the betas are the project's own. The remaining hypotheses are (a) the age band was compiled on a different IDAT-to-beta pipeline than noob (the Xu-538 pipeline-relativity effect, 0.38-0.62), or (b) the donors are below band. (a) is testable by locating the pipeline the band was compiled from.

PROC-N7-01 — End-to-end synthetic simulation - the gauge as wired reads the wrong CpG set

field	value
input	synthetic_patient_generator.py (restored from RETIRED; patched to read IAM_Atlas/IAMAtlasREBUILD.csv; composition_alpha added): 16 healthy + 8 case, whole-blood composition (~immune 0.89, progenitor 0.08, stem_adult 0.03), disease panel 500 CpGs on the immune identity loci, signal 2.0. Each patient is a linear mix of the Atlas class posterior means plus age/sex/batch loadings and noise 0.03. Run through cpq_conductor.run_full.
R1 composition	PASS - deconvolver recovers every class: immune MAE 0.0095, progenitor 0.0145, stem_adult 0.0116, epithelial classes <= 0.0016 (r 0.79-0.98 on the legacy mix).
R2 gauge	FAIL - every synthetic patient, healthy or case, reads immune A 1.12-1.13 ABOVE_BAND / BREACH; case and healthy do not separate.
root cause	cpq_conductor.stage_b_classes computes $H(\beta_{\text{mean}})/H_{\text{min}}$ over the class MARKER UNION (2,952 bimodal CpGs for immune; comment: 'PRODUCTION A-score (GAPE_WEB_v13 + Reproduction Paper v3) ... NOT identity loci'), while its own docstring, SOP s41/s106 and Issue 002 say identity loci (42,134 unimodal CpGs). Confirmed: ge.read on the marker-union β_{mean} returns 1.1303, the conductor's number, to four decimals. The marker union's β_{mean} depends on the SHAPE of the input (Jensen gap 0.257 vs 0.026 on identity loci): a smooth synthetic mixture pulls it to 0.63 (BREACH); real bimodal blood lands at 0.75 (NORMAL).
synthetic healthy (pure mixture of healthy posteriors)	identity 0.9875 marker-union 1.1303 BREACH
real WB healthy 43M	identity 0.874 marker-union 0.957 IN_BAND
real WB healthy 58M	identity 0.807 marker-union 0.960 IN_BAND
real WB RA	identity 0.879 marker-union 1.002 IN_BAND
real tissue adenoma GSM8772492	identity 1.096 (elevated) marker-union 0.932 IN_BAND (missed)
consequence 1	RULING A3 was right about the formula ($H(\beta_{\text{mean}})$ on identity loci) and WRONG in asserting that the wired chain computes it. Corrected in s1.5.
consequence 2	age_reference_matrix.json was compiled (GAPE_WEB_v13 _AGE_REFERENCE) on the marker-union statistic. PROC-CHAIN-01's '7/7 healthy IN_BAND' is the marker-union gauge agreeing with the marker-union band - self-consistency, not conformance. Withdrawn as evidence of a working gauge.
consequence 3	s6a's 'below band' reading compared identity-loci A against a marker-union band: a statistic mismatch, not an offset. The +0.05-0.09 beta shift at identity loci vs the Atlas posterior IS real and remains Phase 1's business.
consequence 4	The identity-loci gauge is the correct statistic - unimodal, passes N7, carries the adenoma signal - and has NO BAND. Phase 1 (rebuild the band on identity-loci $H(\beta_{\text{mean}})$ from GSE87571 through Stage 1) is now a prerequisite for any gauge reading, not a refinement.
consequence 5	The conductor is NOT switched today (a statistic without its band would read every healthy patient below band). Every class-gauge reading now carries gauge_surface = 'marker_union' and n_cpqs so it cannot be mistaken.
verdict	N7 did its job: the synthetic healthy patient exposed a gauge that real blood had been masking. Composition PASS; gauge as wired FAIL; identity-loci gauge PASS on N7 and blocked on its band.

PROC-NILC-01 — The retired second deconvolver, rerun as designed

field	value
input	Walther's own class-marker reference X (1,256 CpGs x 8 classes, built by the engine at load); the 11 Stage-1 test samples; NILC's stated algorithm (unconstrained least squares on X, then simplex projection) run on the SAME X for a fair comparison
conditioning	condition number of X: 47.0 (all 8); 30.6 for the blood sub-problem (immune, progenitor, stem_adult); 8.8 for the epithelial trio. Centred column correlations: progenitor-stem_adult +0.99, immune-progenitor +0.94, immune-stem_adult +0.93.
result	per-patient L1 disagreement 0.55-0.65 on every blood sample (design target < 0.05); NILC places 0.26-0.34 in stem_adult where Walther places 0.005-0.044, and deletes secretory and stromal; on tissue it drives progenitor negative (-0.16 to -0.23). Exactly the 'collapsed on correlated blood mixtures, deleted correct calls' of the cut.
reading	By NILC's own design rule, divergence marks where composition is 'genuinely ill-defined by the atlas reference'. It was: in blood the Atlas cannot separately determine immune, progenitor and stem_adult; Walther's 2-4% stem_adult is chosen by the simplex constraint, not by the data. That is the class that read false BREACH in 6/7 healthy donors (PROC-CHAIN-01). The second deconvolver diagnosed the stem_adult problem in June; it was cut for delivering the message.
recommendation	(1) Whole blood: report progenitor + stem_adult as ONE haematopoietic-progenitor component until the Atlas separates them (kappa of the blood sub-problem must fall below ~10 first). (2) Pre-flight: compute the substrate sub-problem's condition number alongside the Jensen gap. (3) A second deconvolver stays a goal, but its disagreement is a diagnostic to be read, not a gate to be failed.

PROC-WB-IMMUNE-01 — Whole blood × immune — healthy reads in band

field	value
input	betas_cache.pkl for GSM2333901 (58M), GSM2333905 (67F), GSM2333950 (43M), GSM1051533, GSM1051534 (RA-study controls), GSM1051525, GSM1051526 (RA); iamatlas_gauge_identity_loci_v1_0.json; age_reference_matrix.json; IAMAtlasREBUILD.csv + map
operation	(1) deconvolve → require immune fraction ≥ DETECT_FLOOR (it is: 0.80–0.97); (2) β_{mean} over the 42,134 immune identity loci (36,290 present on 450K); (3) $A = H(\beta_{\text{mean}})/0.838889$; (4) place against the immune p10–p90 band at the donor's age; (5) NO input offset (retired)
expected	healthy donors IN BAND; RA cases not required to differ (chronic autoimmune; no prior)
tolerance	IN BAND = p10 ≤ A ≤ p90 of age_reference_matrix immune at nearest decade midpoint
observed 2026-09-19	43M A=0.8739 vs [0.9021, 0.9934] z –2.07 BELOW · 58M 0.8066 vs [0.9083, 1.0029] z –4.03 BELOW · 67F 0.8326 vs [0.9166, 1.0138] z –3.49 BELOW · RA 0.8789/0.8743 and controls 0.8924/0.8673 (ages not in GEO). Composition: epithelial 0.0000–0.0111 all seven.
verdict	OPEN. 3/3 healthy BELOW BAND. Two hypotheses, not separable at n=3: (a) age_reference_matrix (compiled from Hannum/Horvath/Alisch literature β) and the Stage-1 noob path are on different scales — the ~0.1 gap is in the direction the retired 0.055 offset addressed; (b) the donors are below band. Next: run ≥30 public healthy whole-blood 450K/EPIC through the same path; if the median sits ~0.1 below p50 at every decade, (a) is confirmed and the reference is re-anchored to the atlas scale (VAL-053).
do not	compute a group mean and call it the result; the reading is per donor against the band

PROC-WB-COMP-01 — Whole blood — composition is immune, epithelium absent

field	value
input	as PROC-WB-IMMUNE-01, deconvolver only
expected	immune + progenitor + stem_adult ≥ 0.95; epithelial (cycling+secretory+terminal+stromal) ≤ 0.02; residual MAE below tissue (0.108–0.153)
observed 2026-09-19	immune 0.8027–0.9663, progenitor 0–0.174, epithelial 0.0000–0.0111, residual 0.0434–0.0596 — all seven
verdict	PASS on n=7 including two RA cases. This is the substrate specification (README 'Key validated facts'), not a limitation. It says nothing about the immune-ARCHITECTURE channel, which is PROC-WB-IMMUNE-01.

PROC-PLASMA-COMP-01 — Plasma cfDNA — shed epithelium by condition

field	value
input	GSE122126-GPL21145 series matrix (Moss 2018, EPIC): healthy 4, colon 4, breast 3, lung 4, CUP 4, sepsis 22, in-vitro mixes 14; author-processed β (no Stage 1)
operation	deconvolve each sample; epithelial = cycling+secretory+terminal+stromal
expected	healthy ≤ 0.05 ; solid-tumour plasma $>$ healthy; in-vitro mixes recover declared proportions
observed 2026-09-19	healthy 0.030 · colon 0.463 · breast 0.466 · lung 0.010 · CUP 0.293 · sepsis median 0.040 (p75 0.233, max 0.520). Residual MAE tracks epithelial fraction $r=0.98$. Mixes NOT yet scored.
verdict	PASS on direction for colon/breast; lung null recorded; sepsis sheds in a minority. NOT a gauge result — any A on this cohort must use frag/wps/nucl floors (RECON S1), which array data cannot supply. Cell-level tissue-of-origin FAILS: all epithelium $\rightarrow$ gastric references.
caveat	author-processed GEO β without Stage 1; n per cancer 3–4

PROC-FORMULA-01 — A-score aggregation - SOP v1.4.0 (mean of per-CpG H) vs code at HEAD (H of mean beta)

field	value
input	betas_cache.pkl (7 Stage-1 whole blood incl. 2 RA; 4 EPIC colorectal tissue); iamatlas_gauge_identity_loci_v1_0.json; iamatlas_celltype_markers_v0_2.json aggregated to class unions; age_reference_matrix.json
operation	per class {immune, cycling, secretory}, per sample: (a) $H(\beta_{\text{mean}})/H_{\text{min}}$ over identity loci [code at HEAD]; (b) $\text{mean}_i H(\beta_i)/H_{\text{min}}$ over identity loci [SOP v1.4.0 formula on the HEAD loci]; (c),(d) the same two over the discriminative class union. Then Spearman(a,b) and the offset a-b, separately for whole blood and tissue. Then check which aggregation age_reference_matrix.json was compiled with.
expected	if the two aggregations are equivalent on the declared substrate, rank agreement ~ 1 and a constant offset; if s105's mechanism is real, the offset should grow on mixed or absent-class panels; the band must match one of them exactly
observed 2026-09-19	WHOLE BLOOD immune: Spearman +1.000, offset +0.029 $\pm$ 0.002. WHOLE BLOOD cycling/secretory (absent classes): offset +0.16. TISSUE immune: offset +0.24 $\pm$ 0.07. age_reference_matrix A_mean = $H(\beta_{\text{mean}})/H_{\text{min}}$ to 5 decimals in 80/80 cells. Discriminative mean-of-H tracks presence (cycling WB 0.35-0.41 vs tissue 0.57-0.62).
verdict	OPEN - author decision required. On the declared substrate with the class present, formula = constant offset; both valid IF band and runtime agree, and the band is $H(\beta_{\text{mean}})$. On mixed/absent panels $H(\beta_{\text{mean}})$ inflates as s105 says - never read it as architecture (presence gate; glioma-LL-002). The current SOP and the running code disagree and one must be amended: keep $H(\beta_{\text{mean}})$ and scope s105's NEVER, or adopt mean-of-H and recompile the band.
retraction	the first draft of this PROC (same day) read the tissue $H(\beta_{\text{mean}})$ values as a disease ordering; that was composition inflation and is withdrawn (s10).
caveat	n = 11, two substrates; RA-study samples read +0.04 on discriminative $H(\beta_{\text{mean}})$ for cases and controls alike - a GEO cohort offset (CCL-004), not disease

PROC-PLASMA-MIX-01 — Deconvolver against real known mixtures (Moss 2018 Table 6 / Table 8)

field	value
input	GSE122126 GPL21145: in_vitro_mix_9..17 (9 genomic-DNA mixes: leukocytes 85-96% + hepatocytes/lung/neurons/colon 3.5-10%, Supp Data 1 Table 6) and cfdna_mix_1,3,5,6,7 (colon-cancer cfdna CC2 at 0-10% into healthy cfdna, Table 8). Author-processed β .
operation	deconvolve; compare declared spike per tissue to the class the atlas assigns that tissue (Hepatocytes→secretory, Cortical_neurons→terminal, Colon_epithelial_cells→cycling). GEO title→Mix mapping by order, corroborated by max-terminal landing on Mix2 (10% neurons) and max-cycling on Mix4 (10% colon).
expected	declared non-leukocyte fraction recovered (slope ~1, bias ~0); each spiked tissue rises in its own class
tolerance	proposed: per-class MAE ≤ 0.02 against declared; $r \geq 0.9$ per spiked tissue
observed 2026-09-19	TOTAL non-leukocyte: $r = 0.945$, MAE 0.0677, bias +0.0677 (over-reads by ~7 points in every mix). NEURONS→terminal $r = +0.945$ (10% → 6.5%, under-read, ordering correct). HEPATOCYTES→secretory $r = +0.192$, $p = 0.62$ (10% → 1.7%). COLON→cycling $r = +0.101$, $p = 0.80$ (cycling reads 5–10% whether or not colon is present). Cancer-cfdna spike (n=5): epithelial $r = +0.725$, $p = 0.17$ — monotone but not interpretable as recovery (CC2's own load unknown).
verdict	SPLIT. Terminal class recovers a real neuronal spike (PASS on direction, under-reads magnitude). Secretory does NOT register hepatocyte DNA and cycling does NOT register colon DNA — both FAIL on the first real ground truth the deconvolver has faced. Consistent with §7.3 cell routing (all epithelium → gastric) and LESSON-DECONV-01 (reference mismatch): the atlas profiles for liver and colon do not match real liver and colon DNA closely enough to earn weight. This is atlas work, not solver work.
consequence	the plasma colon/breast epithelial fractions of PROC-PLASMA-COMP-01 (0.46) are real departures from healthy (0.03) but their class assignment is not trustworthy; the tissue-of-origin claim is withdrawn until the atlas recovers Table 6.

PROC-ANCHOR-01 — Sealed GSE51032 per-cell anchor, reproduced from raw GEO betas

field	value
input	GSE51032 series matrix (3,009 MB, EPIC, 845 samples) filtered to the 460 in the sealed foundation-cohort CSV (36 breast >10 yr pre-dx + 424 HC); iamatlas_celltype_markers_v0_2.json at repo HEAD and the TRIAL (chrX-removed) copy; H_min per class from iamatlas_gauge_identity_loci_v1_0.json
operation	for each of 115 cell types: per-cell A under (a) $\text{mean}_i H(\text{beta}_i)/H_{\text{min}}$ and (b) $H(\text{beta}_{\text{mean}})/H_{\text{min}}$ over that cell's ~100 discriminative markers; Pearson r and max diff against GSE51032_115celltype_ascores.csv (sealed 2026-05-29)
expected	r = 1.0000, maxdiff = 0.0000 against the sealed CSV, as SOP v1.4.0 s105 reports for the vault module
observed HEAD markers, mean-of-H	r = 1.00000; max diff 0.00004; 112/115 cells r > 0.999 - REPRODUCES THE SEAL
observed HEAD markers, H(beta_mean)	r = 0.450; max diff 1.048; 0/115 - does not
observed TRIAL markers, mean-of-H	r = 0.99957; max diff 0.060; 89/115 - 32 cells shift by up to 0.059 - the chrX removal
observed TRIAL markers, H(beta_mean)	r = 0.449; max diff 1.048; 0/115 - does not
observed GSE51057 (188: 11 case + 177 HC), HEAD markers, mean-of-H	r = 1.00000; max diff 0.00004; 115/115 cells r > 0.999 - REPRODUCES THE SEAL - both halves of the 648-sample foundation cohort
also observed	healthy per-cell A median 0.520 on this surface (s105 says ~0.52); per-cell case-vs-HC d median +0.32, 37/115 cells d > 0.5, largest Glia +1.10, stem_pluri +1.07, Baso +1.01. The seal was made with the REPO-HEAD markers (pre chrX removal): the TRIAL file changes 32 cells by up to 0.059 (largest: Mela) and would need re-sealing.
verdict	PASS for the separation surface: the sealed per-cell anchor is exactly mean-of-H over the v0_2 discriminative markers as SOP v1.4.0 s105 states, reproduced from raw GEO on a stranger's machine. This does NOT adjudicate the gauge (identity loci) - see PROC-FORMULA-01. It does establish that the current repo markers file, not the chrX-removed one, is what the seal rests on; adopting the chrX file re-opens the seal.

PROC-QC-01 — Stage 0-1 intake and calibration gates (SOP v1.4.0 s11-25) - specified; PROC-CAL-01 executes it

field	value
Manifest	raw IDAT pair keyed to demographics and array substrate before any processing (the Planck 'scan-ring' association); SHA-256 on every file
Detection p-value	>5% probes failed = HARD reject; 1-5% = SOFT flag (s15)
Bead count	>=3 per probe (Illumina default); >=5 for high-stakes (s16)
Call rate	>=0.98 for Stage 0 PROCEED (s19); <0.95 after QC = quarantine
Sex check	predicted vs declared sex; mismatch = sample swap or rare karyotype - halt (s18; Planck: pointing-direction sanity check)
Dye bias	noob (s20; Planck: bandpass correction)
Probe-type normalisation	funnorm default (s21)
Batch	ComBat requires >=3 samples per batch (s22)
Bisulfite efficiency	<0.95 HARD quarantine; 0.95-0.98 SOFT (s23); low BS efficiency + low intensity = degraded DNA
Beta sanity	range, bimodality coefficient 0.45-0.55 SOFT flag, cohort FLAG rate >5% = HARD halt (s25)
verdict	SPECIFIED, NOT RUN - every gate above is a pass/fail with a threshold; PROC-CAL-01 will record each on the seven test-data IDATs

Not yet written — and what each needs

PROC	needs
PROC-WB-PROGENITOR-01 / STEM_ADULT-01	the CHANGELOG's suspected breach artifact in peripheral blood; requires the presence-gated class ranking fix (design decision on threshold pending)
PROC-TISSUE-N-COMP-01	any tissue claim must report immune-class A beside the target class (glioma-LL-002); write once CPG-NEW-001 is re-run presence-gated
PROC-URINE-01, PROC-CSF-01	GSE119260 in hand; CSF cohorts need supplementary-file fetch; floors for these substrates not yet established
PROC-N5-01	plate/sentrix position null from raw IDAT filenames for GSE51032/GSE51057 — first thing a methylation referee asks for on the breast pre-dx result

SECTION 9

OPERATING RULES

The lessons that govern how the instrument may be read, promoted from appendix to rules. Each carries its date, source, and supersession status. A rule that is broken produces a wrong number, not a style complaint; several of these were broken in the runs of 2026-09-19 and the resulting wrong numbers are in §10.

ID	Rule	Content	Source	Status
L-5	Physics measures, cohorts only point.	The A-score is intrinsic and self-calibrating. Cohort data enters only as a DIRECTION, never as a baseline. The moment a cohort mean/SD becomes the yardstick the model works only in its own cohort. Primary readout = absolute A per sample against the fixed reference and the age band; group statistics are supporting detail.	LESSONS_collapse_and_KISS_v1	current
CCL-019/020	Direction depends on (class, compartment), not disease.	Specimen, class and panel are three independent dimensions and all three must match the clinical question. Discovered when VAL-061 predicted the wrong sign.	LESSONS_LEARNED.md	current
CCL-039	Marker-tile and full-genome A are two observables.	They move in opposite directions in tumour vs adjacent normal (READ, COAD, COAD-VAL-099). Never compare one to the other.	LESSONS_LEARNED.md	current
CCL-032 / CHK-3.1	Extreme-beta integrity gate is substrate-dependent.	Whole blood 450K reads 26-34% extreme; Stage-1-calibrated EPIC tissue reads 7-24%. The >30% criterion does not transfer to tissue. Thresholds must be stated per substrate.	LESSONS_LEARNED.md; run 2026-09-19	current, amended
LESSON-DE CONV-01	Global NNLS is correct; sparsity is reference mismatch.	Synthetic mixtures recover at MAE 0.000; real blood fits 2.6x worse. Hierarchical within-class refinement invents cells (MAE 0.271 on pure monocytes) and does not ship. Richer per-cell resolution is atlas work, not a solver swap.	SOP v1.3.3 §0.9 (2026-06-21)	current
2026-06-11	Never feed the gauge discriminative markers.	They are bimodal; their mean beta collapses toward 0.5 and the gauge pins the ceiling (all-BREACH).	cpg_gauge_engine.py	current
PRESENCE	A class not in the sample has no A-score.	Presence comes from the deconvolved fraction (DETECT_FLOOR). Scoring an absent class's identity loci produced spurious BREACH calls on stromal (fraction 0.0000) in 4/4 healthy plasma samples, 2026-09-19.	cpg_conductor.py; run 2026-09-19	current
SUBSTRATE	Every report states its substrate; every (class, substrate) has its own floor.	Whole blood = immune-architecture readout; plasma cfDNA = shed-tissue readout on frag/wps/nucl floors; tissue = positive control.	cpg_gauge_engine.py; TEST_DATA_MANIFEST.md	current
glioma-LL-002	A high A on bulk tissue is a heterogeneity marker, not a tumour marker.	Composition alone can move beta-bar. N-comp (report immune-class A alongside the target class) is the binding null for any tissue claim.	LESSONS_LEARNED.md	current
PL-002	The honest record is the best defense.	Amendments after data are observed are labelled as such; outcomes are allowed to be RESTATE, FAIL or VOID.	LESSONS_LEARNED.md	current
PIPELINE	One IDAT-to-beta pipeline per patient; absolute A is pipeline-relative.	Same Xu-538 panel, different IDAT-to-beta pipelines: absolute A baselines 0.38-0.62 across cohorts (within-EPIC-Italy −0.44), direction preserved within pipeline (April 2026). Absolute A is pipeline-relative; this is the leading hypothesis for the 3/3 below-band healthy whole-blood donors (PROC-WB-IMMUNE-01).	GAPE_EDEAR_Reproduction_Paper_v3 (T15, Rule 1-2)	current
FIREWALL	Production subtracts no age/sex/smoking foreground.	SOP s104: such change is the signal itself; age enters via the reference band, sex via chrX marker removal; subtraction modules are test-only and must not be loaded in production config.	SOP v1.4.0 s104	current
SEAL	A measurement that has not beaten the null suite is a candidate, not a claim.	N1-N8; SEALED/RESTATE/RETRACT; PREREG edited after data = auto-RETRACT.	SOP v1.4.0 s80-91	current

ID	Rule	Content	Source	Status
SURFACE	One aggregation per surface, fixed by its reference.	Gauge (identity loci): $H(\text{beta_mean})/H_{\text{min}}$. Separation (discriminative markers): $\text{mean_i } H(\text{beta_i})/H_{\text{min}}$. Never read a statistic against a band compiled the other way.	RULING A3	current
L-10	Reporting rule (extends the s3.2 presence gate)	a class gauge is REPORTED on a substrate only where the class is DETERMINED (κ of the substrate sub-problem $< \sim 10$, or reported jointly with what it cannot be split from) and PRESENT (typical fraction $\geq 5\%$); otherwise composition-only	whole blood: immune reported; progenitor + stem_adult jointly as one haematopoietic-progenitor component; the other five composition-only. Decided from κ and fraction BEFORE any A is computed	author, 2026-09-19, after PROC-CHAIN-01 / PROC-NILC-01; SOP s108

SECTION 10

FALSIFICATION RECORD

Every outcome that did not go the framework's way, kept on the record. A framework that cannot show its failures cannot show that its passes mean anything. Historical VALs are quoted from their sealed OUTCOME files; the 2026-09-19 rows are from this session and include the author's own failed predictions.

Test	What happened	Outcome
VAL-004	bimodality direction reversed vs prereg	RESTATE
VAL-006	chr6 signal did not survive look-elsewhere correction	FAIL
VAL-061	predicted wrong sign; produced CCL-019/020	RESTATE
VAL-102	voided four minutes after sealing for post-hoc accommodation; original seal preserved	VOID
VAL-128	failed opposite to prereg	FAIL
CPG-NEW-001 P1	GSE48684 adenoma read above carcinoma; monotonic progression failed	FAIL
CPG-NEW-001 P4	serrated vs conventional reversed sign between Spain and Finland	FAIL
CPG-NEW-001 N-comp	cycling effect not larger than immune effect on adenoma contrasts	composition-ambiguous
2026-09-19 two-channel	immune-A x epithelial-fraction did not beat epithelial alone (AUC .955 vs .961; LOO .793 vs .828)	prediction FAILED
2026-09-19 WB x immune	3/3 healthy whole-blood donors read BELOW the age band ($z = -2.07, -4.03, -3.49$)	OPEN - scale or biology
PROC-PLASMA-MIX-01 secretory	10% hepatocyte genomic DNA in leukocytes reads 1.7% secretory (Moss Table 6 ground truth)	FAIL
PROC-PLASMA-MIX-01 cycling	colon spike 0-10% not recovered; cycling reads 5-10% regardless ($r=0.10$)	FAIL
Analyst verdict 2026-09-19 (first draft of PROC-FORMULA-01)	'identity-loci $H(\beta_{\text{mean}})$ carries the disease ordering healthy < adenoma < carcinoma' - the tissue ordering is composition inflation (+0.24 Jensen offset on a mixture), not architecture; retracted the same day when the whole-blood Spearman came back +1.000 for both formulas	RETRACTED
SOP v1.4.0 s105 universal claim	'the A-score is NEVER $H(\beta_{\text{mean}})$ ' - on the declared substrate with the class present the two aggregations are rank-identical with a constant offset, and the shipped age band is compiled as $H(\beta_{\text{mean}})$; the NEVER is correct for mixed/absent panels only	OVER-SCOPED
Issue 003 draft v3-v8 s5A.1 claim	' $M_{\text{cell}} = n_{\text{bio}}/\ln 2 = 30.21$; AI transmon $M = 1$ exactly; one statistic across three substrates' - computed with a $\ln 2$ denominator the author does not use; withdrawn	WITHDRAWN
PROC-CHAIN-01 'healthy blood IN_BAND 7/7 = the shipped chain absorbs the offset'	the conductor's gauge and its band are both marker-union statistics; agreement between them is self-consistency, not validation (PROC-N7-01)	WITHDRAWN
RULING A3 'the wired chain computes $H(\beta_{\text{mean}})$ on identity loci'	the wired chain computes it on the marker union (cpg_conductor.stage_b_classes, comment citing GAPE_WEB_v13)	CORRECTED
s6a 'healthy whole blood reads below band'	identity-loci A compared against a marker-union band - statistic mismatch; the beta shift at identity loci is real, the band conclusion was not	RESTATED

Also on the record from this session, as method failures rather than results: the same analyst computed group effect sizes where the framework requires per-sample absolute readings (three times); scored classes the deconvolver had not found in the sample; scored plasma on the methyl floor; and read the pre-atlas SOP as the specification of the post-atlas engine. Each is now a rule in §9 or a row in §1. That is what the rules are for.

ENGINE MAP — WHAT THE RUNNING CHAIN CONTAINS, AND WHAT THIS ISSUE DOES NOT YET COVER

The flowchart at HEAD (flowchart_vKISS.html) lists the twelve stages below (0 through 9, with 4.5 and 4.6 as separate wired stages), plus orchestration and test data. Issue 002 had no chain at all — it was class cards and physics — so every stage is new territory for 003, and this issue documents only the stages that were executed with the canonical files during its preparation. The right-hand column says so, stage by stage. A reader who needs Stage 4.5, 5, 6, 8 or 9 will not find them here; they are the next issue's work, and the files that implement them are named so they can be read now.

Stage	Files at HEAD	Status	What it does (from source)	In Issue 003?
Stage 0 — Intake	stage_0_intake.py, cpg_intake_form.html, questionnaire.json	BUILT	patient folder, questionnaire, substrate declaration	NO
Stage 1 — Calibration	stage_1_idat_calibration.py, stage_1_calibration.py, idat_parse.py, idat_decoder_pure.py	BUILT	raw IDAT pair -> noob-normalised beta (methylprep 1.7.1; pure-python IDAT decoder shipped)	PROC-CAL-01 written, NOT RUN
Stage 2 — Deconvolution	Walther_iam_deconvolver/walther_iam_deconvolver.py, IAM_Atlas/*	BUILT (NILC cut 2026-07-02)	NNLS class + cell fractions; composition/presence; gates no call	YES — §2.4, §3.2, PROC-DECON-01, PROC-PLASMA-MIX-01
Stage 3 — Foreground subtraction	age/sex/smoking modules (test-only)	NOT WIRED by decision (SOP §104)	production subtracts nothing: age via the reference band, sex via chrX marker removal, smoking via Stage-7 bin floors	YES — RECON F1, RULES FIREWALL
Stage 4 — A-score gauge	cpg_gauge_engine.py, Runtime Matrices/A_Scoring_Module/*, test_a_score_canonical.py	BUILT	H(beta_mean)/H_min over identity loci; placement vs age band; severity ladder; brightness CI	YES — §3; brightness CI NOT described
Stage 4.5 — Bidirectional decomposition	Runtime Matrices/Directional Panel/bidirectional_decomposition.py, directional_panels_v1_0.json	BUILT (wired 2026-06-27)	composition-independent directional composite; AD detector (sealed VAL-051 Rule A, 7-CpG immune panel); gate composite > 0.40; AIBL-trained, does not transfer to GIFT	NO
Stage 4.6 — Patient CMB	cpg_patient_cmb.py, Runtime Matrices/cpg_healpix_mapping/*	BUILT (2026-06-29)	per-class departure $z=(\beta-\mu)/sd_{class}$ on a HEALPix sky; absent-tissue panels self-masked	NO
Stage 5 — Second chain	stage_5_second_chain.py, Mahalanobis_healthy_reference/*, Literature_anchors_Report_building/literature_anchors.json	BUILT	fires only on a flag; Mahalanobis adjudicator Option A (age-matched class gauge, presence gate $\geq 3\%$ AND outside [0.95,1.04]); RUN-everything residual matched-filter sweep (breast + immune-alarm maps; AD removed); literature anchors	NO — only the 3% floor is mentioned (RECON D2)
Stage 6 — Cellular age	iam_cellular_age_scoring.py (in CPG_TRIAL_CODE; NOT in repo engine)	BUILT 2026-06-30, calibration pending	cellular age from class A vs age_reference_matrix	NO — and the file is not in the repository (RECON F2)
Stage 7 — Tier	Runtime Matrices/Tier_breakpoints/tier_breakpoints.json	BUILT (edge 1.04->1.01, 2026-07-03)	continuous A -> customer tiers	YES — §3.1, RECON T1-T3

Stage	Files at HEAD	Status	What it does (from source)	In Issue 003?
Stage 8 — Disease matching	Disease Matrix/disease_cell_signature_matrix_v1_13.csv (81 rows, 49 VAL-anchored), disease_origin_cells.json, iamatlas_115_to_matrix_v0_2_mapping.json, Collinearity_Groups/*	BUILT	Route B directional concordance (weighted matcher over SIGNAL cells d >=0.20; STRONG needs dc>=0.70, coverage>=0.40, >=3 cells); Mode 2 cell-of-origin presence; Mode 3 systemic-stress wellness read (never a disease call); specificity gate (NLR axis = NON_SPECIFIC_GENERIC). Patient-match loop verified on 381 breast + 142 colon (VAL-093)	NO
Stage 9 — Report	cpg_report_builder.py, cpg_report_builder_v2.py, build_dashboard_v1.py, cpg_gauge.py (Appendix A1 gauge), Crown Jewel and Patient Strawman/*	BUILT	clinician report; patient straw man on the eight-class grid; crown-jewel reference wall; dashboard	NO
Orchestration	cpg_conductor.py (2026-07, replaces walther_clinical.py), walther_clinical.py (still the wired chain), run_batch.py, preflight.py, bootstrap.sh	BOTH present	conductor = Stage A pure functions; walther_clinical = the full wired chain the report builders still call	PARTIAL — conductor's presence rule only. walther_clinical.py origin-gate fail-open (bare except disables specificity rule) NOT in this issue
Test data	TEST_DATA/TEST_DATA_MANIFEST.md, harness/*, N7	present	11 public IDATs with expected outputs; synthetic harness; N7 chain-integrity	YES — §2.4, PROC-DECON-01

Findings from the 2026-09 review that are recorded here only as reconciliation rows, not yet as sections

The origin-gate fail-open in walther_clinical.py (F3); two stale strings describing the retired gauge, one on the patient report (F4); the gauge's missing sign (F5); two inequivalent A-score definitions (F6); the immune H_min revision history (F7); the seminoma value carried as both 0.67 and 0.755 in different files; the breast pre-diagnostic result (n = 47 cases / 601 controls across GSE51032 + GSE51057; matched-filter $p = +0.058$, CI [+0.001, +0.114]) and its unrun plate-position null; the colorectal cohort test CPG-NEW-001 in full (its P1/P4 failures are in §10, its method and N-random/N-comp results are not here); the two-observables test on GSE48684/GSE139404; and the disease-signature matrix v1.13 with its 81 rows. Each is in the project record (CPG_first_read.md, CPG_LEDGER.md, CPG_breast_CRC_learned.md, OUTCOME_CPG-NEW-001.md) and none has been re-verified for this issue.

SECTION 12

FOR THE CLINICIAN - THE INSTRUMENT IN PLAIN TERMS

Two of the source documents (The Cellular Margin; What Is Astro-Genetics) were written for readers who are not physicists. This section carries their content, corrected where the errata in s10 apply, so a physician reading only this section understands what a departure is and is not.

MARGIN

A cell spends one packet of ATP energy (about 54,000 J/mol) every time it does a unit of ordering work. Thermal noise at 37 C is worth $RT = 8.314 \times 310.15 \text{ J/mol}$. The ratio - how many ATP packets a cell has per packet of noise - is about 21. That is the cellular margin: healthy human cells run roughly twenty-one times above the bare thermal floor. It is a constant of the chemistry, the same for every patient; what varies between patients is how much of that margin the cell is actually using, which is what the A-score reads.

FLOOR

The floor is not a statistical cutoff. It is the minimum disorder a cell of a given type must carry to stay that type - the Landauer cost, $k_B T \ln 2$ per irreversible write, summed over the sites the cell is committed to. Eight cell architectures, eight floors, calculated once and frozen. A patient is read against the floor, not against other patients.

READING

$A = 1.0$ is the commitment line. Healthy cells sit on an age-dependent band near it (immune ~ 0.91 at age 4 rising to ~ 1.00 at 95). Frank cancer cells cluster at $A \sim 1.28\text{-}1.32$; established colon cancer reads ~ 1.15 . Below the band is also a reading - a cell that has locked down (senescence, or the stem-cell reversion seen in seminoma).

SUBSTRATE

The tube decides what can be read. Whole blood carries the immune architecture and nothing else - a healthy person has no epithelial DNA in it, and the instrument correctly says so. Plasma carries what tissues shed. Tissue reads everything present, mixed. Nothing in this document is a diagnosis; a departure is a flag for the clinician to place in context.

LIFE

One way to say what the instrument measures: life is the continuous purchase of organised pattern against the second law, paid for in metabolic work. The A-score is the receipt.

ERRATA

Two earlier clinician-facing drafts stated the floor as kT (it is $kT \ln 2$) and presented the Mahaffey number as a per-patient reading (it is a constant, 20.94; the per-patient quantity is A). Both are corrected here.

The stellar analogy the author uses: white dwarfs and neutron stars sit below a mass limit (Chandrasekhar $1.4 M_{\text{sun}}$, TOV $\sim 2.3 M_{\text{sun}}$) set by physics, not by a survey of stars; the Sun's eventual core at $0.54 M_{\text{sun}}$ reads 0.38 of the limit, Procyon B 0.42, PSR J0740+6620 0.90. The cellular floor is the same kind of number: a limit from the physics, against which each object is read individually.

APPENDIX V - VALIDATION INDEX: EVERY VAL IN THE REPOSITORY, WITH ITS PATH

Built by walking the repository for every path matching VAL-####: 103 identifiers, 1,175 files. Title is the first heading of the outcome (or prereg) file; date, cohorts and SHA presence are read from that file. 'Decision' is filled ONLY where the file states an explicit outcome code or status (36 of 103); otherwise 'see file' - the older outcome files carry the verdict in prose and this index does not paraphrase it. 'Record' says what the repo holds: outcome (87), prereg only (0), files only (16). Seven identifiers exist only under RETIRED_Phase1_PreBuild_Cards. VAL-047, cited in Issue 002 s8.1 as the deployment-readiness validation, has scripts in the repo but no prereg and no outcome file. This index is the map to the evidence, not the evidence; the historical VALs were left to stand as their own record (Issue 003 s10) and none was re-run here except where a PROC says so.

Path prefix BP/ = Biological_Physics/ (tree reorganized 2026-09-19: Testing_and_Code/VAL_PreAtlas, VAL_PostAtlas, RETIRED). Paths are directories; the prereg, outcome and results files sit inside them.

VAL	title	date	cohorts	decision	record	path (repo, 66f37fe)
VAL-001	001 — Per-cell-type A-score fan-out across 115 cell types — Brea	2026-06-03	GSE51032, GSE51057	O1_PRIMAR Y_VALIDATE D	outcome	BP/Testing_and_Code/VAL_PostAtlas/CPG_VAL_001_Breast_per_celltype_fanout
VAL-002	002 — Mahalanobis hyper-volume — Universal departure summary	2026-06-03	GSE51032, GSE51057	O1_PRIMAR Y_VALIDATE D	outcome	BP/Testing_and_Code/VAL_PostAtlas/CPG_VAL_002_Breast_mahalanobis
VAL-003	003 — Per-CpG residual map — CPG_breast_panel_v1 seed	2026-06-03	GSE51032, GSE51057	O1_PRIMAR Y_VALIDATE D	outcome	BP/Testing_and_Code/VAL_PostAtlas/CPG_VAL_003_Breast_residual_map
VAL-004	004 — Loss-of-bimodality count — RESTATED	2026-06-03	GSE51032, GSE51057	O3_RESTAT E_DIRECTIO N_REVERSE D_THEN	outcome	BP/Testing_and_Code/VAL_PostAtlas/CPG_VAL_004_Breast_bimodality
VAL-005	005 — Principal component axes — PC2 T-cell suppression	2026-06-03	GSE51032, GSE51057	O1_PRIMAR Y_VALIDATE D	outcome	BP/Testing_and_Code/VAL_PostAtlas/CPG_VAL_005_Breast_PC_axes
VAL-006	006 — chr6 MHC enrichment — RESTATED	2026-06-03	GSE51032, GSE51057	O4_RESTAT E_INSUFFICI ENT_POWE R_OR_C	outcome	BP/Testing_and_Code/VAL_PostAtlas/CPG_VAL_006_Breast_chr6_mhc
VAL-007	007 — Age-axis foreground subtraction confirms signal robustness	2026-06-03	GSE51032, GSE51057	O1_PRIMAR Y_VALIDATE D	outcome	BP/Testing_and_Code/VAL_PostAtlas/CPG_VAL_007_Breast_age_subtraction
VAL-008	008 + 009 — AIBL Preliminary Outcomes	2026-06-02	GSE153712, GSE51032, GSE51057	see file	outcome	BP/Testing_and_Code/VAL_PreAtlas/CPG_VAL_008_AD_AIBL_per_celltype
VAL-009	009 — Mahalanobis hyper-volume (AIBL)		GSE153712, GSE51032, GSE51057	see file	outcome	BP/Testing_and_Code/VAL_PreAtlas/CPG_VAL_009_AD_AIBL_mahalanobis
VAL-010	010 — AddNeuroMed cross-platform replication	2026-06-02	GSE144858	see file	outcome	BP/Testing_and_Code/VAL_PreAtlas/CPG_VAL_010_AD_AddNeuroMed
VAL-011	011 — Age-axis foreground subtraction	2026-06-02	GSE144858, GSE53740	O3_TISSUE_ FLOOR_DO MINATED	outcome	BP/Testing_and_Code/VAL_PreAtlas/CPG_VAL_011_AD_age_subtraction
VAL-012	012 — Principal-component axes of the AD signal	2026-06-02	GSE153712	see file	outcome	BP/Testing_and_Code/VAL_PreAtlas/CPG_VAL_012_AD_PC_axes
VAL-013	013 — Per-CpG residual map → CPG_ad_panel_v1 candidate	2026-06-02	GSE144858, GSE153712, GSE51032, GSE51057	see file	outcome	BP/Testing_and_Code/VAL_PreAtlas/CPG_VAL_013_AD_residual_map

VAL	title	date	cohorts	decision	record	path (repo, 66f37fe)
VAL-014	014 — GIFT specificity arm (AD vs FTD vs PSP/CBD vs HC)	2026-06-02	GSE53740	see file	outcome	BP/Testing_and_Code/VAL_PreAtlas/CPG_VAL_014_AD_GIFT_specificity
VAL-015	015 — OUTCOME	2026-06-07		PASS	outcome	BP/Testing_and_Code/VAL_PostAtlas/CPG_VAL_015_Immune_Aging_Trajectory
VAL-016	016 — OUTCOME	2026-06-07	GSE51032, GSE51057	DIRECTIONAL	outcome	BP/Testing_and_Code/VAL_PostAtlas/CPG_VAL_016_Cross_Disease_Universal_Alarm
VAL-017	017 — OUTCOME	2026-06-07	GSE141682, GSE40279, GSE50660, GSE51032	NULL	outcome	BP/Testing_and_Code/VAL_PostAtlas/CPG_VAL_017_Inflammaging_Pooled_HC
VAL-018	018 — OUTCOME	2026-06-07	GSE51032, GSE51057	NULL	outcome	BP/Testing_and_Code/VAL_PostAtlas/CPG_VAL_018_Menarche_Immune
VAL-019	019 — OUTCOME	2026-06-07	GSE51057	PASS	outcome	BP/Testing_and_Code/VAL_PostAtlas/CPG_VAL_019_Immune_Bidirectional
VAL-020	020 — OUTCOME (sealed 2026-06-06)	2026-06-06	GSE40279, GSE51032, GSE51057	O_CHAIN_INTEGRITY_PASS	outcome	BP/Testing_and_Code/VAL_PostAtlas/CPG_VAL_020_Immune_Hannum_full_chain
VAL-021	(no heading in file) dir: CPG_VAL_021_Weight_Loss_DEFERRED			see file	files only	BP/Testing_and_Code/VAL_PreAtlas/CPG_VAL_021_Weight_Loss_DEFERRED
VAL-022	022 — OUTCOME	2026-06-07		NULL	outcome	BP/Testing_and_Code/VAL_PostAtlas/CPG_VAL_022_Smoking_Cessation_Reversibility
VAL-037	(no heading in file) dir: validation_runs			see file	files only	BP/Testing_and_Code/VAL_PreAtlas
VAL-038	(no heading in file) dir: validation_runs			see file	files only	BP/Testing_and_Code/VAL_PreAtlas
VAL-039	(no heading in file) dir: validation_runs			see file	files only	BP/Testing_and_Code/VAL_PreAtlas
VAL-040	(no heading in file) dir: validation_runs			see file	files only	BP/Testing_and_Code/VAL_PreAtlas
VAL-041	(no heading in file) dir: validation_runs			see file	files only	BP/Testing_and_Code/VAL_PreAtlas
VAL-042	(no heading in file) dir: validation_runs			see file	files only	BP/Testing_and_Code/VAL_PreAtlas
VAL-043	(no heading in file) dir: validation_runs			see file	files only	BP/Testing_and_Code/VAL_PreAtlas
VAL-044	(no heading in file) dir: validation_runs			see file	files only	BP/Testing_and_Code/VAL_PreAtlas
VAL-045	(no heading in file) dir: validation_runs			see file	files only	BP/Testing_and_Code/VAL_PreAtlas
VAL-046	(no heading in file) dir: validation_runs			see file	files only	BP/Testing_and_Code/VAL_PreAtlas
VAL-047	(no heading in file) dir: validation_runs			see file	files only	BP/Testing_and_Code/VAL_PreAtlas
VAL-050	(no heading in file) dir: val_050_aibl	2026-04-23		see file	outcome	BP/Testing_and_Code/VAL_PreAtlas/val_050_aibl

VAL	title	date	cohorts	decision	record	path (repo, 66f37fe)
VAL-051	(no heading in file) dir: cohort_artifacts	2026-04-23		see file	outcome	BP/Testing_and_Code/VAL_PreAtlas/val_051_ad_directional
VAL-052	(no heading in file) dir: val_052_addneuromed	2026-04-23		see file	outcome	BP/Testing_and_Code/VAL_PreAtlas/val_052_addneuromed
VAL-053	(no heading in file) dir: val_053_sex_panels	2026-04-23		see file	outcome	BP/Testing_and_Code/VAL_PreAtlas/val_053_sex_panels
VAL-054	(no heading in file) dir: val_054_age_bound	2026-04-23		see file	outcome	BP/Testing_and_Code/VAL_PreAtlas/val_054_age_bound
VAL-056	(no heading in file) dir: validation_runs			see file	files only	BP/Testing_and_Code/VAL_PreAtlas
VAL-057	(no heading in file) dir: validation_runs	2026-04-24	GSE53740	see file	outcome	BP/Testing_and_Code/VAL_PreAtlas
VAL-058	(no heading in file) dir: validation_runs	2026-04-24	GSE269244	see file	outcome	BP/Testing_and_Code/VAL_PreAtlas
VAL-059	(no heading in file) dir: validation_runs	2026-04-24	GSE281691, GSE298812	see file	outcome	BP/Testing_and_Code/VAL_PreAtlas
VAL-060	(no heading in file) dir: validation_runs	2026-04-24	TCGA-BRCA, TCGA-COAD, TCGA-LUAD	see file	outcome	BP/Testing_and_Code/VAL_PreAtlas
VAL-061	Outcome — CRC Tissue, Immune-Compartment Reading on TCGA-COAD	2026-04-24	TCGA-COAD	see file	outcome	BP/Testing_and_Code/VAL_PreAtlas
VAL-062	Outcome — CRC Tumor Tissue Cycling-Class Rescore	2026-04-24	TCGA-COAD	see file	outcome	BP/Testing_and_Code/VAL_PreAtlas
VAL-063	Outcome — Lung-EPIC Tissue Arm on TCGA-LUAD HM450	2026-04-24	GSE269244, TCGA-BRCA, TCGA-COAD, TCGA-LUAD	see file	outcome	BP/Testing_and_Code/VAL_PreAtlas
VAL-064	Outcome — HCC-EPIC Tissue Arm on TCGA-LIHC HM450	2026-04-24	GSE269244, GSE298812, TCGA-BRCA, TCGA-COAD	see file	outcome	BP/Testing_and_Code/VAL_PreAtlas
VAL-065	Outcome — prostate-epic Urine Arm Specimen Comparison	2026-04-25	GSE119260	see file	outcome	BP/Testing_and_Code/VAL_PreAtlas
VAL-066	(no heading in file) dir: validation_runs	2026-04-25		see file	outcome	BP/Testing_and_Code/VAL_PreAtlas
VAL-067	(no heading in file) dir: validation_runs	2026-04-25		see file	outcome	BP/Testing_and_Code/VAL_PreAtlas
VAL-068	(no heading in file) dir: validation_runs	2026-04-25		see file	outcome	BP/Testing_and_Code/VAL_PreAtlas
VAL-069	(no heading in file) dir: validation_runs	2026-04-25		see file	outcome	BP/Testing_and_Code/VAL_PreAtlas
VAL-070	(no heading in file) dir: validation_runs			see file	files only	BP/Testing_and_Code/VAL_PreAtlas

VAL	title	date	cohorts	decision	record	path (repo, 66f37fe)
VAL-071	(no heading in file) dir: validation_runs			see file	files only	BP/Testing_and_Code/VAL_PreAtlas
VAL-072	— TCGA-CESC tissue arm OUTCOME	2026-04-25	GSE143752, GSE287994, GSE38266, GSE46306	see file	outcome	BP/Testing_and_Code/VAL_PreAtlas
VAL-073	(no heading in file) dir: validation_runs	2026-04-25		see file	outcome	BP/Testing_and_Code/VAL_PreAtlas
VAL-074	— GSE46306 Farkas 2013 cervical tissue OUTCOME	2026-04-25	GSE46306, GSE68339, GSE99511	O5_NEGATIVE_DIRECTI ON — disease re	outcome	BP/Testing_and_Code/VAL_PreAtlas
VAL-076	(no heading in file) dir: validation_runs	2026-04-25		O3_LBC_NU LL — primary specimen pat	outcome	BP/Testing_and_Code/VAL_PreAtlas
VAL-077	(no heading in file) dir: validation_runs	2026-04-25		O3_LBC_NU LL — largest LBC cohort c	outcome	BP/Testing_and_Code/VAL_PreAtlas
VAL-081	(no heading in file) dir: validation_runs	2026-04-25		O5_NEGATIVE_DIRECTI ON — tumors rea	outcome	BP/Testing_and_Code/VAL_PreAtlas
VAL-082	— GSE62298 AML myeloid arm OUTCOME	2026-04-25	GSE51057, GSE62298	O1_PASS_M YELOID_ARM_AT_BLOOD_LEVEL	outcome	BP/Testing_and_Code/VAL_PreAtlas
VAL-088	— glioma-epic Stage 1 outcome		GSE180683, GSE51057	see file	outcome	BP/Testing_and_Code/VAL_PreAtlas
VAL-089	— glioma-epic Tumor Tissue arm outcome		GSE143843, GSE90496, TCGA-LGG	see file	outcome	BP/Testing_and_Code/VAL_PreAtlas
VAL-090	(no heading in file) dir: validation_runs	2026-04-25		O1_PASS	outcome	BP/Testing_and_Code/VAL_PreAtlas
VAL-091	(no heading in file) dir: validation_runs	2026-04-26		see file	outcome	BP/Testing_and_Code/VAL_PreAtlas
VAL-092	(no heading in file) dir: VAL-092	2026-04-26		see file	outcome	BP/Testing_and_Code/VAL_PreAtlas/VAL-092
VAL-093	(no heading in file) dir: VAL-093	2026-04-26		see file	outcome	BP/Testing_and_Code/VAL_PreAtlas/VAL-093
VAL-094	— Outcome	2026-04-26	GSE51032, GSE51057	see file	outcome	BP/Testing_and_Code/VAL_PreAtlas
VAL-095	— Outcome	2026-04-26	GSE51032, GSE51057	see file	outcome	BP/Testing_and_Code/VAL_PreAtlas/VAL-095
VAL-096	— Outcome	2026-04-26	GSE51032, GSE51057	see file	outcome	BP/Testing_and_Code/VAL_PreAtlas/VAL-096
VAL-097	(no heading in file) dir: VAL-097	2026-04-28		see file	outcome	BP/Testing_and_Code/VAL_PreAtlas/VAL-097

VAL	title	date	cohorts	decision	record	path (repo, 66f37fe)
VAL-098	(no heading in file) dir: VAL-098	2026-04-28		see file	outcome	BP/Testing_and_Code/VAL_PreAtlas/VAL-098
VAL-099	(no heading in file) dir: VAL-099	2026-04-28		see file	outcome	BP/Testing_and_Code/VAL_PreAtlas/VAL-099
VAL-100	(no heading in file) dir: VAL-100	2026-04-28		see file	outcome	BP/Testing_and_Code/VAL_PreAtlas/VAL-100
VAL-101	(no heading in file) dir: VAL-101	2026-04-28		see file	outcome	BP/Testing_and_Code/VAL_PreAtlas/VAL-101
VAL-102	(no heading in file) dir: VAL-102			see file	outcome	BP/Testing_and_Code/VAL_PreAtlas/VAL-102
VAL-106	(no heading in file) dir: VAL-106			O3_CALIBRATION_DEGENERATE	outcome	BP/Testing_and_Code/VAL_PreAtlas/VAL-106
VAL-107	(no heading in file) dir: VAL-107			see file	outcome	BP/Testing_and_Code/VAL_PreAtlas/VAL-107
VAL-108	(no heading in file) dir: VAL-108			O3_CARDIO_EPIC_3SUBTYPE_UNDIFFEREN	outcome	BP/Testing_and_Code/VAL_PreAtlas/VAL-108
VAL-109	(no heading in file) dir: VAL-109			O2_PAH_VASCULAR_TILE_DIFFERENTIATI	outcome	BP/Testing_and_Code/VAL_PreAtlas/VAL-109
VAL-110	(no heading in file) dir: VAL-110			O2_AORTIC_ANY_TILE_DIFFERENTIATING	outcome	BP/Testing_and_Code/VAL_PreAtlas/VAL-110
VAL-111	(no heading in file) dir: VAL-111	2026-04-29	GSE69138, GSE84274, GSE84395	O3_TISSUE_FLOOR_DOMINATED	outcome	BP/Testing_and_Code/VAL_PreAtlas/VAL-111
VAL-112	VAL-112+113 unified outcome — Cardio-epic run-everything	2026-04-29	GSE69138, GSE84274, GSE84395, TCGA-KIRC	see file	outcome	BP/Testing_and_Code/VAL_PreAtlas/VAL-112_113_unified
VAL-113	(no heading in file) dir: VAL-113_caggiano			see file	files only	BP/Testing_and_Code/VAL_PostAtlas/VAL-113_caggiano
VAL-117	(no heading in file) dir: VAL-117_prostateref_calibrate			see file	outcome	BP/Testing_and_Code/VAL_PreAtlas/VAL-117_prostateref_calibrate
VAL-118	— Outcome (Amendment Re-Execution)		GSE269244	see file	outcome	BP/Testing_and_Code/VAL_PreAtlas/VAL-118_prostateref_phaseC
VAL-119	(no heading in file) dir: VAL-119_bladderref_calibrate			see file	outcome	BP/Testing_and_Code/VAL_PreAtlas/VAL-119_bladderref_calibrate
VAL-120	(no heading in file) dir: VAL-120_bladder_stage1_xu538			see file	outcome	BP/Testing_and_Code/VAL_PreAtlas/VAL-120_bladder_stage1_xu538
VAL-121	(no heading in file) dir: VAL-121_bladder_stage2_multiatlas			see file	outcome	BP/Testing_and_Code/VAL_PreAtlas/VAL-121_bladder_stage2_multiatlas
VAL-122	(no heading in file) dir: VAL-122_bladder_stage3_immune			see file	outcome	BP/Testing_and_Code/VAL_PreAtlas/VAL-122_bladder_stage3_immune

VAL	title	date	cohorts	decision	record	path (repo, 66f37fe)
VAL-123	(no heading in file) dir: VAL-123_boccellato_calibrate	2026-05-02		see file	outcome	BP/Testing_and_Code/VAL_PreAtlas/VAL-123_boccellato_calibrate
VAL-124	(no heading in file) dir: VAL-124_esoref_calibrate	2026-05-02		see file	outcome	BP/Testing_and_Code/VAL_PreAtlas/VAL-124_esoref_calibrate
VAL-125	(no heading in file) dir: VAL-125_oeref_calibrate	2026-05-02		see file	outcome	BP/Testing_and_Code/VAL_PreAtlas/VAL-125_oeref_calibrate
VAL-126	(no heading in file) dir: VAL-126_stad_phase_c	2026-05-02		SEALED	outcome	BP/Testing_and_Code/VAL_PreAtlas/VAL-126_stad_phase_c
VAL-127	(no heading in file) dir: VAL-127_esca_phase_c	2026-05-02		SEALED	outcome	BP/Testing_and_Code/VAL_PreAtlas/VAL-127_esca_phase_c
VAL-128	(no heading in file) dir: VAL-128_crohns_blood	2026-05-02		SEALED	outcome	BP/Testing_and_Code/VAL_PreAtlas/VAL-128_crohns_blood
VAL-135	135 — Immune card v1.0 full-chain reproduction of VAL-051 anchor	2026-06-06	GSE153712, GSE40279, GSE50660, GSE51032	O1_PRIMAR Y_VALIDATE D	outcome · RETIRED	BP/Testing_and_Code/Cohorts_and_Manifests/ RETIRED_OFF_SCOPE_VAL_EXPLORATION/ VAL-135_immune_card_v1_full_chain_validation
VAL-136	136 — Smoking-axis subtraction reduces the never-vs-current A_im	2026-06-06	GSE50660	O1_PRIMAR Y_VALIDATE D	outcome · RETIRED	BP/Testing_and_Code/Cohorts_and_Manifests/ RETIRED_OFF_SCOPE_VAL_EXPLORATION/ CPG_VAL_136_GSE50660_smoking_subtraction
VAL-137	137 — Sex-axis subtraction reduces the M-vs-F A_immune contrast	2026-06-06	GSE40279, GSE50660	O3_INVERTE D	outcome · RETIRED	BP/Testing_and_Code/Cohorts_and_Manifests/ RETIRED_OFF_SCOPE_VAL_EXPLORATION/ CPG_VAL_137_GSE40279_sex_subtraction
VAL-138	138 — Age-axis subtraction reduces A_immune's chronological age	2026-06-06	GSE40279, GSE51032, GSE51057	O3_INVERTE D	outcome · RETIRED	BP/Testing_and_Code/Cohorts_and_Manifests/ RETIRED_OFF_SCOPE_VAL_EXPLORATION/ CPG_VAL_138_GSE40279_age_subtraction
VAL-139	139 — Directional outperforms pooled by 8.0× on AIBL AD discrimi	2026-06-06	GSE153712	O1_PRIMAR Y_VALIDATE D	outcome · RETIRED	BP/Testing_and_Code/Cohorts_and_Manifests/ RETIRED_OFF_SCOPE_VAL_EXPLORATION/ CPG_VAL_139_AIBL_pooled_vs_directional
VAL-140	140 — 8-class A-score vector recovers 25.1% of chronological age	2026-06-06	GSE40279	O5_BASELIN E_DOMINAT ED	outcome · RETIRED	BP/Testing_and_Code/Cohorts_and_Manifests/ RETIRED_OFF_SCOPE_VAL_EXPLORATION/ CPG_VAL_140_GSE40279_cellular_age_inversion
VAL-141	141 — Cross-cohort A_immune baseline concordance	2026-06-06	GSE40279, GSE50660	O2_NEAR_T HRESHOLD	outcome · RETIRED	BP/Testing_and_Code/Cohorts_and_Manifests/ RETIRED_OFF_SCOPE_VAL_EXPLORATION/ CPG_VAL_141_cross_cohort_A_immune_concordance

APPENDIX VI — THE CMB → METHYLOME TRANSLATION MAP, SCORED

Written by the author before the chain was built: the 32-section CMB-analysis curriculum walked module by module, each assigned a methylome analog and a status (HAVE · ROADMAP · ADD = should add · REINTERP = translates with reinterpretation · NO = does not translate). The last column was added on 2026-09-19 and records what the chain actually did. Two rows went *against* the map and are the most instructive: a second deconvolver (row 20, built then cut) and de-aging (row 47, built then refused). Row 1 is the frame for the whole chain: the 115 cell types are the harmonic basis, and the deconvolver is the projection onto it. This appendix exists so that a geneticist can see where each tool came from, and a cosmologist can see where the methylome stops behaving like the sky.

Section I — Foundations: what the CMB field actually is

#	CMB module	CPG analog	status	author's note	2026-09-19 status
1	Spherical harmonic decomposition $T(\hat{n}) = \sum a_{lm} Y_{lm}$, including E/B polarization	The genome isn't a sphere, but it's a structured manifold (chromosomes × position × 3D chromatin TADs). Harmonic decomposition becomes (a) per-chromosome Fourier along position, (b) per-TAD modal decomposition, (c) the 115-cell-type basis itself acts as a "feature-space harmonic basis"	REINTERP	Need to write this up — it's the single biggest conceptual move in the rulebook. The 115 cell types are our spherical harmonics. A patient's β profile gets decomposed onto that basis via Walther IAM Deconvolver.	UNWRITTEN as a chapter — but the chain does exactly this: the deconvolver projects a patient's β onto the 115-cell-type basis (Stage 2) and the class gauge reads eight coefficients of that projection. The 'harmonic basis' is the Atlas. Row 1 is the conceptual frame for Part II's Atlas chapters.
2	Monopole, dipole, quadrupole, low-I culture	Monopole = mean genome-wide β . Dipole = age axis (CPG-VAL-007) + sex axis. Quadrupole+ = cellular-compositi on axes (TODO 1.5 PCA). Low-I anomalies = global-genome-level patterns (CIMP+, X-inactivation status, ploidy)	HAVE	Implicit today, need to make explicit.	
3	Cosmic variance — finite-mode irreducibl e noise	Cellular variance — finite number of independent cells contributing to a methylation reading. A pure plasma sample only carries ~10,000 cell-equivalents of cfDNA; you cannot reduce the variance below the cell-count floor. Same structural limit as 2l+1 modes.	ADD	Important to name. Defines fundamental limits on small samples.	NOT BUILT — cfDNA cell-equivalent variance floor never modelled

Section II — Experimental / instrument layer

#	CMB module	CPG analog	status	author's note	2026-09-19 status
4	Frequency coverage (multi-band foreground separatio n)	Multi-substrate methylation (5mC, 5hmC, 5fC, 5caC), plus orthogonal modalities (RNA-seq, ATAC-seq, fragmentomics). CPG v0.1 is single-frequency (5mC); the architecture must declare the multi-substrate slot even though we're not multi-frequency yet.	ROADMAP	Tier 5 of EDEAR Physics Roadmap — 5-substrate multi-assay (already named in the L3 lab plan).	
5	Beam modeling	Probe response function — each 450K/EPIC probe has a sequence context, hybridization efficiency, and adjacent-CpG cross-reactivity. Probe design errors are the methylome's beam errors.	ADD	Not yet explicit in CPG. Should add.	

#	CMB module	CPG analog	status	author's note	2026-09-19 status
6	Bandpass calibration	Bisulfite conversion efficiency + probe-type normalization (Type I vs Type II). Same physics: non-delta-function response in the chemistry domain leaks signal between channels.	REINTERP	Standard preprocessing handles it but we don't TREAT it as bandpass calibration with nuisance parameters. Should formalize.	
7	Gain calibration / drift	Batch effect + plate position + run date drift. Lab-to-lab and run-to-run gain drift is the methylome version of detector gain drift.	HAVE	Standard pre-processing (BMIQ, funnorm, ComBat) — but we should declare it L2 explicitly.	
8	Polarization-angle calibration	Strand-orientation calibration. Methylation has a forward-strand vs reverse-strand asymmetry; bisulfite sequencing can mis-assign. The polarization-rotation analog is real but minor at the array level.	REINTERP	Mostly a sequencing-era concern; arrays largely solved it. Document but low-priority.	
9	Scanning strategy	Cohort sampling design. How patients are sampled across age/sex/site/batch determines what coverage you have and what nulls you can run. Pre-build cohorts were chosen for us; future CPG-acquired cohorts must be designed with "scanning strategy" rigor.	ADD	Important for prospective cohort acquisition.	

Section III — Time-ordered data to maps

#	CMB module	CPG analog	status	author's note	2026-09-19 status
10	TOD analysis	IDAT intensities + control probe traces + bead-pool QC. Before β values exist, there are raw intensities + control signals = our "timestream."	HAVE	Standard tooling handles it.	
11	De-glitching + data cuts	Probe flagging, sample QC dropouts, detection-p-value masking. Failed probes, low-quality samples, contamination flags.	HAVE	Standard. Declare explicitly.	
12	Noise modeling (1/f, atmospheric, correlated detector noise)	Probe-position-on-array correlated noise, plate-edge effects, scanner drift within run. The methylome has its own non-white noise — and CPG hasn't modeled it explicitly.	ADD	Significant gap. Should add a noise-covariance module.	NOT BUILT
13	Destriping / max-likelihood map-making	Quantile normalization / funnorm / BMIQ. All raw-intensity-to- β methods are inverse-problem solvers in this same family.	HAVE	Standard tooling. Map preprocessing.	

#	CMB module	CPG analog	status	author's note	2026-09-19 status
1 4	Transfer functions (pipeline filters the true sky)	The deconvolver itself filters disease signal — it explains away anything that looks like cell composition. We saw this in CPG-VAL-006: chr16/chr17 depleted because their CpGs got "explained away."	ADD	This is exactly what we already discovered empirically and didn't have a name for. Transfer-function thinking is critical — every step of the chain alters the signal in modeled ways.	NAMED in the lessons (CCL-039, glioma-LL-002: the deconvolver 'explains away'); no module

Section IV — Pixelization & spherical computation

#	CMB module	CPG analog	status	author's note	2026-09-19 status
1 5	HEALPix pixelization	CpG manifest as the pixelization — 450K → 850K → EPIC v2 1.6M = increasingly fine pixelization. Each CpG is a pixel. Cell-type space is also pixelized (8 classes → 115 cells → finer with each MCMC refinement).	HAVE	Implicit. Make explicit.	
1 6	Masks + apodization	CpG masks — X-inactivation, imprinting, repeat regions, low-confidence probes, SNP-overlapping probes. Mask apodization = down-weighting near suspect regions rather than hard-cutting.	REINTERP	Hard masking common; soft apodization unusual. Could be a CPG innovation.	
1 7	Cut-sky mode coupling, pseudo-Cl, MASTER	Cut-genome mode coupling. When some CpGs are masked, the per-chromosome correlation function (TODO 2.1) is biased — same MASTER algorithm applies.	ADD	Critical for TODO 2.1 to be rigorous.	

Section V — Foregrounds and component separation (the biggest section)

#	CMB module	CPG analog	status	author's note	2026-09-19 status
1 8	Galactic foreground modeling (sync, dust, free-free, AME, CO, etc.)	Cellular foreground stack: cell composition, age, sex, ancestry, batch, smoking, BMI, medication exposure, plate position. These are our synchrotron+dust+free-free+AME.	REINTERP	We handle cell composition (Walther) + age (CPG-VAL-007). We do NOT yet handle sex, ancestry, batch, smoking explicitly. Big gap.	PARTIAL — composition (Walther) and age (band) handled; sex/smoking/batch layers BUILT and deliberately NOT wired (SOP §104)
1 9	Extragalactic foregrounds (CIB, SZ, point sources)	High-resolution disease-overlapping foregrounds: subclinical inflammation, comorbidities, pharmacological methylation effects (e.g., 5-azacytidine), recent acute illness, vaccination response. These are the methylome's "point sources."	ADD	Almost completely unmodeled. Big gap.	NOT BUILT
2 0	Component-separation methods (Commander, NILC, SMICA, SEVEM, GNILC)	Walther IAM Deconvolver = one method. We need to build at least one INDEPENDENT method (parametric Bayesian vs blind/ILC-style) so we can compare, the way Planck compared Commander/NILC/SMICA/SEVEM.	ROADMAP	Critical addition. No single-method CMB result was ever trusted. Need at least 2 deconvolvers.	REVERSED — NILC built as the second method, then CUT 2026-07-02 (c1be0c3): it collapsed on correlated blood mixtures and deleted correct calls. Single-method today, by decision. The Planck principle this row cites is the one the chain knowingly does not follow.

#	CMB module	CPG analog	status	author's note	2026-09-19 status
21	Foreground residual marginalization in likelihood	Nuisance parameter for residual cell-composition error, residual age error, residual batch error. Right now we point-estimate these. Should marginalize.	ADD	Big upgrade. Required for clean L7 likelihood.	NOT BUILT — point estimates

Section VI — Power-spectrum estimation

#	CMB module	CPG analog	status	author's note	2026-09-19 status
22	Auto vs cross-spectra (split maps, half-mission × half-mission)	Split-cohort cross-correlations. GSE51057 × GSE51032 cross is the methylome's "half-mission × half-mission." We already used this in CPG-VAL-003 (concordance gate).	HAVE	Mark explicit. Could go further with within-cohort odd/even sample splits.	BUILT — GSE51032 × GSE51057 is the foundation-cohort anchor; reproduced 2026-09-19 at $r = 1.00000$
23	Pseudo-Cl on masked sky	Per-chromosome C(d) with masked CpGs handled correctly. Required for TODO 2.1.	ROADMAP	Build with MASTER-style mode-coupling correction.	
24	Quadratic max-likelihood at low l	Per-TAD or per-chromosome-arm maximum-likelihood spectrum estimation. For low-mode-count regimes.	ADD	Niche but worth declaring.	
25	Bandpowers	Genomic distance bandpowers — group C(d) into log-spaced d-bins with proper covariance.	ROADMAP	Part of TODO 2.1.	
26	Beam/calibration nuisance treatment	Probe-response + bisulfite efficiency nuisance treatment.	ADD	Add to L7 likelihood.	

Section VII — Likelihoods and covariances

#	CMB module	CPG analog	status	author's note	2026-09-19 status
27	Low-l likelihoods (non-Gaussian, pixel-space, Gibbs sampling)	Low-frequency disease signal likelihoods — when there are very few discriminating CpGs (rare cancers, early pre-dx), the distribution of A-score isn't Gaussian. Need exact-likelihood treatment.	ADD	Future.	
28	High-l likelihoods (Gaussian bandpower with detailed covariance)	Mahalanobis hyper-volume (CPG-VAL-002) is the seed. Need to extend to per-card formal likelihoods.	ROADMAP	This is L7 of the chain.	
29	Covariance matrices	A-score covariance (have via Ledoit-Wolf in CPG-VAL-002). Need: per-CpG covariance, cross-cohort covariance, nuisance-parameter covariance.	ROADMAP	Tier 2-3 work.	
30	End-to-end simulations	Synthetic patient generation: known cell-type fractions + known disease signal + simulated batch + run through entire chain. Validate that the chain recovers the truth.	ADD	The single biggest missing piece in CPG today. Cannot defend the chain without it.	BUILT — <code>synthetic_patient_harness.py</code> + null N7

#	CMB module	CPG analog	status	author's note	2026-09-19 status
3 1	Null tests	HC permutation, age permutation, sex permutation, cohort splits, plate-position null, batch null, synthetic injection-recovery. Some scattered today, no unified suite.	ROADMAP	L9 of the chain — unified null suite is on the engine-completion sprint.	BUILT — <code>cpg_null_runner.py</code> N1–N8 (found in RETIRED/ and restored to CPG_Engine/ 2026-09-19)
3 2	PTEs and goodness-of-fit + look-elsewhere	Mandatory. Especially for rare biological anomalies.	ADD	Add to outcome.md template.	

Section VIII — Theory prediction engines

#	CMB module	CPG analog	status	author's note	2026-09-19 status
3 3	Boltzmann solvers (CAMB/CLASS)	The IAM physics framework + $H_{\min}$ derivation IS our Boltzmann solver. It predicts the architectural floor against which everything else is measured.	HAVE	Already in place — the Recipe is the methylome's CAMB. Vault content.	BUILT — the Recipe / $H_{\min}$ derivation (vault); 40 floors in <code>cpg_gauge_engine.py</code>
3 4	Recombination physics (RECFAST/HyRec)	Cellular differentiation thermodynamics — the per-class $H_{\min}$ values are frozen at differentiation, analogous to acoustic peaks frozen at recombination.	HAVE	Implicit in IAM.	
3 5	Line-of-sight integration	Lineage tracing through cell development. A methylation pattern in a mature cell reflects integrated history of decisions along the developmental trajectory.	REINTERP	Translatable but speculative.	
3 6	Acoustic-peak phenomenology	The pattern cheat-sheet for methylation features. (See Section XVII below.) This is the rulebook's most teachable section.	ADD	Build a CPG version of the cheat-sheet.	

Section IX — Parameter inference & degeneracies

#	CMB module	CPG analog	status	author's note	2026-09-19 status
3 7	MCMC + posterior geometry	MCMC for per-card posterior tier (replacing threshold-based scoring). Already familiar from IAMAtlas build.	ROADMAP	Engine-completion work.	NOT BUILT — MCMC used for the Atlas and the floors; scoring is threshold + age band
3 8	Banana degeneracies	CIMP-axis degeneracy is visible in Image 2. Other expected ones: age-vs-class-drift, immune-fraction-vs-disease-signal, sex-vs-X-inactivation.	ROADMAP	TODO 2.3 banana degeneracy is the first concrete one.	
3 9	Priors	Disease prior (cancer_prior runtime matrix), family-history prior(family_history_multiplier), age prior. Already implemented but should be treated as proper Bayesian priors with volume awareness.	REINTERP	Have, but informal. Upgrade.	

#	CMB module	CPG analog	status	author's note	2026-09-19 status
40	Profile likelihood vs marginalized posterior	Critical distinction. Profile likelihood = "where is the best-fit cell-type fraction" vs marginalized = "where is the volume in fraction space." They can disagree when nuisance dimensions are high.	ADD	Add to L7.	
41	Multimodal distributions	Disease subtypes are multimodal posteriors. HER2+ vs Basal vs Luminal-A breast cancer don't sit on one mode. Posterior clustering matters.	REINTERP	Recognized but not formal.	
42	Nested sampling and Bayesian evidence	Model comparison: does this patient look more like 'breast pre-dx' or 'pancreatic pre-dx' or 'inflammatory baseline'? Bayesian evidence is the right tool.	ADD	Future engine layer.	
43	Tension metrics	Cross-card tension when two cards both fire on one patient. Need proper tension language, not just "both elevated."	ADD	Operational rule for Stage 6 reports.	

Section X — Lensing: the CMB as a backlight

#	CMB module	CPG analog	status	author's note	2026-09-19 status
44	Lensing smoothing of acoustic peaks	Late-life methylation drift smooths early-life methylation patterns. Aging is the methylome's lensing — it smears the original developmental signal.	REINTERP	Beautiful analogy. Worth a paper.	UNWRITTEN — the AGE BAND is the operational form of 'aging as lensing'; the analogy itself appears nowhere in the corpus and is the cleanest way to explain the band to a physicist
45	Four-point lensing reconstruction	CpG-CpG-CpG-CpG four-point correlations — coordinated methylation changes across 4 loci would reveal coordinated regulatory programs. Bispectrum (TODO 2.2) is 3-point; this would be 4-point.	ADD	Tier 2+ research.	
46	Lensing power spectrum	Aging "drift power spectrum" — measure how aging redistributes power across methylation features. The IAM cellular age clock (TODO 2.4) is the seed.	ROADMAP	TODO 2.4 lives here.	
47	Delensing	De-aging: subtract the age component to recover the disease component. CPG-VAL-007 dipole subtraction is the first instance.	HAVE	Have seed. Generalize.	REVERSED — de-aging (age_axis_foreground.py) BUILT, then §104 ruled annotate-don't-subtract. The chain does the opposite of delensing, on purpose: the methylome's foreground is the patient's own biology. The one place the analogy is deliberately broken.
48	Cross-correlations with external probes	Cross-correlate CPG output with RNA-seq, proteomics, imaging.	ROADMAP	Future multi-omics.	

Section XI — Polarization special topics

#	CMB module	CPG analog	status	author's note	2026-09-19 status
49	E/B separation	Methylated vs hydroxymethylated separation — the methylome's literal polarization. 5mC vs 5hmC is an E/B-like spin decomposition.	REINTERP	Substrate-level; requires oxBS-seq (5-substrate Tier 5).	
50	B-mode taxonomy	5hmC contamination taxonomy — what shows up in a "5hmC signal" can be true 5hmC, residual unconverted 5mC, sequencing error, or biological switch.	ADD	Substrate-dependent.	
51	Reionization bump (low-I EE)	Developmental methylation reprogramming bump — large-scale early-life methylation events leave a "reionization-like" signature in the lowest-frequency modes.	REINTERP	Speculative analogy.	
52	Cosmic birefringence (EB/TB correlations)	Methylation-state rotation during disease (5mC → 5hmC → 5fC → 5caC oxidation cascade is literally a rotation in methylation chemistry-space). Would produce equivalent cross-correlation signatures.	REINTERP	Genuinely novel research direction.	

Section XII — Non-Gaussianity / higher-order statistics

#	CMB module	CPG analog	status	author's note	2026-09-19 status
53	Bispectrum estimators (local, equilateral, orthogonal, etc.)	TODO 2.2 bispectrum on highest-A-score class markers. Different "shapes" of bispectrum probe different biological coupling structures (local = pairwise regulatory networks; equilateral = shared upstream regulator; orthogonal = anti-correlated regulation).	ROADMAP	On the roadmap. The shape taxonomy will need its own translation.	
54	Trispectrum (4-point)	Coordinated 4-CpG methylation events — chromatin remodeling complexes, multi-enhancer hubs.	ADD	Tier 2+.	
55	Minkowski functional s / topology	Genome-wide methylation topology — how methylation is spatially organized within the nucleus (A/B compartments, TADs, LADs).	ADD	Connects to Hi-C data — multi-omics cross.	
56	Phase statistics	CpG phase structure — same amplitude, different ordering = different biology. Phase information is what the bispectrum sees that the power spectrum misses.	ADD	Tier 2+.	

Section XIII — Isotropy, anomalies, preferred axes

#	CMB module	CPG analog	status	author's note	2026-09-19 status
5 7	Statistical isotropy tests (bipolar harmonics, dipole modulation, hemispherical asymmetry, local variance maps, multipole vectors)	CPG-VAL-006 chromosome isotropy is the seed. Extend to bipolar (paternal vs maternal allele asymmetry), dipole modulation, hemispherical asymmetry (genomic half-asymmetry).	HAVE→RO ADMAP	Have seed; extend.	
5 8	Parity asymmetry	Strand-bias parity tests + even/odd-chromosome parity tests.	ADD	Niche but specifically testable.	
5 9	Alignment tests (axis of evil, multipole vectors)	Chromosomal arm alignment — do disease-signal-rich chromosomes share an axis (e.g., chr6 MHC + chr2 immune-related = immune-axis alignment)?	REINTERP	Need to be careful with look-elsewhere.	
6 0	Topology searches (matched circles)	Methylation "matched-pair" searches between distant loci with co-regulated patterns.	REINTERP	Speculative.	
6 1	Look-elsewhere discipline	MANDATORY. Pre-register the test before running it. We've been informal about this.	ADD	Make absolute rule in v1_VAL_Test_Checklist.	BUILT — N8 in the null runner; VAL-006's chr6 signal died under it

Section XIV — Secondary anisotropies

#	CMB module	CPG analog	status	author's note	2026-09-19 status
6 2	Integrated Sachs-Wo lfe effect	Late-life methylation drift integrated along developmental trajectory. Same structure: late-time potential decay → large-angle imprint.	REINTERP	Speculative analogy worth exploring.	
6 3	Rees-Sciama	Nonlinear corrections. Niche.	NO	Skip.	
6 4	Thermal SZ (hot electron Comptonizat ion)	Inflammation-associated methylation shifts — chronic inflammation reshapes the methylome the way hot cluster gas reshapes the CMB.	REINTERP	Real biological analog.	
6 5	Kinetic SZ (bulk velocities)	Clonal hematopoiesis bulk-velocity-like signal — large clones with shared methylation drift create a kSZ-like signature.	REINTERP	Specific to CHIP/AML cards.	
6 6	Patchy reionization	Patchy senescence / cellular replacement in aging tissues.	REINTERP	Aging-related.	
6 7	Cosmic infrared background	Background methylation noise from healthy tissue turnover.	ADD	Worth modeling.	
6 8	Rayleigh scattering	Frequency-dependent scattering.	NO	Skip.	

#	CMB module	CPG analog	status	author's note	2026-09-19 status
69	Spectral distortions (μ , y -type)	Methylation spectral distortions: systematic deviations from the IAM-predicted information spectrum at specific scales = signatures of disease-specific energy injection into the methylome.	REINTERP	Genuinely novel research direction.	

Section XV — Data combination & external probes

#	CMB module	CPG analog	status	author's note	2026-09-19 status
70	BAO combination	Multi-substrate combination (RNA-seq, fragmentomics, ATAC-seq, proteomics) as the methylome's BAO — breaks degeneracies CMB-only / methylation-only cannot resolve.	ROADMAP	Tier 5.	
71	Supernovae (distance ladder)	Cellular age clock as the distance ladder — calibrate biological age via known-age reference cohorts.	ROADMAP	TODO 2.4.	
72	BBN cross-check	Mendelian inheritance + germline methylation cross-checks — early-life patterns set bounds on later patterns.	REINTERP	Niche but principled.	
73	Large-scale structure	Cross-cohort population structure. UK Biobank scale eventually.	ROADMAP	Far future.	
74	Joint likelihoods (handle shared covariance)	Joint likelihood across cards for one patient — handle covariance between immune-class drift detected by breast-epic AND immune-epic.	ADD	Stage 5 architecture.	

Section XVI — Modern computational methods

#	CMB module	CPG analog	status	author's note	2026-09-19 status
75	Public compressed likelihoods	Per-card compressed-likelihood release. ACT-DR6 style. Future.	ADD	Once cards stabilize.	
76	Emulators	Neural-network emulator for IAMAtlas reconstruction at very large CpG coverage. Speed optimization.	ADD	Tier 4+.	
77	Simulation-based inference	Synthetic-patient SBI — when L7 likelihood is intractable.	ROADMAP	Pairs with end-to-end sims.	
78	Blinding	Mandatory for confirmation cards. Apply pipeline before seeing case labels. Unblind only after seal.	ADD	Add to v1_VAL_Test_Checklist.	PARTIAL — PREREG sealed before data opens; no formal blinding step
79	Reproducibility discipline	Already absolute — the reproducibility quadruple is our version.	HAVE	We have this.	BUILT — bit-identical Stage 1, deconvolver MAE 0.0004, anchors $r = 1.00000$; the PROC kit

APPENDIX VII — THE COMPLETION SPRINT, SCORED

The spring-2026 plan to take the chain "from C- to A" in 12–15 sessions: Phase A (L9 nulls) → B (L4 foregrounds) → C (L5 correlation structure) → D (L6 covariance) → E (L7–L8 likelihood and inference) → F (audit). Scored against the repository on 2026-09-19. The author's verdict: *"Too ambitious too quick. It ended up harming rather than helping. These should have been worked on long after the bones were trusted."*

phase	deliverable	what happened
A1	cpg_null_runner.py, nulls N1–N8	BUILT, live (CPG_Engine/CPG_Null_Runner/). Every post-Atlas CPG-VAL carries null_results.json. The one phase built on trusted bones; still the best-engineered part of the chain.
A2	synthetic patient generator	BUILT (RETIRED/.../Synthetic_Patient_Generator/; harness in CPG_Engine/report_builders/). N7 depends on it — to be restored beside the null runner.
A3	Family A VALs through the nulls	DONE (CPG-VAL-001...007 all carry null results).
B1	foreground registry	never built
B2	NILC second deconvolver	BUILT → CUT 2026-07-02 (c1be0c3): collapsed on correlated blood mixtures, deleted correct calls. Single method today, by decision.
B3/B4	age / sex / smoking foreground modules	BUILT → REFUSED (SOP §104): the methylome's foreground is the patient's own biology — annotate, never subtract. In RETIRED/.../IAM_Cellular_Age/.
C1	C(d) genomic-distance correlation, "acoustic peaks"	never built
C2	bispectrum	never built
C3	banana degeneracy mapper	never built ("I never got my banana degeneracy")
D1/D2	sim-based and per-CpG covariance	never built
E1	cellular age clock	BUILT as a trained clock (v1), later judged wrong (not physics); replaced by the band inversion (v3), which does not yet calibrate. Not reportable.
E2/E3	per-card likelihood + MCMC posteriors	never built; scoring remains threshold + age band
F1–F3	chain audit, v2 report, v2 checklist	report iterated v2 → v10; checklist v2 not written; the audit is what the 2026-09 PROC series is doing now

The lesson

Phase A worked because it tested a chain that already existed. Everything after it was built on bones not yet trusted: an atlas still going flat (see IAM_Atlas/IAMAtlas_FLATNESS_LESSON.md), a band compiled from nine pipelines, a gauge formula that changed four times between 2026-06-11 and 07-01. So B was built and torn out, and C–E had nothing to stand on. The right order was A → trust the bones (Stage 1 bit-identical; deconvolver conformant; anchors reproducing; the band rebuilt from one cohort through one pipeline — Phase 1 of the September plan) → then the correlation structure, the degeneracies, the Bayesian layer. The sprint had the right pieces in the wrong order. Its own sign-off question 5 named the risk — "the discipline doesn't hold when the first inconvenient result comes back" — and the discipline did hold for individual VALs (VAL-102 voided in four minutes, seal preserved); it did not hold for the sequence.

What the sprint still gets right, for later: C1 (two-point correlation per class with MASTER-style mode coupling), C3 (2D posterior shape for A-score pairs — the CIMP axis is the first), and E2/E3 (a per-card likelihood marginalised over composition and age, replacing thresholds) remain the correct next layer once Phase 1 has re-derived the bands. Not before.

FUTURE GOALS — WHAT IS WORTH THE EFFORT, IN ORDER

Drawn from the scored translation map (Appendix VI) and the scored sprint (Appendix VII): only items not yet built, kept only where the data and tools in hand can support them, each with the gate it waits on. The order is the lesson of Appendix VII applied — nothing above the bands is scheduled until the bands are rebuilt from one cohort through one pipeline. "Not now" is a list, not a refusal; "does not translate" is the author's own ruling.

gate	goal	from	why it is realistic	needs
GATE 0 — before anything else	Rebuild the eight class BANDS from one healthy cohort through one pipeline	sprint lesson; RECON B1	GSE87571 (n~730 healthy whole blood, ages 14–94, raw 450K IDATs) through Stage 1 as shipped. Replaces a band compiled from nine papers on nine pipelines. Resolves the stem_adult false alarm and the identity-loci offset in one step.	6 GB download; ~6 h of Stage 1; a sealed PREREG first
GATE 0	Presence rule: report a class gauge only where DETERMINED ($\kappa < \sim 10$, or joint) and PRESENT ($\geq 5\%$); whole blood → immune + one haematopoietic-progenitor component	PROC-NILC-01; author's risk-over-reward rule	The Atlas reference columns for the two classes are $r = +0.99$ at the deconvolver's markers; the split between them is chosen by the simplex constraint, and the stem_adult band then flags healthy donors. One joint haematopoietic-progenitor component removes the false alarm without touching any floor. κ of the substrate sub-problem joins the Jensen gap as a pre-flight number. stem_adult has carried zero findings in nine VALs; its band is $n = 4\text{--}32$ sorted HSC.	a config change in the conductor + the band rebuild of Phase 1
GATE 0	Wire the Stage 0 ↔ Stage 1 intensity hand-off	map rows 10–11; PROC-STAGE0-01	methy prep already exposes negative-control probes, per-probe detection p, bead counts and X/Y intensities; Stage 0 only has to receive them. Turns four DEFERRED QCs into real ones.	half a day; no new data
GATE 0	Reproduce the breast pre-diagnostic anchor from raw IDATs	sprint F1	GSE51057 (329 IDATs) through the whole chain against the rebuilt band. The first result a reader will check; it should stand on the corrected reference.	2.8 GB; ~3 h
after GATE 0	Cellular variance as cosmic variance	map row 3	A cfDNA sample carries a finite number of cell-equivalents; the variance floor that sets is computable from the deconvolved composition and the read depth. Names the fundamental limit on small plasma samples. Analytic; no new data.	a derivation and one figure
after GATE 0	Transfer function of the chain	map row 14; CCL-039	Inject a known per-CpG signal into synthetic patients (N6) and measure what fraction survives each stage — the deconvolver 'explains away' composition-like signal and this quantifies how much. The synthetic generator and null runner already exist.	N6/N7 runs; one table per stage
after GATE 0	Second, independent deconvolver — the right way this time	map row 20; sprint B2	Planck never trusted a single component-separation method. NILC failed on correlated blood mixtures; a parametric Bayesian (Commander-style) deconvolver on the same Atlas is the natural second, and it fails differently. Agreement within tolerance becomes a gate; disagreement becomes a flag, not a deletion. Lesson from the first attempt (RECON D3): NILC ran unconstrained GLS on the bimodal cell-type marker pool and was ill-conditioned on blood's near-parallel immune/progenitor/stem_adult columns; the replacement must solve on Walther's class-marker ruler with a simplex prior, so that it differs in principle rather than in stability.	1–2 sessions; the Atlas; Moss 2018 mixtures as ground truth
after GATE 0	Nuisance marginalisation in the gauge	map row 21	Composition and age enter as point estimates today. Propagating the deconvolver's residual and the band's sd into the gauge A gives every reading an uncertainty — the posterior sd the MCMC atlas already carries but the chain drops at Stage 4.	half a session; no new data
after the anchor	Banana degeneracy — 2D posterior shape for A-score pairs	map row 38; sprint C3 ('I never got my banana degeneracy')	For pairs such as immune × cycling or terminal × stem_pluri, map the 2D case-vs-HC distribution; the CIMP axis is the first known one. Cheap once bands are trusted; meaningless before.	one session; the foundation cohort

gate	goal	from	why it is realistic	needs
after the anchor	C(d): two-point correlation of residuals vs genomic distance, per class	map rows 17, 23, 25; sprint C1	The methylome's power spectrum. Look for characteristic-scale features (the 'acoustic peaks' of the map), with MASTER-style correction for masked CpGs. Needs CHR/MAPINFO on every CpG (the manifest is in IAM_Atlas/external_manifests).	one session; the Atlas + anchor cohort
after the anchor	Per-card likelihood, marginalised, with MCMC posteriors	map rows 28, 37; sprint E2/E3	Replace threshold-plus-band scoring with a proper posterior over per-card parameters, nuisance-marginalised. emcee is already in the toolchain from the Atlas build. This is what a Planck reader will expect L7 to be.	1–2 sessions; everything above
after the anchor	Formal blinding for confirmation VALs	map row 78	Apply the pipeline before case labels are seen; unblind only after the seal. PREREG already does half of this; the other half is a procedural rule and a script flag.	a checklist change
substrates	Urine, CSF, and the within-patient tissue/plasma/urine trio	Issue 003 §7	Cohorts already located and partly downloaded (GSE119260 four men × three substrates; GSE292312 and GSE269403 CSF). Each is one experiment with a declared prediction: the class present in the shed tissue should appear in its fluid and be absent from the same patient's blood.	data in hand; PREREG each
not now	Bispectrum / trispectrum; Minkowski functionals; isotropy and alignment tests; spectral distortions	map rows 45, 53–60, 69	Real analogs, genuinely novel, and every one of them needs a trusted two-point function first. Listed so they are not forgotten; not scheduled.	after C(d)
not now	Multi-substrate 5mC/5hmC (E/B separation); multi-omics cross-correlation	map rows 4, 49, 70	Tier 5. Needs oxBS or matched RNA-seq/ATAC on the same samples — data the project does not hold.	new data
does not translate	Rees–Sciama; Rayleigh scattering	map rows 63, 68	The author's own <i>X</i> rows; kept for completeness.	—

PART II — THE CHAIN, STAGE BY STAGE (FORTHCOMING)

Issue 003 documents the engine and its evidence. Part II, written once the chain carries its seal, will teach it: one chapter per stage, each giving the purpose, the cosmology it borrows, what was tried first and why it failed, the runtime files it reads, the procedure that confirms it, and what a researcher trained on bootstrapping needs to know before touching MCMC output or a Mahalanobis distance. The outline as it stands:

chapter	content
Opening — the translation map	Appendix VI as narrative: what the CMB gave us, row by row, and the two places we stopped taking it (a second deconvolver; de-aging).
Three kinds of file	FLOOR ($H_{\min}$: physics, 40 numbers) / RULER (identity loci: where the gauge reads) / BAND (age reference: a cohort statistic). Every 'healthy reads wrong' case of 2026 traced to a band.
Stage 0 — intake	Array type from the IDAT header; SHA-256 with re-transmission detection; the fail-open we closed.
Stage 1 — calibration	noob via methylprep; bit-identical on 11/11; the 6-7% normalisation gain between atlas-source betas and Stage 1 output and why the band absorbs it.
Stage 2 — component separation	Walther NNLS against the Atlas; presence, not gating; why NILC was cut.
Stage 3 — the foreground we refused to subtract	Built age/sex/smoking layers; SOP s104; a galactic foreground is a separate source, the methylome's is the patient.
Stage 4 — the gauge	$H(\text{beta_mean})/H_{\min}$ on identity loci, placed in the band; two surfaces (s106); four formula changes in three weeks and the measurement that settled them.
Stage 4.5 — bidirectional	Shannon H is symmetric about 0.5; VAL-050's null and VAL-051's recovery.
Stage 4.6 — the patient sky	$z = (\text{beta} - \mu)/\sigma$ on HEALPix; the residual map against a reference with per-pixel uncertainty; Plate 05.
Stage 5 — Mahalanobis Option A	Eight coefficients as a vector; distance from the age-band centroid; why v1 put healthy at $z \sim -22$.
Stage 6 — cellular age	The band run backwards; three attempts; why it does not calibrate yet.
The Atlas (several chapters)	Per-class MCMC over 115 cell types; the flatness lesson (R-hat misled twice, in opposite directions); the brightness posterior; why a sphere.
What we built before the bones were trusted	Appendix VII as narrative; the order that should have been.
For the bootstrapper	MCMC vs resampling; what a posterior sd buys that a bootstrap CI does not; when a Mahalanobis distance is and is not a p-value.

MASTER PREDICTIONS TABLE — all G-2026 filings

Every dated prediction in this publication consolidated into one index. Each is numbered, dated, specific, and falsifiable. Entries marked PENDING will be revisited in subsequent issues as outcome data becomes available.

Class	ID	Filed	Status	Claim summary
Terminal	G-2026-P006	Originally f	CONFIRMED (partial)	In longitudinal cohorts with archived CSF samples and subsequent Alzheimer's disease diagnosis, the terminal-class A-score will show elevation above 1...
Terminal	G-2026-P015	April 2026	PENDING	In longitudinal cohorts of patients with subsequently-diagnosed amyotrophic lateral sclerosis (ALS), the terminal-class A-score from CSF cfDNA will sh...
Secretory	G-2026-P003	Originally f	PENDING	Among PSA-screened prostate cancer patients with matched methylation + WPS + DELFI profiles at diagnosis and at 24 months, the combined A-score trajec...
Secretory	G-2026-P004	Originally f	PENDING	In asbestos-exposed occupational cohorts with archived serial blood samples, the secretory-class combined A-score will show elevation above 1.05 at le...
Immune	G-2026-P010	April 2026	PENDING	In prospective cohorts of patients with known clonal hematopoiesis of indeterminate potential (CHIP) and archived serial blood samples, the immune-cla...
Immune	G-2026-P011	April 2026	PENDING	In patients receiving immune checkpoint inhibitor therapy, the immune-class A-score trajectory will distinguish responders from non-responders within ...
Progenitor	G-2026-P014	April 2026	PENDING	In MDS patients receiving hypomethylating agent (HMA) therapy, the progenitor-class combined A-score trajectory at 3 months post-initiation will ident...
Cycling	G-2026-P001	April 2026	PENDING	In a prospective screening cohort of average-risk adults aged 45-75 with archived serial blood samples and colonoscopy outcomes, the cycling-class com...
Cycling	G-2026-P002	April 2026	PENDING	In longitudinal cohorts of Barrett's esophagus patients with known progression outcomes, the cycling-class combined A-score will show elevation above ...
Stromal	G-2026-P012	April 2026	PENDING	In prospective IPF cohorts with archived serial blood samples and known progression outcomes, the stromal-class combined A-score trajectory slope will...
Adult Stem	G-2026-P013	April 2026	PENDING	In prospective CHIP/CCUS cohorts with archived serial blood samples and known progression outcomes, the stem_adult A_active (2/5, methyl + frag only) ...
Adult Stem	G-2026-P014	April 2026	PENDING	Trajectory slope ($\Delta A/\Delta t$) on the saturating substrates (nucl, fuzz, wps) during the pre-saturation window will correlate with clinical outcome in adult...
Adult Stem	G-2026-P015	April 2026	PENDING	Multi-class A-score divergence signature — adult-stem class flatlined at ceiling while an adjacent architecture class (immune, cycling, or secretory d...
Pluripotent	G-2026-P005	Originally l	PENDING	In cryptorchidism patients and post-orchietomy TGCT survivors monitored prospectively with serial cfDNA sampling, the stem_pluri class will show a CH...
Pluripotent	G-2026-P016	April 2026	PENDING	In iPSC reprogramming protocols, the stem_pluri class combined A-score across at least four substrates measured at colony harvest will predict downstr...

Class	ID	Filed	Status	Claim summary
Pluripotent	G-2026-P017	April 2026	PENDING	In TGCT patients undergoing bleomycin-etoposide-cisplatin (BEP) chemotherapy, the trajectory of the stem_pluri A_methyl signal during the first two cy...

DATA SOURCES — PRIMARY CITATIONS

All β values from primary publications. No synthetic data. No proprietary datasets. All scripts and data pipelines are reproducible and archived at the Zenodo DOI.

Five-Substrate Framework

- [Li et al. 2024 Genome Med — MESA test \(4 substrates, colorectal, AUC 0.931, n=690\)](#)
- [Cristiano et al. 2019 Nature — DELFI fragmentomics \(7 cancers, AUC 0.940, n=208\)](#)
- [Mathios et al. 2022 Nat Commun — DELFI 2-year pre-diagnostic window](#)
- [Snyder et al. 2016 Cell — WPS original \(15/15 tissue types\)](#)
- [Doebley et al. 2022 Nat Commun — Griffin nucleosome occupancy \(breast, n=139\)](#)
- [Corces et al. 2018 Science — TCGA ATAC-seq pan-cancer \(23 types\)](#)
- [Esfahani et al. 2022 Cancer Discovery — nucleosome fuzziness \(prostate PDX, n=26\)](#)

Reference Cell Methylation (Per-Class H_{min})

- [Lister et al. 2009 Nature — hESC H1 \(stem_pluri reference\)](#)
- [Lister et al. 2013 Science — frontal cortex neuron \(terminal, global floor 0.756500\)](#)
- [Kozlenkov et al. 2014 Hum Mol Genet — cortical neuron mature](#)
- [Movassagh et al. 2011 NEJM — adult cardiomyocyte \(terminal\)](#)
- [Roadmap Epigenomics Consortium 2015 Nature — 127 reference epigenomes](#)
- [Hannum et al. 2013 Mol Cell — blood methylation aging \(GSE40279, immune\)](#)
- [Moss et al. 2018 Nat Commun — cfDNA tissue-of-origin atlas \(cfDNA % per class\)](#)

Cancer Validation (G-008, 27/28 TCGA types)

- [TCGA Pan-Cancer Atlas — 28 methylation datasets, all cancer types](#)
- [Ceccarelli et al. 2016 Cell — LGG/GBM \(terminal, largest \$\Delta A\$ in dataset\)](#)
- [TCGA Research Network 2013 NEJM — AML, Ley et al., n=200 \(immune, primary \$\beta\$ source\)](#)
- [Chapuy et al. 2018 Nat Med — DLBCL \(immune, n=48\)](#)
- [Radovich et al. 2018 Cancer Cell — TCGA Thymoma \(immune, n=120\)](#)
- [TCGA Research Network 2012 Nature — BRCA, n=825 \(secretory, primary \$\beta\$ source\)](#)
- [TCGA Research Network 2015 Cell — PRAD prostate molecular taxonomy, n=333 \(secretory\)](#)
- [TCGA Research Network 2017 Cancer Cell — PAAD pancreatic adenocarcinoma \(secretory\)](#)
- [TCGA Research Network 2017 Cell — LIHC hepatocellular carcinoma \(secretory\)](#)
- [Zheng et al. 2016 Cancer Cell — TCGA ACC adrenocortical carcinoma \(secretory\)](#)
- [Robertson et al. 2017 Cancer Cell — UVM \(stem_adult\)](#)
- [All TCGA citations per-cancer-type in individual card tables](#)

Cycling Epithelial Cancers — 14 TCGA Types (Card #5)

- [TCGA Research Network 2012 Nature — COAD + READ colorectal, n=276](#)
- [TCGA Research Network 2014 Nature — LUAD lung adenocarcinoma, n=230](#)
- [TCGA Research Network 2012 Nature — LUSC lung squamous, n=178](#)
- [TCGA Research Network 2014 Nature — BLCA bladder urothelial, n=131](#)
- [TCGA Research Network 2011 Nature — OV ovarian serous, n=489](#)
- [TCGA Research Network 2014 Nature — STAD stomach adenocarcinoma, n=295](#)
- [TCGA Research Network 2017 Nature — CESC cervical squamous, n=228](#)
- [TCGA Research Network 2015 Nature — HNSC head/neck squamous, n=504](#)
- [TCGA Research Network 2013 Nature — KIRC kidney clear cell, n=418](#)
- [TCGA Research Network 2016 NEJM — KIRP kidney papillary, n=161](#)
- [TCGA Research Network 2015 Cell — SKCM skin cutaneous melanoma, n=477](#)

- [TCGA Research Network 2014 Cell — THCA thyroid carcinoma, n=496](#)
- [TCGA Research Network 2013 Nature — UCEC endometrial carcinoma, n=373](#)

Stromal-Class Cancers — Sarcomas & Mesothelioma (Card #6)

- [Abeshouse et al. 2017 Cell — TCGA SARC soft-tissue sarcoma, n=206 \(adult\)](#)
- [Hmeliak et al. 2018 Cancer Discov — TCGA MESO mesothelioma, n=74](#)
- [Crompton et al. 2014 Cancer Discov — Ewing sarcoma pediatric, n=112 WGS](#)
- [Shern et al. 2014 Cancer Discov — rhabdomyosarcoma pediatric, n=147](#)
- [Tirode et al. 2014 Cancer Discov — Ewing sarcoma aggressive subtype, n=299](#)

Adult-Stem Cancers — HSC-Origin AML, BCC, MCC, Cholangiocarcinoma (Card #7)

- [Adelman et al. 2019 Cancer Discov — HSC-enriched aging methylation, n=5-7 per age](#)
- [Beerman et al. 2013 Cell Stem Cell — HSC aging methylation \(GSE44117\)](#)
- [Harms et al. 2015 Cancer Res — Merkel cell carcinoma MCPyV-negative, n=49](#)
- [Farshidfar et al. 2017 Cell Reports — TCGA CHOL cholangiocarcinoma, n=38](#)
- [Ley et al. 2013 NEJM — TCGA AML \(progenitor-lineage\), n=200 \(distinct from HSC-origin\)](#)

Pluripotent-Stem Cancers — TGCT (Card #8)

- [Shen et al. 2018 Cell Reports — TCGA TGCT integrated molecular characterization, n=137](#)
- [Killian et al. 2016 Genome Research — TGCT methylation reprogramming \(pure histology, n=130\)](#)

Non-Cancer Disease Validation & Pre-Malignant States

- [De Jager et al. 2014 Nat Neurosci — ROSMAP AD methylation, n=740 \(terminal\)](#)
- [Shireby et al. 2022 Brain — Brains for Dementia Research AD, n=631 \(terminal\)](#)
- [Lunnon et al. 2014 Nat Neurosci — AD four brain regions \(terminal\)](#)
- [Terman 2004 Age \(Dordr\) — cardiac aging, oxidative stress \(terminal\)](#)
- [Langston & Ballard 1983 N Engl J Med — MPTP Parkinson's disease \(terminal\)](#)
- [Wang et al. 2014 Nat Neurosci — PD methylation changes \(terminal\)](#)
- [Ahrens et al. 2013 Nat Commun — NAFLD hepatocyte methylation \(secretory\)](#)
- [Volkmar et al. 2012 EMBO J — pancreatic \$\beta\$ -cell T2D \(secretory\)](#)
- [Cruickshanks et al. 2013 Nat Cell Biol — senescent fibroblasts \(stromal\)](#)
- [Steensma et al. 2015 Blood — CHIP clonal hematopoiesis definition \(progenitor\)](#)
- [Malcovati et al. 2017 Blood — CCUS diagnostic criteria \(progenitor\)](#)
- [Jiang et al. 2020 Cell Death Dis — MDS whole-genome methylation progression](#)

Progenitor-Lineage Cancers (Pediatric & Neural Progenitor Origin)

- [Nordlund et al. 2013 Genome Biol — pediatric ALL methylome, n=764 \(B-ALL + T-ALL\)](#)
- [Figueroa et al. 2013 J Clin Invest — integrated epigenetic analysis pediatric ALL](#)
- [Northcott et al. 2017 Nature — medulloblastoma whole-genome, n=491 WGS + 1,256 methylation](#)
- [Capper et al. 2018 Nature — DNA methylation CNS tumor classifier, n>2,800](#)
- [Schwalbe et al. 2017 Lancet Oncol — medulloblastoma methylation subgrouping](#)

Vertebrate Lifespan (Nature Aging NATAGING-A13702)

- [Lowe et al. 2018 Genome Biol — 42 mammalian species methylation](#)
- [Lu et al. 2023 Nature Aging — 185-species pan-mammalian clock](#)
- [Haqhani et al. 2023 Science — mammalian methylation networks](#)
- [Wang et al. 2020 Cell Reports — Labrador retriever aging, n=104 dogs](#)
- [Wilkinson et al. 2021 Nat Commun — bat longevity](#)
- [Tacutu et al. 2018 Nucleic Acids Res — AnAge lifespan database](#)

IAM Framework Foundations

- [Mahaffey 2026 — Thermodynamic Operating Constraints \(cell thermodynamics paper\)](#)
- [Landauer 1961 IBM J Res Dev — irreversibility and heat generation](#)
- [Jacobson 1995 Phys Rev Lett — Einstein equation of state from thermodynamics](#)
- [Zenodo DOI 10.5281/zenodo.19547624 — IAM-Validation repository](#)
- [GitHub: hmahaffeyges/IAM-Validation — all scripts archived](#)

Phase B Validation Cohorts — Post-Breach Trajectory Predictions

- [Stupp et al. 2005 NEJM — concurrent radiochemotherapy GBM protocol](#)
- [Gilbert et al. 2013 J Clin Oncol — RTOG 0525 dose-intensification trial \(n=833 GBM with MGMT stratification\). G-2026-P025 reanalysis target.](#)
- [DiNardo et al. 2020 NEJM — VIALE-A azacitidine+venetoclax in older AML \(n=431\). G-2026-P026b target.](#)
- [Taniguchi et al. 2021 Nat Med — GALAXY CRC MRD post-resection ctDNA \(n=1,000+ stage II-III\). G-2026-P027 target.](#)
- [Cristiano et al. 2019 Nature — DELFI lung cancer detection and pre-diagnostic window reanalysis framework for G-2026-P028 \(EGFR-mutant LUAD osimertinib resistance prediction\).](#)
- [OSIC IPF Biobank — Open Source Imaging Consortium, prospective IPF cohort with serial imaging and archived serum. G-2026-P012 target.](#)
- [Detels et al. 2012 Int J Epidemiol — MACS \(Multicenter AIDS Cohort Study\). 40+ year HIV longitudinal cohort with archived blood samples. G-2026-P026 target.](#)
- [Adimora et al. 2018 J Womens Health — WIHS \(Women's Interagency HIV Study\). Companion HIV cohort to MACS. G-2026-P026 target.](#)
- [Hardy et al. 2020 BMJ Open — ACTG \(AIDS Clinical Trials Group\) archived serial blood. G-2026-P026 target for ART response trajectory.](#)
- [Sudlow et al. 2015 PLoS Med — UK Biobank cohort design and baseline characterization \(n~500,000\). G-2026-P035, P038 target.](#)
- [LITMUS Consortium — Liver Investigation: Testing Marker Utility in Steatohepatitis. European NAFLD/NASH biomarker cohort, n=2,000+. G-2026-P032 target.](#)
- [Krop et al. 2019 Clin Cancer Res — BRCA DCIS molecular stratification framework informing G-2026-P030.](#)
- [COG \(Children's Oncology Group\) — archived pediatric sarcoma protocols \(Ewing VDC-IE, rhabdomyosarcoma VAC\). G-2026-P034 target.](#)

Post-Breach Trajectory Validation Cohorts (G-2026-P023 through P040)

- [Stupp et al. 2005 NEJM — GBM concurrent chemoradiation protocol \(Stupp protocol\)](#)
- [Gilbert et al. 2013 J Clin Oncol — RTOG 0525 dose-intensification trial \(n=833 GBM\)](#)
- [DiNardo et al. 2020 NEJM — VIALE-A azacitidine+venetoclax trial \(n=431 AML\)](#)
- [Kotani et al. 2023 Nat Med — GALAXY CRC prospective ctDNA MRD cohort](#)
- [Cristiano et al. 2019 Nature — DELFI 7-cancer AUC 0.940 validation \(cross-reference\)](#)
- [OSIC \(Open Source Imaging Consortium\) IPF Biobank](#)
- [MACS \(Multicenter AIDS Cohort Study\) — NIH-funded longitudinal HIV cohort](#)
- [WIHS \(Women's Interagency HIV Study\) — NIH-funded longitudinal HIV cohort](#)
- [ACTG \(AIDS Clinical Trials Group\) — NIH HIV therapeutics trials network](#)
- [LITMUS \(Liver Investigation: Testing Marker Utility in Steatohepatitis\)](#)
- [UK Biobank — n=500,000 prospective cohort with longitudinal blood and outcomes](#)
- [CHIP / CCUS / MDS — Steensma 2015 Blood \(CHIP definition, cross-reference above\)](#)
- [ELN 2022 risk classification — Döhner et al. Blood 2022 \(AML prognostic\)](#)
- [DELFI validation — Mathios et al. 2022 Nat Commun \(pre-diagnostic window, cross-reference\)](#)
- [ROSMAP longitudinal AD cohort — De Jager 2014 Nat Neurosci \(cross-reference above\)](#)
- [TCGA ATAC-seq / Corces 2018 — chromatin accessibility reanalysis target \(G-2026-P023\)](#)

Multi-Class Drift Cascade (VAL-037 through VAL-046, April 2026)

- [Zeng et al. 2026 Nat Cancer — pan-cancer plasma cfDNA compendium \(n=1,294, 14 types\). VAL-038 primary source.](#)
- [Liu et al. 2020 Ann Oncol — GRAIL multi-cancer detection with tissue-of-origin localization \(n=6,689\). VAL-041 deconvolution reference.](#)
- [Kadota et al. 2014 Am J Respir Crit Care Med — lung adenocarcinoma distance-annotated methylation series. VAL-039 lung source.](#)
- [Teschendorff et al. 2016 Genome Med — breast adjacent-normal field defect methylation outliers. VAL-039 breast source.](#)

- [Shen et al. 2005 Cancer Res — MGMT promoter methylation field defect in sporadic CRC. VAL-039 colon source.](#)
- [Damaschke et al. 2017 Cancer Epidemiol Biomarkers Prev — prostate zonal methylation gradient. VAL-039 prostate source.](#)
- [Villanueva et al. 2015 Hepatology — HCC methylation-based prognosis with cirrhotic gradient. VAL-039 HCC source.](#)
- [Kang et al. 2008 Am J Pathol — gastric carcinoma intestinal metaplasia field effect. VAL-039 gastric source.](#)
- [De Jager et al. 2014 Nat Neurosci — ROSMAP AD cortex methylation \(n=708\). VAL-040 primary source.](#)
- [Shireby et al. 2022 Brain — Brains for Dementia Research AD cortex methylation \(n=1,408\). VAL-040 primary source.](#)
- [Nabais et al. 2021 Genome Biol — peripheral blood AD meta-analysis \(n=3,424\). VAL-040 immune-class AD source.](#)
- [Lunnon et al. 2014 Nat Neurosci — entorhinal cortex AD methylation. VAL-040 temporal cortex source.](#)
- [Widschwendter et al. 2021 Cell Rep Med — WID-CIN cervical progression \(n=2,254\). VAL-042 cervical source.](#)
- [Jammula et al. 2020 Gastroenterology — Barrett's esophagus methylation progression subtypes. VAL-042 esophageal source.](#)
- [Jerónimo et al. 2008 Clin Cancer Res — prostate PIN→HGPIN→metastatic progression. VAL-042 prostate source.](#)
- [Luo et al. 2014 Gastroenterology — colon adenoma-carcinoma sequence methylation pathways. VAL-042 colon source.](#)
- [Yoshizato et al. 2020 Blood — CHIP→MDS→AML clonal evolution. VAL-042 AML source.](#)
- [Pal et al. 2016 Cancer Res — canine mammary tumor methylation. VAL-043 canine source.](#)
- [Beck et al. 2020 Vet Comp Oncol — canine diffuse large B-cell lymphoma methylation. VAL-043 canine source.](#)
- [Decker et al. 2015 PLoS Genet — canine bladder TCC BRAF V600E mutation. VAL-043 canine source.](#)
- [Parikh et al. 2019 Nat Med — CRC FOLFOX resistance ctDNA serial. VAL-044 CRC source.](#)
- [Stover et al. 2018 J Clin Oncol — metastatic TNBC ctDNA trajectory. VAL-044 BRCA source.](#)
- [Cabel et al. 2018 Ann Oncol — melanoma anti-PD-1 ctDNA monitoring. VAL-044 melanoma source.](#)
- [Shen et al. 2018 Cell — plasma cfDNA methylome pan-cancer. VAL-045 TGCT reference.](#)
- [Killian et al. 2016 Cell Rep — seminoma hypomethylation and germ cell tumor epigenetics. VAL-045 seminoma source.](#)
- [Kresovich et al. 2019 JNCI — Sister Study pre-diagnostic breast cancer \(n=2,776\). VAL-046 capstone source.](#)
- [Hillary et al. 2020 Clin Epigenetics — UK Biobank methylation pre-diagnostic. VAL-046 UK Biobank source.](#)
- [Horvath et al. 2014 Genome Biol — Health ABC liver methylation aging. VAL-046 Health ABC source.](#)
- [Hou et al. 2012 Am J Epidemiol — Nurses' Health LINE-1 pre-diagnostic colorectal. VAL-046 Nurses source.](#)
- [Horvath 2015 Aging — Rotterdam Study aging-methylation. VAL-046 Rotterdam source.](#)
- [Alisch et al. 2012 Genome Res — pediatric age-associated methylation. Healthy baseline reference source.](#)
- [Horvath 2013 Genome Biol — 353-CpG multi-tissue epigenetic clock. Healthy baseline reference source.](#)

GLOSSARY

Definitions of the technical terms used throughout this publication, organized by category: Physics & Derivation, Biology & Substrates, Statistical & Validation, and Framework & Clinical Terms. Every term traces to either a published primary source or a derived quantity in the framework.

PHYSICS & DERIVATION

A-score	Dimensionless ratio $A = H(\text{value}) / H_{\min}(\text{class, substrate})$. Formula is unfloored — no $\max()$ clamp. The central measurement in GAPE. $A = 1.00$ marks the architectural commitment point (reference $H_{\min}$). Healthy cells typically read $A \approx 0.95\text{--}1.00$ due to normal biological variation around the MCMC-derived $H_{\min}$. $A > 1.00$ means an accessible entropy gap (C3) has opened above the floor — standard cancer elevation. $A < 0.95$ is inversion territory (documented in seminoma, senescent fibroblasts, aged HSCs). See Section 2.1b for the full axis treatment.
A_ceiling	The maximum A-score physically achievable on a given class×substrate pairing: $A_{\text{ceiling}} = 1 / H_{\min}(\text{class, substrate})$. Reached when $\beta = 0.5$ (maximum binary entropy $H = 1.0$). No sample biology can produce a higher A on that substrate.
A_combined	AUC-weighted combined A-score across all five substrates: $A_{\text{combined}} = \Sigma(AUC_i \times A_i) / \Sigma(AUC_i)$. The headline score. Reduces measurement noise by roughly $\sqrt{5}$ across five correlated physical windows.
A_active	AUC-weighted mean over only the non-saturated substrates for a specific sample. A substrate is flagged saturated when A is within 0.005 of its ceiling. A_{active} is the signal that continues to respond after some substrates have pinned at ceiling — central for chemotherapy response trajectories and reserve-depletion signals.
C1, C2, C3	Three-component decomposition of Shannon entropy. C1 = universal Landauer floor (0.756500, neuron reference). C2 = class-specific architecture overhead above C1. C3 = accessible gap above $H_{\min} = \max(0, H - H_{\min})$. Healthy cells: C3 ≈ 0 . Cancer: C3 = 8-15%.
f_C3	Fraction of total entropy in the accessible gap: $f_{C3} = C3 / H(\beta)$. The component on which every therapeutic intervention operates. Rising f_{C3} under treatment indicates non-response. Declining f_{C3} indicates response.
H(β) — Shannon entropy	Shannon binary entropy of a Bernoulli(β) variable: $H(\beta) = -\beta \cdot \log_2(\beta) - (1-\beta) \cdot \log_2(1-\beta)$. $H(0) = H(1) = 0$. $H(0.5) = 1$ (maximum).
H_min	The class-specific and substrate-specific minimum Shannon entropy consistent with maintained cell identity. Derived from the Landauer cost of DNMT1 maintenance across 19.6M CpG sites per cell division. Calibrated via MCMC posteriors (G-002 methyl, G-003b four other substrates).
H_MIN_GLOBAL	The lowest $H_{\min}$ across all classes: 0.756500, from frontal cortex neuron (Lister 2013). Anchors the global C1 Landauer floor.
Landauer cost	Thermodynamic minimum energy to erase one bit of information: $E = k_B \times T \times \ln(2)$. At 37°C (310.15 K): 2.97×10^{21} J per bit. Landauer 1961 IBM J Res Dev.
Mahaffey Number	Dimensionless metabolic sensitivity parameter $n_{\text{bio}} = \Delta G_{\text{ATP}} / (R \cdot T_{\text{body}})$. Value 20.94 at 37°C. Absolute class-specific values await G-007 MCMC.
Reserve (clinical meaning)	The C3 headroom remaining between a cell's current A-score and its class-substrate ceiling. Declining reserve under treatment indicates the epigenomic capacity for further response is being consumed.

BIOLOGY & SUBSTRATES

β (beta value)	Fraction methylated at a CpG site or averaged across an architecture-class locus panel. $\beta \in [0,1]$. Primary input to the methylation $H(\beta)$.
cfDNA	Cell-free DNA circulating in plasma. Shed from apoptotic and necrotic cells across all tissues, weighted by tissue turnover. ~70% immune, ~12% cycling epithelial, ~8% secretory in standard venous draw.
CpG / CpH	CpG: a cytosine followed by a guanine on the same strand — the standard site of mammalian DNA methylation. CpH: cytosine followed by any other base — mostly unmethylated in somatic cells but methylated in neurons and embryonic stem cells.
CHIP / CCUS / MDS	Clonal Hematopoiesis of Indeterminate Potential: age-associated expansion of HSC clones carrying leukemia-associated mutations; typically benign. Clonal Cytopenias of Undetermined Significance: CHIP with unexplained cytopenias. Myelodysplastic Syndromes: frank malignant transformation. The framework's prediction G-2026-P013 targets the CHIP→MDS transition.
DELFI	DNA Evaluation of Fragments for early Interception. Cristiano et al. 2019 Nature. Measures short-fragment fraction (100-150bp / total) in cfDNA WGS. Substrate 5 in the five-substrate framework.
DNMT1	DNA methyltransferase 1. Maintenance enzyme that copies methylation pattern from parent to daughter strand during DNA replication. Finite fidelity (~1 error per $10^5\text{--}10^6$ CpG sites per division) — the biological rate-limit on $H_{\min}$.
Fragment size	cfDNA fragment length distribution. Healthy fragments cluster around 166 bp (one nucleosome + linker). Tumor-derived cfDNA over-represents 100-150 bp fragments due to altered nucleosome spacing. Substrate 5.
Hypomethylation (seminoma)	Global reduction of methylation β toward zero. In seminoma, β falls toward 0.17-0.20 as malignant germ cells revert toward the primordial germ cell (PGC) state. Drives the Seminoma Hypomethylation Inversion.

Nucleosome occupancy	Mean probability that a given genomic position is wrapped in a histone octamer. Measured via ATAC-seq or MESA. Substrate 2.
Nucleosome fuzziness	Positional precision of nucleosomes relative to their reference positions (0 = precise, 1 = maximally fuzzy). Measured via NucleoATAC. Substrate 3.
PGC	Primordial Germ Cell. The most pluripotent somatic cell state known, with near-zero methylation. Seminomas revert toward this state, producing the inversion signature.
Pluripotent / GCNIS / TGCT	Pluripotent: capable of differentiating into all three germ layers. GCNIS: Germ Cell Neoplasia In Situ, the preinvasive precursor to all postpubertal TGCT. TGCT: Testicular Germ Cell Tumor.
WPS (Windowed Protection Score)	Fraction of cfDNA reads whose endpoints fall within a nucleosome-protected window at a given genomic position. Snyder et al. 2016 Cell. Substrate 4.
Yamanaka factors	Four transcription factors (OCT4, SOX2, KLF4, c-MYC) that reprogram somatic cells back to induced pluripotent state (Yamanaka 2006). Work by driving methylation toward high-entropy pluripotent configuration. Pharmacologic dose-response is non-monotone — excess dose produces the Differentiation Dose Inversion (Section 2.4).
IDH mutation / 2-hydroxyglutamate rate	Isocitrate dehydrogenase (IDH1/IDH2) mutations occur in ≥80% of lower-grade gliomas and produce the oncometabolite 2-hydroxyglutarate, which inhibits TET enzymes. Result: the hyper-methylator phenotype (G-CIMP) characteristic of LGG. Key driver of terminal-class post-breach hyperentropy signature.
TET enzymes	Ten-Eleven Translocation methylcytosine dioxygenases (TET1/2/3) actively demethylate 5-methylcytosine. Inhibition by oncometabolites (e.g., 2-hydroxyglutarate in IDH-mutant tumors) drives aberrant hypermethylation. Counterpart to DNMT1 maintenance methylation.
G-CIMP	Glioma CpG Island Methylator Phenotype. Defined by Ceccarelli 2016 Cell in TCGA LGG/GBM analysis. Hyper-methylation pattern characteristic of IDH-mutant gliomas. Framework reads G-CIMP as terminal-class A-score elevation driven by methylation redistribution rather than global loss.
MGMT promoter methylation	Methylation status of the O6-methylguanine-DNA methyltransferase (MGMT) gene promoter. MGMT-methylated GBM responds better to temozolomide. Framework prediction G-2026-P025: MGMT status stratifies post-breach A_active trajectory in GBM on Stupp protocol.
T-cell exhaustion / PD-1	Progressive loss of effector function in chronically stimulated T cells. Marked by sustained PD-1 expression and compromised DNMT1 maintenance fidelity. Framework predicts detectable immune-class A-score elevation in chronic viral infection (HIV), senescent CD8+ populations, and tumor-infiltrating lymphocytes in checkpoint-blockade-eligible cancers.

STATISTICAL & VALIDATION

AUC	Area Under the receiver operating Curve. Single-substrate discrimination performance for binary classification (cancer vs healthy). Range 0.5 (chance) to 1.0 (perfect).
Bootstrap CI	Confidence interval derived by resampling. For H_min posteriors, leave-one-reference-out bootstrap confirms no single dataset drives the result. Mean Δ between full-data and leave-one-out posteriors: 0.168% across 40 class×substrate H_min values (VAL-031).
MCMC	Markov Chain Monte Carlo. Statistical method for sampling from posterior distributions. G-002 (methylation H_min): 17 chains, R-hat < 1.001, 8×10^5 samples. G-003b (four other substrates): same rigor.
Pre-diagnostic window	The time interval between a detectable framework signal and clinical diagnosis. Mathios 2022 Nat Commun established ~24 months for DELFI in lung cancer. Framework extension predicts 12-18 months for cycling-class solid tumors.
R-hat	Gelman-Rubin diagnostic for MCMC chain convergence. R-hat < 1.01 is convergence; < 1.001 indicates very tight convergence. All framework H_min posteriors satisfy R-hat < 1.001.
Sensitivity / specificity	Sensitivity: fraction of true positives correctly identified. Specificity: fraction of true negatives correctly identified. Framework targets high specificity (low false-positive rate) to support clinical serial monitoring rather than one-shot screening.

FRAMEWORK & CLINICAL

Architecture class	Cell-type grouping defined by shared thermodynamic floor (H_min). Eight classes: cycling, secretory, immune, progenitor, stromal, terminal, stem_adult, stem_pluri. Thermodynamically defined, not histologically.
BREACH	Tier assignment when $A \geq 1.10$. Indicates epigenomic architecture maintenance has failed beyond what metabolic intervention can restore. Structural transition.
Cross-class propagation	A cancer arising in one class producing signal in another class — e.g., breast metastasis to bone marrow producing immune-class elevation. Scenario 5.5; prediction G-2026-P022.
DETECTABLE	Tier assignment when $1.05 \leq A < 1.07$. Clinical workup recommended per class-specific differential.
Divergence signature	A multi-substrate pattern where substrates move in different directions — e.g., seminoma's A_methyl decreasing while A_nucl, A_wps, A_fuzz, A_frag increase. Primary detection signal for Pluripotent-class malignancy.

Floor	$H_{\min}(\text{class, substrate})$ — the minimum entropy consistent with the cell maintaining its class identity. Derived, not fitted.
Inversion	A named signature where A moves in the opposite direction from standard cancer elevation. Three identified: Seminoma Hypomethylation (Pluripotent), Differentiation Dose (Pluripotent research), Niche Depletion (Adult Stem). See Section 2.4.
MARGINAL	Tier assignment when $1.01 \leq A < 1.05$. Monitor; repeat at next scheduled timepoint.
Mahaffey value	Historical name for $H_{\min}$ in early drafts. Now standardized as $H_{\min}$. The class-and-substrate-specific entropy floor.
NORMAL	Tier assignment when $A < 1.01$. Consistent with age-expected healthy.
Saturation (runtime)	Sample-specific substrate flag: A within 0.005 of A_{ceiling} . Saturated substrates are excluded from A_{active} computation.
Saturation (structural)	Class-level property: a substrate whose ceiling is itself below BREACH. Cannot discriminate disease severity above BREACH on that substrate. Detection must rely on substrates whose ceilings exceed BREACH.
Substrate	One of five physical measurement windows: methylation, nucleosome occupancy, nucleosome fuzziness, WPS, fragment size. All five measure departures from the same underlying cellular identity information through different physical lenses.
Tier system	A-score thresholds: NORMAL ($A < 1.01$), MARGINAL (1.01-1.05), DETECTABLE (1.05-1.07), URGENT (1.07-1.10), BREACH ($A \geq 1.10$). All thresholds framework-derived; none fit to cancer data.
Trajectory (dA/dt)	Rate of change in A-score over serial samples. The framework's strongest clinical signal — a single elevated A may be age-consistent, but a rising A is not.
URGENT	Tier assignment when $1.07 \leq A < 1.10$. Workup indicated. Corresponds roughly to the Warburg transition in cycling-class cancer.
Warburg transition	Metabolic shift from OxPhos to aerobic glycolysis in malignant cells. Corresponds to $A \approx 1.07$ URGENT threshold. Past this point, metabolic levers cannot restore the cell to floor because glycolytic commitment is structural.
Warburg boundary (post-breach)	Qualitative post-breach zone boundary at $A \approx 1.15$, separating the metabolic-window zone (metabolic intervention still works) from the structural-only zone (glycolytic program locked in, glucose becomes neutral or harmful). See Section 2.6.
Glucose inversion	Qualitative post-breach boundary at $A \approx 1.25$ where adding glucose actively accelerates disease progression rather than supporting cellular function. Past this point, TPN/high-glucose supportive care is counterproductive. See Section 2.6.
Point of no return	Qualitative post-breach boundary at $A \approx 1.40+$ where epigenomic reserve (f_{C3} headroom) is depleted and no currently-available structural intervention can recover cellular function. Framework use shifts from therapeutic to prognostic. See Section 2.6.
Metabolic window zone	Post-breach zone at $A \approx 1.10$ -1.15 where metabolic intervention (DCA, 2-DG, ketosis) can still push cells back toward OxPhos. The widest therapeutic window post-breach.
Structural-only zone	Post-breach zone at $A \approx 1.15$ -1.25. Metabolic interventions have failed; structural interventions (DNMTi, HDACi, synthetic lethality, reprogramming) are the primary lever. Glucose becomes neutral to harmful.
Palliative range zone	Post-breach zone at $A \approx 1.25$ -1.40. Aggressive combination therapy with palliative intent; clinical trial eligibility becomes the framework's primary therapeutic route. Glucose supportive care actively harmful.
End of life zone	Post-breach zone at $A \approx 1.40+$. Cellular reserve depleted; framework use shifts from therapy to prognosis. Comfort care, symptom management, family-centered planning.
Post-breach trajectory	The time course of A-score and substrate divergence patterns past the ceiling at $A = 1.10$. Every class card has a Post-Breach Trajectory subsection with class-specific Known / Unknown / Test structure.
Z-score (age-adjusted)	$Z = (A_{\text{observed}} - A_{\text{predicted}}(\text{age, class})) / \sigma_{\text{cohort}}$. Standardized departure from the age-expected baseline. $ Z < 1$: age-expected. $ Z > 3$: strong signal warranting workup.

CLINICAL & TREATMENT PROTOCOLS

BEP	Bleomycin-Etoposide-Cisplatin. Standard chemotherapy regimen for advanced testicular germ cell tumors (TGCT). 95% cure rate at stage I-II. Framework prediction G-2026-P017 targets BEP response trajectory via A_{methyl} toward healthy hESC reference.
DCIS	Ductal Carcinoma In Situ. Non-invasive breast lesion diagnosed in ~50,000 US women annually. Current medicine cannot distinguish indolent from active DCIS at diagnosis; framework prediction G-2026-P030 targets this stratification with baseline A_{active} .
FOLFOX / FOLFIRI	Standard first-line chemotherapy regimens for metastatic colorectal cancer. FOLFOX: 5-FU + leucovorin + oxaliplatin. FOLFIRI: 5-FU + leucovorin + irinotecan.
G-CIMP	Glioma CpG Island Methylator Phenotype. Hypermethylator subtype of LGG/GBM defined by Ceccarelli 2016, characterized by IDH mutation and widespread promoter hypermethylation. Drives the terminal-class A-score elevation signal for glioma.
IDH mutation / 2-hydroxyglutamate	Isocitrate dehydrogenase 1/2 mutations produce the oncometabolite 2-hydroxyglutarate (2-HG), which inhibits α -ketoglutarate-dependent enzymes including TET demethylases. Drives G-CIMP hypermethylator phenotype in LGG ($\geq 80\%$ of cases) and a subset of GBM.

MGMT	O6-methylguanine-DNA methyltransferase. DNA repair enzyme. MGMT promoter methylation silences the gene, rendering GBM cells sensitive to temozolomide. Stratification biomarker for Stupp protocol response.
NASH / NAFLD	Non-Alcoholic Steatohepatitis / Non-Alcoholic Fatty Liver Disease. Secretory-class non-cancer failure mode. Framework reads NAFLD at $A \approx 1.015$ - 1.055 (MARGINAL to DETECTABLE) — detectable but pre-malignant. G-2026-P032 tests NAFLD-to-HCC progression via LITMUS cohort.
PD-1 / T-cell exhaustion	Programmed cell death protein 1. Immune checkpoint receptor upregulated on chronically activated T cells. T-cell exhaustion is the functional state characterized by high PD-1, loss of effector function, and compromised DNMT1 fidelity — directly relevant to HIV/AIDS prediction G-2026-P026.
RECIST	Response Evaluation Criteria in Solid Tumors. Radiographic response assessment standard. Framework predictions frequently benchmarked against RECIST at 6-month follow-up.
Stupp protocol	Standard GBM treatment: concurrent radiation + temozolomide, followed by six cycles of adjuvant temozolomide. Stupp 2005 NEJM. Framework prediction G-2026-P025 targets post-Stupp A_active trajectory for progression-free survival prediction.
TET enzymes	Ten-Eleven Translocation methylcytosine dioxygenases (TET1, TET2, TET3). Active DNA demethylation enzymes that convert 5-methylcytosine to 5-hydroxymethylcytosine. Inhibited by 2-hydroxyglutarate in IDH-mutant tumors, driving hypermethylation. TET2 mutations are the second most common CHIP/AML driver after DNMT3A.
Yamanaka factors	Oct4, Sox2, Klf4, and c-Myc (OSKM). Four transcription factors that convert somatic cells to induced pluripotent stem cells (iPSCs). Framework prediction G-2026-P016 targets successful-vs-aberrant reprogramming discrimination via the Differentiation Dose Inversion.

RESEARCH COHORTS & BIOBANKS

ACTG (AIDS Clinical Trials Group)	NIH-funded clinical trials network for HIV/AIDS therapeutics. Archived serial blood samples across multiple trials provide the longitudinal cohort for framework prediction G-2026-P026 (immune-class A trajectory under ART).
GALAXY (CRC MRD)	Japanese prospective CRC cohort with $n=1,000+$ stage II-III patients, serial cfDNA post-resection, and 2-year recurrence outcomes. Target cohort for framework prediction G-2026-P027 (post-resection A_active vs ctDNA-positive/negative classification).
LITMUS (NAFLD)	Liver Investigation: Testing Marker Utility in Steatohepatitis. European NAFLD Biomarkers Consortium, $n=2,000+$ biopsy-confirmed patients with longitudinal serum archives. Target for G-2026-P032 (NAFLD-to-HCC progression).
MACS / WIHS	Multicenter AIDS Cohort Study / Women's Interagency HIV Study. Longitudinal HIV cohorts with blood archives spanning 40+ years. Target for G-2026-P026 (immune-class A trajectory correlation with CD4+ recovery under ART).
OSIC (IPF Biobank)	Open Source Imaging Consortium IPF Biobank. Target for framework prediction G-2026-P012 (combined A-score trajectory distinguishing progressive from stable IPF with $AUC \geq 0.80$, outperforming serial FVC).
ROSMAP	Religious Orders Study and Memory Aging Project. Longitudinal AD cohort (De Jager 2014, $n=740$) with blood and brain samples archived plus subsequent AD diagnosis outcomes. Target for G-2026-P024 (MCI-to-AD conversion trajectory).
RTOG 0525	Radiation Therapy Oncology Group 0525 dose-intensification trial ($n=833$ GBM patients with MGMT stratification and serial imaging). Target for framework prediction G-2026-P025 (post-Stupp A_active trajectory for 6-month PFS prediction).
TCGA / GDC	The Cancer Genome Atlas / Genomic Data Commons. 28 cancer types methylation datasets, 23 cancer types ATAC-seq (Corces 2018). Primary source for all G-008 cancer validations in the framework.
UK Biobank	Prospective cohort of ~500,000 UK adults with baseline blood samples and health outcome follow-up. Target for G-2026-P035 (HSC aging to hematologic malignancy prediction) and G-2026-P038 (CHIP progression to MDS).
VIALE-A	Phase III trial of azacitidine + venetoclax vs azacitidine alone in older AML patients unfit for intensive chemotherapy ($n=431$). DiNardo 2020 NEJM. Target for G-2026-P026b (A_active decline slope predicting 12-month overall survival).

CASCADE TERMS (VAL-037 through VAL-046)

Adjacent-normal tissue (architecturally drifted)	Tissue immediately neighboring a tumor that appears histologically normal but sits at $\Delta A = +0.036$ above true-healthy reference on average across 24 TCGA cancer types (VAL-037). "Adjacent normal" in a pathology report is not architecturally healthy — it is drifted. Elevation extends 5-10 cm from the tumor margin (VAL-039 spatial gradient).
Architectural recovery	Observed decrease in A-score toward NORMAL tier in patients responding to cancer therapy (VAL-044, 5/5 clinical trials). Complete responders approach $A \approx 1.00$. Non-responders remain elevated. The clinical monitoring axis complementary to architectural drift.
Capstone finding (VAL-046)	The central multi-class drift hypothesis test: do future-cancer patients show baseline architectural elevation before clinical diagnosis? Result: across seven cohort-cancer combinations (Sister Study $n=2,776$ + six others), future-cancer participants show mean $\Delta A = +0.014$ above matched controls, detectable 2-5 years pre-diagnosis, across ≥ 2 architecture classes.

Class-universal inversion (pluripotent)	VAL-045 refinement: because $H_{\min_methyl} = 0.9822$ sits very close to the Shannon ceiling (1.000), the pluripotent methylation window above floor is too narrow to accommodate upward departure. All TGCT histologies land in A_{methyl} inversion territory; seminoma is the extreme case (divergence $2.1\times$ others). Specificity is in divergence magnitude, not direction.
Field-effect gradient	Spatially graded departure of A-score as a function of distance from the tumor. VAL-039 confirmed monotonic decay across 6 cancers: tumor $\rightarrow$ near-adjacent $\rightarrow$ far-adjacent (5-10 cm) $\rightarrow$ true-healthy, with far-adjacent tissue remaining elevated $\Delta A = +0.025$ above true-healthy.
Healthy baseline reference tables	80-cell reference (8 architecture classes $\times$ 10 age decades) compiled from Hannum 2013, Horvath 2013, Roadmap 2015, Moss 2018, Lister 2013, Alisch 2012. Each cell contains β_{mean} , β_{sd} , n_{samples} , A_{mean} , A_{sd} , and p10/p25/p50/p75/p90 percentiles. Anchors age-stratified interpretation of any patient A-score. Only terminal class crosses MARGINAL threshold ($A \geq 1.01$) within typical healthy lifespan (age 80-89).
Honest negative	A pre-specified prediction that fails in a way that confirms the framework's own prior finding in negative form. VAL-038 is the framework's honest negative: Spearman $\rho = -0.02$ between tissue-level predicted ΔA and bulk plasma alteration rate, confirming VAL-002's finding that plasma depends on shedding kinetics, not architectural departure, so deconvolution is required for plasma-based scoring. Failure is a feature that defines framework boundaries.
Multi-class drift signature	Coordinated elevation of A-score across ≥ 2 architecture classes in the same patient, observed in both future-cancer cohorts (VAL-046) and Alzheimer's disease (VAL-040). The framework's pre-diagnostic clinical signal: three classes elevated simultaneously (e.g., immune + secretory + stromal) is more specific than any single-class elevation alone.
Pre-diagnostic window	The 2-5 year interval before clinical diagnosis during which VAL-046 found detectable multi-class architectural elevation in future-cancer participants. Clinical implication: GAPE may function as a susceptibility flag years before conventional diagnostics, analogous to the role troponin plays for cardiac state.
Tissue-of-origin deconvolution	Computational method (Moss 2018, Liu 2020) that resolves bulk plasma cfDNA into per-tissue β values using tissue-specific methylation markers. When followed by per-tissue A-score against class-specific $H_{\min}$ (VAL-041), correctly identifies the primary cancer site in 10 of 10 cases (100% top-1 localization). The clinical bridge between honest bulk-plasma null (VAL-038) and functional plasma-based cancer detection.

GLOSSARY — CMB AND CHAIN TERMS

For a reader who knows bootstrapping but not MCMC, and has never met a Mahalanobis hull or a HEALPix sky. Each definition is taken from the source file named in brackets.

Cosmic Methylome Background (CMB) — The patient's per-CpG map, named for its cosmological twin. The chain is built on the Planck pipeline stage for stage: raw intensities -> calibration -> an all-sky map -> component separation -> a statistic against a derived reference -> null tests -> sky rendering. The one deliberate departure is foreground subtraction (see Foreground). [Biological_Physics/README.md]

Brilliance / brightness — The Atlas's per-CpG posterior mean and sd for a class - a reference map with a per-pixel uncertainty, which is what makes the methylome a CMB problem. 'Brilliance' is the author's term for the posterior surface; Stage 4.6 (brightness comparison) reads a patient's residual against it. [cpg_patient_cmb.py; IAMAtlasREBUILD.csv *_mean / *_sd columns]

z-departure — $z = (\text{beta_patient} - \mu_{\text{class}}) / \max(\sigma_{\text{class}}, 0.005)$ per CpG: how many posterior standard deviations the patient sits from the healthy class at that locus. The residual map of one observation against the reference. Median $|z|$ over a class was used for assessability; the deconvolver presence gate is the correct gate. [cpg_patient_cmb.py]

HEALPix / NSIDE / Mollweide — HEALPix: the equal-area sphere pixelisation used for CMB maps. NSIDE 128 -> 196,608 pixels. CpGs are laid on the sphere in genomic order (chr1 ... chrY, MAPINFO within chromosome), i-th CpG -> pixel floor($i \times \text{npix} / n$); ~2.46 CpGs per pixel, averaged. Mollweide is the equal-area full-sky projection every plate uses. Deterministic: same atlas + manifest -> byte-identical mapping. [IAM_Atlas/healpix_mapping/README]

Component separation — Cosmology's name for pulling a mixed signal apart into its sources; here, deconvolving a bulk sample into cell-type fractions (Stage 2). The Walther NNLS deconvolver is the one method in the chain; NILC (the Planck needlet method) was tried as a second and cut on 2026-07-02 after it collapsed on correlated blood mixtures. [walther_iam_deconvolver.py; commit c1be0c3]

Foreground — In cosmology, a separate physical source (the Galaxy) subtracted to reveal the CMB. In the methylome the 'foreground' - age, sex, smoking - is the patient's own biology, so it is annotated, not subtracted. The subtraction layers were built and deliberately left unwired (SOP s104). [RETIRED/.../IAM_Cellular_Age/]

Matched filter — Stage 5 second chain (SOP 8.2): the Pearson correlation between a patient's per-CpG departure and a disease's sealed residual template, with a bootstrap CI. Fires only when Stage 8 Route B flags; a template with no mean shift correctly does not match. [stage_5_second_chain.py]

Mahalanobis departure (Option A) — One number per patient: $\text{distance} = \sqrt{\sum_c ((A_c - \mu_c(\text{age})) / \sigma_c(\text{age}))^2}$ over the classes assessable in the substrate, where $\mu_c = A_{\text{mean}}(\text{class}, \text{age})$ and $\sigma_c = (A_{\text{mean}} - A_{\text{p10}}) / 1.2816$ from the age band. Answers 'is the whole profile off' when no single class is. A healthy patient sums to ~0. Alarm by n-adaptive chi-square. Supersedes v1_0 ($\mu=1.0$, $\sigma=0.02$ on the 115 separation A's, which put healthy at $z \sim -22$). [iamatlas_mahalanobis_scoring.py]

Null runner / the eight nulls — The methylome's CMB null tests: no CPG-VAL is sealed until its declared nulls pass. N1 label permutation (1000x); N2 age-strata permutation; N3 sex-strata; N4 cohort split-half replication with consistent sign; N5 plate / array-position; N6 injection-recovery (inject a known signal into HC, recover it); N7 end-to-end simulation (synthetic patients with known truth through L1-L8); N8 look-elsewhere (Bonferroni/FDR). Required for every VAL: N1, N4; for end-to-end claims: N7. 'A VAL that doesn't declare its nulls is not a sealed VAL - it's a preliminary analysis.' [CPG_Engine/CPG_Null_Runner/cpg_null_runner.py]

Synthetic patient — A constructed beta vector (or Stage-4 dict) with known class fractions and known per-cell A's, everything else absent. The ground truth for N6/N7: the chain must recover what was put in. [report_builders/synthetic_patient_harness.py; TEST_DATA/harness/]

PREREG / seal / RESTATE — Every validation declares its prediction, direction, pass conditions and nulls in a PREREG before data are opened; the file is sealed by SHA-256. Outcomes are recorded as pre-registered codes (O1_PRIMARY_VALIDATED ... O3_INVERTED, NULL, DIRECTIONAL). A change after seeing data is an AMENDMENT, dated; a re-reading is RESTATE; a post-hoc accommodation VOIDs the VAL with the seal preserved (VAL-102). [SOP Part IV; Testing_and_Code/VAL_INDEX.csv]

Chain of custody (L1-L9) — The audit grading of the chain: nine links from raw IDAT (L1) through calibration (L2-3), component separation (L4), scoring (L5-6), nulls (L7), report (L8-9), each with its own evidence requirement. The SOP is organised by stage (operational order) and by link (audit order). [SOP s2, s3]

Identity loci vs discriminative markers — Two CpG sets, two surfaces. Identity loci: where a healthy class sits within ± 0.05 of $H_{\text{min_beta}}$ (unimodal) - the GAUGE reads beta_mean here. Discriminative markers: ~100 one-vs-rest CpGs per cell type (bimodal by construction) - the SEPARATION statistic reads mean-of-per-CpG-H here. Using one set with the other's formula is the 2026-06-11 all-BREACH bug. [iamatlas_gauge_identity_loci_v1_0.json; iamatlas_celltype_markers_v0_2.json; s105-s106]

Jensen gap — $H(\text{beta_mean}) - \text{mean}_i H(\text{beta}_i) \geq 0$ (Jensen's inequality, H concave). Near zero on a unimodal panel (identity loci in blood: 0.013-0.033); large on a bimodal one (marker panels: ~0.7) or a tissue mixture (~0.2). The gauge refuses any panel whose gap exceeds 0.05 - the runtime form of s105. [cpg_kit.gauge_A]

Presence / DETECT_FLOOR — A class is present in a sample when its deconvolved fraction $\geq$ DETECT_FLOOR (0.01 in the conductor; a 3% gate in the Mahalanobis adjudicator - RECON D2). An A-score for an absent class is entropy computed over cells that are not there and is not a reading. [cpg_conductor.py]

Commitment line (A = 1.0) — The point where the class's measured entropy equals its floor. It is NOT where healthy sits: healthy is the age-matched band (immune 0.906 at age 4 -> 1.000 at 95). 'A = 1.0 at the healthy floor' is the pre-band (Issue 002-era) definition. [cpg_gauge_engine.py docstring]

Cellular age — Stage 6: invert the age band's beta_mean(age) curve per class to the age at which healthy beta_mean equals the patient's. v1 was a Horvath-style trained clock (not physics); v2 inverted A with the wrong formula (Jensen). v3 is the Recipe's inversion; built, calibration pending, not reportable. [iam_cellular_age_scoring.py]

Bidirectional decomposition — Stage 4.5. Shannon H is symmetric about beta = 0.5, so a disease that pushes some CpGs up and others down leaves beta_mean unmoved and the pooled gauge reads null (VAL-050, d = +0.08). A directional per-CpG composite recovered d = +0.62 (VAL-051); Stage 4.5 runs that sealed formula per patient. [bidirectional_decomposition.py]

Flatness (the Atlas lesson) — The v0.1 rebuild lesson: the old MCMC script showed R-hat ~1.01 (looks perfect) while every cell type in a class came out identical - a flat, broken atlas; the fixed script shows R-hat 1.4-2.6 (looks alarming) while the cell types are distinct and the atlas is right. Convergence diagnostics cannot see the failure that matters; judge a rebuilt class by the distinctness test on per_celltype.csv, never by R-hat alone. [IAM_Atlas/IAMAtlas_FLATNESS_LESSON.md]

Xu-538 — The 538-CpG breast-cancer panel of Xu et al. 2019 used as the per-patient scoring surface in VAL-047 (pre-Atlas); 'Xu-538 relativity' names the finding that a panel's absolute A depends on the pipeline that produced the betas, so bands compiled on one pipeline do not transfer to another without re-derivation. [VAL-047 record; s6a]

GLOSSARY — CHAIN LINKS

One line per runtime file the conductor touches, tagged by what kind of thing it holds. The tags matter more than the names: a **FLOOR** is physics and is one number; a **RULER** says where the gauge reads and is derived from the Atlas; a **BAND** is a cohort statistic and is only as good as its n and pipeline; **CODE** runs; **DATA** is read. Every "healthy reads wrong" case found in September 2026 traced to a **BAND**, never to a **FLOOR** (RECON B1).

file	kind	what it is
IAM_Atlas/IAMAtlasREBUILD.csv(.xz)	DATA	The reference sky: 483,092 CpGs x 115 cell types, per-class MCMC posterior mean, sd and CI per CpG. Every patient is read against it. 2026-05-28, frozen.
IAMAtlasREBUILD_celltype_to_class.json	DATA	115 cell types -> 8 architecture classes. The deconvolver's cell fractions are summed to class fractions through it.
cpg_gauge_engine.py H_MIN_TABLE	FLOOR	The forty floors: 8 classes x 5 substrates (methyI, nucl, fuzz, wps, frag). Physics; one number each; frozen 2026-04-06, MCMC-confirmed. Never a cohort.
iamatlas_gauge_identity_loci_v1_0.json	RULER	Which CpGs the gauge reads, per class: the loci where healthy cells of that class sit within +/-0.05 of H_min_beta. Derived from the Atlas, frozen. Counts terminal 57,247 ... stromal 2,294.
iamatlas_celltype_markers_v0_2.json	RULER	The separation surface: ~100 one-vs-rest discriminative CpGs per cell type (bimodal by construction). Read with mean-of-per-CpG-H, never H(beta_mean) (s105). chrX markers removed (RULING M1b).
age_reference_matrix.json	BAND	What healthy looks like at each age, per class: A and beta mean/sd/p10-p90 in 10 age bins. A cohort table compiled 2026-05-28 from nine papers; n is uneven (immune ~100/bin, stem_adult 28). The gauge's placement is read against this. Phase 1 rebuilds it from one cohort, one pipeline.
tier_breakpoints.json	BAND	The severity ladder: NORMAL 0.95-1.01, ELEVATED 1.01-1.07, SIGNIFICANTLY_ELEVATED 1.07-1.10, BREACH >= 1.10; Warburg line 1.07. Onset moved 1.04 -> 1.01 on 2026-07-03.
mahalanobis_healthy_reference_v2_0_age_matched_derived.json	BAND	Stage 5's reference: derived from the age band ($\mu = A_{\text{mean}}(\text{class}, \text{age})$, sigma from p10). v0_5 / v1_0 retired and refused by the module.
walther_iam_deconvolver.py	CODE	Stage 2: NNLS composition against the Atlas. Returns cell and class fractions and a residual. Gates nothing; presence (fraction >= DETECT_FLOOR) decides which classes may be scored.
cpg_conductor.py	CODE	The orchestrator (2026-07; replaces walther_clinical.py). Pairs every per-cell A with its deconvolved fraction; a class below DETECT_FLOOR 0.01 is absent, not a reading. Runs Stages 2 -> 4 -> 4.5 -> 5 -> 6.
iamatlas_a_scoring.py	CODE	The separation statistic: mean_i H(beta_i)/H_min over each cell type's discriminative markers -> 115 per-cell A's for the disease matcher. Guarded by test_a_score_canonical.py (separation surface only).
cpg_gauge_engine.py	CODE	Stage 4: $A = H(\text{beta_mean})/H_{\text{min}}$ over identity loci, placed in the age band (BELOW/IN/ABOVE) and on the severity ladder. Carries the FLOOR table and the SATURATION_MARGIN.
bidirectional_decomposition.py + directional_panels_v1_0.json	CODE	Stage 4.5, the fix for entropy's symmetry: H is symmetric about 0.5, so a disease pushing CpGs both ways leaves beta_mean unmoved and the gauge reads null (VAL-050 d=+0.08). Scores direction per CpG with the sealed VAL-051 composite (d=+0.62). Immune panel only in v1.0.
cpg_patient_cmb.py + IAM_Atlas/healpix_mapping/	CODE	Stage 4.6: the patient's per-CpG $z = (\text{beta} - \mu)/\text{sigma}$ against the Atlas posterior, on the HEALPix NSIDE-128 sky (Plate 05). The residual map of one observation against a reference map with per-pixel uncertainty - the Planck workflow.
iamatlas_mahalanobis_scoring.py	CODE	Stage 5, Option A: the eight class-gauge A's as a vector, distance from the age-band centroid with the band's covariance, n-adaptive chi-square alarm. Answers 'is the whole profile off' where any single class may be in band.
iam_cellular_age_scoring.py	CODE	Stage 6: the band run backwards - at what age does healthy beta_mean(class) equal this patient's? Third attempt: v1 was a trained clock (wrong), v2 inverted the wrong formula (Jensen). Built; calibration pending; reads 4 yr for adults today because the band curves are too flat to invert - not reportable.
stage_0_intake.py	CODE	Stage 0: manifest, IDAT header vs declared array, SHA-256 integrity with re-transmission detection, control-probe / detection-p / bead / call-rate / sex QC (the last four await the Stage 1 hand-off), PROCEED / PENALTY / QUARANTINE gate. Fail-open on 0.1 quarantines closed 2026-09-19.

file	kind	what it is
stage_1_idat_calibration.py	CODE	Stage 1: raw IDAT pair -> noob-calibrated beta via methylprep (pandas < 2). Bit-identical to the project cache on 11/11 (PROC-CAL-01). Owns any raw-array alignment; no per-class input offset is applied downstream.
RETIRED/.../IAM_Cellular_Age/ (age_/sex_/smoking_axis_foreground.py, *_layer.csv)	RETIRE D	Stage 3 foreground subtraction, built and deliberately not wired (SOP s104): a galactic foreground is a separate source; the methylome 'foreground' is the patient's own biology - annotated, never subtracted.

[Issue 003 note, RECON M2] The entry above is Issue 002's and is superseded. The Mahaffey number of the cell is 20.94 (author, 2026-09-19), defined as the ATP ordering budget per event over $k_B T$ at body temperature; n_{bio} was an early QAPE-era form of the ratio and is no longer used. H_{min} is not the Mahaffey number: H_{min} is the per-class entropy floor in bits, forty values across eight classes and five substrates.

A FINAL NOTE

Heath W. Mahaffey · Entiat, Washington · April 2026

This research is personal, and my step brother Marcus is often on my mind as I continue working on it.

He was diagnosed with an aggressive liver tumor. He had a transplant. Within six months the tumor was back in the new liver, and a few months after that he was gone. He left behind a wife and three small children. He never said goodbye to any of them. He could not bring himself to accept that he was dying, so he spent his last months sedated in a hospital, fighting, and the goodbye never happened.

I don't know whether an earlier, honest signal about his remaining reserve would have changed his choice. That was his choice to make. But he was never given the chance to make it with good information, and his wife and his three small children were not given the chance to have that conversation with him while he was still there to have it. That is the part I cannot stop thinking about.

I just hope that what is in these pages might help prevent what happened to his family from happening to another family with three small children who don't get the chance to say goodbye to their dad.

The hope is that this framework can give patients and their families honest information about the reserve remaining. Not to recommend when to stop. Not to replace the doctor's judgment. To report, honestly, what the physics says — so the person living it can decide how they want to spend the time.

For Marcus.

— HWM

IAMPerformance · GAPE Issue 002 · Patents pending 64/012,720 and 64/014,568